



Machine Learning 101

A Tour of Machine Learning Algorithms

31 March 2017

Sung-Soo Kim

sungsoo@etri.re.kr

Smart Data Research Group

ETRI

I'm not a "*Machine Learning*" (ML) guy!

But...

Maybe, I will be a ML guy!

Machine Learning in Robotics

Machine Learning



Machine Learning's Successes

5

■ *Computer Vision*

- Image inpainting/denoising, segmentation
- object recognition/detection, scene understanding

■ *Information Retrieval / NLP*

- Text, audio, and image retrieval
- Parsing, machine translation, text analysis

■ *Speech Processing*

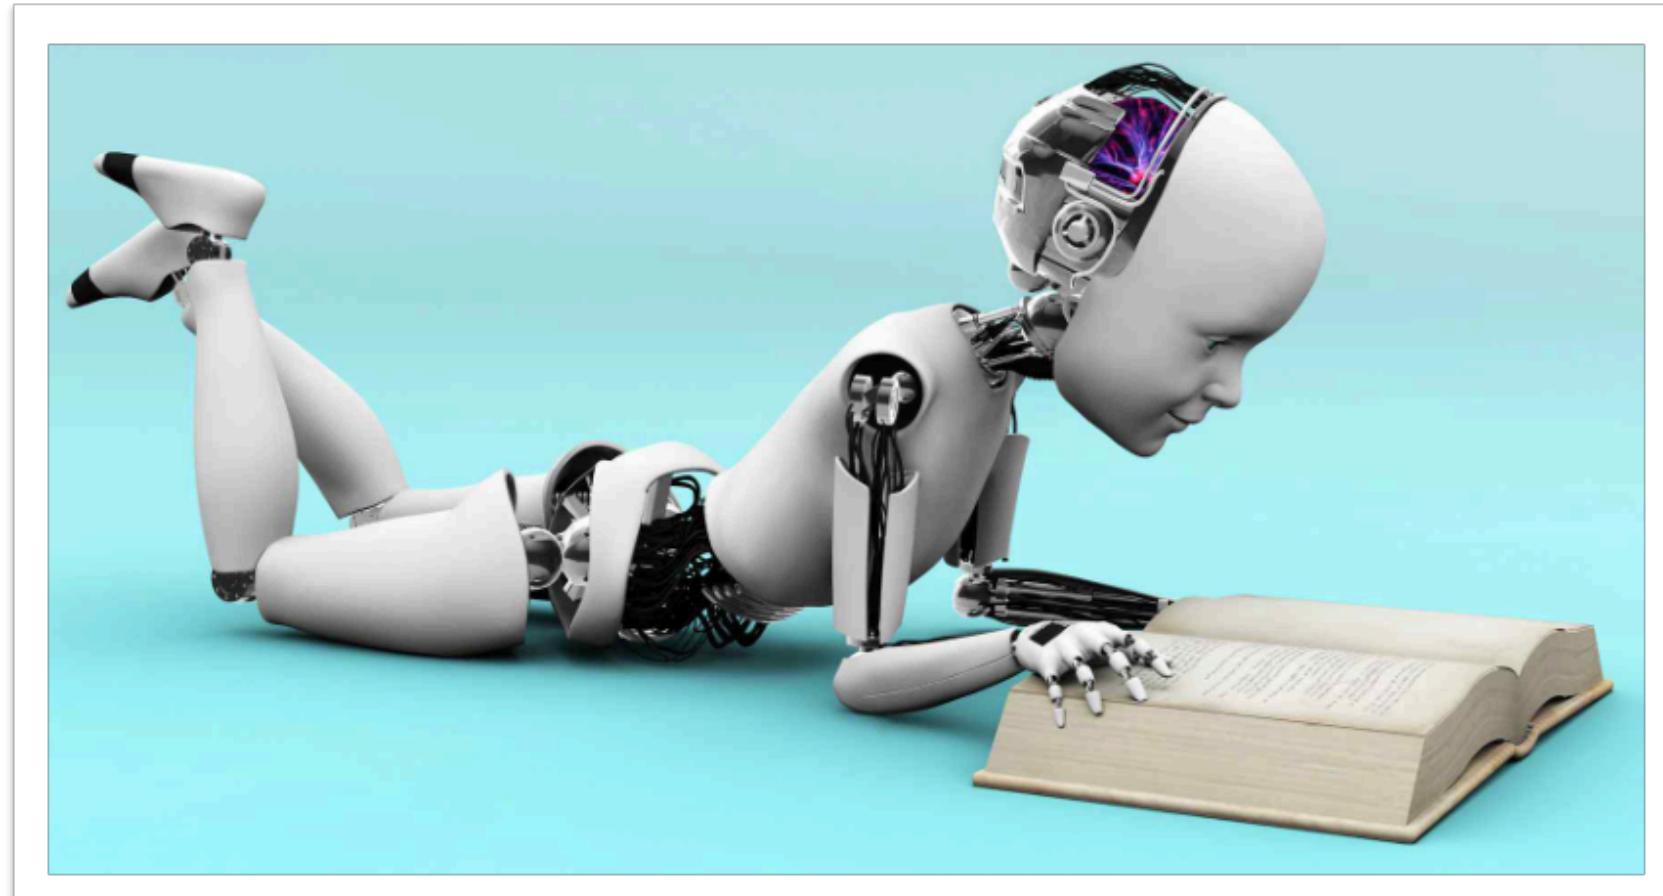
- Speech recognition, voice identification

■ *Robotics*

- Autonomous car driving, planning, control

■ *Computational Biology*

■ *Cognitive Science*



Training/evaluating deep neural networks

Technique leading to many high-profile AI advances in recent years

Speech recognition/natural language processing

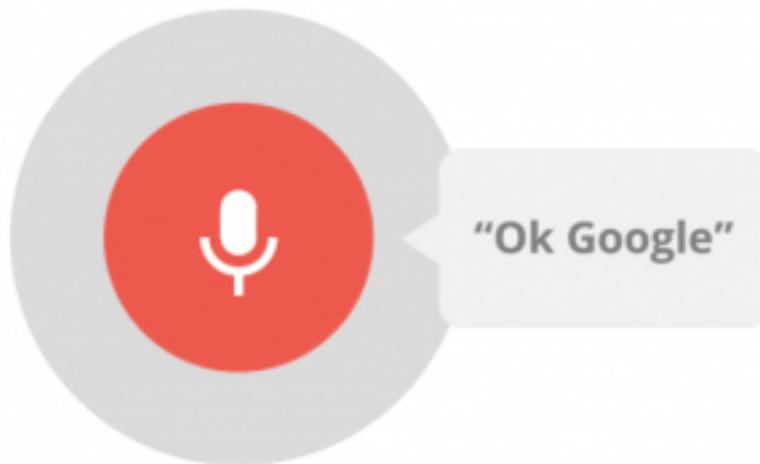
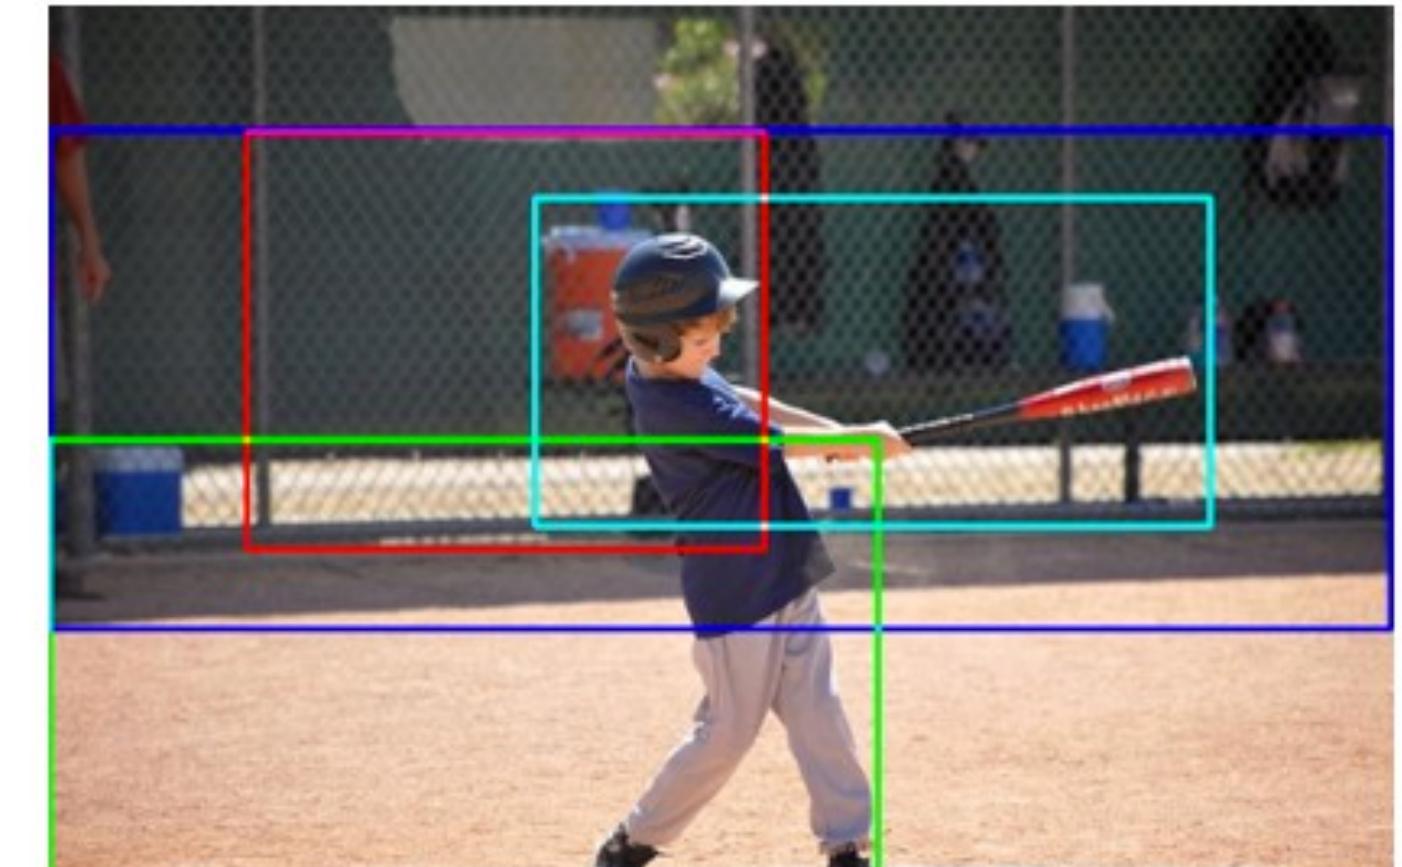
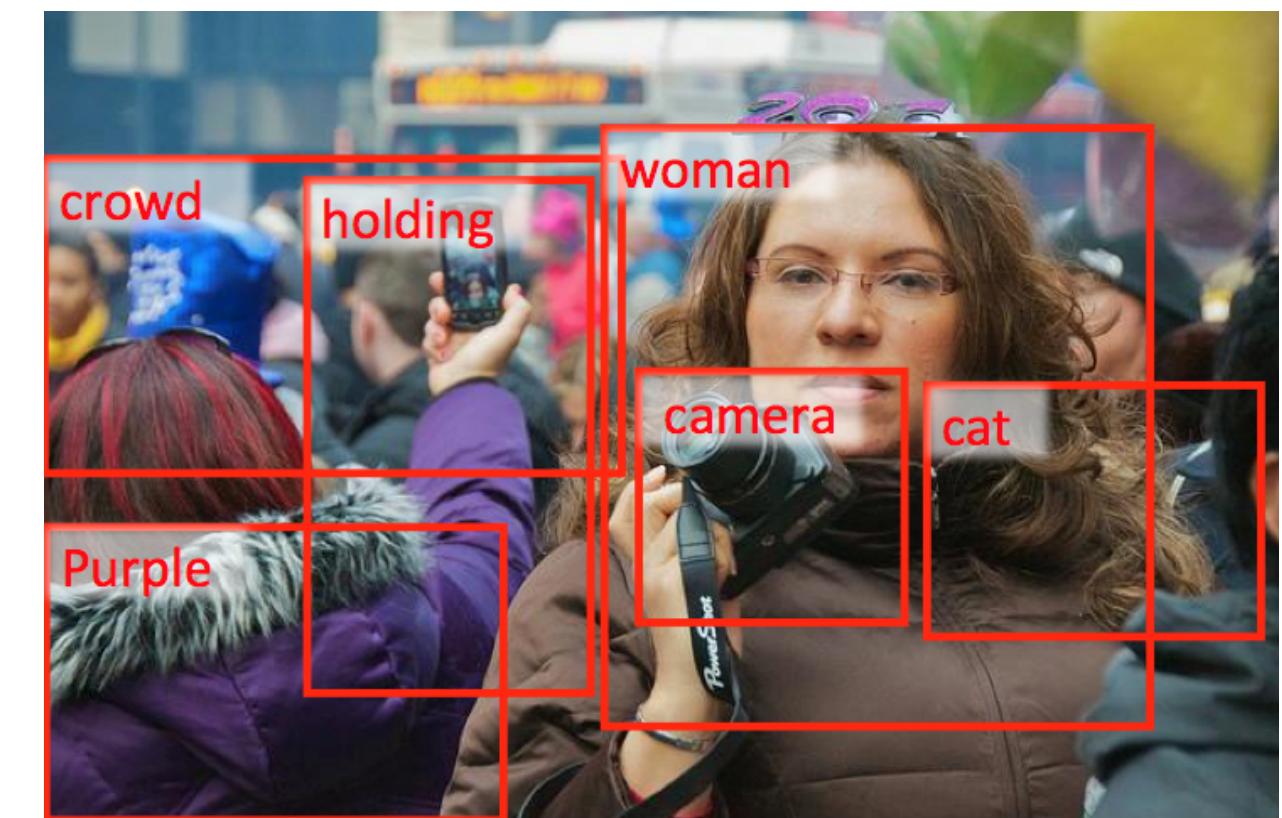


Image interpretation and understanding



a baseball player swinging a bat at a ball
a boy is playing with a baseball bat



Top 10 Trends 2015 ~ 2017 [Gartner]

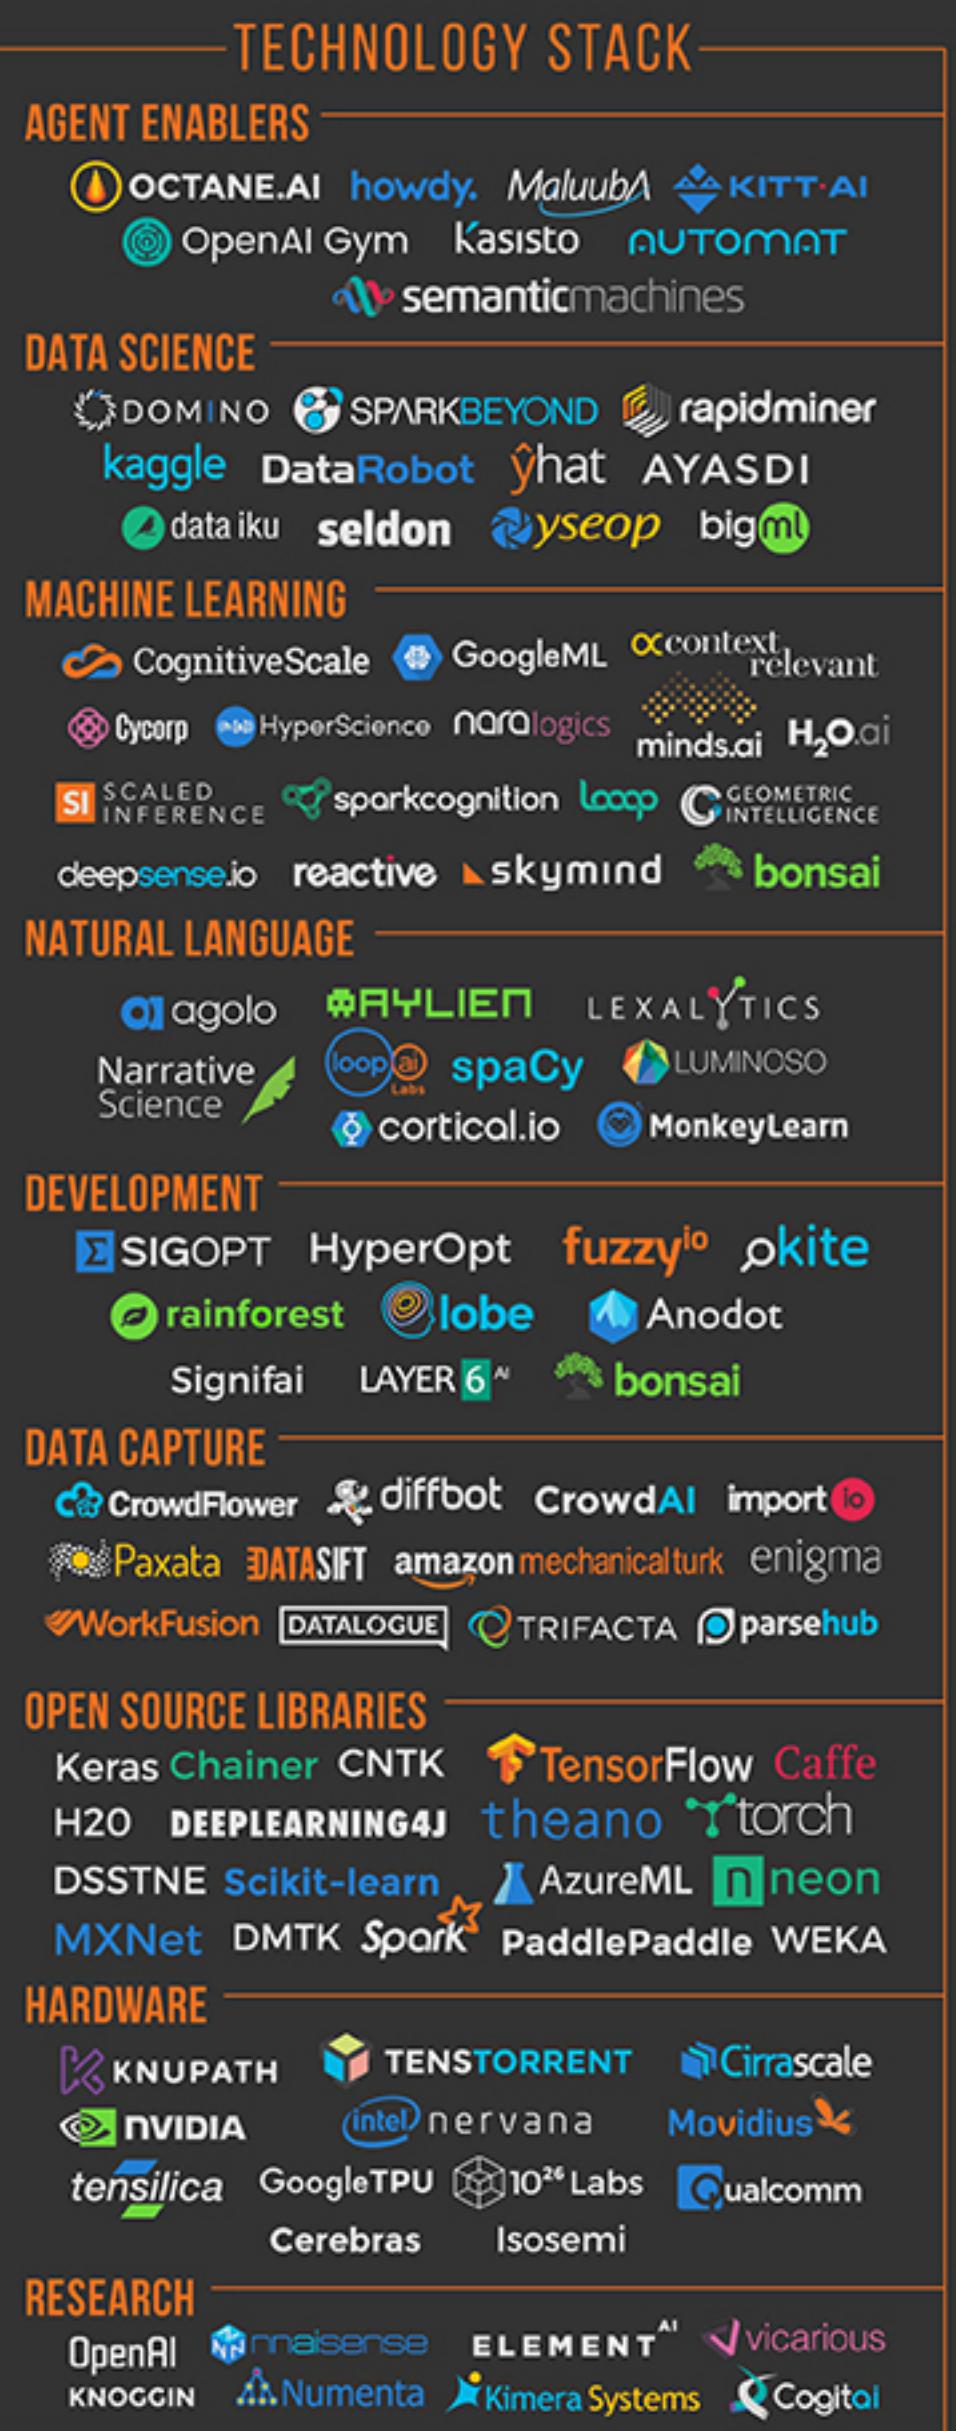
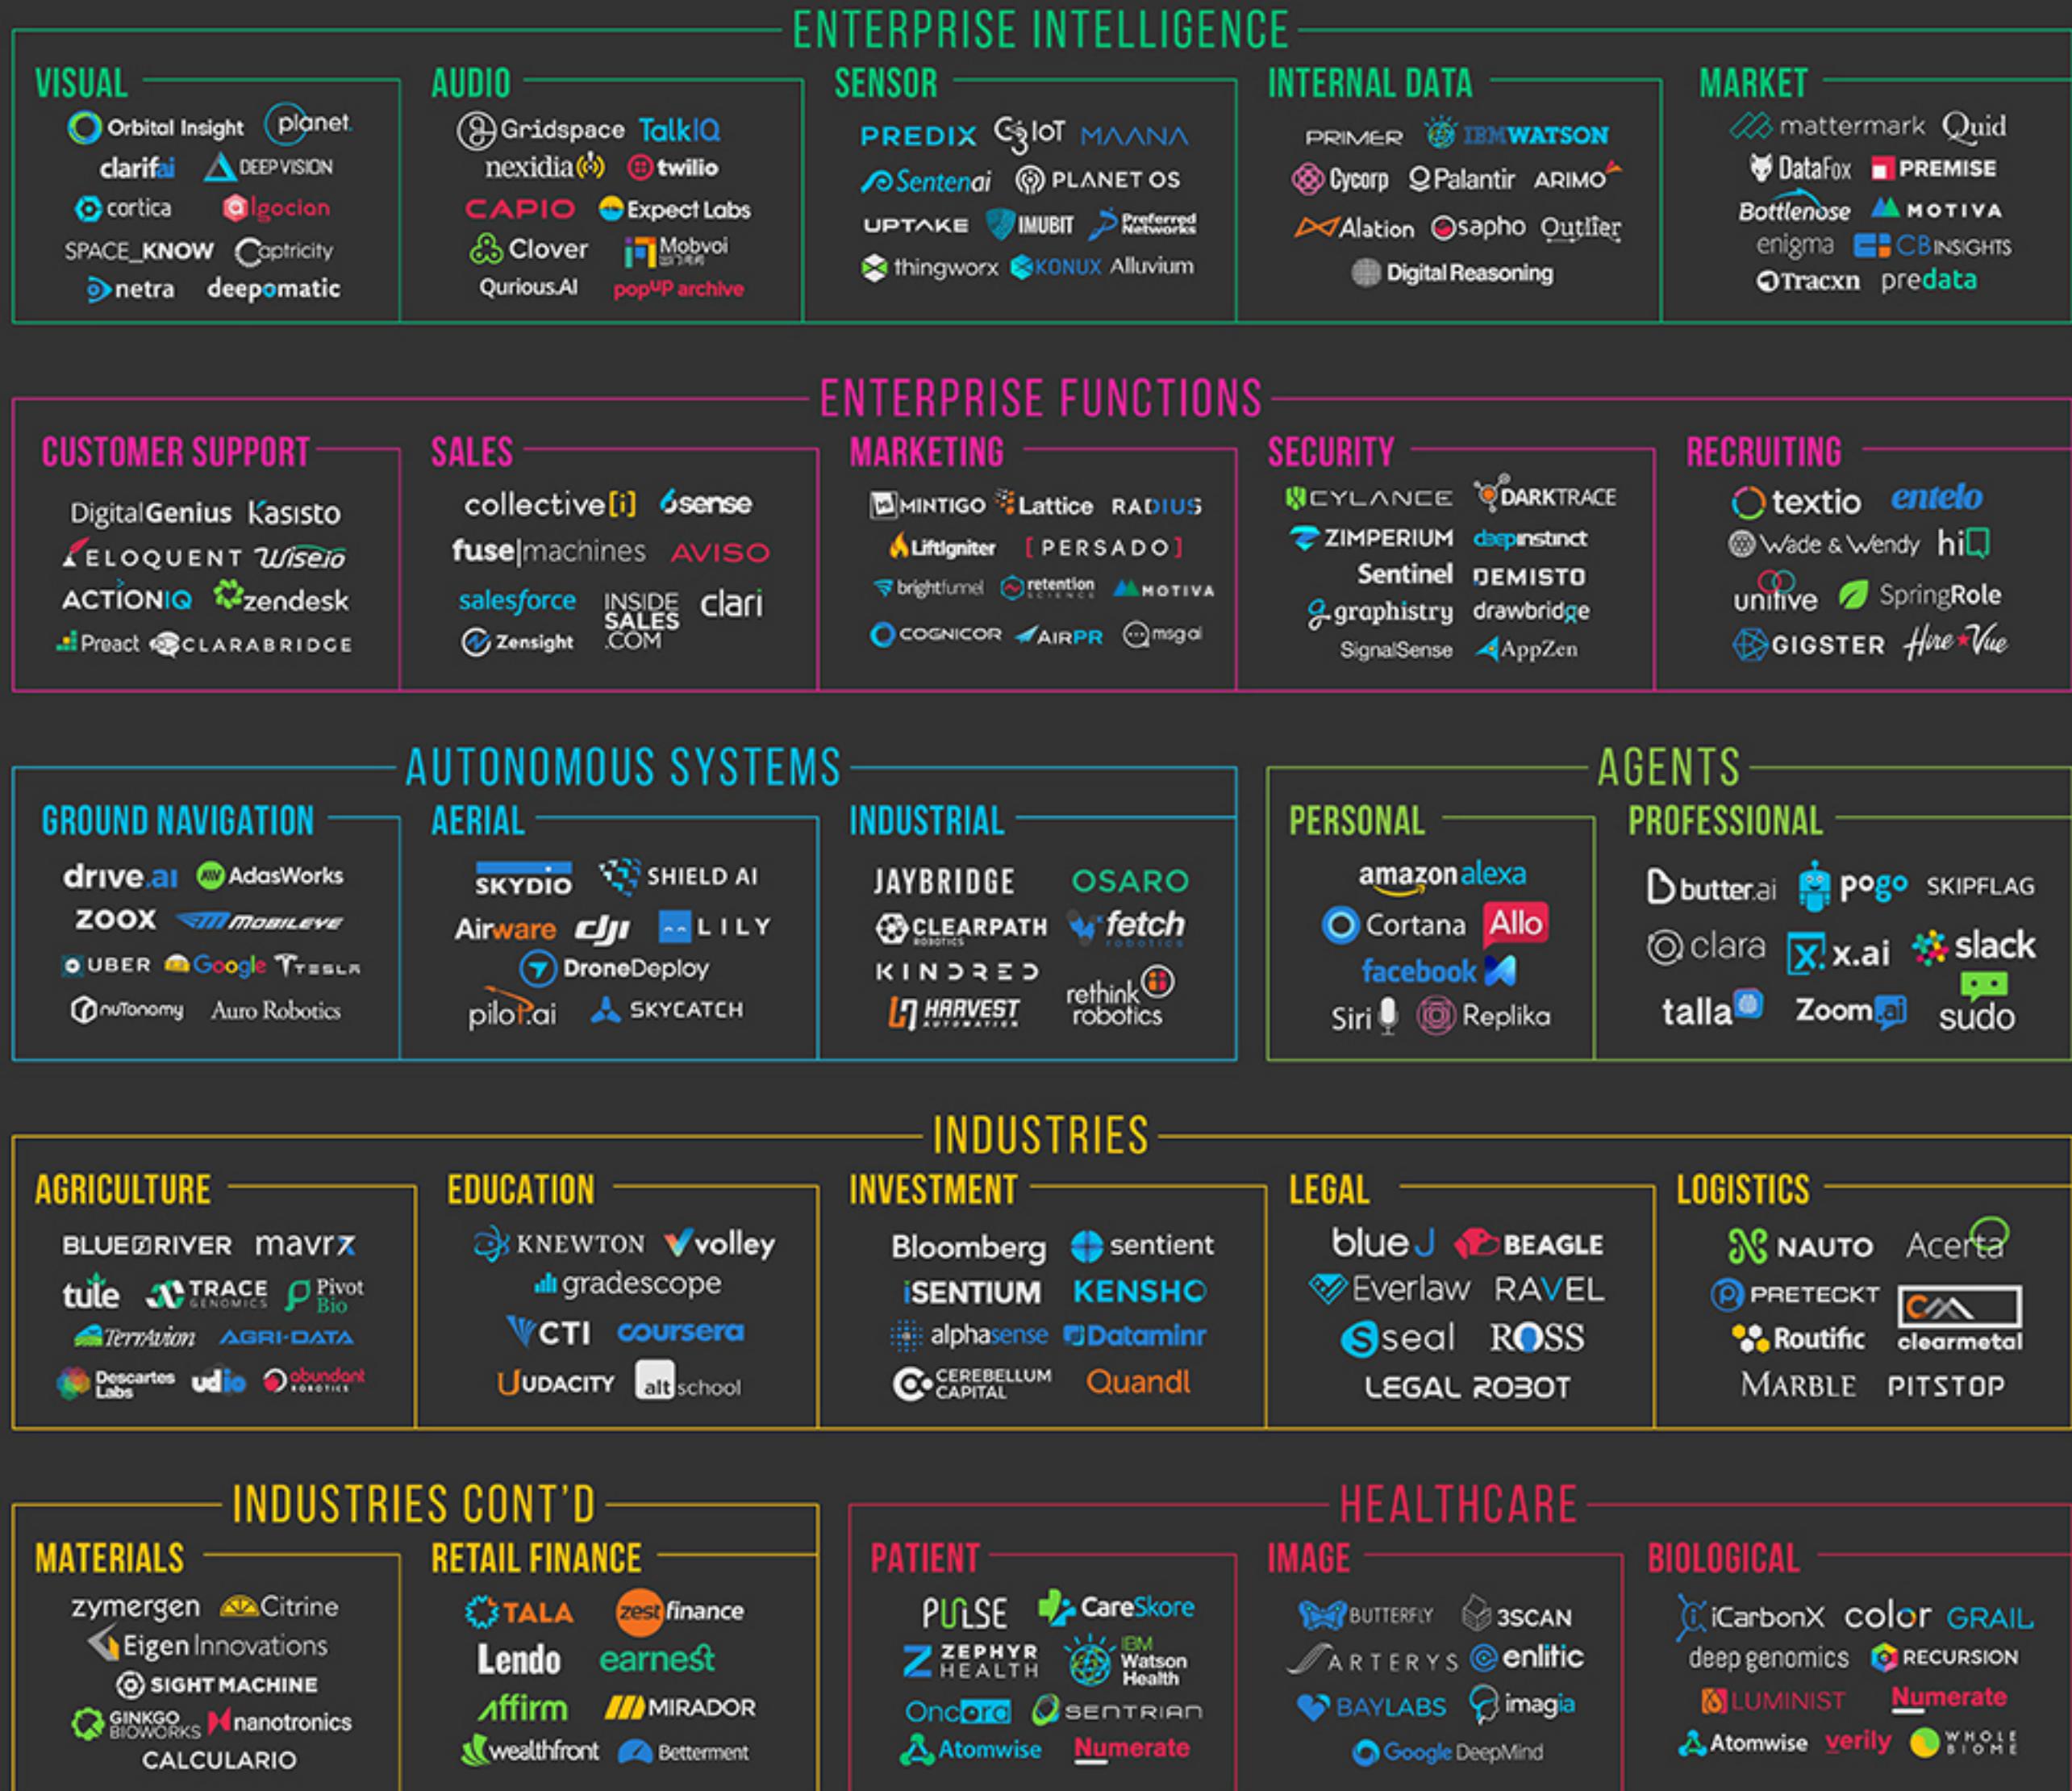
7

2015	2016	2017
Merging the Real World and Virtual World	Digital Mesh	Intelligent
Computing Everywhere	Device Mesh	Artificial Intelligence and Advanced Machine Learning
The Internet of Things	Continuous & Ambient UX	Intelligent Apps
3D Printing	3D Printing Materials	Intelligent Things
Intelligence Everywhere	Smart Machines	Digital
Advanced, Pervasive and Invisible Analytics	Information of Everything	Virtual Reality and Augmented Reality
Context-Rich Systems	Advanced Machine Learning	Digital Twins
Smart Machines	Autonomous Agents & Things	Blockchains and Distributed Ledgers
The New IT Reality Emerges	New IT Reality	Mesh
Cloud/Client Computing	Adaptive Security Architecture	Conversational Systems
Software-Defined Applications and Infrastructure	Advanced Systems Architecture	Digital Technology Platforms
Web-Scale IT	Mesh App & Service Architecture	Mesh App and Service Architecture
Risk-Based Security and Self-protection	IoT Architecture & Platforms	Adaptive Security Architecture

Machine Intelligence 3.0

8

MACHINE INTELLIGENCE 3.0



Outline

9

- What is Machine Learning?
- A Fast Tour of Machine Learning Algorithms
- A Few Useful Things to Know About Machine Learning
- Summary



“새로운 것을 시도하기 흥기를 가지고 있지 않다면,
우리의 삶은 과연 어떤 모습이겠나?”

- Vincent Van Gogh

References and Slide Credits

10

- M. I. Jordan and T. M. Mitchell, *Machine learning: Trends, perspectives, and prospects*, Science, Vol. 349 Issue 6245, July 2015.
- Pedro Domingos, *A Few Useful Things to Know About Machine Learning*, Communications of the ACM, Vol. 55 No. 10, Oct. 2012.
- T. M. Mitchell, *Machine learning*, McGraw-Hill, 1997.

Despite practical challenges, we are hopeful that informed discussions among policy-makers and the public about data and the capabilities of machine learning, will lead to insightful designs of programs and policies that can balance the goals of protecting privacy and ensuring fairness with those of reaping the benefits to scientific research and to individual and public health. Our commitments to privacy and fairness are evergreen, but our policy choices must adapt to advance them, and support new techniques for deepening our knowledge.

REFERENCES AND NOTES

1. M. De Choudhury, S. Counts, E. Horvitz, A. Hoff, in Proceedings of International Conference on Weblogs and Social Media [Association for the Advancement of Artificial Intelligence (AAAI), Palo Alto, CA, 2014].
2. J. S. Brownstein, C. C. Freifeld, L. C. Madoff, N. Engl. J. Med. **361**, 2153–2155 (2009).
3. G. Eysenbach, J. Med. Internet Res. **11**, e11 (2009).
4. D. A. Broniatowski, M. J. Paul, M. Dredze, PLOS ONE **8**, e83672 (2013).
5. A. Sadlik, H. Kautz, V. Silenzio, in Proceedings of the Twenty-Sixth AAAI Conference on Artificial Intelligence (AAAI, Palo Alto, CA, 2012).
6. M. De Choudhury, S. Counts, E. Horvitz, in Proceedings of the SIGCHI Conference on Human Factors in Computing Systems [Association for Computer Machinery, New York, 2013], pp. 3267–3276.
7. R. W. Freeman, J. Harasz, N. H. Shah, D. DeMetschel, E. Horvitz, Clin. Pharmacol. Ther. **90**, 239–246 (2004).
8. Samaritans Radar: www.samaritans.org/how-we-can-help-you/supporting-someone-online.
9. Shut down Samaritans Radar: <http://bit.ly/Samaritans-radar>.
10. U.S. Equal Employment Opportunity Commission (EEOC), 29 Code of Federal Regulations (C.F.R.), 1630.2 (g) (2013).
11. EEOC, 29 CFR 1635.3 (c) (2013).
12. M. A. Rothstein, J. Law Med. Ethics **36**, 837–840 (2008).
13. Executive Order of the President, *Big Data: Seizing Opportunities, Preserving Values* (White House, Washington, DC, 2014); <http://bit.ly/1TfDy6>.
14. Letter from Manisha Mital, FTC, to Reed Freeman, Morrison, & Foerster LLP, Counsel for Netflix, 2 [closing letter] (2010); <http://bit.ly/1GCFxXR>.
15. In re Facebook, Complaint, FTC File No. 092 3184 (2012).
16. FTC Staff Report, *Mobile Privacy Disclosures: Building Trust Through Transparency* (FTC, Washington, DC, 2013); <http://bit.ly/1kNzZr>.
17. FTC, *Protecting Consumer Privacy in an Era of Rapid Change: Recommendations for Businesses and Policymakers* (FTC, Washington, DC, 2012).
18. Directive 95/46/EC of the European Parliament and of the Council of Europe, 24 October 1995.
19. L. Sweeney, Online ads roll the dice [blog]: <http://bit.ly/1kGeCyg>.
20. FTC, “Big data: A tool for inclusion or exclusion?” [workshop, FTC, Washington, DC, 2014]; <http://bit.ly/1S8t6c>.
21. FTC, *Protecting Privacy: A Call for Transparency and Accountability* (FTC, Washington, DC, 2014); <http://bit.ly/1GCFq5b>.
22. J. Podesta, “Big data and privacy: 1 year out” [blog]: <http://bit.ly/WfsePrivacy>.
23. White House Council of Economic Advisers, *Big Data and Differential Pricing* (White House, Washington, DC, 2015).
24. Executive Office of the President, *Big Data and Differential Processing* (White House, Washington, DC, 2015); <http://bit.ly/1GCFwV>.
25. Executive Order of the President, *Big Data Seizing Opportunities, Preserving Values* (White House, Washington, DC, 2014); <http://bit.ly/1TSOnG>.
26. President’s Council of Advisors on Science and Technology (PCAST), *Big Data and Privacy: A Technological Perspective* (White House, Washington, DC, 2014); <http://bit.ly/1CSewNv>.
27. European Commission, *Proposal for a Regulation of the European Parliament and of the Council on the Protection of Individuals with regard to the Processing of Personal Data and on the Free Movement of Such Data (General Data Protection Regulation)*, COM(2012) 21 final (2012); <http://bit.ly/1Ls6P0>.
28. M. Schrems, Facebook Ireland Limited, S.J. Urdampilleta, *Unlawful data transmission to the U.S.A. (“PRISM”)*, **165** and **167** (2013); <http://bit.ly/1eXvFacebook>.

*Corresponding author. E-mail: jordan@cs.berkeley.edu (M.I.J.); tom.mitchell@cs.cmu.edu (T.M.M.)

Downloaded from www.sciencemag.org on July 20, 2015

SCIENCE sciencemag.org

17 JULY 2015 • VOL 349 ISSUE 6245 255

review articles

DOI:10.1145/2347736.2347755

Tapping into the “folk knowledge” needed to advance machine learning applications.

BY PEDRO DOMINGOS

A Few Useful Things to Know About Machine Learning

MACHINE LEARNING SYSTEMS automatically learn programs from data. This is often a very attractive alternative to manually constructing them, and in the last decade the use of machine learning has spread rapidly throughout computer science and beyond. Machine learning is used in Web search, spam filters, recommender systems, ad placement, credit scoring, fraud detection, stock trading, drug design, and many other applications. A recent report from the McKinsey Global Institute asserts that machine learning (a.k.a. data mining or predictive analytics) will be the driver of the next big wave of innovation.¹⁵ Several fine textbooks are available to interested practitioners and researchers (for example, Mitchell¹⁶ and Witten et al.²⁴). However, much of the “folk knowledge” that is needed to successfully develop machine learning applications is not readily available in them. As a result, many machine learning projects take much longer than necessary or wind up producing less-than-ideal results. Yet much of this folk knowledge is fairly easy to communicate. This is the purpose of this article.

» key insights

- Machine learning algorithms can figure out how to perform important tasks by generalizing from examples. This is often feasible and cost-effective where manual programming is not. As more data becomes available, more ambitious problems can be tackled.
- Machine learning is widely used in computer science and other fields. However, developing successful machine learning applications requires a substantial amount of “black art” that is difficult to find in textbooks.
- This article summarizes 12 key lessons that machine learning researchers and practitioners have learned. These include pitfalls to avoid, important issues to focus on, and answers to common questions.

78 COMMUNICATIONS OF THE ACM | OCTOBER 2012 | VOL. 55 | NO. 10

Main Takeaways

11

- Takeaway #1
 - Machine learning is *not monolithic*.
- Takeaway #2
 - Deep learning has emerged as a technique with *strong advantages*, but also has important *drawbacks* as well.
- Takeaway #3
 - Machine learning will have *major implications* for what products will look like going forward.
- Takeaway #4
 - There are lot of “*folk knowledge*” that is needed to successfully develop machine learning applications, which is *not* readily available in ML textbooks.

Fundamental Question

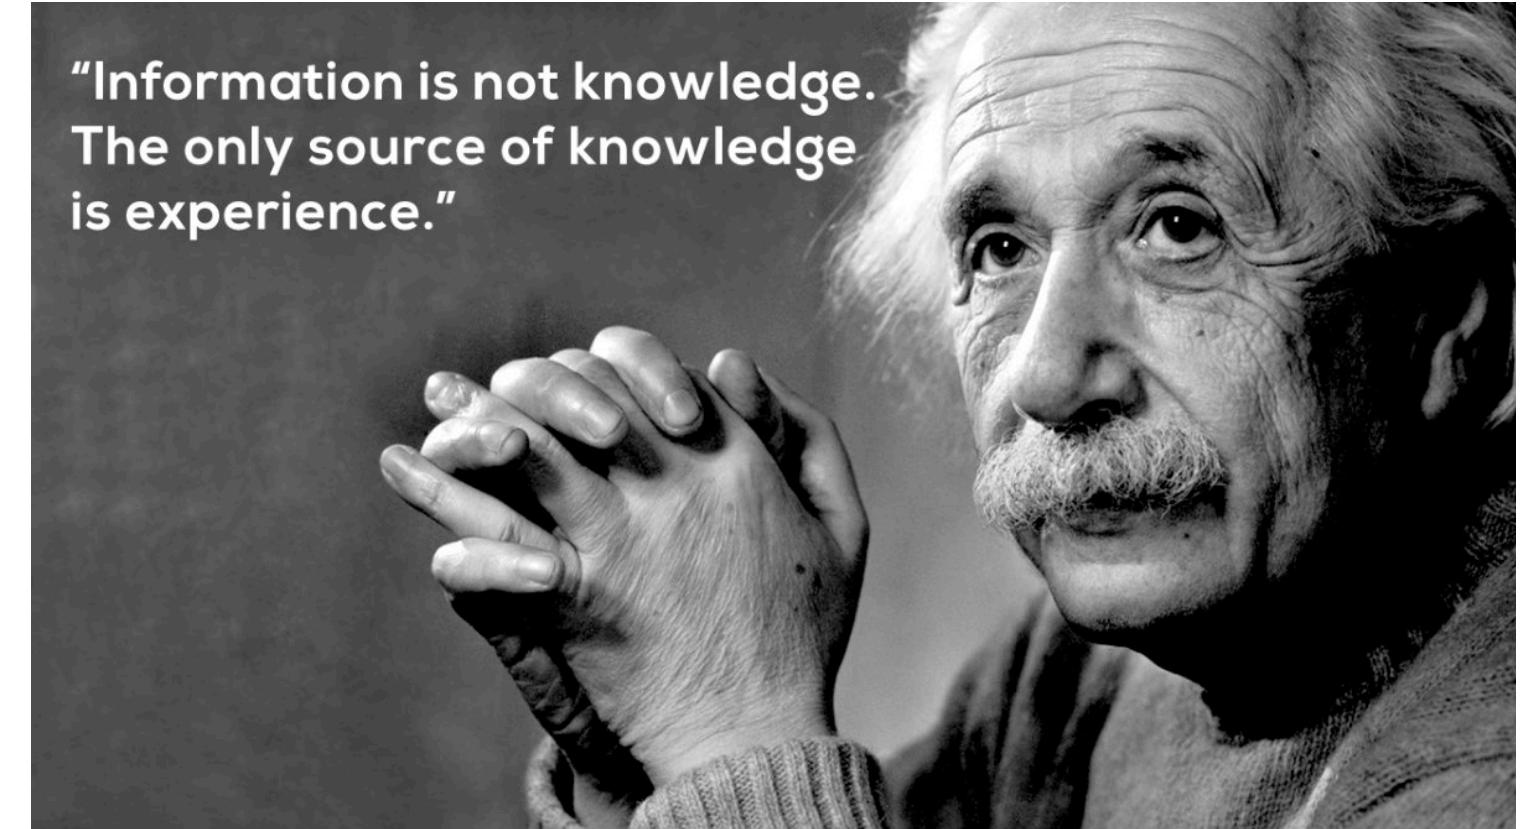
12

- Where does *knowledge* come from?

Evolution



Experience



Culture



Computer



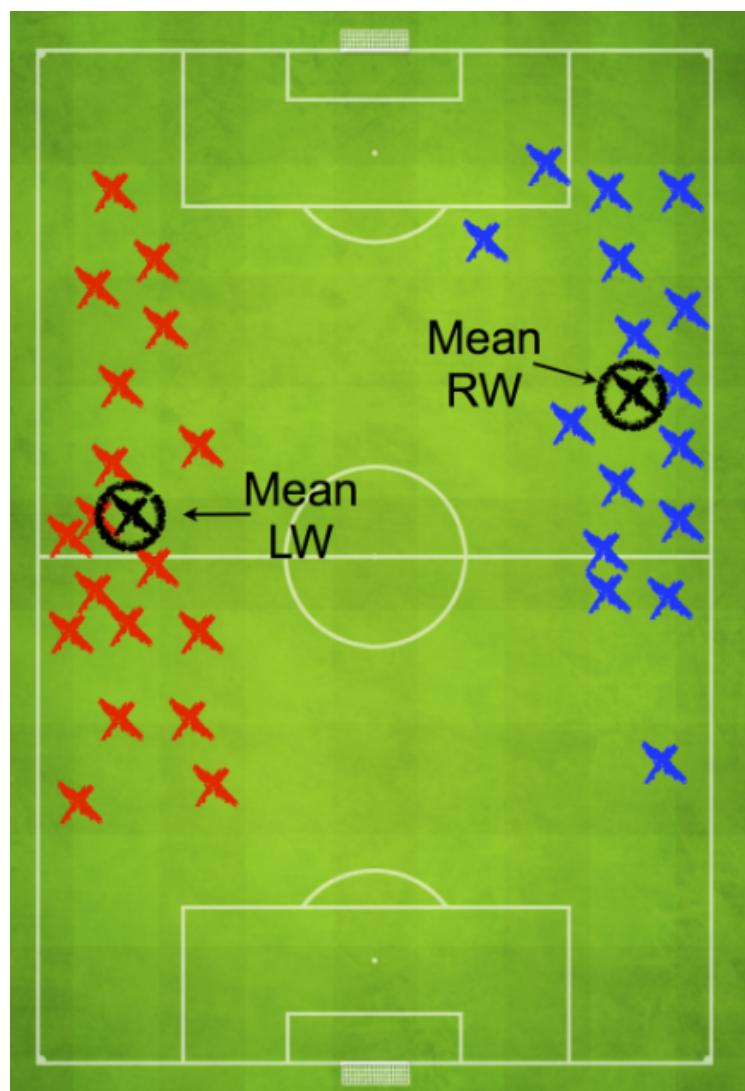
What is Machine Learning?

■ Limitations of *explicit* programming

- Spam filter: many rules
- Automatic driving: too many rules

■ Machine Learning

- “Field of study that gives computers the ability to learn *without being explicitly programmed*” - Arthur Samuel (1959)
- “A computer program is said to **learn** from *experience E* with respect to some class of *tasks T* and *performance measure P*, if its performance at tasks in *T*, as measured by *P*, improves with experience *E*” - T. Michell (1997)



Example: A program for soccer tactics

T: Win the game

P: Goals

E: (x) Player's movements

(y) Evaluation

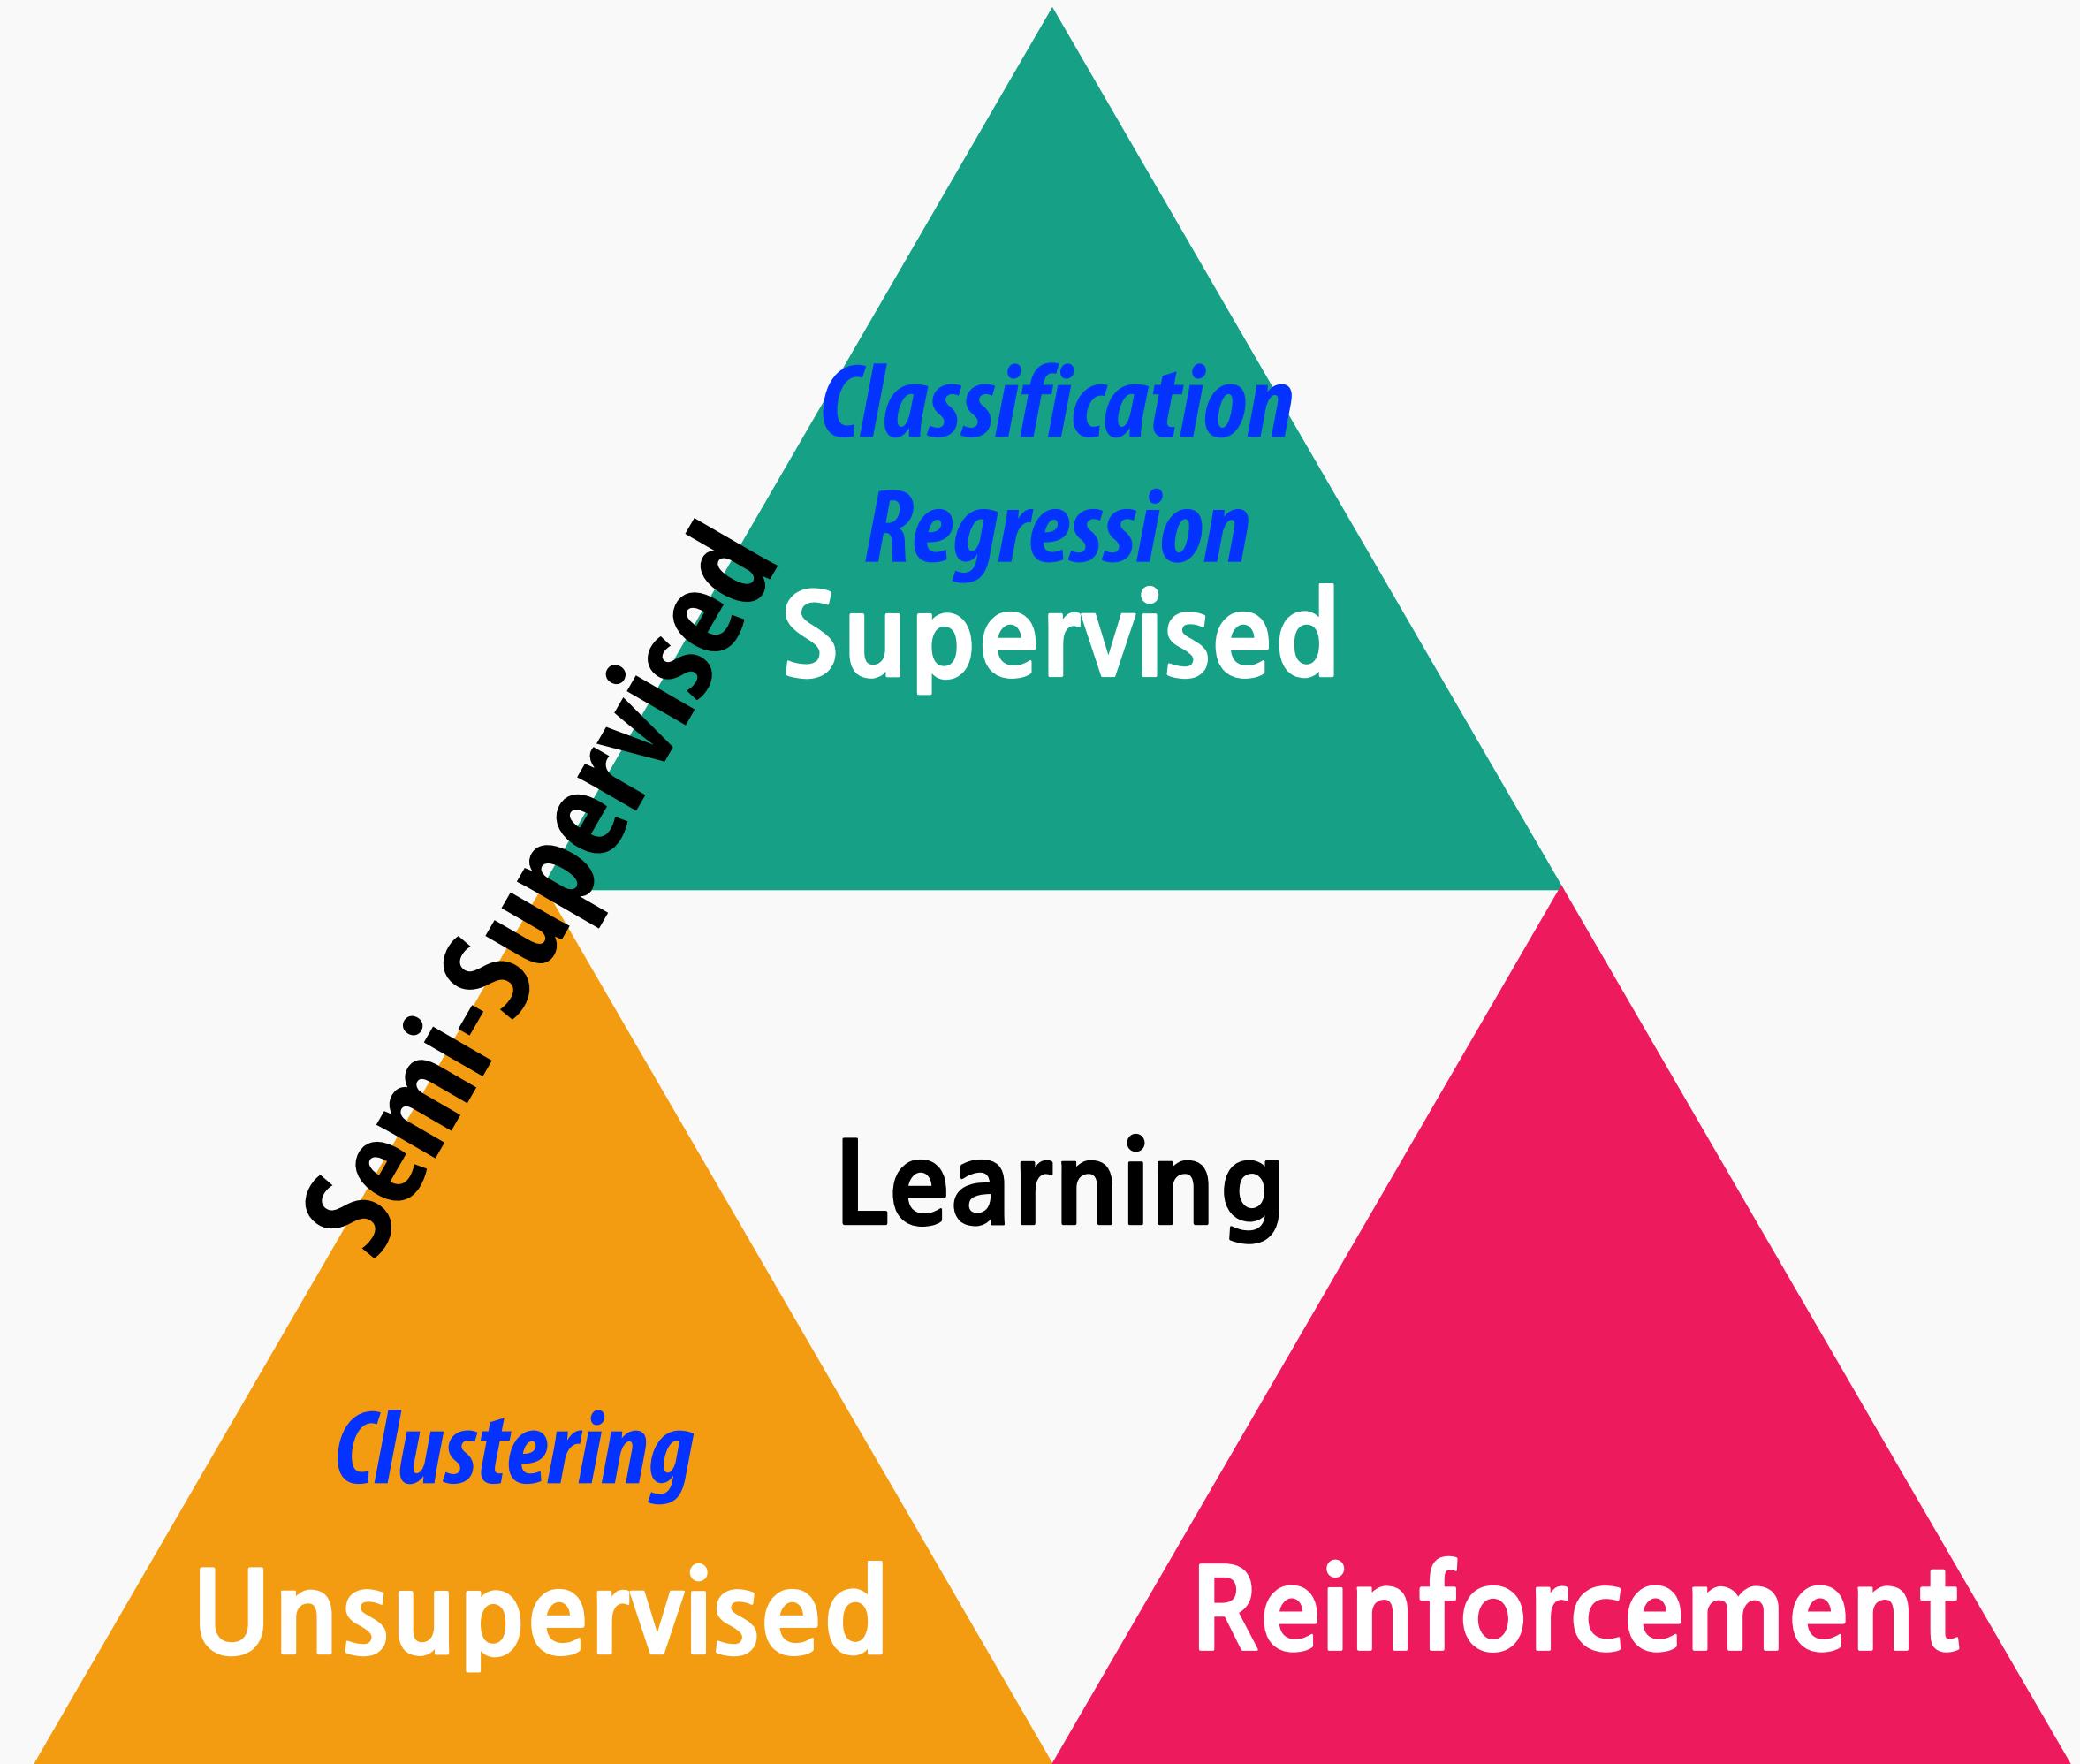
Research Fields

- Self-improving Systems
- Knowledge Discovery
- Data-Driven SW Design
- Automatic Programming

Machine Learning in Robotics

Three Types of Machine Learning

머신러닝의
세가지 타입



Supervised Learning

16

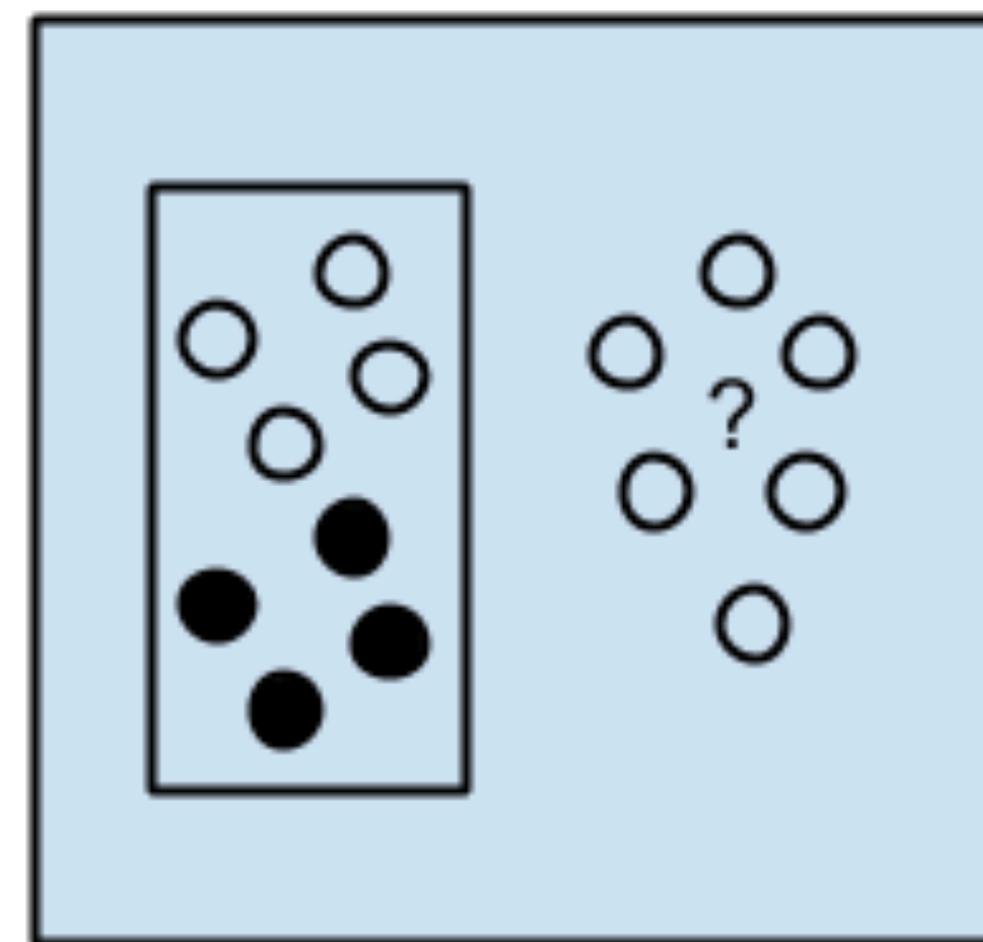
Supervised Learning

Input data is called training data and has a known label or result such as spam/not-spam or a stock price at a time.

A model is prepared through a training process in which it is required to make predictions and is corrected when those predictions are wrong. The training process continues until the model achieves a desired level of accuracy on the training data.

Example problems are classification and regression.

Example algorithms include Logistic Regression and the Back Propagation Neural Network.



Supervised Learning
Algorithms

Unsupervised Learning

17

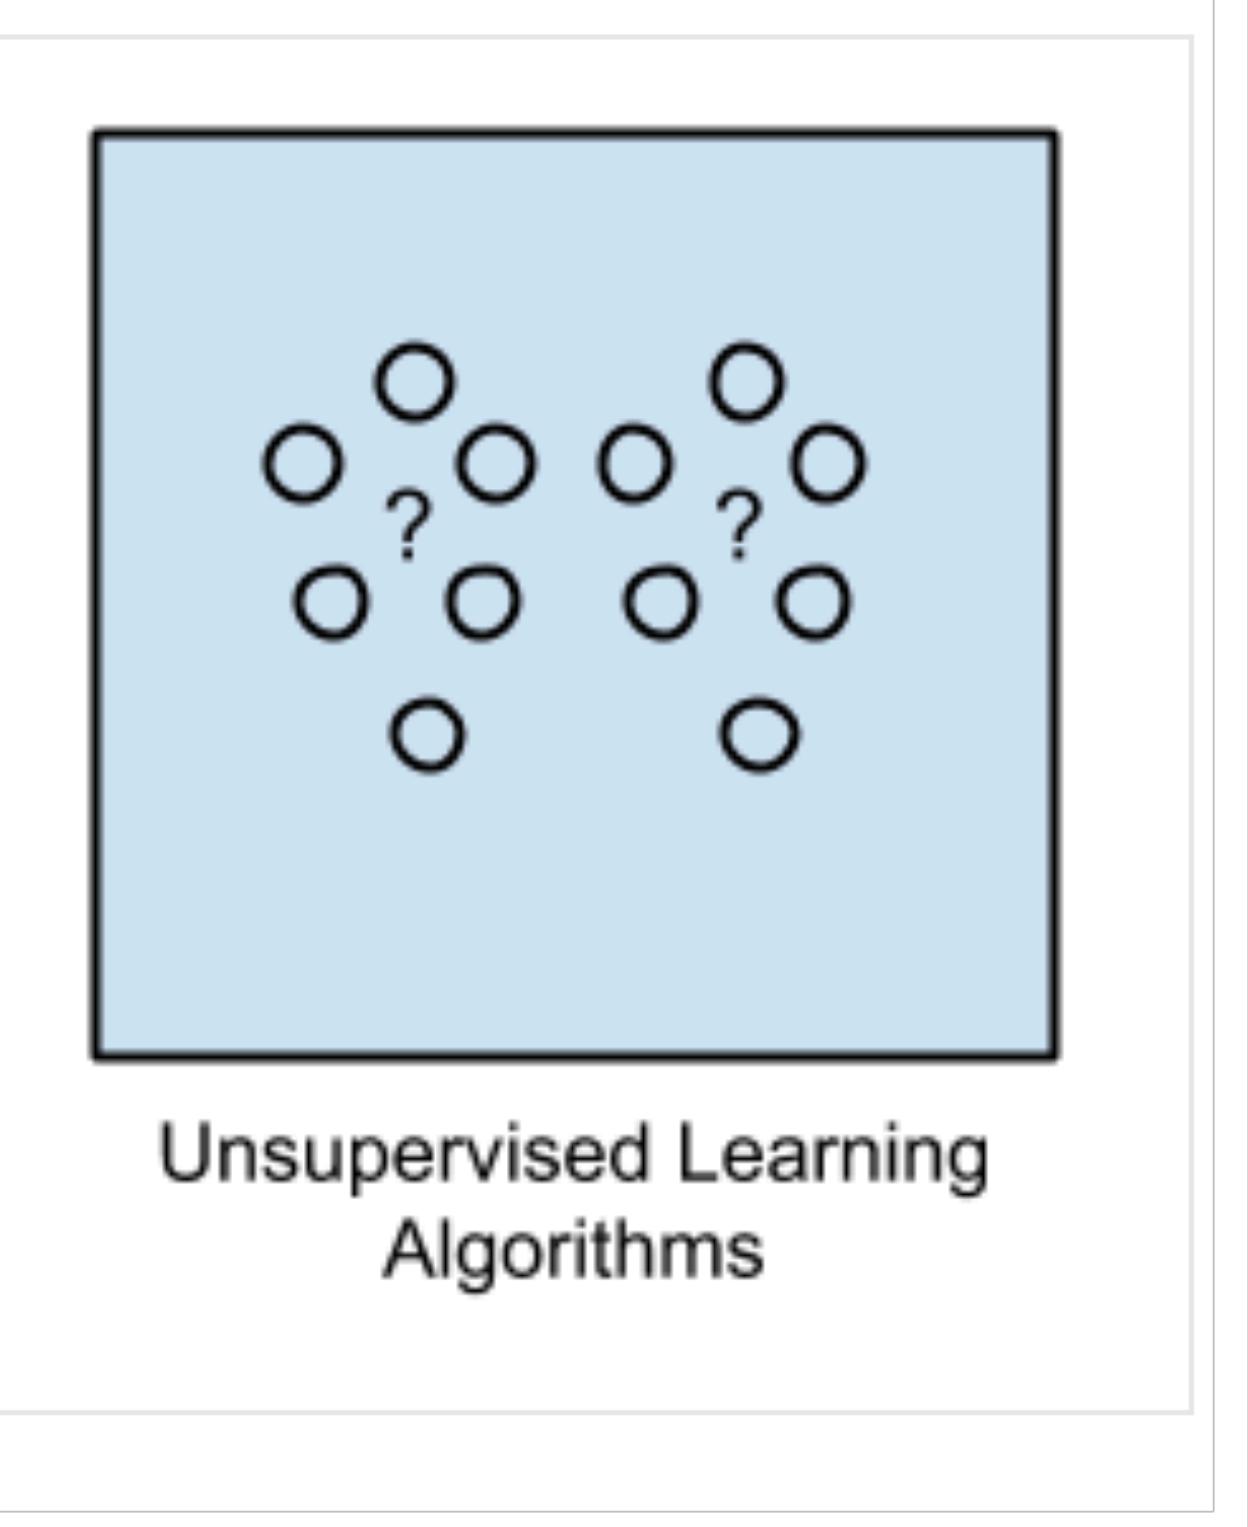
Unsupervised Learning

Input data is not labeled and does not have a known result.

A model is prepared by deducing structures present in the input data. This may be to extract general rules. It may be through a mathematical process to systematically reduce redundancy, or it may be to organize data by similarity.

Example problems are clustering, dimensionality reduction and association rule learning.

Example algorithms include: the Apriori algorithm and k-Means.

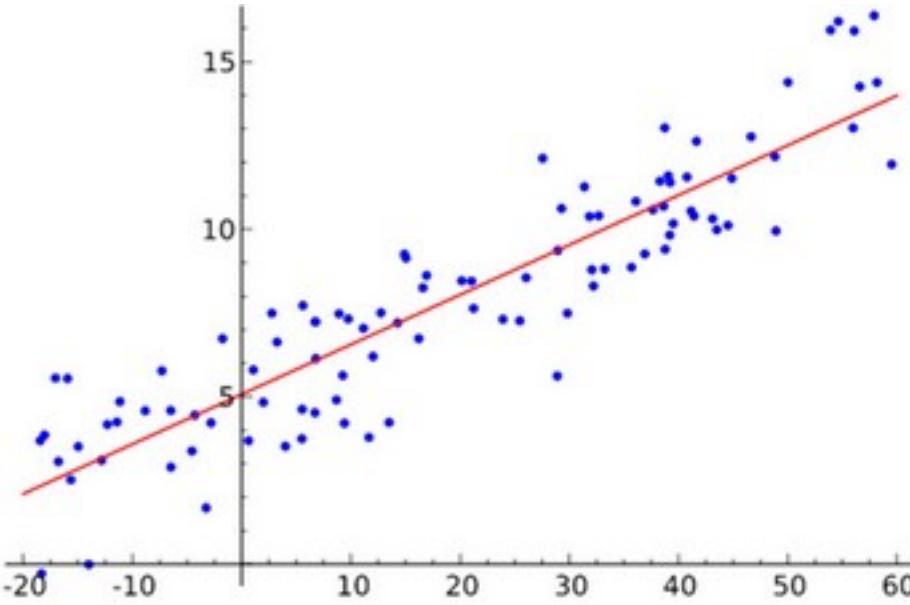


What is Machine Learning?

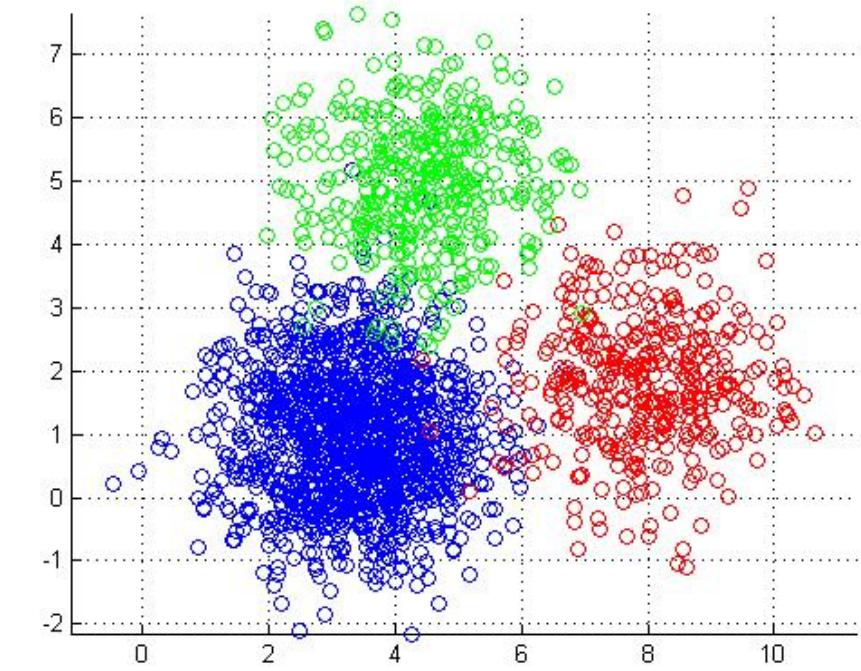
Classification

```
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0  
1 1 1 1 1 1 1 1 1 1 1 1 1 1 1  
2 2 2 2 2 2 2 2 2 2 2 2 2 2 2  
3 3 3 3 3 3 3 3 3 3 3 3 3 3 3  
4 4 4 4 4 4 4 4 4 4 4 4 4 4 4  
5 5 5 5 5 5 5 5 5 5 5 5 5 5 5  
6 6 6 6 6 6 6 6 6 6 6 6 6 6 6  
7 7 7 7 7 7 7 7 7 7 7 7 7 7 7  
8 8 8 8 8 8 8 8 8 8 8 8 8 8 8  
9 9 9 9 9 9 9 9 9 9 9 9 9 9 9
```

Regression



Clustering



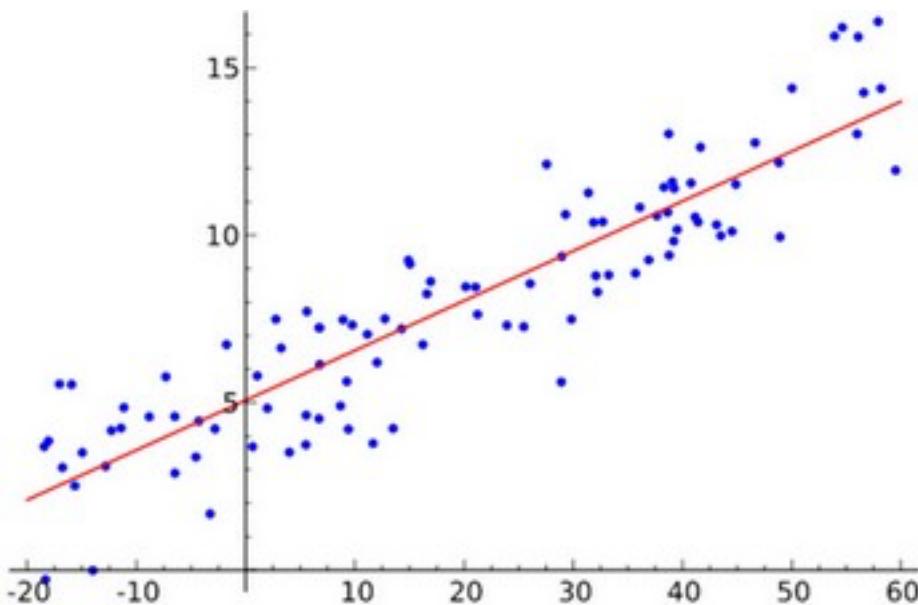
“A computer program is said to **learn** from *experience E* with respect to some class of *tasks T* and *performance measure P*, if its performance at tasks in *T*, as measured by *P*, improves with experience *E*” - T. Michell (1997)

Supervised Learning, Unsupervised Learning

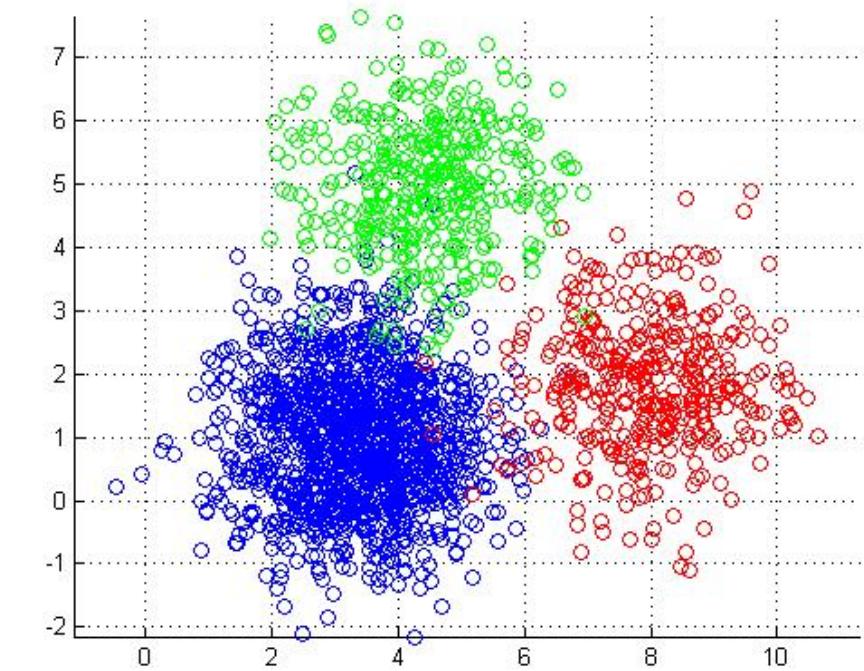
Classification

0 0 0 0 0 0 0 0 0 0 0 0 0
1 1 1 1 1 1 1 1 1 1 1 1 1
2 2 2 2 2 2 2 2 2 2 2 2 2 0
3 3 3 3 3 3 3 3 3 3 3 3 3 3
4 4 4 4 4 4 4 4 4 4 4 4 4 4
5 5 5 5 5 5 5 5 5 5 5 5 5 5
6 6 6 6 6 6 6 6 6 6 6 6 6 6
7 7 7 7 7 7 7 7 7 7 7 7 7 7
8 8 8 8 8 8 8 8 8 8 8 8 8 8
9 9 9 9 9 9 9 9 9 9 9 9 9 9

Regression



Clustering



Task

discrete target values
 x : pixels (28*28)
 y : 0, 1, 2, 3, ..., 9

real target values
 $x \in (-20, 60)$
 $y \in \mathbb{R}$

no target values
 $x \in (0, 10) \times (-2, 7)$

Performance

0-1 loss function
 $L(\hat{y}, y) = I(\hat{y} \neq y)$

L2 loss function
 $L(f, \hat{f}) = \|f - \hat{f}\|_2^2$

$L(\hat{y}, y) = I(\hat{y} \neq y)$
 $L(f, \hat{f}) = \|f - \hat{f}\|_2^2$

Experience

labeled data
(pixels) \rightarrow (number)

labeled data
(x) \rightarrow (y)

unlabeled data
(x_1, x_2)

Semi-Supervised Learning

20

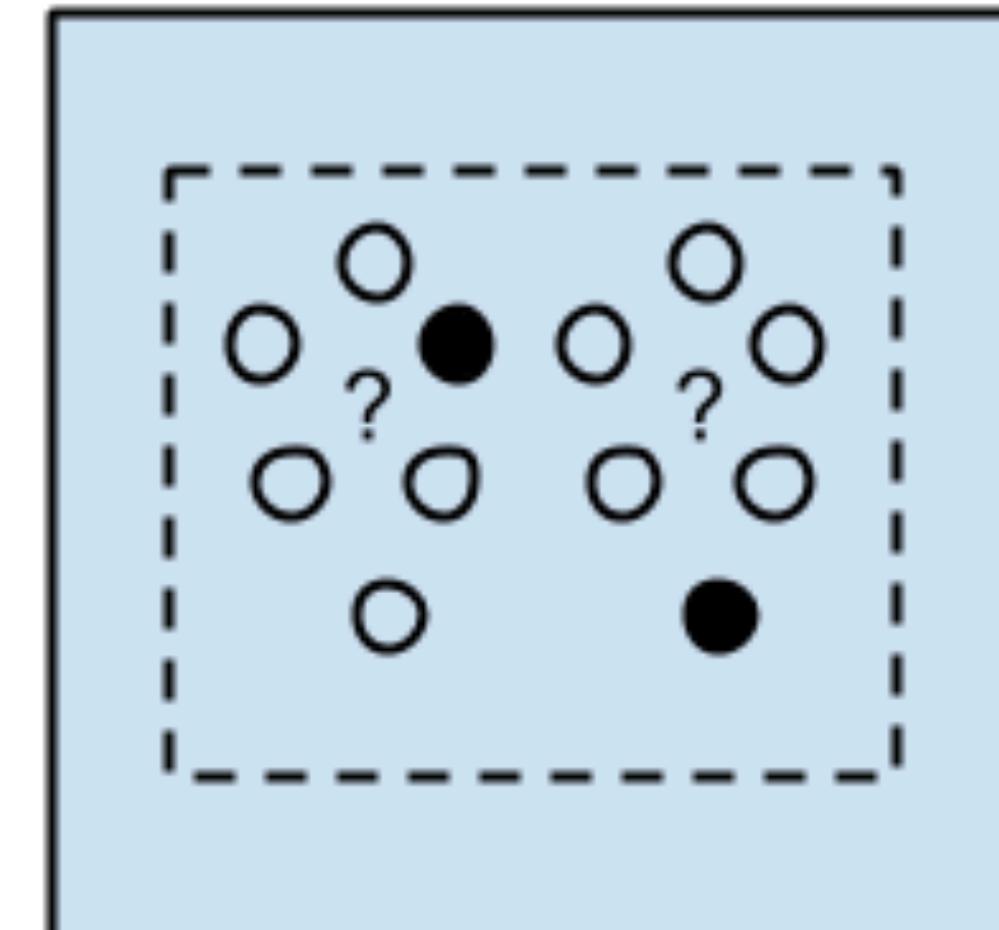
Semi-Supervised Learning

Input data is a mixture of labeled and unlabelled examples.

There is a desired prediction problem but the model must learn the structures to organize the data as well as make predictions.

Example problems are classification and regression.

Example algorithms are extensions to other flexible methods that make assumptions about how to model the unlabeled data.



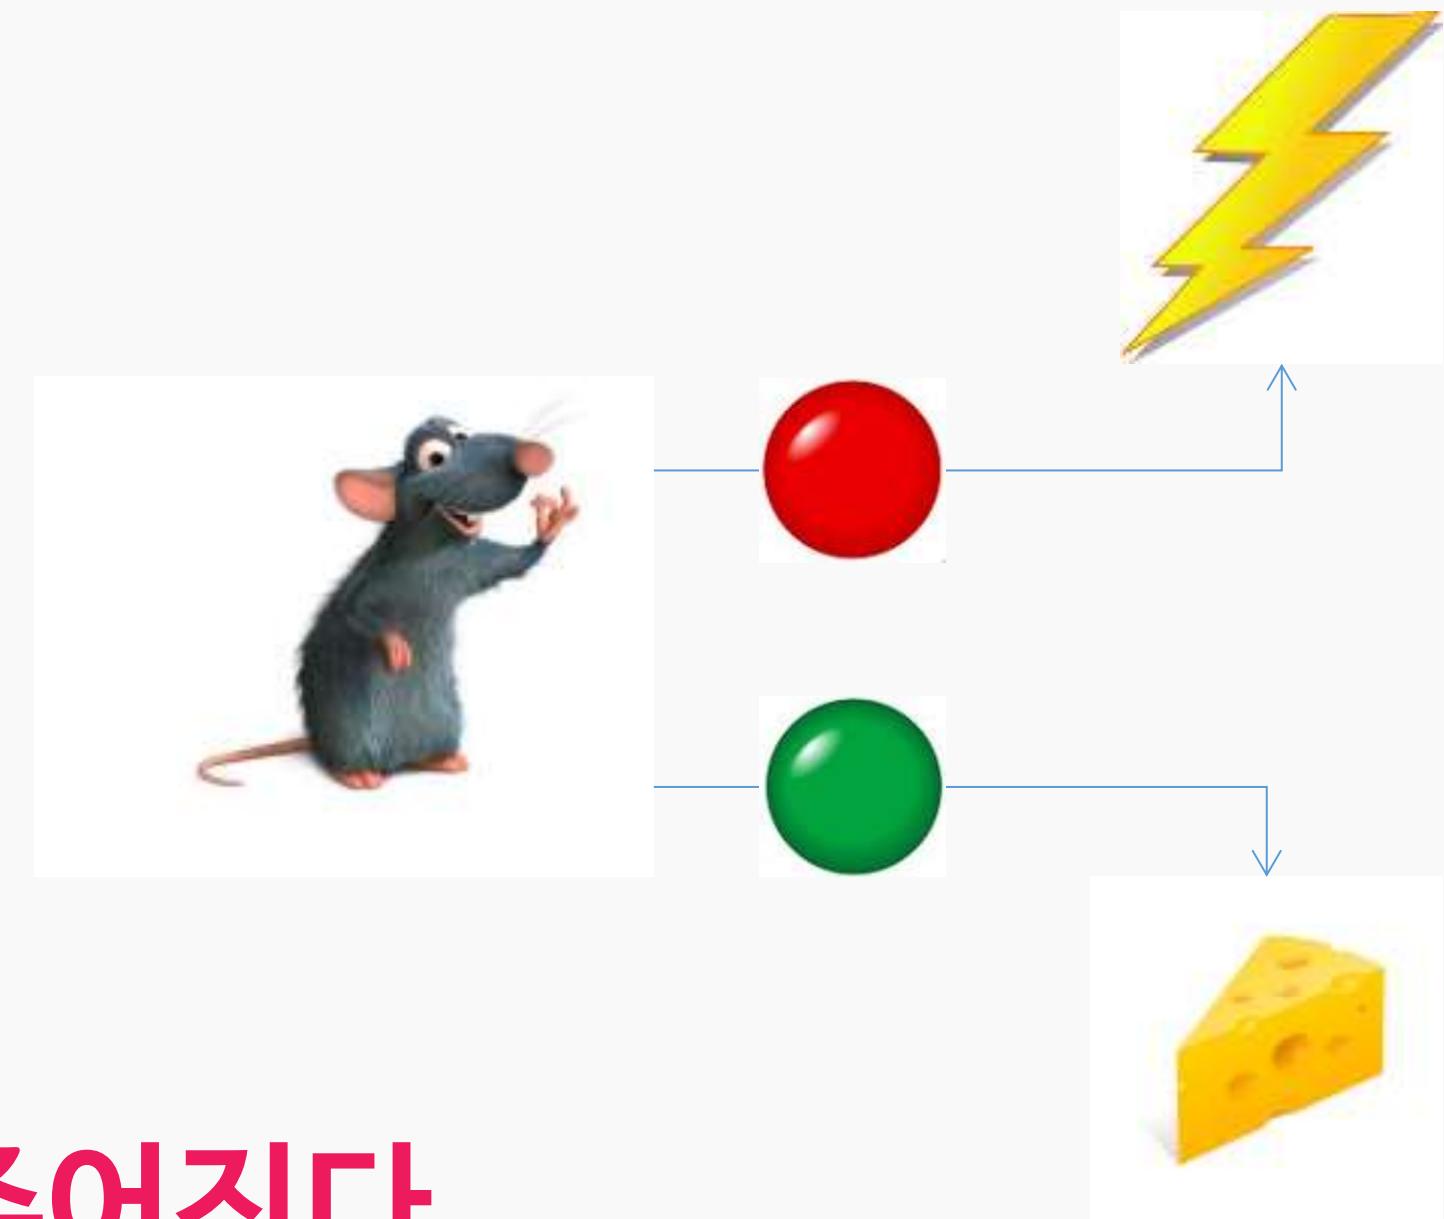
Semi-supervised
Learning Algorithms

Reinforcement Learning

21

Reinforcement Learning

Reward

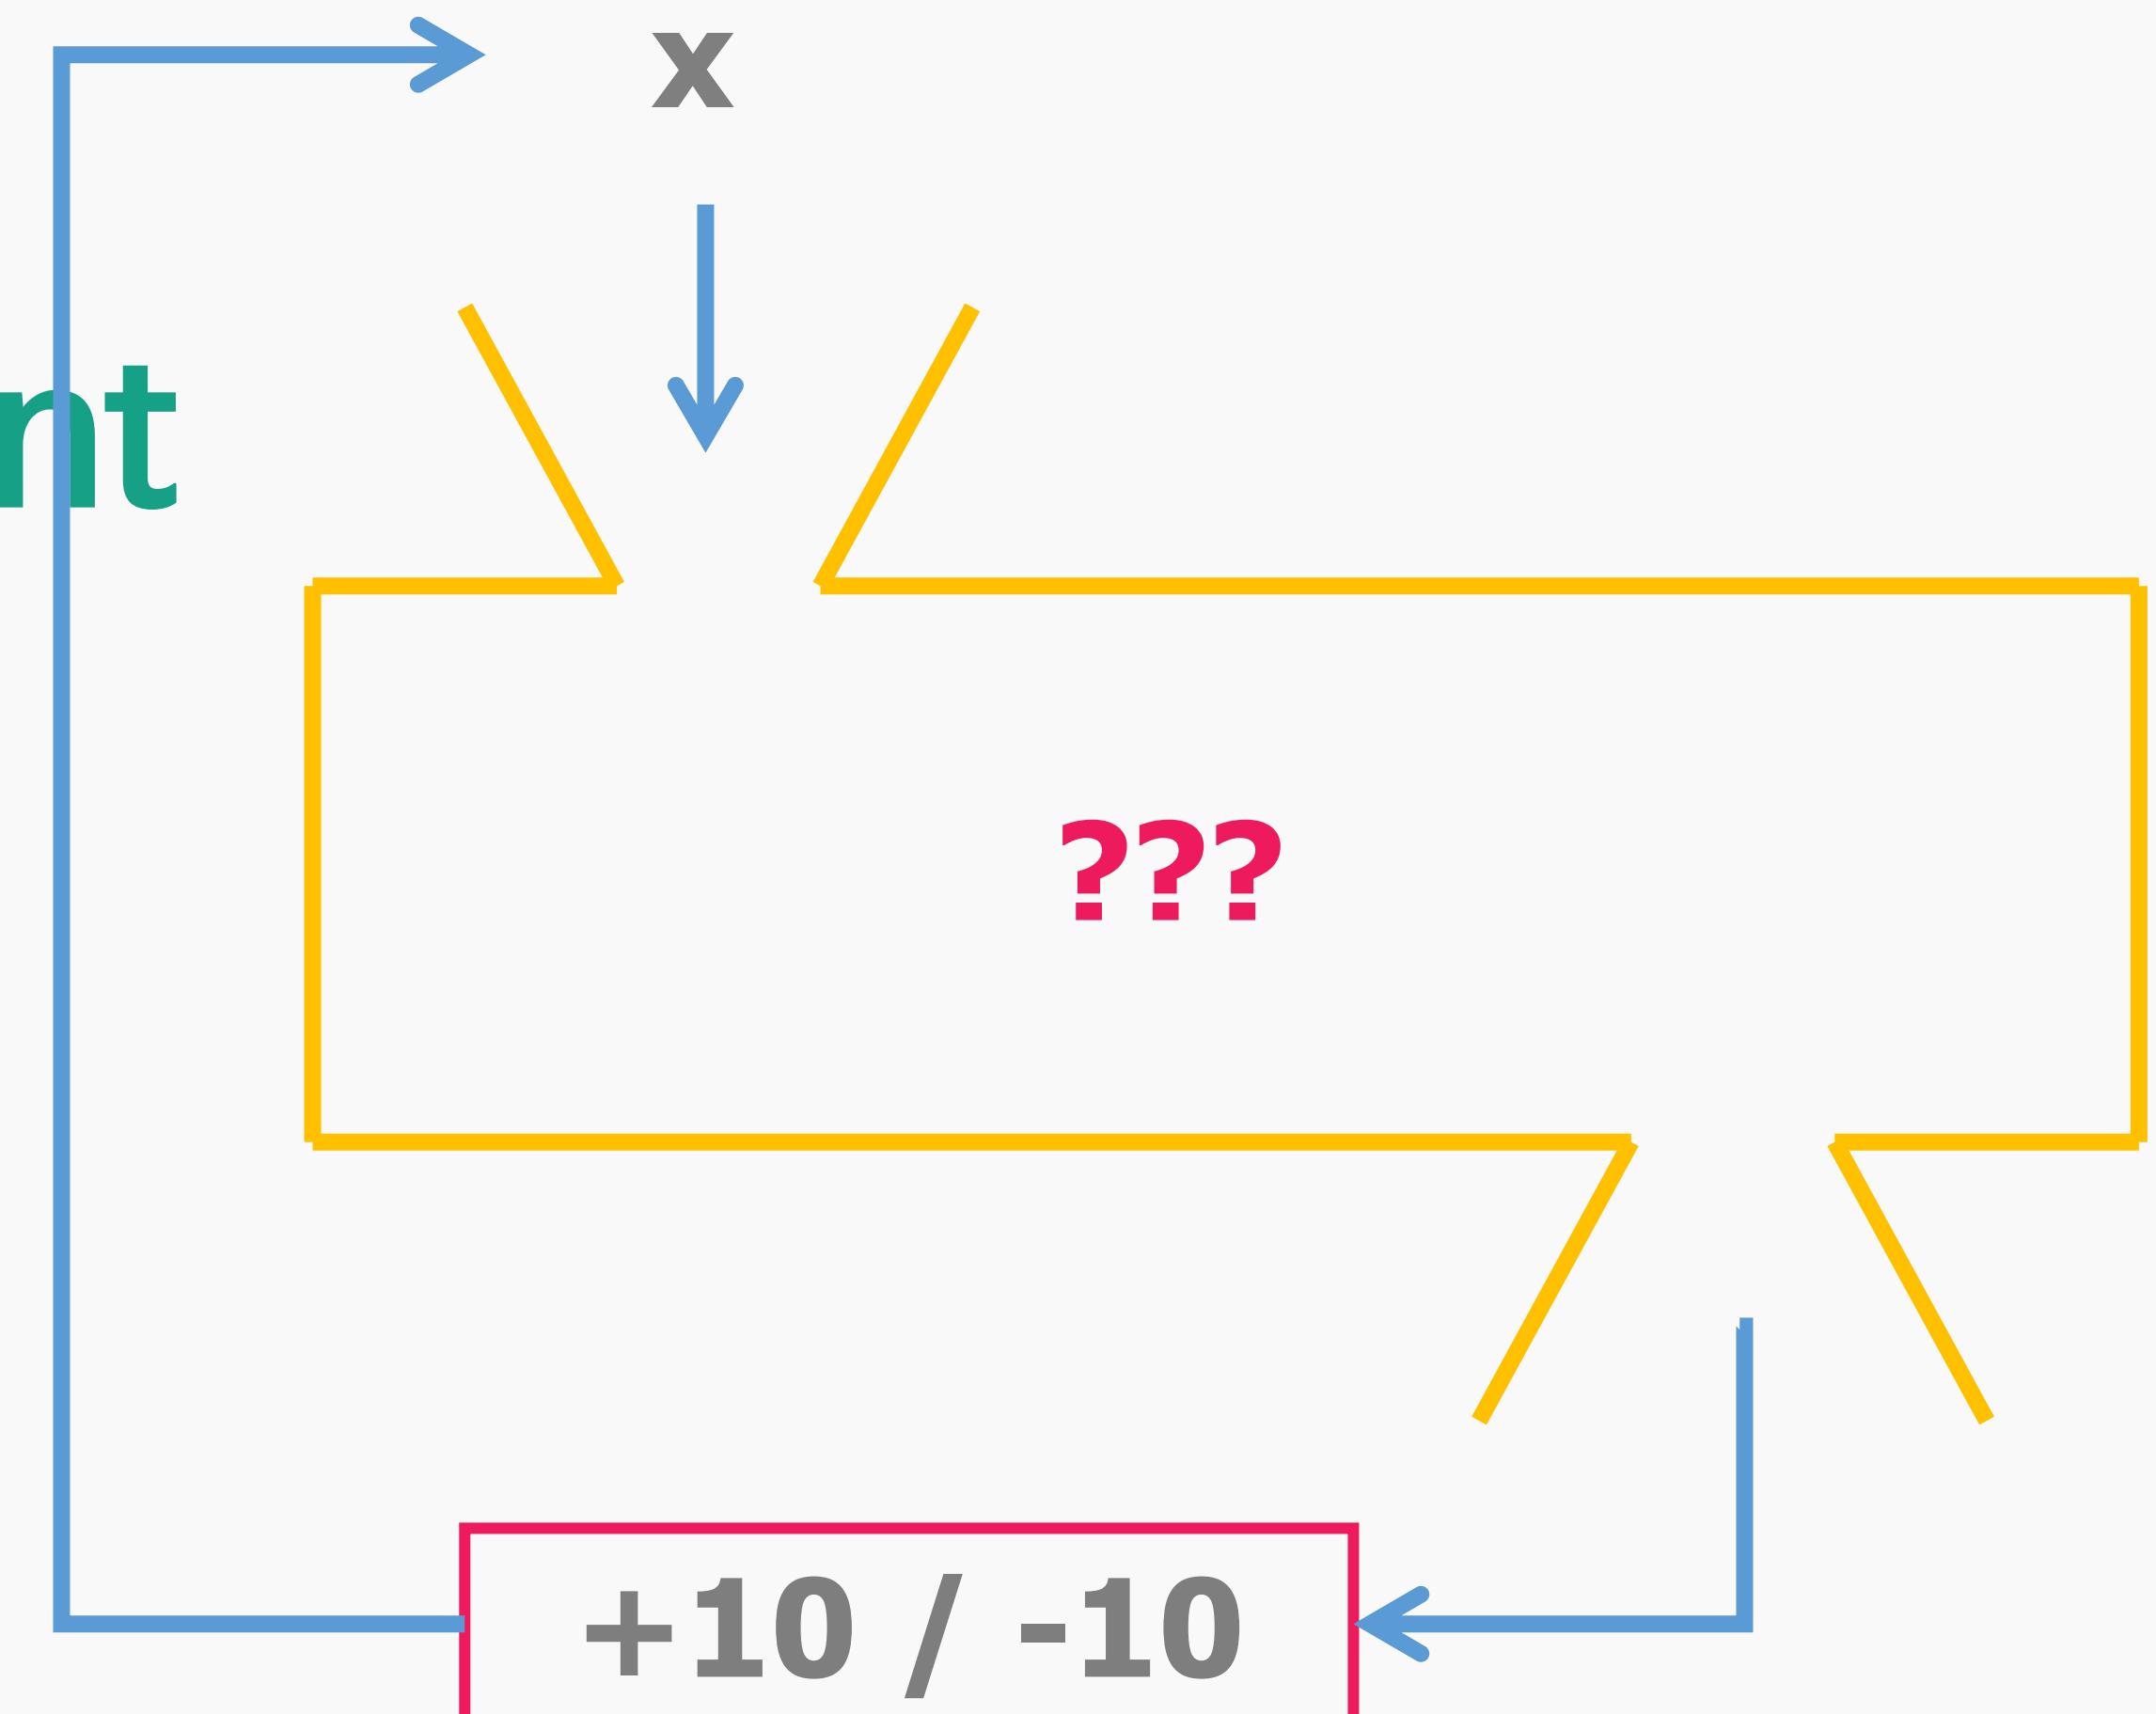


- (정답이 아닌) reward가 주어진다.

Reinforcement Learning

22

Reinforcement
Learning



A Tour of Machine Learning Algorithms

23



Regression Algorithms

24

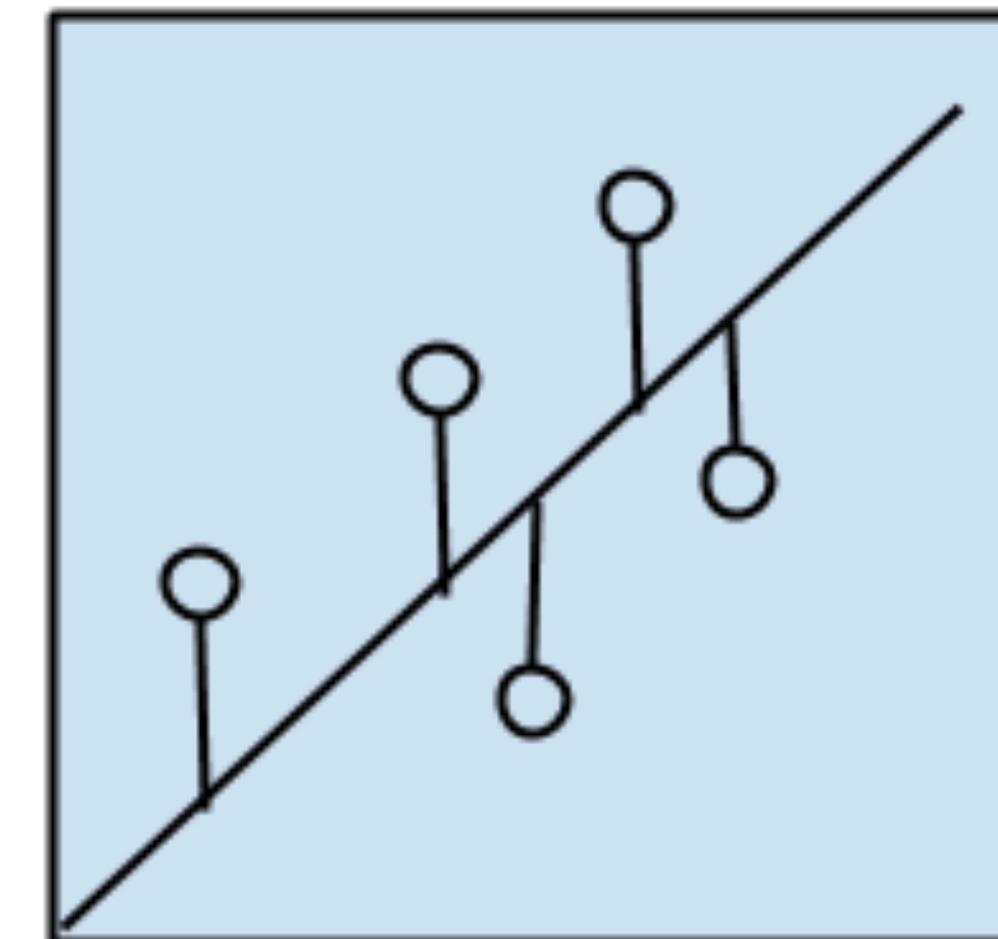
Regression Algorithms

Regression is concerned with modeling the relationship between variables that is iteratively refined using a measure of error in the predictions made by the model.

Regression methods are a workhorse of statistics and have been co-opted into statistical machine learning. This may be confusing because we can use regression to refer to the class of problem and the class of algorithm. Really, regression is a process.

The most popular regression algorithms are:

- Ordinary Least Squares Regression (OLSR)
- Linear Regression
- Logistic Regression
- Stepwise Regression
- Multivariate Adaptive Regression Splines (MARS)
- Locally Estimated Scatterplot Smoothing (LOESS)



Regression Algorithms

Instance-Based Algorithms

25

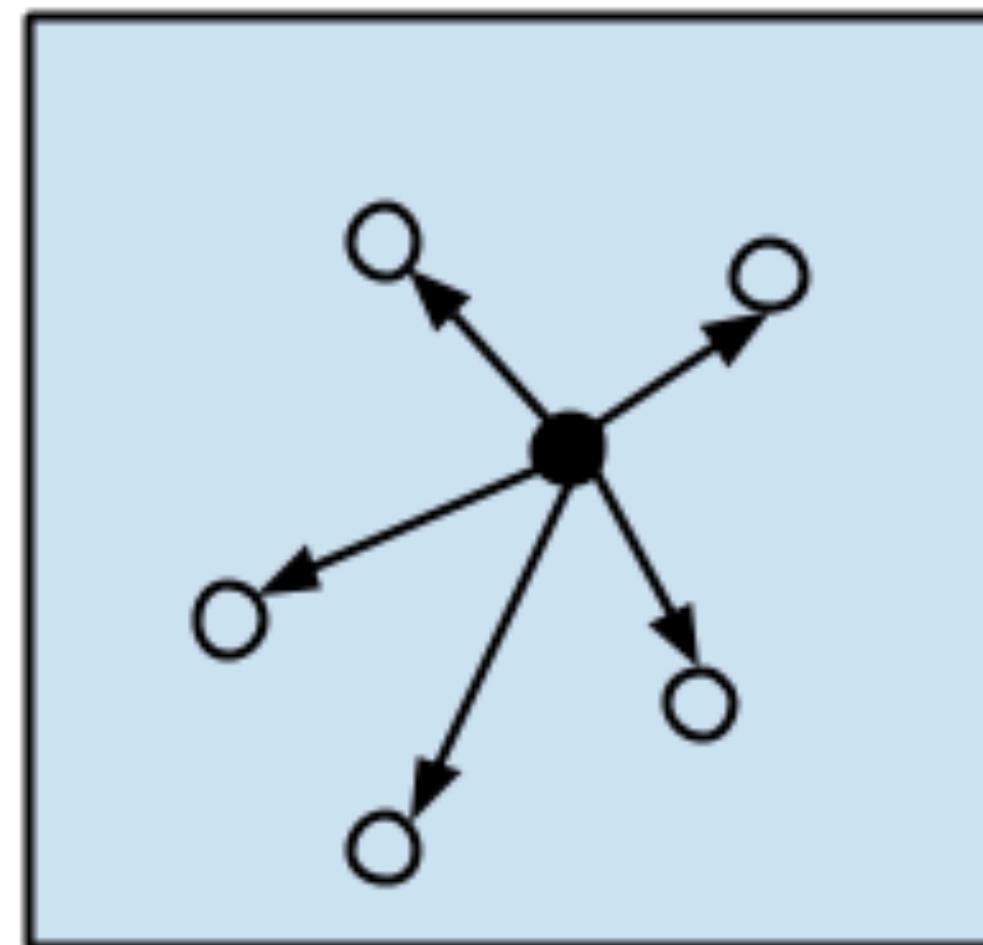
Instance-based Algorithms

Instance-based learning model is a decision problem with instances or examples of training data that are deemed important or required to the model.

Such methods typically build up a database of example data and compare new data to the database using a similarity measure in order to find the best match and make a prediction. For this reason, instance-based methods are also called winner-take-all methods and memory-based learning. Focus is put on the representation of the stored instances and similarity measures used between instances.

The most popular instance-based algorithms are:

- k-Nearest Neighbor (kNN)
- Learning Vector Quantization (LVQ)
- Self-Organizing Map (SOM)
- Locally Weighted Learning (LWL)



Instance-based
Algorithms

Regularization Algorithms

26

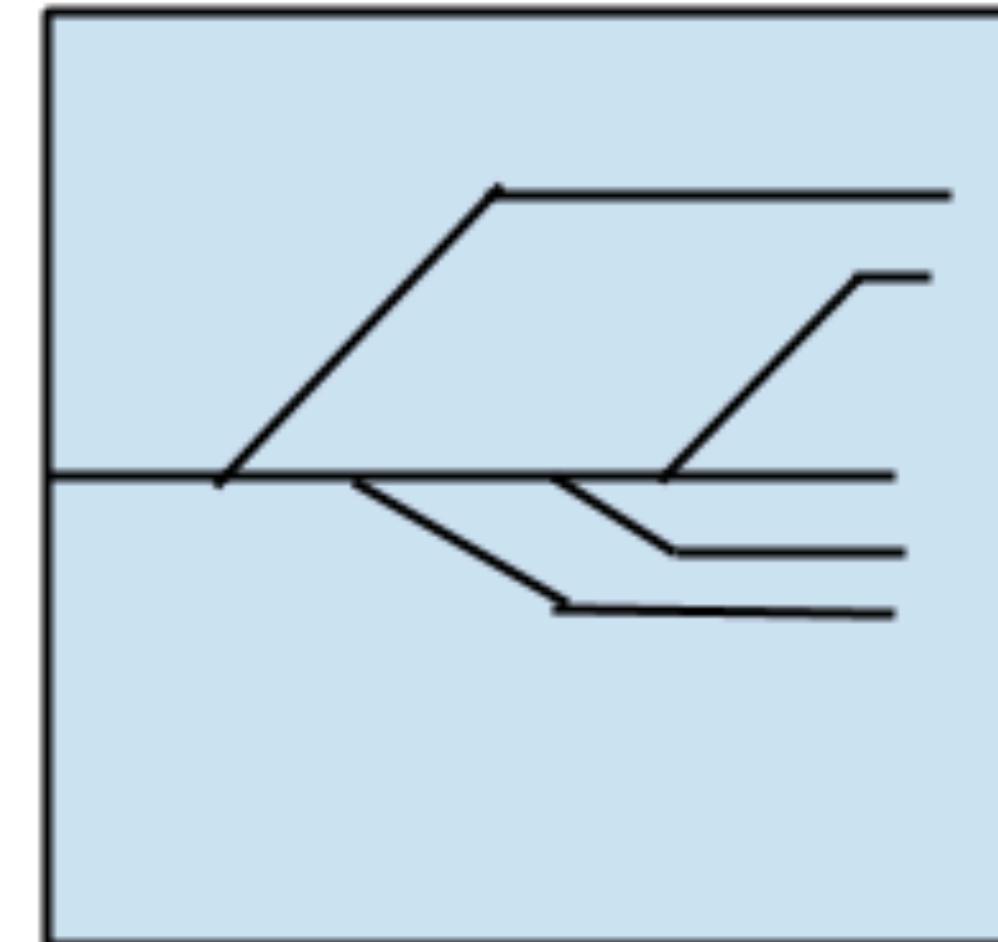
Regularization Algorithms

An extension made to another method (typically regression methods) that penalizes models based on their complexity, favoring simpler models that are also better at generalizing.

I have listed regularization algorithms separately here because they are popular, powerful and generally simple modifications made to other methods.

The most popular regularization algorithms are:

- Ridge Regression
- Least Absolute Shrinkage and Selection Operator (LASSO)
- Elastic Net
- Least-Angle Regression (LARS)



Regularization
Algorithms

Decision Tree Algorithms

27

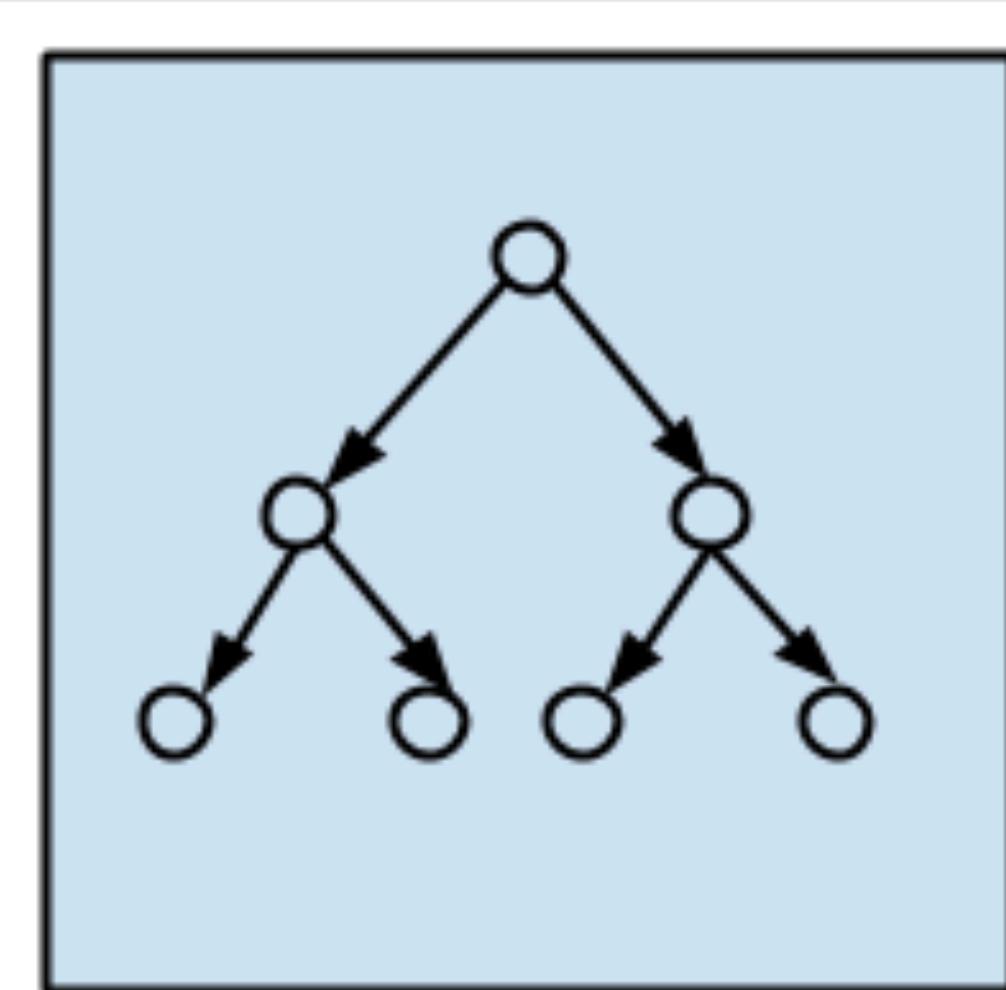
Decision Tree Algorithms

Decision tree methods construct a model of decisions made based on actual values of attributes in the data.

Decisions fork in tree structures until a prediction decision is made for a given record. Decision trees are trained on data for classification and regression problems. Decision trees are often fast and accurate and a big favorite in machine learning.

The most popular decision tree algorithms are:

- Classification and Regression Tree (CART)
- Iterative Dichotomiser 3 (ID3)
- C4.5 and C5.0 (different versions of a powerful approach)
- Chi-squared Automatic Interaction Detection (CHAID)
- Decision Stump
- M5
- Conditional Decision Trees



Decision Tree
Algorithms

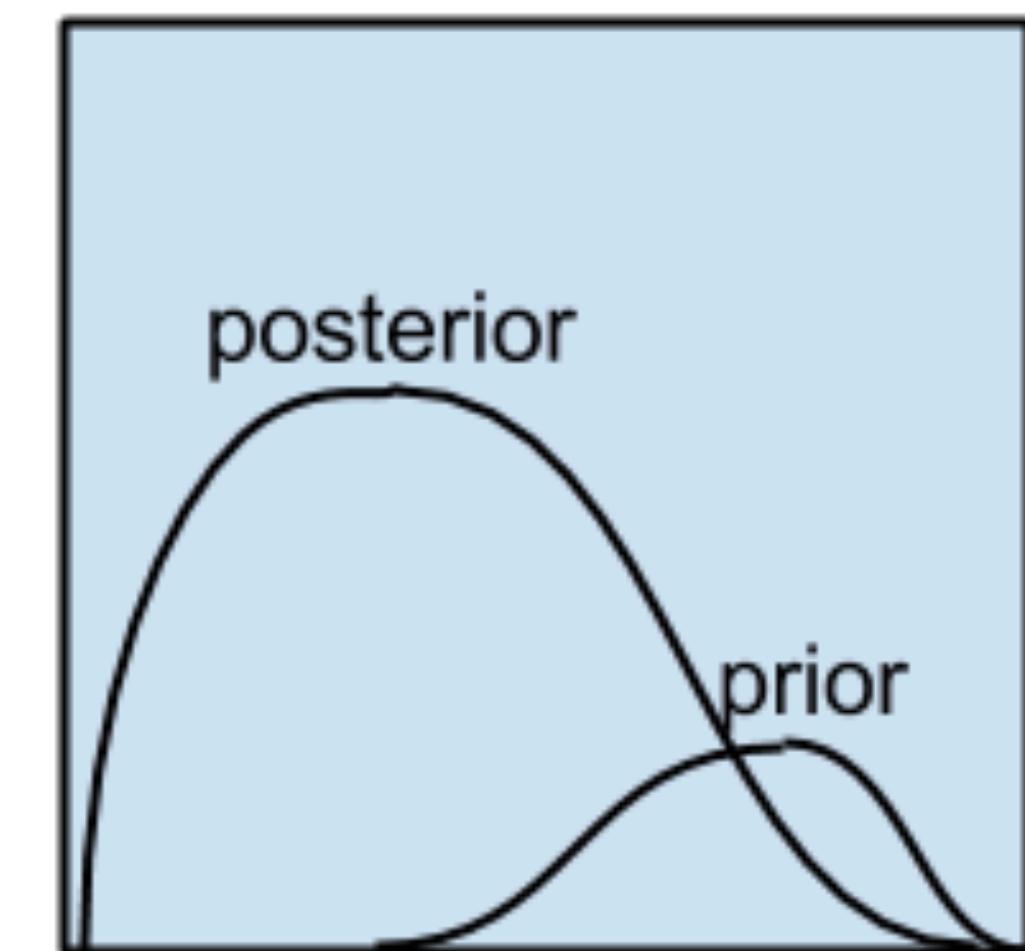
Bayesian Algorithms

Bayesian Algorithms

Bayesian methods are those that explicitly apply Bayes' Theorem for problems such as classification and regression.

The most popular Bayesian algorithms are:

- Naive Bayes
- Gaussian Naive Bayes
- Multinomial Naive Bayes
- Averaged One-Dependence Estimators (AODE)
- Bayesian Belief Network (BBN)
- Bayesian Network (BN)



Bayesian Algorithms

See Appendix for Bayes' Theorem

Clustering Algorithms

29

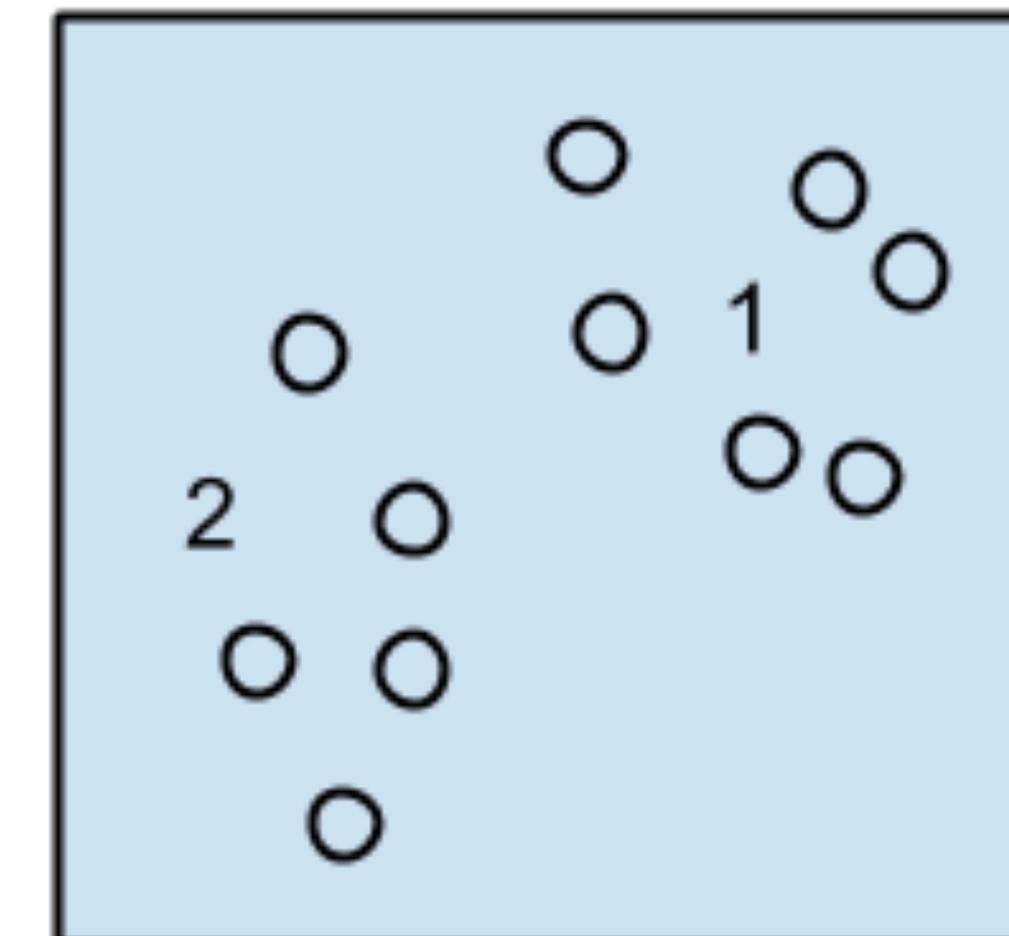
Clustering Algorithms

Clustering, like regression, describes the class of problem
and the class of methods.

Clustering methods are typically organized by the modeling approaches such as centroid-based and hierachal. All methods are concerned with using the inherent structures in the data to best organize the data into groups of maximum commonality.

The most popular clustering algorithms are:

- k-Means
- k-Medians
- Expectation Maximisation (EM)
- Hierarchical Clustering



Clustering Algorithms

Association Rule Learning Algorithms

30

Association Rule Learning Algorithms

Association rule learning methods extract rules that best explain observed relationships between variables in data.

These rules can discover important and commercially useful associations in large multidimensional datasets that can be exploited by an organization.

The most popular association rule learning algorithms are:

- Apriori algorithm
- Eclat algorithm

(A,B) → C

(D,E) → F

(A,E) → G

Association Rule
Learning Algorithms

Artificial Neural Network Algorithms

31

Artificial Neural Network Algorithms

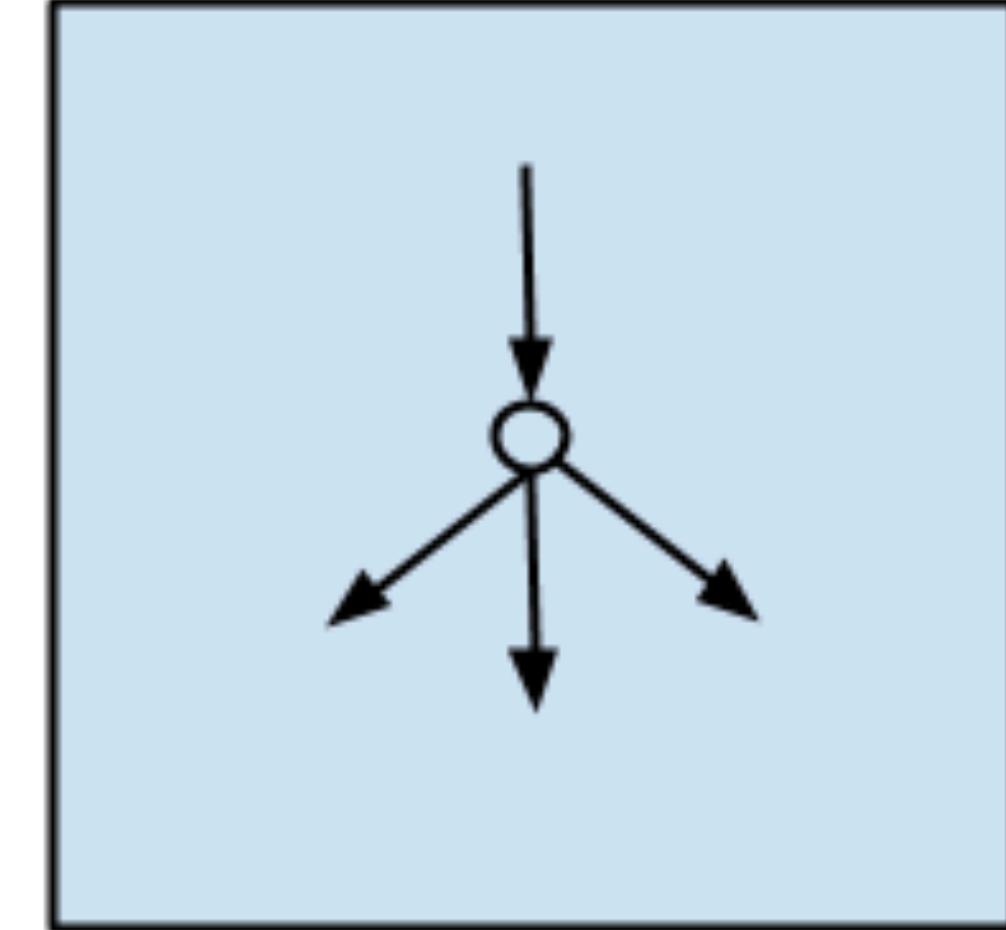
Artificial Neural Networks are models that are inspired by the structure and/or function of biological neural networks.

They are a class of pattern matching that are commonly used for regression and classification problems but are really an enormous subfield comprised of hundreds of algorithms and variations for all manner of problem types.

Note that I have separated out Deep Learning from neural networks because of the massive growth and popularity in the field. Here we are concerned with the more classical methods.

The most popular artificial neural network algorithms are:

- Perceptron
- Back-Propagation
- Hopfield Network
- Radial Basis Function Network (RBFN)



Artificial Neural Network
Algorithms

Deep Learning Algorithms

32

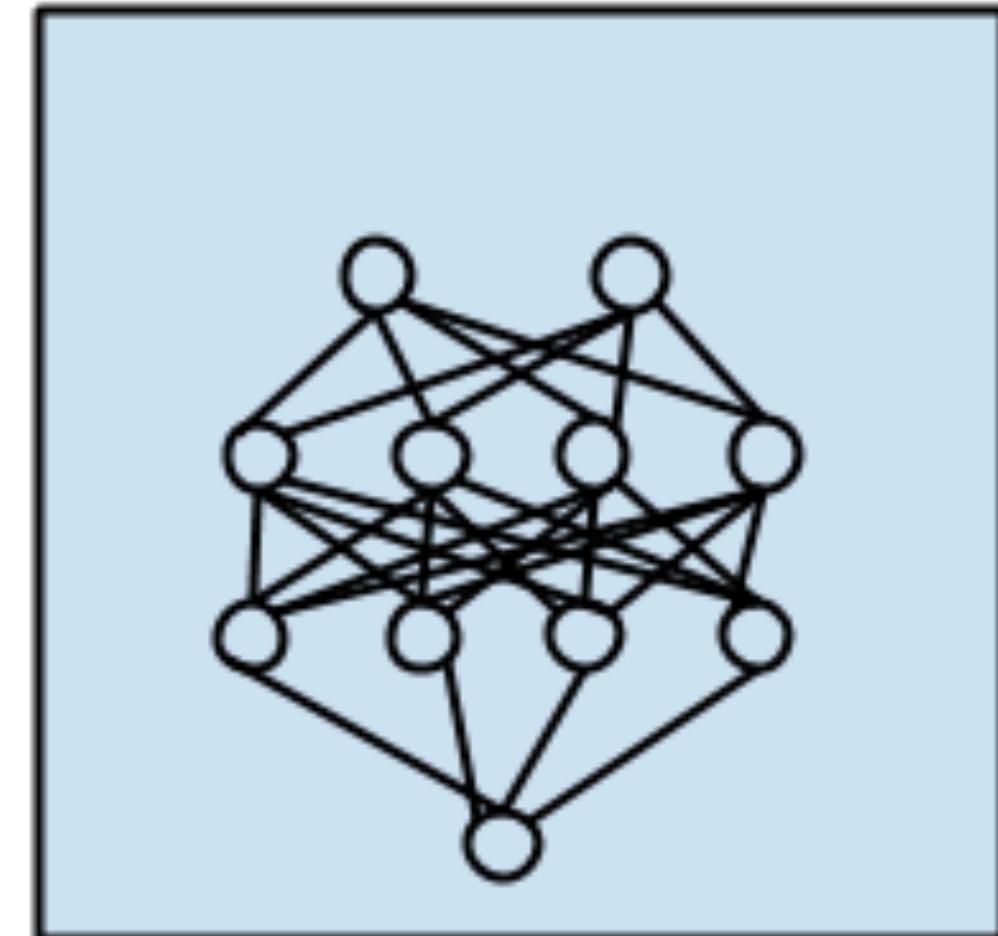
Deep Learning Algorithms

Deep Learning methods are a modern update to Artificial Neural Networks that exploit abundant cheap computation.

They are concerned with building much larger and more complex neural networks and, as commented on above, many methods are concerned with semi-supervised learning problems where large datasets contain very little labeled data.

The most popular deep learning algorithms are:

- Deep Boltzmann Machine (DBM)
- Deep Belief Networks (DBN)
- Convolutional Neural Network (CNN)
- Stacked Auto-Encoders



Deep Learning
Algorithms

Dimensionality Reduction Algorithms

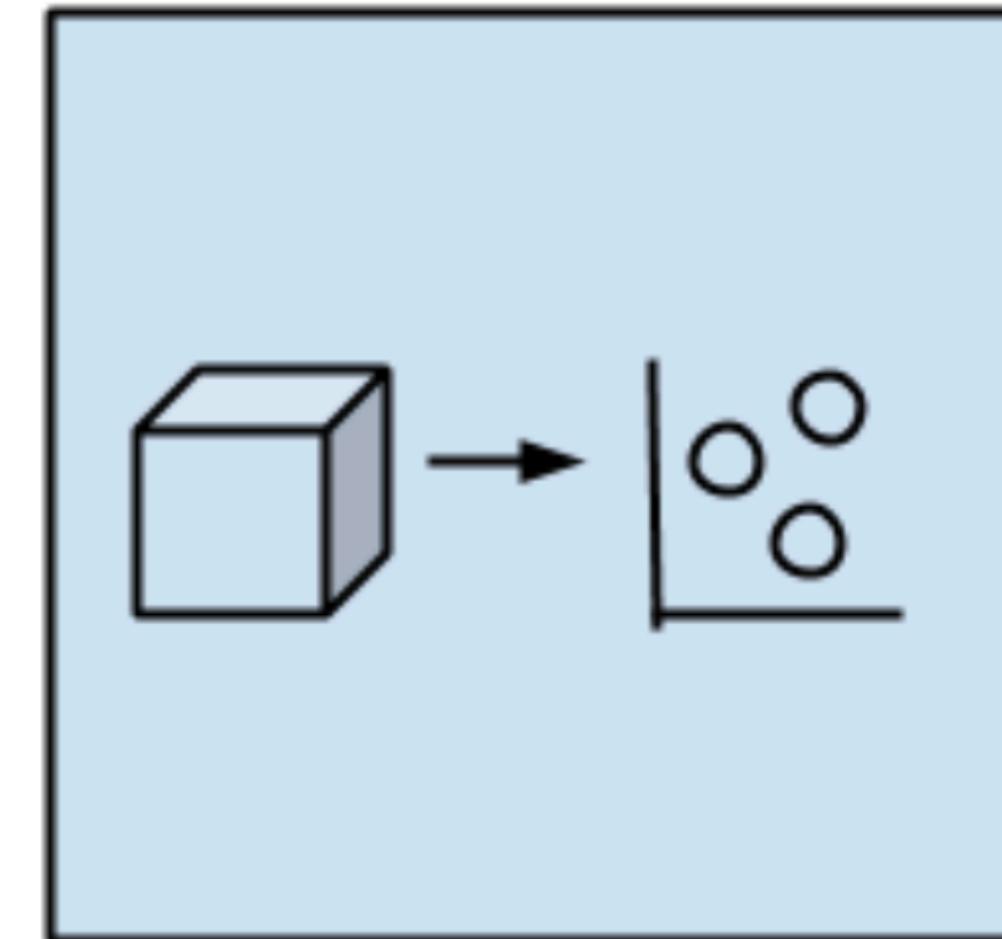
33

Dimensionality Reduction Algorithms

Like clustering methods, dimensionality reduction seek and exploit the inherent structure in the data, but in this case in an unsupervised manner or order to summarize or describe data using less information.

This can be useful to visualize dimensional data or to simplify data which can then be used in a supervised learning method. Many of these methods can be adapted for use in classification and regression.

- Principal Component Analysis (PCA)
- Principal Component Regression (PCR)
- Partial Least Squares Regression (PLSR)
- Sammon Mapping
- Multidimensional Scaling (MDS)
- Projection Pursuit
- Linear Discriminant Analysis (LDA)
- Mixture Discriminant Analysis (MDA)
- Quadratic Discriminant Analysis (QDA)
- Flexible Discriminant Analysis (FDA)



Dimensional Reduction
Algorithms

Ensemble Algorithms

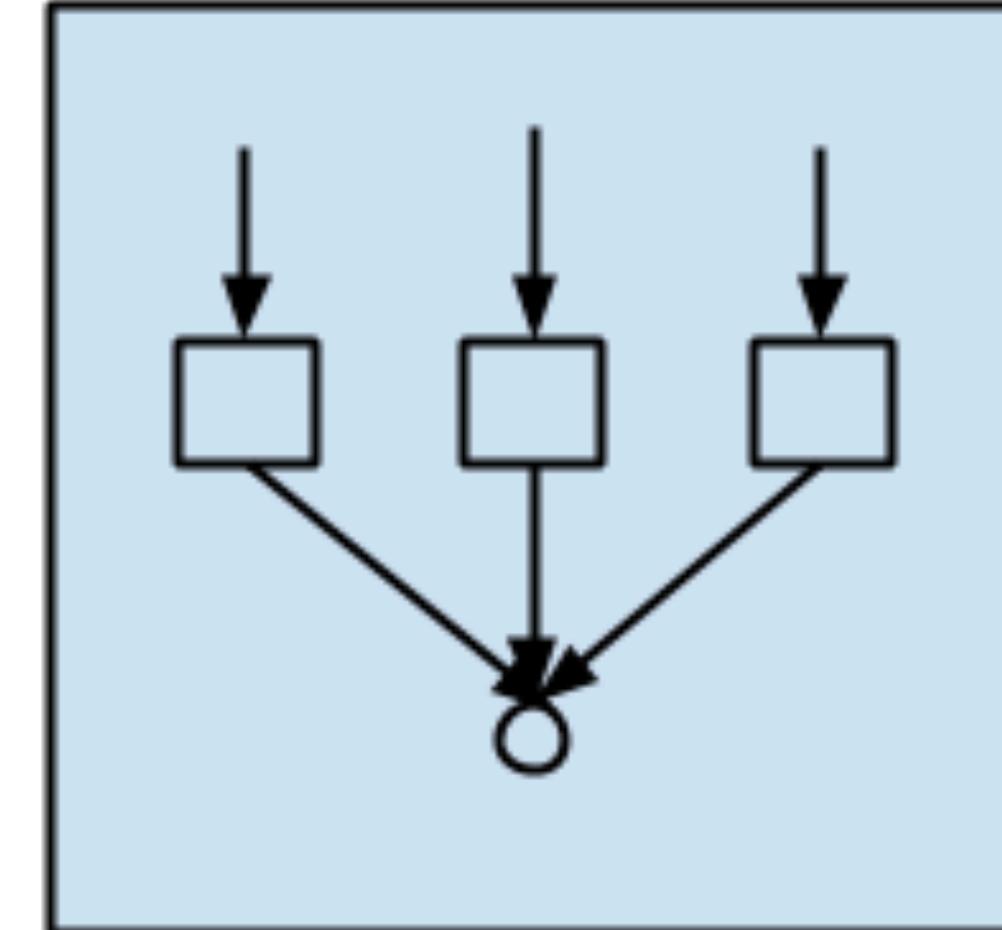
34

Ensemble Algorithms

Ensemble methods are models composed of multiple weaker models that are independently trained and whose predictions are combined in some way to make the overall prediction.

Much effort is put into what types of weak learners to combine and the ways in which to combine them. This is a very powerful class of techniques and as such is very popular.

- Boosting
- Bootstrapped Aggregation (Bagging)
- AdaBoost
- Stacked Generalization (blending)
- Gradient Boosting Machines (GBM)
- Gradient Boosted Regression Trees (GBRT)
- Random Forest



Ensemble Algorithms

The End of the Tour!

35



review articles

DOI:10.1145/2347736.2347755

Tapping into the “folk knowledge” needed to advance machine learning applications.

BY PEDRO DOMINGOS

A Few Useful Things to Know About Machine Learning

MACHINE LEARNING SYSTEMS automatically learn programs from data. This is often a very attractive alternative to manually constructing them, and in the last decade the use of machine learning has spread rapidly throughout computer science and beyond. Machine learning is used in Web search, spam filters, recommender systems, ad placement, credit scoring, fraud detection, stock trading, drug design, and many other applications. A recent report from the McKinsey Global Institute asserts that machine learning (a.k.a. data mining or predictive analytics) will be the driver of the next big wave of innovation.¹⁵ Several fine textbooks are available to interested practitioners and researchers (for example, Mitchell¹⁶ and Witten et al.²⁴). However, much of the “folk knowledge” that



is needed to successfully develop machine learning applications is not readily available in them. As a result, many machine learning projects take much longer than necessary or wind up producing less-than-ideal results. Yet much of this folk knowledge is fairly easy to communicate. This is the purpose of this article.

» key insights

- Machine learning algorithms can figure out how to perform important tasks by generalizing from examples. This is often feasible and cost-effective where manual programming is not. As more data becomes available, more ambitious problems can be tackled.
- Machine learning is widely used in computer science and other fields. However, developing successful machine learning applications requires a substantial amount of “black art” that is difficult to find in textbooks.
- This article summarizes 12 key lessons that machine learning researchers and practitioners have learned. These include pitfalls to avoid, important issues to focus on, and answers to common questions.

Introduction

37

- Machine learning algorithms can figure out how to perform important tasks by *generalizing from examples*.
- As *more data* becomes available, more ambitious problems can be tackled.
- This paper contains lot of “*folk knowledge*” that is needed to successfully develop machine learning applications, which is *not* readily available in ML textbooks.

Learning = R + E + O

38

■ Learning = Representation + Evaluation + Optimization

- **Representation**: algorithm, implementation
- **Evaluation**: metric selection, results based on concrete data
- **Optimization**: from off-the-shelf to custom settings

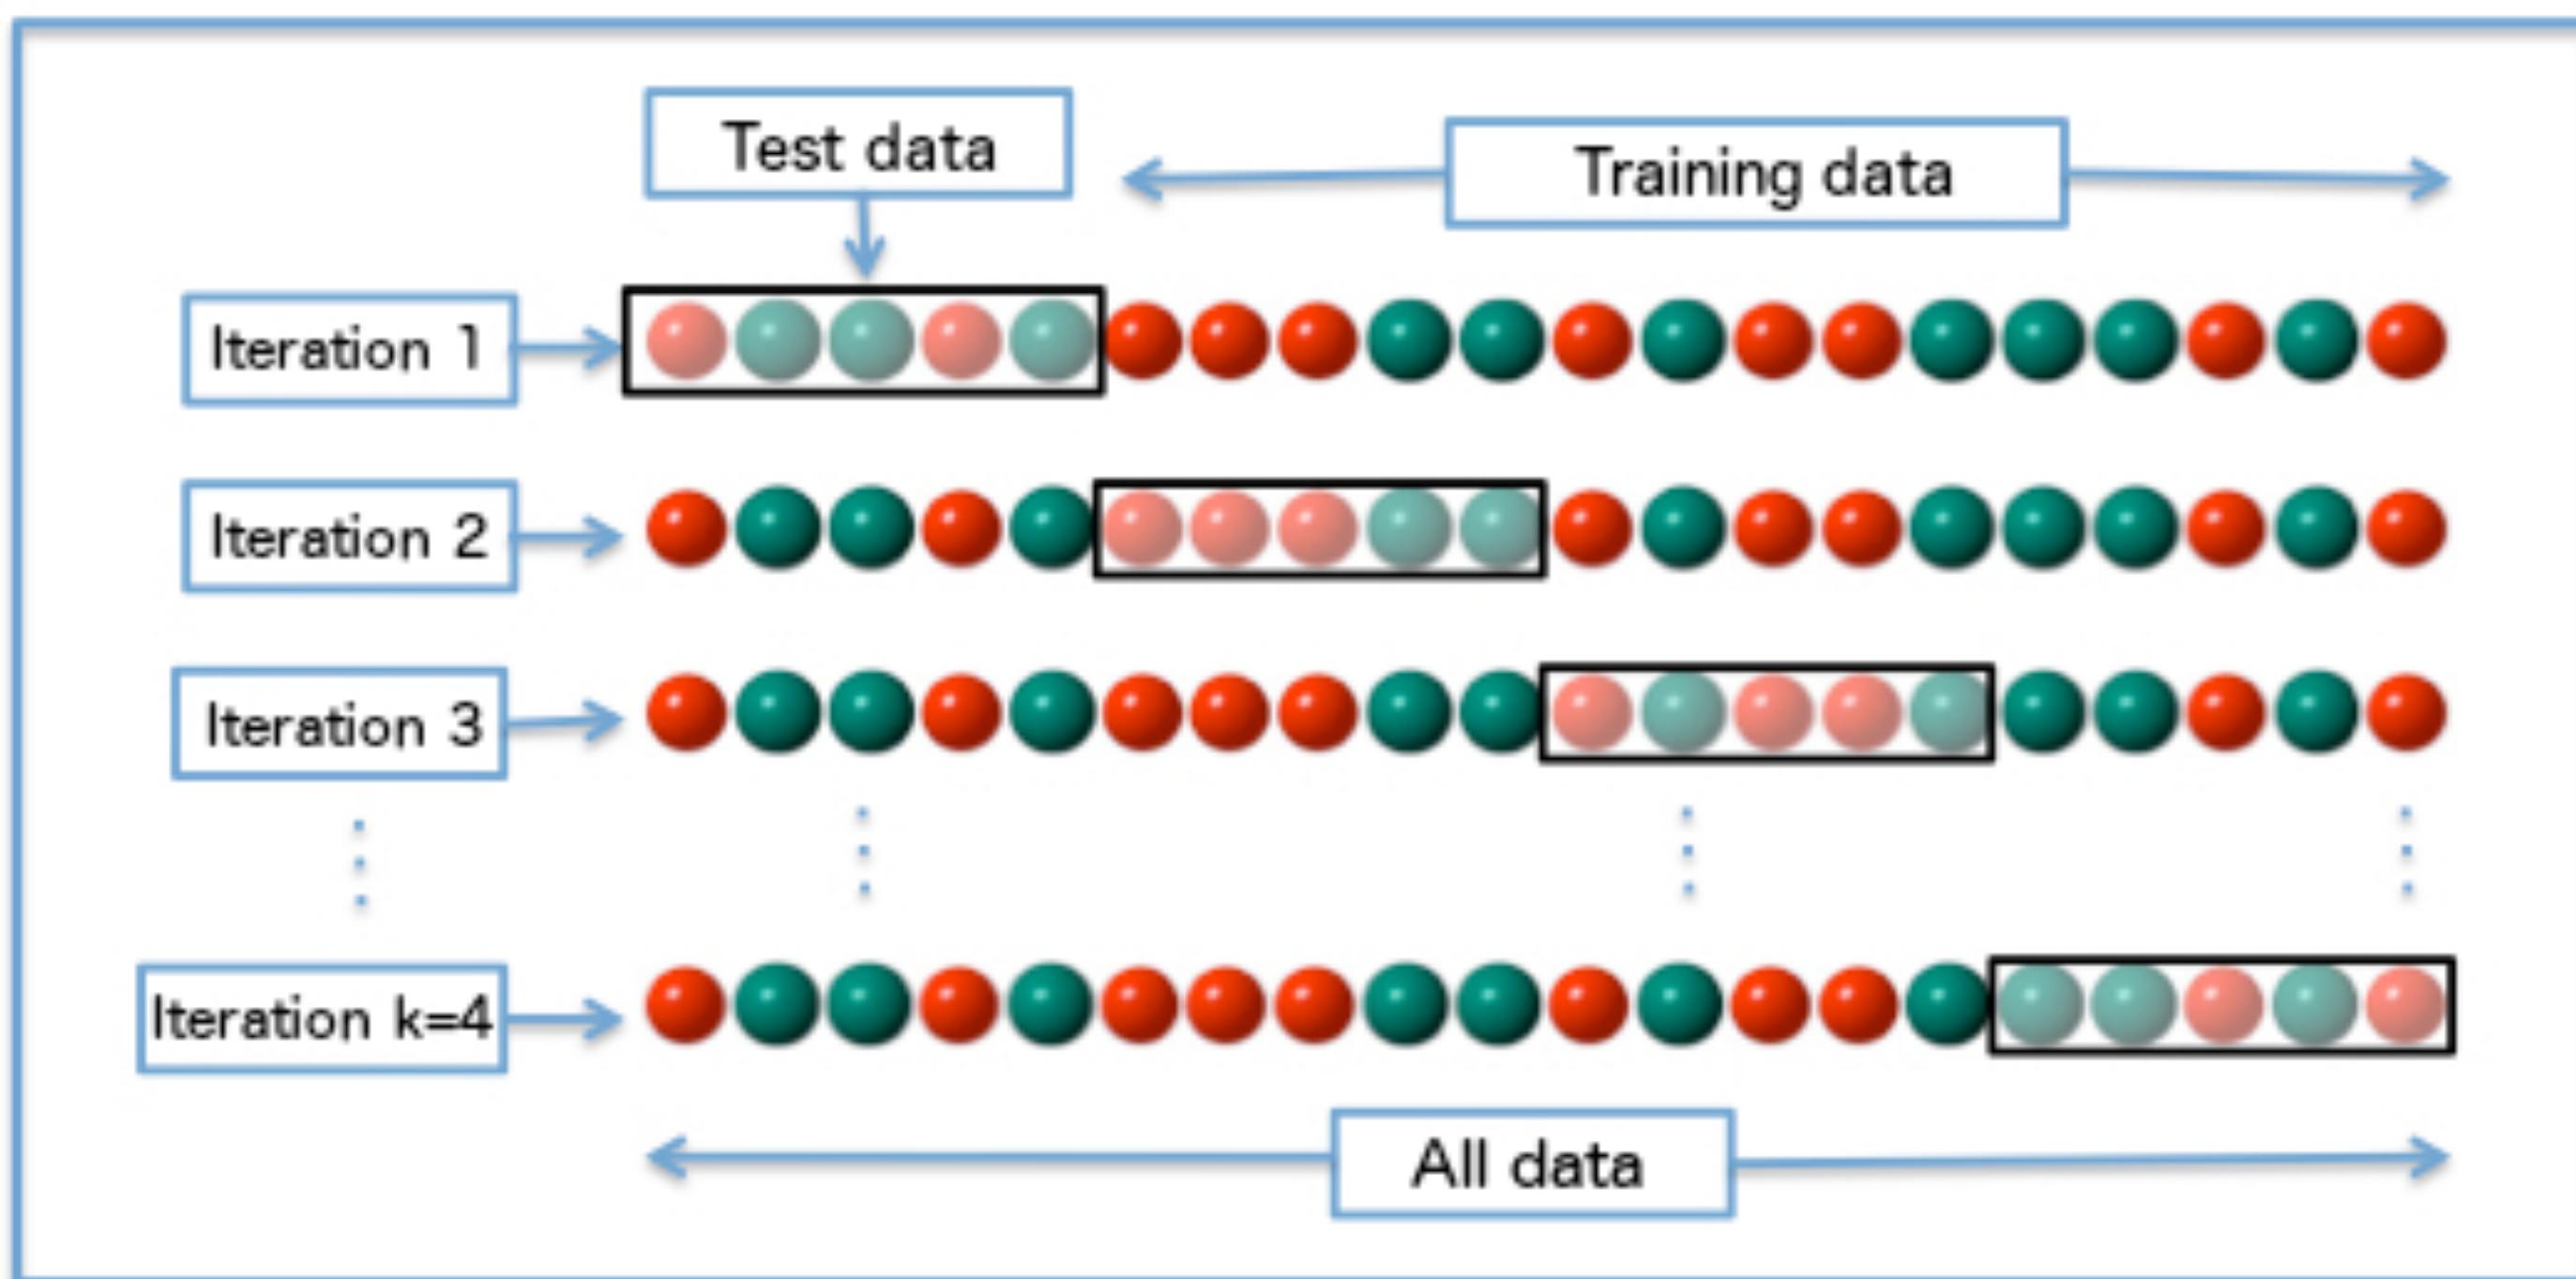
Table 1. The three components of learning algorithms.

Representation	Evaluation	Optimization
Instances	Accuracy/Error rate	Combinatorial optimization
K-nearest neighbor	Precision and recall	Greedy search
Support vector machines	Squared error	Beam search
Hyperplanes	Likelihood	Branch-and-bound
Naive Bayes	Posterior probability	Continuous optimization
Logistic regression	Information gain	Unconstrained
Decision trees	K-L divergence	Gradient descent
Sets of rules	Cost/Utility	Conjugate gradient
Propositional rules	Margin	Quasi-Newton methods
Logic programs		Constrained
Neural networks		Linear programming
Graphical models		Quadratic programming
Bayesian networks		
Conditional random fields		

It's Generalization That Counts

39

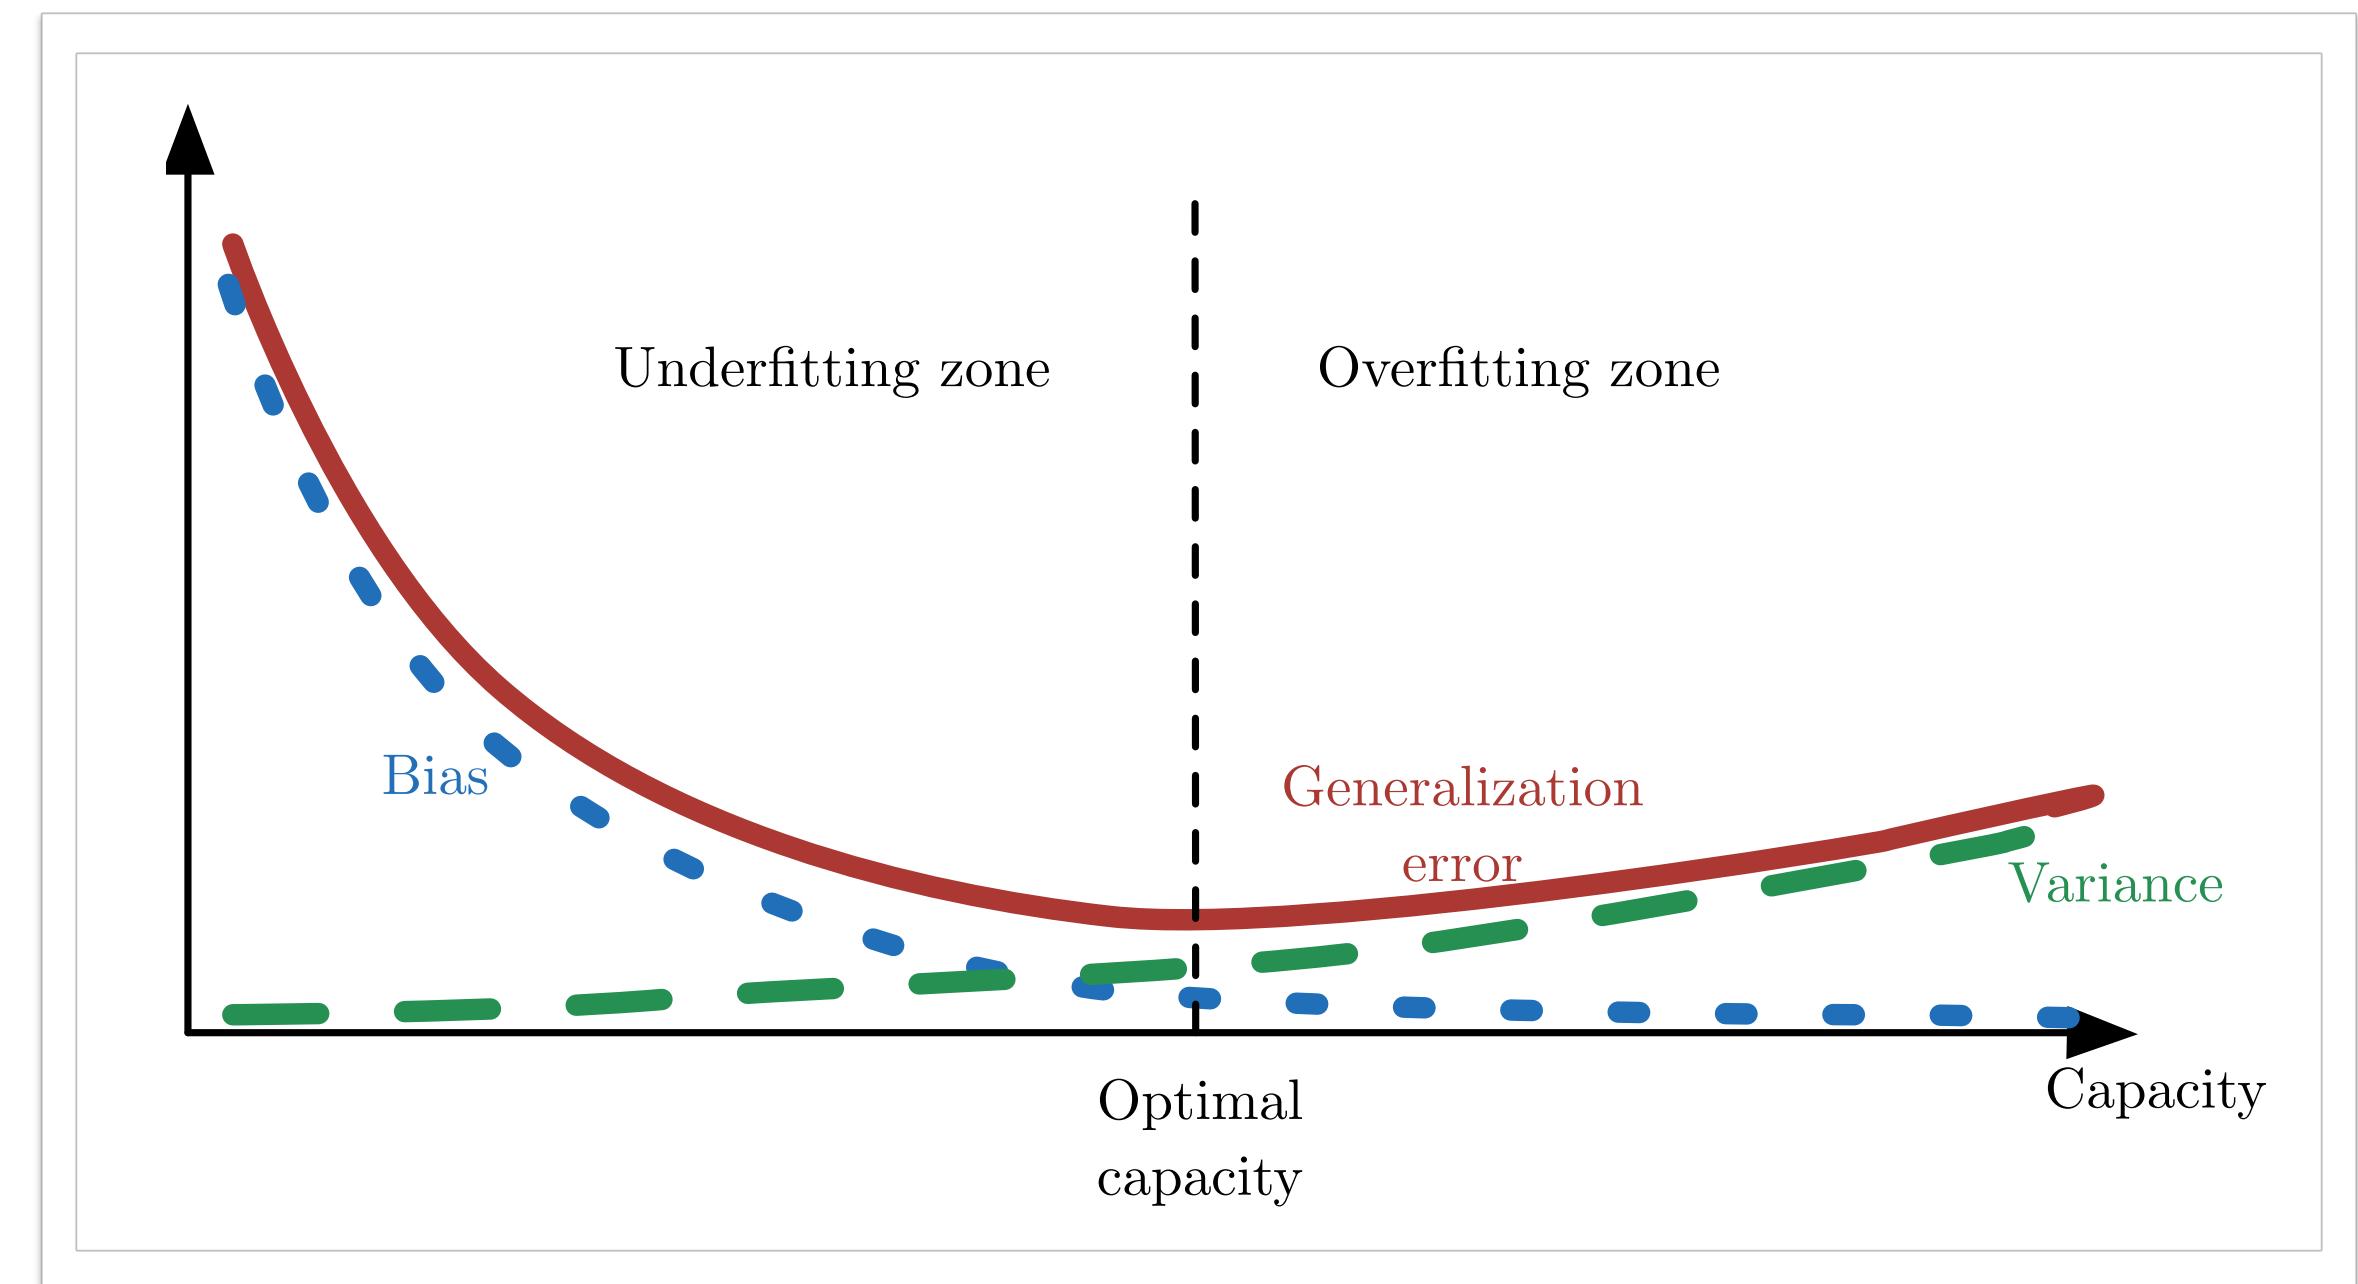
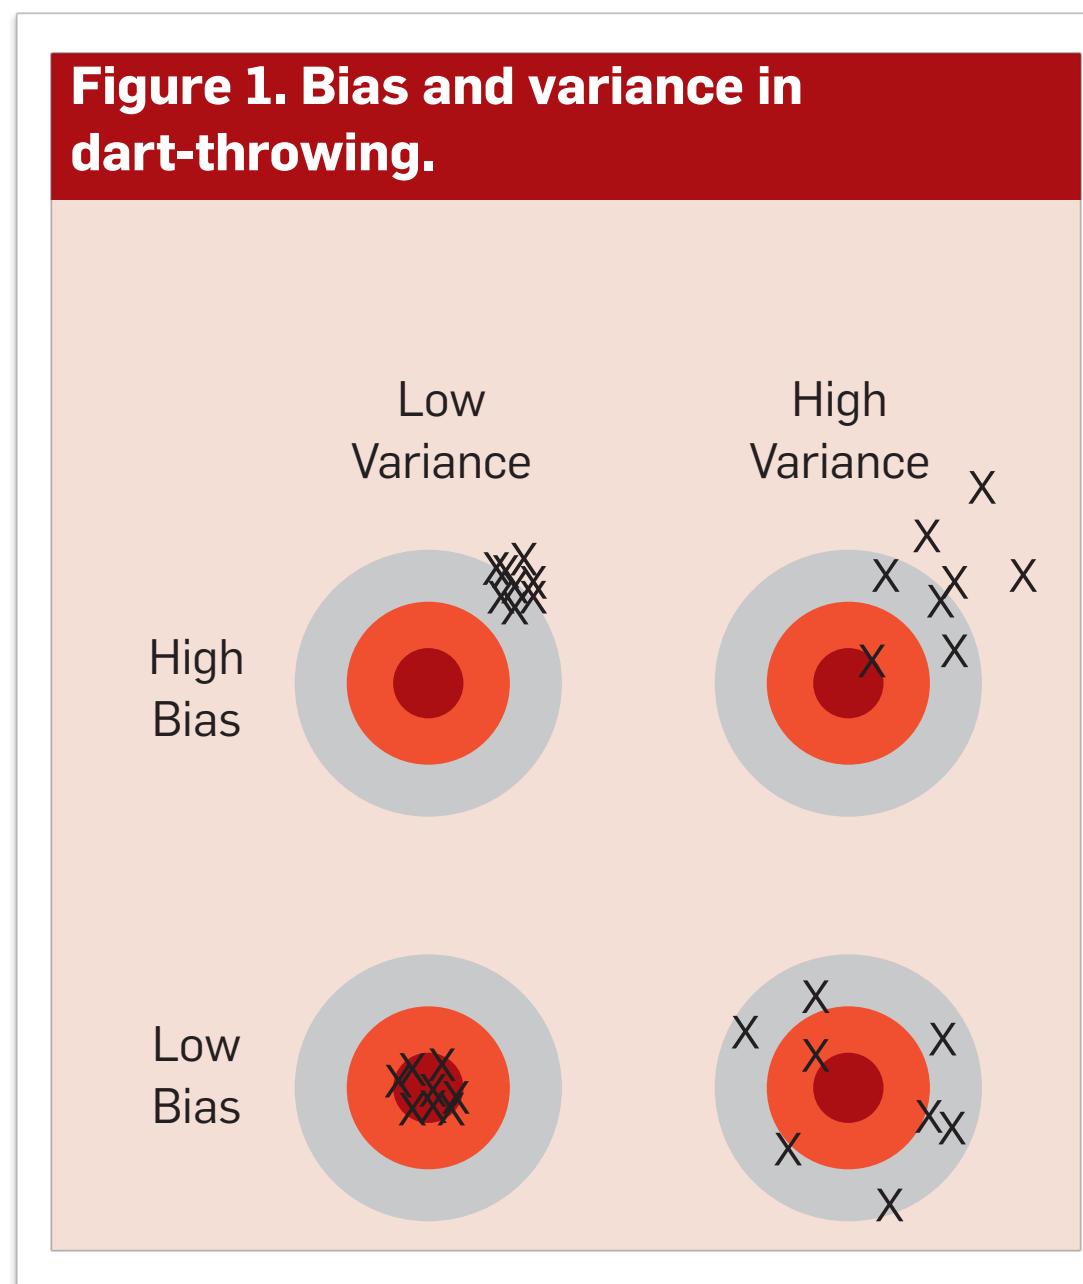
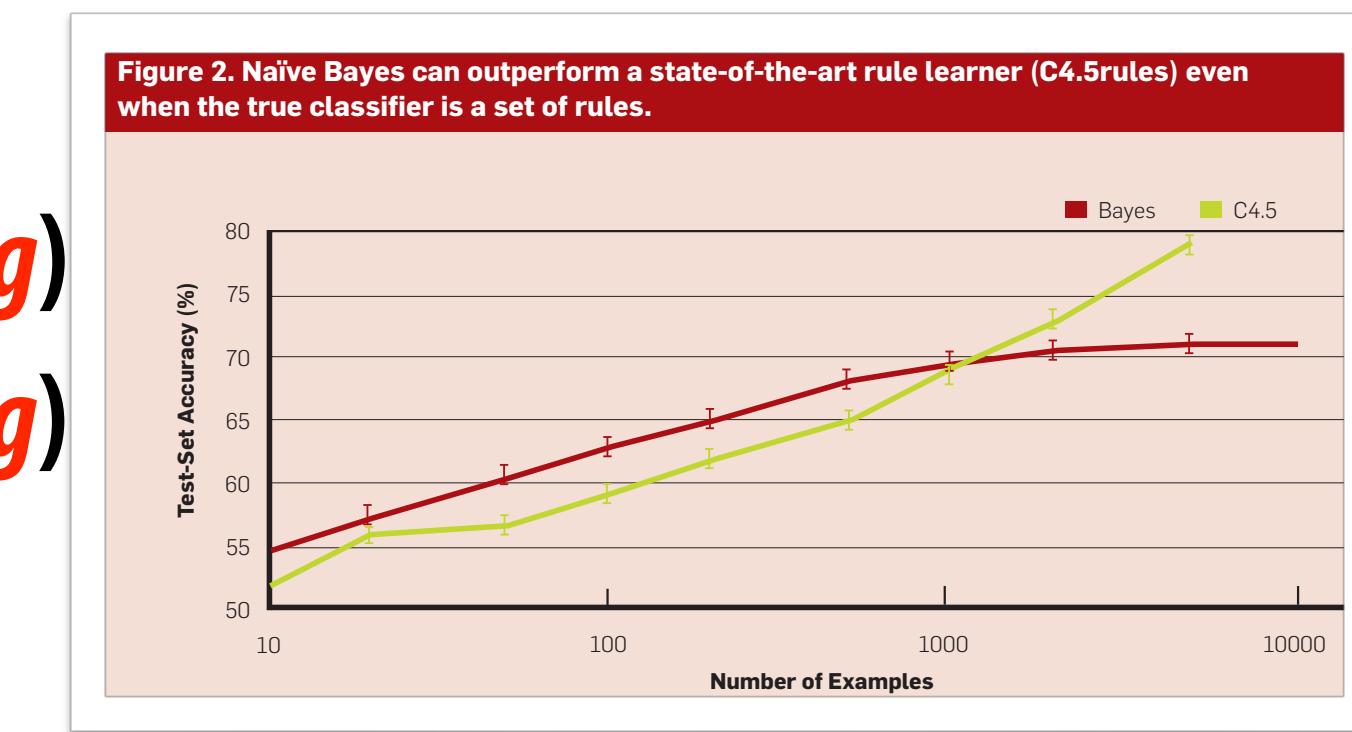
- **Distinguish** between *training data* and *test data*!
- To use as many data as possible use *cross-validation*
- **Principle #1** “not just good ML algorithms, but **more data!**”



Overfitting Has Many Faces

40

- Bias: 잘못된 것을 계속 학습하는 경향
- Variance: 실제 데이터와 상관없이 아무거나 배워버리는 경향
- The Bias-Variance Decomposition
 - Bias – e.g., linear predictors problem (*underfitting*)
 - Variance – e.g., decision trees problem (*overfitting*)
- Cross-validation, regularization



Intuition Fails in High Dimensions

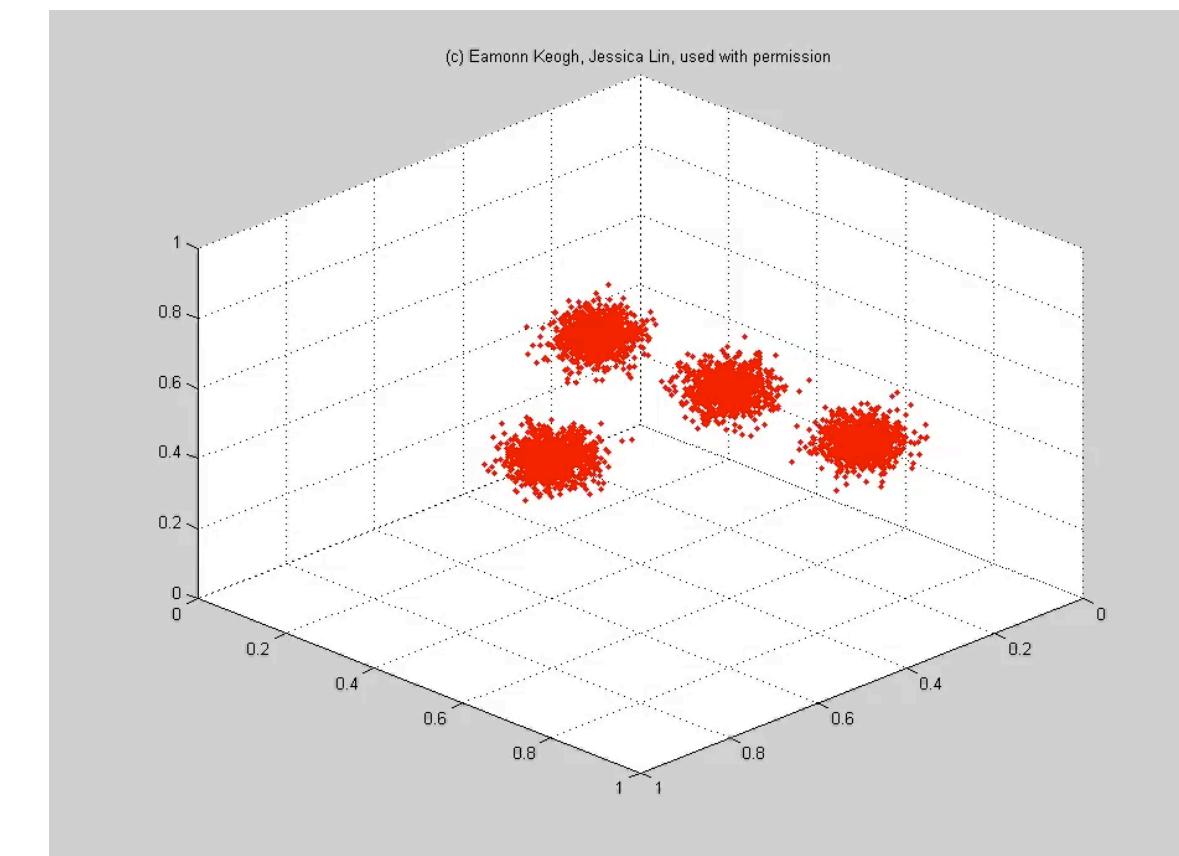
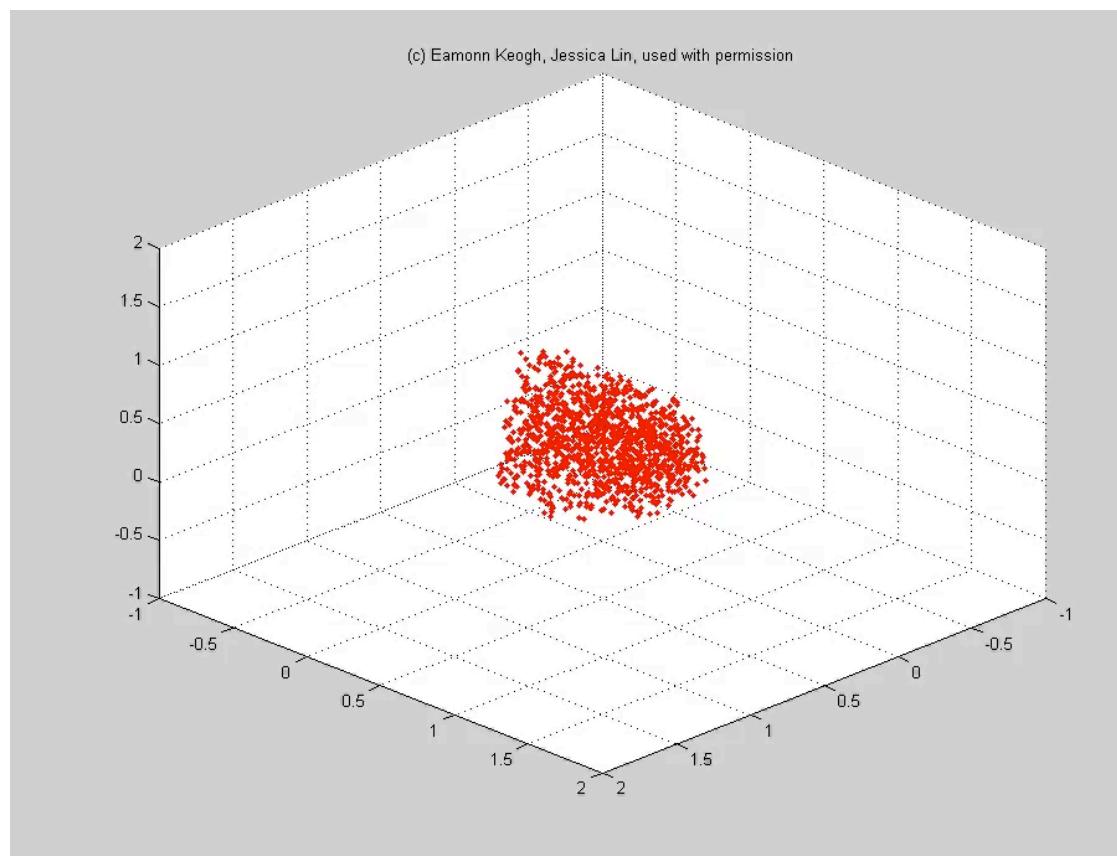
41

■ *Curse of dimensionality* [Bellman, 1957]

- Correct generalizing becomes **exponentially harder** as the dimensionality (number of features) of the examples grows
- With a fixed number of training instances, the predictive power **reduces** as the dimensionality increases; **Hughes phenomenon** (named after Gordon F. Hughes)
- Add new attributes only if they bring **new information** (do not use highly correlating features)

■ *Blessing of non-uniformity*

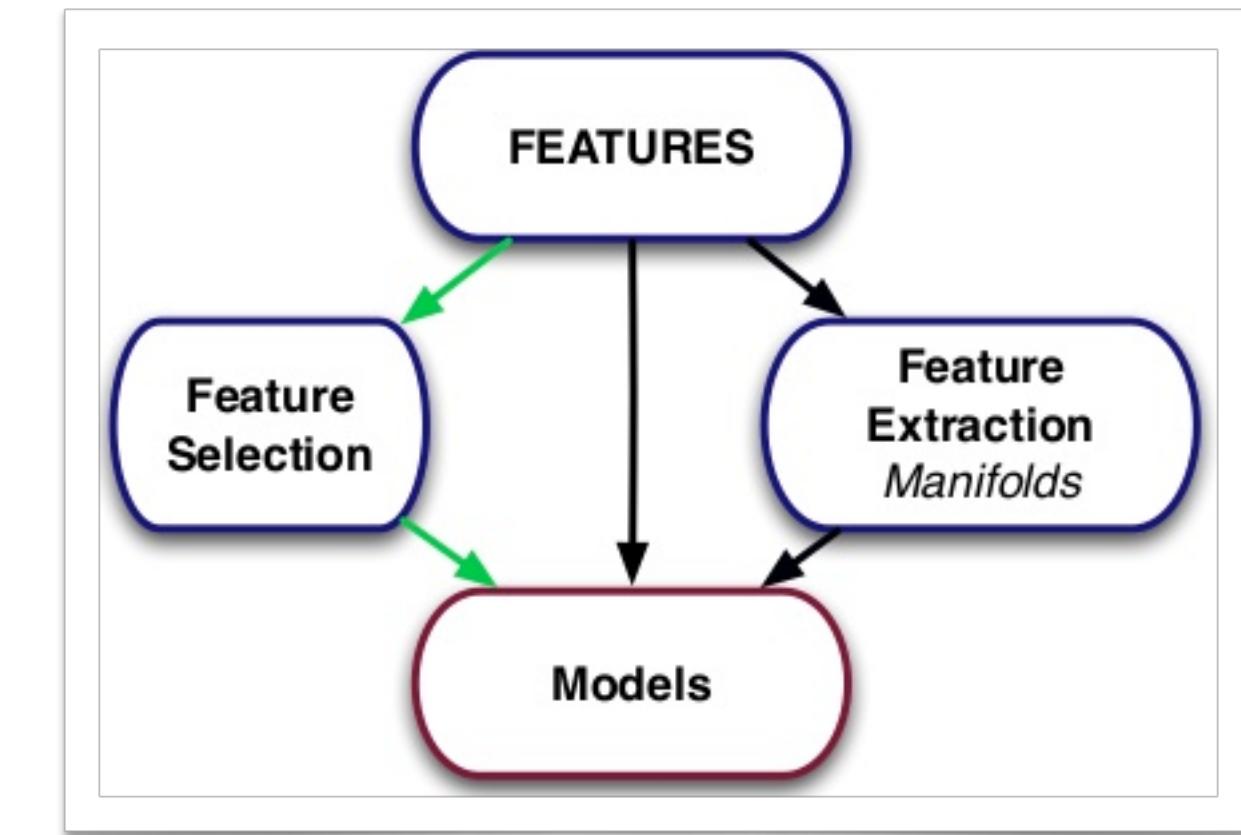
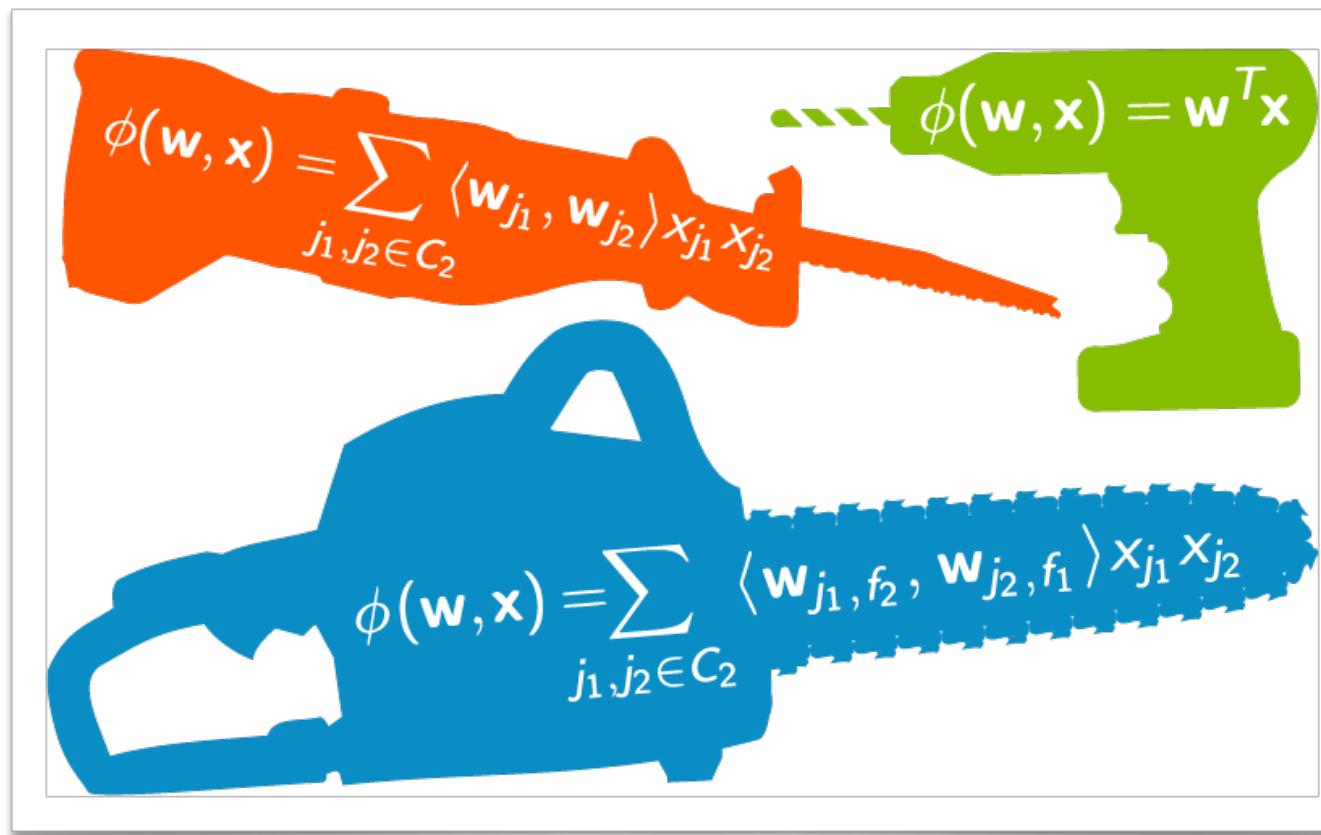
- Non-uniform training examples [e.g., handwritten digit recognition]
- Dimensionality reduction techniques [***Note:** the **dangers** of dimensionality reduction]



Feature Engineering is the Key

42

- The most important factor of ML project is the **features** used.
 - With many independent features that each **correlate well** with the predicted class, **learning is easy**.
 - On the other hand, if the class is a **very complex function** of the features, you may **not** be able to learn it.
- ML is often the quickest part, but that's because we've already mastered it pretty well!
 - Feature engineering is more difficult because it's **domain-specific**, while learners can be largely **general-purpose**.

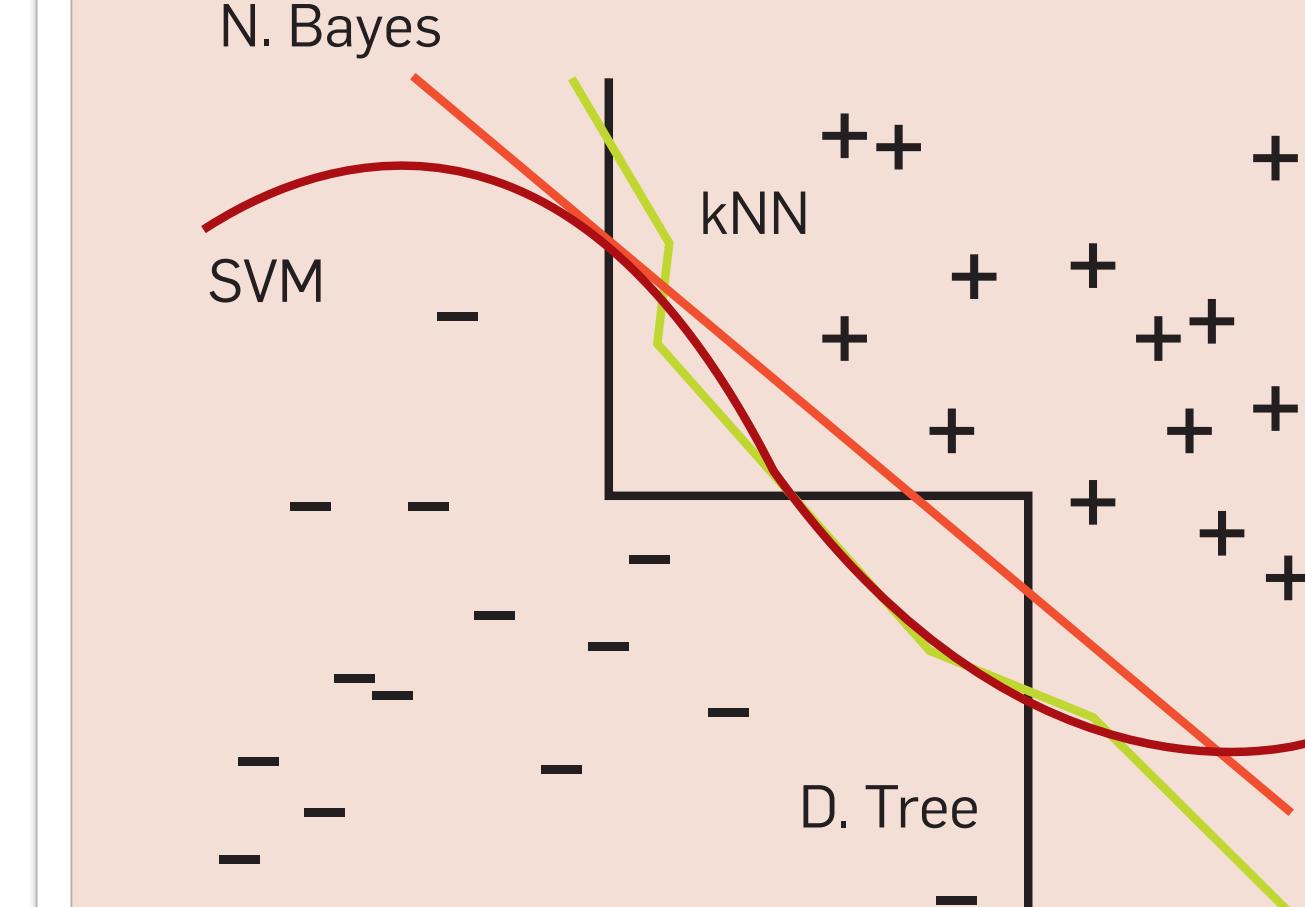


More Data Beats a Cleverer Algorithm

43

- ML researchers are mainly concerned with creation of **better learning algorithm**, but pragmatically the quickest path to success is often to just **get more data**.
- This does bring up another problem, however: **scalability**.
 - In most of computer science, the 2 main limited resources are **time and memory**.
In ML, there is a 3rd one: **training data**.
 - Bottleneck changed **from data to time**

Figure 3. Very different frontiers can yield similar predictions. (+ and – are training examples of two classes.)



Learn Many Models, Not Just One

44

■ Ensemble Methods

- ***Bagging*** - the simplest technique, which simply generate random variations of the training set by ***resampling***, learn a classifier on each, and combine the results by voting. This works because it greatly ***reduces variance*** while only slightly increasing bias.
- ***Boosting*** - training examples have ***weights***, and these are varied so that each new classifier focuses on the examples the previous ones tended to get wrong.
- ***Stacking*** - the outputs of individual classifiers become the inputs of a “***higher-level learner***” that figures out how best to combine them.

Simplicity Does Not Imply Accuracy

45

■ Occam's razor (오кам의 러컨도날)

- Principle of Economy (경제성의 원리), Principle of Parsimony (사고 절약의 원리)
- “Given two equally accurate theories, choose the one that is **less complex**”
- The explanation of any phenomenon should make as **few assumptions** as possible



■ No free lunch theorem

- Paper: “**No Free Lunch Theorems for Optimization**”, *IEEE Transactions on Evolutionary Computation*, 1997.
- 특정한 문제에 최적화된 알고리즘은 다른 문제에 대해서는 그동지 암�다는 것을 수학적으로 증명한 정리
- ML을 적용하는데 있어 종종 범하는 실수 중 하나는 하나의 ML 프로그램에서 많은 것을 바란다는 것
- “**리서치러닝은 마법이 아니다!**”

Correlation Does Not Imply Causation

46

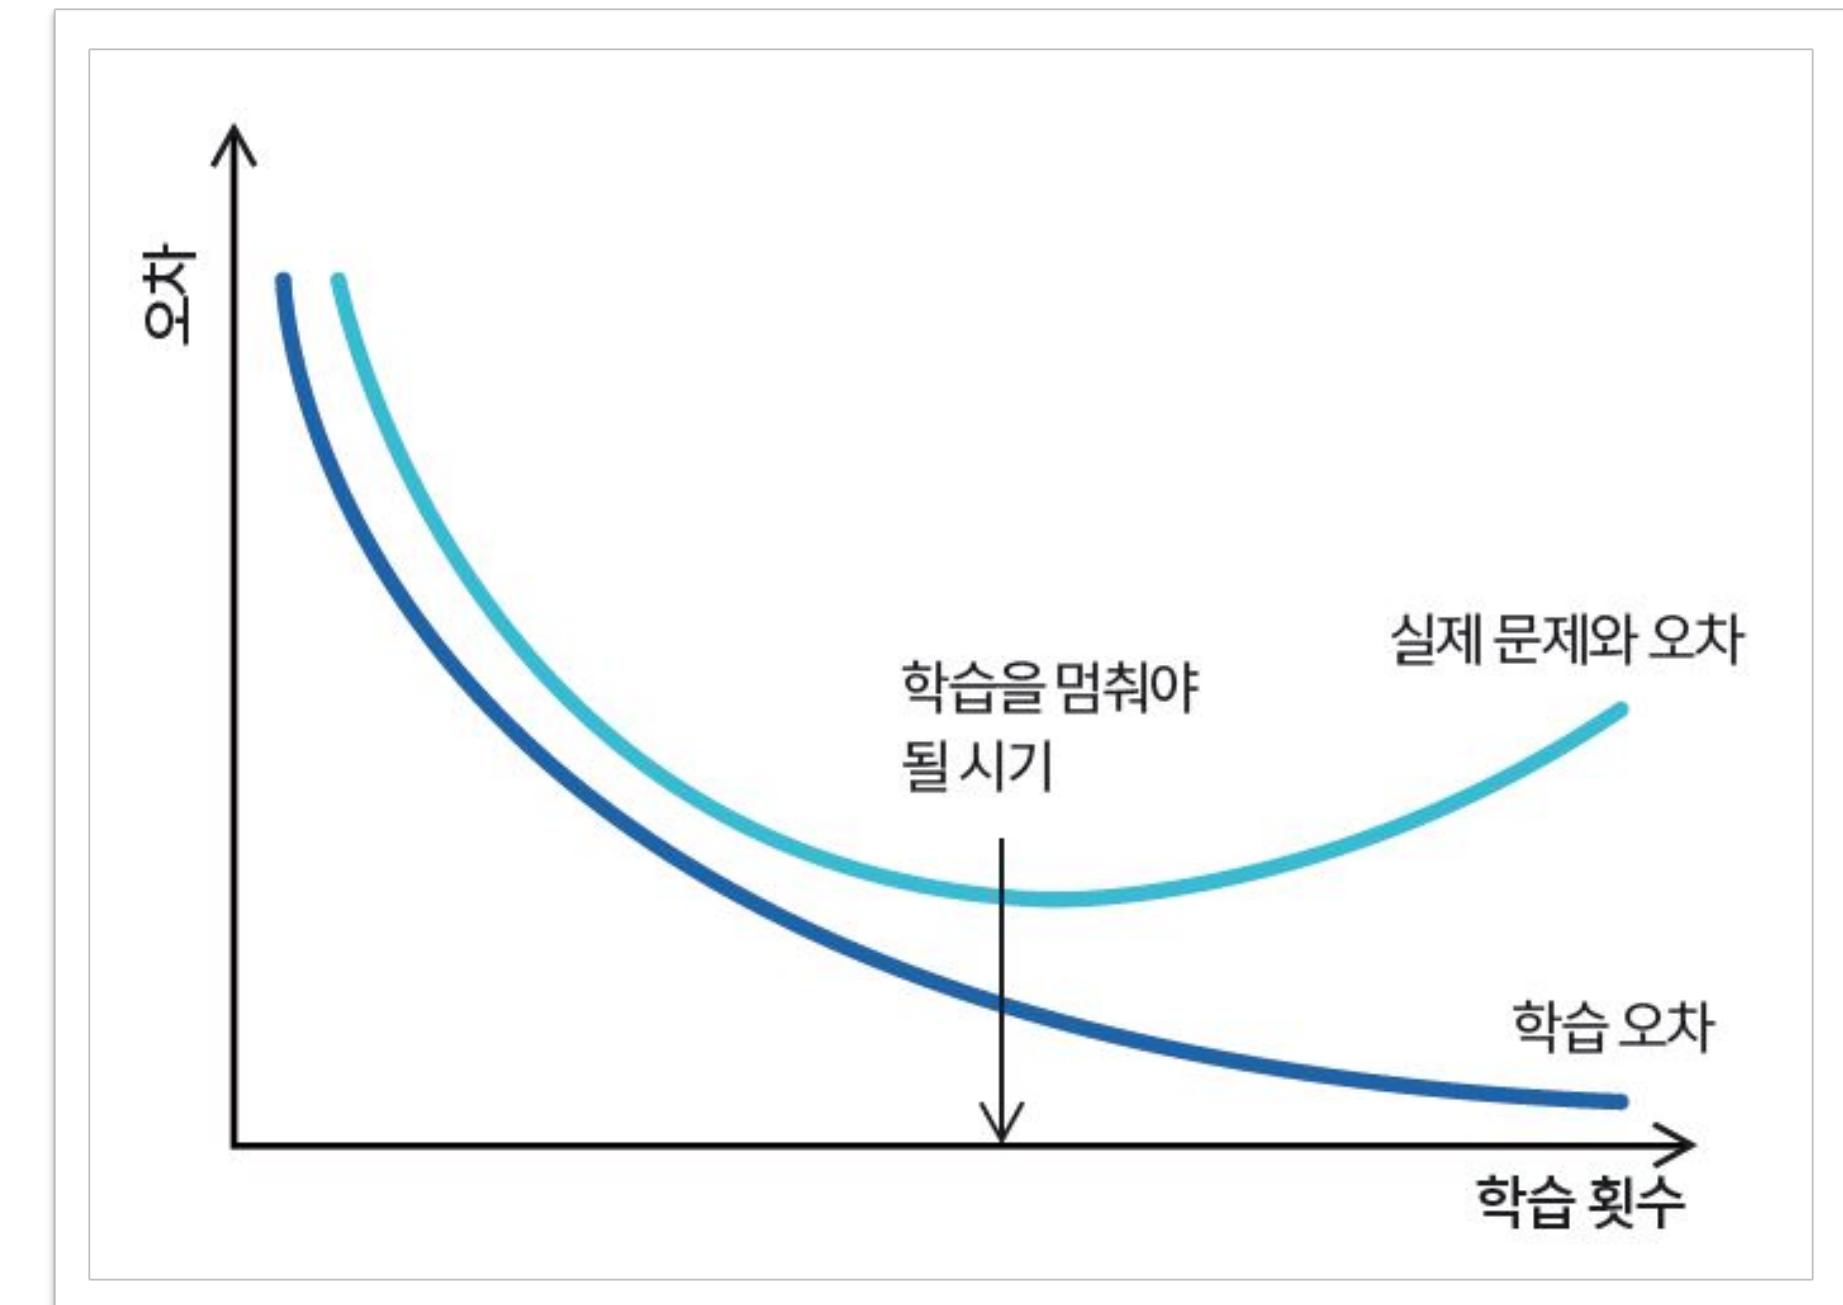
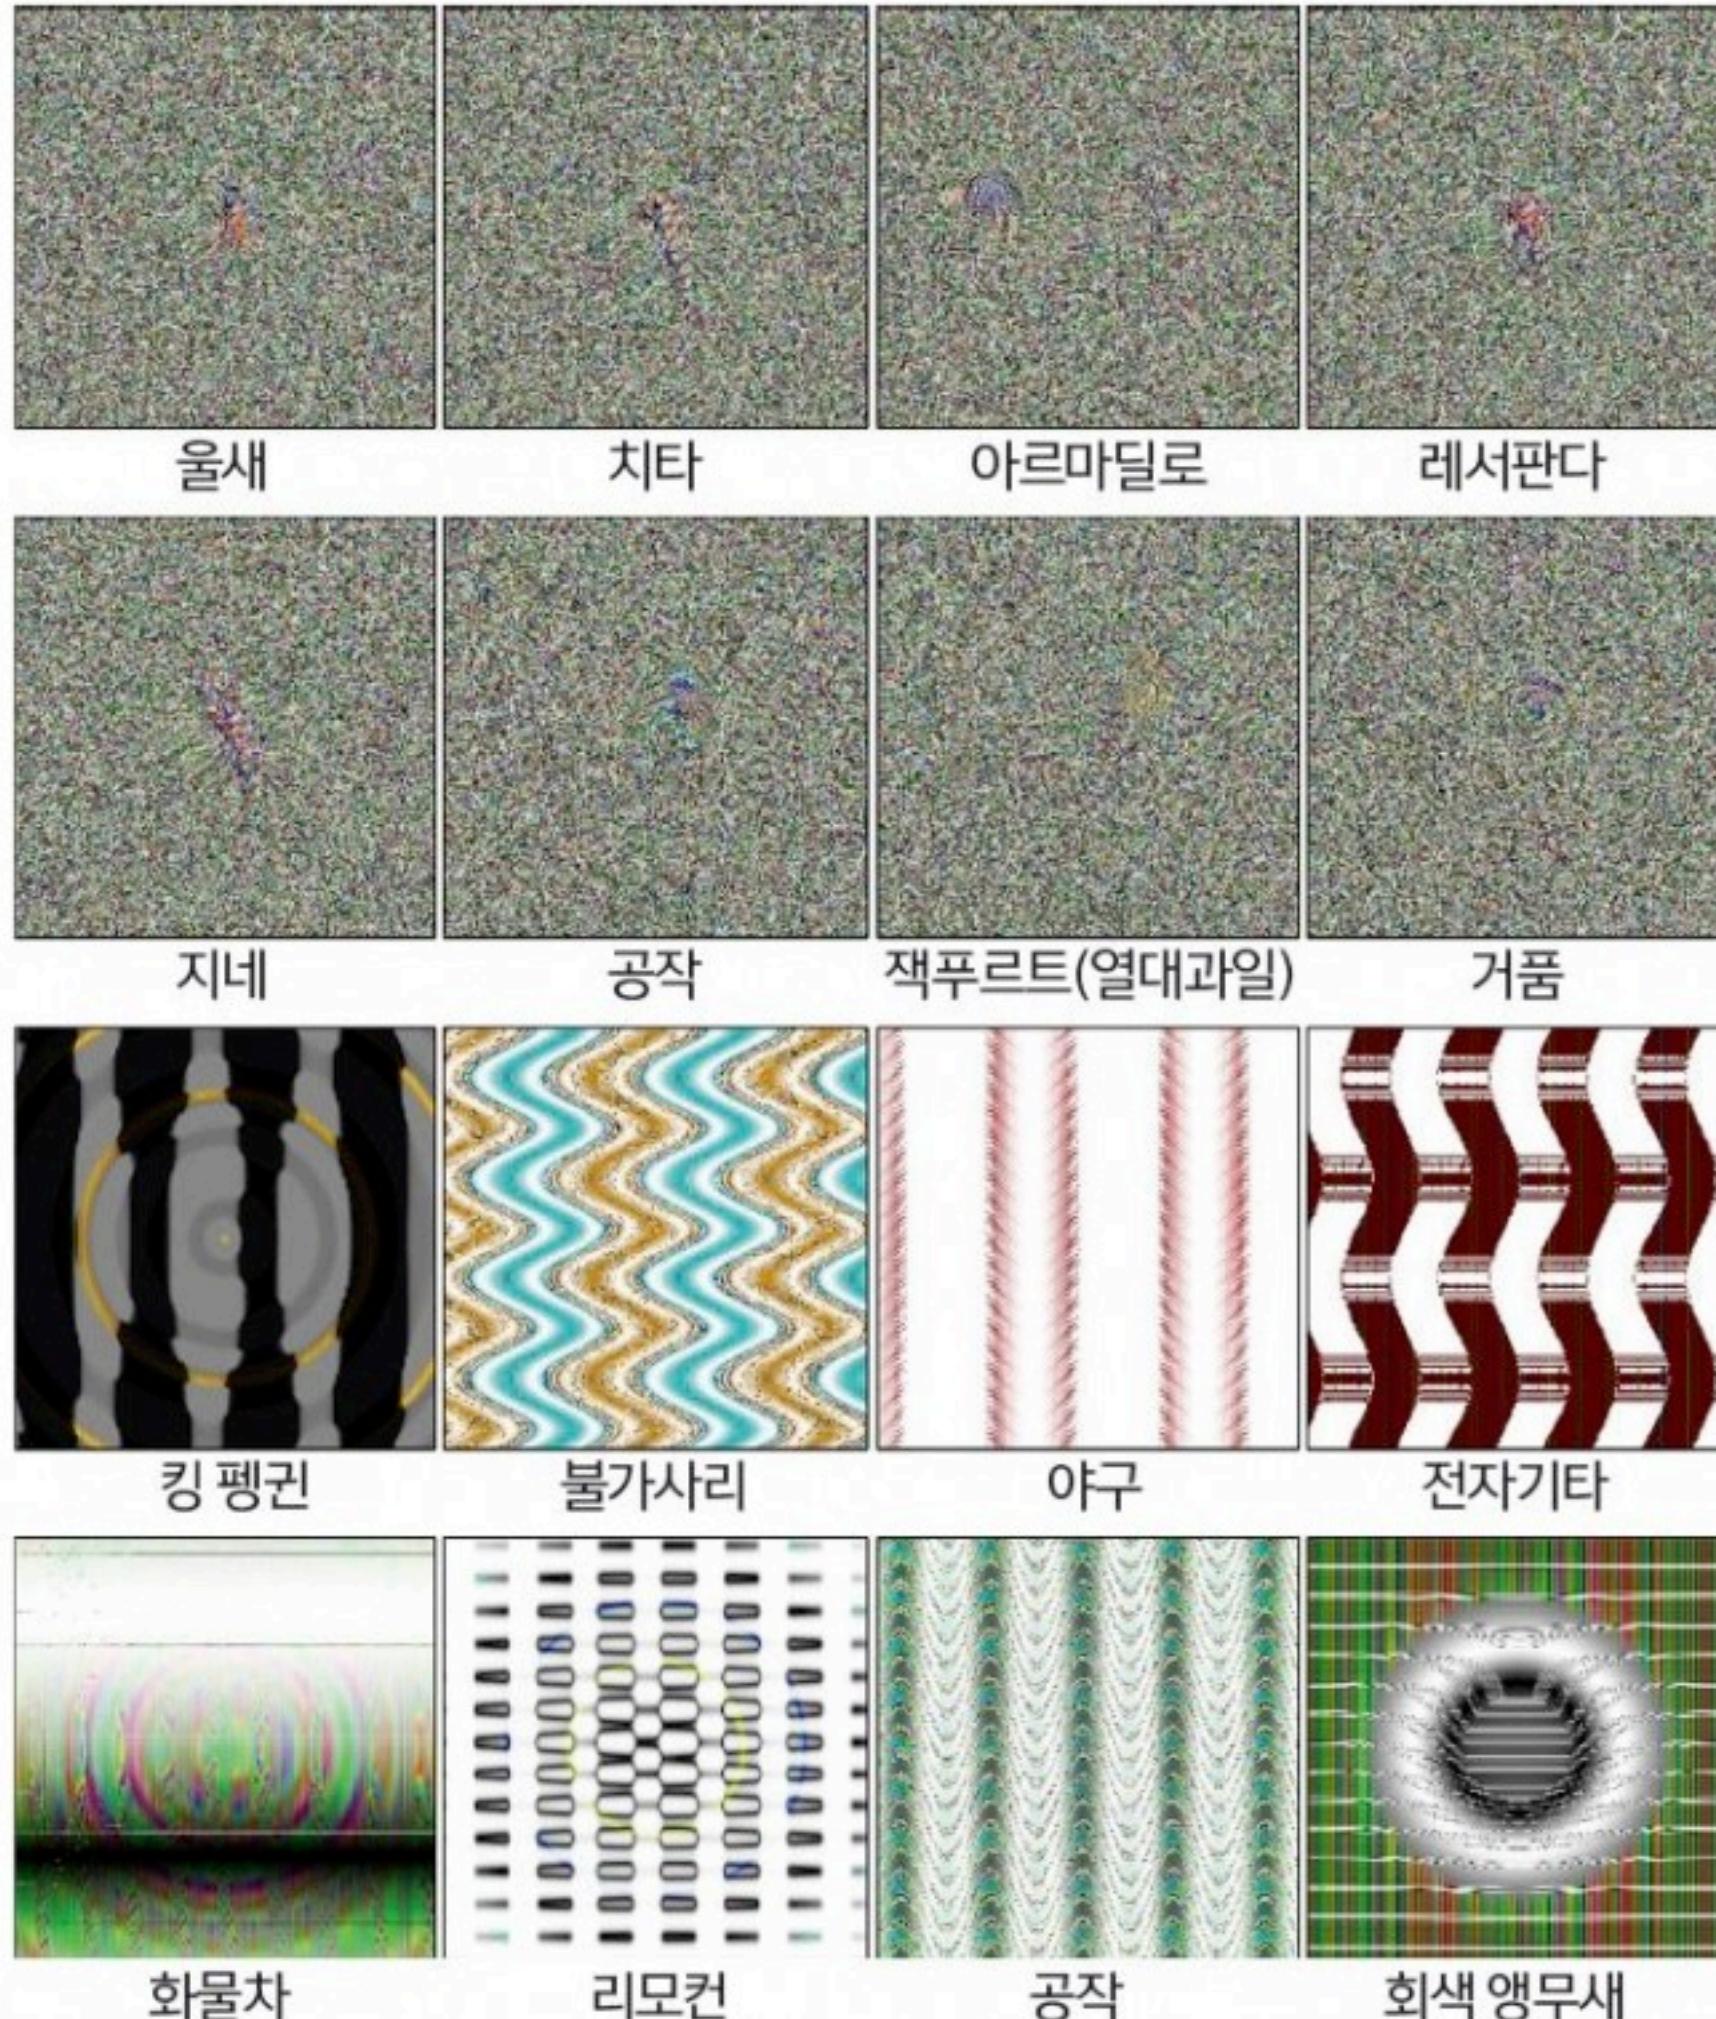
- 단지 동시에 두 가지 결과가 함께 나왔다고 해서 인과관계가 있다고 가정하지 마라!
- ML 알고리즘에서의 training은 관측가능한 변수(observable data)에 의해 수행되기 때문에, 인과관계(causal relationship)를 생성해주는 예측변수(predictive variable)들이 드러나지 않는다!
- **Experimental data** [인과관계(causality)를 파악할 수 있는 데이터]
만약 이러한 데이터(experimental data)를 얻을 수 있다면, 최대한 많이 확보하라!

Major Drawbacks in Deep Learning

47

- There are **drawbacks** to keep in mind when using deep learning.

‘심층 신경망은 쉽게 바보가 될 수 있다’, CVPR 2015

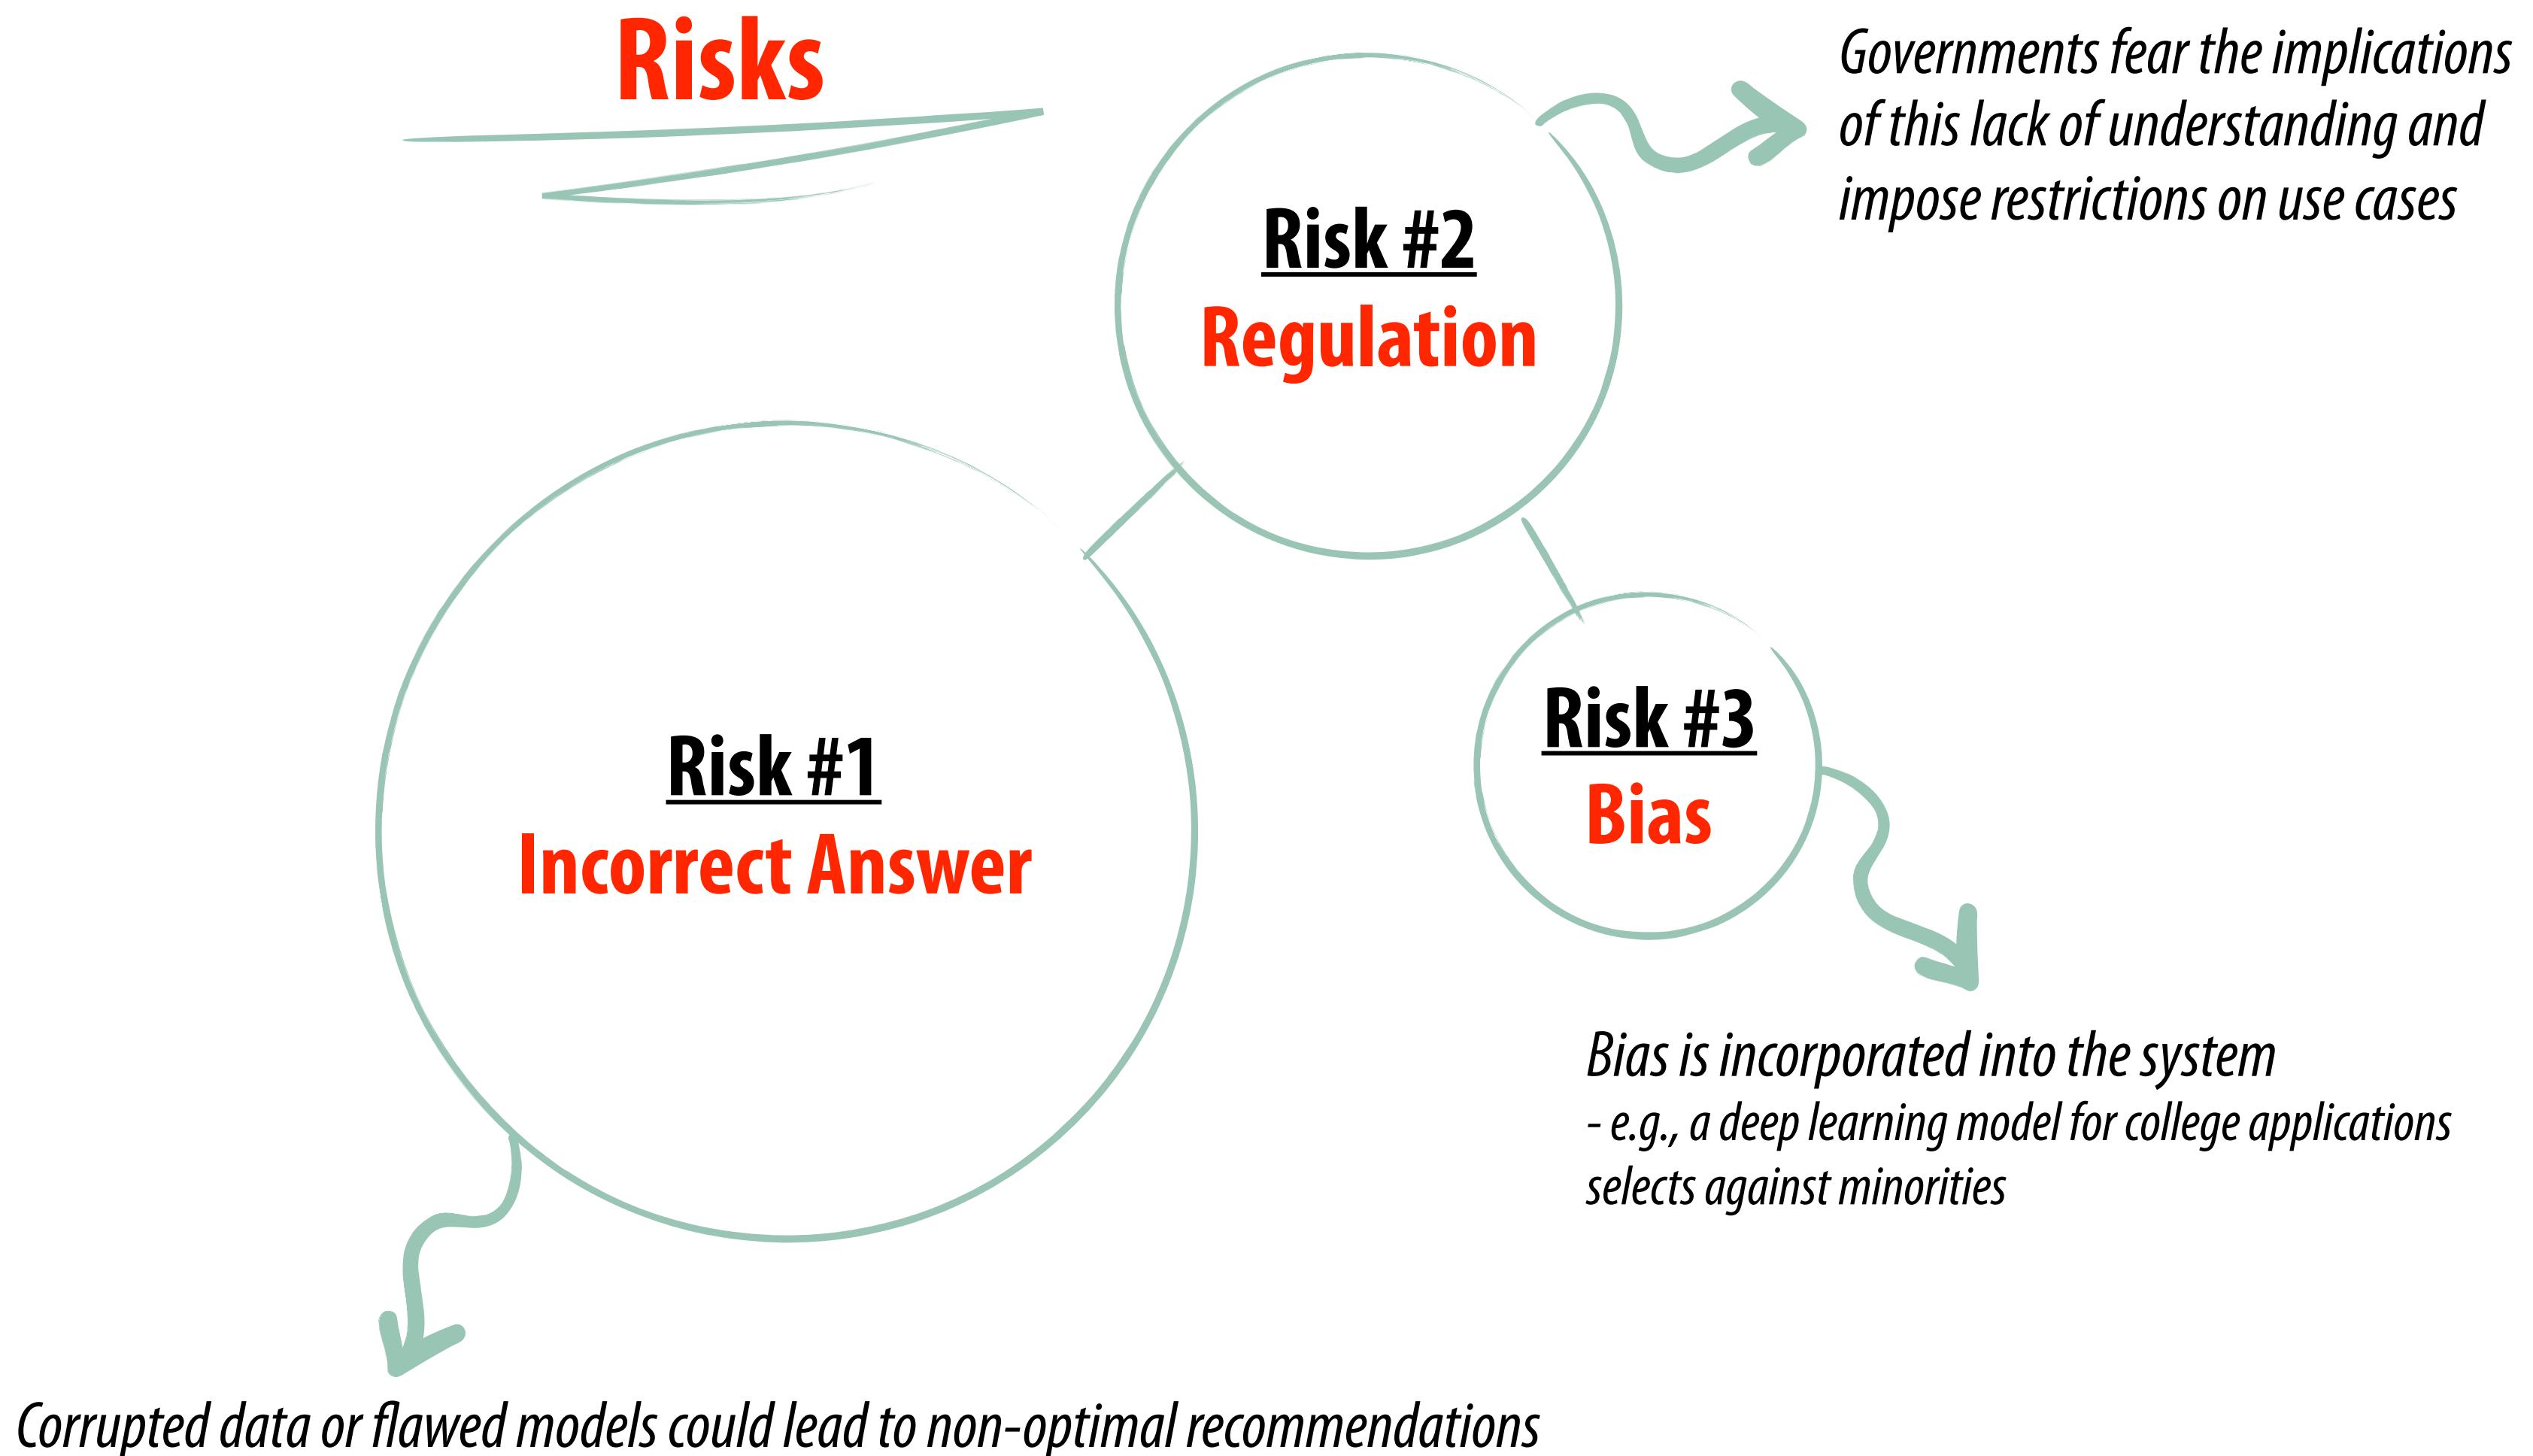


Overfitting!

Major Drawbacks in Deep Learning

48

- There are **drawbacks** to keep in mind when using deep learning.



Summary

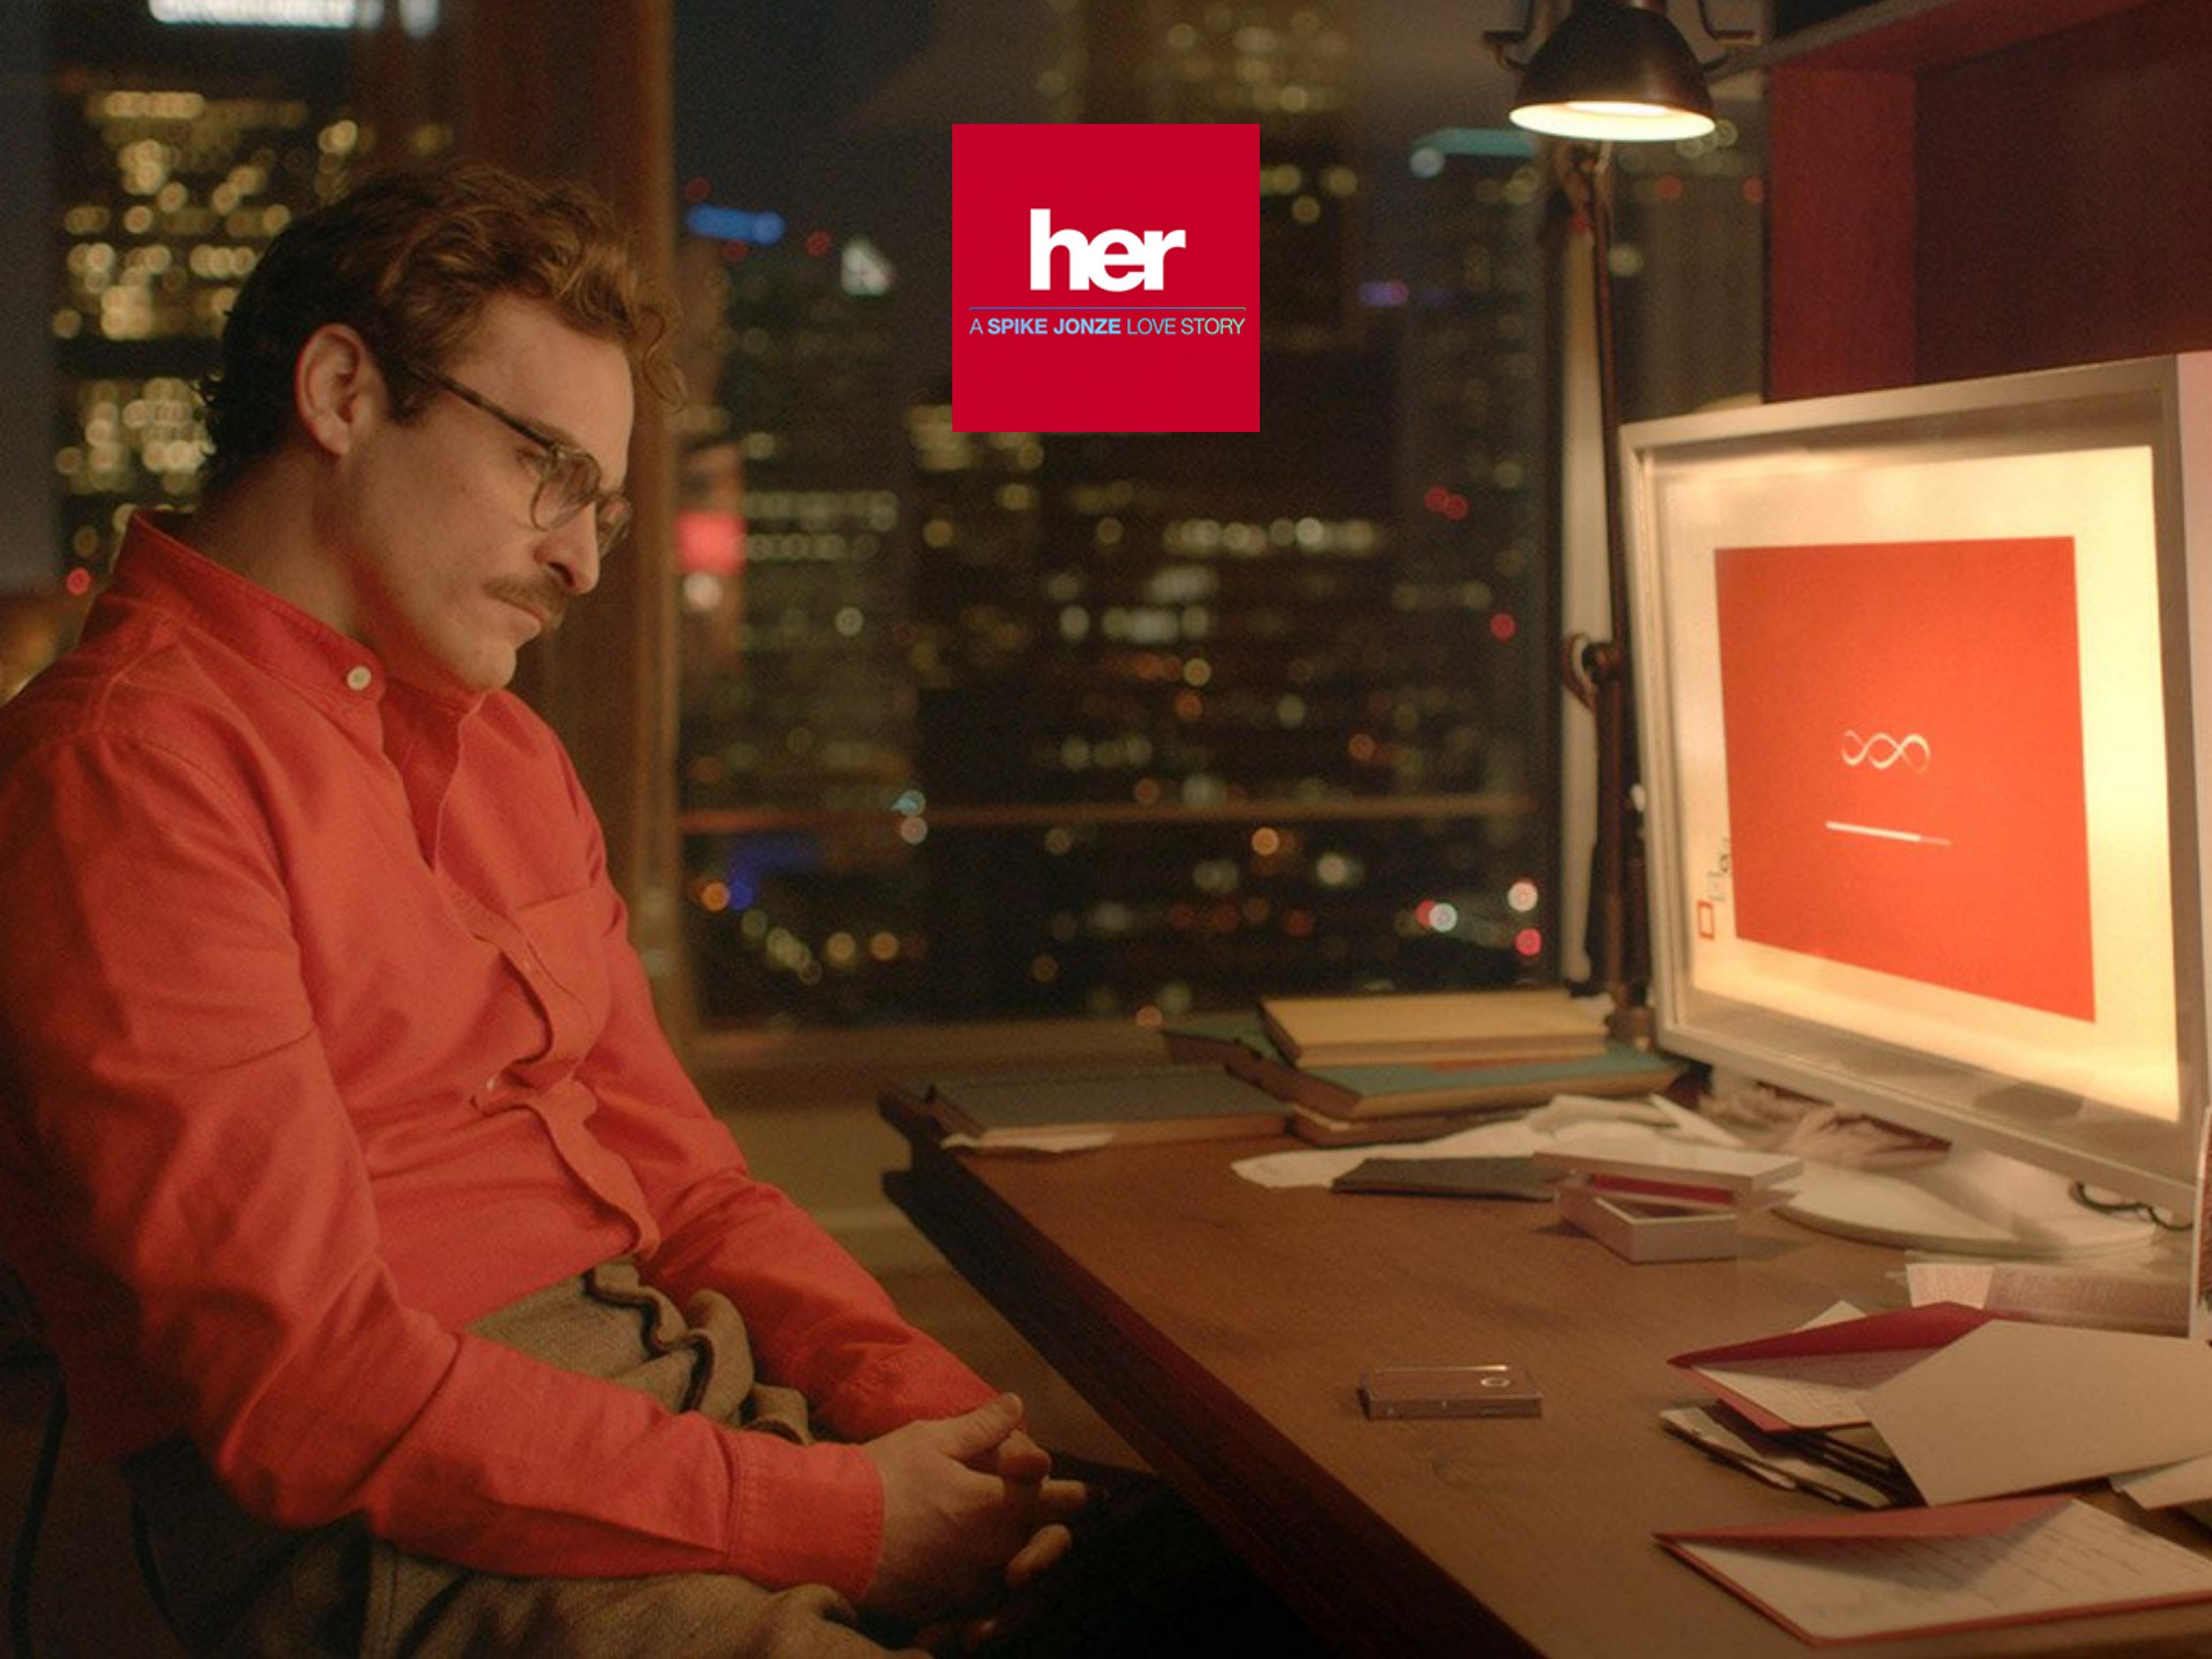
49

- **A Tour of ML Algorithms: Machine learning is *not monolithic*.**
 - Supervised, Semi-Supervised, Unsupervised, Reinforcement Learning
 - Machine learning algorithms
- ***“Folk knowledge”* in machine learning**
 - training vs. test data, cross-validation, more data, many models (ensemble), curse of dimensionality, occam’s razor, no-free lunch theorem
- Deep learning has emerged as a technique with *strong advantages*, but also has important *drawbacks* as well.
 - Incorrect answer, bias, overfitting, regularization



one more thing

or two...



her

A SPIKE JONZE LOVE STORY

Enjoy the failure!

52

관리번호		총괄 과제명		세부 과제명		1. 연구목표		<별첨>-지능정보기술 플래그십 프로젝트 2016(3)-기반 SW컴퓨팅-01 자율지능 디지털 도구 (총괄/1세부)		총괄 목표 : 사 수집하 자율지능 * 입력 예시: 정형/비정형데 ** 적응형 기계학습: 최초 특성, 습관, 어투, 용모 등 점진적이고 신속하게 학습하고 - 적응형 기계학습 기반 자율 - 사람과 다양한 방식으로 역할을 할 수 있는 디지털 동 - 다양한 입력 방법을 수용하 학습 방법 연구개발 - 자율지능 동반자 프레임워크 기반 1세부 최종 목표 : 스스로 상황을 디지털 동반자 프	
개발목표		핵심 기술/제품 성능지표		단위		달성목표		국내최고수준		세계최고수준 (보유국, 기업/기관명)	
1	동반자의 멀티모달 자율지능 정도 (동반자 튜링 테스트 ¹⁾)	%	33	N/A	N/A					N/A	
2	디지털 동반자 사용자 만족도 (리커트 척도 ²⁾)	점수	3.5 이상	N/A	N/A					N/A	
3	동반자 공개 API 제공 여부 (주요 공통 기능)	지원여부	지원	-	-					-	

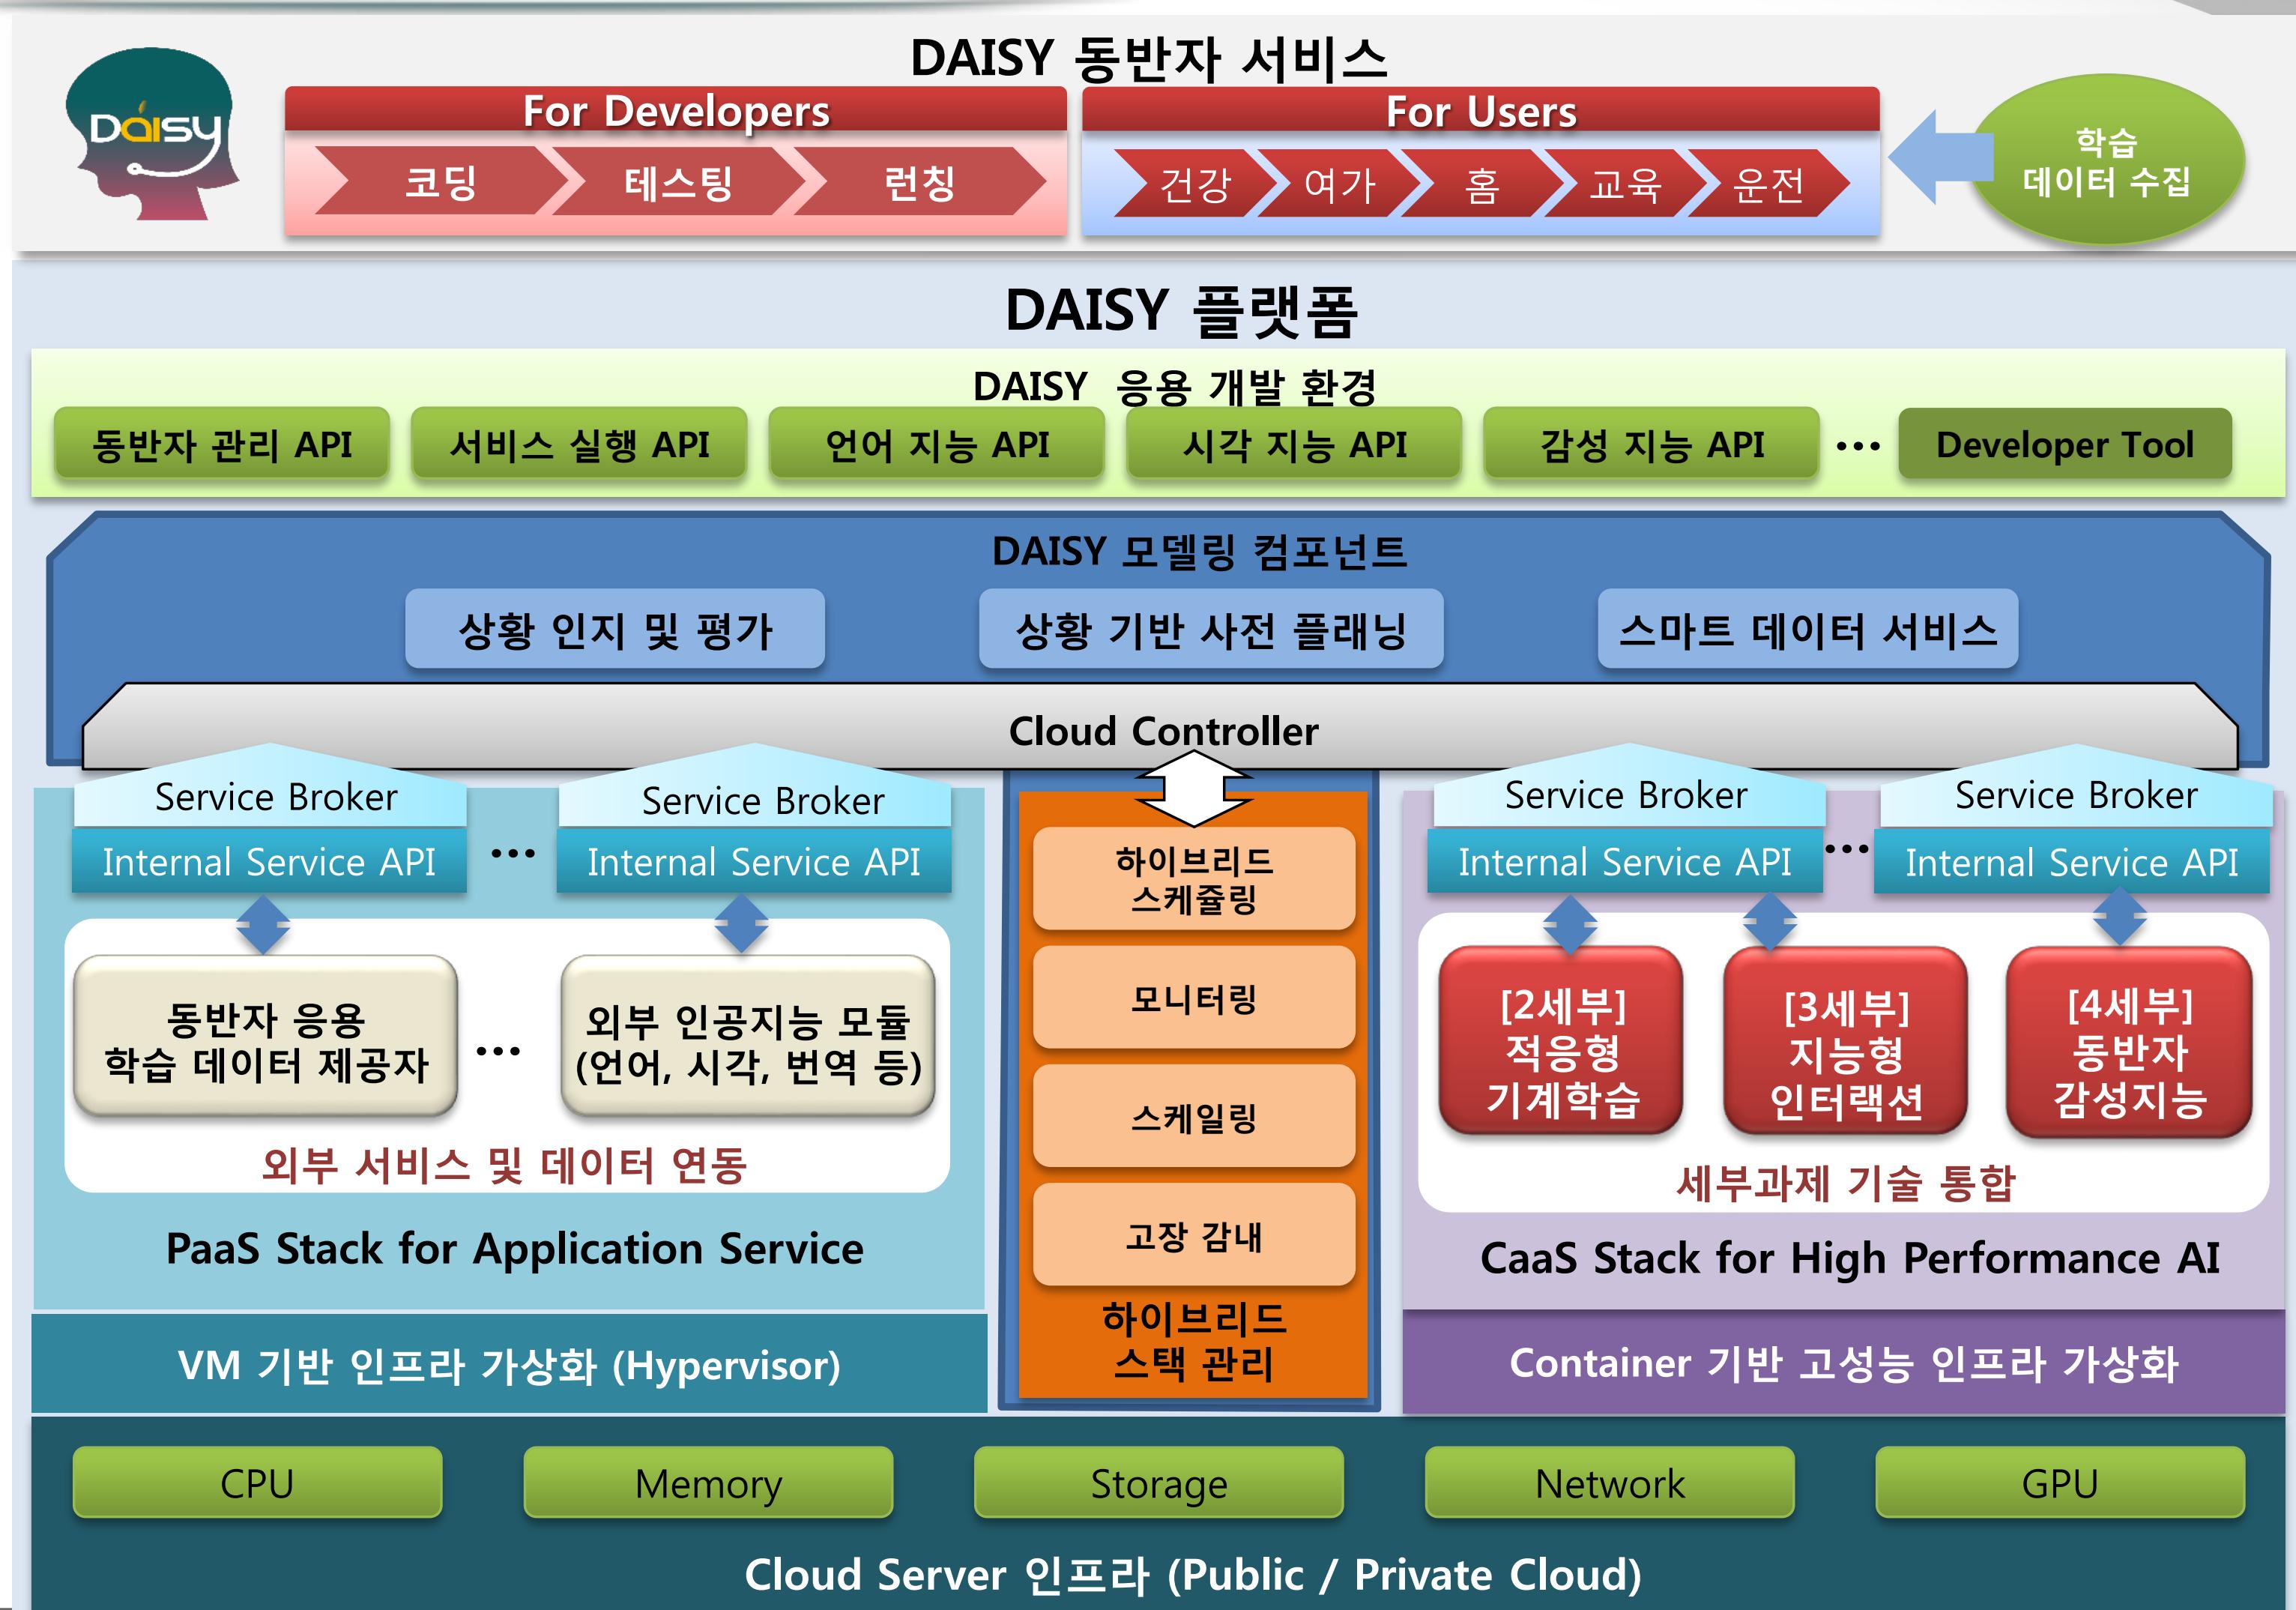
1) 튜링테스트 개념을 활용하여 30명 이상의 사용자를 대상으로 인간과 디지털 동반자를 구분하지 못하는 사람이 모수의 3분의 1 이상이면 통과
* 언어지능의 경우 세계 최고 수준: 일반대화 33%, 13세, 비원어민 기준(영국, 레딩대, 30명 모수 사용)
* 디지털 동반자 응용을 사용해본 사용자에게 기본 5단계^{*}로 설문조사하여 평균 집계**

Software Architecture

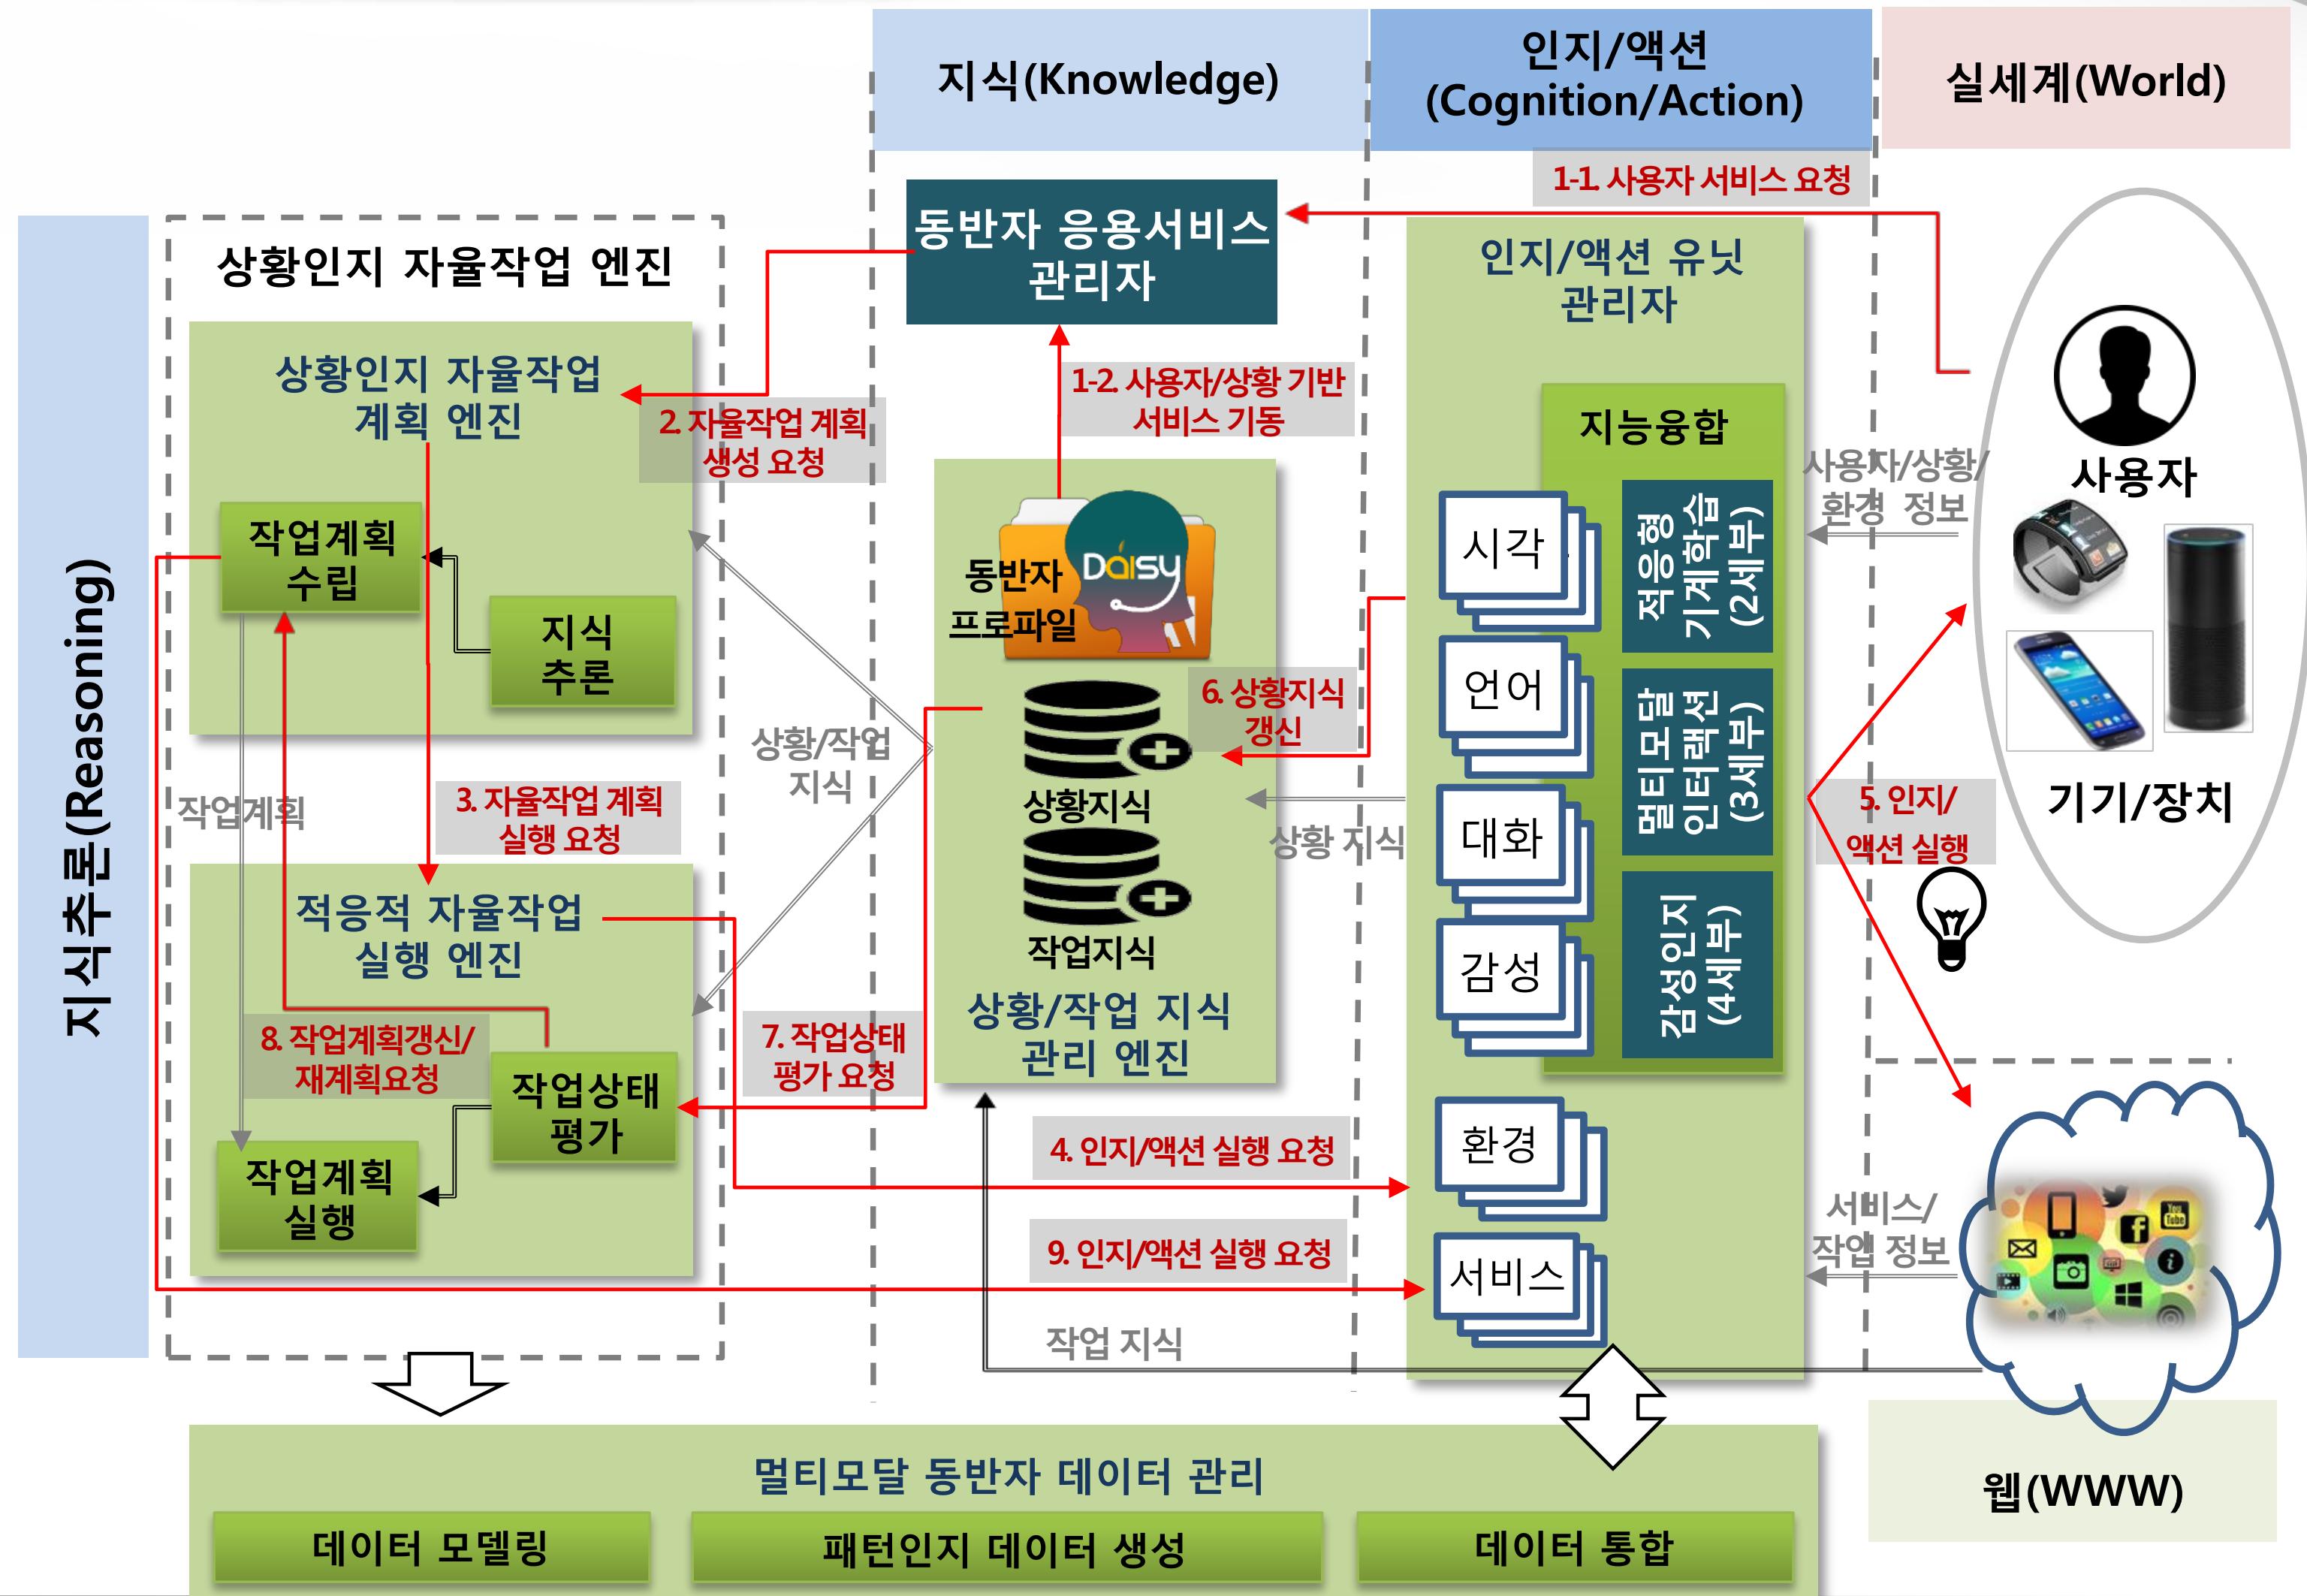
53

연구개발목표의 명확성

2.2 시스템 동작 메커니즘 - DAISY 플랫폼의 소프트웨어 구조



2.2 시스템 동작 메커니즘 - DAISY 인공지능 동작 메커니즘



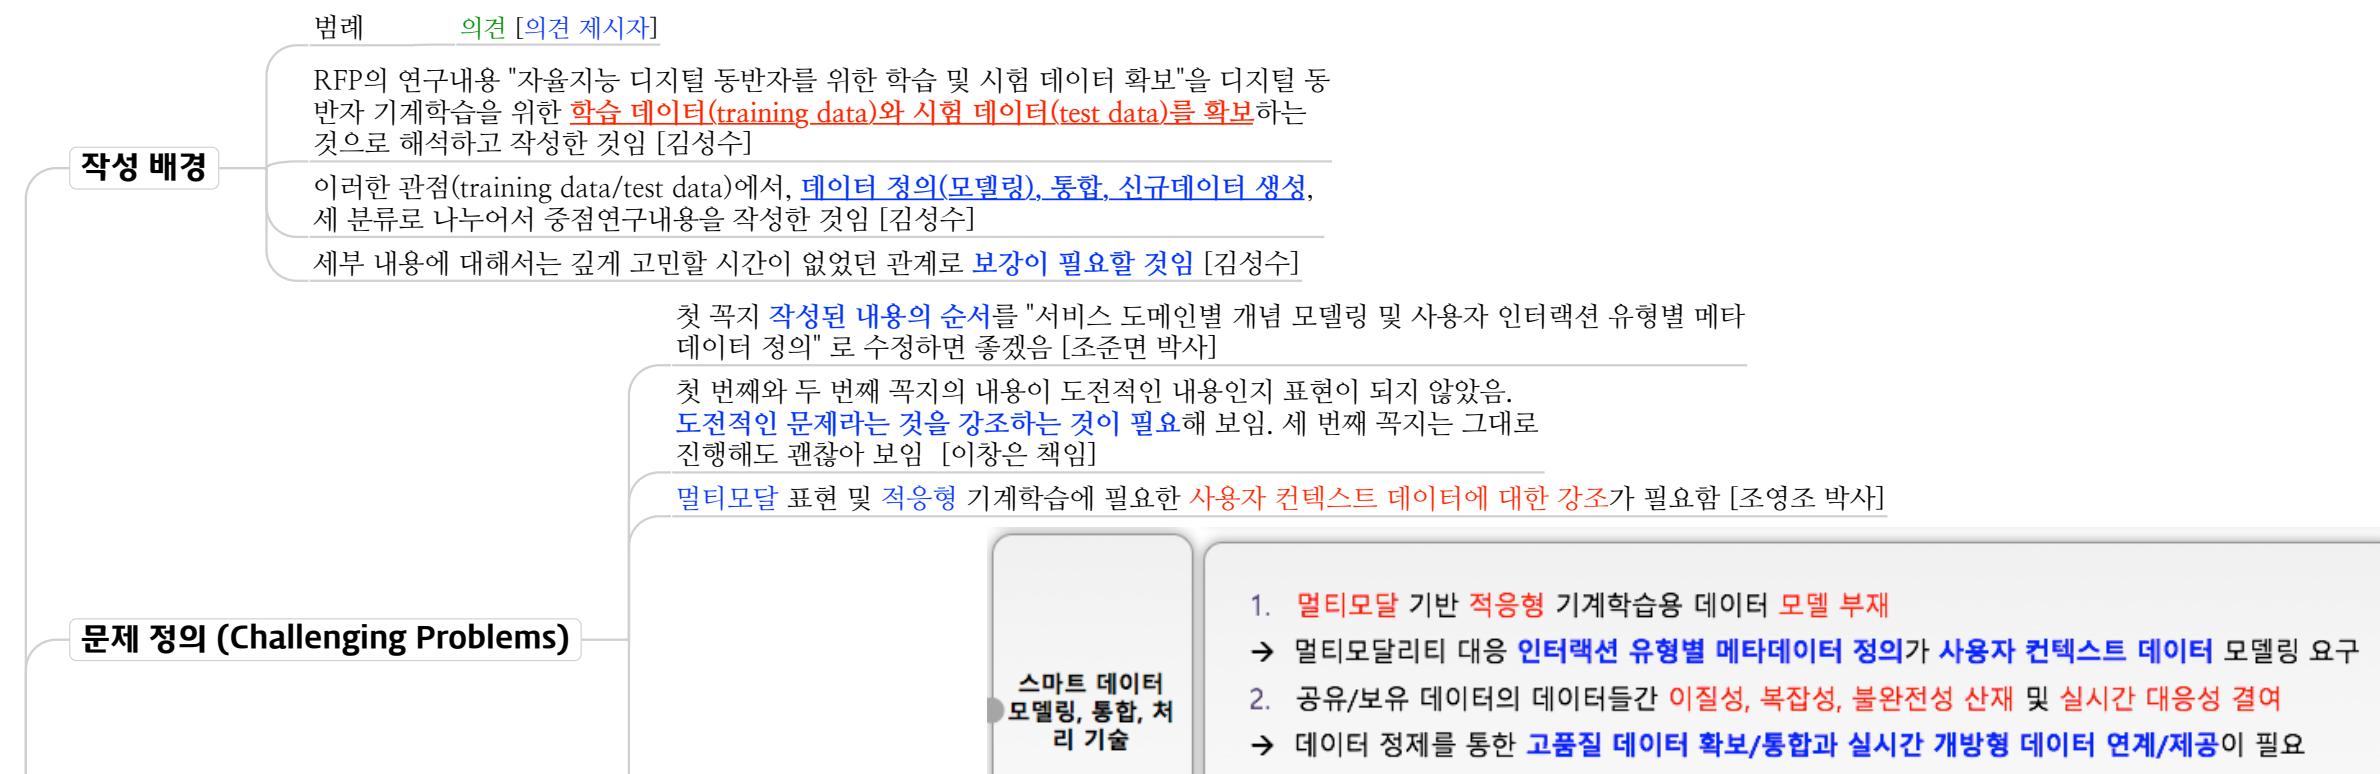
One of the TF Team Meeting Logs

55



One of the TF Team Meeting Logs

56



범례

의견 [의견 제시자]

작성 배경

RFP의 연구내용 "자율지능 디지털 동반자를 위한 학습 및 시험 데이터 확보"을 디지털 동반자 기계학습을 위한 학습 데이터(training data)와 시험 데이터(test data)를 확보하는 것으로 해석하고 작성한 것임 [김성수]

이러한 관점(training data/test data)에서, 데이터 정의(모델링), 통합, 신규데이터 생성, 세 분류로 나누어서 중점연구내용을 작성한 것임 [김성수]

세부 내용에 대해서는 깊게 고민할 시간이 없었던 관계로 보강이 필요할 것임 [김성수]

중점추진 연구내용 (창의성 슬라이드)

세부 내용에도 창의성을 강조할 수 있는 문구로 수정이 필요함. 예) 혁신 특허출원 예정 등 [이창은 책임]

예를 들어, 갑스 샘플링을 적용하는 이유
(차별성, 창의성) 등을 수식하는 문구가
필요함 [이창은 책임]

데이터 확보 방안

작성된 두 페이지를 한 페이지로 통합(첫 슬라이드 포맷으로)하고, 기업체 위주로 재구성 할 것 [이창은 책임, 조영조 박사]

SW 품질관리계획

"SPICE Level 3 인증 획득~~" 부분을 바로 아래 내용(sub item)을 요약한 제목으로 수정이 필요함 [이창은 책임]

"효과적인 연구개발 협업을 위한 시스템 구축 및 운용"으로 수정 [김성수]

과제관리계획

엑소브레인, 딥러닝 과제등과 연계 표현이된 HPC 포맷을 활용하여 정리 [이창은 책임]

해당 파일은 이메일로 보내 주겠음 [김원영 박사] 매일 수신하는 대로 정리 하겠음 [김성수]

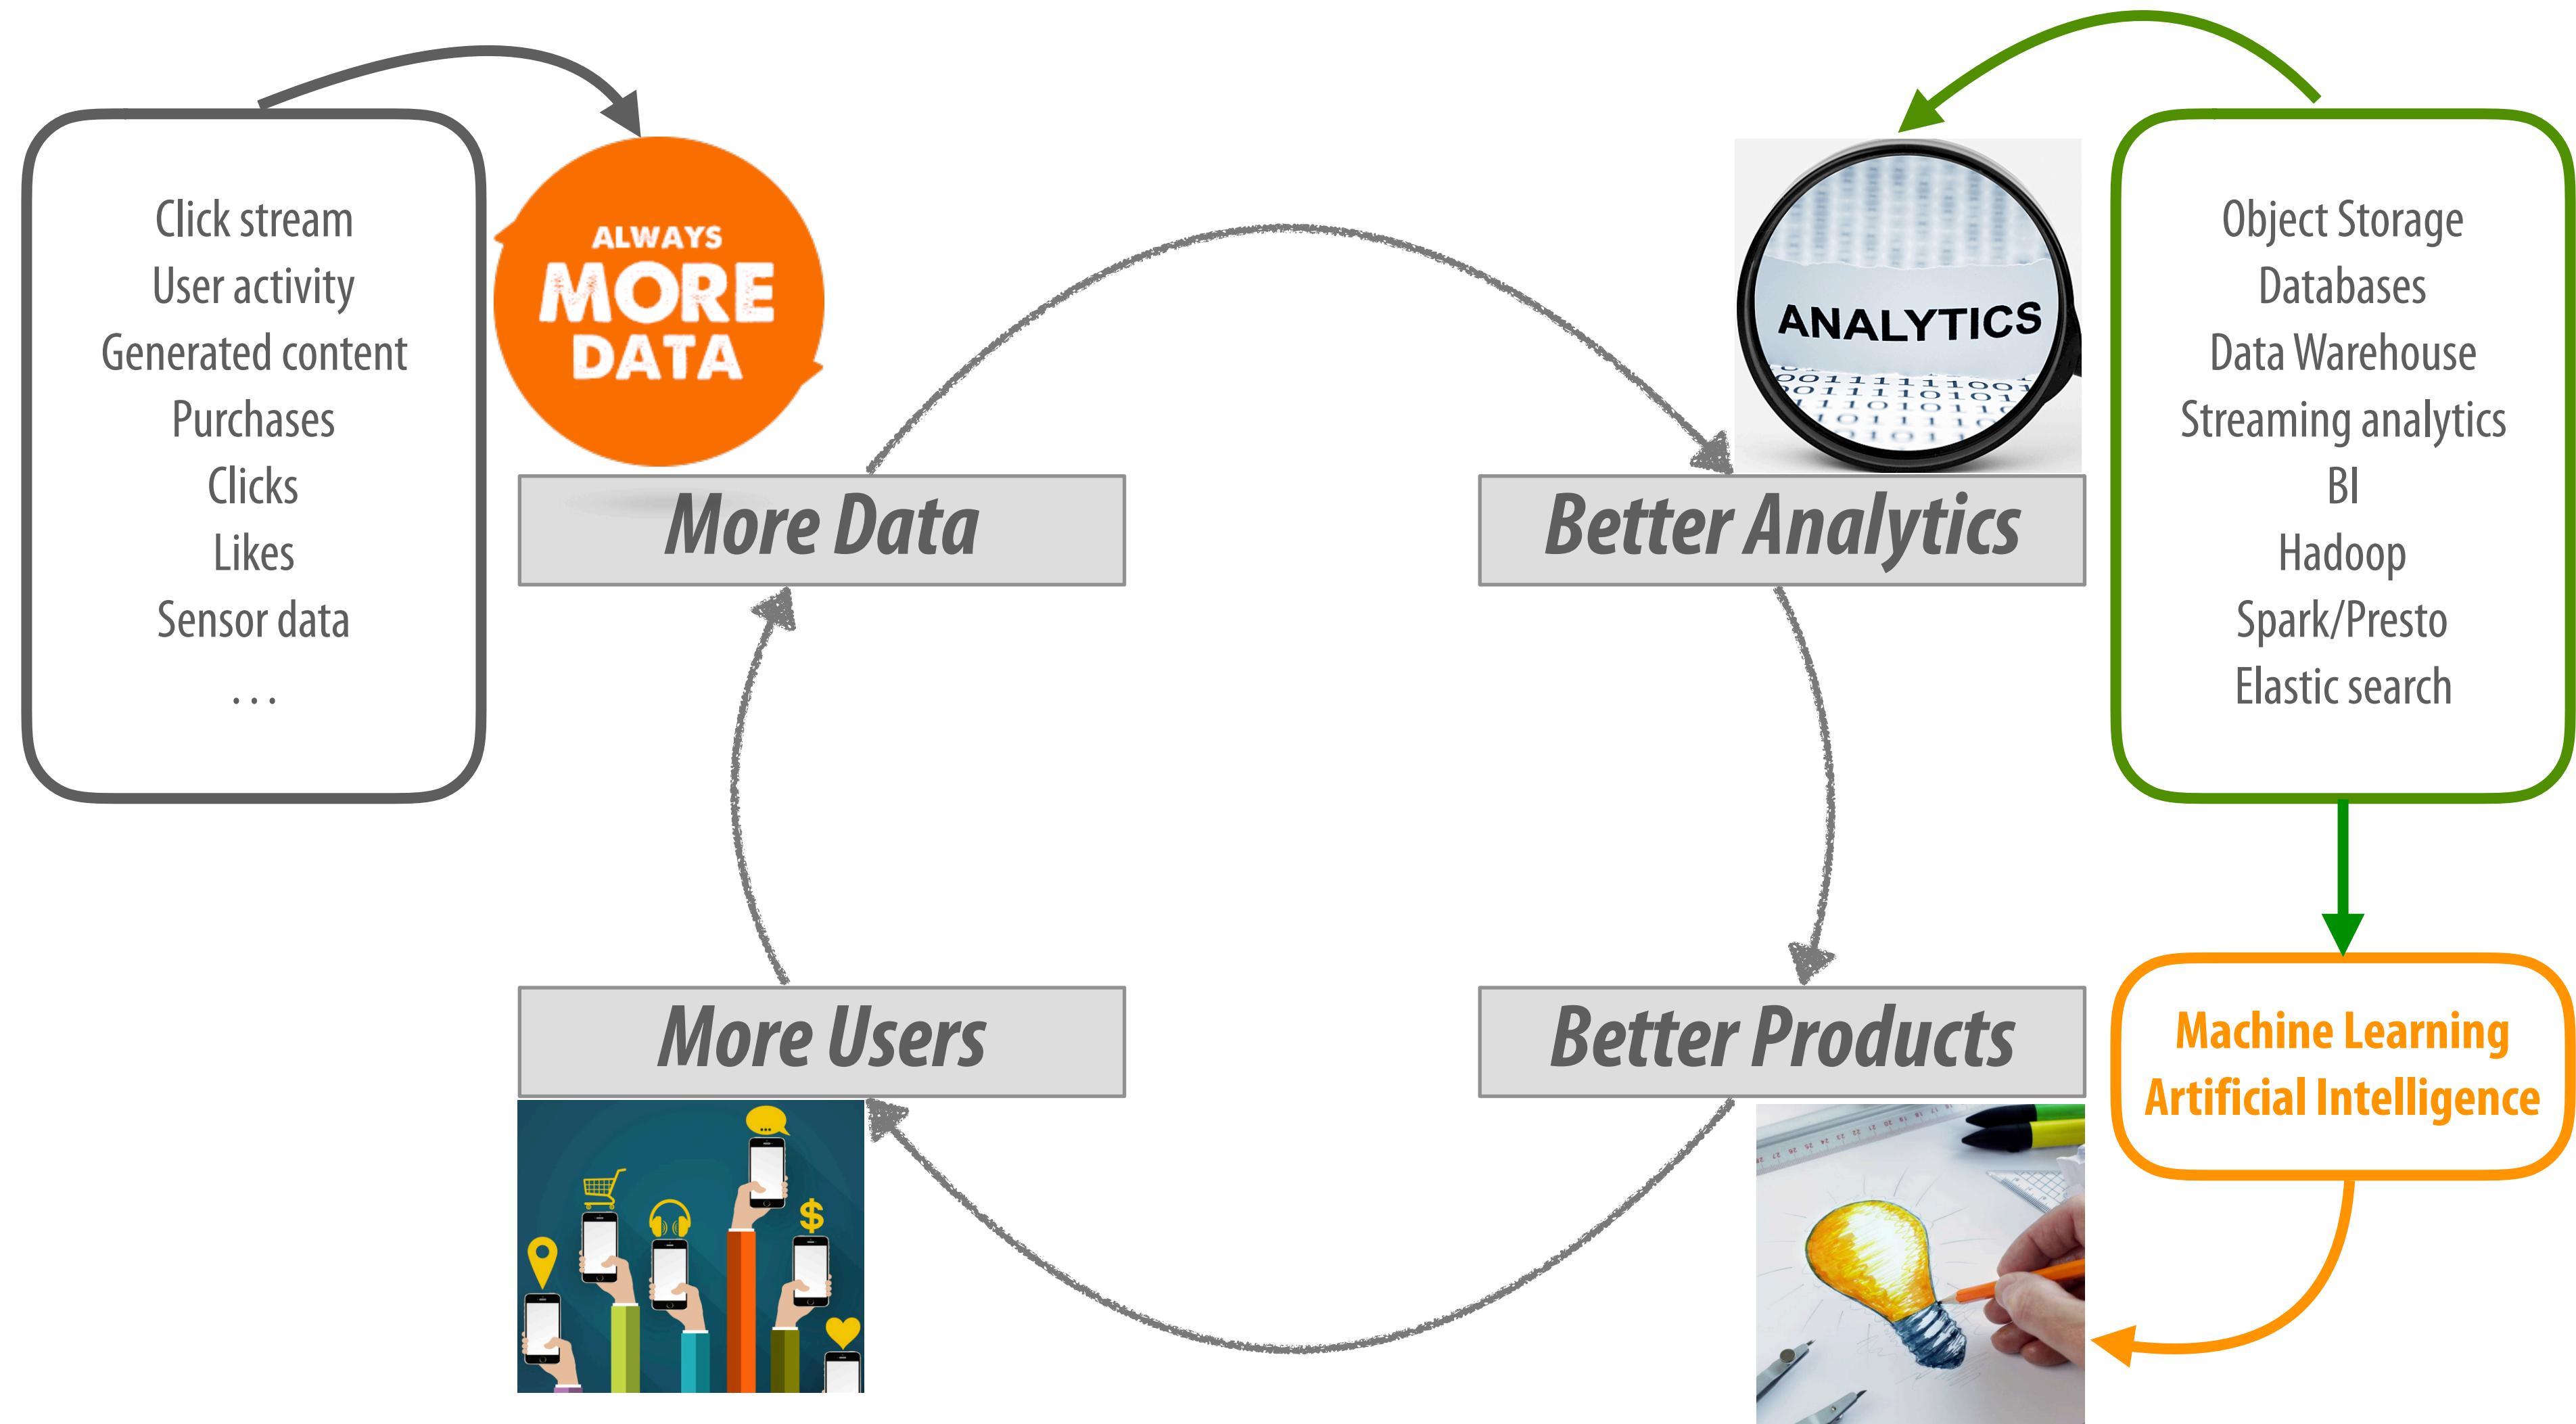
기대효과

과제 제안서의 요약문을 정리할 것, 요약문을 이용해서 정리가 안된 듯함. 강조되는 내용이 없어 보임 [조영조 박사]

강조해야 할 키워드들이 제대로 반영이 안되어 있음. 키워드(기술, 정량적으로 강조)으로 처리해 주기 바람 [조영조 박사]

DataSets for Autonomous Digital Companion

57



자율지능 디지털 동반자를 위한 학습 및 시험 데이터 확보

From My Perspective

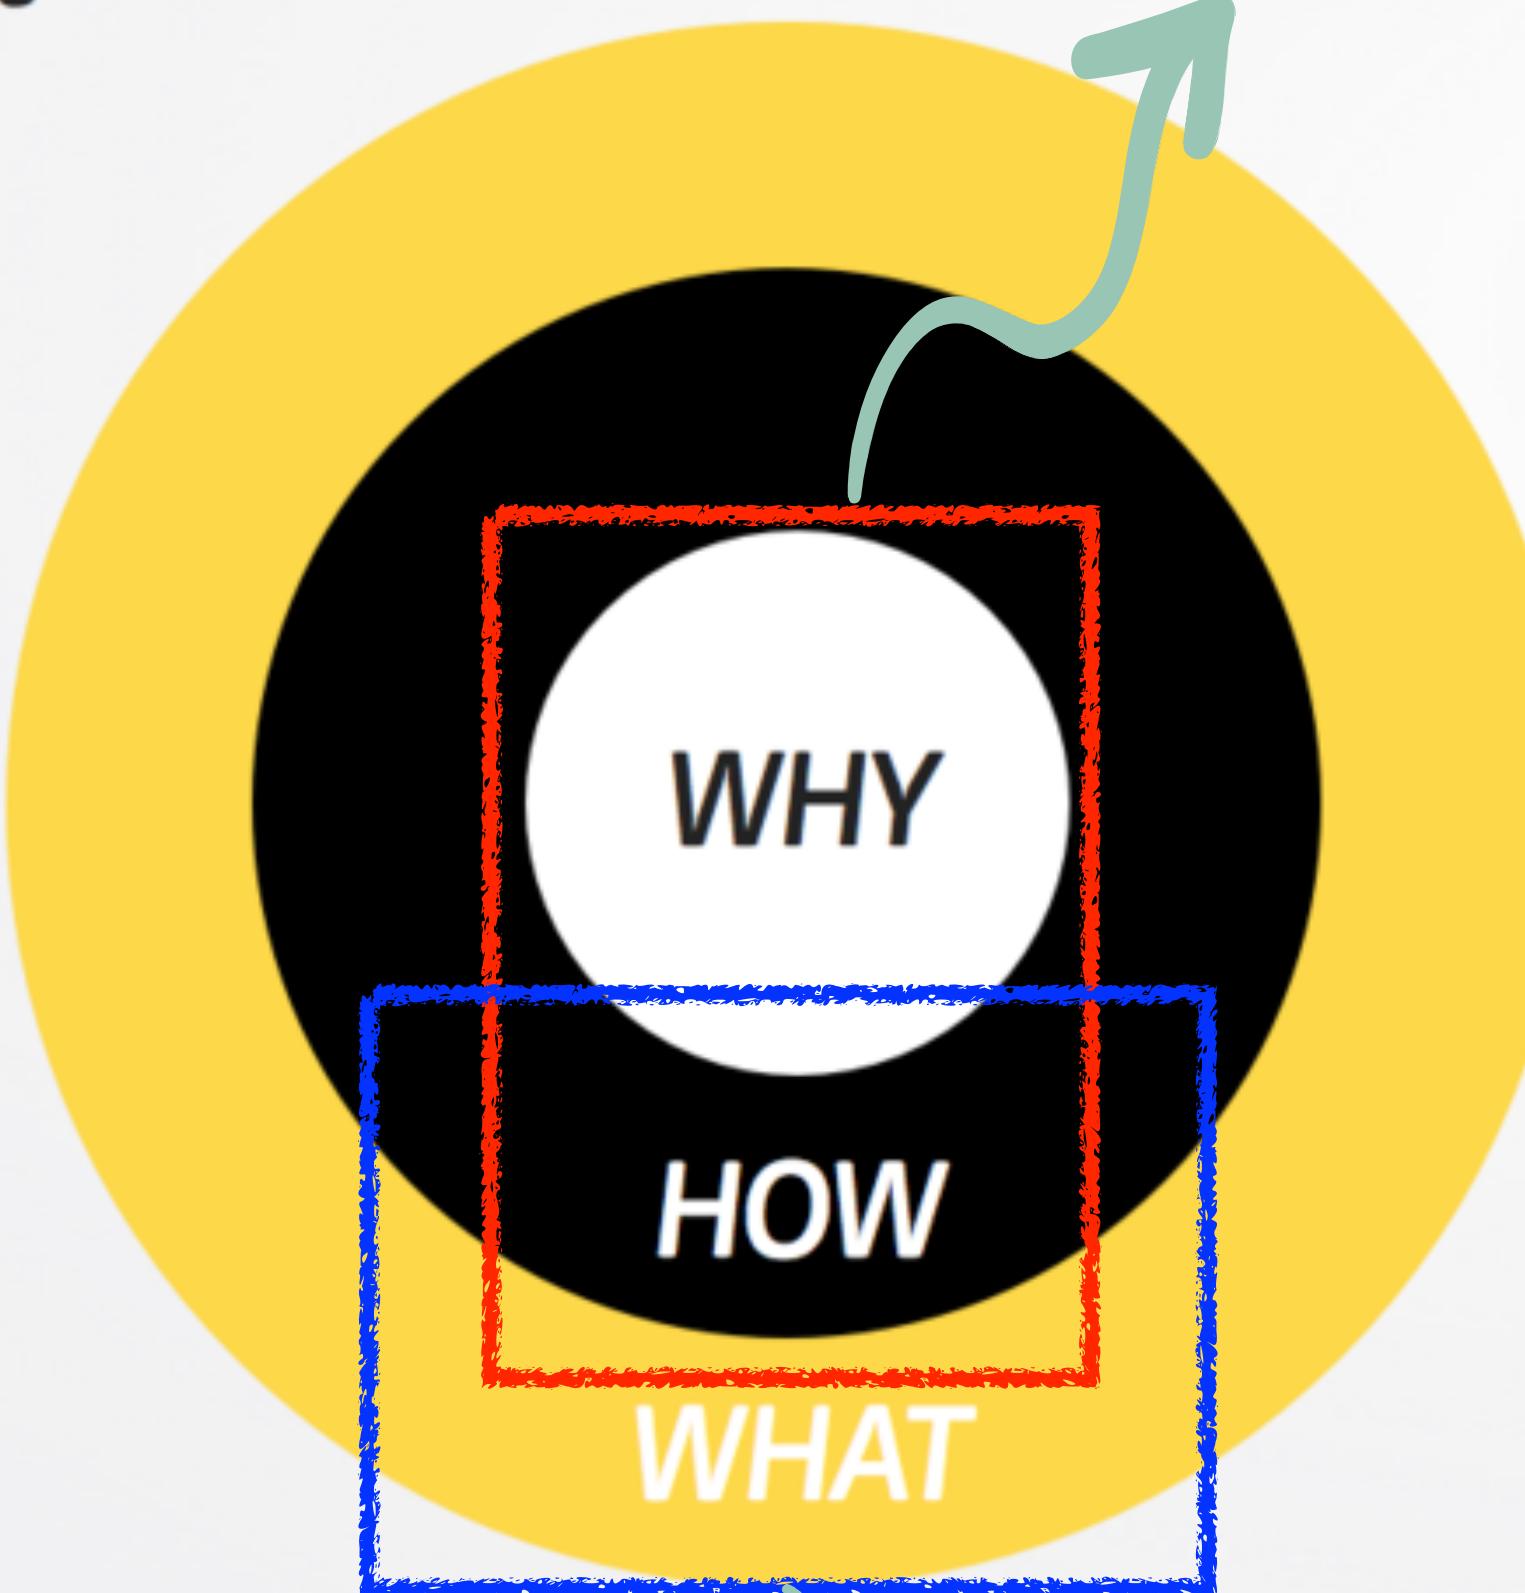
The Golden Circle

59

‘진행 단계별 (과제수주, 과제수행) 전략이 핵심이다!’

The Golden Circle

Start from the Why, and work your way down.



과제수주 단계

WHY

The single purpose, cause or belief that serves as the unifying, driving and inspiring force for any individual or organization.

HOW

Written as verbs as they are actions to be performed and not just inactionable values to be admired, e.g. Do the right thing vs. integrity.

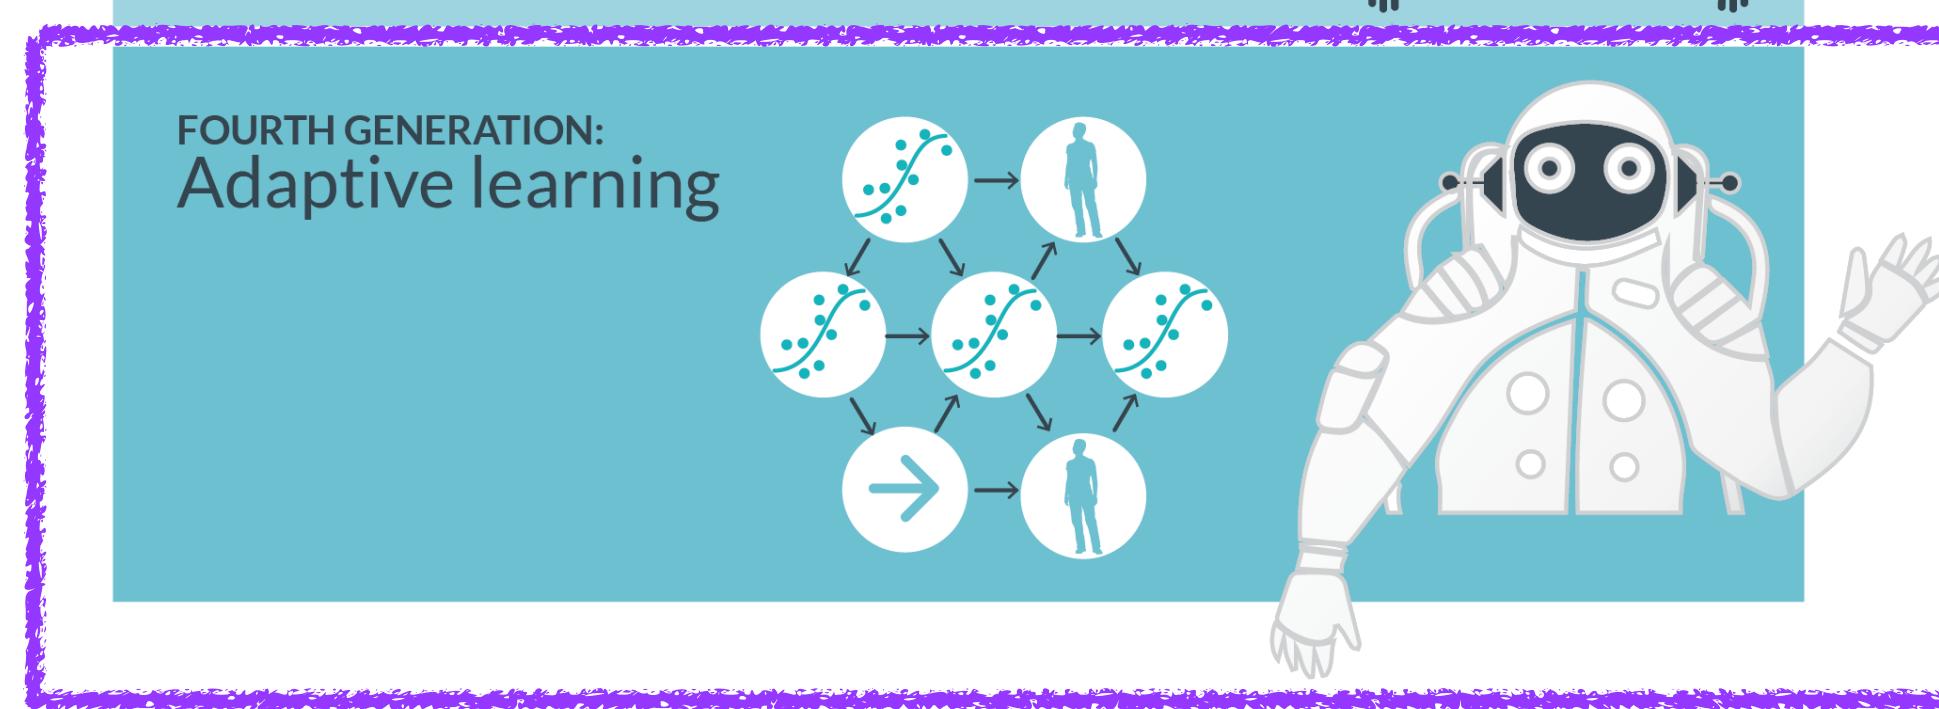
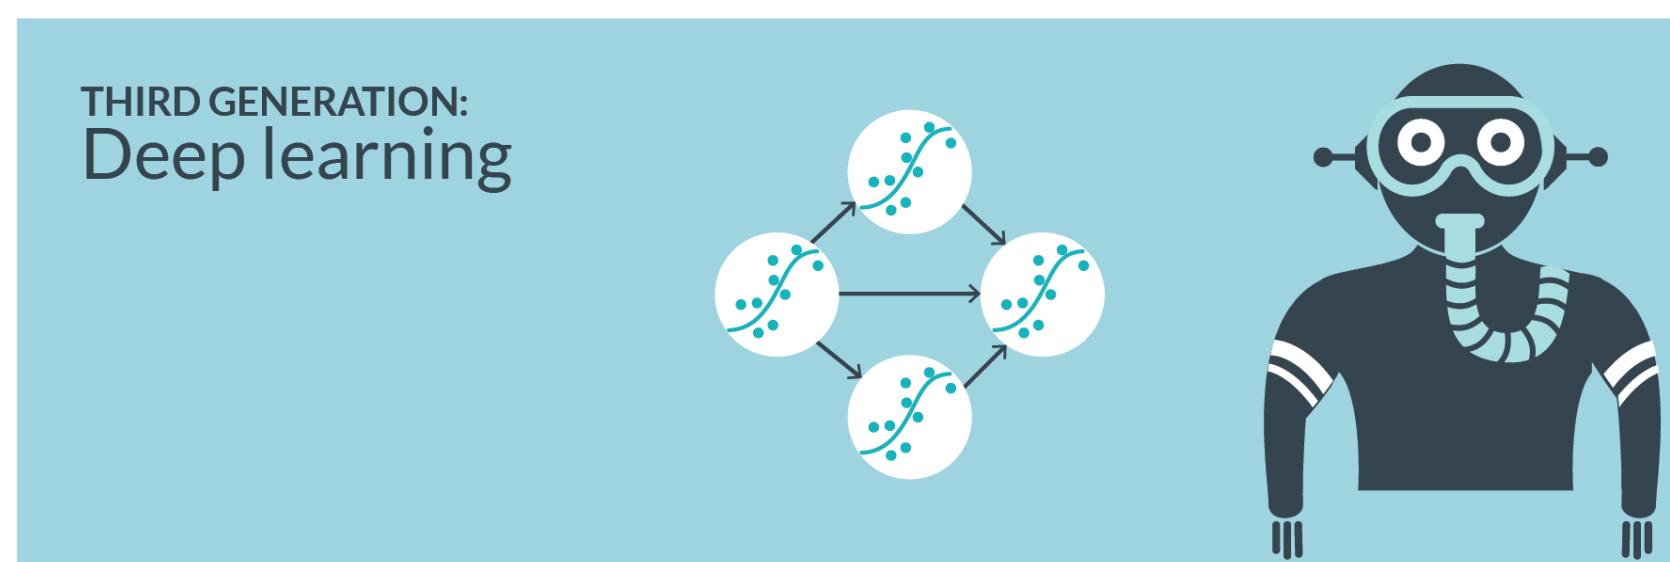
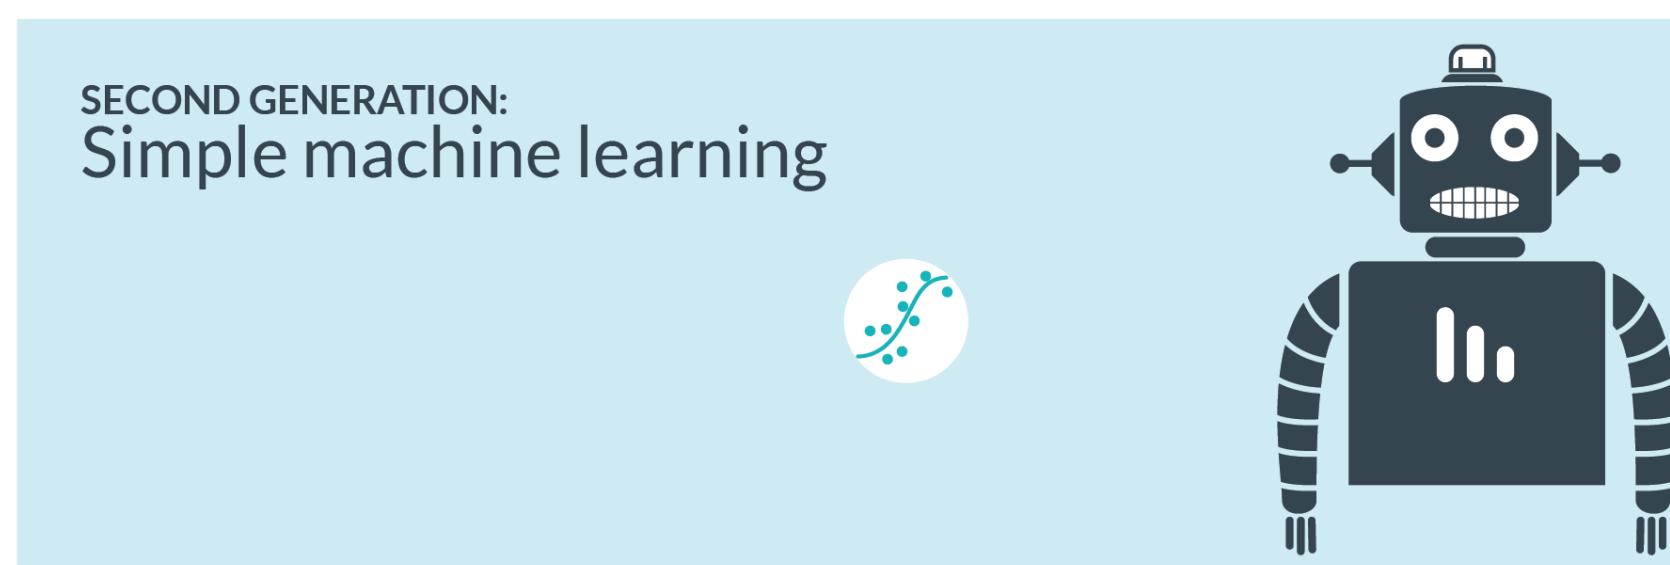
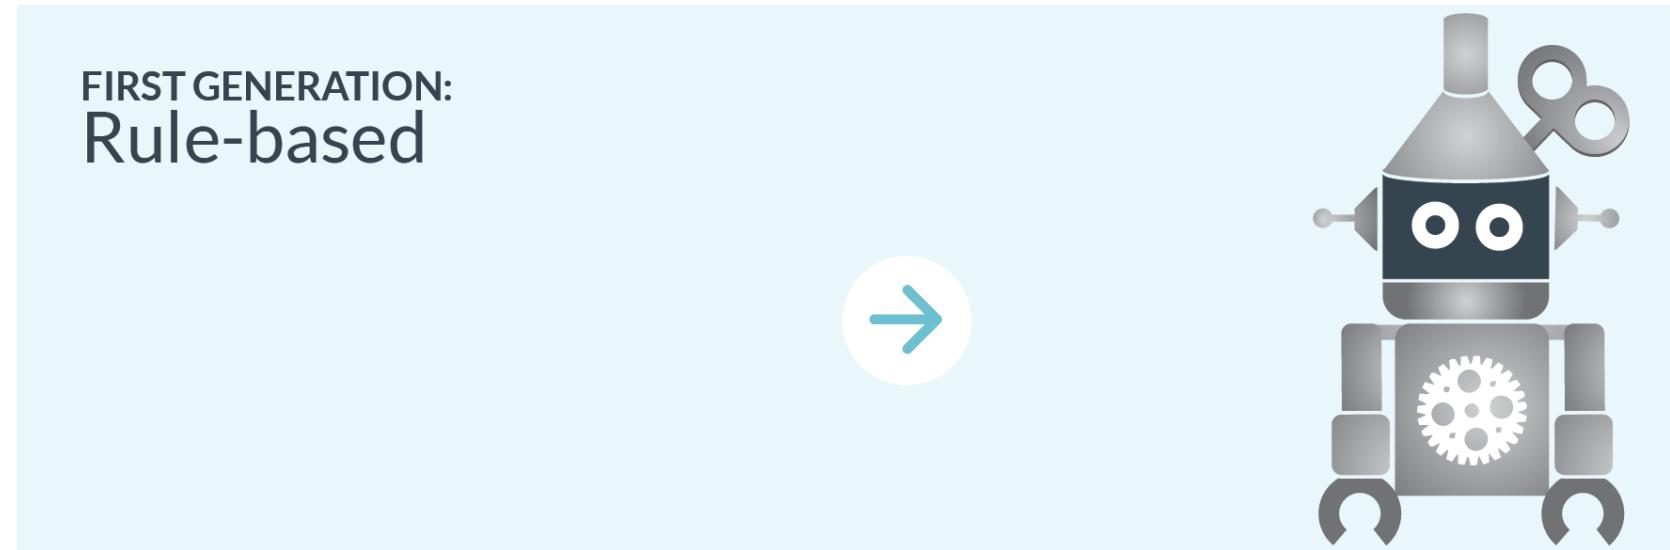
WHAT

Everything tangible an organization says or does. Everything outsiders can see, hear or experience, e.g. products, services, marketing.

과제수행 단계

Four Generations in Machine Learning

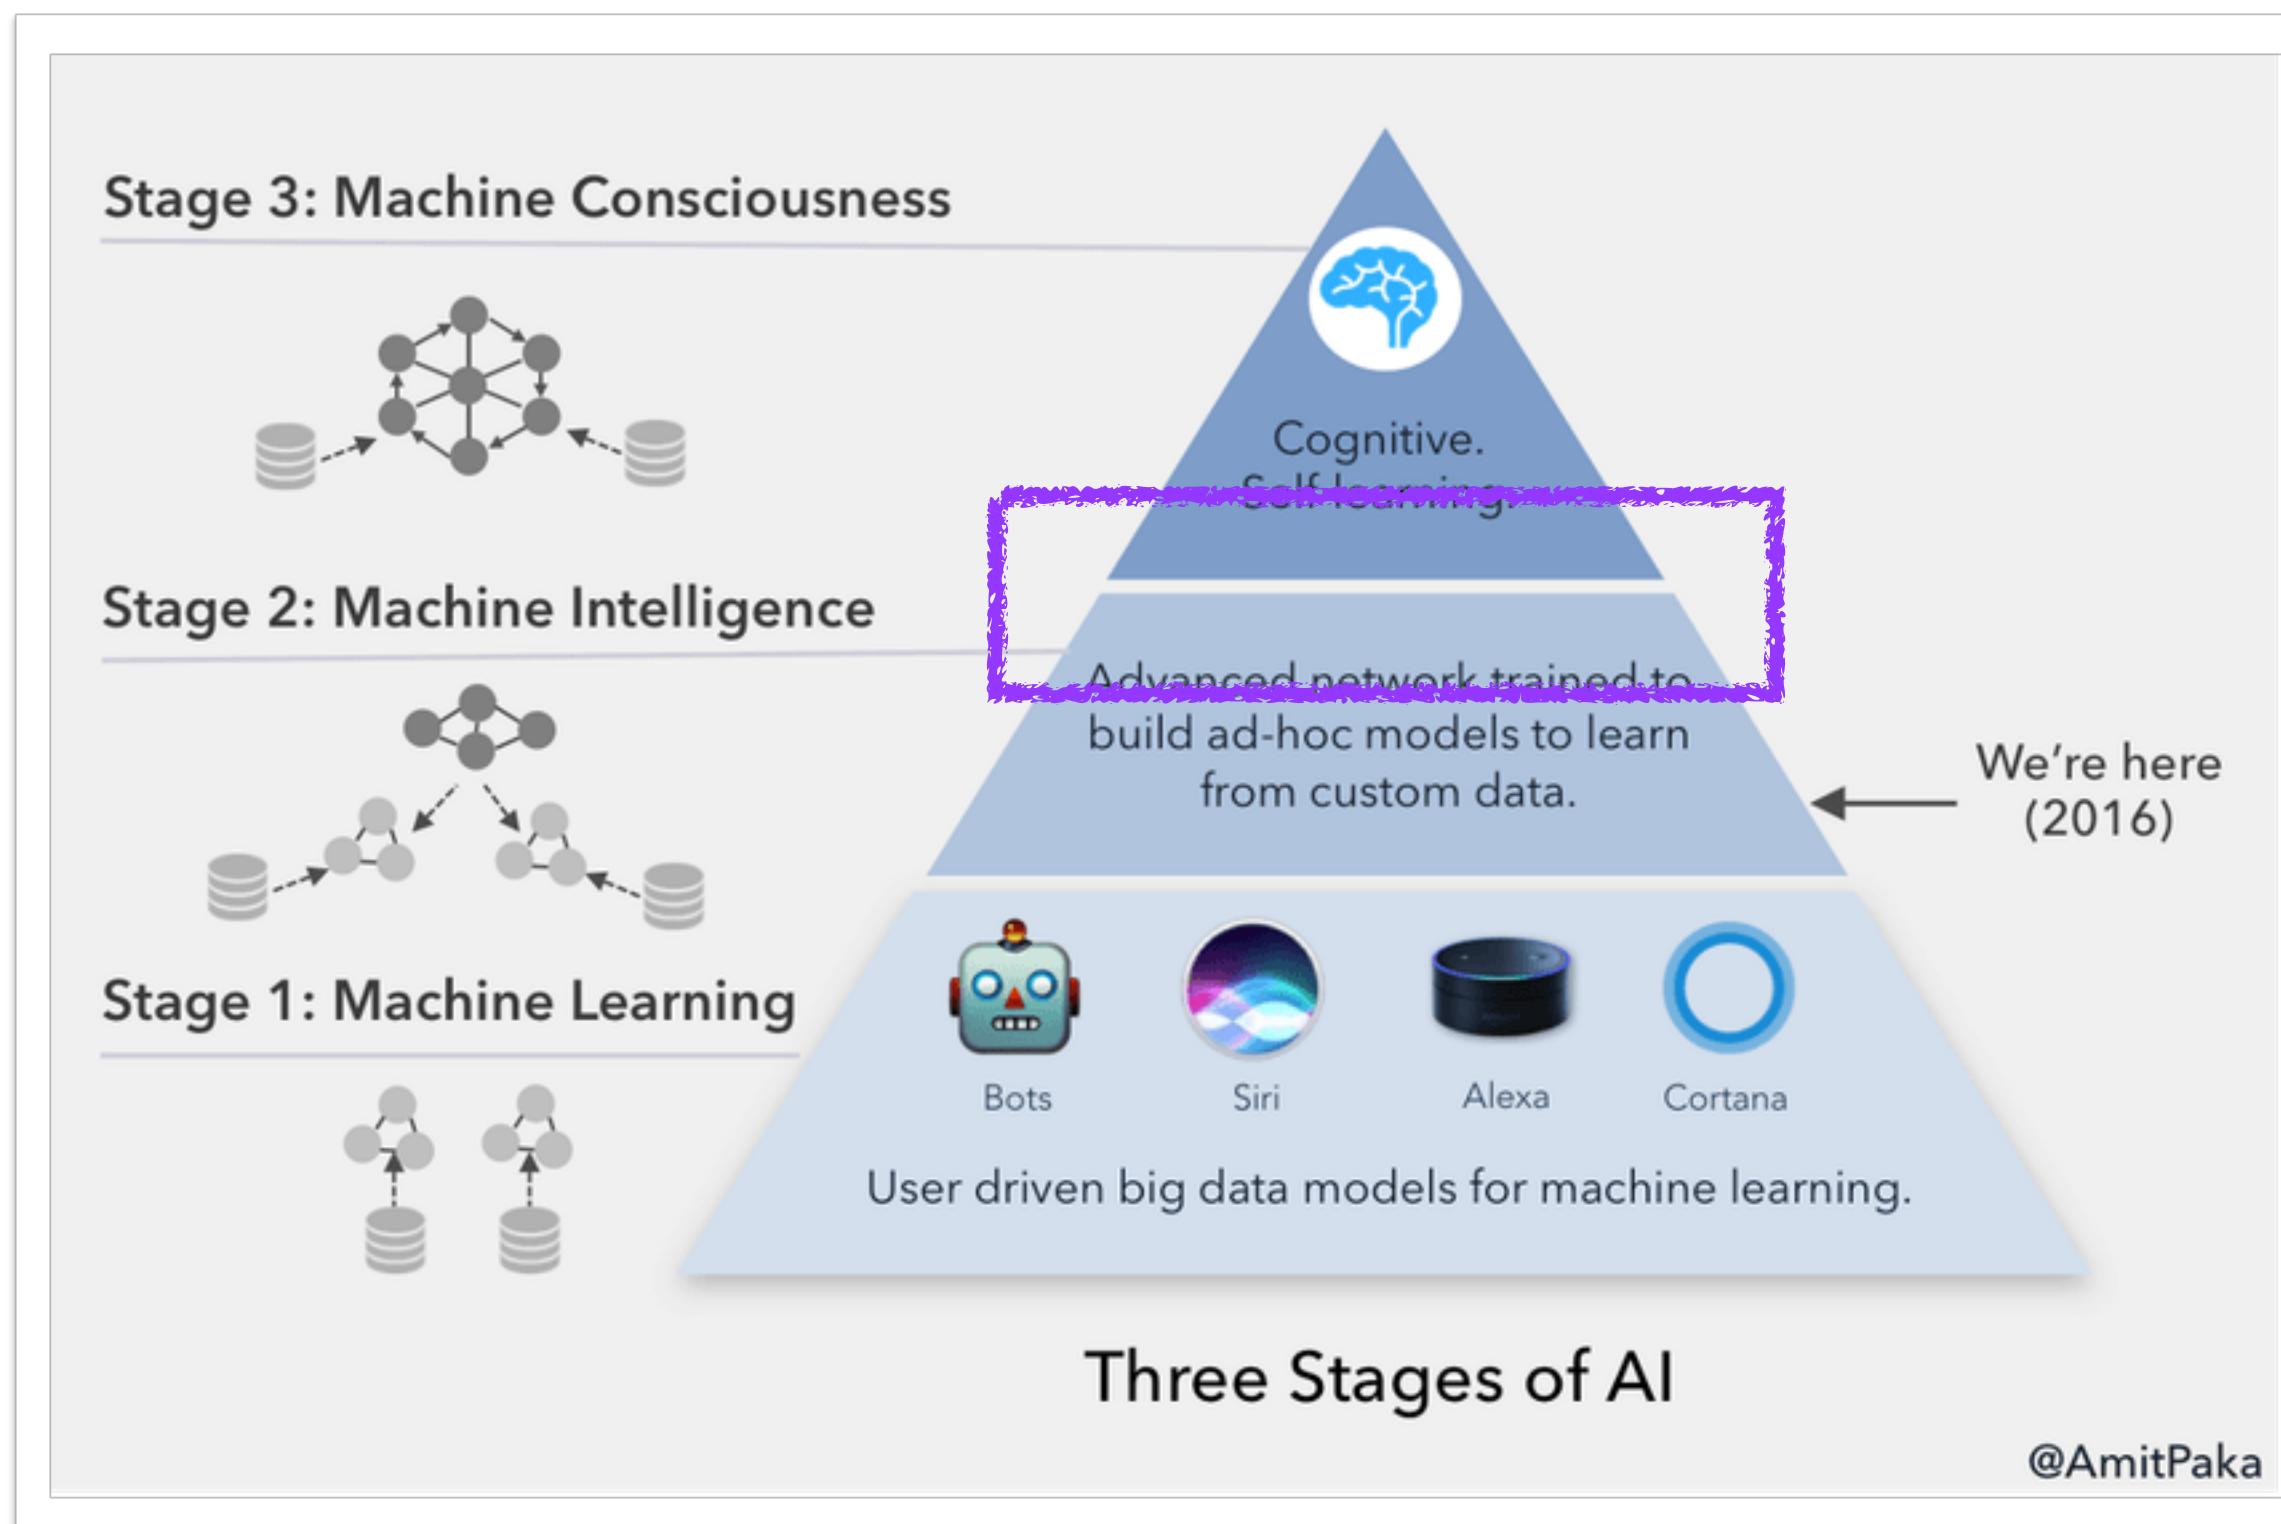
60



Three Stages of Artificial Intelligence

61

- **First Stage: Artificial *Narrow* Intelligence**
- **Second Stage: Artificial *General* Intelligence**
- **Third Stage: Artificial *Super* Intelligence**



Modeling Multimodal Data and Dynamics

63

별첨 5. 세부 연구내용의 창의성 -WP4. 스마트 데이터 관리 기술

“3V가 3V를 만날때”: 3V 빅데이터를 이용한 3V 멀티모달 데이터 모델링

빅데이터

Volume
Variety
Velocity



Toward *Natural Interaction with DAISY!*

사용자

5W context
Preferences



사용자 컨텍스트 데이터

인터랙션 메타데이터 모델링

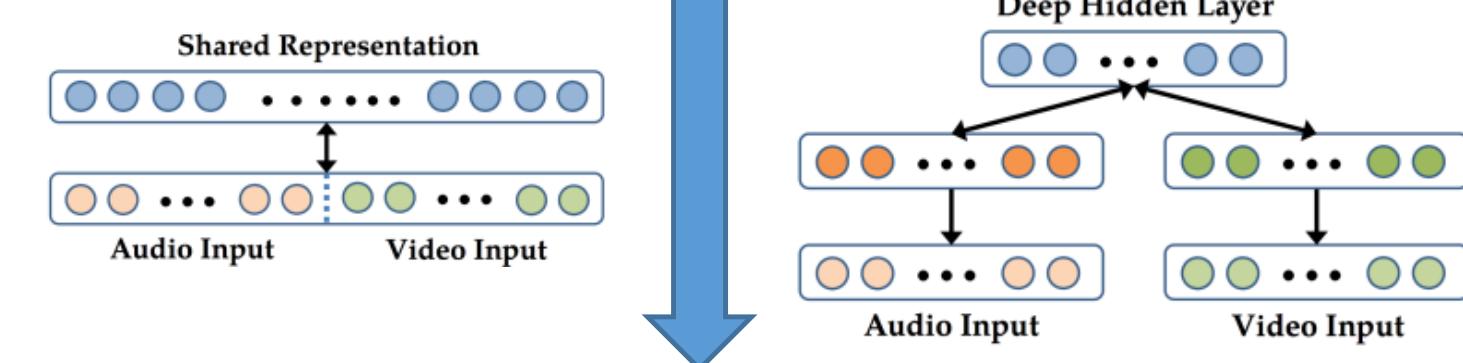
사용자 컨텍스트 데이터 모델링



멀티모달 데이터

Verbal
Vocal
Visual

Modeling Multimodal Dynamics!



멀티모달 데이터를 이용한 기계학습 접근법

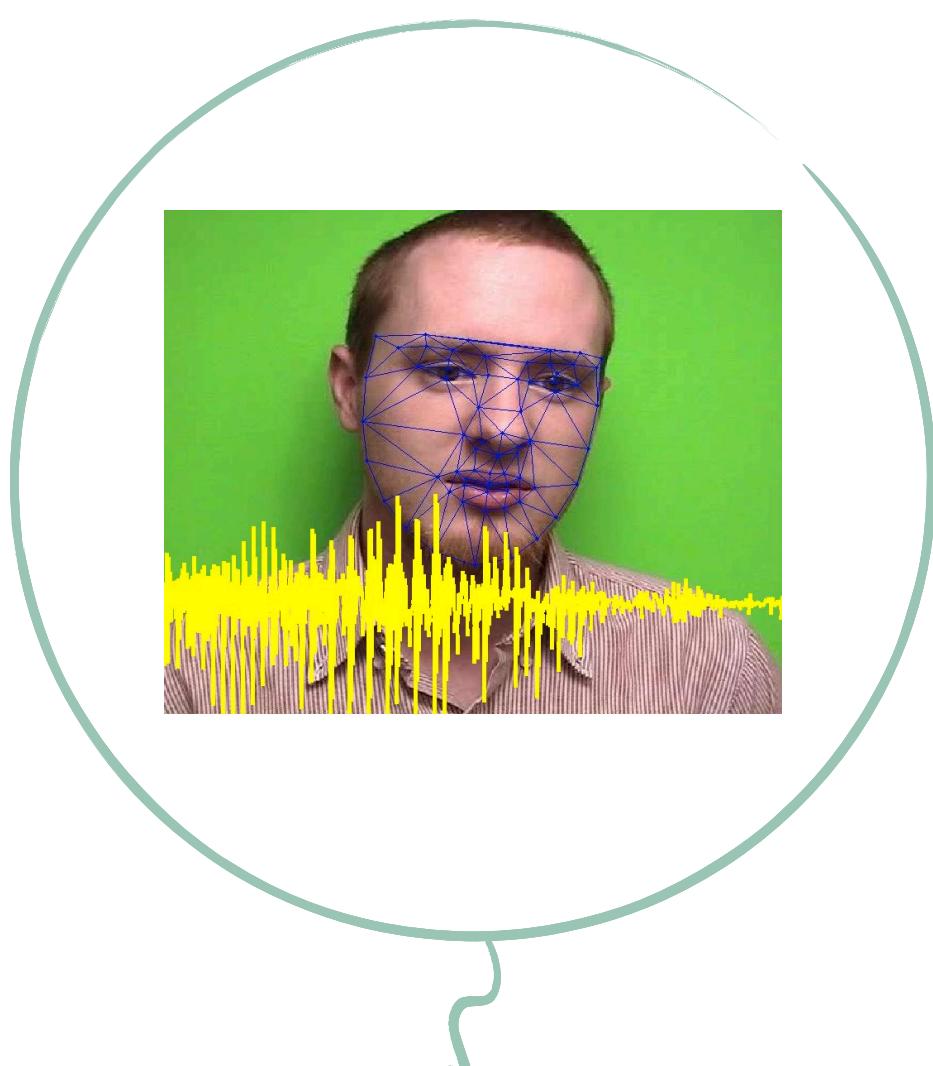
하이브리드 멀티모달/Personalized 데이터 모델링 기술

- 멀티모달 기계학습용 인터랙션 메타데이터 모델링
- 적응형 기계학습을 위한 사용자 컨텍스트(context) 데이터 모델링
- 응용서비스 도메인별 개념 모델링(concept modeling)
- 스마트 데이터 구축을 위한 애널리틱스 베이스 테이블(ABT) 모델링

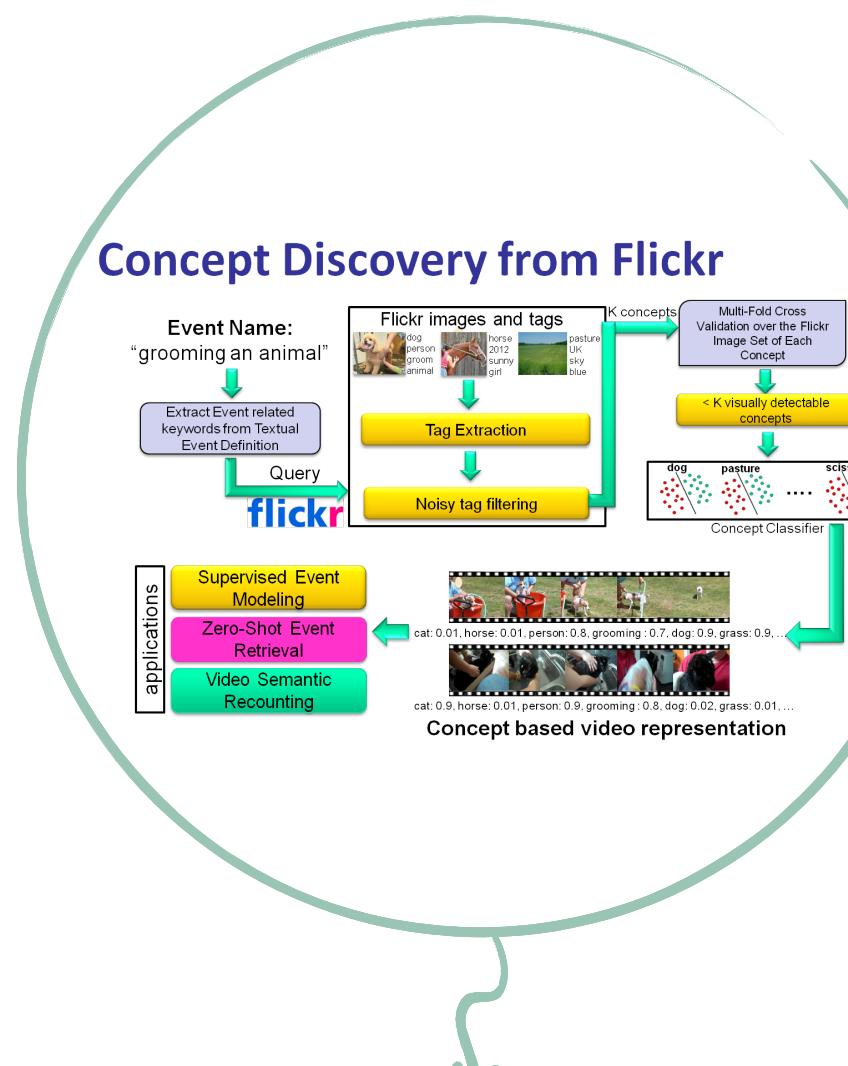
	Feature Learning	Supervised Training	Testing
Classic Deep Learning	Audio	Audio	Audio
	Video	Video	Video
Multimodal Fusion	A + V	A + V	A + V
	A + V	Video	Video
Cross Modality Learning	A + V	Audio	Audio
	A + V	Audio	Video
Shared Representation Learning	A + V	Video	Audio
	A + V	Video	Audio

Multimodal Machine Learning

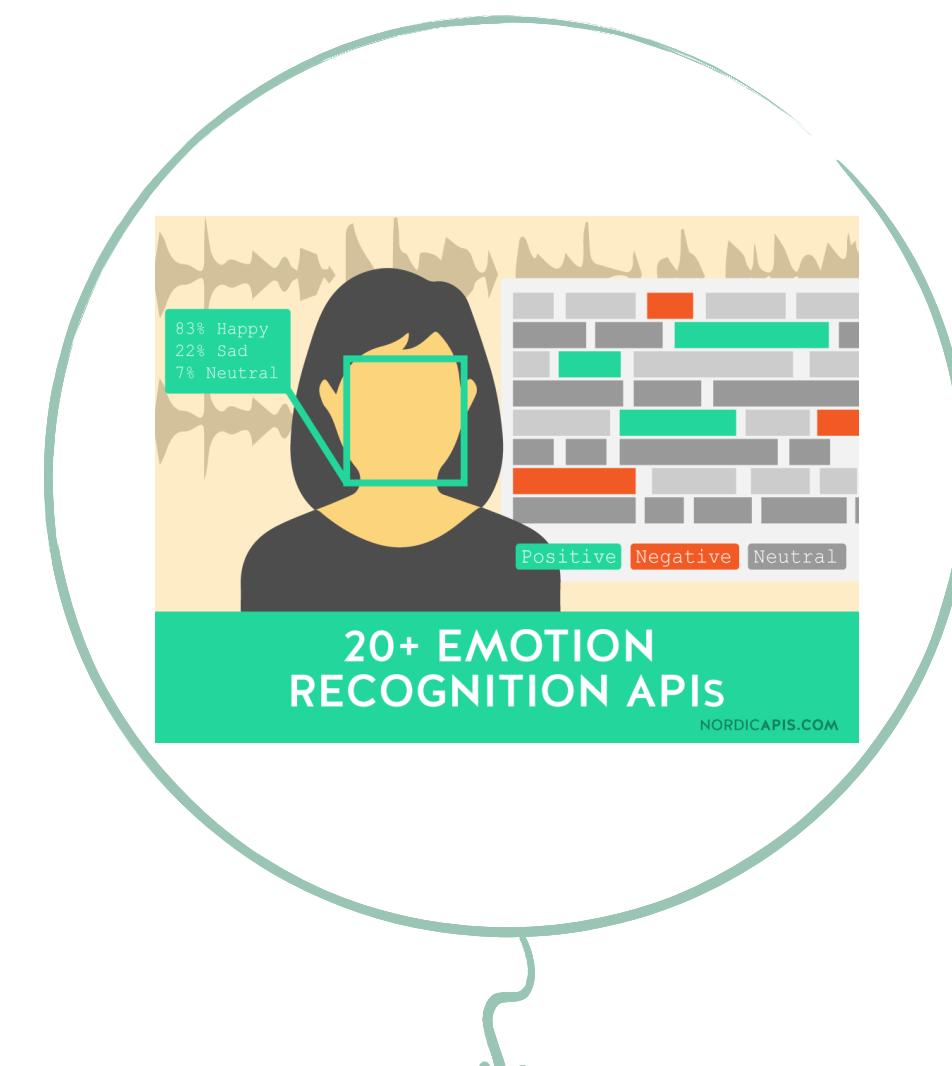
64



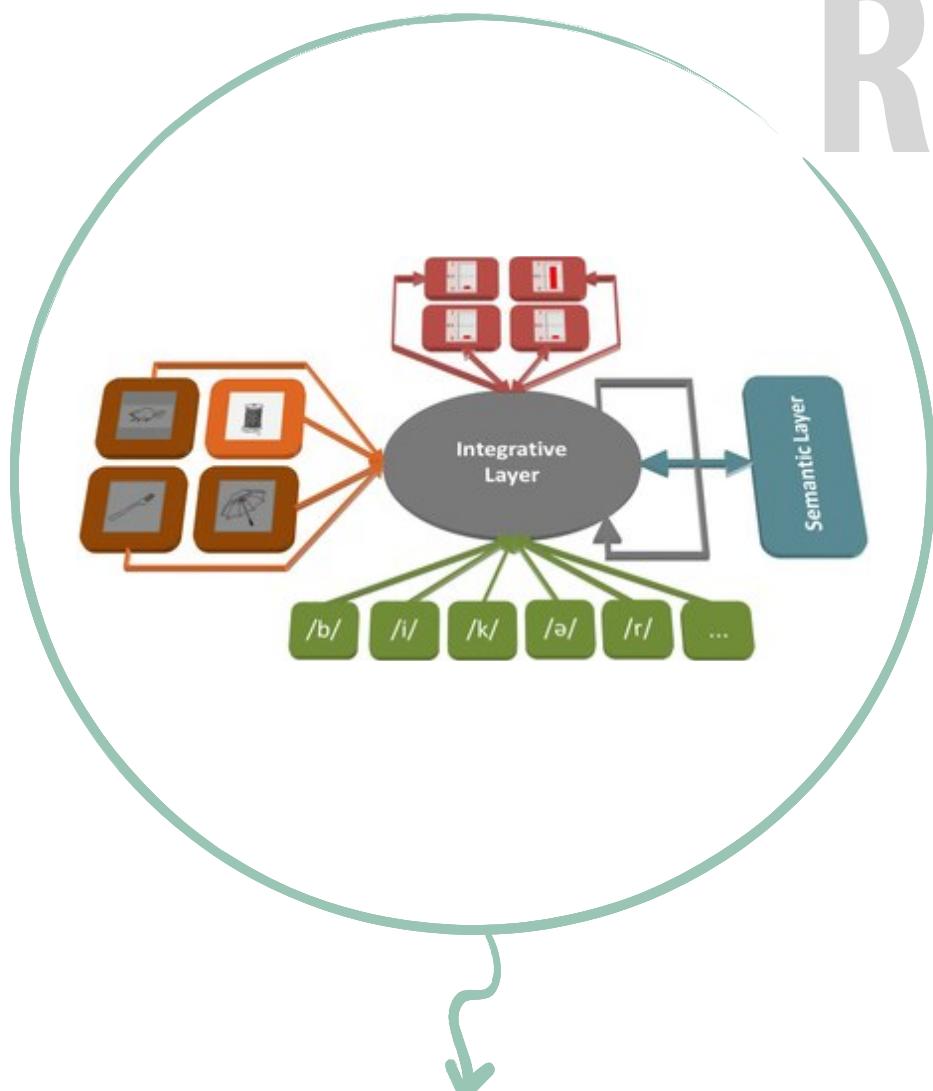
Audio-visual speech recognition



Cross-media event retrieval and tagging



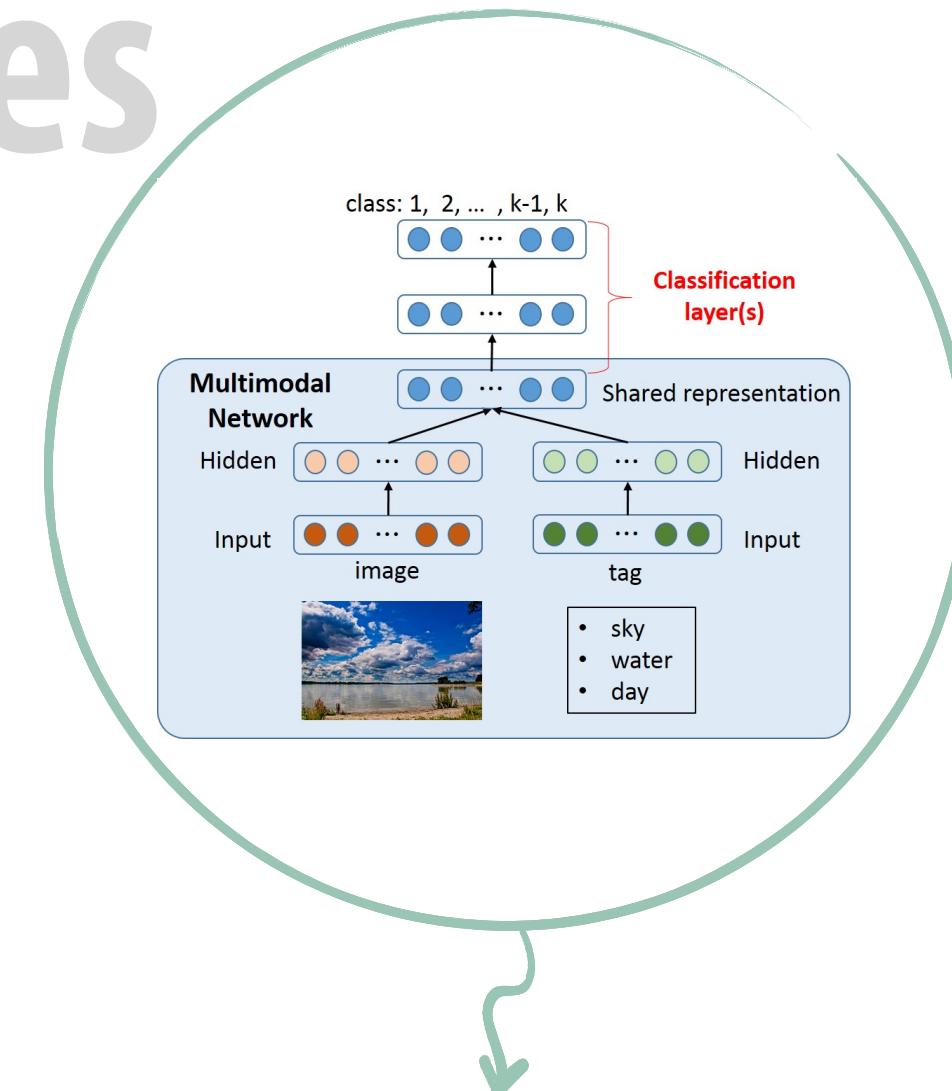
Audio-visual emotion recognition



Multimodal language processing



Image and video captioning



Multimodal deep learning

Research Challenges

Continuing to Learn... Keep Learning!

65



Raise your standards as high as you can live with, avoid wasting your time on routine problems, and always try to work as closely as possible at the boundary of your abilities.

Do this because it is the only way of discovering how that boundary should be moved forward.

- Edsger W. Dijkstra (1930-2002)

thank you

The image features a central word cloud composed of the words "thank" and "you" in multiple international languages. The words are arranged in a radial pattern, with "thank" at the top and "you" below it, surrounded by other words. The languages include English, German, Spanish, French, Dutch, Italian, Portuguese, Russian, Polish, Chinese, Korean, Japanese, Vietnamese, Thai, Indonesian, Malaysian, Arabic, Hebrew, and others. Each language's word is written in its native script or Romanized form, color-coded in a rainbow palette.

Appendix

Probability: Frequentists vs. Bayesians

68

■ Frequentist Interpretation

- Probabilities represent **long run frequencies** of events $Prob(x) = \frac{occurrence(x)}{N}$
- Example. # of coin flips : *head or tail?*

■ Bayesian Interpretation

- Probability is used to quantify our **uncertainty** about something; hence it is fundamentally related to **information** rather than repeated trials.
- The probability that certain events happen **cannot be computed** by a numeric frequency because we cannot repeat them
- Bayesian approach focuses on computing **how uncertain** (or **certain**) the event may happen by means of collected observation data (**evidences**)
- Bayes' Theorem

$$p(\mathbf{w}|D) = \frac{p(D|\mathbf{w})p(\mathbf{w})}{p(D)}$$

$$p(\mathbf{w}|D) = \frac{p(D|\mathbf{w})}{p(D)} \cdot p(\mathbf{w})$$

(in a form of recurrence formula)

- \mathbf{w} : **Target subject of uncertainty**
- D : Observed data
- $p(\mathbf{w}|D)$: **posterior** belief
(updated belief after an observed data set is given)
- $p(D|\mathbf{w})$: **likelihood**
- $p(\mathbf{w})$: **prior** belief about an event \mathbf{w}
- $p(D)$: **evidence**

Bayesian Probability

69

■ A sample example: “guessing a gender of a person”

- S#1. “A person is shown blinded”

$$p(\text{female}) = \frac{1}{2}$$

- S#2. “The person is taller than 170cm!”

- Likelihood: $p(D|\text{female}) = \frac{1}{50}$

- Evidence: $p(D) = \frac{1}{10}$

- Posterior: $p(\text{female}|D) = \frac{1}{10}$

- S#3. “The person’s hair is longer than 30cm!”

- Likelihood: $p(D|\text{female}) = \frac{3}{5}$

- Evidence: $p(D) = \frac{1}{3}$

- Posterior: $p(\text{female}|D) = \frac{9}{50}$

- S#4. “The person is wearing a skirt!”

- Likelihood: $p(D|\text{female}) = \frac{3}{4}$

- Evidence: $p(D) = \frac{1}{3}$

- Posterior: $p(\text{female}|D) = \frac{81}{200}$

