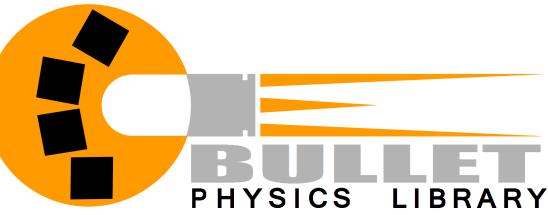




SIGGRAPH 2015

Xroads of Discovery



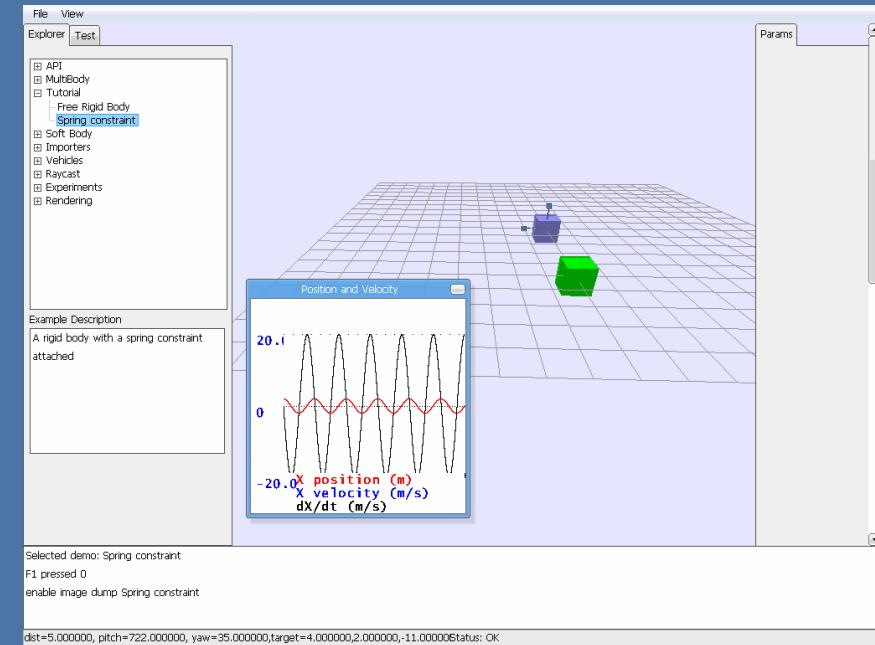
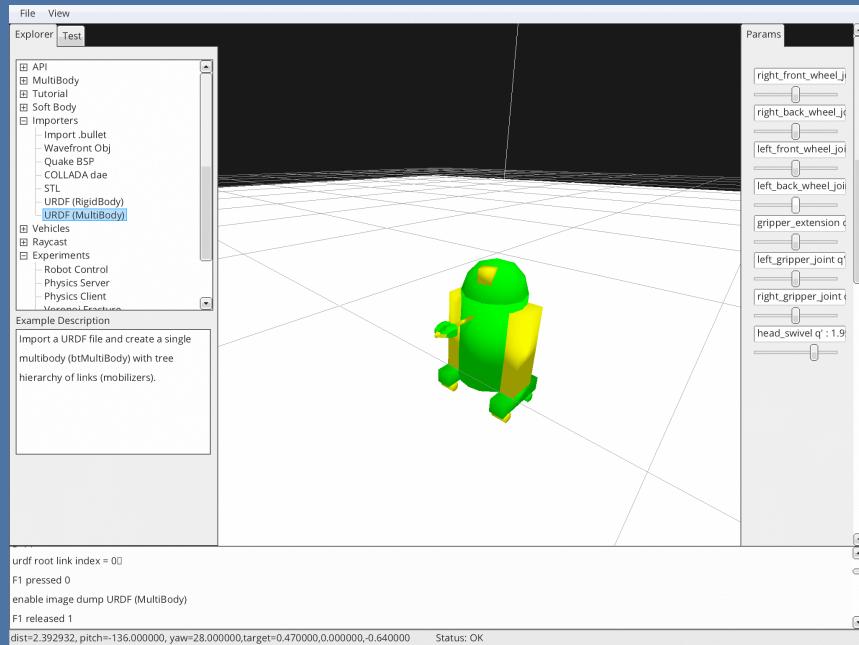
SIGGRAPH 2015
Xroads of Discovery

The 42nd International Conference and Exhibition
on Computer Graphics and Interactive Techniques

Bullet Physics Simulation
OpenCL GPU Acceleration of rigid body
dynamics and collision detection

Erwin Coumans
Google Inc.

Bullet Physics Example Browser



ExampleBrowser --enable_experimental_opengl

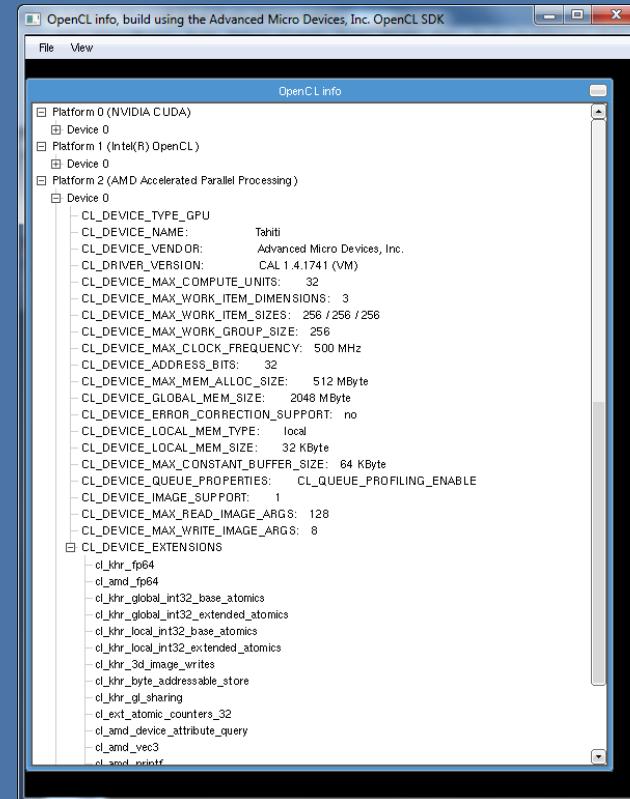
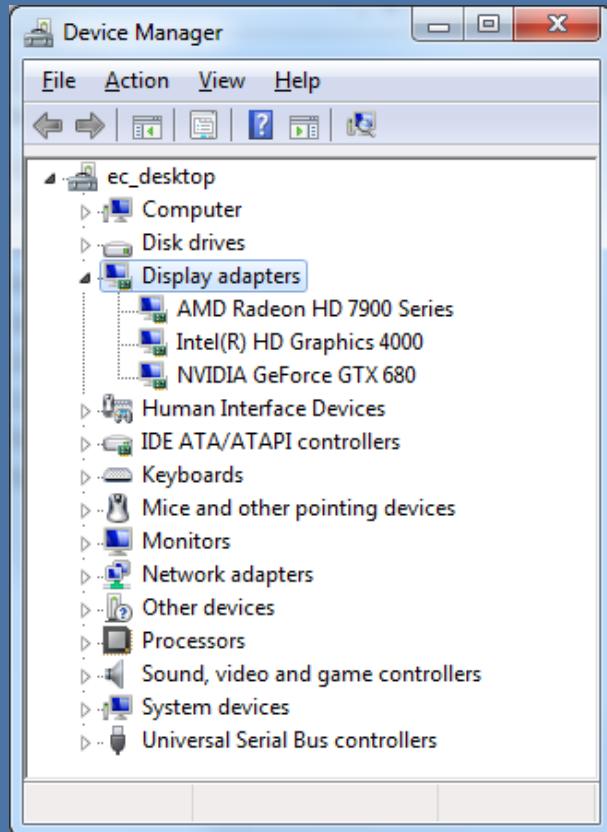
github.com/bulletphysics/bullet3



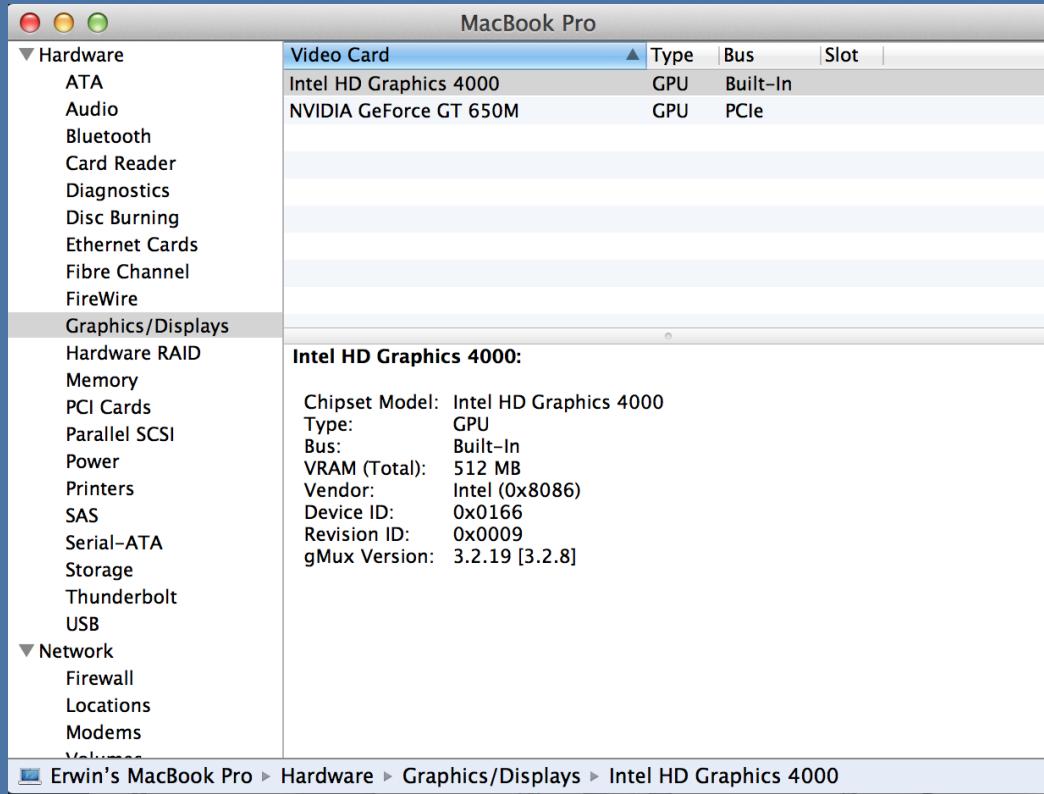
GNU Make



Windows OpenCL Devices



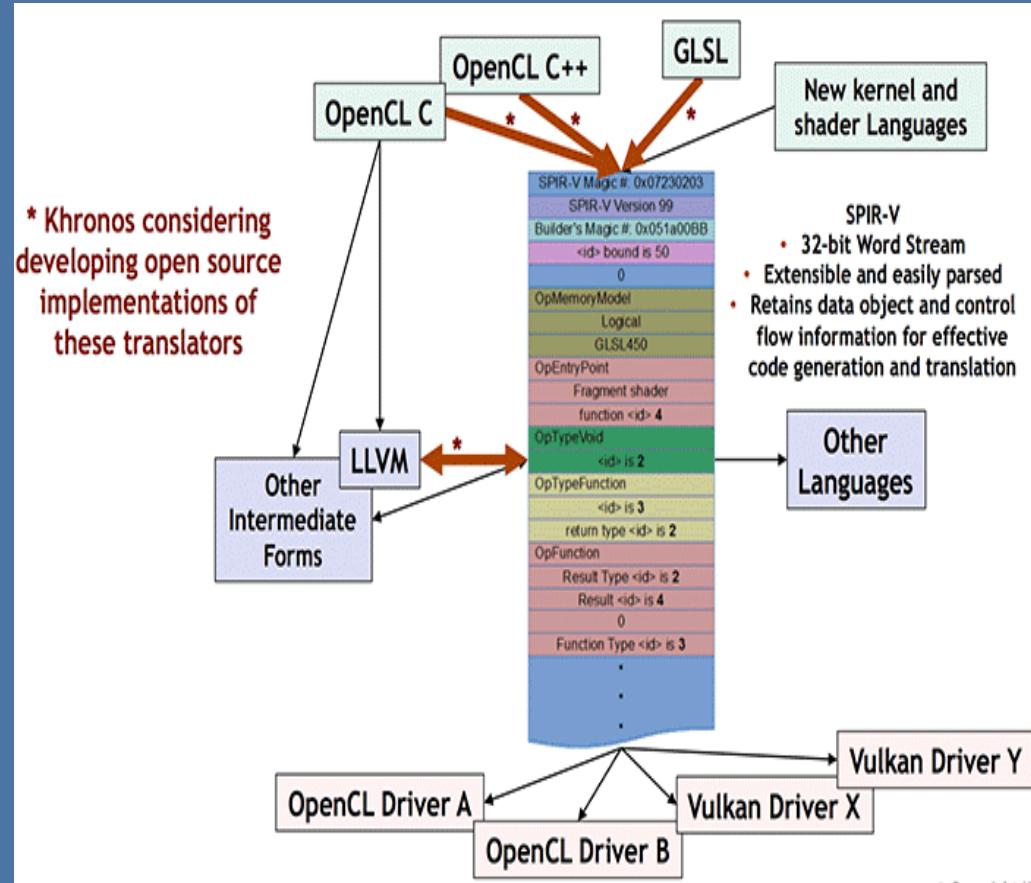
Mac OSX OpenCL Devices



OpenCL, SPIR-V, Vulkan

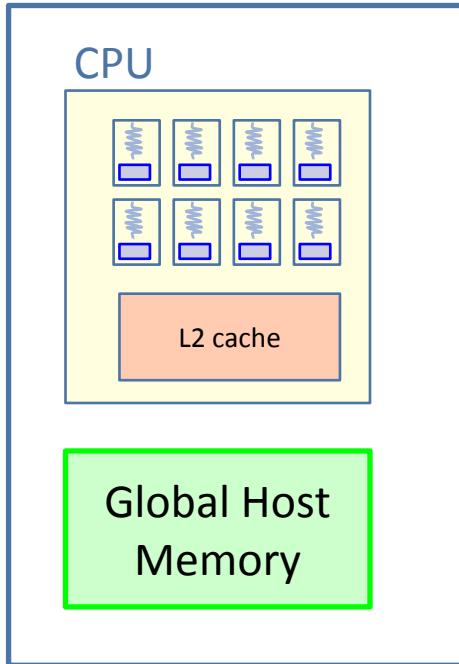


SPIR	SPIR 1.2	SPIR 2.0	SPIR-V
LLVM Interaction	Uses LLVM 3.2	Uses LLVM 3.4	100% Khronos defined Round-trip lossless conversion
Compute Constructs	Metadata/Intrinsics	Metadata/Intrinsics	Native
Graphics Constructs	No	No	Native
Supported Language Feature Set	OpenCL C 1.2	OpenCL C 1.2 OpenCL C 2.0	OpenCL C 1.2 / 2.0 OpenCL C++ GLSL
OpenCL Ingestion	OpenCL 1.2 Extension	OpenCL 2.0 Extension	OpenCL 2.1 Core
Vulkan Ingestion	-	-	Vulkan Core

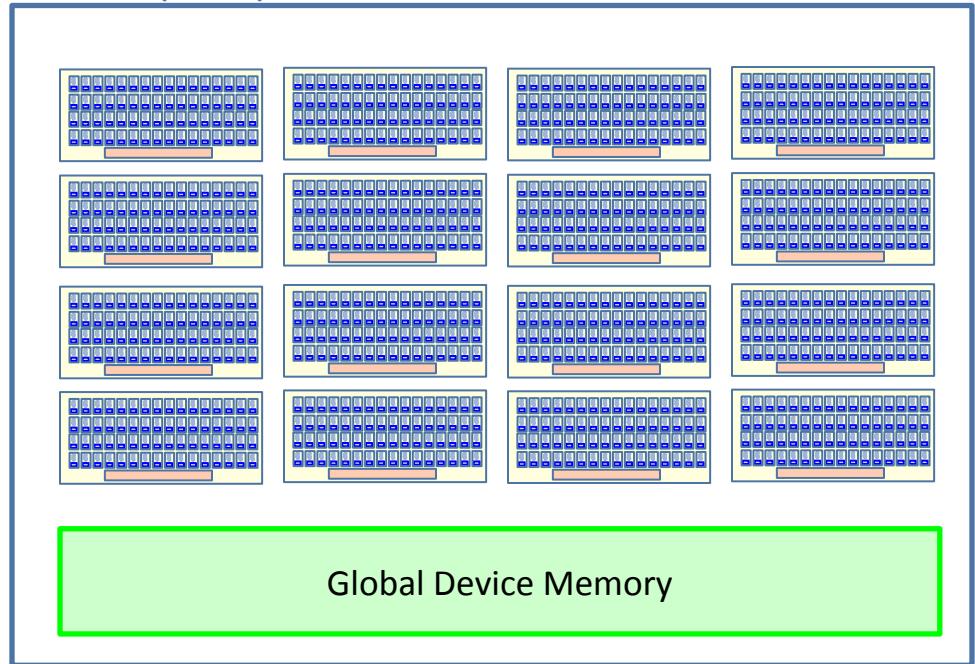


Host and Device

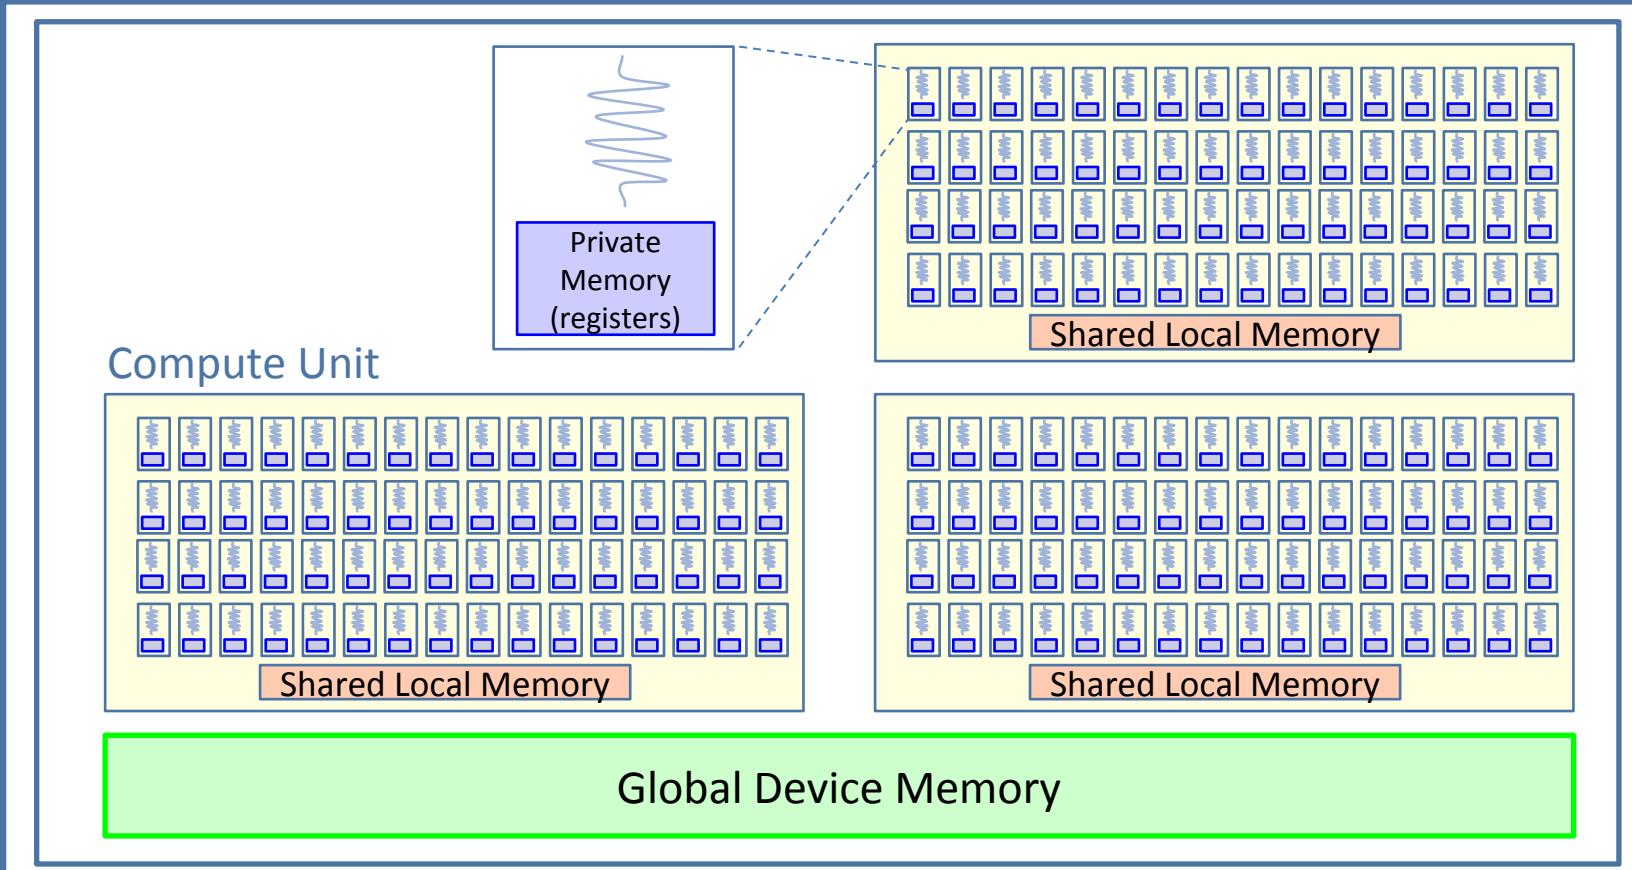
Host



Device (GPU)



GPU in a Nutshell



Refactor or Rewrite from Scratch

Sequential to
Parallel Algorithms

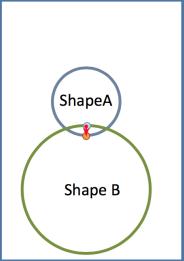
Physics API suitable
for GPU? Callbacks,
Extensibility etc.

C++ virtuals,
templates to C99

Array of Structures
to
Structure of Arrays

Rigid Body Simulation Loop

point on A
point on B
normal B->A
distance



Collision Detection
Contact,
Time of Impact

Forward Dynamics
Compute Inertia, Forces
and Accelerations



Numerical Time Integration
Update Velocity, Position



$$v_{t+\Delta t} = v_t + a\Delta t = v_t + \frac{F_{ext} + F_c}{m} \Delta t = v_t + \frac{F_{ext}}{m} \Delta t + \frac{\text{Impulse}_c}{m}$$

F_c could be constraint forces such as contact, friction, joints

$$x_{t+\Delta t} = x_t + v_{t+\Delta t} \Delta t$$

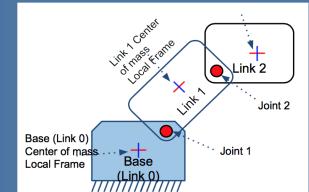
$$p_{correction} = M_{\text{effective}} \Delta v_{\text{desired}} \cdot n$$

$$p_{correction} = \frac{-\Delta v_{ab} \cdot n}{n \cdot n \left(\frac{1}{M_a} + \frac{1}{M_b} \right) + \left(\frac{r_a \times n}{I_a} \right) \times r_a + \left(\frac{r_b \times n}{I_b} \right) \times r_b}$$

$$A\lambda + b \geq 0$$

$$\lambda \geq 0$$

$$\lambda(A\lambda + b) = 0$$



Base (Link 0)

Center of mass

Local Frame

Link 1

Center

of mass

Local Frame

Link 2

Joint

2

Joint

1

Base

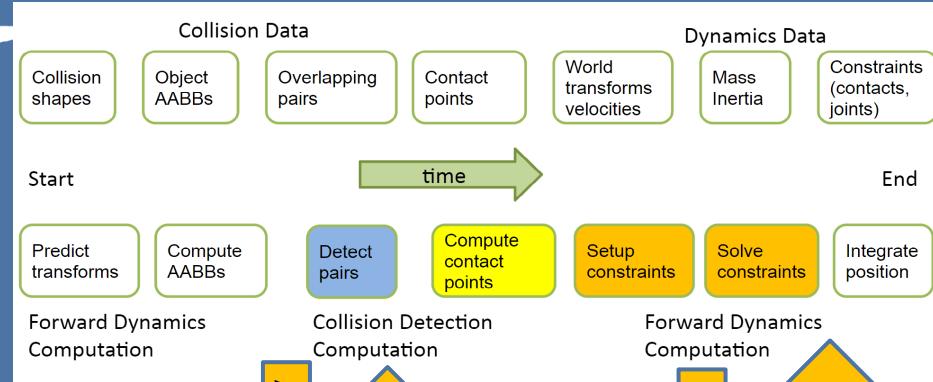
(Link 0)

Refactor: GPU as Co-Processor

Host Memory
Gather and
Scatter is SLOW

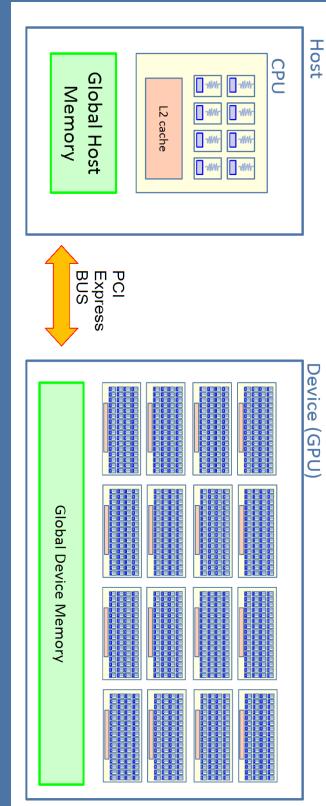
pointers to
indices

PCI BUS latency
& work item
dispatch is SLOW



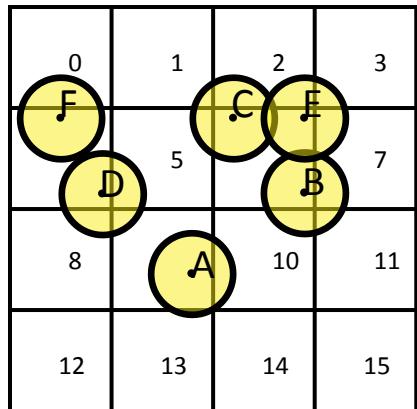
GPU Pair
Detection

CUDA
Constraint
Solver



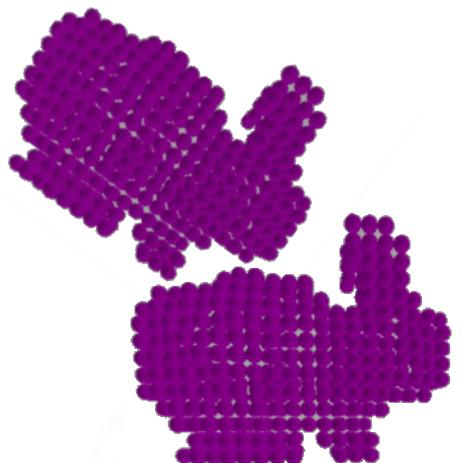
1st Rewrite: GPU Rigid Body Pipeline (~2008)

Detect Contact pairs



Uniform grid

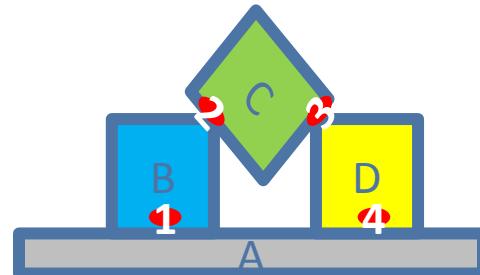
Compute contact points



Spherical Voxelization

Setup Contact constraints

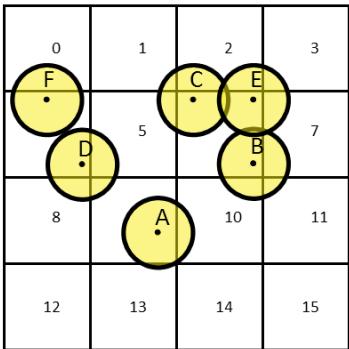
Solve constraint s



A	B	C	D
1	1	3	3
4	2	2	4

CPU batch and GPU solve
(Slow PCI BUS dispatch)

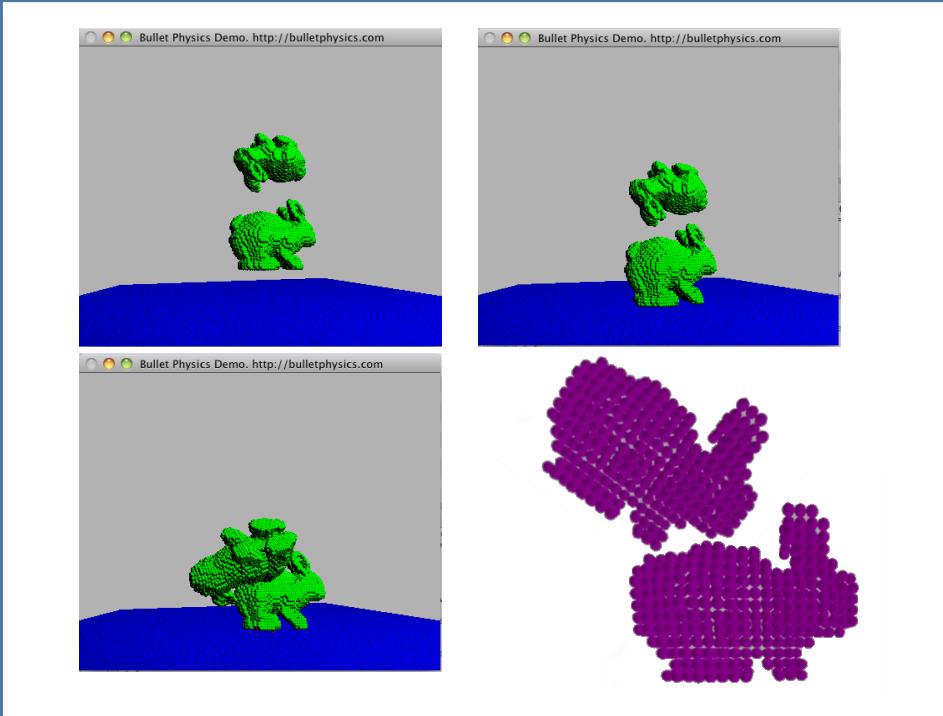
Parallel Primitives: Radix Sort, Prefix Scan



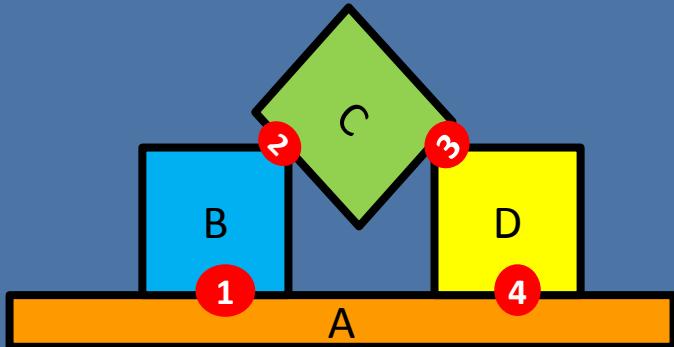
Cell Index	Cell Start
0	
1	
2	
3	
4	0
5	
6	2
7	
8	
9	5
10	
11	
12	
13	
14	
15	

Array Index	Unsorted Cell ID, Particle ID	Sorted Cell ID Particle ID
0	9, A	4, D
1	6, B	4, F
2	6, C	6, B
3	4, D	6, C
4	6, E	6, E
5	4, F	9, A

GPUs love Fine Grain Parallelism



Reordering Constraints



A	B	C	D
1	1		
	2	2	
		3	3
4			4



	A	B	C	D
Batch 0	1	1	3	3
Batch 1	4	2	2	4

CPU sequential batch creation

```
while( nIdxSrc ) {
    nIdxDst = 0;      int nCurrentBatch = 0;
    for(int i=0; i<N_FLG/32; i++) flg[i] = 0; //clear flag
    for(int i=0; i<nIdxSrc; i++) {
        int idx = idxSrc[i];  btAssert( idx < n );
        //check if it can go
        int aIdx = cs[idx].m_bodyAPtr & FLG_MASK;    int bIdx = cs[idx].m_bodyBPtr & FLG_MASK;
        u32 aUnavailable = flg[ aIdx/32 ] & (1<<(aIdx&31));u32 bUnavailable = flg[ bIdx/32 ] & (1<<(bIdx&31));
        if( aUnavailable==0 && bUnavailable==0 )  {
            flg[ aIdx/32 ] |= (1<<(aIdx&31));    flg[ bIdx/32 ] |= (1<<(bIdx&31));
            cs[idx].getBatchIdx() = batchIdx;
            sortData[idx].m_key = batchIdx; sortData[idx].m_value = idx;
            nCurrentBatch++;
            if( nCurrentBatch == simdWidth ) {
                nCurrentBatch = 0;
                for(int i=0; i<N_FLG/32; i++) flg[i] = 0;
            }
        }
        else  {
            idxDst[nIdxDst++] = idx;
        }
    }
    swap2( idxSrc, idxDst ); swap2( nIdxSrc, nIdxDst );
    batchIdx++;
}
```

2st Rewrite: GPU Rigid Body Pipeline (~2012)

Detect Contact pairs

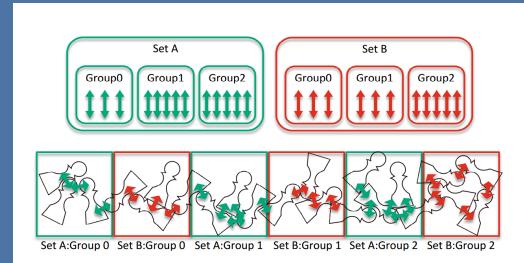
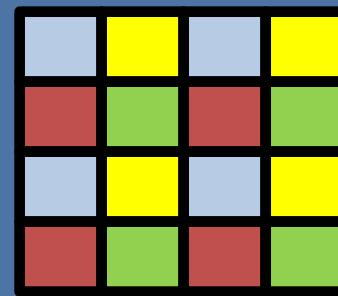
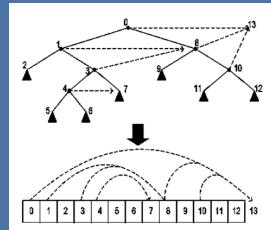
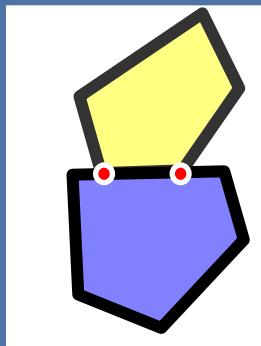
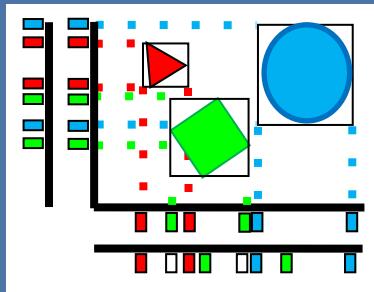
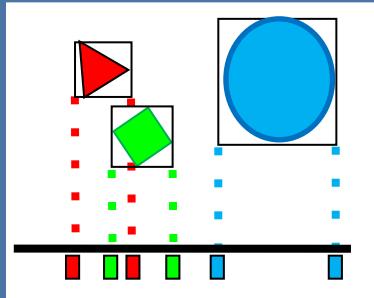


Compute contact points

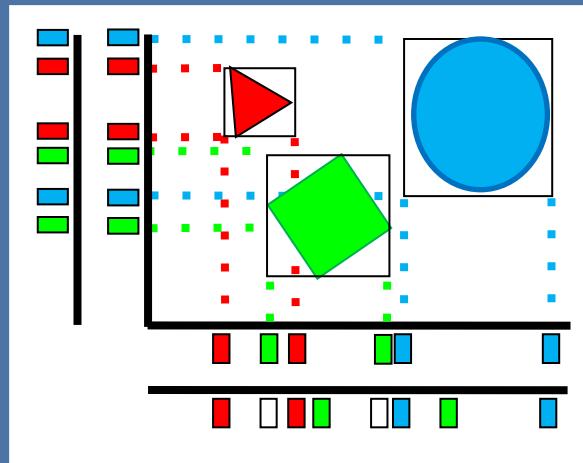
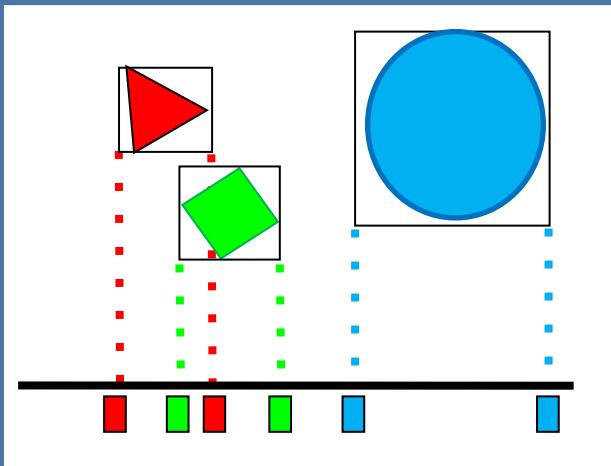


Setup Contact constraints

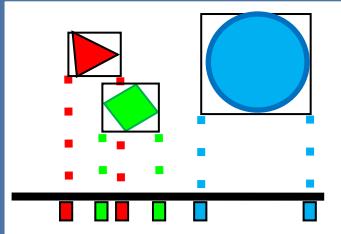
Solve constraint s



GPU Pair Search

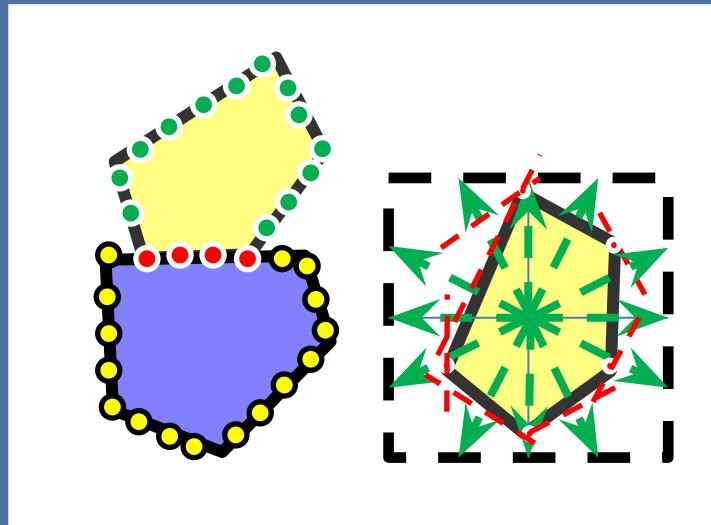


Optimizations in GPU 1-axis SAP



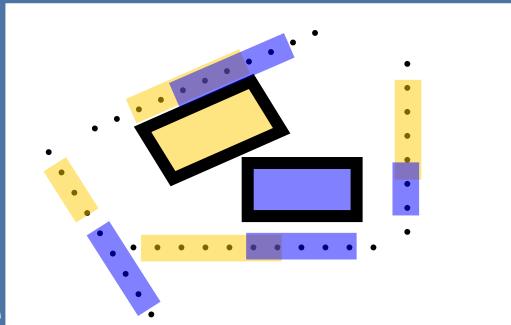
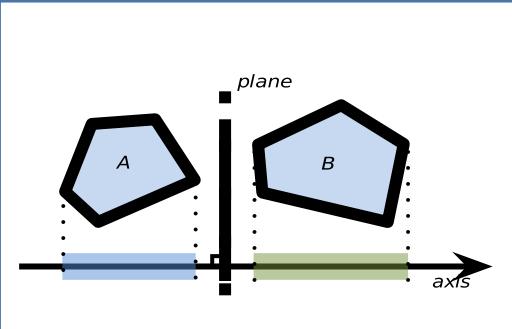
- Local memory
 - blocks to fetch AABBs and re-use them within a workgroup (requires a barrier)
 - Reduce global atomic operations
 - Private memory to accumulate overlapping pairs (append buffer)
 - Local atomics
 - Determine early exit condition for all work items within a workgroup
 - Load balancing
 - One work item per object, multiple work items for large objects
- See `opencl/gpu_broadphase/kernels/sapFast.cl` and `sap.cl`
(contains un-optimized and optimized version of the kernel for comparison)

Cubemap: Convex Heightfield Contact

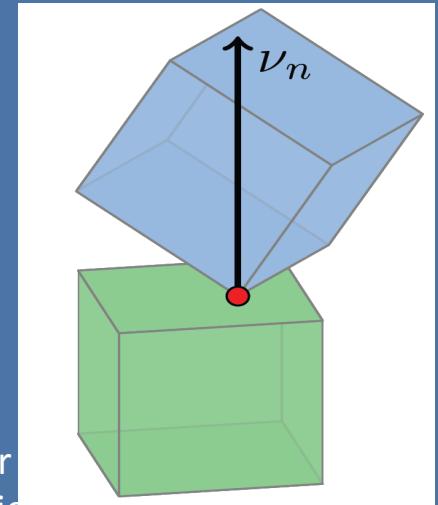


GPU Separating Axis Test

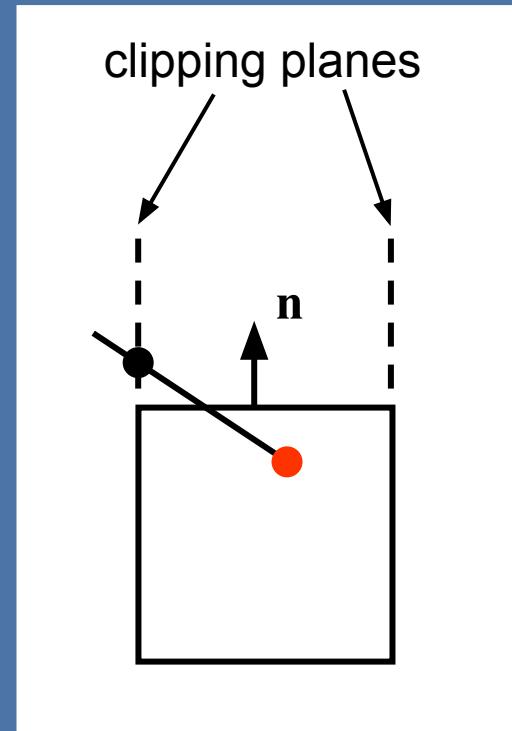
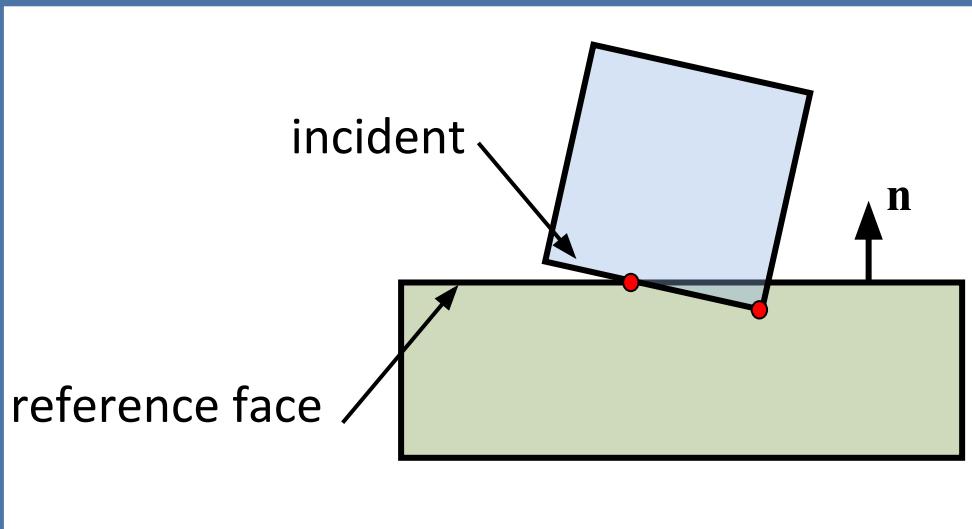
- Face normal A
- Face normal B
- Edge-edge normal



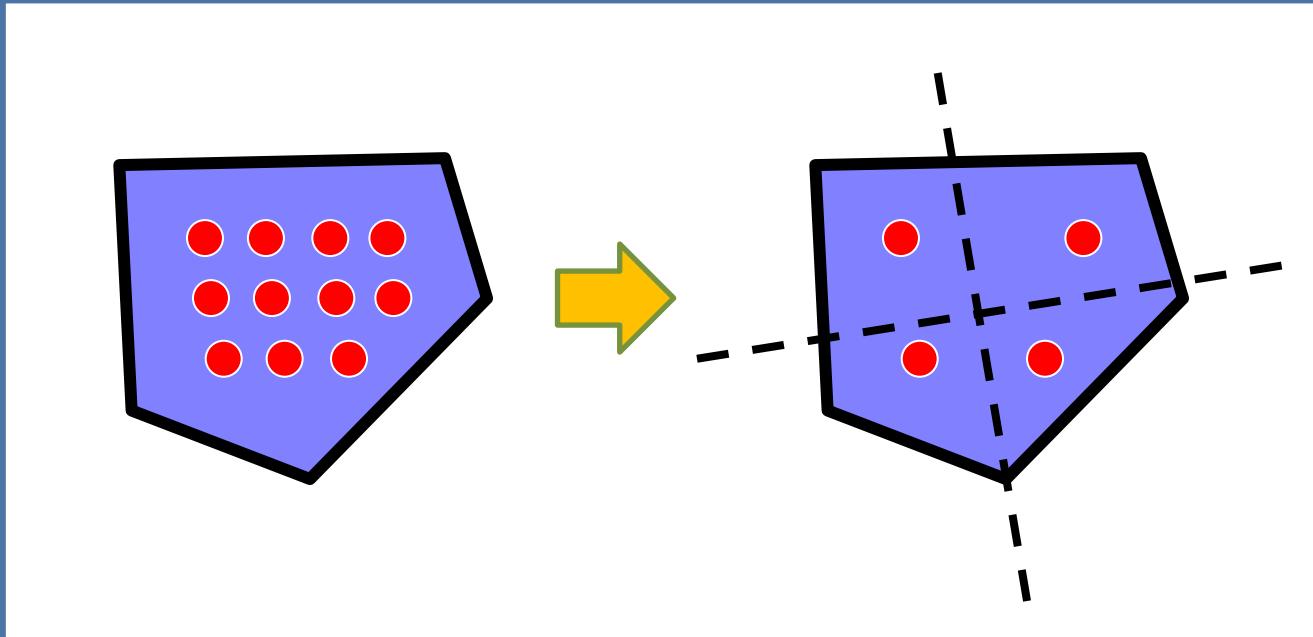
- Use one axis per dimension, resulting in one axis per dimension for each shape.
- Precise solution and faster than height field approximation for low-resolution convex shapes
- See [opencl/gpu_sat/kernels/sat.cl](#)



GPU Sutherland Hodgeman Clipping



GPU Contact Reduction



SAT Convex Collision Refactor

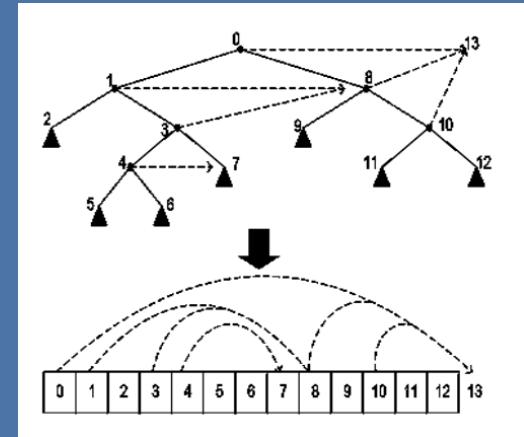
Find Potential Overlapping Pairs
For each convex-convex pair
 findSeparatingAxis
 For each SAT pass
 findClippingFaces
 clipHullHull
 contactReduction
For each compound pair
 Search potential overlapping children
 For each pair of overlapping children
 findSeparatingAxis
 For each SAT pass
 findClippingFaces
 clipHullHull
 contactReduction
For each concave pair
 For each triangle-convex overlap
 findSeparatingAxis
 findClippingFaces
 clipHullHull, contactReduction



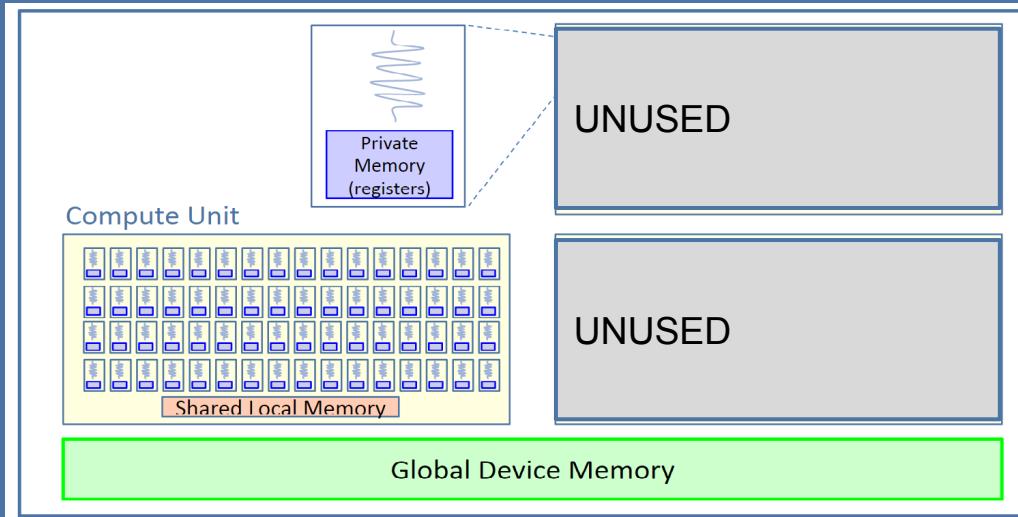
Find Potential Overlapping Pairs
For each compound pair
 Find Overlapping Compound Pairs
For each concave pair
 Find Overlapping tri-convex Pairs
For each Pair
 findSeparatingAxisKernel
 findConcaveSeparatingAxisKernel
For each SAT pass
 clipHullHullKernel
 clipCompoundsHullHullKernel
 clipHullHullConcaveConvexKernel
 findClippingFacesKernel
 clipFacesAndContactReductionKernel
 newContactReductionKernel

GPU trimesh and compound shapes

- Create skip indices for faster traversal
- Create subtrees that fit in Local Memory
- Stream subtrees for entire wavefront/warp
- Quantize Nodes
 - 16 bytes/node

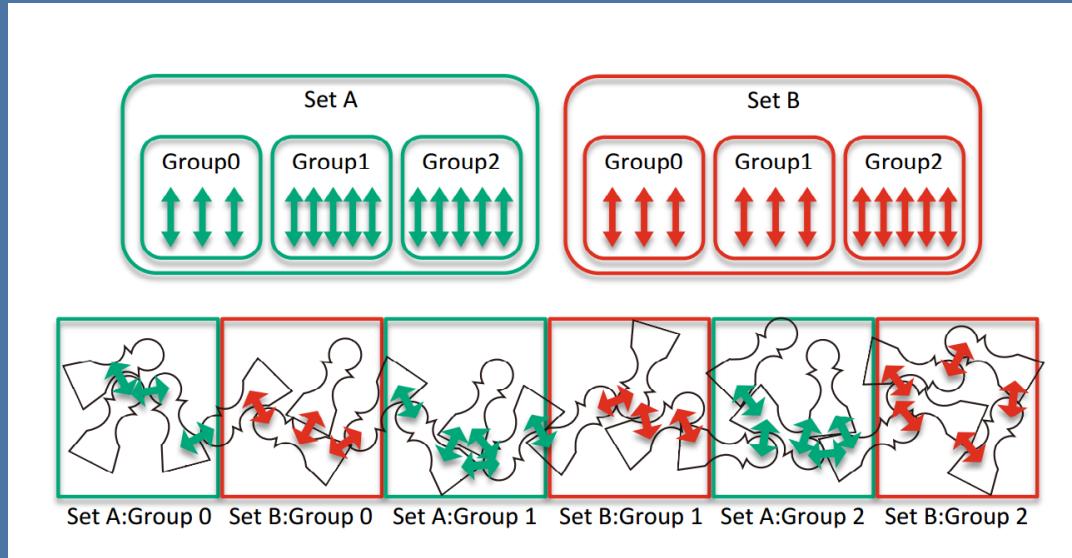


Naïve GPU batch creation



- Use a single Compute Unit
- All threads in the Compute Unit synchronize the locking of bodies using atomics and barriers
- Didn't scale well for larger scale simulations ($>\sim 30k$)

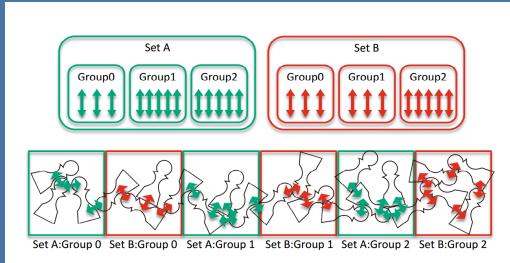
2-way Batch Generation



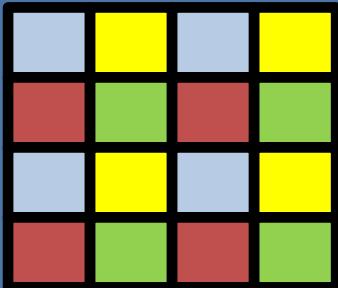
See "A parallel constraint solver for a rigid body simulation", Takahiro Harada, <http://dl.acm.org/citation.cfm?id=2077378.2077406>
Source code at `opencl\gpu_rigidbody\kernels\solveContact.cl` and other `solve*.cl`

2-way Batch Generation

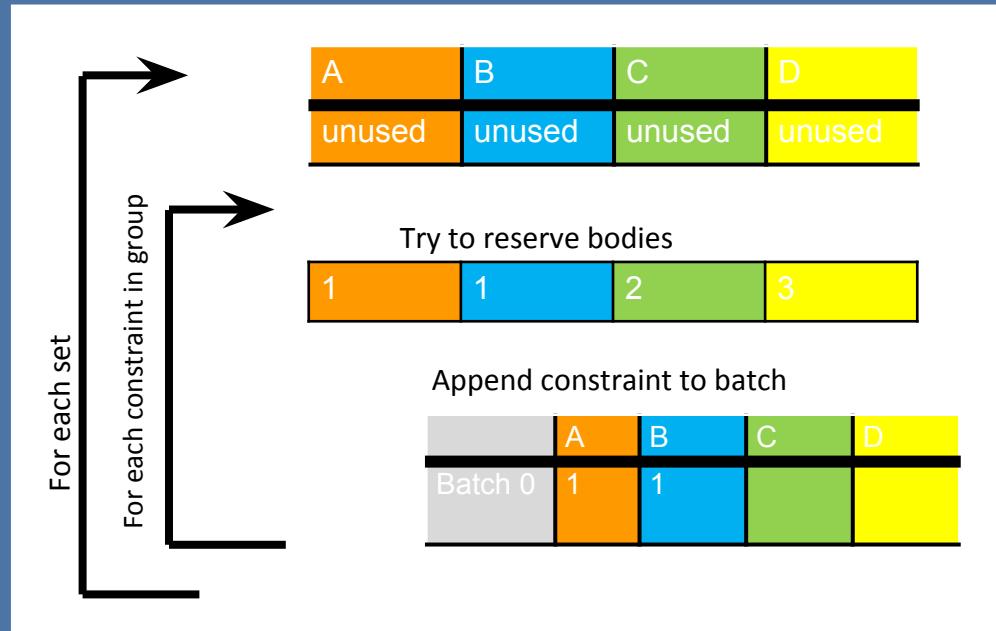
1: Create Disjoint Sets



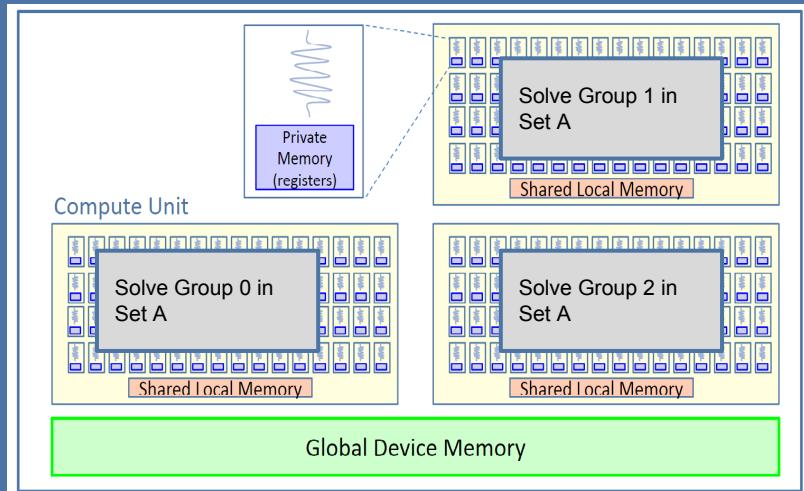
1: Spatial Sort to create Disjoint Sets, or alternative method



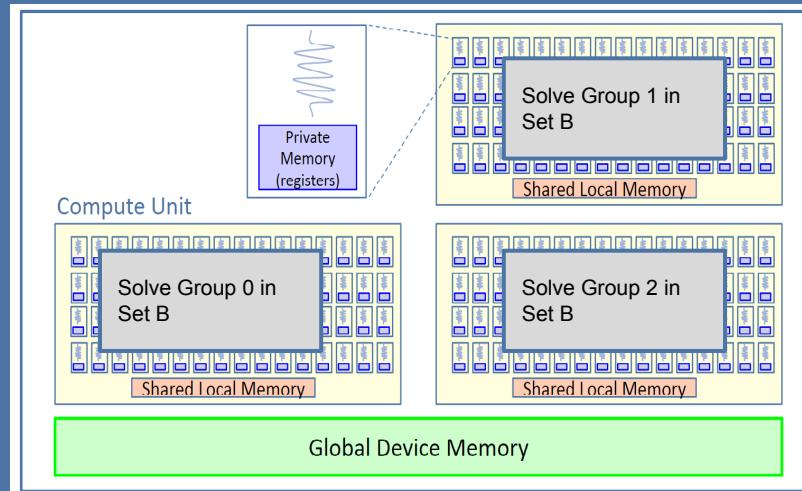
2: For each set, Compute Units solve groups in parallel



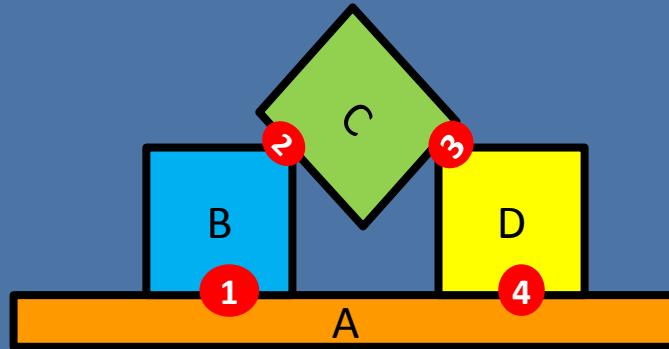
2-way Batch Generation



Host
SYNC
(enqueue
work)

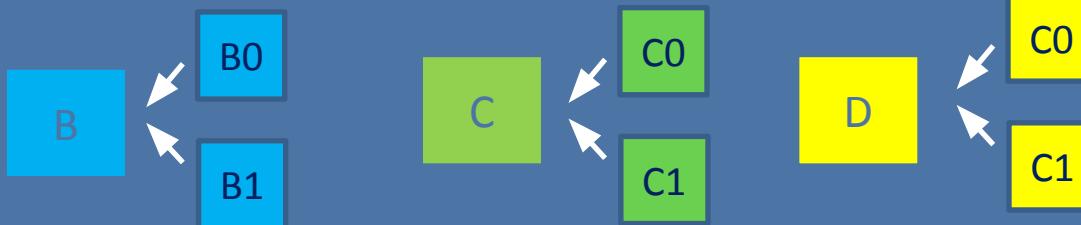


Mass Splitting+Jacobi \cong PGS



A	B0	B1	C0	C1	D1	D1	A
1	1	2	2	3	3	4	4

Parallel Jacobi



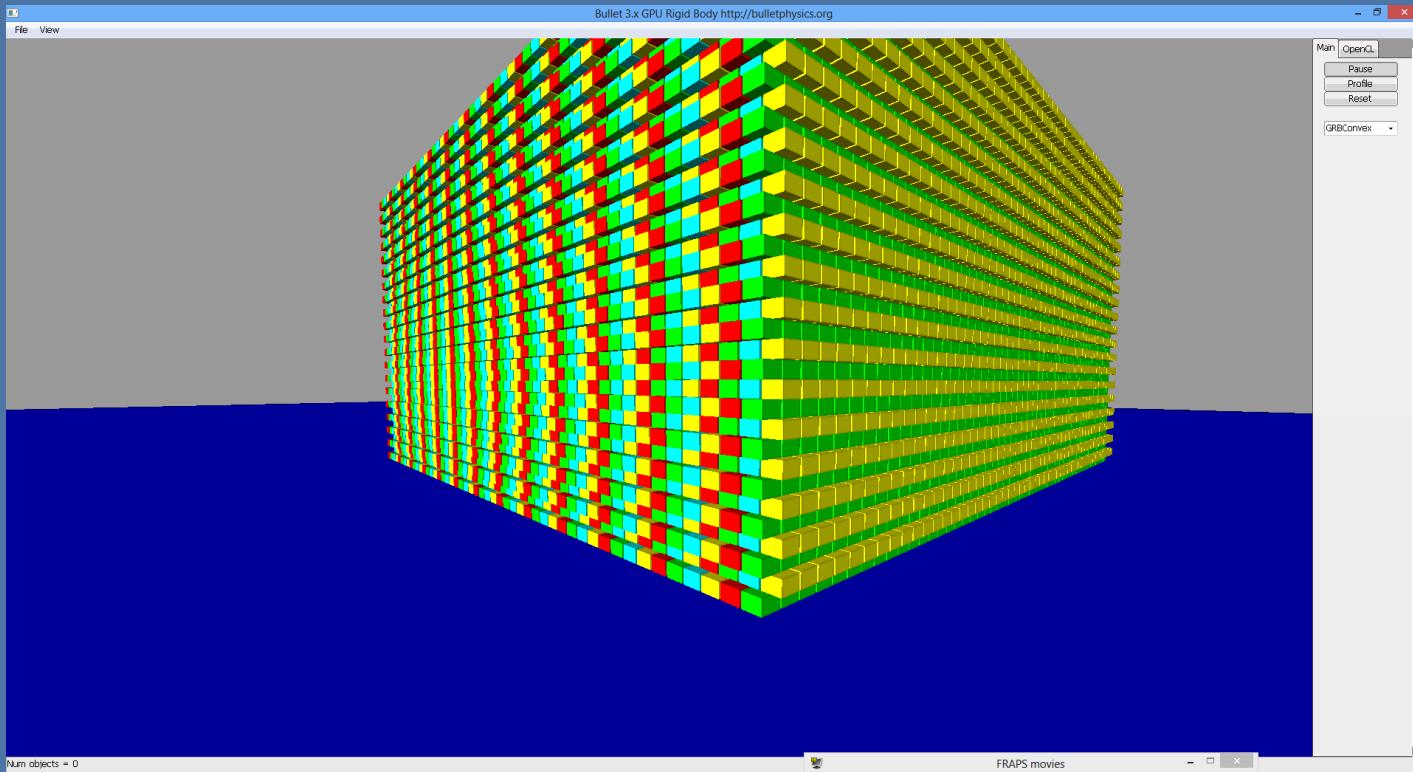
Averaging velocities

- See “Mass Splitting for Jitter-Free Parallel Rigid Body Simulation” by Tonge et. al.

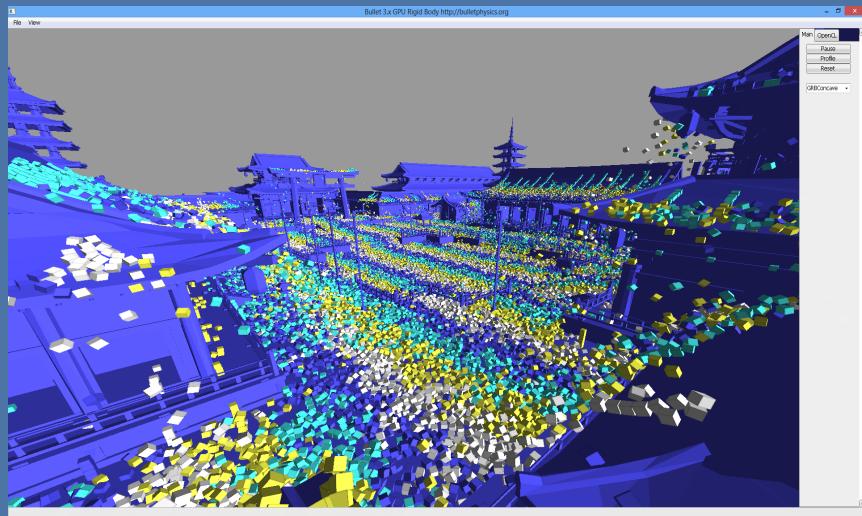
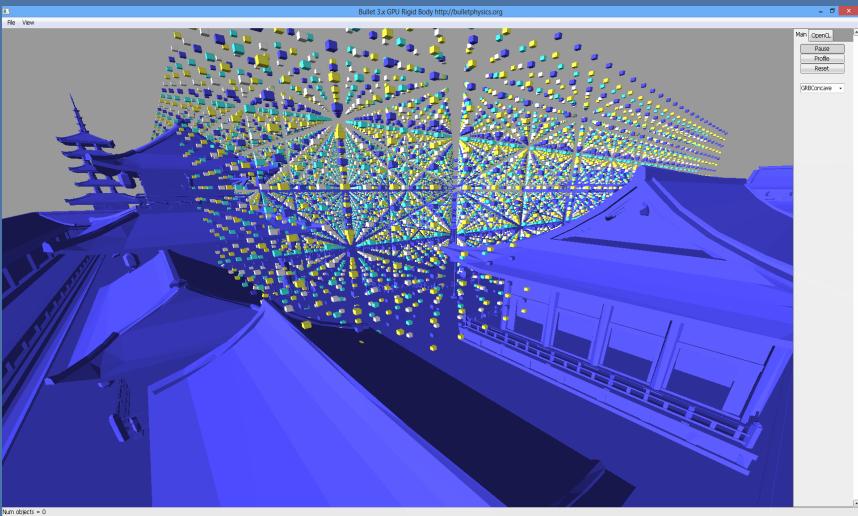
Over 50 OpenCL™ kernels

AddOffsetKernel	AverageVelocitiesKernel	BatchSolveKernelContact	BatchSolveKernelFriction	ClearVelocitiesKernel	ContactToConstraintKernel
ContactToConstraintSplitKernel	CopyConstraintKernel	CountBodiesKernel	CreateBatches	CreateBatchesNew	FillFloatKernel
FillInt2Kernel	FillIntKernel	FillUnsignedIntKernel	LocalScanKernel	PrefixScanKernel	ReorderContactKernel
SearchSortDataLowerKernel	SearchSortDataUpperKernel	SetSortDataKernel	SolveContactJacobiKernel	SolveFrictionJacobiKernel	SortAndScatterKernel
SortAndScatterSortDataKernel	StreamCountKernel	StreamCountSortDataKernel	SubtractKernel	TopLevelScanKernel	UpdateBodyVelocitiesKernel
bvhTraversalKernel	clipCompoundsHullHullKernel	clipFacesAndContactReductionKernel	clipHullHullConcaveConvexKernel	clipHullHullKernel	computePairsKernel
computePairsKernelTwoArrays	copyAabbsKernel	copyTransformsToVBOKernel	extractManifoldAndAddContactKernel	findClippingFacesKernel	findCompoundPairsKernel
findConcaveSeparatingAxisKernel	findSeparatingAxisKernel	flipFloatKernel	initializeGpuAabbsFull	integrateTransformsKernel	newContactReductionKernel
processCompoundPairsKernel	scatterKernel				

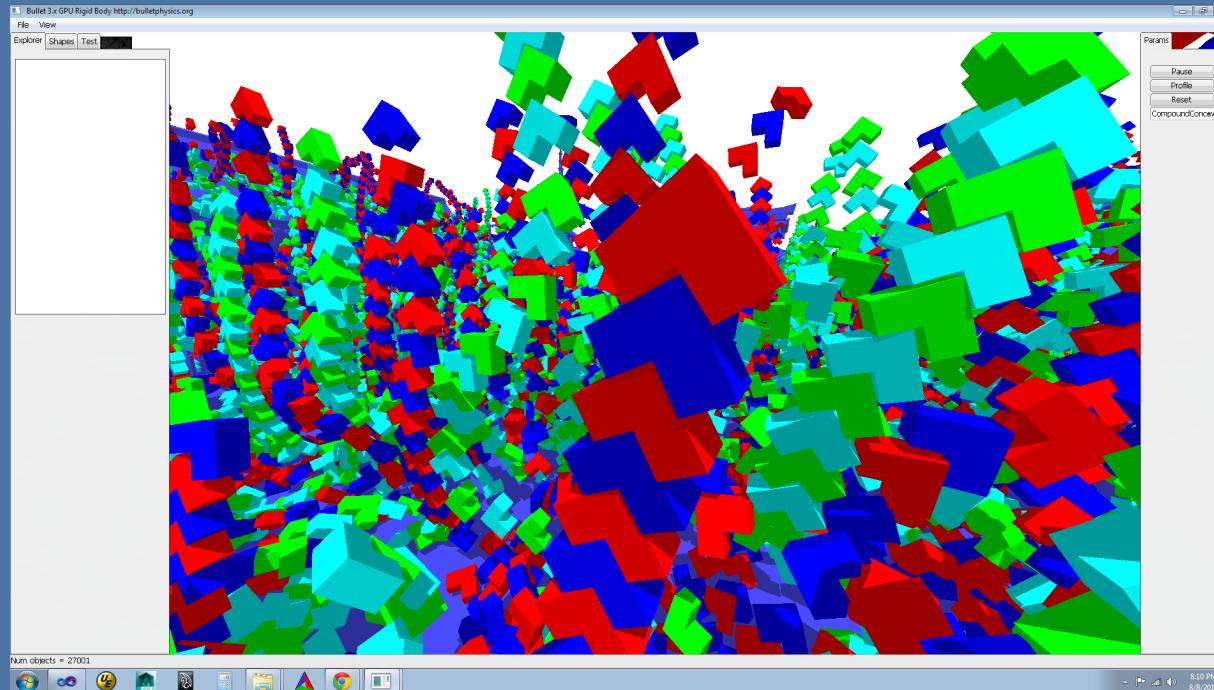
100k Box Stack at 60FPS



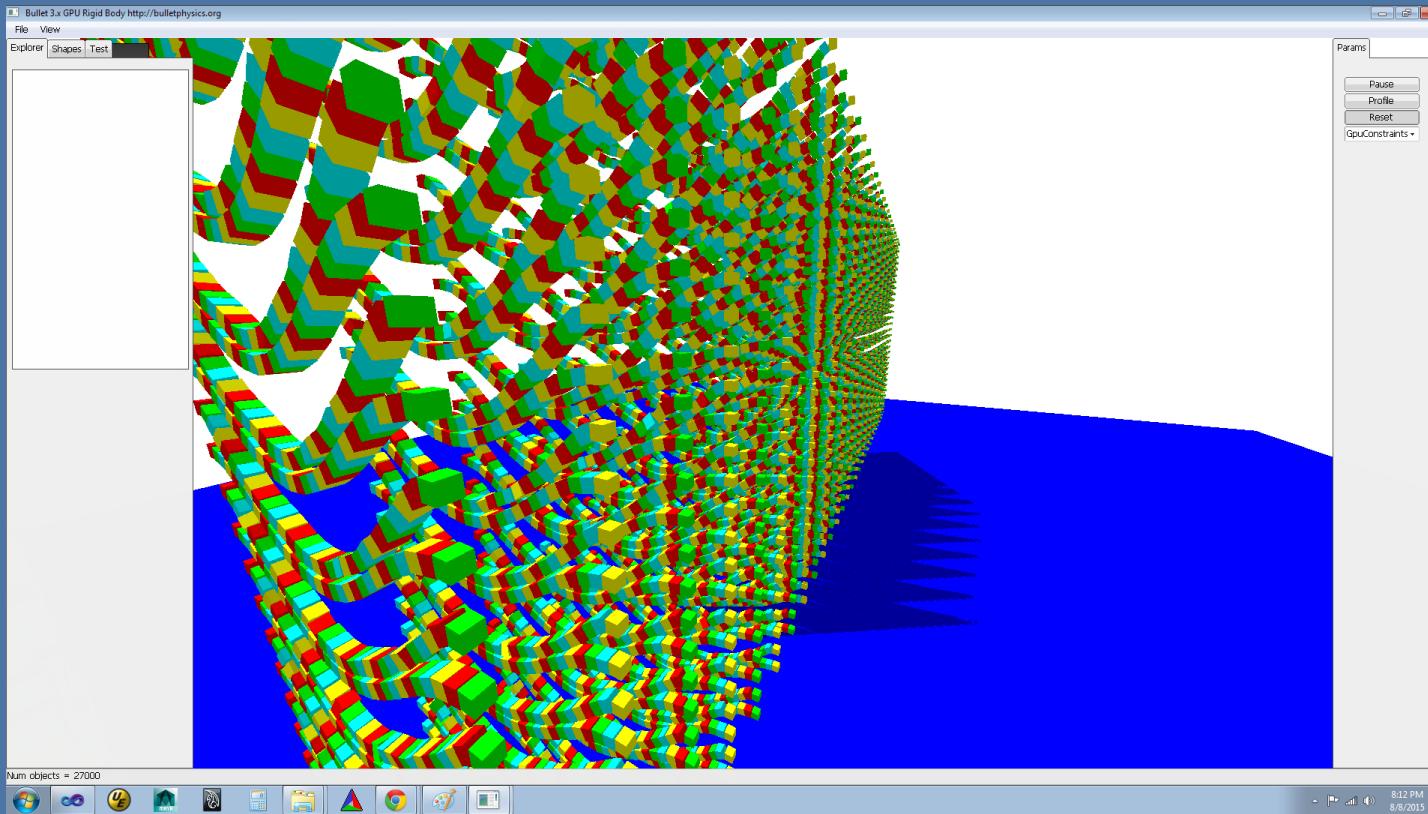
Concave Static Triangle Mesh



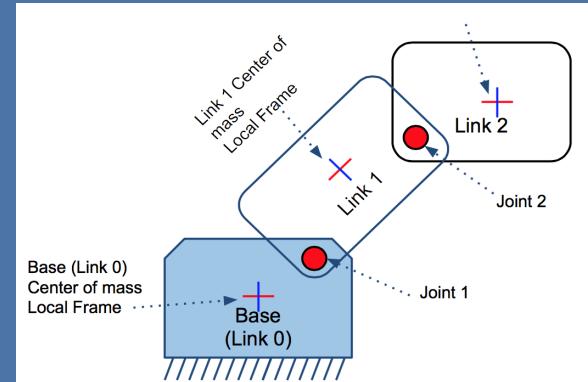
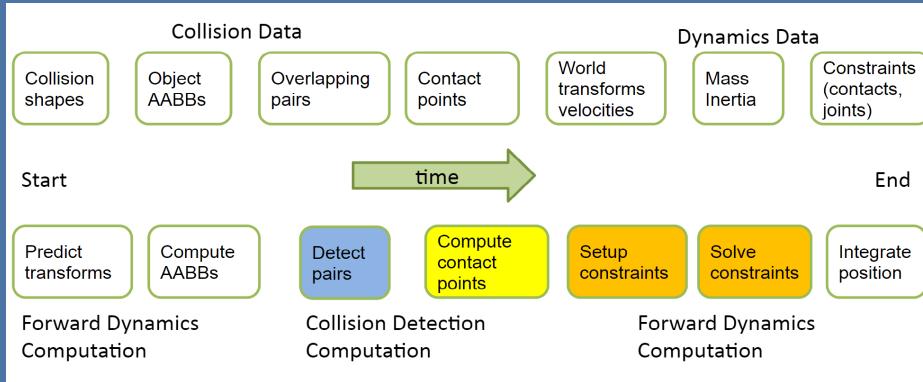
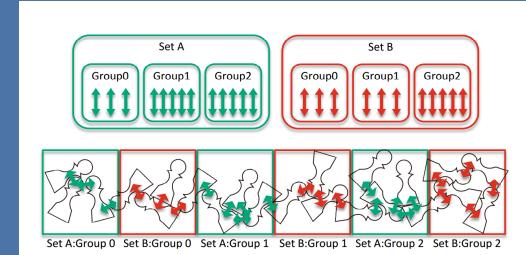
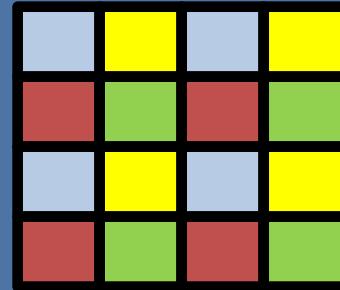
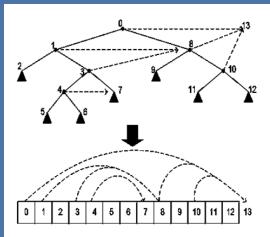
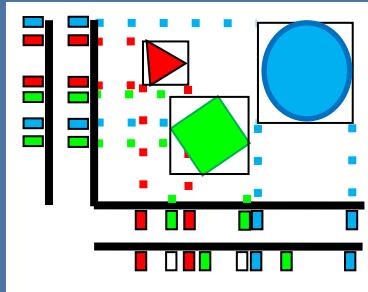
GPU Compound Collision



GPU Hinge Joints



Ongoing: uniting API and code bases



Thank you! Questions?

See <https://github.com/bulletphysics>

ExampleBrowser --enable_experimental_opencl

https://github.com/bulletphysics/bullet3/tree/old_demos
App_Bullet3_OpenCL_Demos_clew



MULTITHREADING FOR VISUAL EFFECTS

Martin Watt • Erwin Coumans • George ElKoura • Ronald Henderson
Manuel Kraemer • Jeff Lait • James Reinders

CRC
Taylor & Francis Group

