

# Tutorial: Deep Reinforcement Learning

David Silver, Google DeepMind

# Outline

Introduction to Deep Learning

Introduction to Reinforcement Learning

Value-Based Deep RL

Policy-Based Deep RL

Model-Based Deep RL

# Reinforcement Learning in a nutshell

RL is a general-purpose framework for decision-making

- ▶ RL is for an **agent** with the capacity to **act**
- ▶ Each **action** influences the agent's future **state**
- ▶ Success is measured by a scalar **reward** signal
- ▶ Goal: **select actions to maximise future reward**

# Deep Learning in a nutshell

DL is a general-purpose framework for representation learning

- ▶ Given an **objective**
- ▶ Learn **representation** that is required to achieve objective
- ▶ Directly from **raw inputs**
- ▶ Using minimal domain knowledge

# Deep Reinforcement Learning: $AI = RL + DL$

We seek a single agent which can solve any human-level task

- ▶ RL defines the objective
- ▶ DL gives the mechanism
- ▶  $RL + DL = \text{general intelligence}$

## Examples of Deep RL @DeepMind

- ▶ **Play** games: Atari, poker, Go, ...
- ▶ **Navigate** worlds: 3D worlds, Labyrinth, ...
- ▶ **Control** physical systems: manipulate, walk, swim, ...
- ▶ **Interact** with users: recommend, optimise, personalised, ...

# Outline

Introduction to Deep Learning

Introduction to Reinforcement Learning

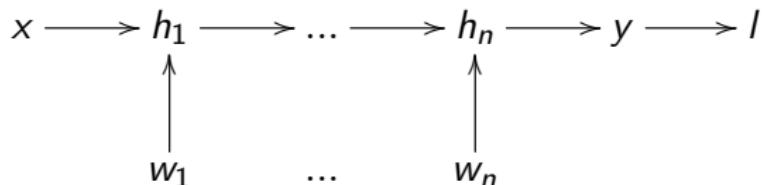
Value-Based Deep RL

Policy-Based Deep RL

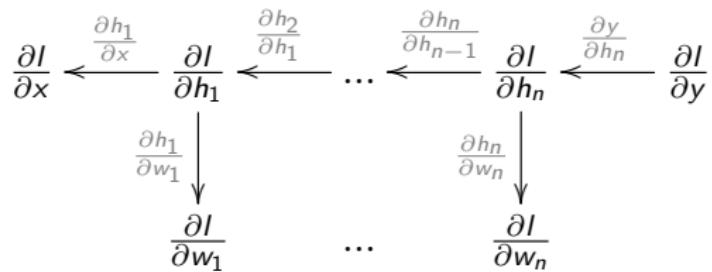
Model-Based Deep RL

# Deep Representations

- A **deep representation** is a composition of many functions



- Its gradient can be **backpropagated** by the chain rule



# Deep Neural Network

A **deep neural network** is typically composed of:

- ▶ Linear transformations

$$h_{k+1} = Wh_k$$

- ▶ Non-linear activation functions

$$h_{k+2} = f(h_{k+1})$$

- ▶ A loss function on the output, e.g.

- ▶ Mean-squared error  $l = ||y^* - y||^2$
- ▶ Log likelihood  $l = \log \mathbb{P}[y^*]$

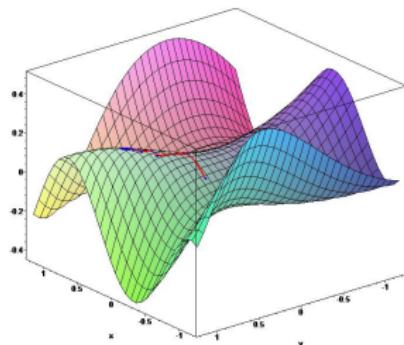
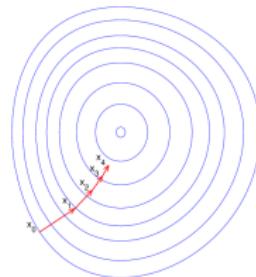
# Training Neural Networks by Stochastic Gradient Descent

- ▶ Sample gradient of expected loss  $L(\mathbf{w}) = \mathbb{E}[l]$

$$\frac{\partial l}{\partial \mathbf{w}} \sim \mathbb{E} \left[ \frac{\partial l}{\partial \mathbf{w}} \right] = \frac{\partial L(\mathbf{w})}{\partial \mathbf{w}}$$

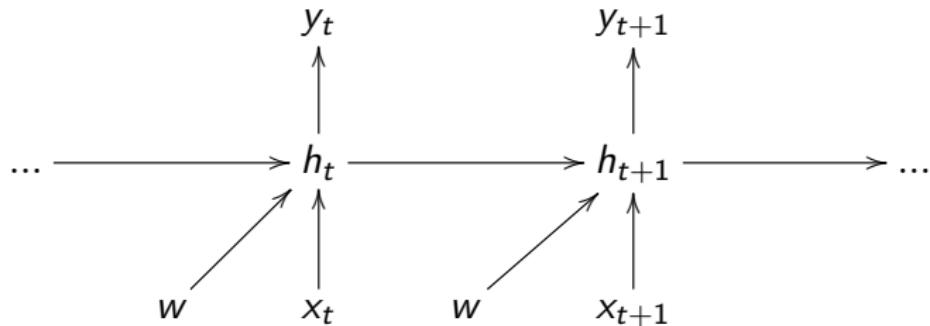
- ▶ Adjust  $\mathbf{w}$  down the sampled gradient

$$\Delta \mathbf{w} \propto \frac{\partial l}{\partial \mathbf{w}}$$

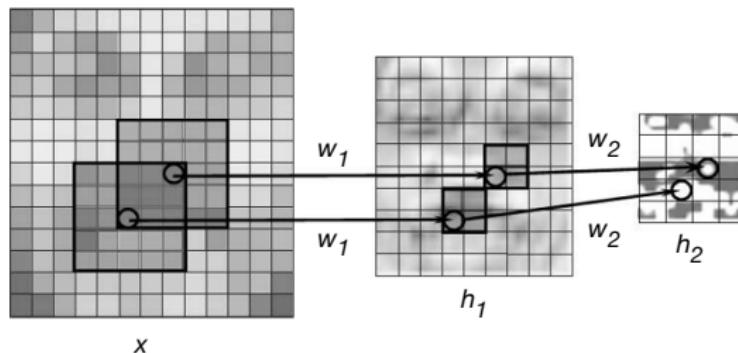


# Weight Sharing

Recurrent neural network shares weights between time-steps



Convolutional neural network shares weights between local regions



# Outline

Introduction to Deep Learning

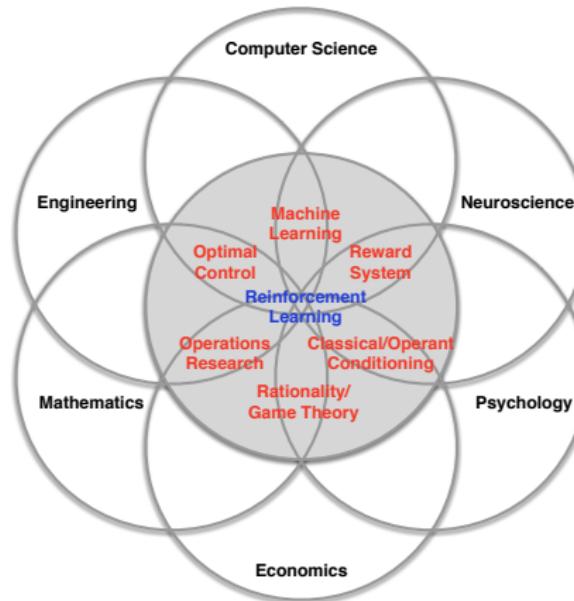
Introduction to Reinforcement Learning

Value-Based Deep RL

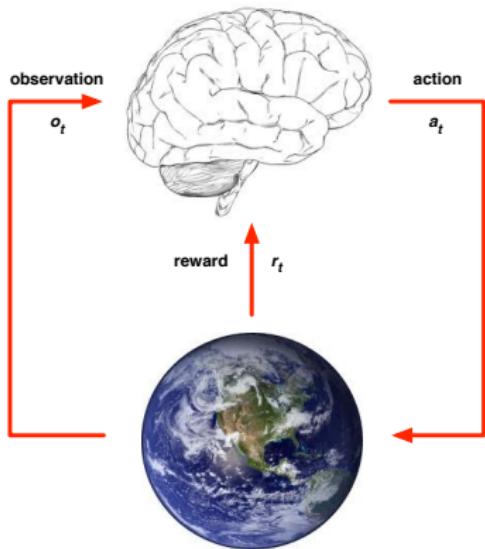
Policy-Based Deep RL

Model-Based Deep RL

# Many Faces of Reinforcement Learning



# Agent and Environment



- ▶ At each step  $t$  the agent:
  - ▶ Executes action  $a_t$
  - ▶ Receives observation  $o_t$
  - ▶ Receives scalar reward  $r_t$
- ▶ The environment:
  - ▶ Receives action  $a_t$
  - ▶ Emits observation  $o_{t+1}$
  - ▶ Emits scalar reward  $r_{t+1}$

# State



- ▶ Experience is a sequence of observations, actions, rewards

$$o_1, r_1, a_1, \dots, a_{t-1}, o_t, r_t$$

- ▶ The **state** is a summary of experience

$$s_t = f(o_1, r_1, a_1, \dots, a_{t-1}, o_t, r_t)$$

- ▶ In a fully observed environment

$$s_t = f(o_t)$$

# Major Components of an RL Agent

- ▶ An RL agent may include one or more of these components:
  - ▶ **Policy**: agent's behaviour function
  - ▶ **Value function**: how good is each state and/or action
  - ▶ **Model**: agent's representation of the environment

# Policy

- ▶ A **policy** is the agent's behaviour
- ▶ It is a map from state to action:
  - ▶ Deterministic policy:  $a = \pi(s)$
  - ▶ Stochastic policy:  $\pi(a|s) = \mathbb{P}[a|s]$

# Value Function

- ▶ A **value function** is a prediction of future reward
  - ▶ “How much reward will I get from action  $a$  in state  $s$ ?”
- ▶  **$Q$ -value function** gives expected total reward
  - ▶ from state  $s$  and action  $a$
  - ▶ under policy  $\pi$
  - ▶ with discount factor  $\gamma$

$$Q^\pi(a|s) = \mathbb{E} [r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots | s, a]$$

# Value Function

- ▶ A **value function** is a prediction of future reward
  - ▶ “How much reward will I get from action  $a$  in state  $s$ ?”
- ▶  **$Q$ -value function** gives expected total reward
  - ▶ from state  $s$  and action  $a$
  - ▶ under policy  $\pi$
  - ▶ with discount factor  $\gamma$

$$Q^\pi(a|s) = \mathbb{E} [r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots | s, a]$$

- ▶ Value functions decompose into a Bellman equation

$$Q^\pi(a|s) = \mathbb{E}_{s',a'} [r + \gamma Q^\pi(a'|s') | s, a]$$

# Optimal Value Functions

- ▶ An optimal value function is the maximum achievable value

$$Q^*(s, a) = \max_{\pi} Q^{\pi}(s, a) = Q^{\pi^*}(s, a)$$

## Optimal Value Functions

- ▶ An optimal value function is the maximum achievable value

$$Q^*(s, a) = \max_{\pi} Q^{\pi}(s, a) = Q^{\pi^*}(s, a)$$

- ▶ Once we have  $Q^*$  we can act optimally,

$$\pi^*(s) = \operatorname{argmax}_a Q^*(s, a)$$

# Optimal Value Functions

- ▶ An optimal value function is the maximum achievable value

$$Q^*(s, a) = \max_{\pi} Q^{\pi}(s, a) = Q^{\pi^*}(s, a)$$

- ▶ Once we have  $Q^*$  we can act optimally,

$$\pi^*(s) = \operatorname{argmax}_a Q^*(s, a)$$

- ▶ Optimal value maximises over all decisions. Informally:

$$\begin{aligned} Q^*(s, a) &= r_{t+1} + \gamma \max_{a_{t+1}} r_{t+2} + \gamma^2 \max_{a_{t+2}} r_{t+3} + \dots \\ &= r_{t+1} + \gamma \max_{a_{t+1}} Q^*(s_{t+1}, a_{t+1}) \end{aligned}$$

# Optimal Value Functions

- ▶ An optimal value function is the maximum achievable value

$$Q^*(s, a) = \max_{\pi} Q^{\pi}(s, a) = Q^{\pi^*}(s, a)$$

- ▶ Once we have  $Q^*$  we can act optimally,

$$\pi^*(s) = \operatorname{argmax}_a Q^*(s, a)$$

- ▶ Optimal value maximises over all decisions. Informally:

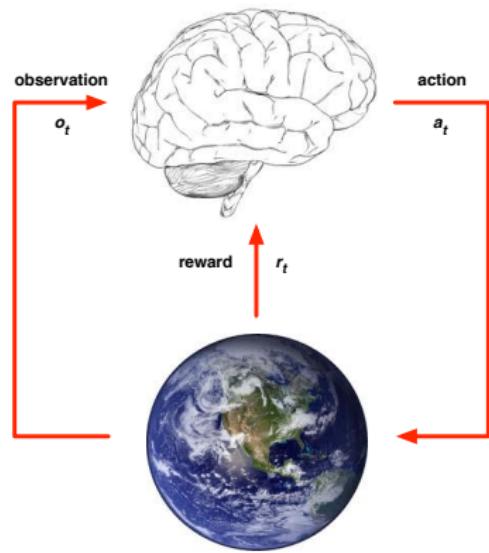
$$\begin{aligned} Q^*(s, a) &= r_{t+1} + \gamma \max_{a_{t+1}} r_{t+2} + \gamma^2 \max_{a_{t+2}} r_{t+3} + \dots \\ &= r_{t+1} + \gamma \max_{a_{t+1}} Q^*(s_{t+1}, a_{t+1}) \end{aligned}$$

- ▶ Formally, optimal values decompose into a Bellman equation

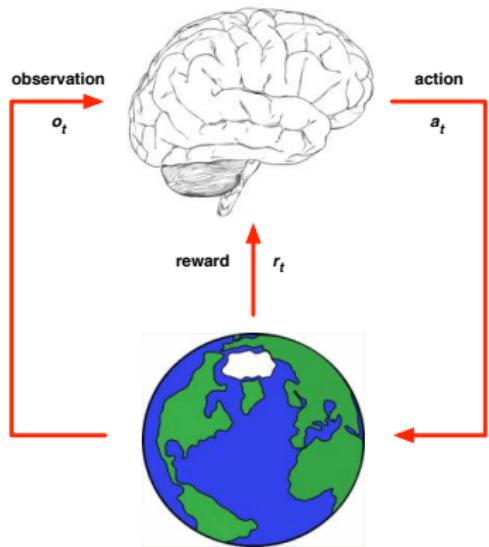
$$Q^*(s, a) = \mathbb{E}_{s'} \left[ r + \gamma \max_{a'} Q^*(s', a') \mid s, a \right]$$

# Value Function Demo

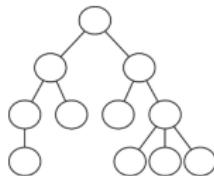
# Model



# Model



- ▶ **Model** is learnt from experience
- ▶ Acts as proxy for environment
- ▶ Planner interacts with model
- ▶ e.g. using lookahead search



# Approaches To Reinforcement Learning

## Value-based RL

- ▶ Estimate the **optimal value function**  $Q^*(s, a)$
- ▶ This is the maximum value achievable under any policy

## Policy-based RL

- ▶ Search directly for the **optimal policy**  $\pi^*$
- ▶ This is the policy achieving maximum future reward

## Model-based RL

- ▶ Build a model of the environment
- ▶ Plan (e.g. by lookahead) using model

# Deep Reinforcement Learning

- ▶ Use deep neural networks to represent
  - ▶ Value function
  - ▶ Policy
  - ▶ Model
- ▶ Optimise loss function by stochastic gradient descent

# Outline

Introduction to Deep Learning

Introduction to Reinforcement Learning

Value-Based Deep RL

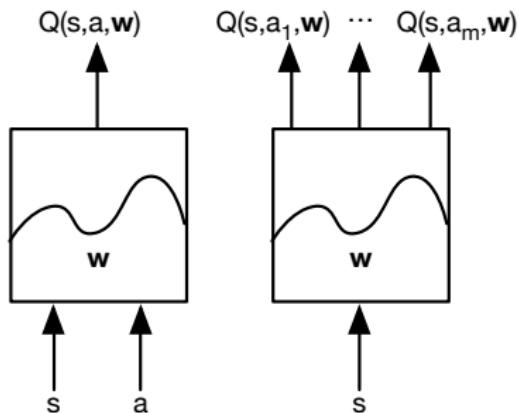
Policy-Based Deep RL

Model-Based Deep RL

# Q-Networks

Represent value function by **Q-network** with weights  $\mathbf{w}$

$$Q(s, a, \mathbf{w}) \approx Q^*(s, a)$$



# Q-Learning

- Optimal Q-values should obey Bellman equation

$$Q^*(s, a) = \mathbb{E}_{s'} \left[ r + \gamma \max_{a'} Q(s', a')^* \mid s, a \right]$$

- Treat right-hand side  $r + \gamma \max_{a'} Q(s', a', \mathbf{w})$  as a target
- Minimise MSE loss by stochastic gradient descent

$$l = \left( r + \gamma \max_a Q(s', a', \mathbf{w}) - Q(s, a, \mathbf{w}) \right)^2$$

# Q-Learning

- Optimal Q-values should obey Bellman equation

$$Q^*(s, a) = \mathbb{E}_{s'} \left[ r + \gamma \max_{a'} Q(s', a')^* \mid s, a \right]$$

- Treat right-hand side  $r + \gamma \max_{a'} Q(s', a', \mathbf{w})$  as a target
- Minimise MSE loss by stochastic gradient descent

$$l = \left( r + \gamma \max_a Q(s', a', \mathbf{w}) - Q(s, a, \mathbf{w}) \right)^2$$

- Converges to  $Q^*$  using table lookup representation

# Q-Learning

- ▶ Optimal Q-values should obey Bellman equation

$$Q^*(s, a) = \mathbb{E}_{s'} \left[ r + \gamma \max_{a'} Q(s', a')^* \mid s, a \right]$$

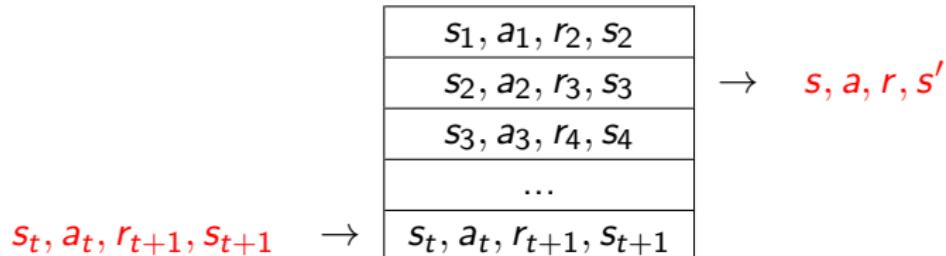
- ▶ Treat right-hand side  $r + \gamma \max_{a'} Q(s', a', \mathbf{w})$  as a target
- ▶ Minimise MSE loss by stochastic gradient descent

$$l = \left( r + \gamma \max_a Q(s', a', \mathbf{w}) - Q(s, a, \mathbf{w}) \right)^2$$

- ▶ Converges to  $Q^*$  using table lookup representation
- ▶ But **diverges** using neural networks due to:
  - ▶ Correlations between samples
  - ▶ Non-stationary targets

# Deep Q-Networks (DQN): Experience Replay

To remove correlations, build data-set from agent's own experience

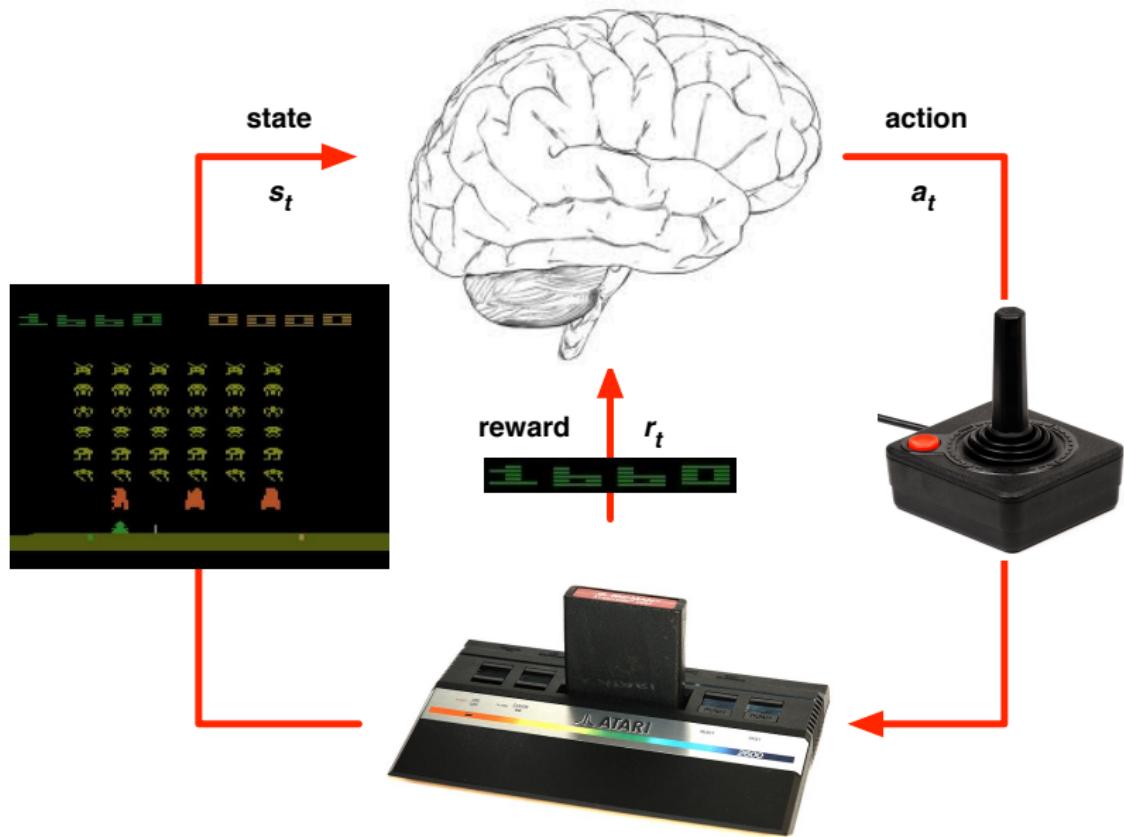


Sample experiences from data-set and apply update

$$l = \left( r + \gamma \max_{a'} Q(s', a', \mathbf{w}^-) - Q(s, a, \mathbf{w}) \right)^2$$

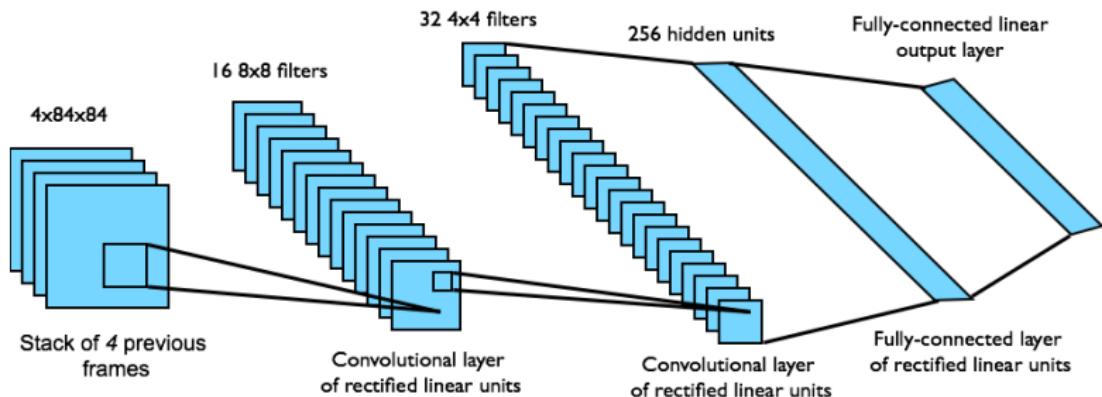
To deal with non-stationarity, target parameters  $\mathbf{w}^-$  are held fixed

# Deep Reinforcement Learning in Atari



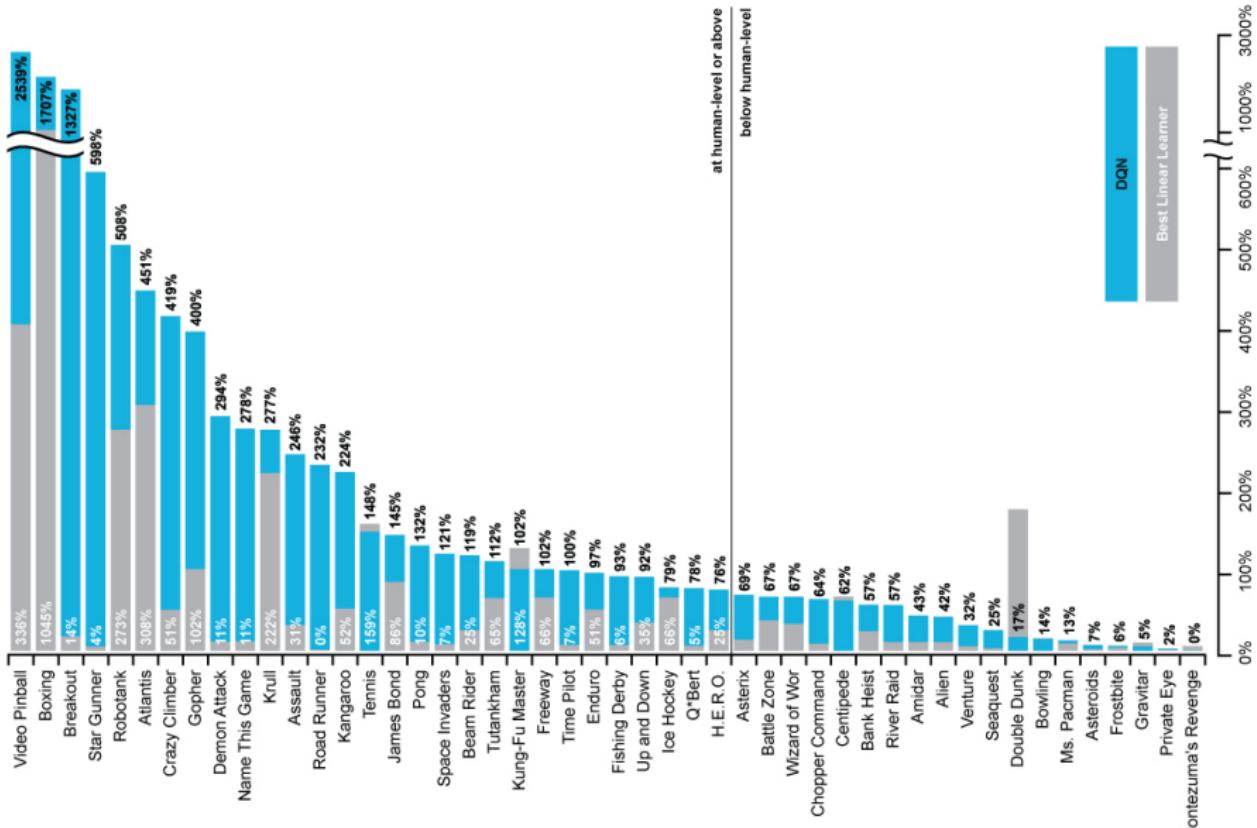
# DQN in Atari

- ▶ End-to-end learning of values  $Q(s, a)$  from pixels  $s$
- ▶ Input state  $s$  is stack of raw pixels from last 4 frames
- ▶ Output is  $Q(s, a)$  for 18 joystick/button positions
- ▶ Reward is change in score for that step



Network architecture and hyperparameters fixed across all games

# DQN Results in Atari



# DQN Atari Demo

DQN paper

[www.nature.com/articles/nature14236](http://www.nature.com/articles/nature14236)

DQN source code:

[sites.google.com/a/deepmind.com/dqn/](https://sites.google.com/a/deepmind.com/dqn/)



## Improvements since Nature DQN

- ▶ Double DQN: Remove upward bias caused by  $\max_a Q(s, a, \mathbf{w})$ 
  - ▶ Current Q-network  $\mathbf{w}$  is used to **select** actions
  - ▶ Older Q-network  $\mathbf{w}^-$  is used to **evaluate** actions

$$l = \left( r + \gamma \underset{a'}{\operatorname{argmax}} Q(s', a', \mathbf{w}, \mathbf{w}^-) - Q(s, a, \mathbf{w}) \right)^2$$

# Improvements since Nature DQN

- ▶ **Double DQN:** Remove upward bias caused by  $\max_a Q(s, a, \mathbf{w})$ 
  - ▶ Current Q-network  $\mathbf{w}$  is used to **select** actions
  - ▶ Older Q-network  $\mathbf{w}^-$  is used to **evaluate** actions

$$l = \left( r + \gamma \underset{a'}{\operatorname{argmax}} Q(s', a', \mathbf{w}), \mathbf{w}^- \right) - Q(s, a, \mathbf{w}) \right)^2$$

- ▶ **Prioritised replay:** Weight experience according to surprise
  - ▶ Store experience in priority queue according to DQN error

$$\left| r + \gamma \underset{a'}{\max} Q(s', a', \mathbf{w}^-) - Q(s, a, \mathbf{w}) \right|$$

# Improvements since Nature DQN

- ▶ **Double DQN:** Remove upward bias caused by  $\max_a Q(s, a, \mathbf{w})$ 
  - ▶ Current Q-network  $\mathbf{w}$  is used to **select** actions
  - ▶ Older Q-network  $\mathbf{w}^-$  is used to **evaluate** actions

$$l = \left( r + \gamma \operatorname{argmax}_{a'} Q(s', a', \mathbf{w}), \mathbf{w}^- \right) - Q(s, a, \mathbf{w}) \right)^2$$

- ▶ **Prioritised replay:** Weight experience according to surprise
  - ▶ Store experience in priority queue according to DQN error

$$\left| r + \gamma \max_{a'} Q(s', a', \mathbf{w}^-) - Q(s, a, \mathbf{w}) \right|$$

- ▶ **Duelling network:** Split Q-network into two channels
  - ▶ Action-independent **value function**  $V(s, v)$
  - ▶ Action-dependent **advantage function**  $A(s, a, \mathbf{w})$

$$Q(s, a) = V(s, v) + A(s, a, \mathbf{w})$$

# Improvements since Nature DQN

- ▶ **Double DQN:** Remove upward bias caused by  $\max_a Q(s, a, \mathbf{w})$ 
  - ▶ Current Q-network  $\mathbf{w}$  is used to **select** actions
  - ▶ Older Q-network  $\mathbf{w}^-$  is used to **evaluate** actions

$$l = \left( r + \gamma \underset{a'}{\operatorname{argmax}} Q(s', a', \mathbf{w}, \mathbf{w}^-) - Q(s, a, \mathbf{w}) \right)^2$$

- ▶ **Prioritised replay:** Weight experience according to surprise
  - ▶ Store experience in priority queue according to DQN error

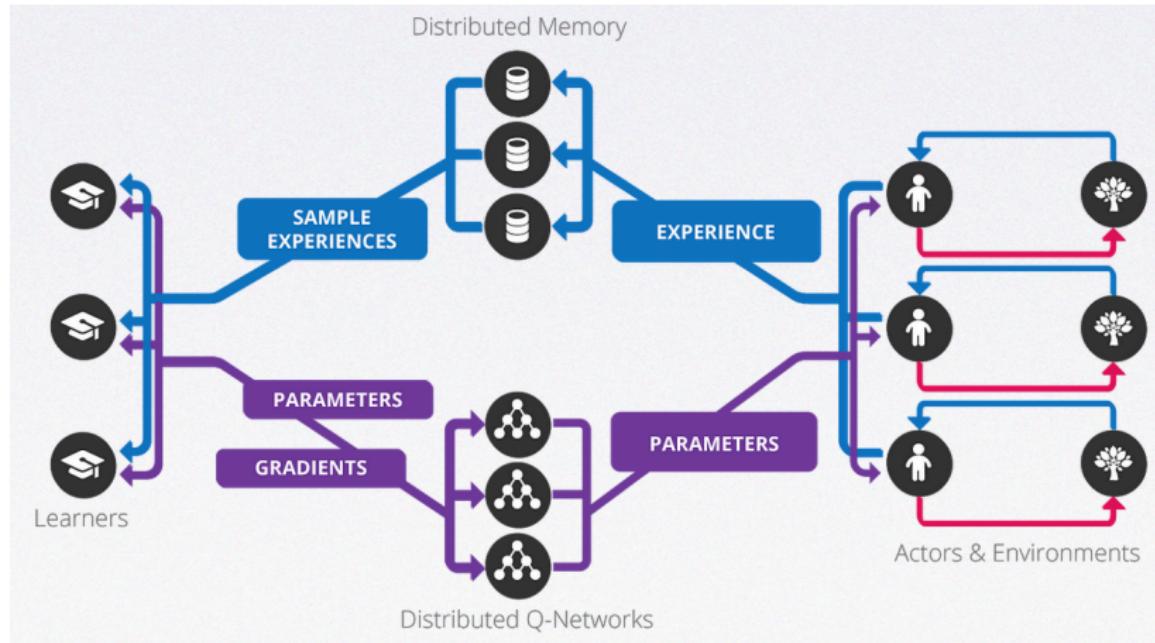
$$\left| r + \gamma \underset{a'}{\max} Q(s', a', \mathbf{w}^-) - Q(s, a, \mathbf{w}) \right|$$

- ▶ **Duelling network:** Split Q-network into two channels
  - ▶ Action-independent **value function**  $V(s, v)$
  - ▶ Action-dependent **advantage function**  $A(s, a, \mathbf{w})$

$$Q(s, a) = V(s, v) + A(s, a, \mathbf{w})$$

Combined algorithm: 3x mean Atari score vs Nature DQN

# Gorila (General Reinforcement Learning Architecture)



- ▶ 10x faster than Nature DQN on 38 out of 49 Atari games
- ▶ Applied to recommender systems within Google

# Asynchronous Reinforcement Learning

- ▶ Exploits multithreading of standard CPU
- ▶ Execute many instances of agent in parallel
- ▶ Network parameters shared between threads
- ▶ Parallelism decorrelates data
  - ▶ Viable alternative to experience replay
  - ▶ Parallelism decorrelates data
- ▶ Similar speedup to Gorila - on a single machine!

# Outline

Introduction to Deep Learning

Introduction to Reinforcement Learning

Value-Based Deep RL

Policy-Based Deep RL

Model-Based Deep RL

# Deep Policy Networks

- ▶ Represent policy by deep network with weights  $\mathbf{u}$

$$a = \pi(a|s, \mathbf{u}) \text{ or } a = \pi(s, \mathbf{u})$$

- ▶ Define objective function as total discounted reward

$$L(\mathbf{u}) = \mathbb{E} [r_1 + \gamma r_2 + \gamma^2 r_3 + \dots | \pi(\cdot, \mathbf{u})]$$

- ▶ Optimise objective end-to-end by SGD
- ▶ i.e. Adjust policy parameters  $\mathbf{u}$  to achieve more reward

# Policy Gradients

How to make high-value actions more likely:

- ▶ The gradient of a stochastic policy  $\pi(a|s, \mathbf{u})$  is given by

$$\frac{\partial L(\mathbf{u})}{\partial \mathbf{u}} = \mathbb{E} \left[ \frac{\partial \log \pi(a|s, \mathbf{u})}{\partial \mathbf{u}} Q^\pi(s, a) \right]$$

# Policy Gradients

How to make high-value actions more likely:

- ▶ The gradient of a stochastic policy  $\pi(a|s, \mathbf{u})$  is given by

$$\frac{\partial L(\mathbf{u})}{\partial \mathbf{u}} = \mathbb{E} \left[ \frac{\partial \log \pi(a|s, \mathbf{u})}{\partial \mathbf{u}} Q^\pi(s, a) \right]$$

- ▶ The gradient of a deterministic policy  $a = \pi(s)$  is given by

$$\frac{\partial L(\mathbf{u})}{\partial \mathbf{u}} = \mathbb{E} \left[ \frac{\partial Q^\pi(s, a)}{\partial a} \frac{\partial a}{\partial \mathbf{u}} \right]$$

- ▶ if  $a$  is continuous and  $Q$  is differentiable

# Actor-Critic Algorithm

- ▶ Estimate value function  $Q(s, a, \mathbf{w}) \approx Q^\pi(s, a)$
- ▶ Update policy parameters  $\mathbf{u}$  by stochastic gradient ascent

$$\frac{\partial I}{\partial \mathbf{u}} = \frac{\partial \log \pi(a|s, \mathbf{u})}{\partial \mathbf{u}} Q(s, a, \mathbf{w})$$

or

$$\frac{\partial I}{\partial \mathbf{u}} = \frac{\partial Q(s, a, \mathbf{w})}{\partial a} \frac{\partial a}{\partial \mathbf{u}}$$

## Asynchronous Advantage Actor-Critic (A3C)

- ▶ Estimate state-value function

$$V(s, \mathbf{v}) \approx \mathbb{E} [r_{t+1} + \gamma r_{t+2} + \dots | s]$$

- ▶ Q-value estimated by an  $n$ -step sample

$$Q(s_t, a_t) = r_{t+1} + \gamma r_{t+2} \dots + \gamma^{n-1} r_{t+n} + \gamma^n V(s_{t+n}, \mathbf{v})$$

# Asynchronous Advantage Actor-Critic (A3C)

- ▶ Estimate state-value function

$$V(s, \mathbf{v}) \approx \mathbb{E} [r_{t+1} + \gamma r_{t+2} + \dots | s]$$

- ▶ Q-value estimated by an  $n$ -step sample

$$Q(s_t, a_t) = r_{t+1} + \gamma r_{t+2} \dots + \gamma^{n-1} r_{t+n} + \gamma^n V(s_{t+n}, \mathbf{v})$$

- ▶ Actor is updated towards target

$$\frac{\partial I_u}{\partial \mathbf{u}} = \frac{\partial \log \pi(a_t | s_t, \mathbf{u})}{\partial \mathbf{u}} (Q(s_t, a_t) - V(s_t, \mathbf{v}))$$

- ▶ Critic is updated to minimise MSE w.r.t. target

$$I_v = (Q(s_t, a_t) - V(s_t, \mathbf{v}))^2$$

# Asynchronous Advantage Actor-Critic (A3C)

- ▶ Estimate state-value function

$$V(s, \mathbf{v}) \approx \mathbb{E} [r_{t+1} + \gamma r_{t+2} + \dots | s]$$

- ▶ Q-value estimated by an  $n$ -step sample

$$Q(s_t, a_t) = r_{t+1} + \gamma r_{t+2} \dots + \gamma^{n-1} r_{t+n} + \gamma^n V(s_{t+n}, \mathbf{v})$$

- ▶ Actor is updated towards target

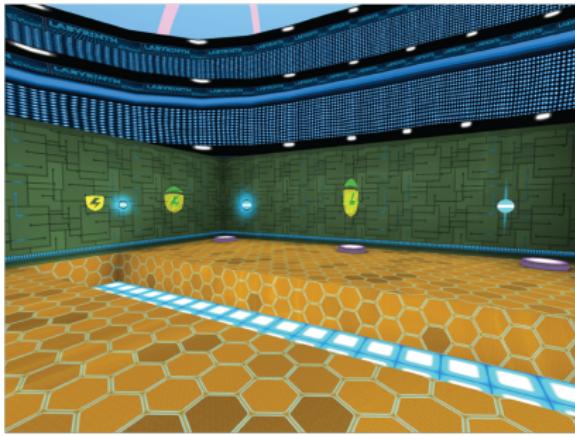
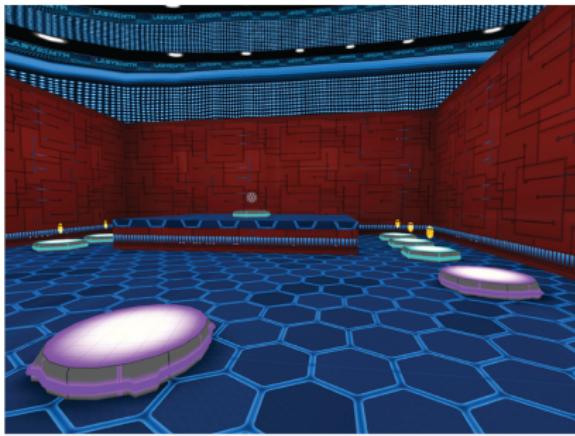
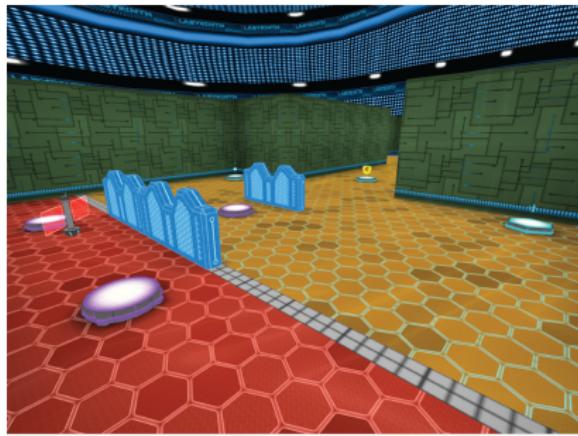
$$\frac{\partial I_u}{\partial \mathbf{u}} = \frac{\partial \log \pi(a_t | s_t, \mathbf{u})}{\partial \mathbf{u}} (Q(s_t, a_t) - V(s_t, \mathbf{v}))$$

- ▶ Critic is updated to minimise MSE w.r.t. target

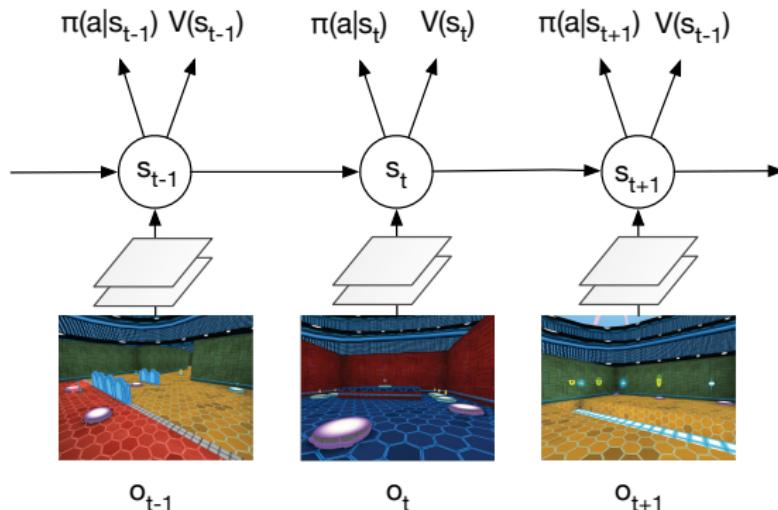
$$I_v = (Q(s_t, a_t) - V(s_t, \mathbf{v}))^2$$

- ▶ 4x mean Atari score vs Nature DQN

# Deep Reinforcement Learning in Labyrinth



# A3C in Labyrinth



- ▶ End-to-end learning of softmax policy  $\pi(a|s_t)$  from pixels
- ▶ Observations  $o_t$  are raw pixels from current frame
- ▶ State  $s_t = f(o_1, \dots, o_t)$  is a recurrent neural network (LSTM)
- ▶ Outputs both value  $V(s)$  and softmax over actions  $\pi(a|s)$
- ▶ Task is to collect apples (+1 reward)

# A3C Labyrinth Demo

Demo:

[www.youtube.com/watch?v=nMR5mjCFZCw&feature=youtu.be](https://www.youtube.com/watch?v=nMR5mjCFZCw&feature=youtu.be)

Labyrinth source code (coming soon):

[sites.google.com/a/deepmind.com/labyrinth/](https://sites.google.com/a/deepmind.com/labyrinth/)

# Deep Reinforcement Learning with Continuous Actions

How can we deal with high-dimensional continuous action spaces?

- ▶ Can't easily compute  $\max_a Q(s, a)$ 
  - ▶ Actor-critic algorithms learn without max
- ▶ Q-values are differentiable w.r.t  $a$ 
  - ▶ Deterministic policy gradients exploit knowledge of  $\frac{\partial Q}{\partial a}$

# Deep DPG

DPG is the continuous analogue of DQN

- ▶ **Experience replay**: build data-set from agent's experience
- ▶ **Critic** estimates value of current policy by DQN

$$l_w = \left( r + \gamma Q(s', \pi(s', \mathbf{u}^-), \mathbf{w}^-) - Q(s, a, \mathbf{w}) \right)^2$$

To deal with non-stationarity, targets  $\mathbf{u}^-$ ,  $\mathbf{w}^-$  are held fixed

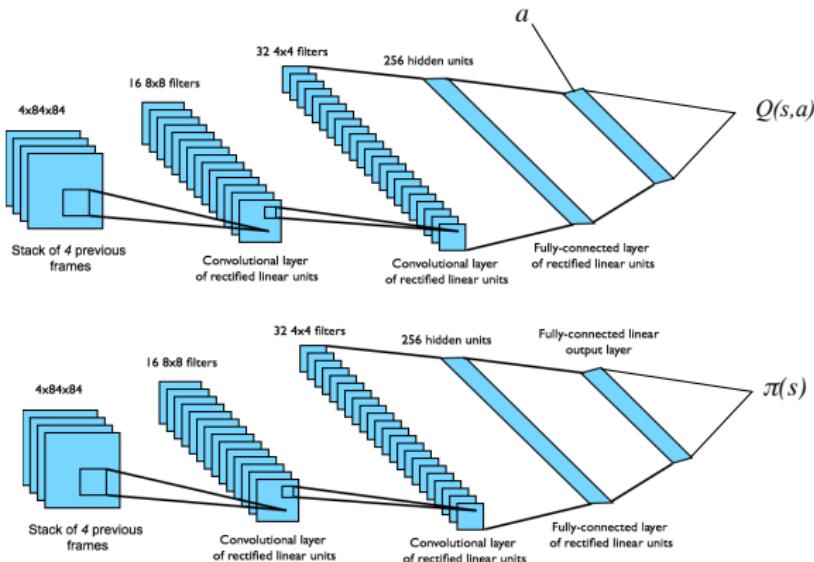
- ▶ **Actor** updates policy in direction that improves  $Q$

$$\frac{\partial l_u}{\partial \mathbf{u}} = \frac{\partial Q(s, a, \mathbf{w})}{\partial a} \frac{\partial a}{\partial \mathbf{u}}$$

- ▶ In other words critic provides loss function for actor

# DPG in Simulated Physics

- ▶ Physics domains are simulated in MuJoCo
- ▶ End-to-end learning of control policy from raw pixels  $s$
- ▶ Input state  $s$  is stack of raw pixels from last 4 frames
- ▶ Two separate convnets are used for  $Q$  and  $\pi$
- ▶ Policy  $\pi$  is adjusted in direction that most improves  $Q$

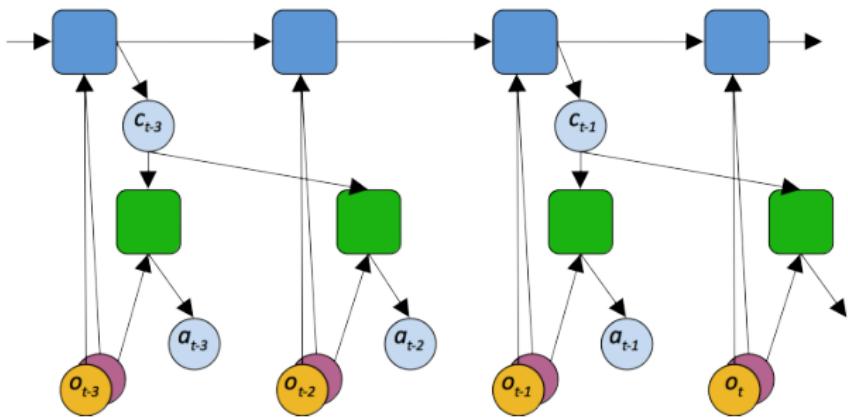


# DPG in Simulated Physics Demo

- ▶ Demo: DPG from pixels

# A3C in Simulated Physics Demo

- ▶ Asynchronous RL is viable alternative to experience replay
- ▶ Train a hierarchical, recurrent locomotion controller
- ▶ Retrain controller on more challenging tasks



# Fictitious Self-Play (FSP)

Can deep RL find Nash equilibria in multi-agent games?

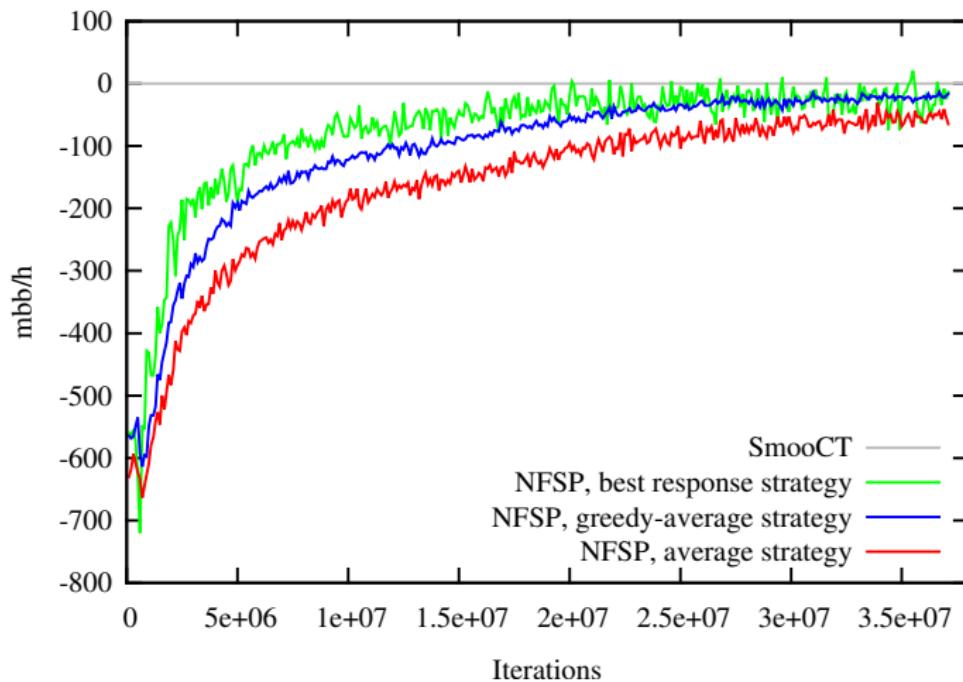
- ▶ Q-network learns “best response” to opponent policies
  - ▶ By applying DQN with experience replay
  - ▶ c.f. fictitious play
- ▶ Policy network  $\pi(a|s, \mathbf{u})$  learns an average of best responses

$$\frac{\partial I}{\partial \mathbf{u}} = \frac{\partial \log \pi(a|s, \mathbf{u})}{\partial \mathbf{u}}$$

- ▶ Actions a sample mix of policy network and best response

# Neural FSP in Texas Hold'em Poker

- ▶ Heads-up limit Texas Hold'em
- ▶ NFSP with raw inputs only (no prior knowledge of Poker)
- ▶ vs SmooCT (3x medal winner 2015, handcrafted knowlege)



# Outline

Introduction to Deep Learning

Introduction to Reinforcement Learning

Value-Based Deep RL

Policy-Based Deep RL

Model-Based Deep RL

# Learning Models of the Environment

- ▶ Demo: generative model of Atari
- ▶ Challenging to plan due to compounding errors
  - ▶ Errors in the transition model compound over the trajectory
  - ▶ Planning trajectories differ from executed trajectories
  - ▶ At end of long, unusual trajectory, rewards are totally wrong

# Deep Reinforcement Learning in Go

What if we have a perfect model? e.g. game rules are known

AlphaGo paper:

[www.nature.com/articles/nature16961](http://www.nature.com/articles/nature16961)

AlphaGo resources:

[deepmind.com/alphago/](http://deepmind.com/alphago/)



# Conclusion

- ▶ General, stable and scalable RL is now possible
- ▶ Using deep networks to represent value, policy, model
- ▶ Successful in Atari, Labyrinth, Physics, Poker, Go
- ▶ Using a variety of deep RL paradigms