



**Barcelona
Supercomputing
Center**

Centro Nacional de Supercomputación



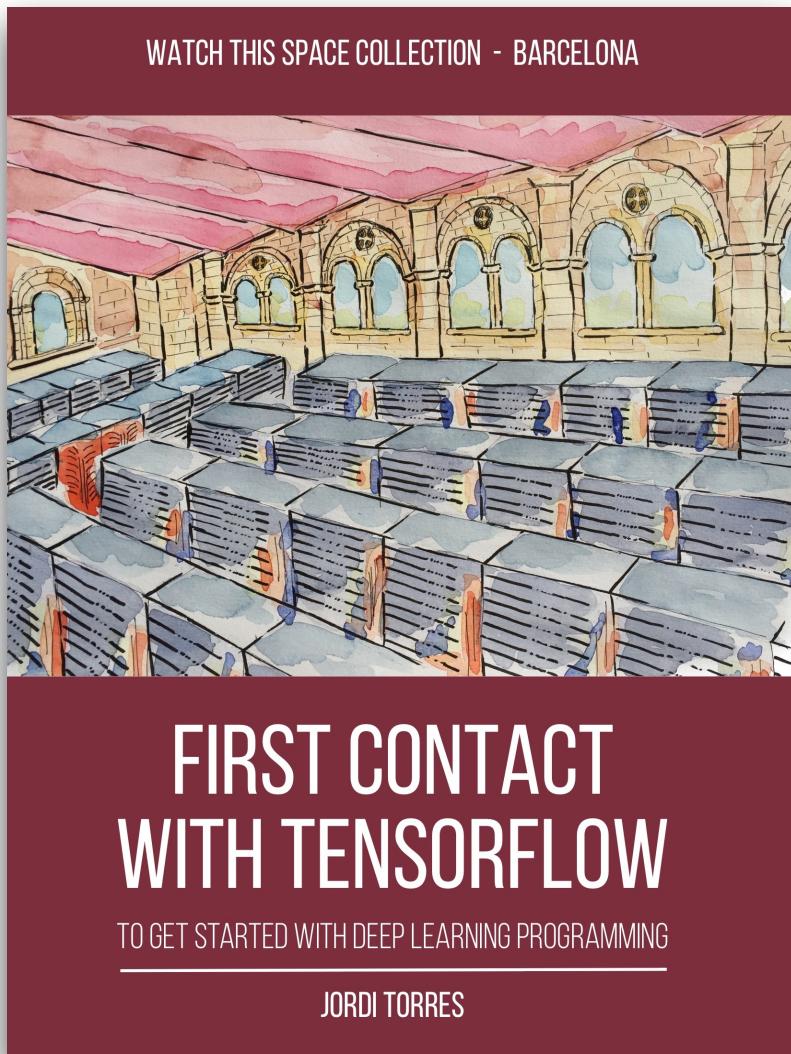
First Contact with TensorFlow

Part 1: Basics

PRACE Advanced Training Centre Course: Big Data Analytics
Autonomic Systems and eBusiness Platforms
Computer Science Department
Jordi Torres

10/02/2016

Complete Course Content:



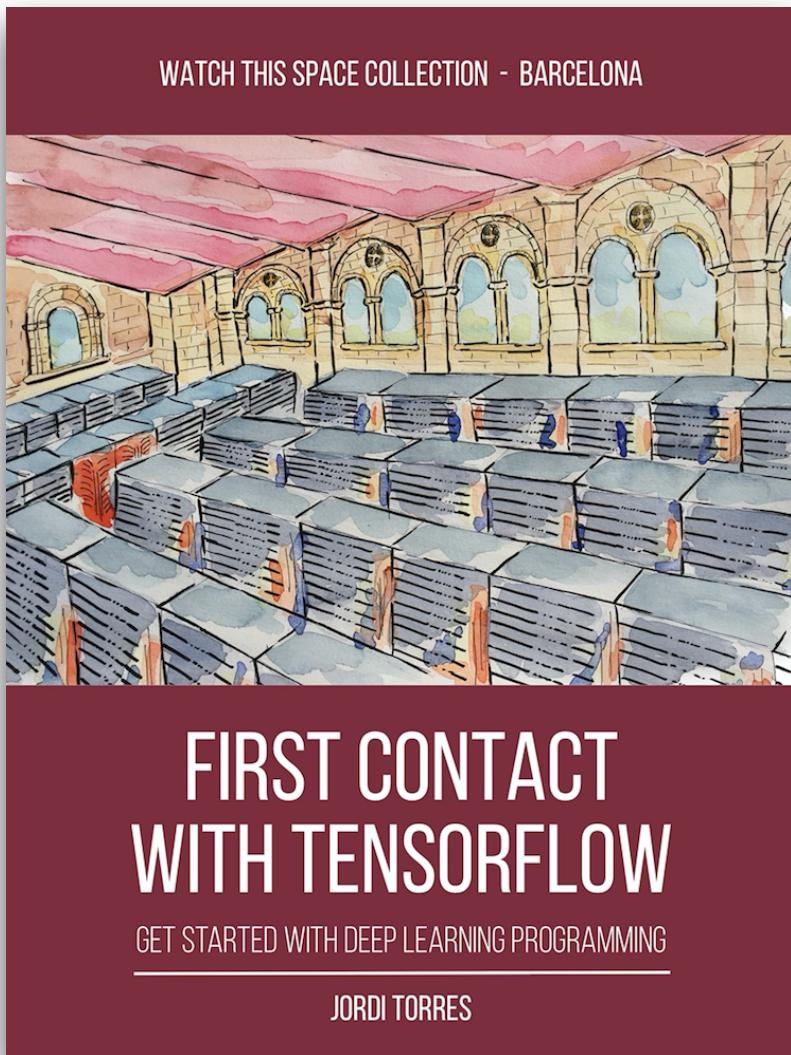
Part1: Basics

- Introduction
- 1. TensorFlow Basics
- 2. Linear Regression
- 3. Clustering

Part2: Advanced

- 4. Single Layer Neural Network
- 5. Multi-layer Neural Networks
- 6. TensorFlow and GPUs
- Closing

Today:

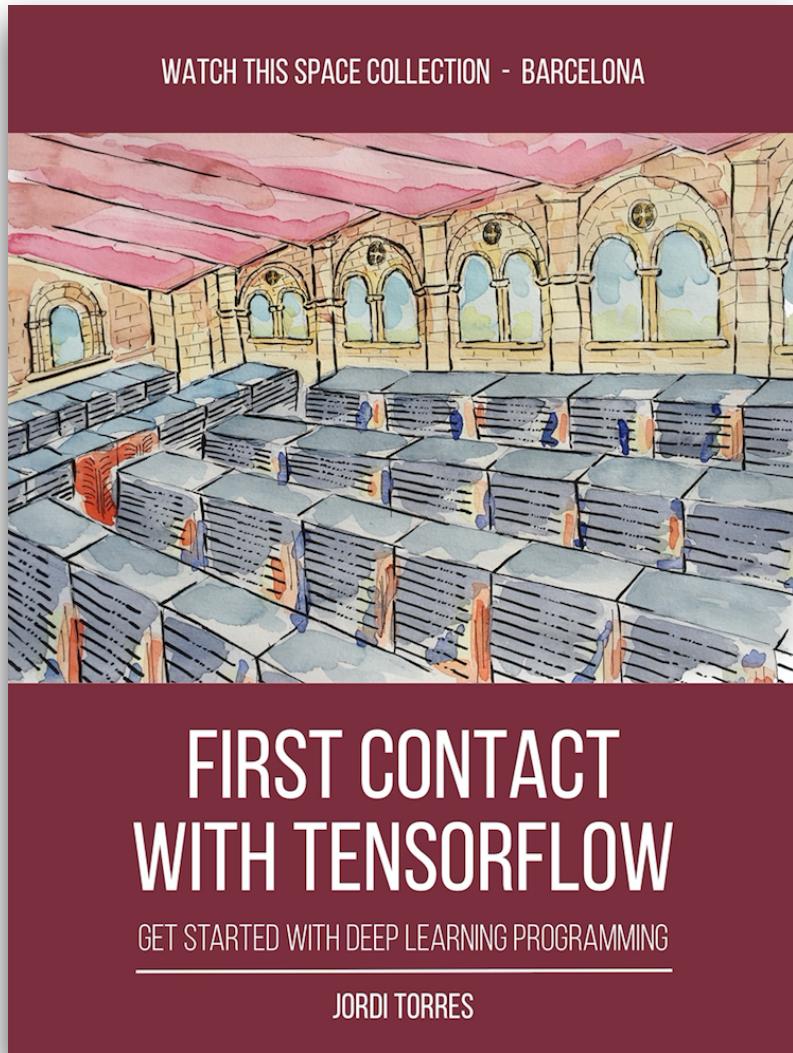


Part1: Basics

Introduction

1. TensorFlow Basics
2. Linear Regression
3. Clustering

Content:



Introduction

1. TensorFlow Basics
 2. Linear Regression
 3. Clustering
 4. Single Layer Neural Network
 5. Multi-layer Neural Networks
 6. TensorFlow and GPUs
- Closing

Global Challenge Insight Report

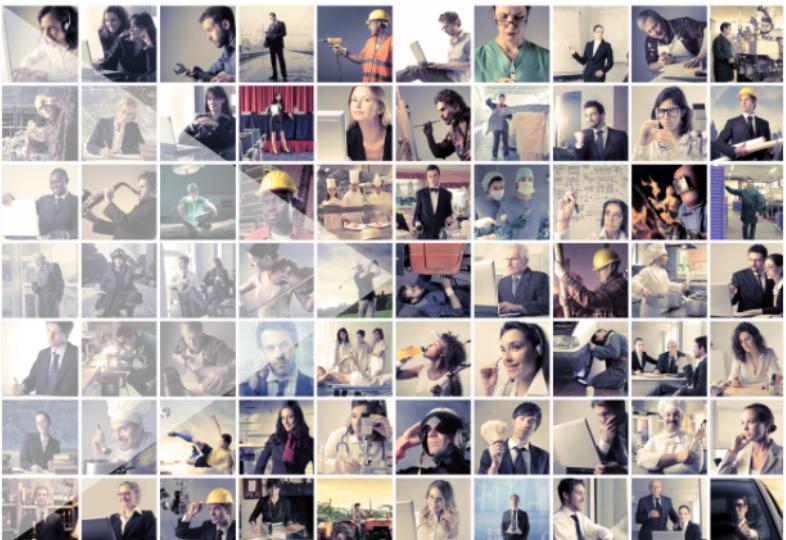
WORLD
ECONOMIC
FORUM

COMMITTED TO
IMPROVING THE STATE
OF THE WORLD

The Future of Jobs

Employment, Skills and Workforce Strategy for the Fourth Industrial Revolution

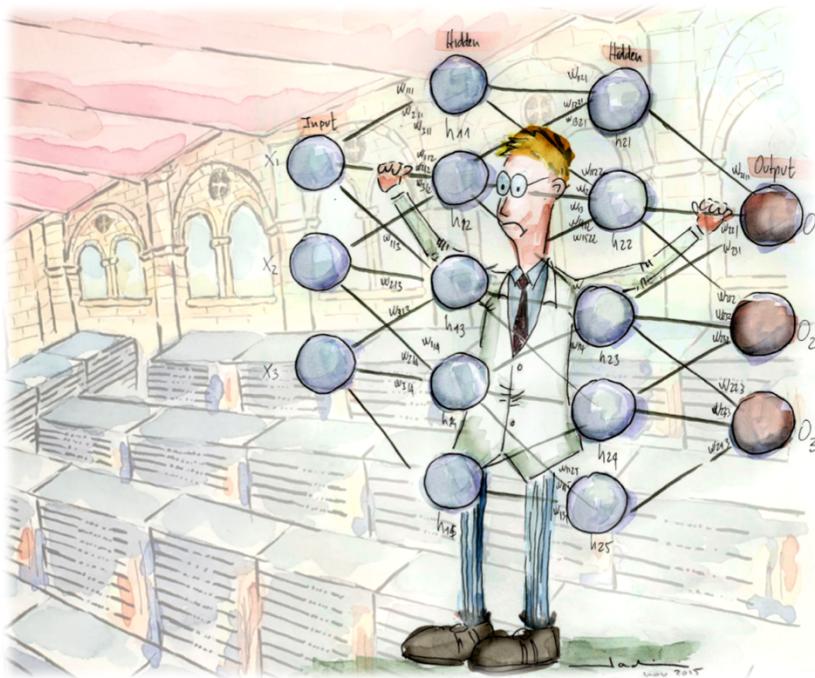
January 2016



“The Fourth Industrial Revolution, which includes developments in previously disjointed fields such as artificial intelligence and machine-learning, robotics, nanotechnology, 3-D printing, and genetics and biotechnology, will cause widespread disruption not only to business models but also to labour markets over the next five years, with enormous change predicted in the skill sets needed to thrive in the new landscape.”

~~Future?~~ Present!

New type of Computing is required



Giving computers a greater ability to understand information, and to learn, to reason, and act upon it

Data Analytics Today

Applications

Artificial neural network

Latent Dirichlet Allocation

Gaussian process regression

Support vector machines

Random Forests

Linear classifiers

k-nearest neighbor

Bayesian networks

Naive Bayes

Hidden Markov models

Unsupervised learning

Expectation-maximization alg.

K-means algorithm

DBSCAN

Deep learning

Linear regression

Logistic regression

...

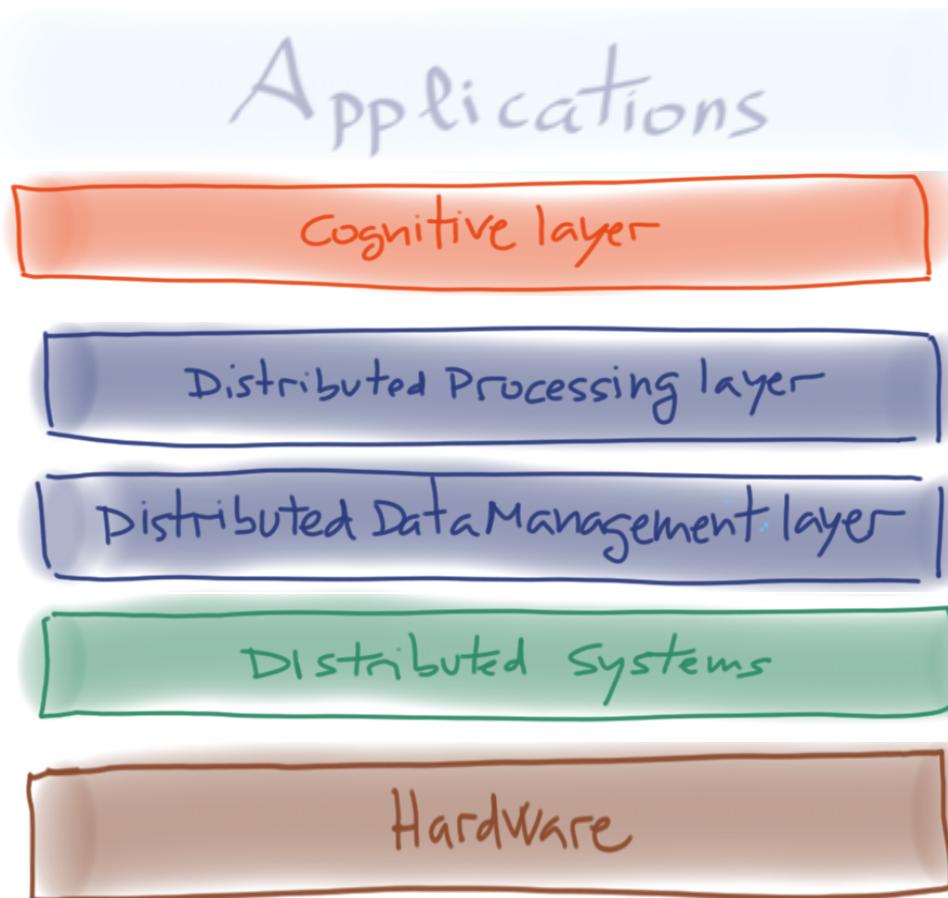
Distributed Processing layer

Distributed Data Management layer

Distributed Systems

Hardware

Advanced Analytics will be provided by the “system”



Meanwhile . . .

DEEP LEARNING	STREAM ANALYTICS	BIG DATA ML	DATA MINING	ML as a Service
Caffe	Theano	Apache S4	Spark MLlib	Amazon ML Big ML
Convnet.js	PyLearn2	Summingbird	Mahout	Predictis AlchemyAPI
Keras	DL4J	SAMOA	Flink ML	Google Predict. Diffbot
Torch7	Lasagne			IBM Watson Azure ML

. . . and TensorFlow!

Attendance background?

“The course has a practical nature, and therefore I reduced the theoretical part as much as possible, assuming that the attendance has some basic understanding about Machine Learning”

Machine Learning?

- Field of Computer Science that evolved from the study of pattern recognition and computational learning theory into Artificial Intelligence.
- Its goal is to give computers the ability to learn without being explicitly programmed.
- For this purpose, Machine Learning uses mathematical/statistical techniques to construct models from a set of observed data rather than have specific set of instructions entered by the user that define the model for that set of data.

Why Deep Learning?

- Conventional Machine Learning techniques were limited in their ability to process natural data in their raw form;
for example, to identify objects in images or transcribe speech into text.

- Increasingly, these applications make use of a class of techniques called Deep Learning.

These methods are based on “deep” multi-layer neural networks with many hidden layers that attempt to model high-level abstractions in data.

- Right now, a research in diverse types of Deep Learning’s networks is being developed

Deep Convolutional Nets have brought breakthroughs in processing images and speech: we will consider this type in this course.

Why is Deep Learning taking off now?

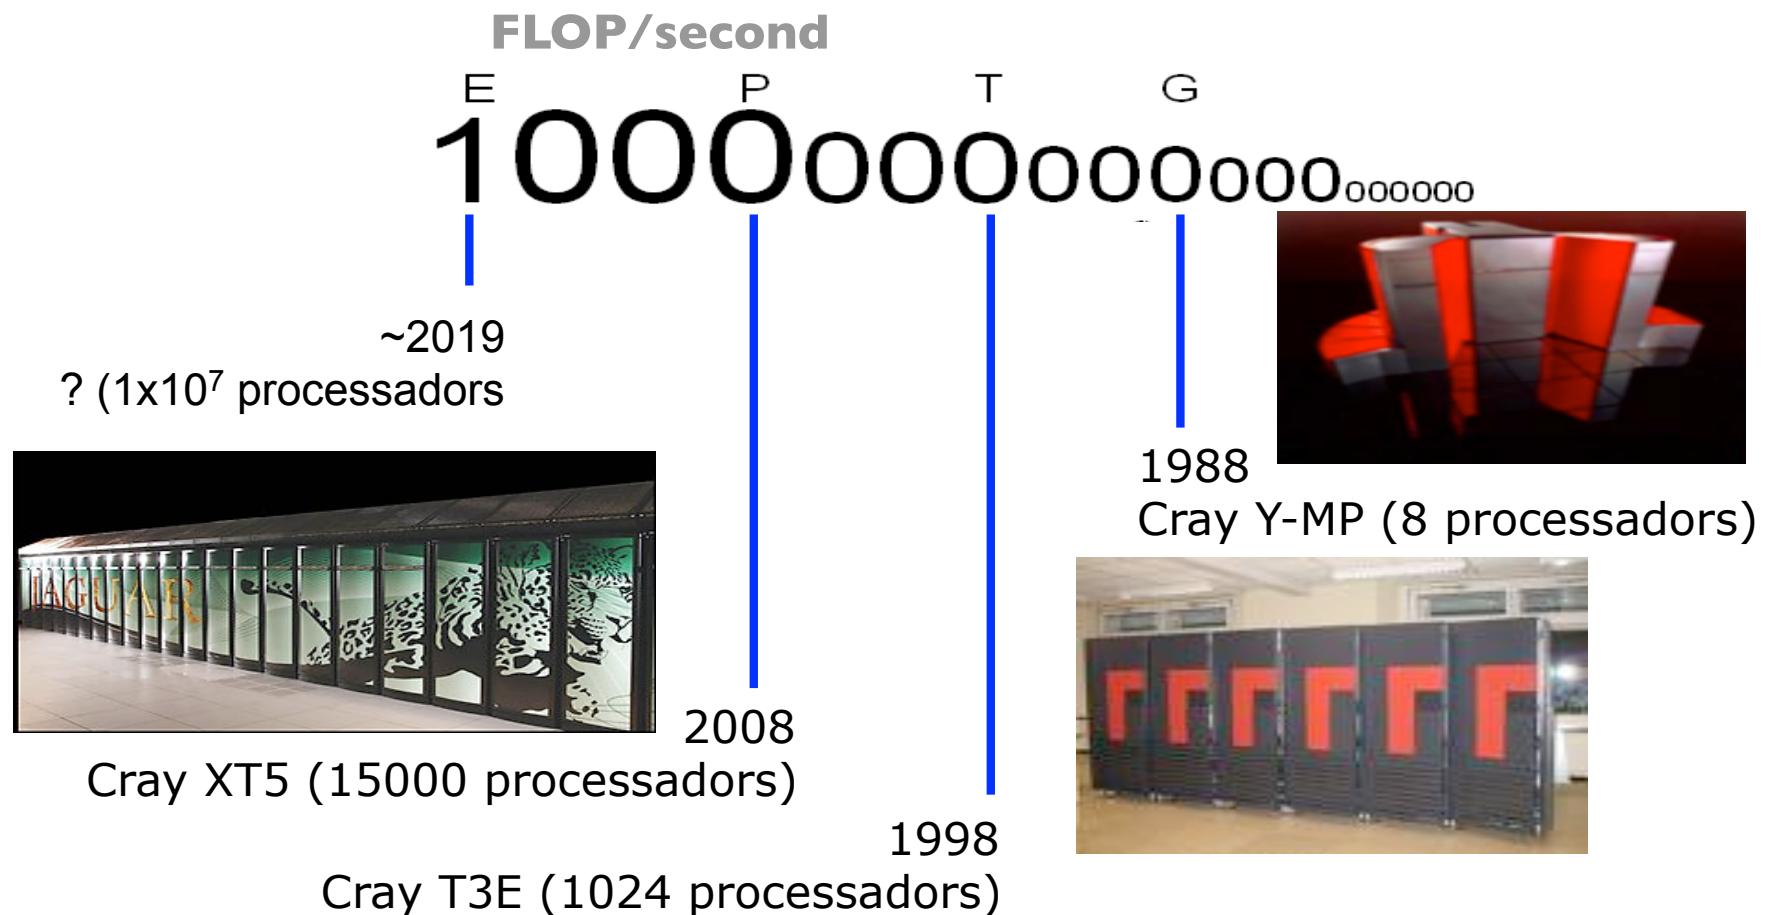
- It is known that many of the ideas used in Deep Learning have been around for decades.
- One of the key drivers of its current progress is clearly the huge deluge of data available today. Thanks to the advent of Big Data these models can be “trained” by exposing them to large data sets that were previously unavailable.
- But another not less important driver is the computation power available today. As an example, due to the raising of GPUs, Deep Learning’s community started shifting to GPUs (TensorFlow works with GPUs).

Data deluge

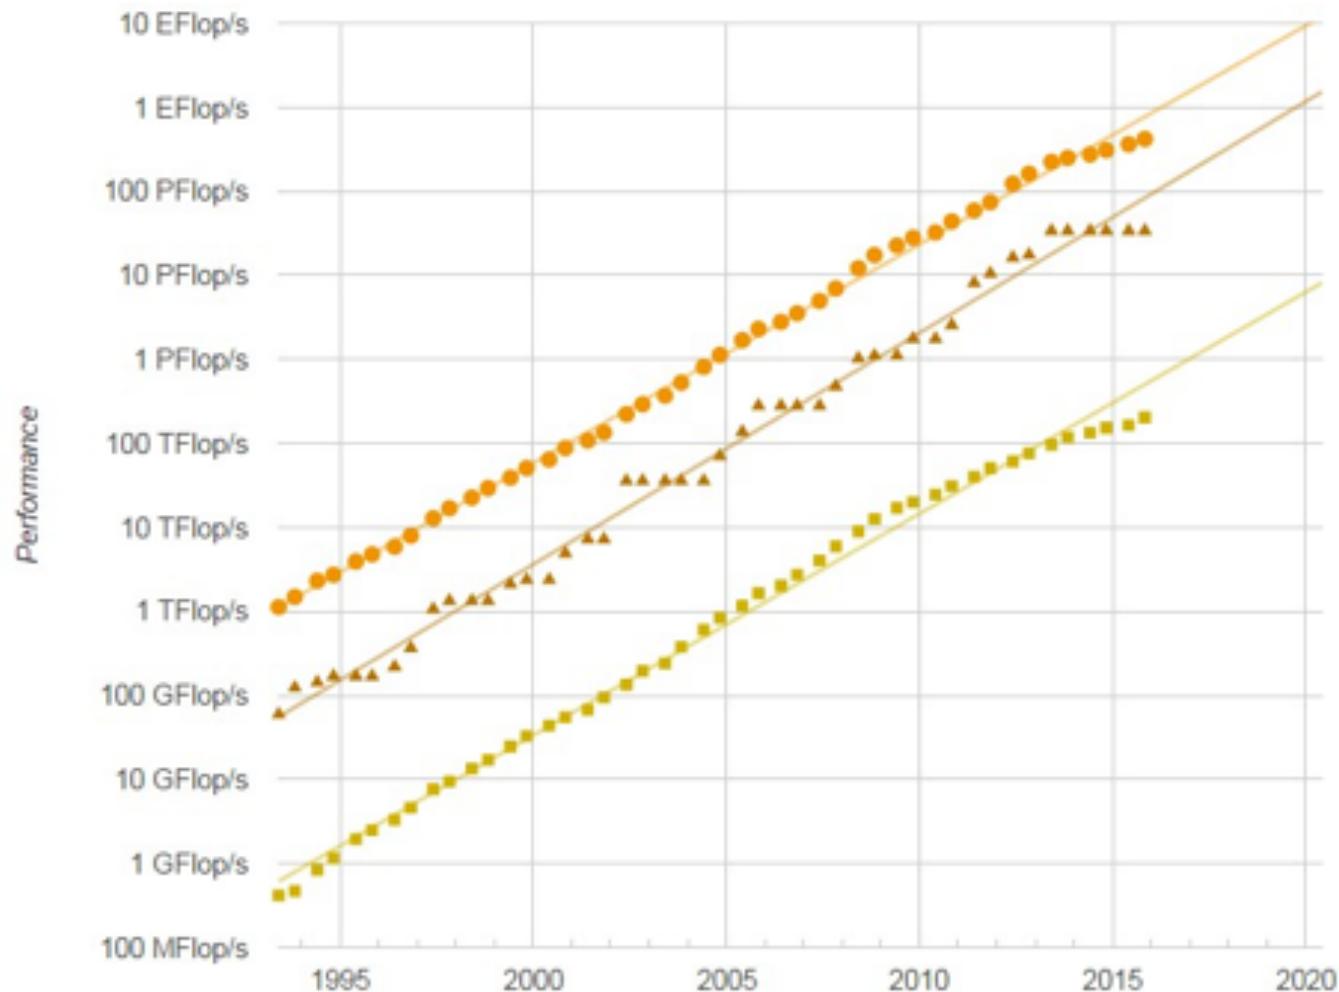


SOURCE: <http://www.cray.com/Assets/Images/urika/edge/analytics-infographic.html>

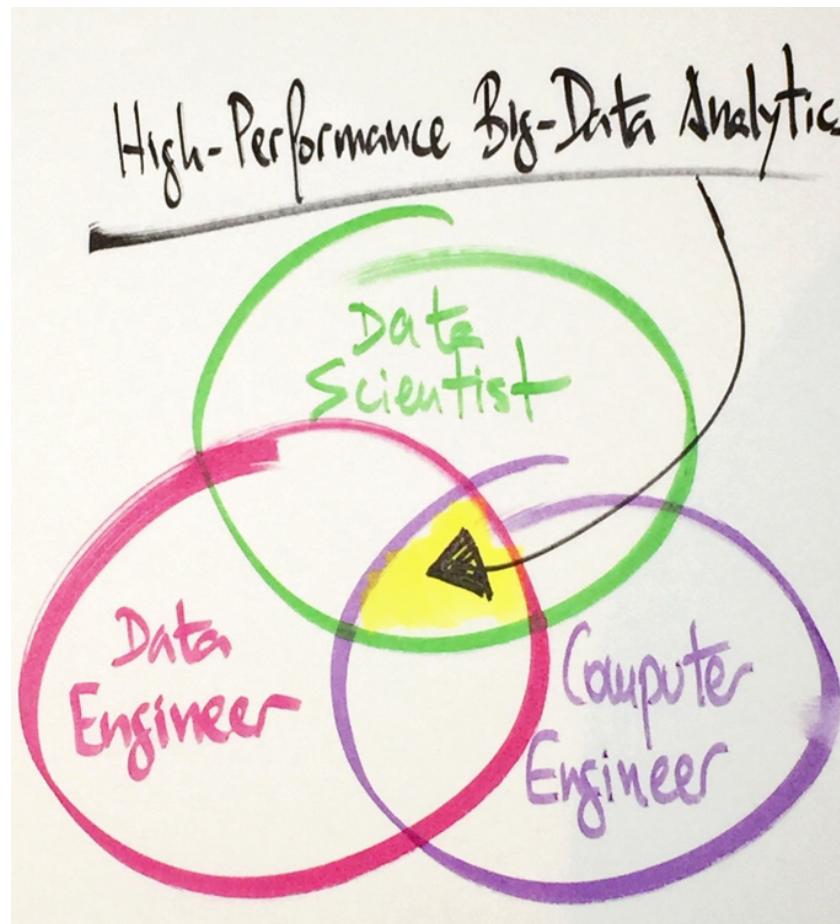
Evolution of Computers



Top500 list



New requirements!





*Foreword: **Oriol Vinyals**
Research Scientist en Google Brain*

“I would like to thank the author, whom I have the pleasure of knowing, the effort to disseminate this technology. He had written this book (first Spanish version) in record time, two months after it was announced the open source project.”

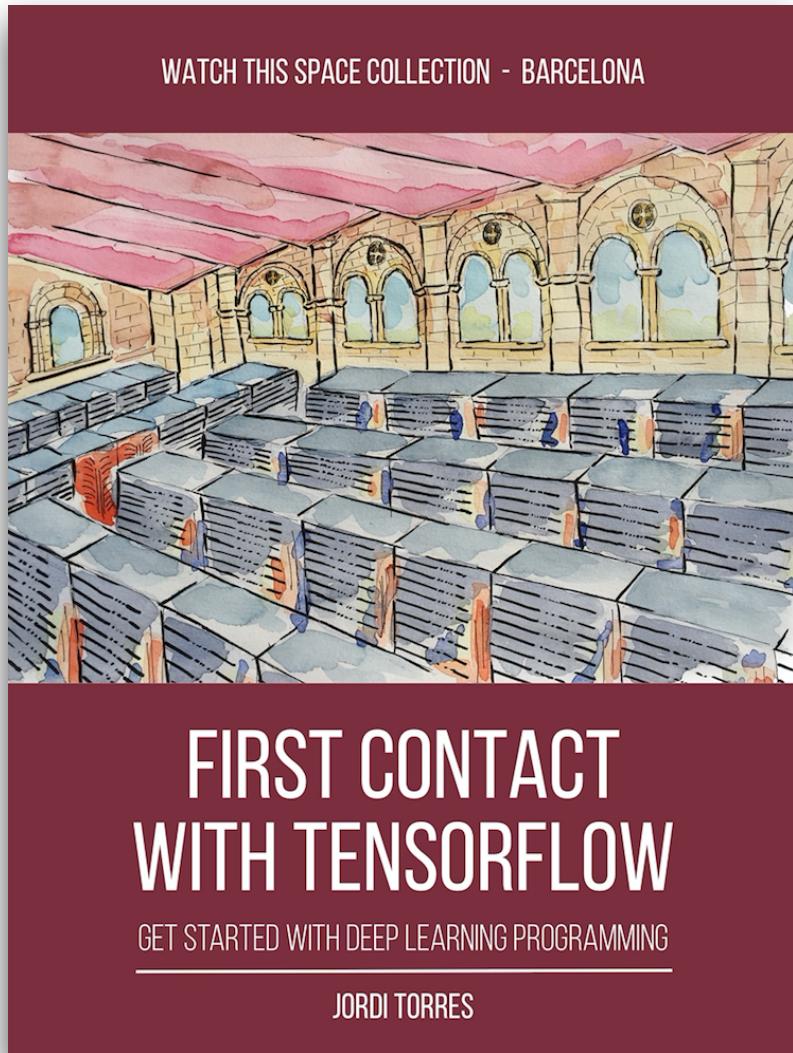
Learn by doing!

*"Tell me and I forget.
Teach me and I remember.
Involve me and I learn"*

- Benjamin Franklin

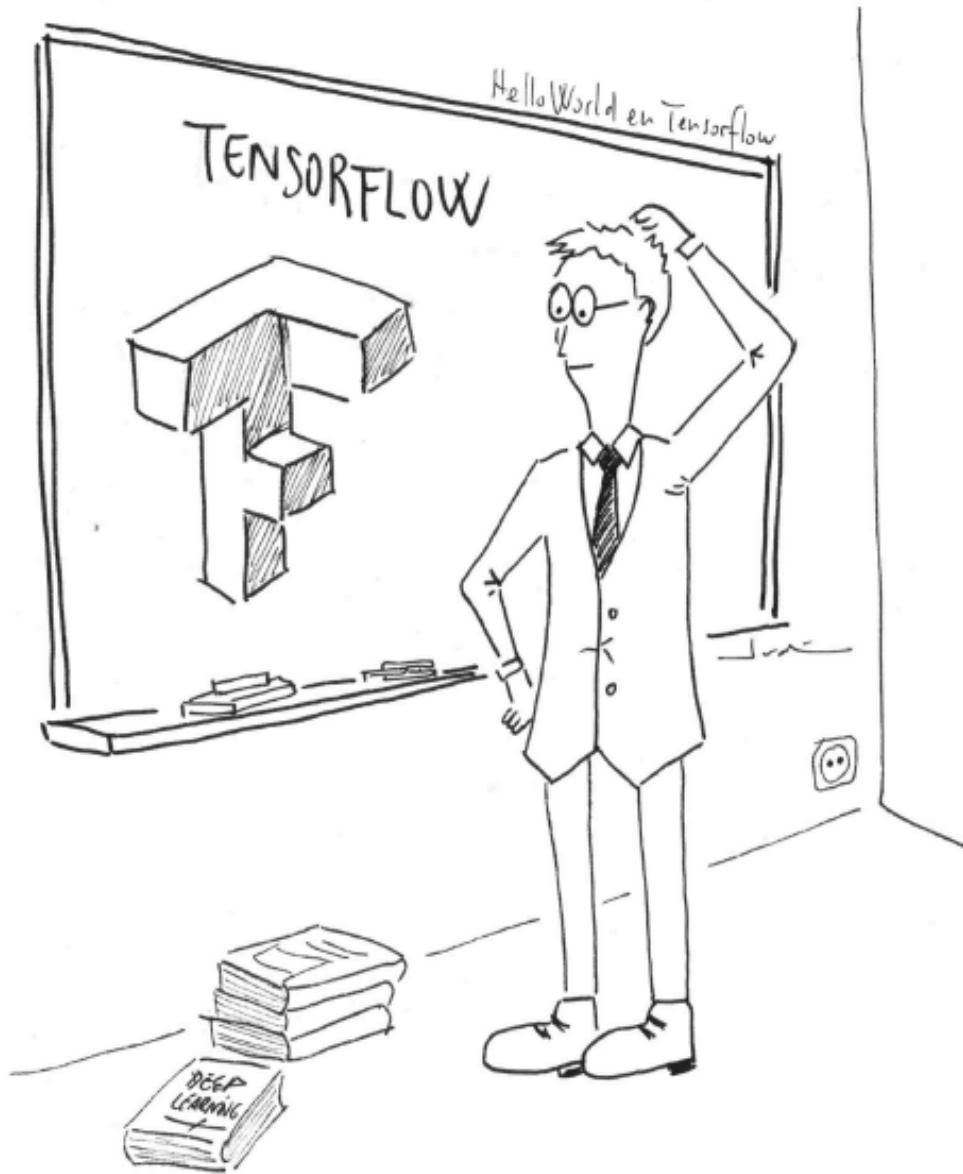


Content:



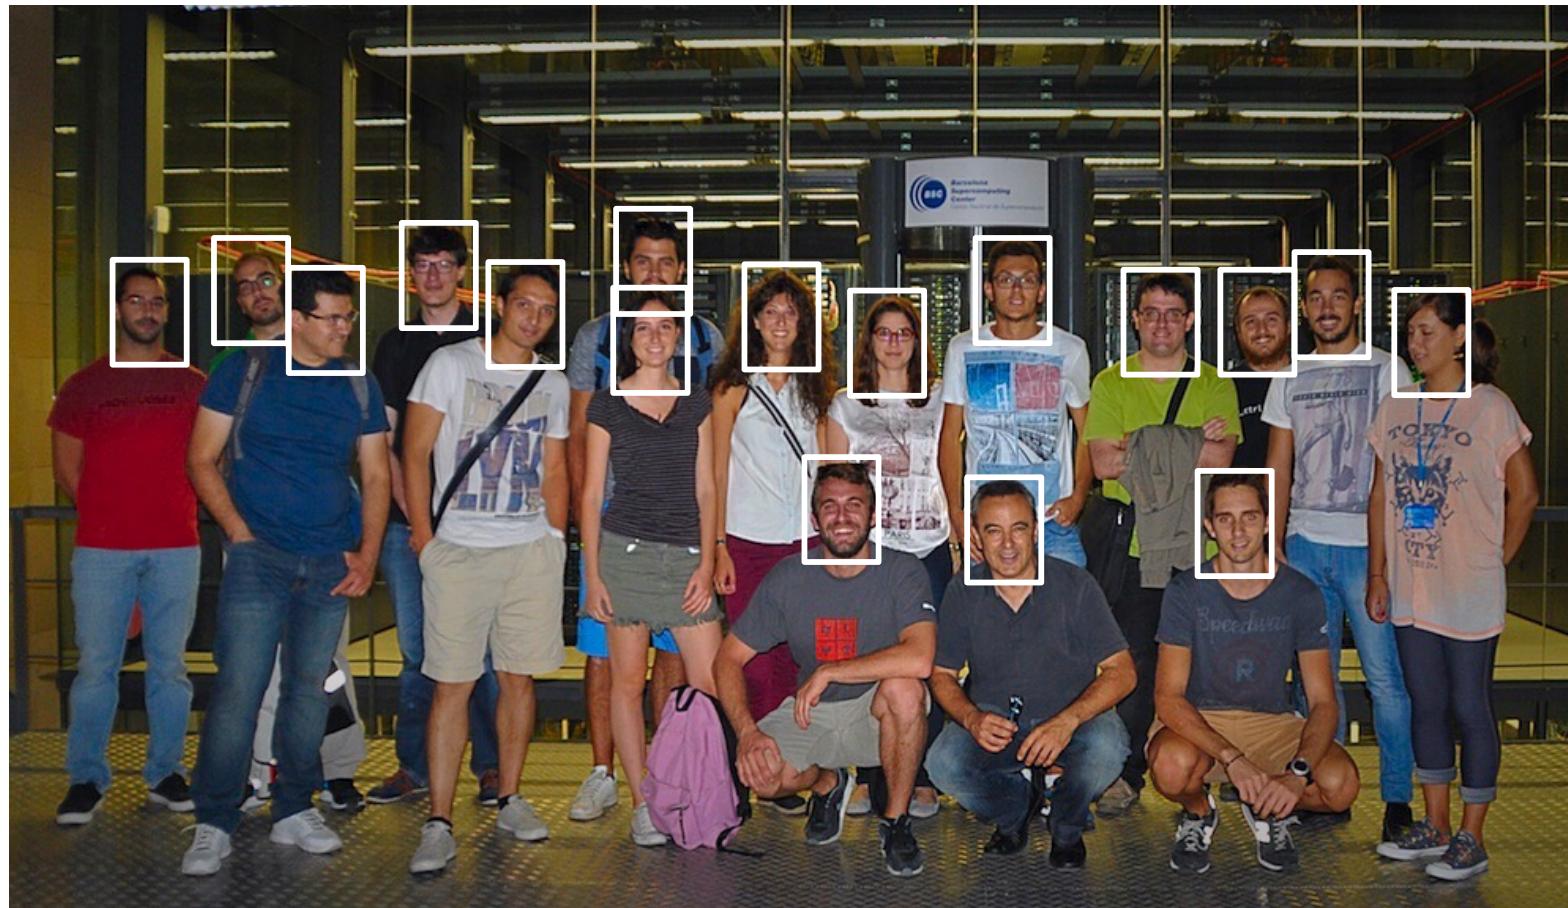
Introduction

1. TensorFlow Basics
 2. Linear Regression
 3. Clustering
 4. Single Layer Neural Network
 5. Multi-layer Neural Networks
 6. TensorFlow and GPUs
- Closing



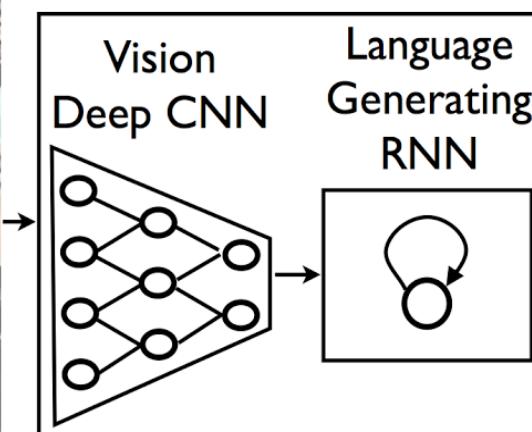
A few years ago ...

A breakthrough in ML suddenly enabled computers to recognize objects shown in photographs with unprecedented accuracy.



The question now is . . .

Whether machines can make another leap, by learning to make sense of what's actually going on in such images.

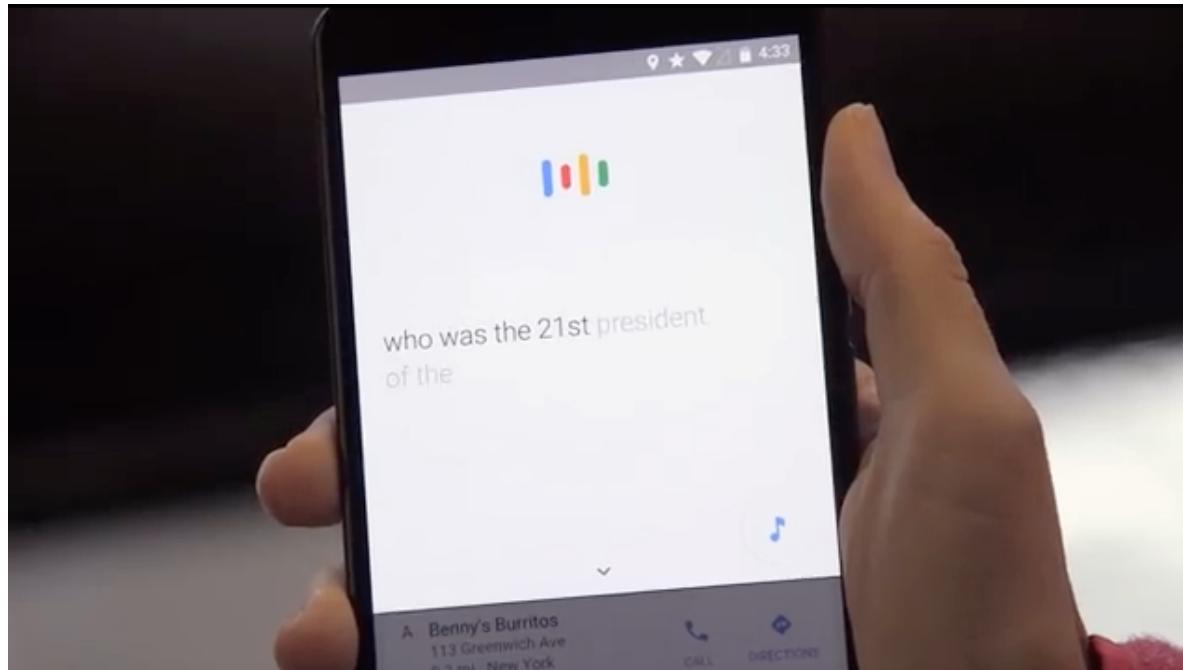


A group of people shopping at an outdoor market.

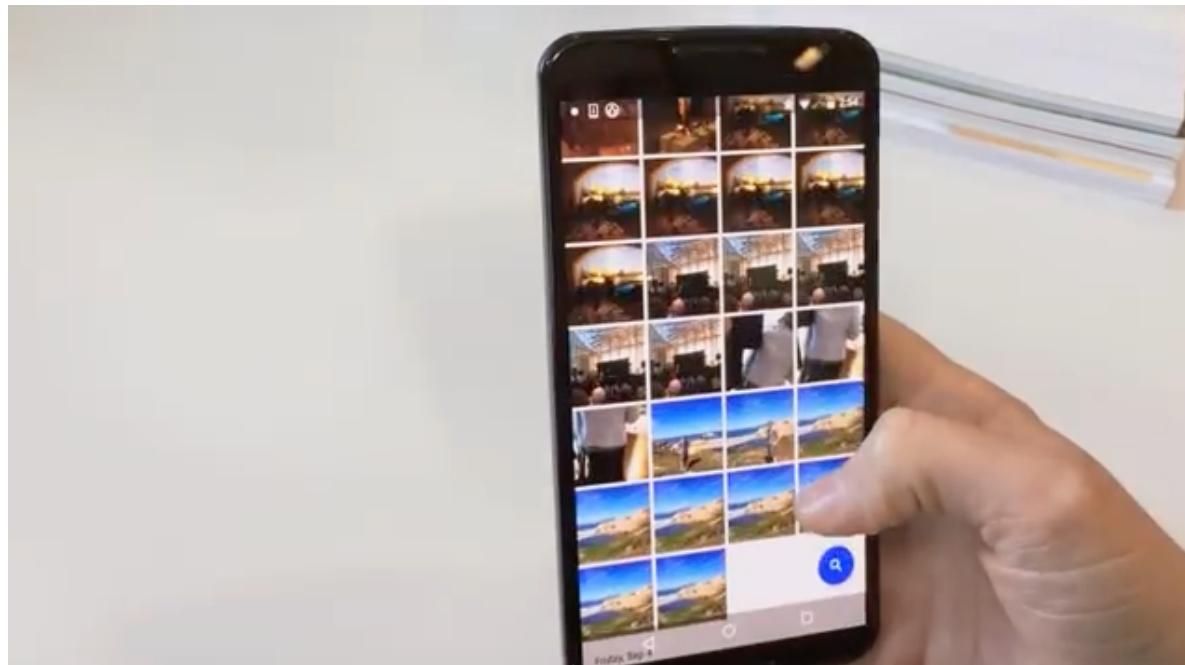
There are many vegetables at the fruit stand.

Source: Oriol Vinyals – Research Scientist at Google Brain

TensorFlow at Google



TensorFlow at Google



TensorFlow at Google



TensorFlow at Google

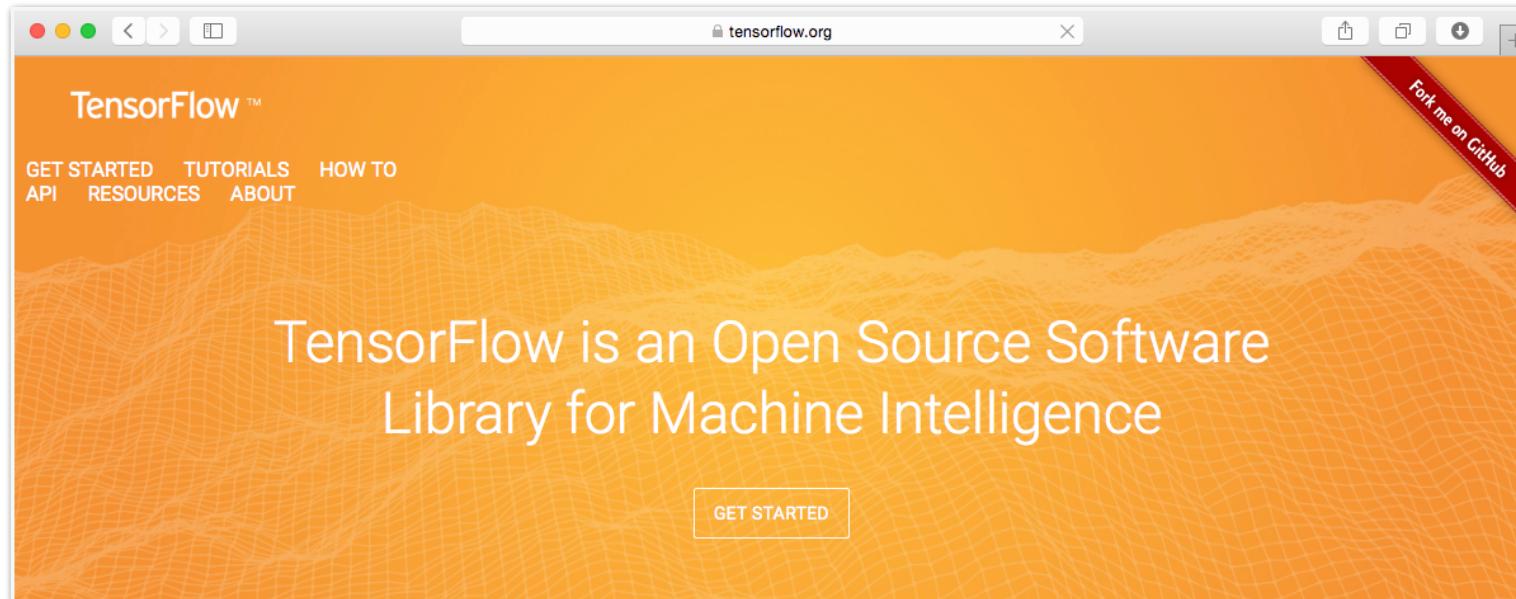
The screenshot shows a Google search results page with the following details:

- Search Query:** Tensorflow BARCELONA
- Number of Results:** Aproximadament 12.200 resultats (0,36 segons)
- Language Preference:** Consell: Cerca només resultats en català. Podeu especificar l'idioma de la cerca a Preferències
- Result 1:** **Barcelona TensorFlow Meetup (Barcelona) - Meetup**
www.meetup.com/.../Barcelona-TensorFlow-Me... ▾ Tradueix aquesta pàgina
We are waiting to have more members in this meetup before to start organizing activities. During this waiting period the activities will be held with the support of ...
dl., 1 febr. Barcelona convolucionada ... Sala d'actes de la FIB del ...
- Result 2:** **Nuevo libro de TensorFlow | Prof. Jordi Torres – UPC & BSC**
www.jorditorres.org/como-funcional-tensorflow/ ▾ Tradueix aquesta pàgina
20 gen. 2016 - Se acaba de publicar el libro "HELLO WORLD EN TENSORFLOW ...
Barcelona Spark Meetup Announces Changes in Leadership Team to ...
- Result 3:** **Introducción práctica al Deep Learning con TensorFlow de ...**
www.jorditorres.org/introduccion-practica-al-de... ▾ Tradueix aquesta pàgina
6 des. 2015 - el próximo febrero empezarán las actividades de un nuevo meetup en Barcelona dedicado a TensorFlow a través del cual pueden seguir los ...

TensorFlow?

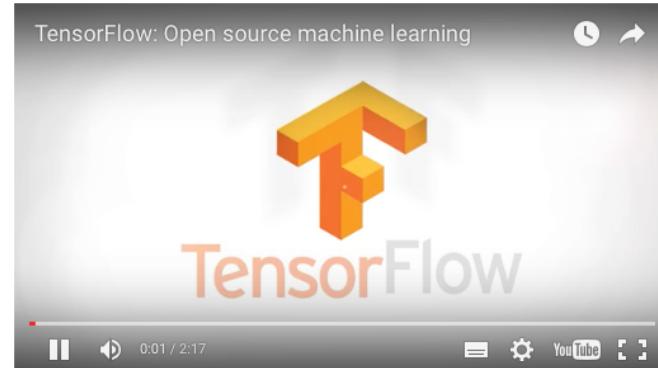


source: <https://www.youtube.com/watch?v=CMdHDHEuOUE>



About TensorFlow

TensorFlow™ is an open source software library for numerical computation using data flow graphs. Nodes in the graph represent mathematical operations, while the graph edges represent the multidimensional data arrays (tensors) communicated between them. The flexible architecture allows you to deploy computation to one or more CPUs or GPUs in a desktop, server, or mobile device with a single API. TensorFlow was originally developed by researchers and engineers working on the Google Brain Team within Google's Machine Intelligence research organization for the purposes of conducting machine learning and deep neural networks research, but the system is general enough to be applicable in a wide variety of other domains as well.

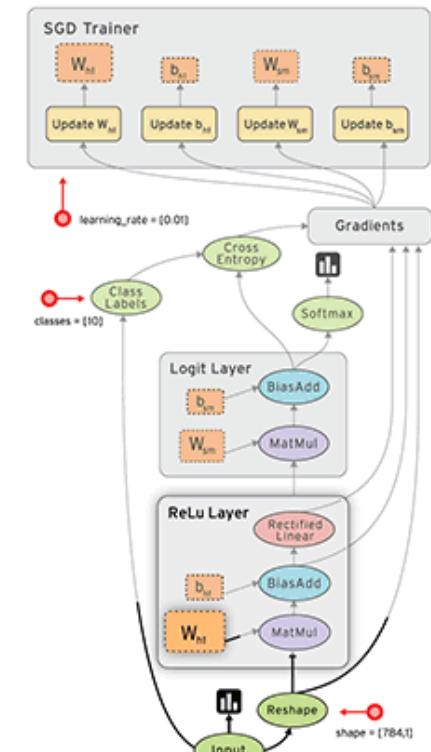


About TensorFlow

- TensorFlow was originally developed by the Google Brain Team within Google's Machine Intelligence research organization **for the purposes of conducting machine learning and deep neural networks research, but the system is general enough to be applicable in a wide variety of other domains as well.**
- Open source software library for numerical computation using data flow graphs.
- The flexible architecture allows you to deploy computation to one or more CPUs or GPUs in a desktop, server, or mobile device with a single API.

Data Flow Graph?

- Describe mathematical computation with a directed graph of nodes & edges.
 - ◆ Nodes in the graph represent mathematical operations,
 - ◆ Edges describe the i/o relationships between nodes.
 - ◆ Data edges carry dynamically-sized multidimensional data arrays, or **tensors**.
- The flow of tensors through the graph is where TensorFlow gets its name. Nodes are assigned to computational devices and execute asynchronously and in parallel once all the tensors on their incoming edges becomes available.



Before to start: Python basics

- Line indentation

Python has no mandatory statement termination characters and blocks are specified by indentation

Statements that expect an indentation level end in a colon (:).

The number of spaces in the indentation is variable, but all statements within the block must be indented the same amount.

Python will advise you if there is a unclear line indentation with the following warning:

IndentationError: unexpected indent

- Comments

Start with the pound (#) sign and are single-line

Before to start: Python basics

- Variables, operators and data types

```
>>> IntegerVar = 10
>>> IntegerVar += 10
>>> print IntegerVar
20
>>> StringVar = "Welcome"
>>> StringVar += " to Barcelona"
>>> print StringVar
Welcome to Barcelona
```

Before to start: Python basics

- Variables, operators and data types (cont.)

```
>>> IntegerVar, StringVar = StringVar,  
IntegerVar  
>>> print IntegerVar  
Welcome to Barcelona  
>>> print StringVar  
20  
>>> help (StringVar)  
Help on int object:  
class int(object)  
| int(x=0) -> int or long  
| int(x, base=10) -> int or long  
|
```

Before to start: Python basics

- Data Structures available in Python are lists, tuples and dictionaries.

Lists are like one-dimensional arrays and we can also have lists of other lists or tuples.
There are number of methods for lists (methods follow the list name). Tuples are immutable one-dimensional array.

```
>>> sampleList = [1,2,3,4,5,6,7,8]
>>> print (sampleList[1])
2
>>> sampleTuple = (1,2,3,4,5,6,7,8)
>>> print (sampleTuple)
(1, 2, 3, 4, 5, 6, 7, 8)
>>> print sampleList
[1, 2, 3, 4, 5, 6, 7, 8]
>>> sampleList.append(9)
>>> print sampleList
[1, 2, 3, 4, 5, 6, 7, 8, 9]
>>> sampleTuple.append(9)
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
AttributeError: 'tuple' object has no attribute 'append'
```

Before to start: Python basics

- Data Structures available in Python are lists, tuples and dictionaries (cont.)

Python dictionaries are associative arrays that has keys to obtain the elements of the dictionary. Dictionaries aren't exactly based on an index, there is no particular order in dictionaries (we could add a key: value and, it will appear in random places). A big caution here is that you cannot create different values with the same key (Python will just overwrite the value of the duplicate keys).

```
>>> myDicc = {'someItem': 2, 'otherItem': 20}  
>>> print(myDicc['otherItem'])  
20
```

Before to start: Python basics

- Data Structures available in Python are lists, tuples and dictionaries (cont.)

Python lists/dictionaries/tuples can be of any type, so you can mix them in. Variables can point to functions. The index of the first element is 0 and negative numbers count from the end towards the beginning (-1 is the last element).

```
>>> Example = [1, "BCN", ["another", "list"], {"and","a","tuple"}]
>>> print Example
[1, 'BCN', ['another', 'list'], set(['a', 'and', 'tuple'])]
>>> myfunction = len
>>> print myfunction (Example)
4
>>> print Example[-1]
set(['a', 'and', 'tuple'])
>>> print Example[3]
set(['a', 'and', 'tuple'])
```

Before to start: Python basics

- Flow control statements: **if**

Often partnered with the if statement are else if and else.

However, Python's else if is shortened into elif.

After every conditional we have a colon.

Next, we could proceed to a new line with number of spaces to tell Python we only want this code to be run when the previous conditional is satisfied.

```
>>> a = 20
>>> if a >= 22:
...     print("Paris")
... elif a >= 21:
...     print("Londres")
... else:
...     print("Barcelona")
...
Barcelona
>>>
```

Before to start: Python basics

- Flow control statements: **for, while**

Loops in Python are very versatile and simple.

```
>>> for a in range(1,4):
...     print (a)
...
1
2
3
>>> a = 1
>>> while a < 4:
...     print (a)
...     a+=1
...
1
2
3
>>> a = 20
```

Before to start: Python basics

- Functions

We can define a function by using the keyword def before your function name. Optional arguments are set in the function declaration after the mandatory arguments by being assigned a default value. Functions can return a tuple (and using tuple unpacking you can effectively return multiple values).

```
>>> def someFunction():
...     print("Barcelona")
...
>>> someFunction()
Barcelona
>>>
>>> def fib(n):      # write Fibonacci series up to n
...     a, b = 0, 1
...     while b < n:
...         print b,
...         a, b = b, a+b
...
>>> fib(1000)
1 1 2 3 5 8 13 21 34 55 89 144 233 377 610 987
```

Before to start: Python basics

- Lambda Functions

Python supports the creation of anonymous functions (i.e. functions that are not bound to a name) at runtime, using a construct called "lambda". This is not exactly the same as lambda in functional programming languages

```
>>> fibonacci = (lambda x: 1 if x <= 2 else
                  fibonacci(x-1) + fibonacci(x-2))
>>> fibonacci(10)
55
>>> foo = [2, 18, 9, 22, 17, 24, 8, 12, 27]
>>> for i in filter(lambda x: x % 3 == 0, foo):
    print (i)
18
9
24
12
27
```

Before to start: Python basics

- Classes

Python supports a limited form of multiple inheritance in classes.

```
>>> class Calculator(object):
...     #define class to simulate a calculator
...     def __init__(self):
...         #start with zero
...         self.current = 0
...     def add(self, amount):
...         #add number to current
...         self.current += amount
...     def getCurrent(self):
...         return self.current
...
>>> myCalc = Calculator() # make myCalc into a Calculator object
>>> myCalc.add(2) #use myCalc's new add method derived from Calculator class
>>> print(myCalc.getCurrent())
2
>>> myCalc.add(2)
>>> print(myCalc.getCurrent())
```

In the previous example the first part defines a Class. With def __init__(self) it is created a new instance of this class. The second part shows how to use this class in Python.

Before to start: Python basics

- Importing

External libraries are used with the `import [libname]` keyword. We can also use `from [libname] import [funcname]` for individual functions.

```
>>> from random import randint
>>> randomint = random.randint(1, 100)
>>> print randomint
84
>>> randomint = random.randint(1, 100)
>>> print randomint
44
```

Before to start: Python basics

- Read/Write files

Python uses the following syntax for read/write files:

```
>>> f = open("test.txt","w") #opens file with name of "test.txt"
>>> f.write("Barcelona, ")
>>> f.write("is the best city of the world.")
>>> f.write("With an excellent weather.")
>>> f.close()
>>>
>>> f = open("test.txt","r") #opens file with name of "test.txt"
>>> print(f.read())
Barcelona, is the best city of the world.With an excellent weather.
>>> f.close
```

Before to start: Python basics

- Iterators

Python define a object type for taking each item of something, one after another. Any time you use a loop, explicit or implicit, to go over a group of items, that is iteration

An iterable is an object that has an `__iter__` method which returns an iterator, or which defines a `__getitem__` method that can take sequential indexes starting from zero.

An iterator is an object with a `next` (Python 2) or `__next__` (Python 3) method.

```
>>> vec = [1, 2, 3]
>>> it = iter(vec)
>>> next(it)
1
>>> next(it)
2
>>> next(it)
3
>>> type(vec)
<class 'list'>
>>> type(it)
<class 'list_iterator'>
```

In this example vec iterable and it es a iterador.

Before to start: Python basics

- More information if necessary?

<https://docs.python.org>

Let's start!

```
# Ubuntu/Linux 64-bit  
$ sudo apt-get install python-pip python-dev python-virtualenv  
  
# Mac OS X  
$ sudo easy_install pip  
$ sudo pip install --upgrade virtualenv
```

Available at: https://www.tensorflow.org/versions/master/get_started/os_setup.html#download-and-setup [Accessed: 16/12/2015].

Virtualenv

```
$ virtualenv --system-site-packages ~/tensorflow  
  
$ source ~/tensorflow/bin/activate # si se usa bash  
$ source ~/tensorflow/bin/activate.csh # si se usa csh  
(tensorflow)$
```

Available at: https://www.tensorflow.org/versions/master/get_started/os_setup.html#download-and-setup [Accessed: 16/12/2015].

How to start?

```
# Ubuntu/Linux 64-bit, CPU only:  
(tensorflow)$ pip install --upgrade https://storage.googleapis.com/tensorflow/linux/cpu/tensorflow-0.5.0-cp27-none-linux_x86_64.whl  
  
# Mac OS X, CPU only:  
(tensorflow)$ pip install --upgrade https://storage.googleapis.com/tensorflow/mac/tensorflow-0.5.0-py2-none-any.whl
```

Available at: https://www.tensorflow.org/versions/master/get_started/os_setup.html#download-and-setup [Accessed: 16/12/2015].

Deactivate

```
(tensorflow)$ deactivate
```

Available at: https://www.tensorflow.org/versions/master/get_started/os_setup.html#download-and-setup [Accessed: 16/12/2015].

My first code

```
import tensorflow as tf

a = tf.placeholder("float")
b = tf.placeholder("float")

y = tf.mul(a, b)

sess = tf.Session()

print sess.run(y, feed_dict={a: 3, b: 3})
```

Available: <https://github.com/jorditorresBCN/TutorialTensorFlow>

Hands-on 1

1. Download the code from:
<https://github.com/jorditorresBCN/TutorialTensorFlow>
2. Follow teacher's instructions

math operations for manipulate the tensors

Operation	Description
tf.add	sum
tf.sub	subtraction
tf.mul	multiplication
tf.div	division
tf.mod	module
tf.abs	return the absolute value
tf.neg	return negative value

math operations for manipulate the tensors

tf.sign	return the sign
tf.inv	returns the inverse
tf.square	calculates the square
tf.round	returns the nearest integer
tf.sqrt	calculates the square root
tf.pow	calculates the power
tf.exp	calculates the exponential
tf.log	calculates the logarithm
tf.maximum	returns the maximum
tf.minimum	returns the minimum
tf.cos	calculates the cosine
tf.sin	calculates the sine

math operations on matrices

Operation	Description
<code>tf.diag</code>	returns a diagonal tensor with a given diagonal values
<code>tf.transpose</code>	returns the transposes of the argument
<code>tf.matmul</code>	returns a tensor product of multiplying two tensors listed as arguments
<code>tf.matrix_determinant</code>	returns the determinant of the square matrix specified as an argument
<code>tf.matrix_inverse</code>	returns the inverse of the square matrix specified as an argument

Operations & Kernels

- An operation has a name and represents an abstract computation (e.g., “matrix multiply”, or “add”). An operation can have attributes
- A kernel is a particular implementation of an operation that can be run on a particular type of device (e.g., CPU or GPU).
- A TensorFlow binary defines the sets of operations and kernels available

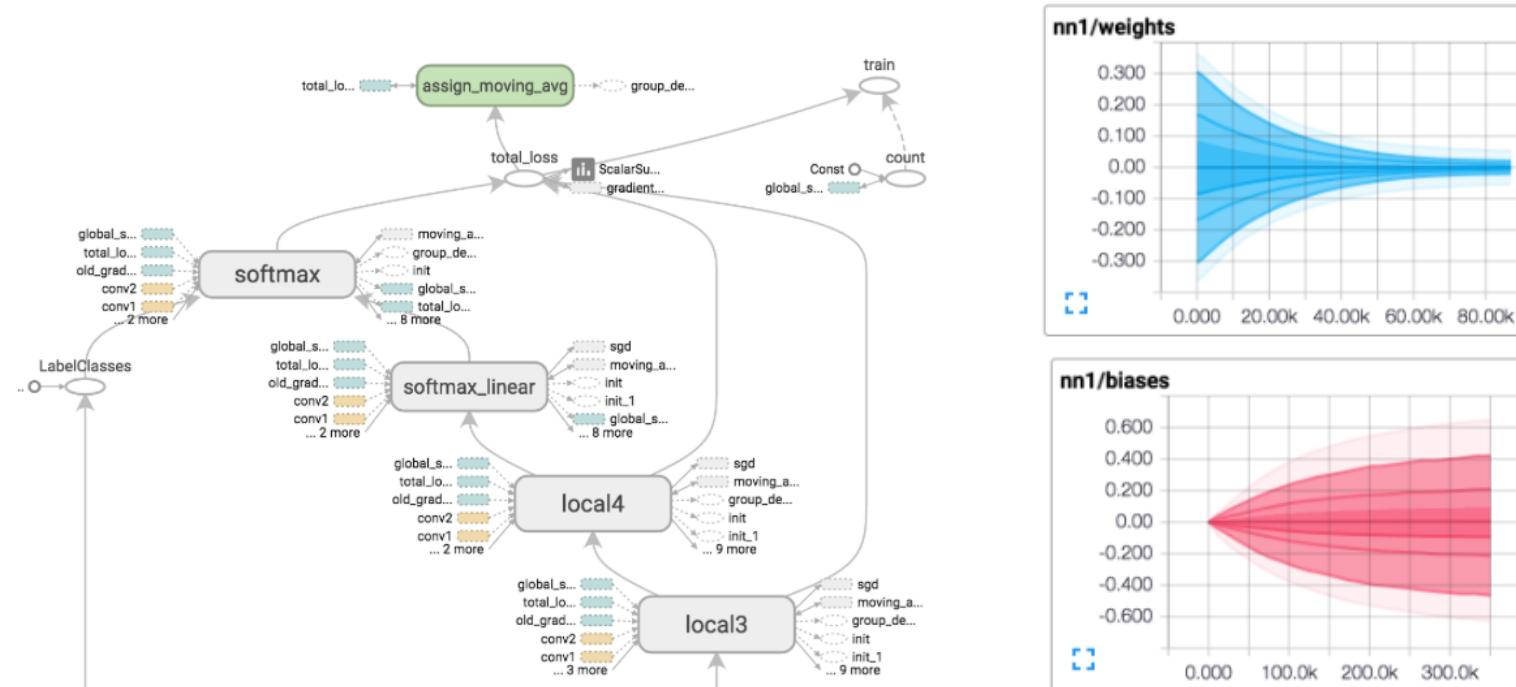
kernels

Operations groups	Operations
Maths	Add, Sub, Mul, Div, Exp, Log, Greater, Less, Equal
Array	Concat, Slice, Split, Constant, Rank, Shape, Shuffle
Matrix	MatMul, MatrixInverse, MatrixDeterminant
Neuronal Network	SoftMax, Sigmoid, ReLU, Convolution2D, MaxPool
Checkpointing	Save, Restore
Queues and sincronizations	Enqueue, Dequeue, MutexAcquire, MutexRelease
Flow control	Merge, Switch, Enter, Leave, NextIteration

TensorFlow: Large-scale machine learning on heterogeneous systems, (2015). [Online]. Available at: <http://download.tensorflow.org/paper/whitepaper2015.pdf> [Accessed: 20/12/2015].

TensorBoard: Visualization of graph structures and summary statistics

- In order to help users understand the structure of their computation graphs and also to understand the overall behavior of machine learning models.



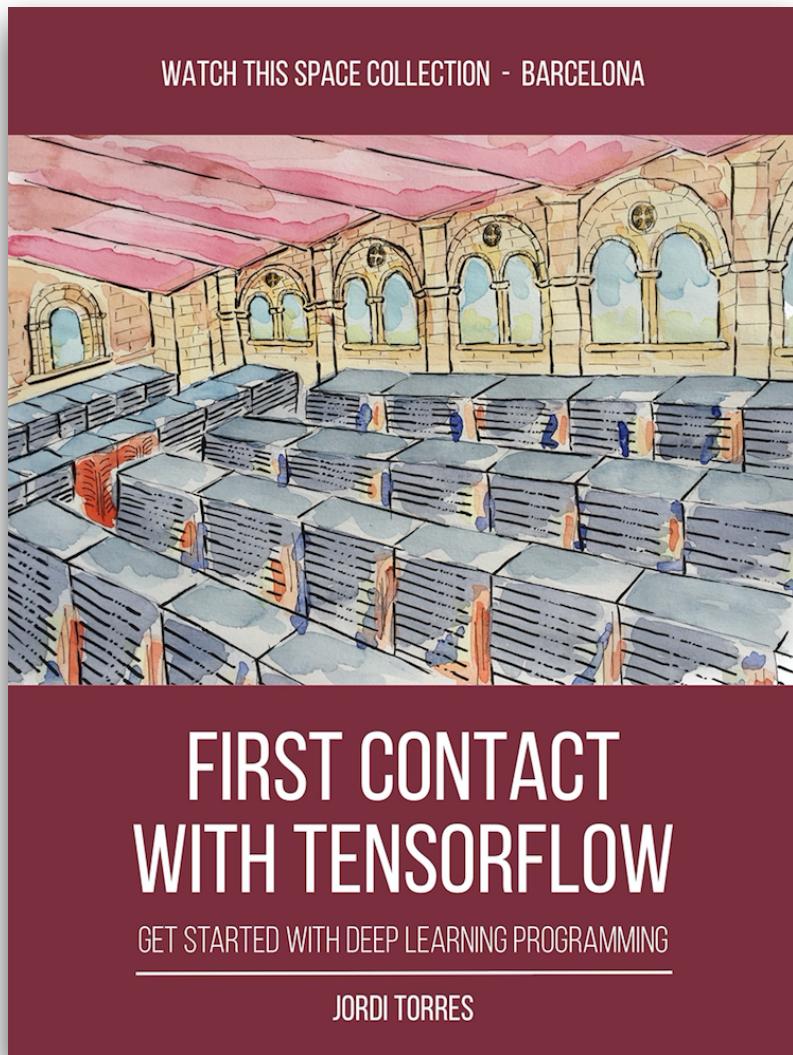
Tensorboard

- TensorFlow included functions to debug and optimize programs in a visualization tool called TensorBoard.
- The data displayed with TensorBoard module are generated during the execution of TensorFlow and stored in trace files whose data are obtained from the summary operations
- The way we can invoke it is very simple: <http://localhost:6006/>

(out of scope of this presentation)

- More information: TensorFlow, (2016) TensorBoard: Graph Visualization. [Online]. Available at: https://www.tensorflow.org/versions/master/how_tos/graph_viz/index.html [Accessed: 02/01/2016].

Content:



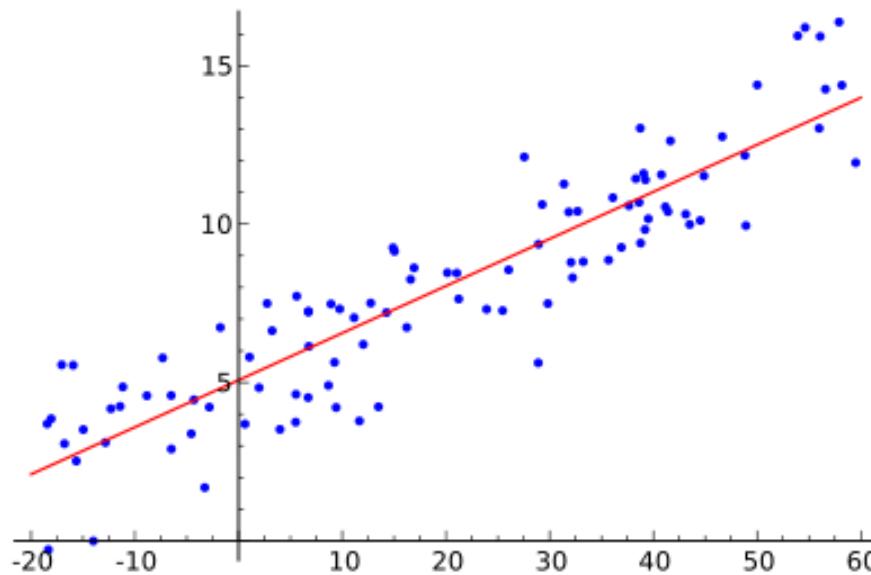
Introduction

1. TensorFlow Basics
2. Linear Regression
3. Clustering
4. Single Layer Neural Network
5. Multi-layer Neural Networks
6. TensorFlow and GPUs
- Closing



Linear Regression

- Wikipedia: “In **statistics**, **linear regression** is an approach for modeling the relationship between a scalar **dependent variable** y and one or more **explanatory variables** (or independent variables) denoted X . “



Linear Regression

- Linear regression is a statistical technique used to measure the relationship between variables.
- a simple linear regression can be $y = W * x + b$.
- Case Study: $y = 0.1 * x + 0.3$

Case study

For this, we use a simple Python program that creates data in two-dimensional space

```
import numpy as np

num_puntos = 1000
conjunto_puntos = []
for i in xrange(num_puntos):
    x1= np.random.normal(0.0, 0.55)
    y1= x1 * 0.1 + 0.3 + np.random.normal(0.0, 0.03)
    conjunto_puntos.append([x1, y1])

x_data = [v[0] for v in conjunto_puntos]
y_data = [v[1] for v in conjunto_puntos]
```

Available: <https://github.com/jorditorresBCN/TutorialTensorFlow>

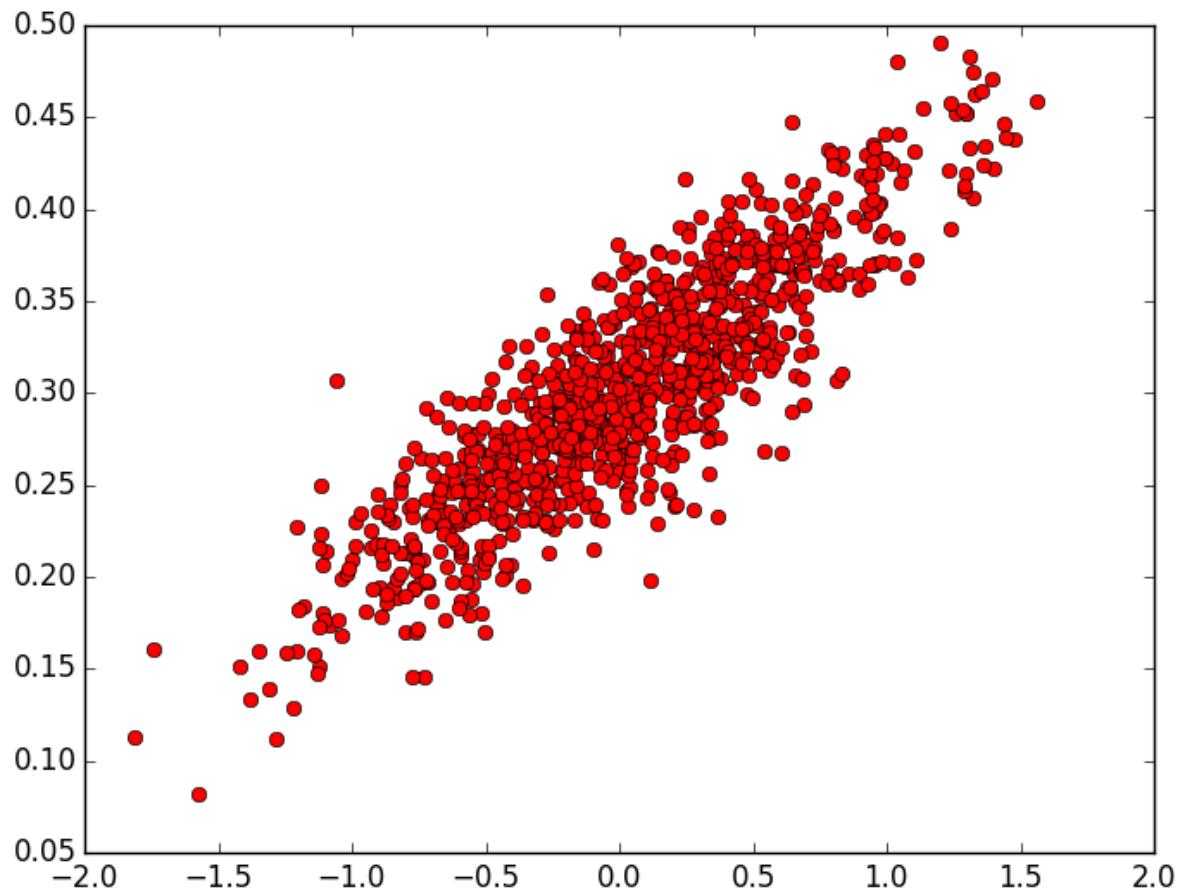
Case study

```
import matplotlib.pyplot as plt

plt.plot(x_data, y_data, 'ro', label='Original data')
plt.legend()
plt.show()
```

Available: <https://github.com/jorditorresBCN/TutorialTensorFlow>

Case study



We will ask TensorFlow looking for the line that best fits these points

```
import tensorflow as tf

W = tf.Variable(tf.random_uniform([1], -1.0, 1.0))
b = tf.Variable(tf.zeros([1]))
y = W * x_data + b

loss = tf.reduce_mean(tf.square(y - y_data))
optimizer = tf.train.GradientDescentOptimizer(0.5)
train = optimizer.minimize(loss)

init = tf.initialize_all_variables()

sess = tf.Session()
sess.run(init)

for step in xrange(16):
    sess.run(train)
```

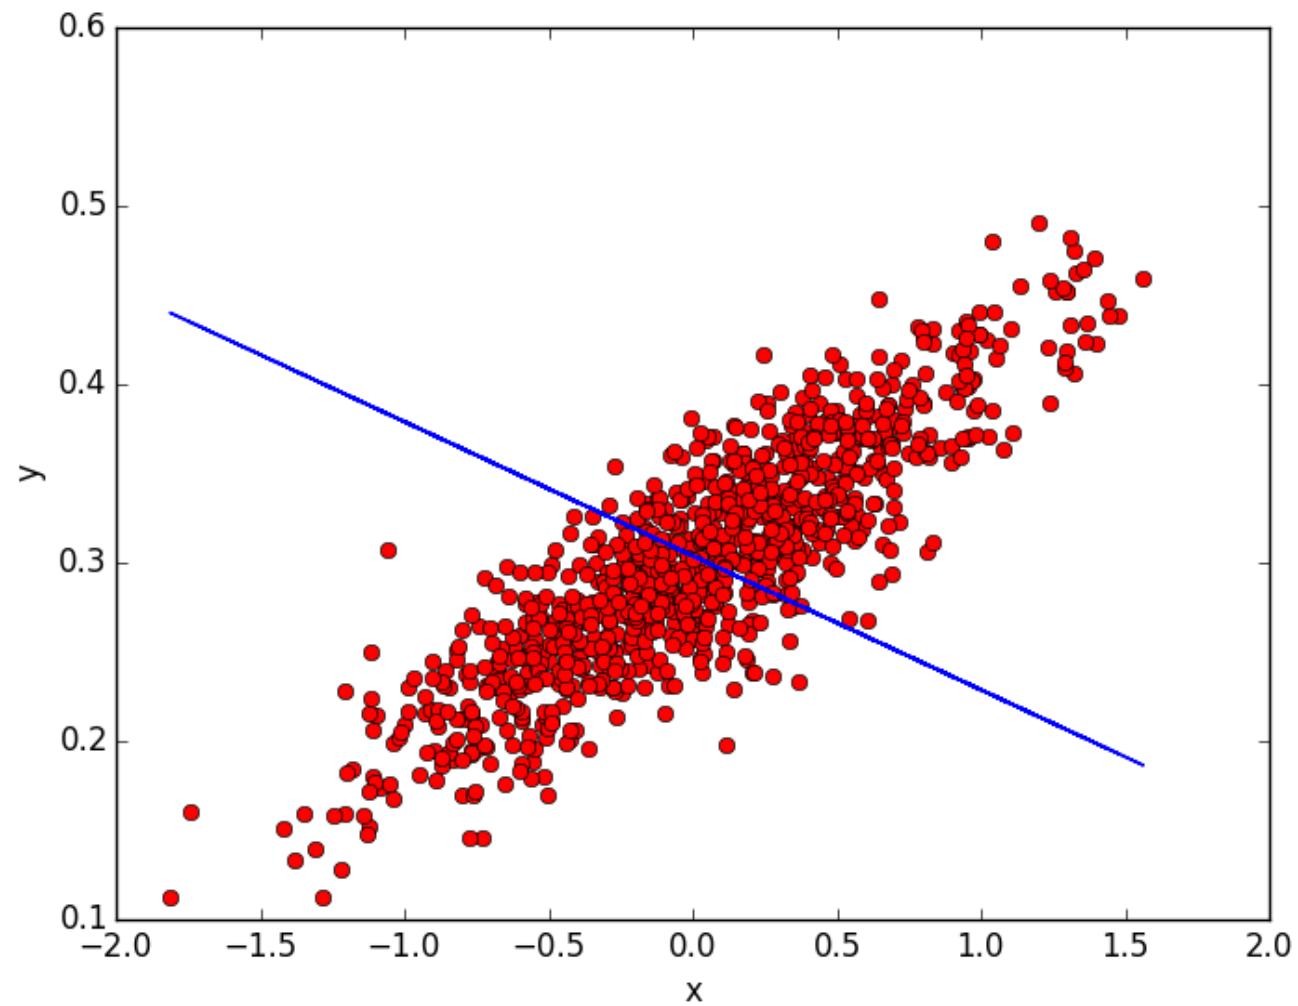
Available: <https://github.com/jorditorresBCN/TutorialTensorFlow>

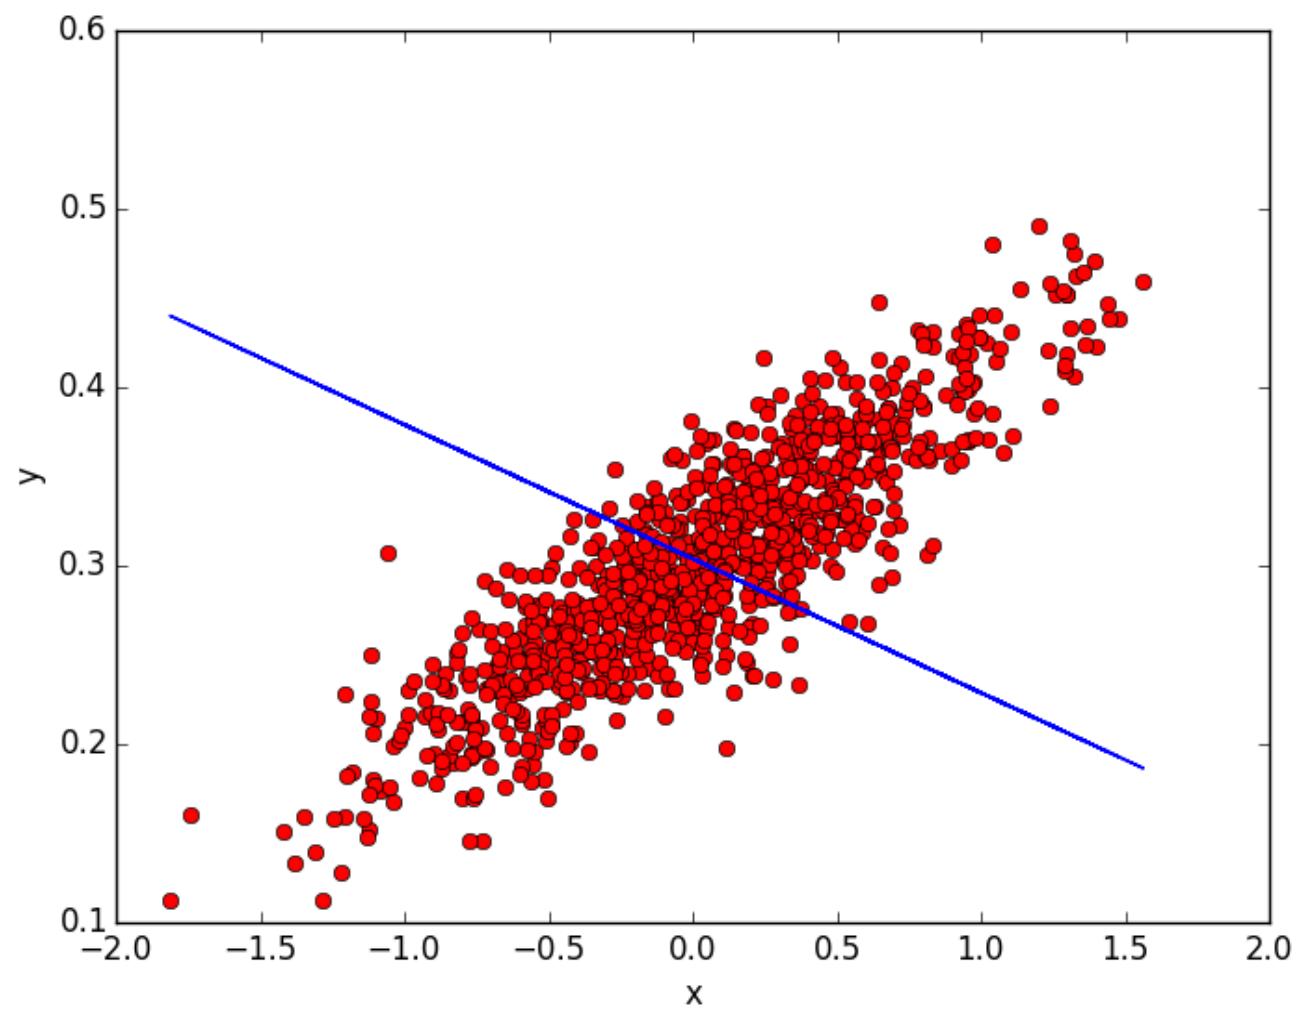
Case Study: Running the Algorithm

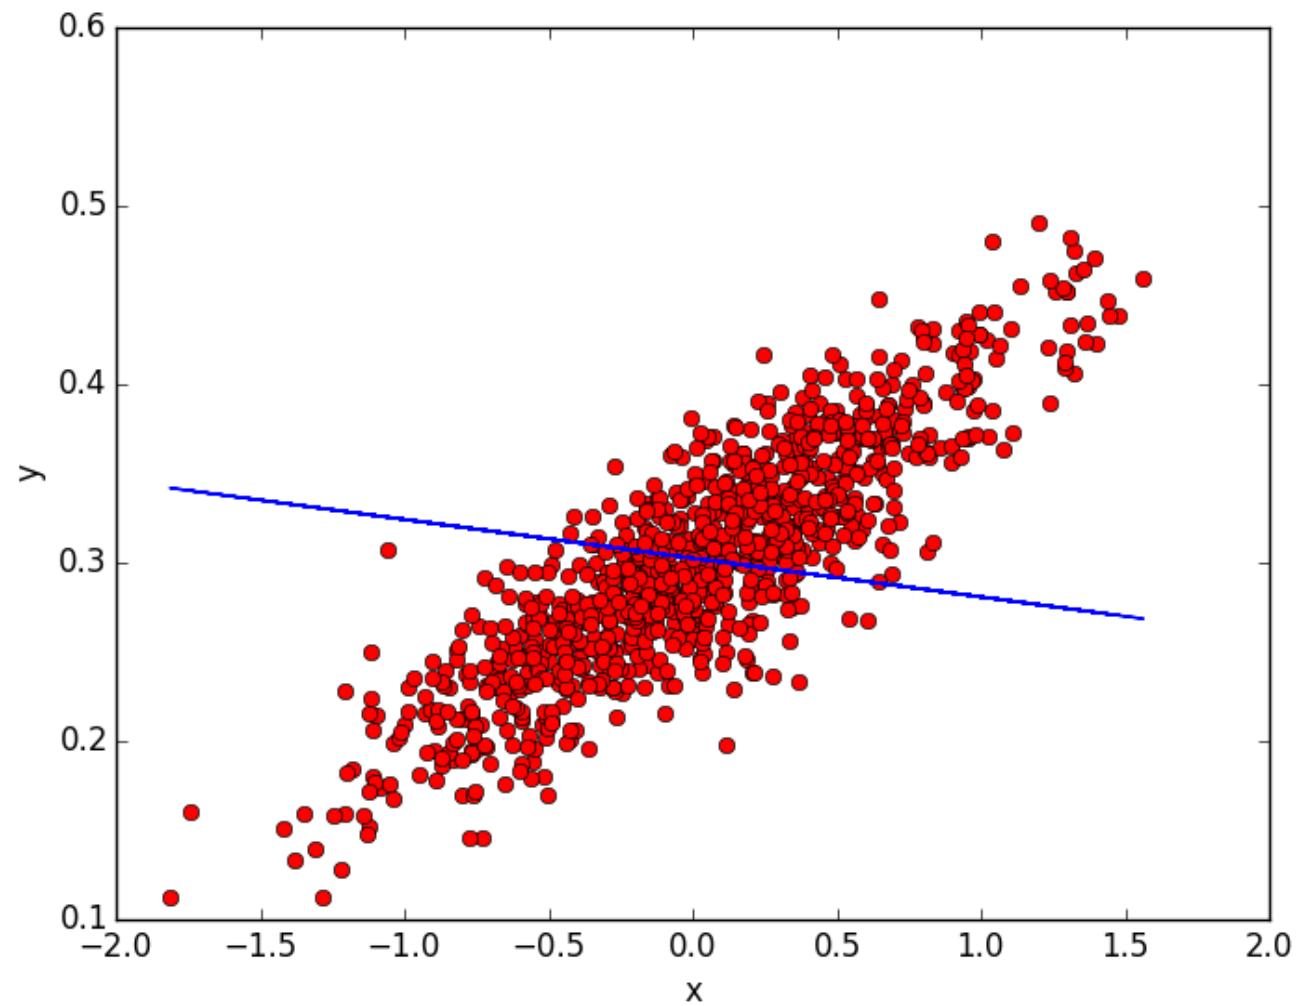
```
for step in xrange(16):
    sess.run(train)

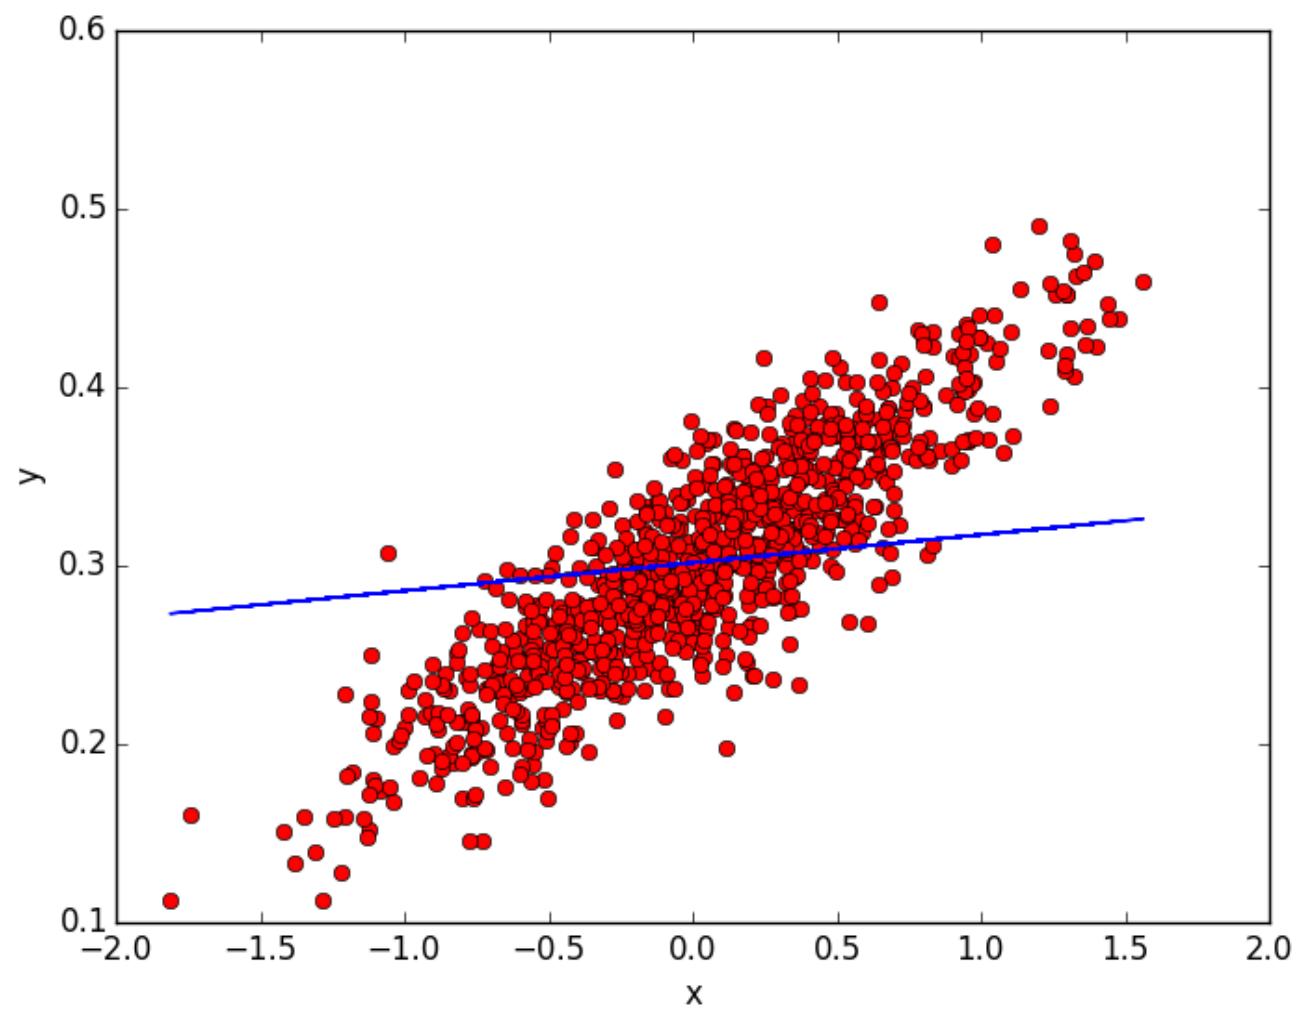
    plt.plot(x_data, y_data, 'ro')
    plt.plot(x_data, sess.run(W) * x_data + sess.run(b))
    plt.xlabel('x')
    plt.xlim(-2,2)
    plt.ylim(0.1,0.6)
    plt.ylabel('y')
    plt.legend()
    plt.show()
```

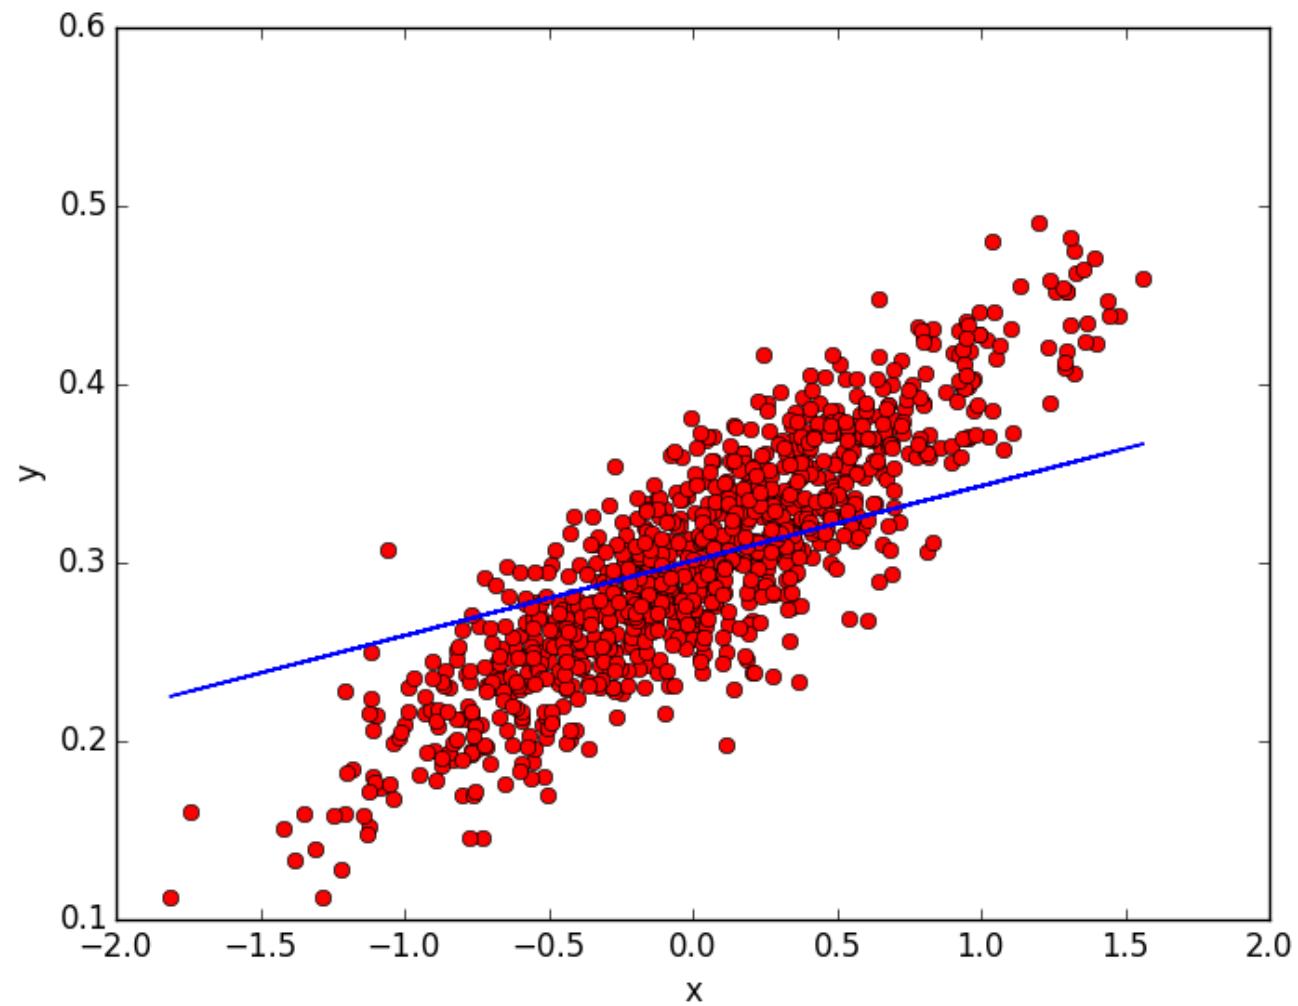
Disponible en: <https://github.com/jorditorresBCN/TutorialTensorFlow>

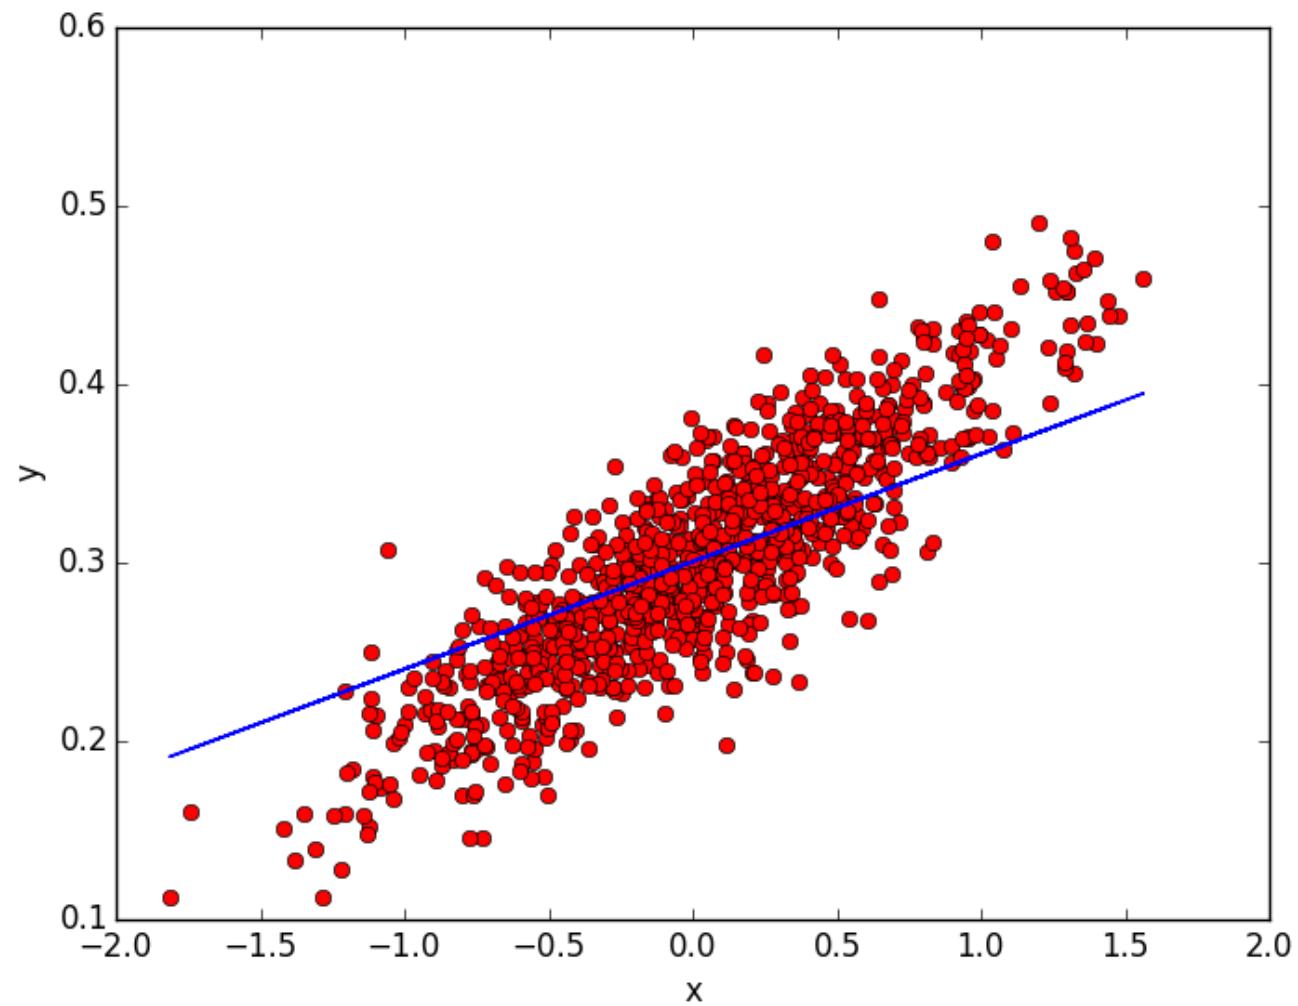


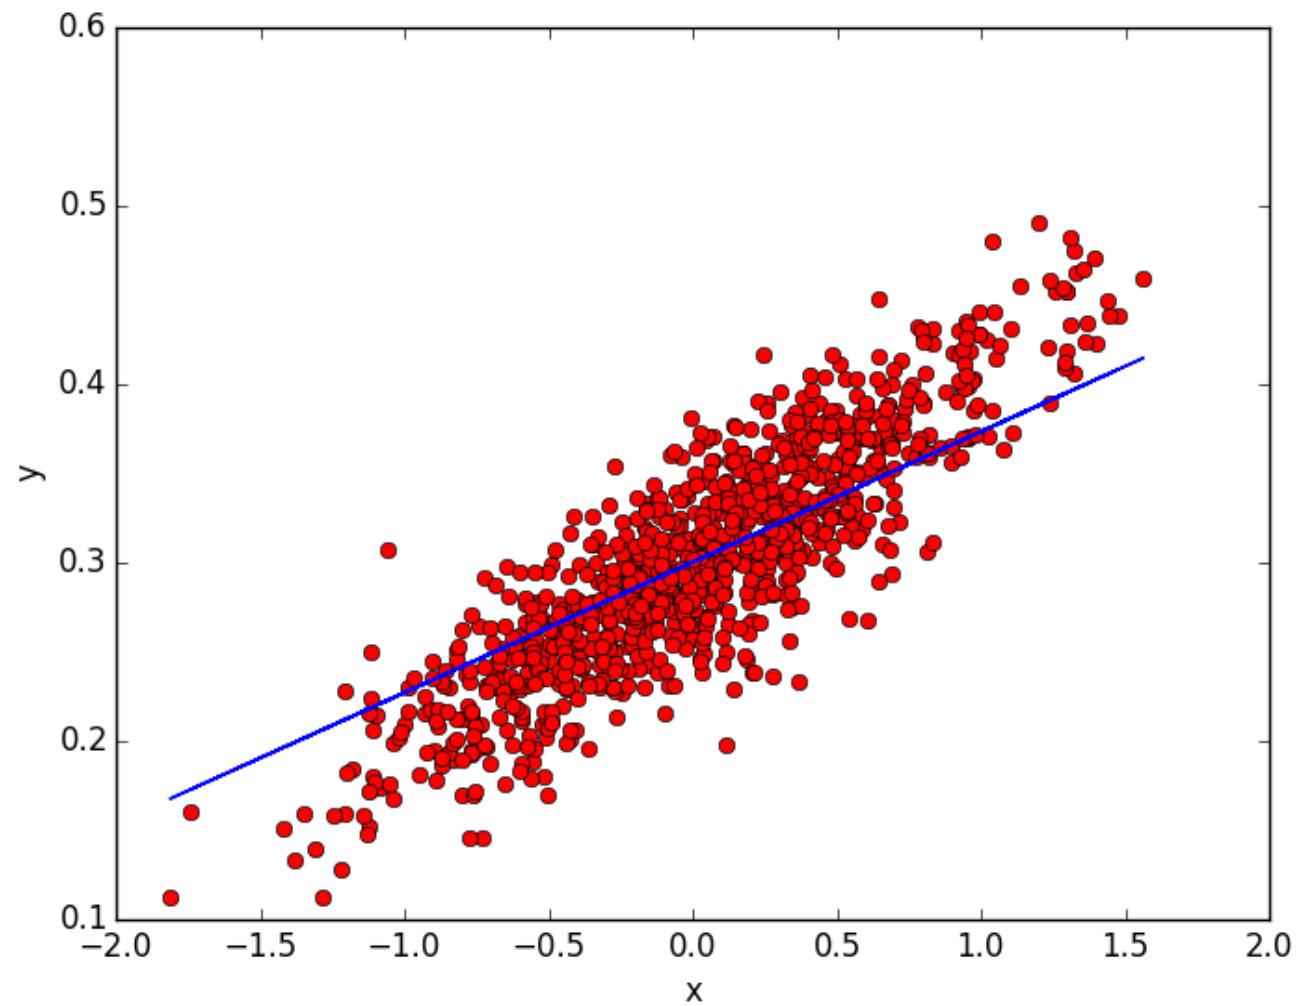


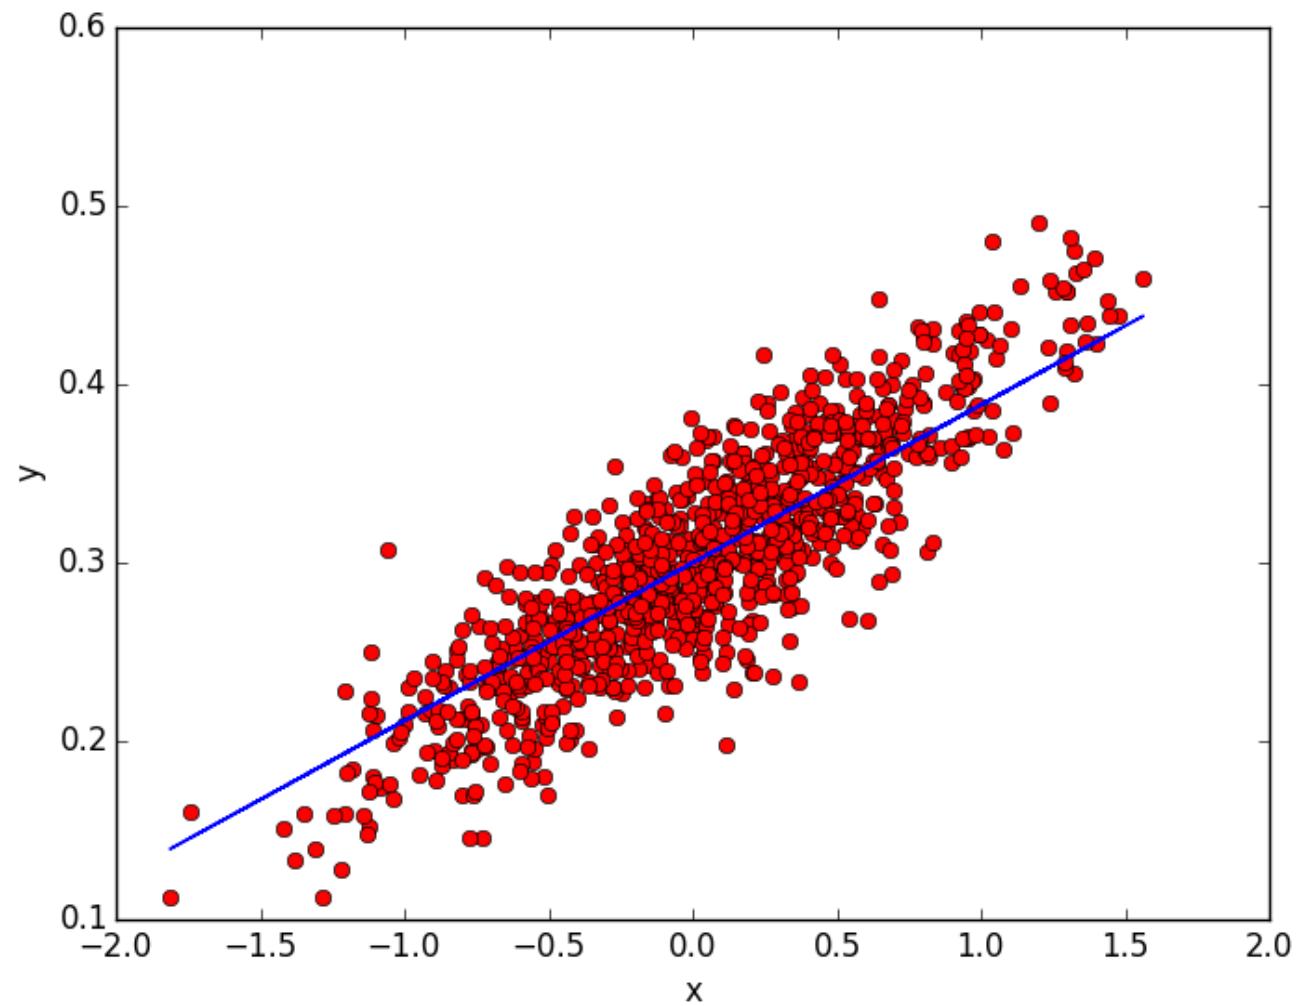


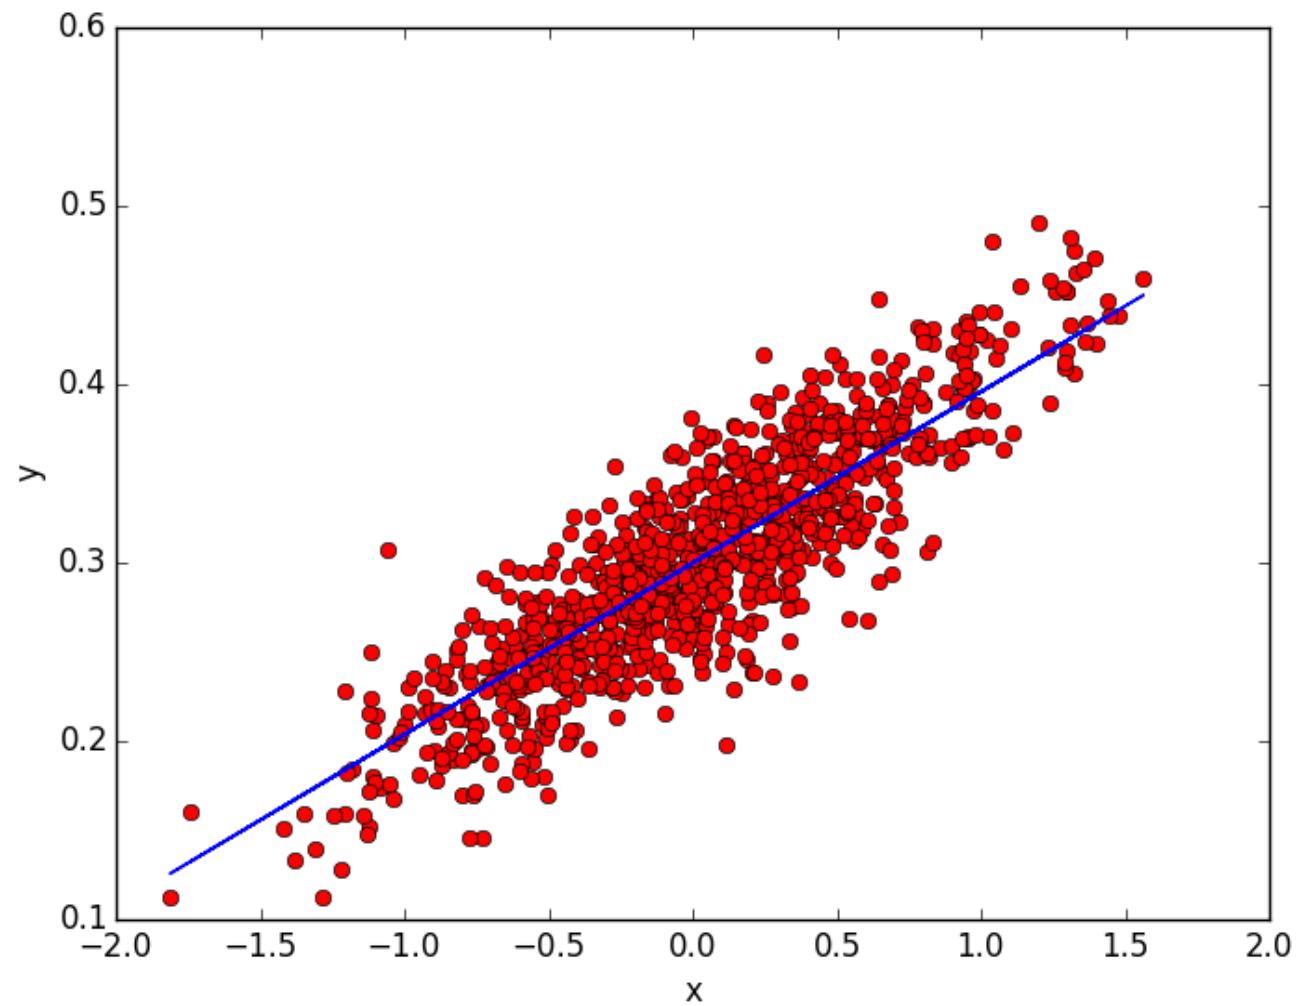


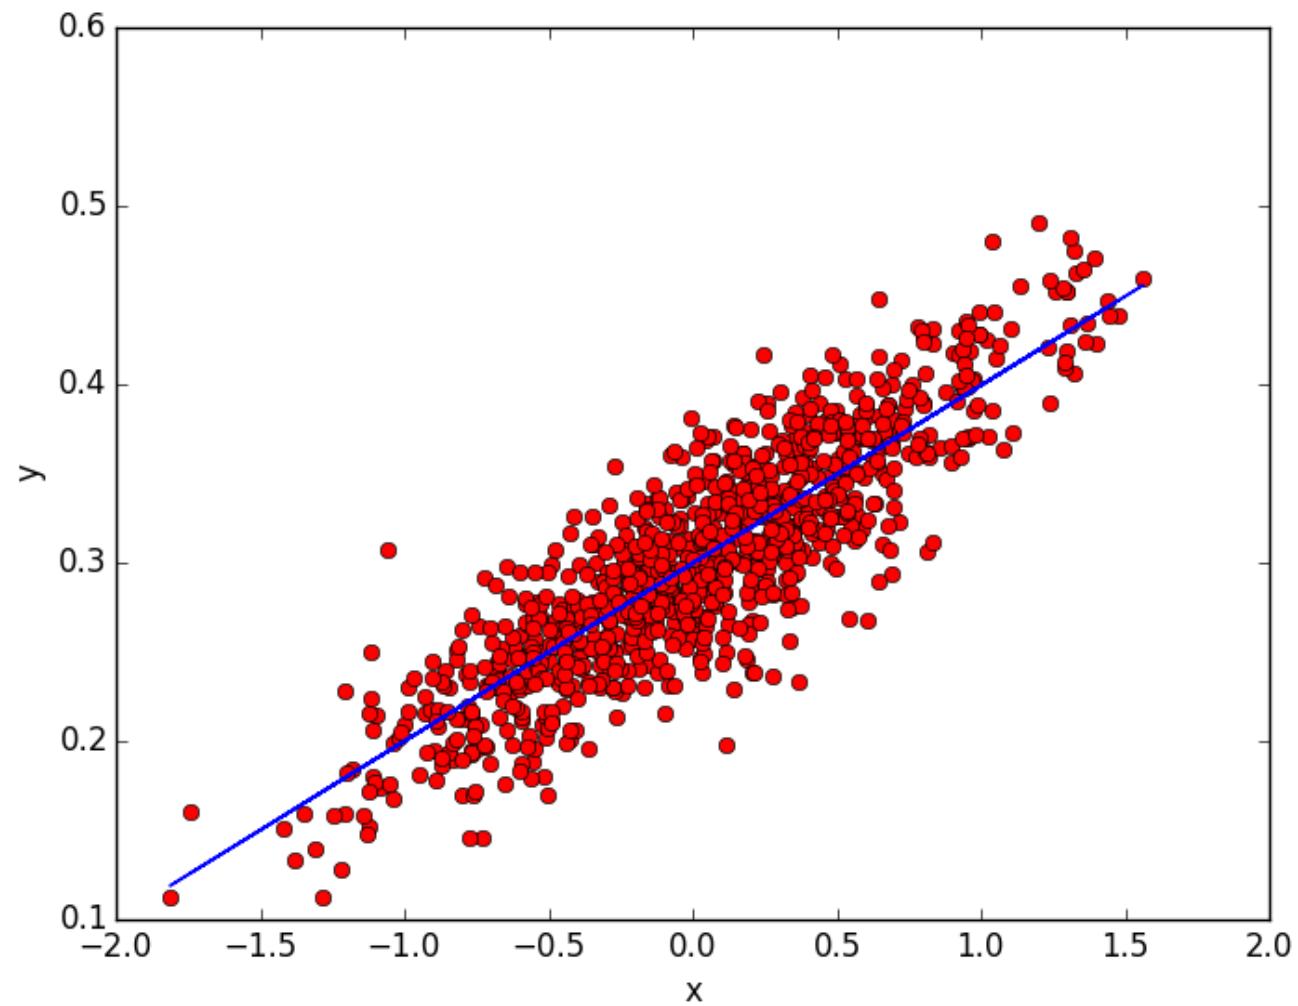


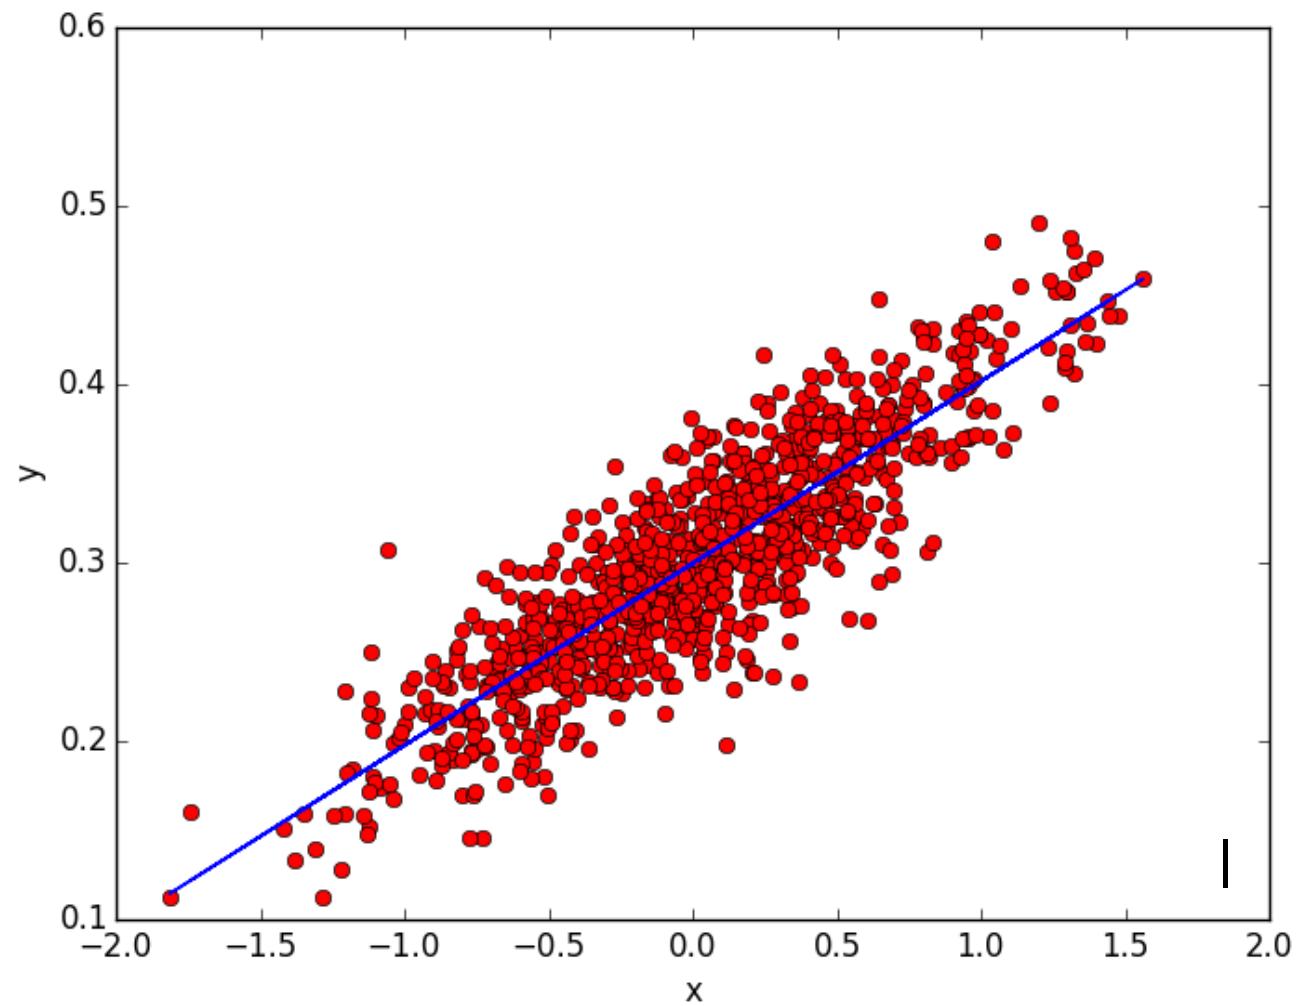












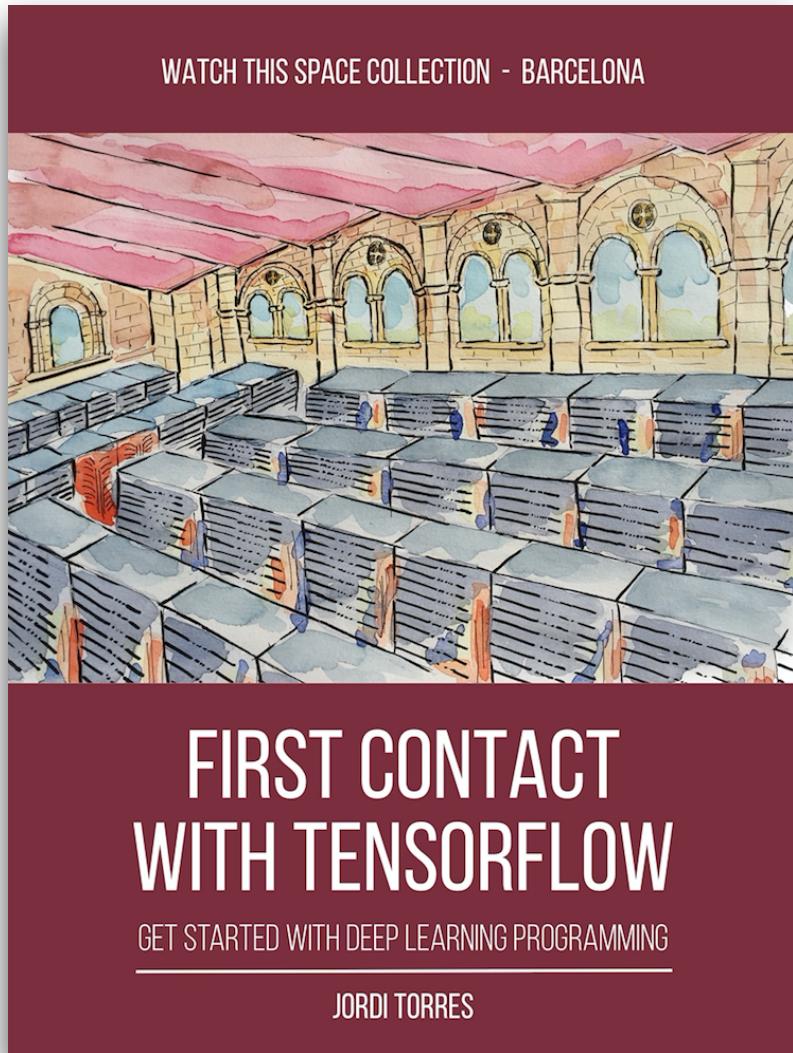
```
print(step, sess.run(W), sess.run(b))
```

```
(0, array([-0.04841119], dtype=float32), array([ 0.29720169], dtype=float32))
(1, array([-0.00449257], dtype=float32), array([ 0.29804006], dtype=float32))
(2, array([ 0.02618564], dtype=float32), array([ 0.29869056], dtype=float32))
(3, array([ 0.04761609], dtype=float32), array([ 0.29914495], dtype=float32))
(4, array([ 0.06258646], dtype=float32), array([ 0.29946238], dtype=float32))
(5, array([ 0.07304412], dtype=float32), array([ 0.29968411], dtype=float32))
(6, array([ 0.08034936], dtype=float32), array([ 0.29983902], dtype=float32))
(7, array([ 0.08545248], dtype=float32), array([ 0.29994723], dtype=float32))
```

Hands-on 2

1. Download the code from:
<https://github.com/jorditorresBCN/TutorialTensorFlow>
2. Follow teacher's instructions

Content:



Introduction

1. TensorFlow Basics
 2. Linear Regression
 - 3. Clustering**
 4. Single Layer Neural Network
 5. Multi-layer Neural Networks
 6. TensorFlow and GPUs
- Closing



Basic data type: Tensor

Tensors can be considered a dynamically-sized multidimensional data arrays.

Type in TensorFlow	Type in Python	Description
DT_FLOAT	tf.float32	Floating point of 32 bits
DT_INT16	tf.int16	Integer of 16 bits
DT_INT32	tf.int32	Integer of 32 bits
DT_INT64	tf.int64	Integer of 64 bits
DT_STRING	tf.string	<i>String</i>
DT_BOOL	tf.bool	<i>Boolean</i>

Tensor Shape, Rank and Dimension

Shape	Rank	Dimension Number
[]	0	0-D
[D0]	1	1-D
[D0, D1]	2	2-D
[D0, D1, D2]	3	3-D
...
[D0, D1, ... Dn]	n	n-D

Example: Tensor rank 2

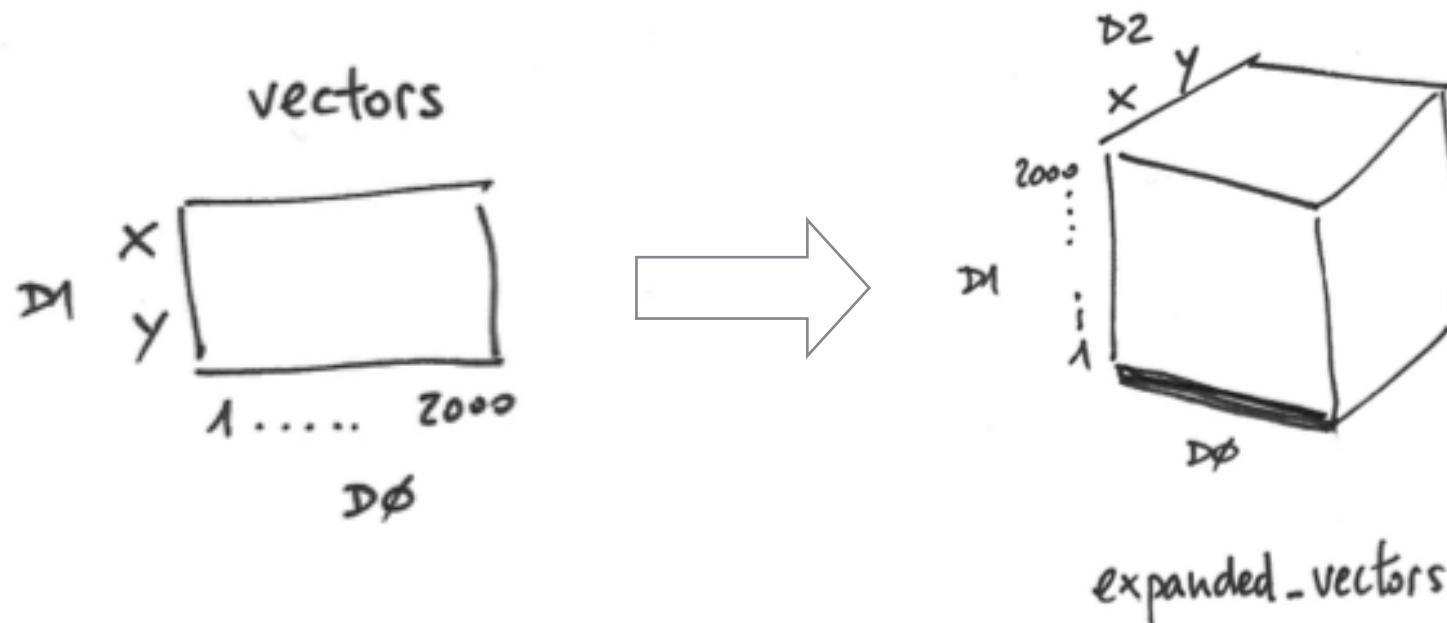
```
t = [[1, 2, 3], [4, 5, 6], [7, 8, 9]]
```

Tensor Transformations

Operación	Descripción
tf.shape	Para saber la forma de un tensor
tf.size	Para saber el número de elementos de un tensor
tf.rank	Para saber el rango de un tensor
tf.reshape	Permite cambiar la forma de un tensor manteniendo los mismos
tf.squeeze	Permite borrar del tensor dimensiones de tamaño 1
tf.expand_dims	Permite insertar una dimensión al tensor
tf.slice	Extrae una porción del tensor
tf.split	Divide el tensor en varios tensores a lo largo de una dimensión
tf.tile	Crea un nuevo tensor replicando un tensor múltiples veces
tf.concat	Concatena tensores en una de las dimensiones
tf.reverse	Invierte una determinada dimensión de un tensor
tf.transpose	Permuta dimensiones de un tensor
tf.gather	Recolecta porciones de acuerdo a un índice

Tensor Transformation example:

```
vectors = tf.constant(conjunto_puntos)
extended_vectors = tf.expand_dims(vectors, 0)
```



Data Storage in TensorFlow

1. Data files

example: https://github.com/jorditorresBCN/LibroTensorFlow/blob/master/input_data.py

2. Memory:

- Constants *tf.constant(...)*
- Variables *tf.Variable(...)*

3. Python code : Placeholders

- *feed_dict* parameter

Constants generation

Operation	Description
<code>tf.zeros_like</code>	Crea un <i>tensor</i> con todos los elementos inicializados a 0
<code>tf.ones</code>	Crea un <i>tensor</i> con todos los elementos inicializados a 1
<code>tf.ones_like</code>	Crea un <i>tensor</i> con todos los elementos inicializados a 1
<code>tf.fill</code>	Crea un <i>tensor</i> con todos los elementos inicializados a un valor <i>scalar</i> pasado como
<code>tf.constant</code>	Crea un <i>tensor</i> de constantes con los elementos indicados como argumento

Tensor random generation

Operation	Description
<code>tf.random_normal</code>	Valores <i>random</i> con una distribución normal, con una media y desviación estándar indicadas como argumentos
<code>tf.truncated_normal</code>	Igual que la anterior pero eliminando aquellos valores cuya magnitud es más de 2 veces la desviación estándar.
<code>tf.random_uniform</code>	Valores <i>random</i> con una distribución uniforme de un rango indicado en los argumentos
<code>tf.random_shuffle</code>	Mezclados aleatoriamente los elementos de un <i>tensor</i> en su primera dimensión
<code>tf.set_random_seed</code>	Establece la semilla aleatoria a nivel de grafo

Placeholders

- With calls `Session.run()` or `Tensor.eval()`

```
import tensorflow as tf

a = tf.placeholder("float")
b = tf.placeholder("float")

y = tf.mul(a, b)

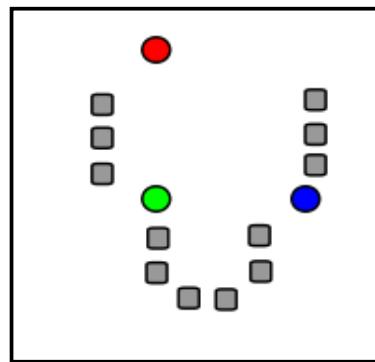
sess = tf.Session()

print sess.run(y, feed_dict={a: 3, b: 3})
```

K-means algorithm

(Example from Wikipedia)

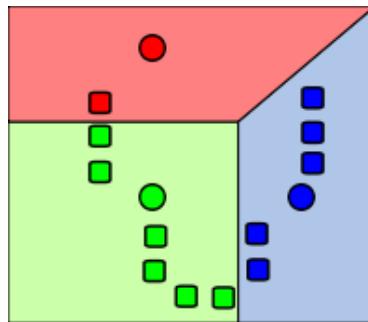
1. k initial "means" (in this case $k=3$) are randomly generated within the data domain (shown in color)



K-means algorithm

(Example from Wikipedia)

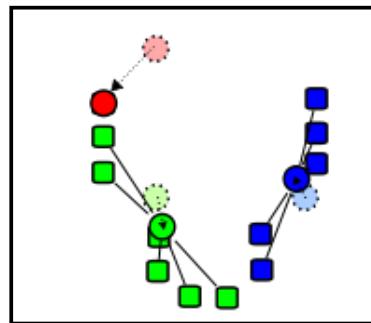
2. k clusters are created by associating every observation with the nearest mean. The partitions here represent the [Voronoi diagram](#) generated by the means.



K-means algorithm

(Example from Wikipedia)

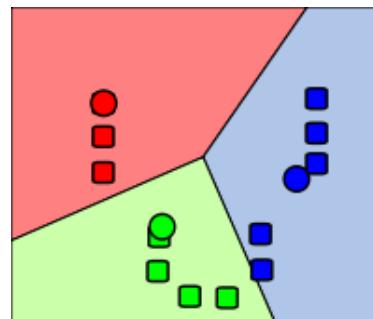
3. The **centroid** of each of the k clusters becomes the new mean. k clusters are created by associating every observation with the nearest mean.



K-means algorithm

(Example from Wikipedia)

4. Steps 2 and 3 are repeated until convergence has been reached.



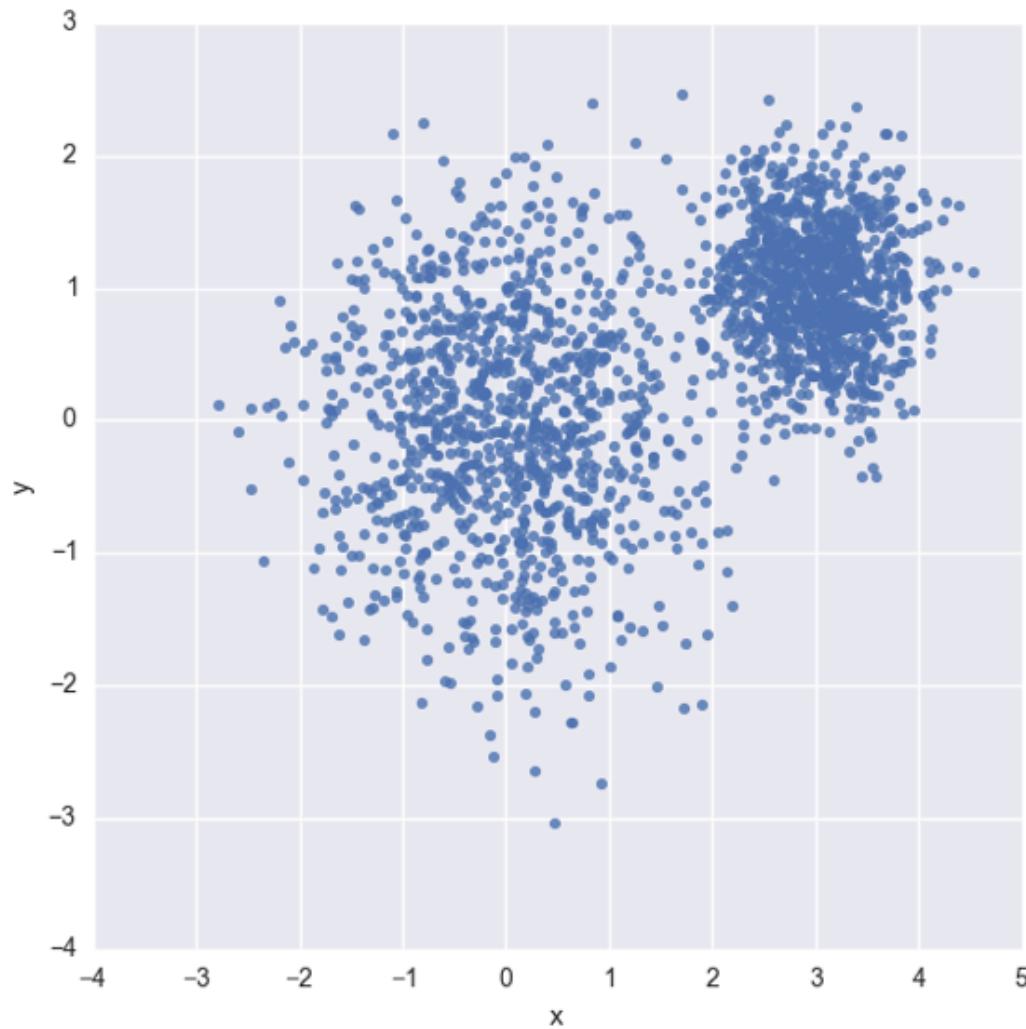
K-means algorithm

As it is a heuristic algorithm, there is no guarantee that it will converge to the global optimum, and the result may depend on the initial clusters. As the algorithm is usually very fast, it is common to run it multiple times with different starting conditions.

K-means: Case study

```
num_puntos = 2000
conjunto_puntos = []
for i in xrange(num_puntos):
    if np.random.random() > 0.5:
        conjunto_puntos.append([np.random.normal(0.0, 0.9),
                               np.random.normal(0.0, 0.9)])
    else:
        conjunto_puntos.append([np.random.normal(3.0, 0.5),
                               np.random.normal(1.0, 0.5)])
```

K-means: Case study



K-means in TensorFlow

```
vectores = tf.constant(conjunto_puntos)
k = 4
centroides = tf.Variable(tf.slice(tf.random_shuffle(vectores),[0,0],[k,-1]))
expanded_vectors = tf.expand_dims(vectores, 0)
expanded_centroides = tf.expand_dims(centroides, 1)

assignments = tf.argmin(tf.reduce_sum(tf.square(tf.sub(expanded_vectors,
            expanded_centroides)), 2), 0)

means = tf.concat(0, [tf.reduce_mean(tf.gather(vectores,
    tf.reshape(tf.where( tf.equal(assignments, c)),[1,-1])),
    reduction_indices=[1]) for c in xrange(k)])

update_centroides = tf.assign(centroides, means)

init_op = tf.initialize_all_variables()

sess = tf.Session()
sess.run(init_op)

for step in xrange(100):
    _, centroid_values, assignment_values = sess.run([update_centroides,
        centroides, assignments])
```

K-means: Case study



K-means: Case study

- Plot code

```
data = {"x": [], "y": [], "cluster": []}

for i in xrange(len(assignment_values)):
    data["x"].append(conjunto_puntos[i][0])
    data["y"].append(conjunto_puntos[i][1])
    data["cluster"].append(assignment_values[i])

df = pd.DataFrame(data)
sns.lmplot("x", "y", data=df,
           fit_reg=False, size=6,
           hue="cluster", legend=False)

plt.show()
```

K-means in TensorFlow: in more detail!

- Create a Tensor

```
vectors = tf.constant(conjunto_puntos)
```

- Create a k initial "means"

```
k = 4
centroides = tf.Variable(tf.slice(tf.random_shuffle(vectors),
[0,0],[k,-1]))
```

```
print vectors.get_shape()
print centroides.get_shape()

TensorShape([Dimension(2000), Dimension(2)])
TensorShape([Dimension(4), Dimension(2)])
```

K-means in TensorFlow: in more detail!

```
expanded_vectors = tf.expand_dims(vectors, 0)
expanded_centroides = tf.expand_dims(centroides, 1)

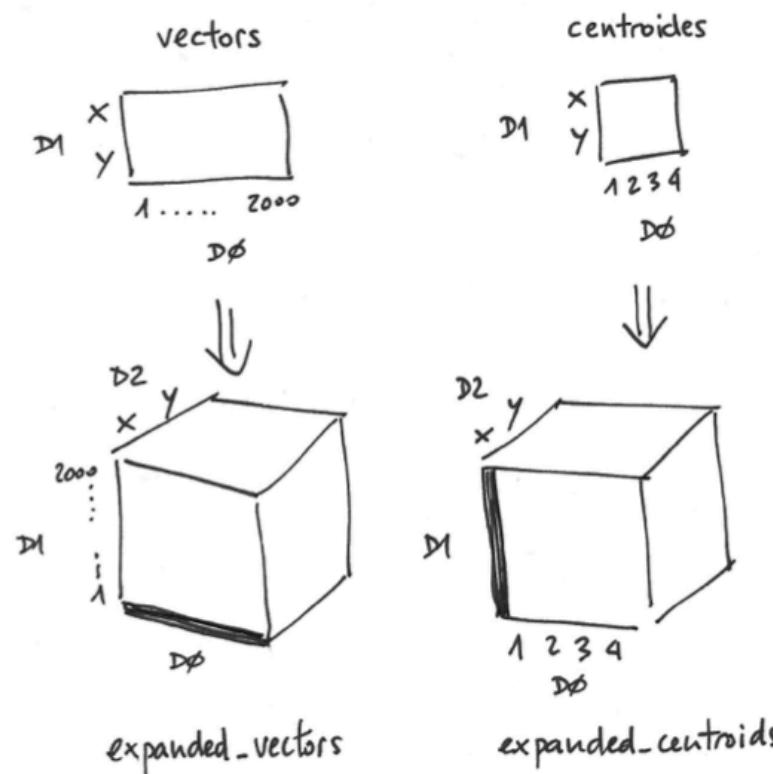
assignments =
tf.argmin(tf.reduce_sum(tf.square(tf.sub(expanded_vectors,
expanded_centroides)), 2), 0)
```

- “Squared Euclidian Distance”

$$d^2(\text{vector}, \text{centroide}) = (\text{vector}_x - \text{centroide}_x)^2 + (\text{vector}_y - \text{centroide}_y)^2$$

K-means in TensorFlow: in more detail!

```
expanded_vectors = tf.expand_dims(vectors, 0)  
expanded_centroides = tf.expand_dims(centroides, 1)
```



K-means in TensorFlow: in more detail!

- “Squared Euclidian Distance”

$$d^2(\text{vector}, \text{centroide}) = (\text{vector}_x - \text{centroide}_x)^2 + (\text{vector}_y - \text{centroide}_y)^2$$

```
diff=tf.sub(expanded_vectors, expanded_centroides)
sqr= tf.square(diff)
distances = tf.reduce_sum(sqr, 2)
assignments = tf.argmin(distances, 0)
```

```
diff -> TensorShape([Dimension(4), Dimension(2000), Dimension(2)])
sqr -> TensorShape([Dimension(4), Dimension(2000), Dimension(2)])
distance -> TensorShape([Dimension(4), Dimension(2000)])
assignments -> TensorShape([Dimension(2000)])
```

(*) ***Shape broadcasting***

Tensor operations

Operación	Descripción
tf.reduce_sum	Calcula la suma de los elementos a lo largo de una de las dimensiones
tf.reduce_prod	Calcula el producto de los elementos a lo largo de una de las dimensiones
tf.reduce_min	Calcula el mínimo de los elementos a lo largo de una dimensión
tf.reduce_max	Calcula el máximo de los elementos a lo largo de una dimensión
tf.reduce_mean	Calcula la media de los elementos a lo largo de una dimensión

Tensor operations

Operación	Descripción
tf.argmin	Retorna el índice del elemento con el valor menor en la dimensión del <i>tensor</i> indicada
tf.argmax	Retorna el índice del elemento con el valor máximo en la dimensión del <i>tensor</i> indicada

K-means in TensorFlow: in more detail!

- New centroids

```
means = tf.concat(0, [tf.reduce_mean(tf.gather(vectors,  
tf.reshape(tf.where( tf.equal(assignments, c)),[1,-1])),  
reduction_indices=[1]) for c in xrange(k)])
```

means tensor is equal to the concatenation of k tensors that correspond to the mean value of cluster elements.

Detailed explanation:

Detalle de cada una de las operaciones TensorFlow que intervienen en el cálculo del valor medio de los puntos pertenecientes a un *cluster*:

- Con `tf.equal` se obtiene un tensor booleano (Dimension(2000)) que indica (con valor “true”) las posiciones donde el valor del tensor `assignment` coincide con el cluster *c* (uno de los *k*), del que en aquel momento estamos calculando el valor medio de sus puntos.
- Con `tf.where` se construye un tensor (Dimension (1) x Dimension(2000)) con la posición donde se encuentran los valores “true” en el tensor booleano recibido como parámetro. Es decir, una lista con las posiciones de estos.
- Con `tf.reshape` se construye un tensor (Dimension (2000) x Dimension(1)) con los índices de los puntos en el tensor `vectors` que pertenecen a este cluster *c*.
- Con `tf.gather` se construye un tensor (Dimension (1) x Dimension (2000) x Dimension(2)) que reúne las coordenadas de los puntos que forman el cluster *c*.
- Con `tf.reduce_mean` se construye un tensor (Dimension (1) x Dimension(2)) que contiene el valor medio de todos los puntos que pertenecen a este cluster *c*.

K-means in TensorFlow: in more detail!

```
update_centroides = tf.assign(centroides, means)

init_op = tf.initialize_all_variables()

sess = tf.Session()
sess.run(init_op)

for step in xrange(100):
    _, centroid_values, assignment_values = sess.run(
        [update_centroides,centroides, assignments])
```

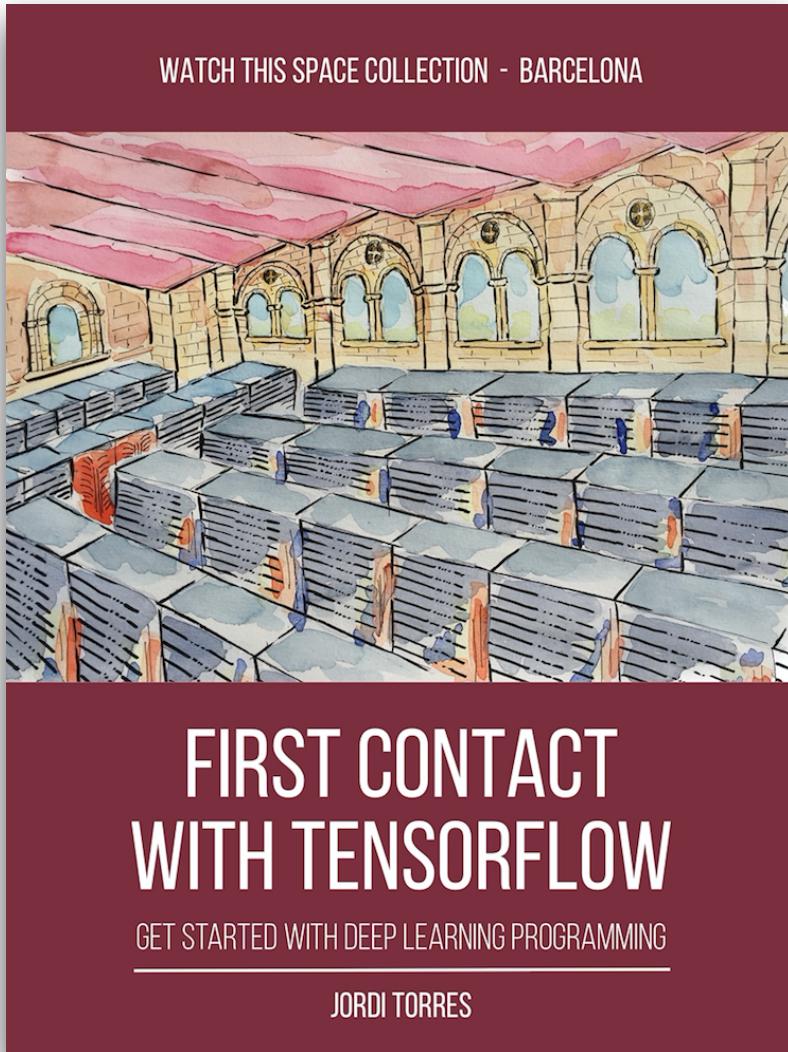
Hands-on 3

1. Download **kmeans.py** from:

<https://github.com/jorditorresBCN/TutorialTensorFlow>

2. Try with diff number of clusters and elements
3. Follow teacher's instructions

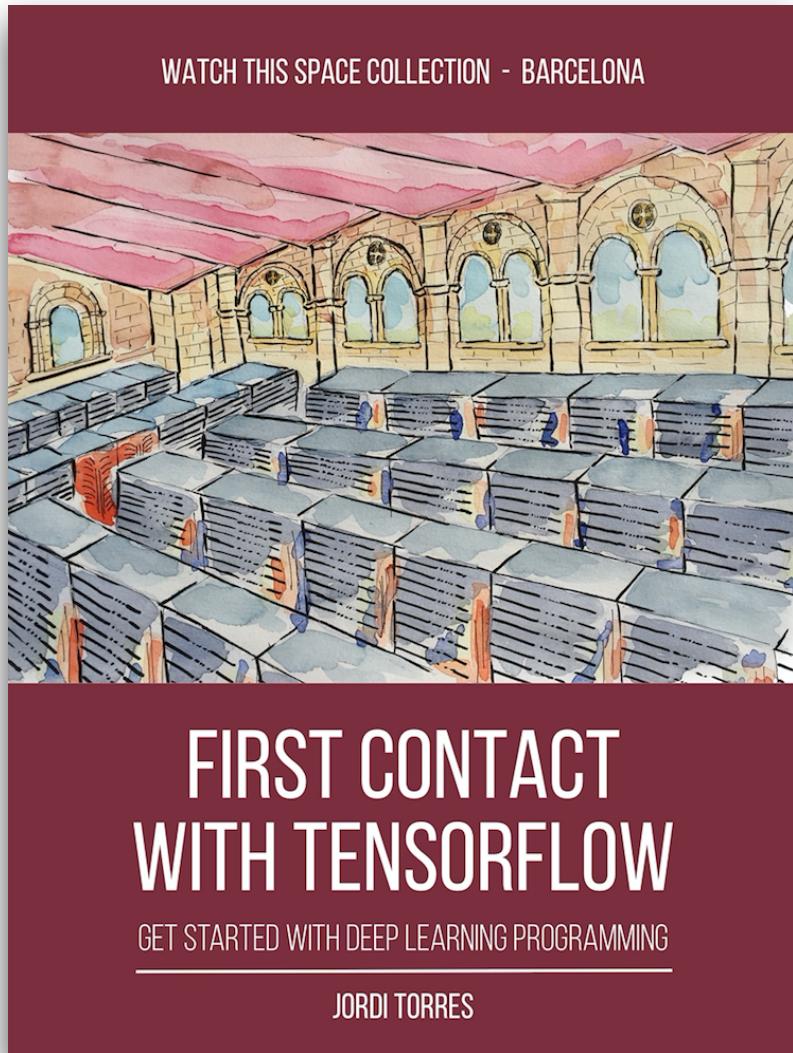
Bonus: advanced content



Part2: Advanced

- 4. Single Layer Neural Network
 - 5. Multi-layer Neural Networks
 - 6. TensorFlow and GPUs
- Closing

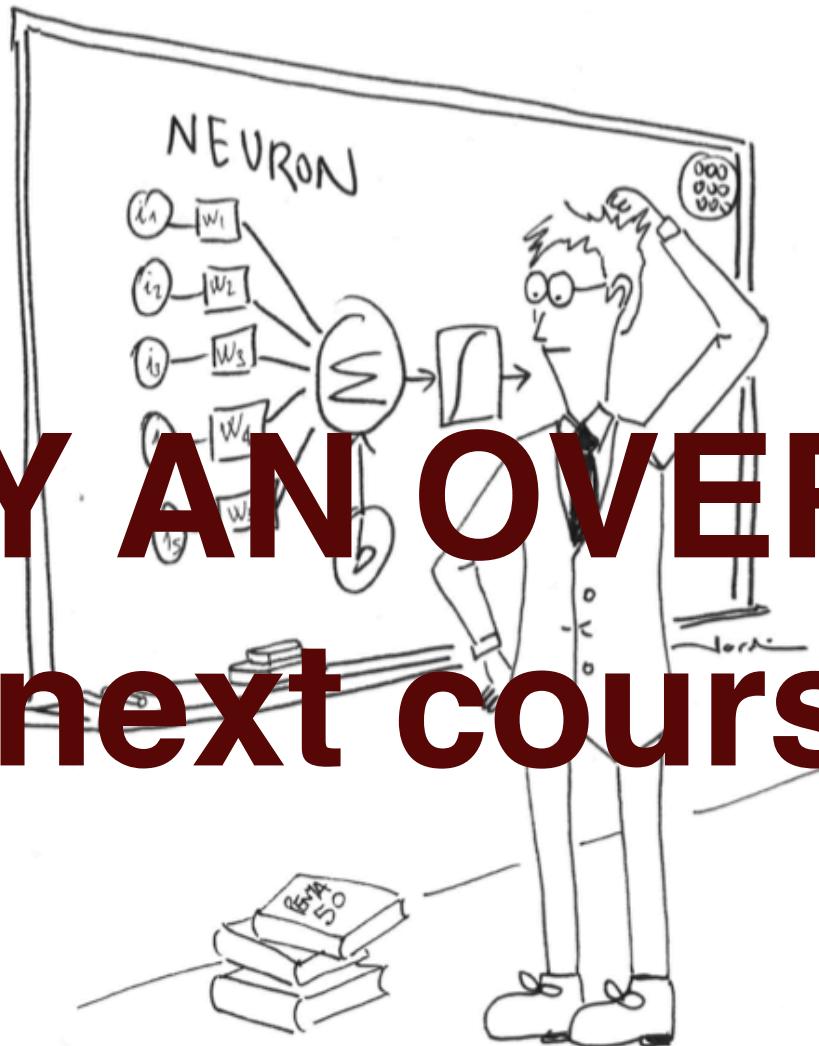
Content:



Introduction

1. TensorFlow Basics
2. Linear Regression
3. Clustering
- 4. Single Layer Neural Network**
5. Multi-layer Neural Networks
6. TensorFlow and GPUs
- Closing

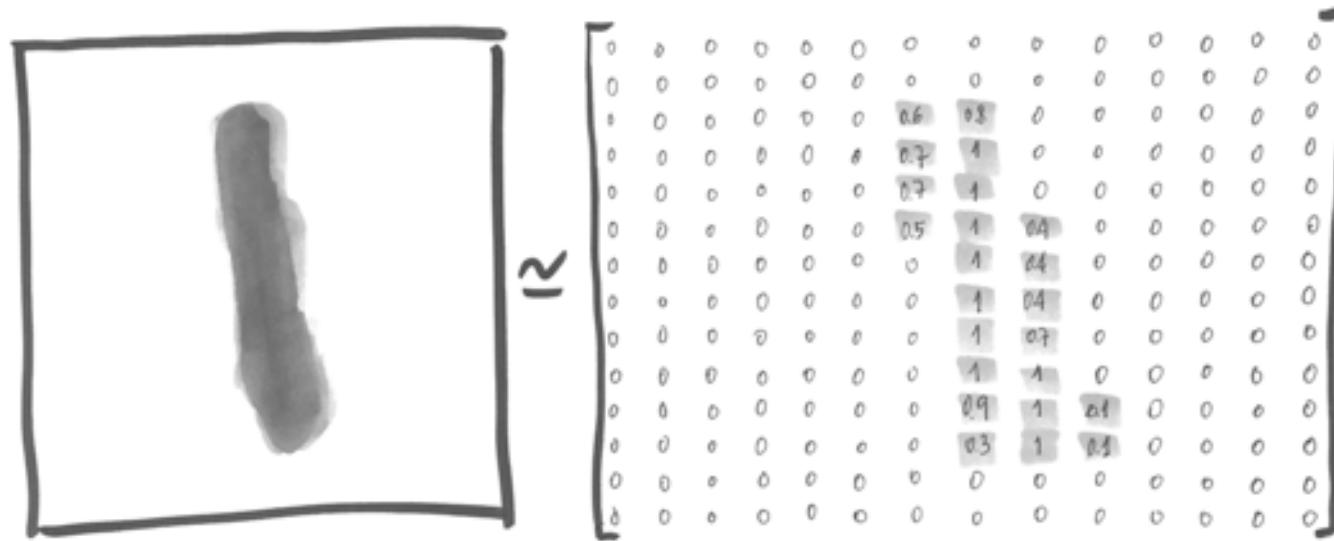
**ONLY AN OVERVIEW
(next course)**

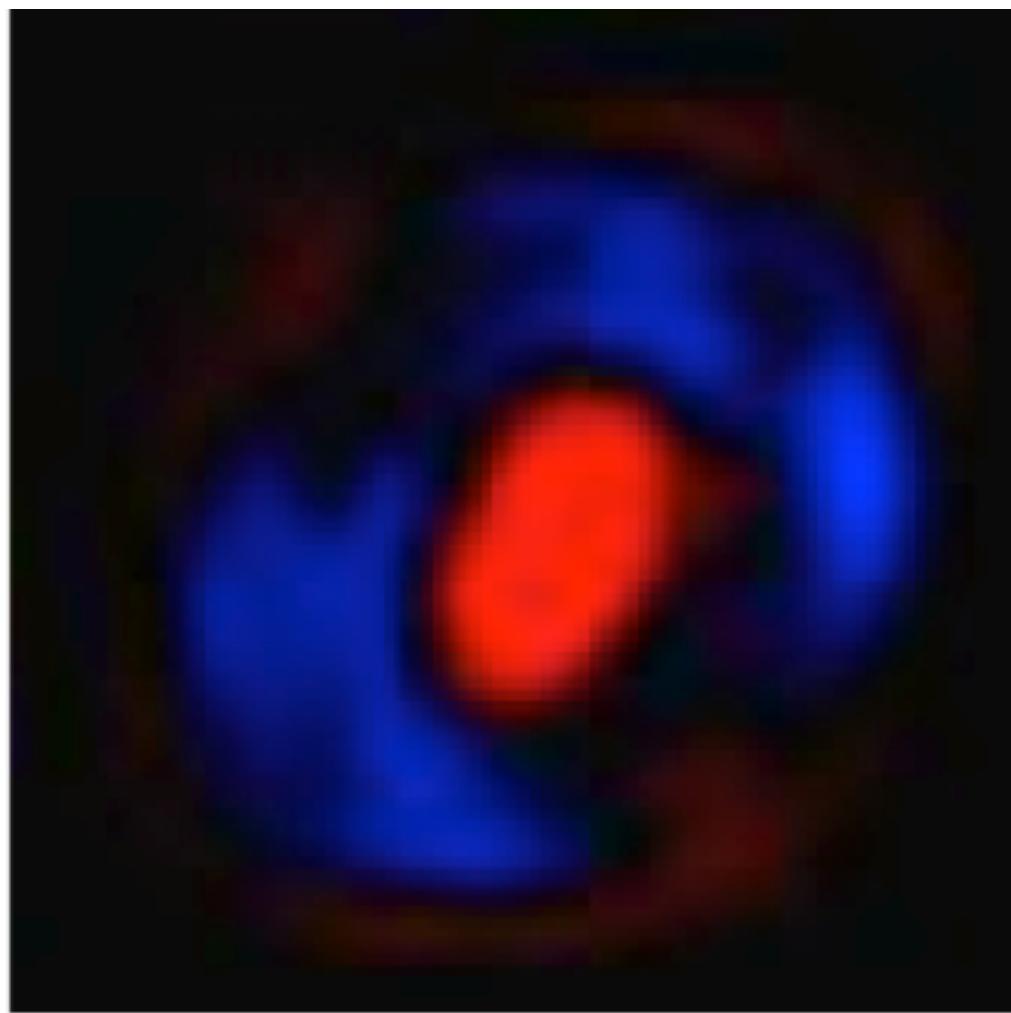


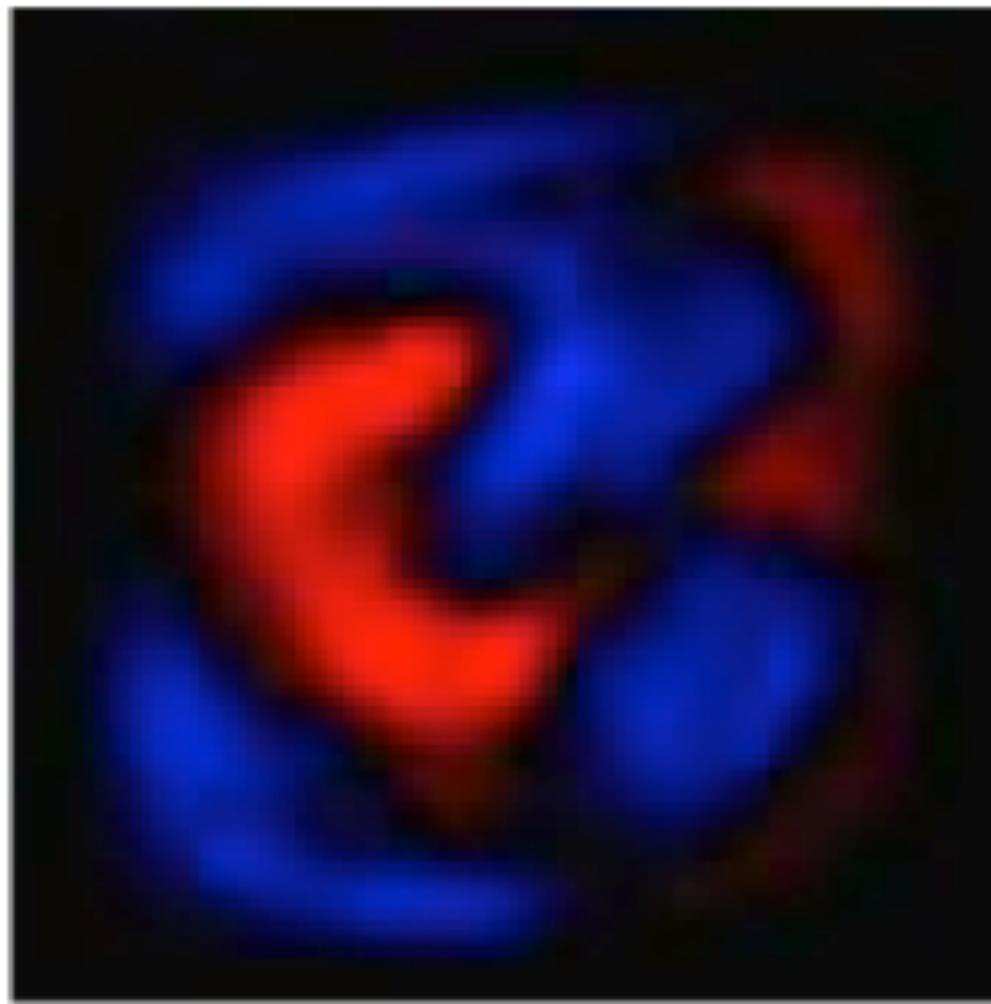
“Hello World” = MNIST

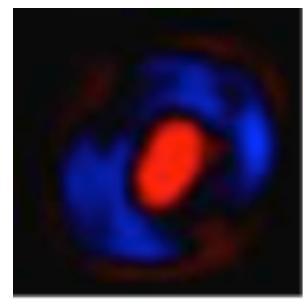


“Hello World” = MNIST

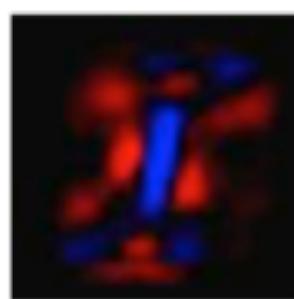




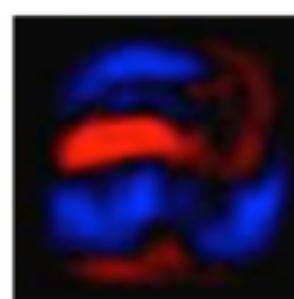




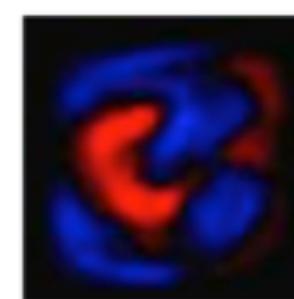
0



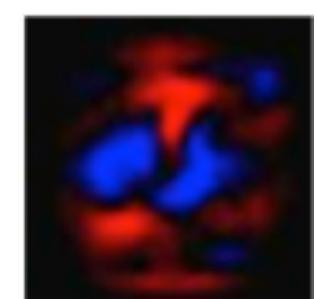
1



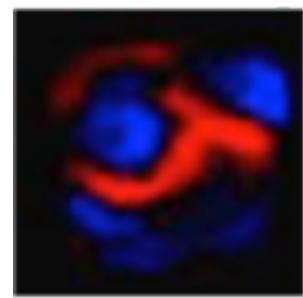
2



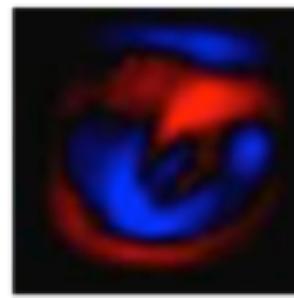
3



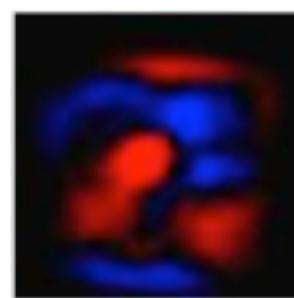
4



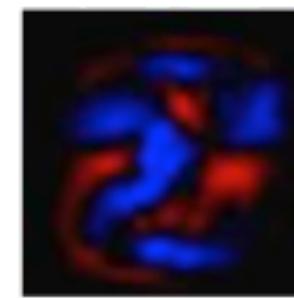
5



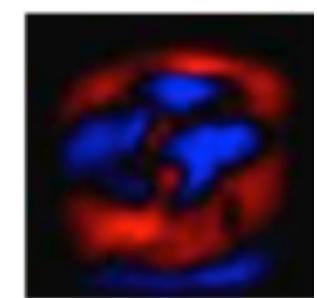
6



7



8



9

$$\text{softmax}(x)_i = \frac{\exp(x_i)}{\sum_j \exp(x_j)} = \frac{e^{x_i}}{\sum_j e^{x_j}}$$

$$\begin{bmatrix} y \\ 0 \end{bmatrix} = \text{softmax} \left(\begin{bmatrix} W & X \\ 0 & b \end{bmatrix} \right)$$

The diagram shows a vector y of size 10, which is the softmax output of a linear transformation. This transformation is represented by a matrix multiplication $\begin{bmatrix} W & X \\ 0 & b \end{bmatrix}$. The matrix has two columns: W (size 10x784) and X (size 784x1). A plus sign (+) indicates that the bias b (size 10x1) is added to the result of the multiplication $W \cdot X$.

Single Layer Neural Network

```
import input_data
mnist = input_data.read_data_sets("MNIST_data/", one_hot=True)

import tensorflow as tf
x = tf.placeholder("float", [None, 784])
W = tf.Variable(tf.zeros([784,10]))
b = tf.Variable(tf.zeros([10]))

matm=tf.matmul(x,W)
y = tf.nn.softmax(tf.matmul(x,W) + b)
y_ = tf.placeholder("float", [None,10])

cross_entropy = -tf.reduce_sum(y_*tf.log(y))
train_step = tf.train.GradientDescentOptimizer(0.01).minimize(cross_entropy)

sess = tf.Session()
sess.run(tf.initialize_all_variables())

for i in range(1000):
    batch_xs, batch_ys = mnist.train.next_batch(100)
    sess.run(train_step, feed_dict={x: batch_xs, y_: batch_ys})
    correct_prediction = tf.equal(tf.argmax(y,1), tf.argmax(y_,1))
    accuracy = tf.reduce_mean(tf.cast(correct_prediction, "float"))
    print sess.run(accuracy, feed_dict={x: mnist.test.images, y_: mnist.test.labels})
```

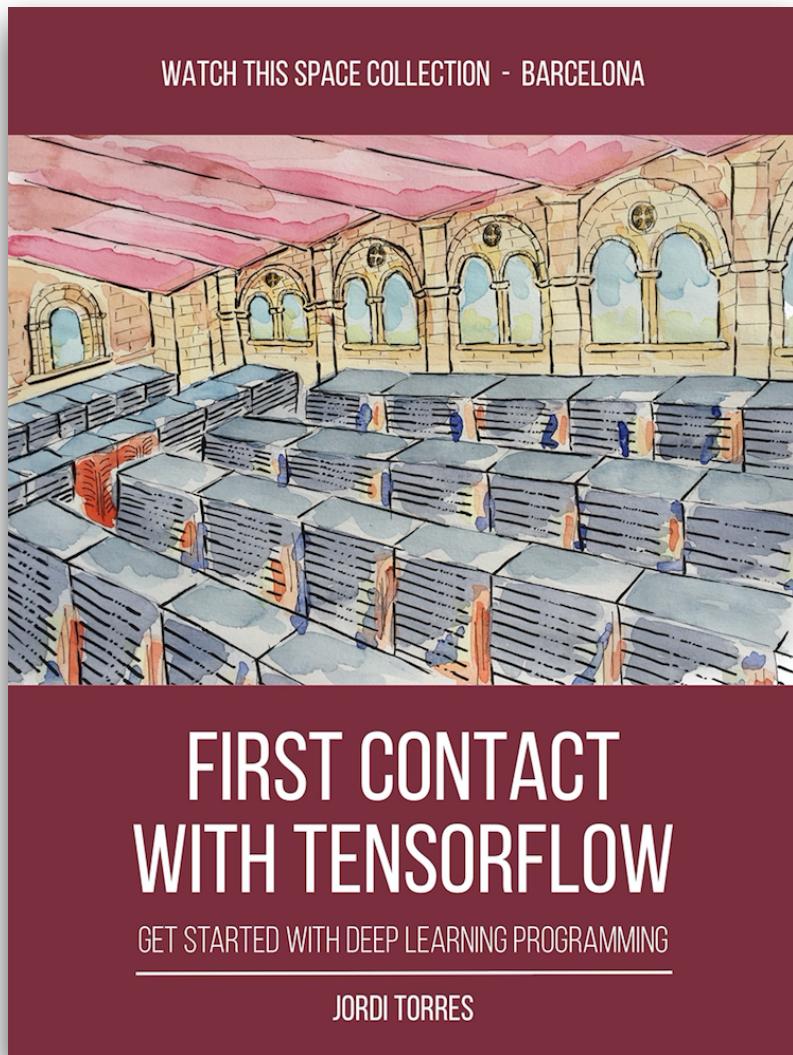
Hands-on 4

1. Download **redneuronalsimple.py** from:

<https://github.com/jorditorresBCN/TutorialTensorFlow>

2. Follow teacher's instructions

Content:



Introduction

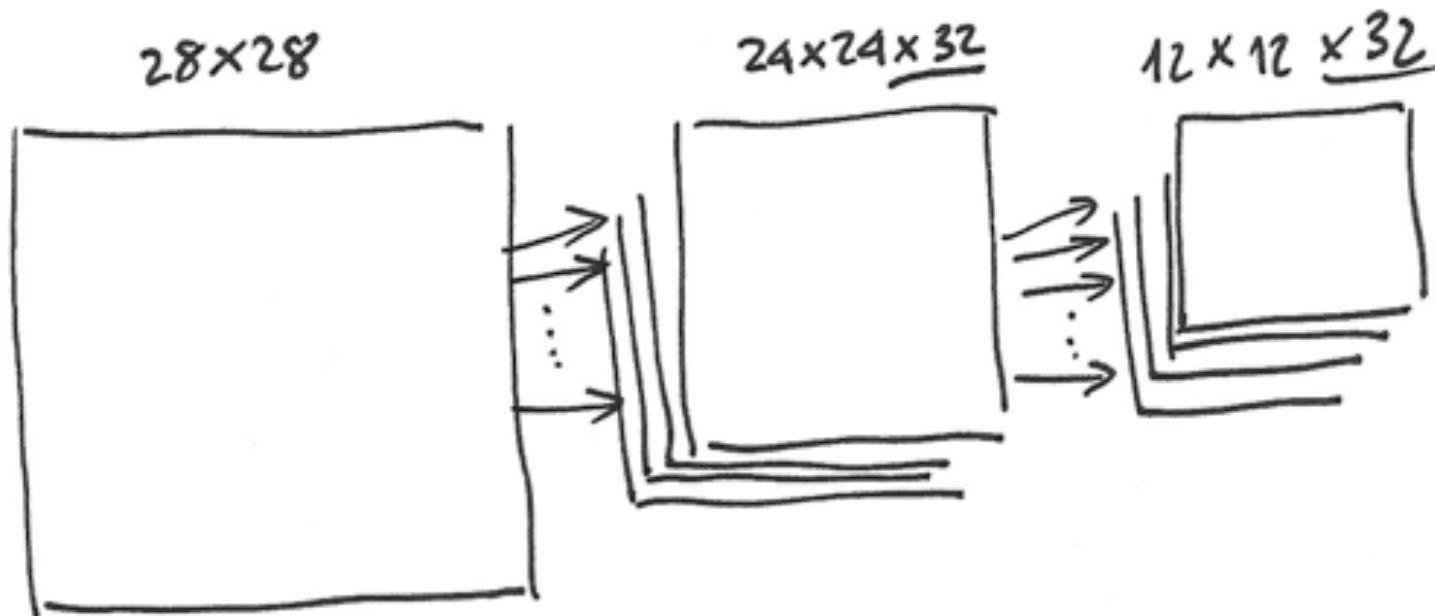
1. TensorFlow Basics
 2. Linear Regression
 3. Clustering
 4. Single Layer Neural Network
 5. Multi-layer Neural Networks
 6. TensorFlow and GPUs
- Closing

ONLY AN OVERVIEW (next course)



Convolutional Neural Networks

- Convolution
- Pooling



```

import input_data
mnist = input_data.read_data_sets('MNIST_data', one_hot=True)
import tensorflow as tf

x = tf.placeholder("float", shape=[None, 784])
y_ = tf.placeholder("float", shape=[None, 10])

x_image = tf.reshape(x, [-1, 28, 28, 1])
print "x_image="
print x_image

def weight_variable(shape):
    initial = tf.truncated_normal(shape, stddev=0.1)
    return tf.Variable(initial)

def bias_variable(shape):
    initial = tf.constant(0.1, shape=shape)
    return tf.Variable(initial)

def conv2d(x, W):
    return tf.nn.conv2d(x, W, strides=[1, 1, 1, 1],
                       padding='SAME')

def max_pool_2x2(x):
    return tf.nn.max_pool(x, ksize=[1, 2, 2, 1],
                          strides=[1, 2, 2, 1], padding='SAME')

W_conv1 = weight_variable([5, 5, 1, 32])
b_conv1 = bias_variable([32])

h_conv1 = tf.nn.relu(conv2d(x_image, W_conv1) + b_conv1)
h_pool1 = max_pool_2x2(h_conv1)

W_conv2 = weight_variable([5, 5, 32, 64])
b_conv2 = bias_variable([64])

h_conv2 = tf.nn.relu(conv2d(h_pool1, W_conv2) + b_conv2)
h_pool2 = max_pool_2x2(h_conv2)

import tensorflow as tf
tf.reset_default_graph()

W_fc1 = weight_variable([7 * 7 * 64, 1024])
b_fc1 = bias_variable([1024])

h_pool2_flat = tf.reshape(h_pool2, [-1, 7*7*64])
h_fc1 = tf.nn.relu(tf.matmul(h_pool2_flat, W_fc1) + b_fc1)

keep_prob = tf.placeholder("float")
h_fc1_drop = tf.nn.dropout(h_fc1, keep_prob)

W_fc2 = weight_variable([1024, 10])
b_fc2 = bias_variable([10])

y_conv = tf.nn.softmax(tf.matmul(h_fc1_drop, W_fc2) + b_fc2)

cross_entropy = -tf.reduce_sum(y_*tf.log(y_conv))
train_step = tf.train.AdamOptimizer(1e-4).minimize(cross_entropy)
correct_prediction = tf.equal(tf.argmax(y_conv, 1),
                             tf.argmax(y_, 1))
accuracy = tf.reduce_mean(tf.cast(correct_prediction, "float"))

sess = tf.Session()
sess.run(tf.initialize_all_variables())

for i in range(200):
    batch = mnist.train.next_batch(50)
    if i%10 == 0:
        train_accuracy = sess.run(accuracy, feed_dict={
            x:batch[0], y_: batch[1], keep_prob: 1.0})
        print("step %d, training accuracy %g" % (i, train_accuracy))
    sess.run(train_step, feed_dict={x: batch[0], y_: batch[1],
                                   keep_prob: 0.5})

print("test accuracy %g" % sess.run(accuracy, feed_dict={
    x: mnist.test.images, y_: mnist.test.labels, keep_prob:
    1.0}))

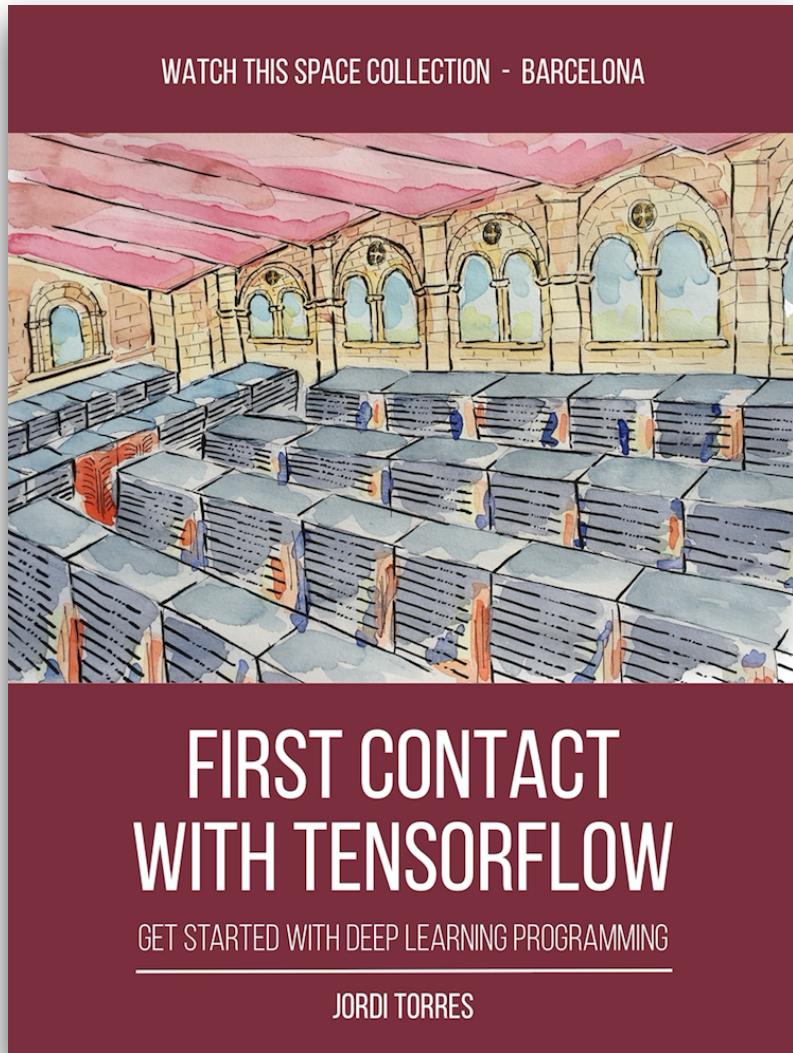
```

91% → 99,2%

Hands-on 5

1. Download **CNN.py** from:
<https://github.com/jorditorresBCN/TutorialTensorFlow>
2. Follow teacher's instructions

Content:



Introduction

1. TensorFlow Basics
2. Linear Regression
3. Clustering
4. Single Layer Neural Network
5. Multi-layer Neural Networks
6. TensorFlow and GPUs

Closing

**ONLY AN OVERVIEW
(next course)**



```

import numpy as np
import tensorflow as tf
import datetime

#Processing Units logs
log_device_placement = True

#num of multiplications to perform
n = 10

#create random large matrix
A = np.random.rand(1e4, 1e4).astype('float32')
B = np.random.rand(1e4, 1e4).astype('float32')

# Creates a graph to store results
c1 = []
c2 = []

def matpow(M, n):
    if n < 1: #Abstract cases where n < 1
        return M
    else:
        return tf.matmul(M, matpow(M, n-1))

with tf.device('/gpu:0'):
    a = tf.constant(A)
    b = tf.constant(B)
    #compute A^n and B^n and store results in c1
    c1.append(matpow(a, n))
    c1.append(matpow(b, n))

with tf.device('/cpu:0'):
    sum = tf.add_n(c1) #Addition of all elements in c1, i.e. A^n + B^n

t1_1 = datetime.datetime.now()
with tf.Session(config=tf.ConfigProto(log_device_placement=log_device_placement)) as sess:
    # Runs the op.
    sess.run(sum)
t2_1 = datetime.datetime.now()

#GPU:0 computes A^n
with tf.device('/gpu:0'):
    #compute A^n and store result in c2
    a = tf.constant(A)
    c2.append(matpow(a, n))

#GPU:1 computes B^n
with tf.device('/gpu:1'):
    #compute B^n and store result in c2
    b = tf.constant(B)
    c2.append(matpow(b, n))

with tf.device('/cpu:0'):
    sum = tf.add_n(c2) #Addition of all elements in c2, i.e. A^n + B^n

t1_2 = datetime.datetime.now()
with tf.Session(config=tf.ConfigProto(log_device_placement=log_device_placement)) as sess:
    # Runs the op.
    sess.run(sum)
t2_2 = datetime.datetime.now()

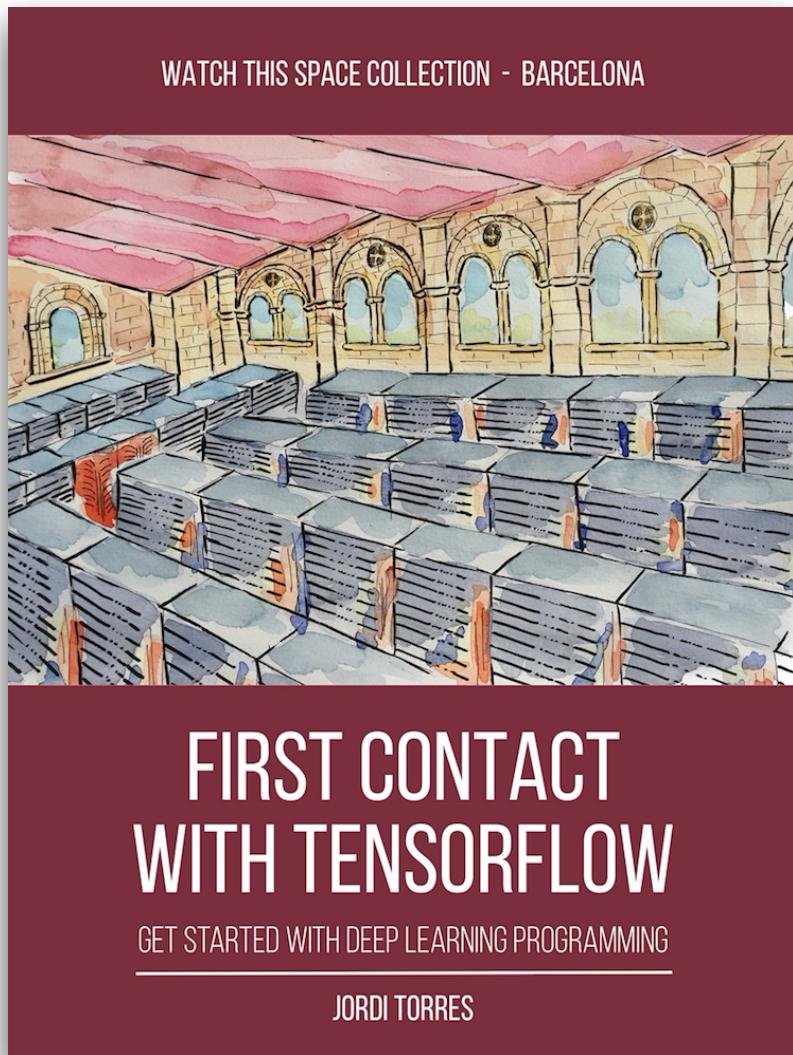
print "Single GPU computation time: " + str(t2_1-t1_1)
print "Multi GPU computation time: " + str(t2_2-t1_2)

```

Hands-on 6

1. Download **GPUs.py** from:
<https://github.com/jorditorresBCN/TutorialTensorFlow>
2. Follow teacher's instructions

Content:



Introduction

1. TensorFlow Basics
2. Linear Regression
3. Clustering
4. Single Layer Neural Network
5. Multi-layer Neural Networks
6. TensorFlow and GPUs

Closing

Kyle McDonald

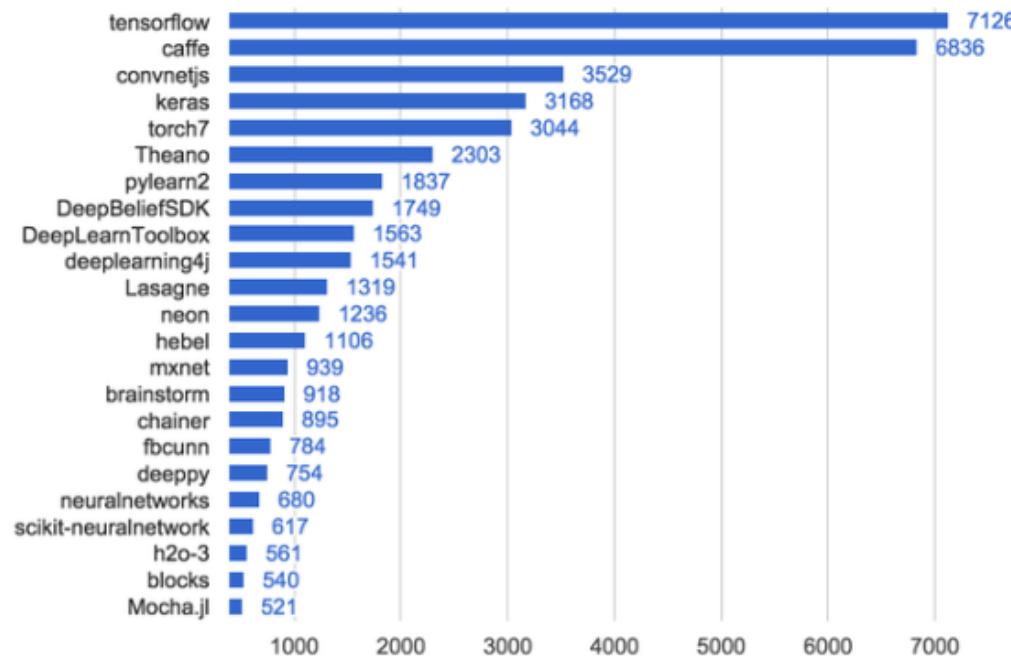
@kcimc



Siguiendo

2010-2014: a new deep learning toolkit is released every 47 days. 2015: every 22 days. tensorflow & caffe top github

Ver traducción



RETWEETS ME GUSTA
151 157



0:05 - 11 nov. 2015



...



[ABOUT US](#)
[CONTACT](#)
[PITCH US](#)
[BOOKMARKS](#)

DATA

Cognitive Computing: Benefits and Challenges of the Next Tech Revolution

Cognitive Computing is going to transform and improve our lives. But it also presents challenges that we need to be conscious of in order to make the best use of this technology.

Co-authored by: Mateo Valero and Jordi Torres

More information:

