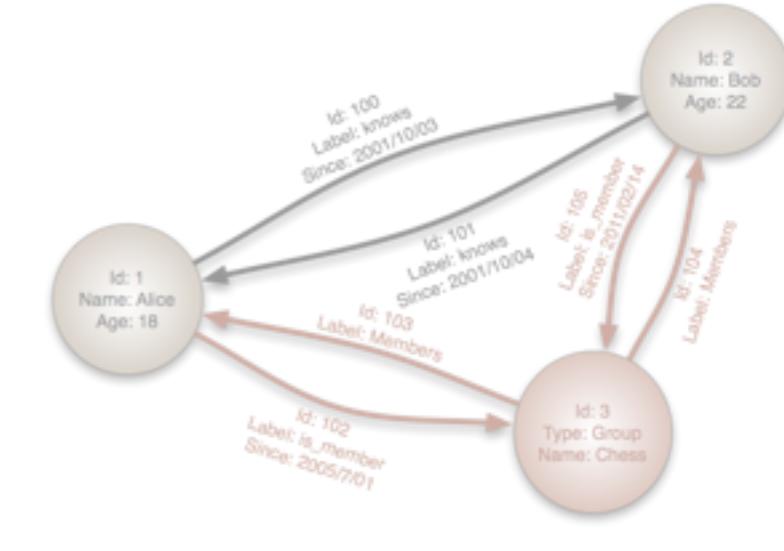
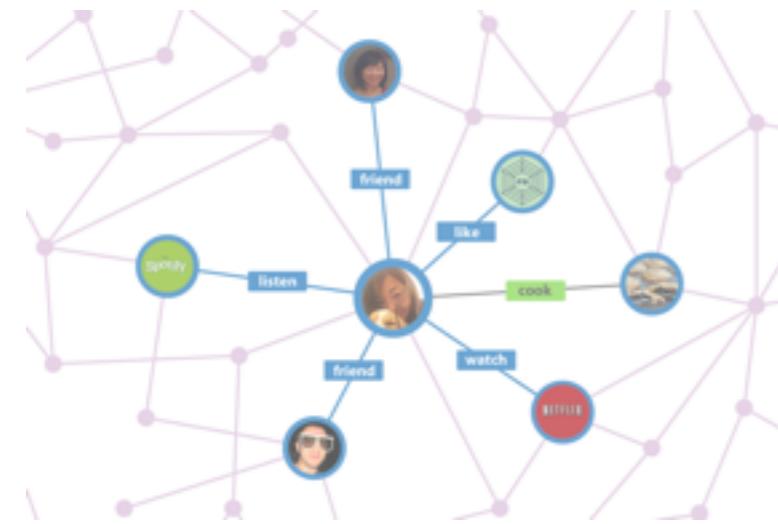
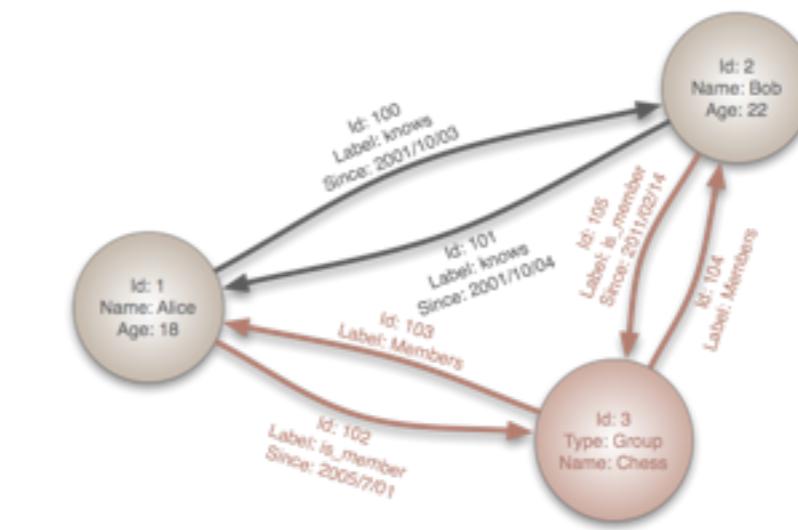


# Graph Databases Use Cases

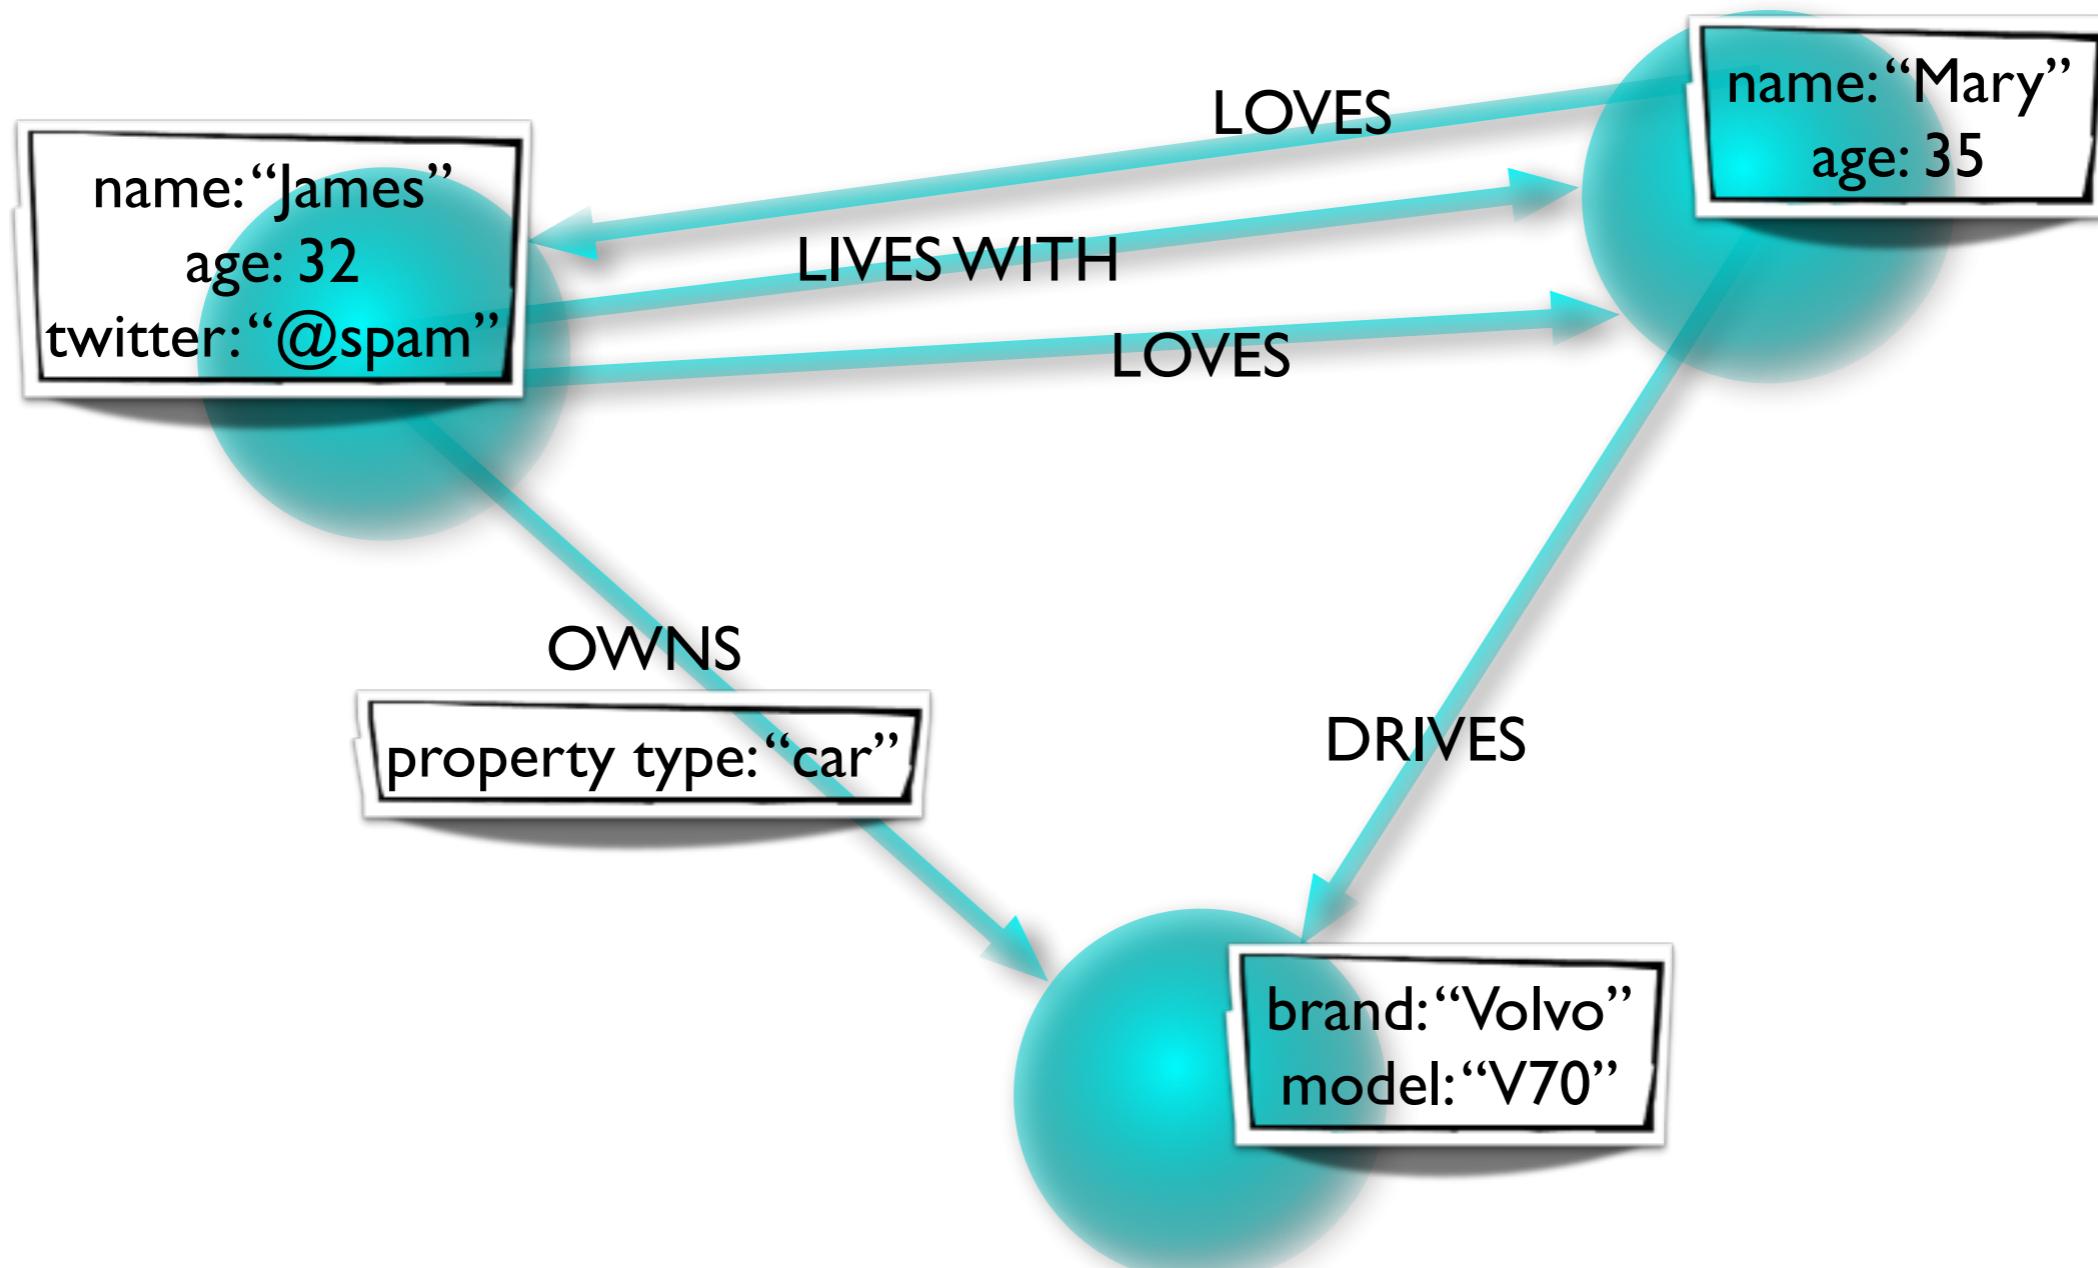




# What's a Graph?



# Graph data model



# Relational Tables

## SQL Join Hell (1)

Customer		
<b>Id</b>	Name	<b>Address</b>
1	Robert	3
2	Lars	7
3	Michael	23



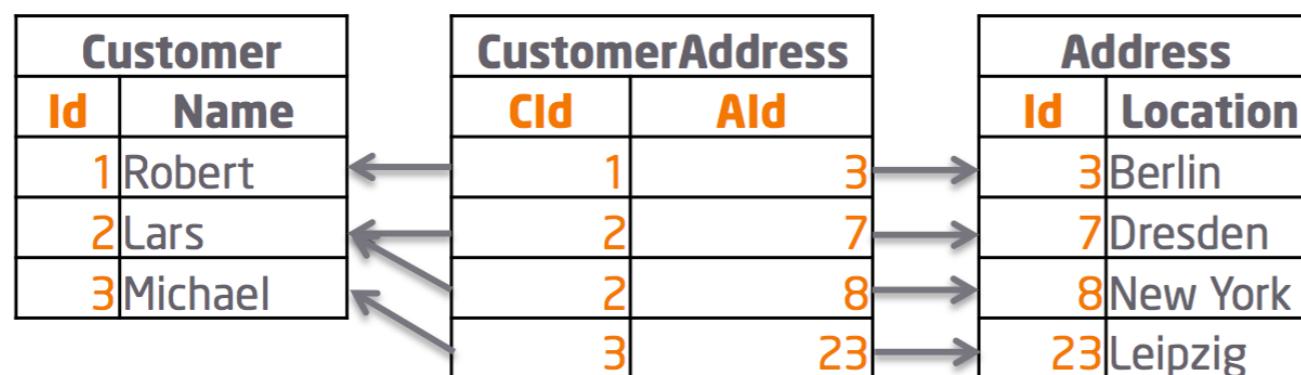
1:1 Relationship

Customer		Address		
<b>Id</b>	Name	<b>Id</b>	<b>Customer</b>	Location
1	Robert	3	1	Berlin
2	Lars	4	2	Munich
3	Michael	7	3	Dresden
		23	23	Leipzig



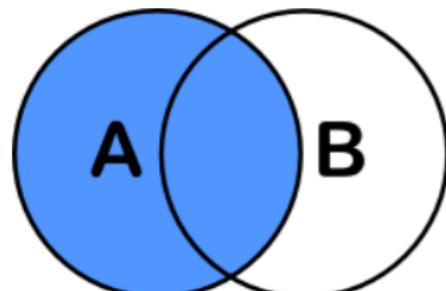
1:n Relationship

Customer		CustomerAddress		Address	
<b>Id</b>	Name	<b>CId</b>	<b>AId</b>	<b>Id</b>	Location
1	Robert	1	3	3	Berlin
2	Lars	2	7	7	Dresden
3	Michael	2	8	8	New York
		3	23	23	Leipzig

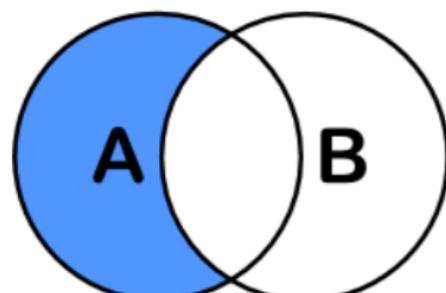


m:n Relationship

# Join this way...

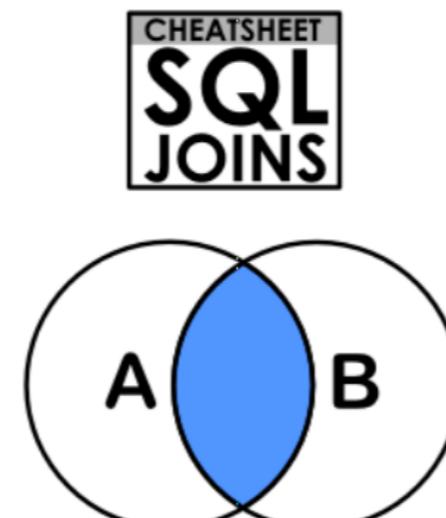


```
SELECT <auswahl>
FROM tabelleA A
LEFT JOIN tabelleB B
ON A.key = B.key
```

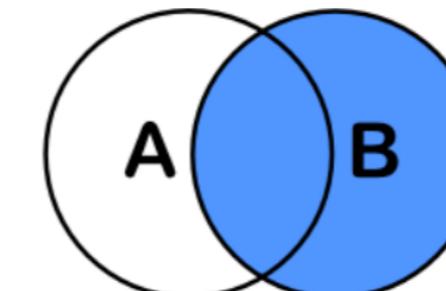
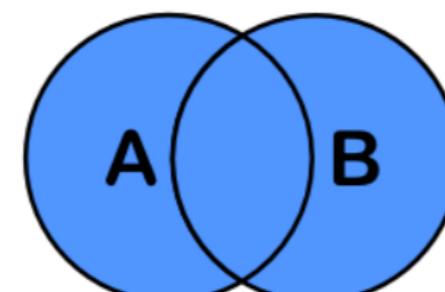


```
SELECT <auswahl>
FROM tabelleA A
LEFT JOIN tabelleB B
ON A.key = B.key
WHERE B.key IS NULL
```

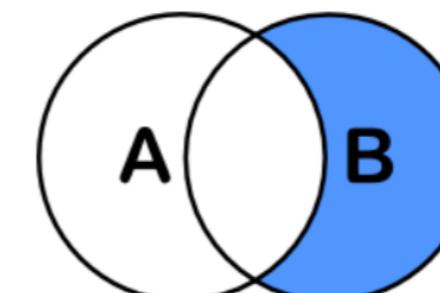
```
SELECT <auswahl>
FROM tabelleA A
FULL OUTER JOIN tabelleB B
ON A.key = B.key
```



```
SELECT <auswahl>
FROM tabelleA A
INNER JOIN tabelleB B
ON A.key = B.key
```

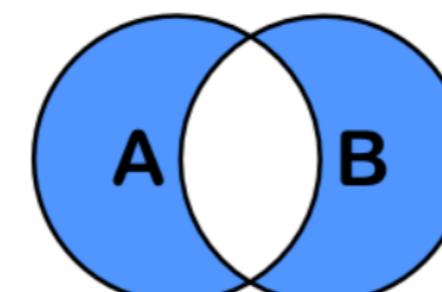


```
SELECT <auswahl>
FROM tabelleA A
RIGHT JOIN tabelleB B
ON A.key = B.key
```



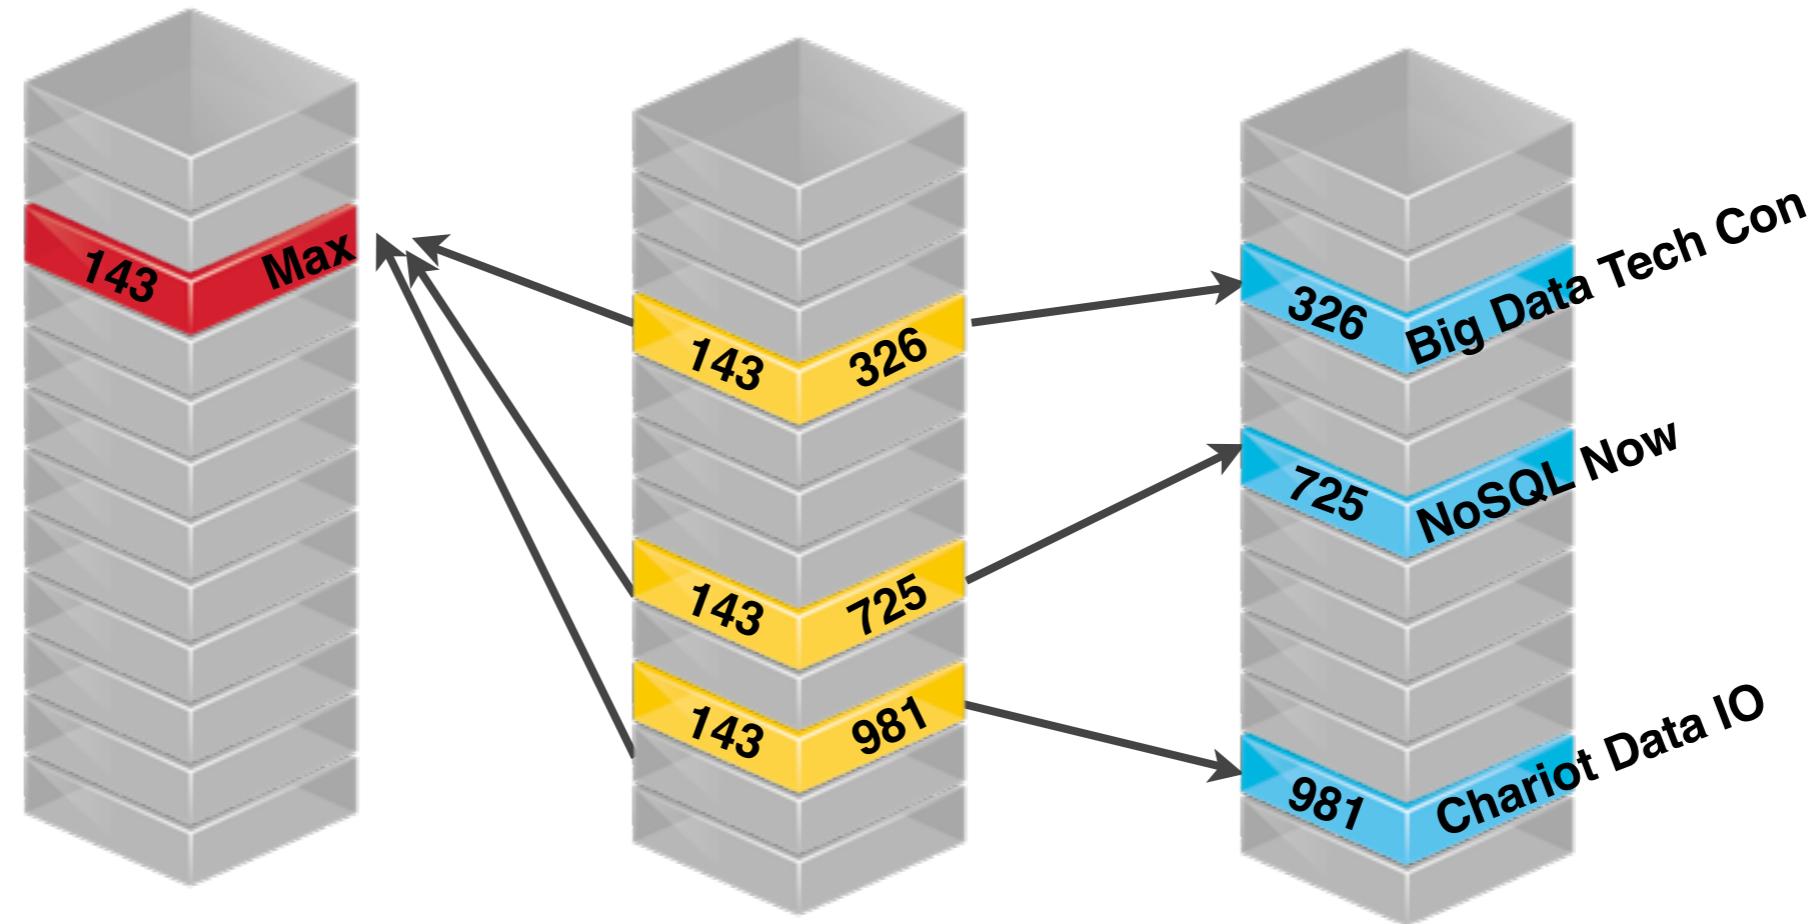
```
SELECT <auswahl>
FROM tabelleA A
RIGHT JOIN tabelleB B
ON A.key = B.key
WHERE A.key IS NULL
```

```
SELECT <auswahl>
FROM tabelleA A
FULL OUTER JOIN tabelleB B
ON A.key = B.key
WHERE A.key IS NULL
OR B.key IS NULL
```



# The Problem

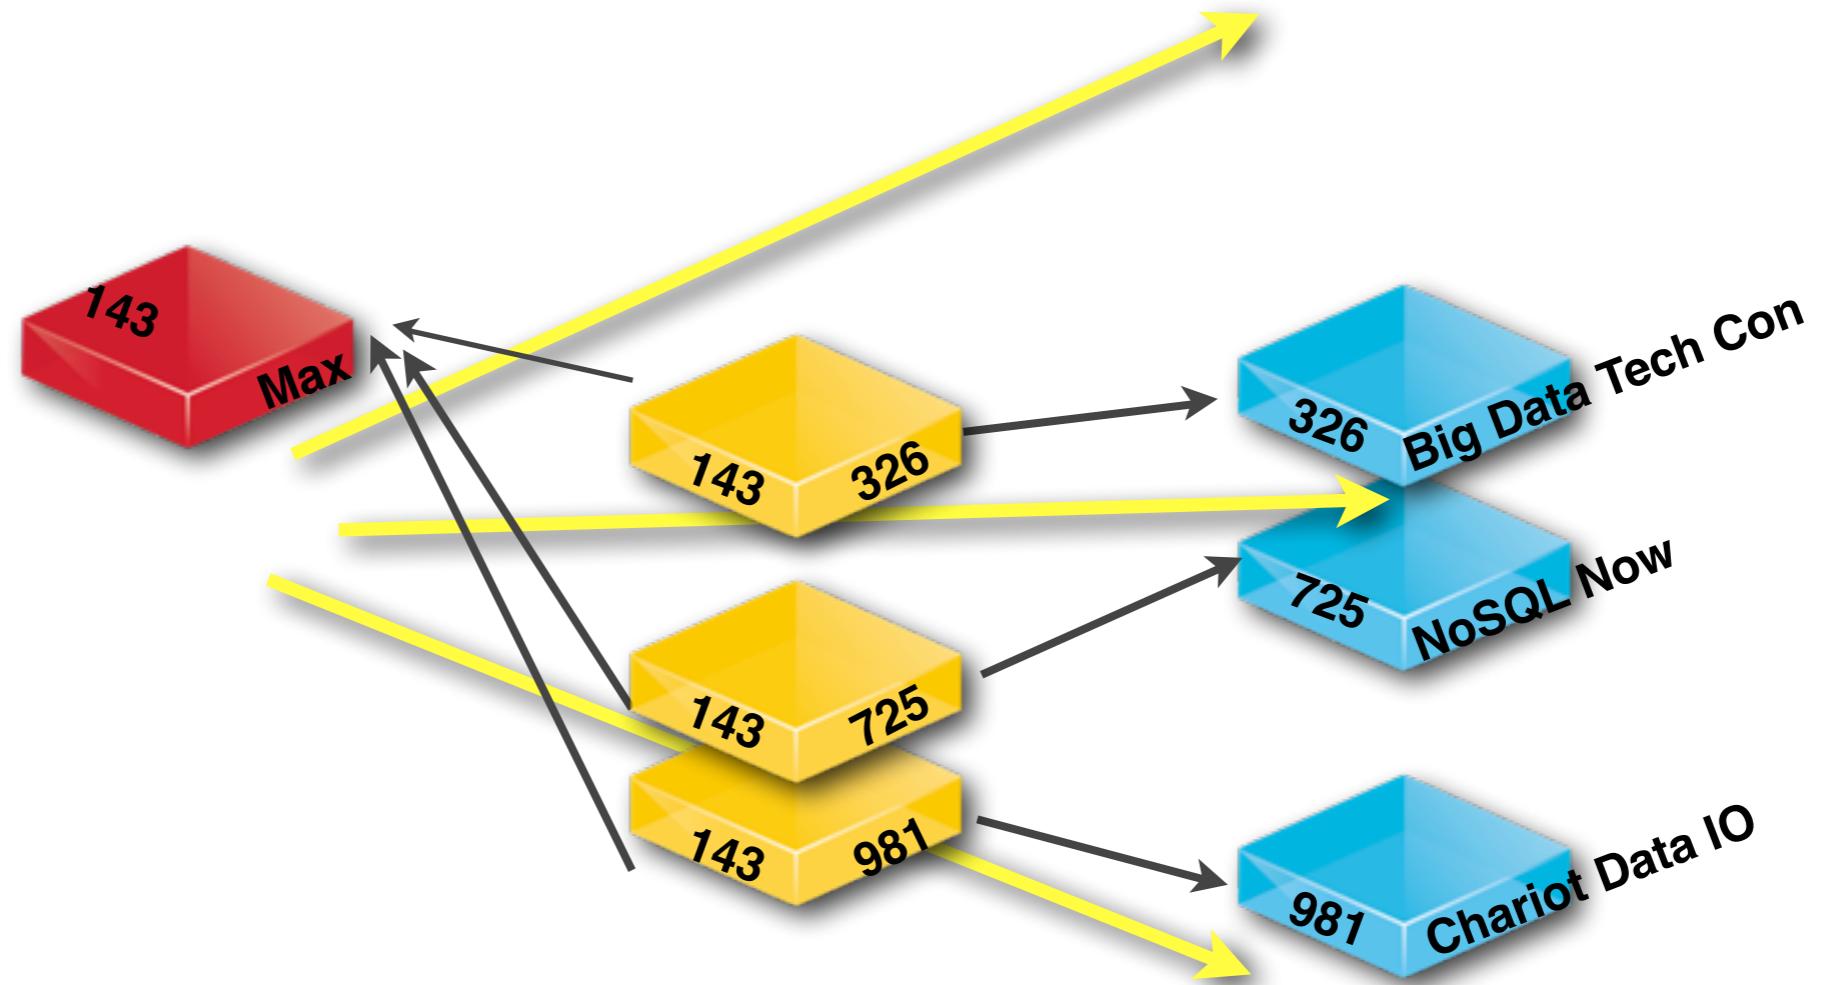
- all JOINs are executed **every time you query (traverse) the relationship**
- executing a JOIN means to **search for a key in another table**
- with Indices executing a JOIN means to **lookup a key**
- B-Tree Index:  $O(\log(n))$
- more entries => more lookups => slower JOINs



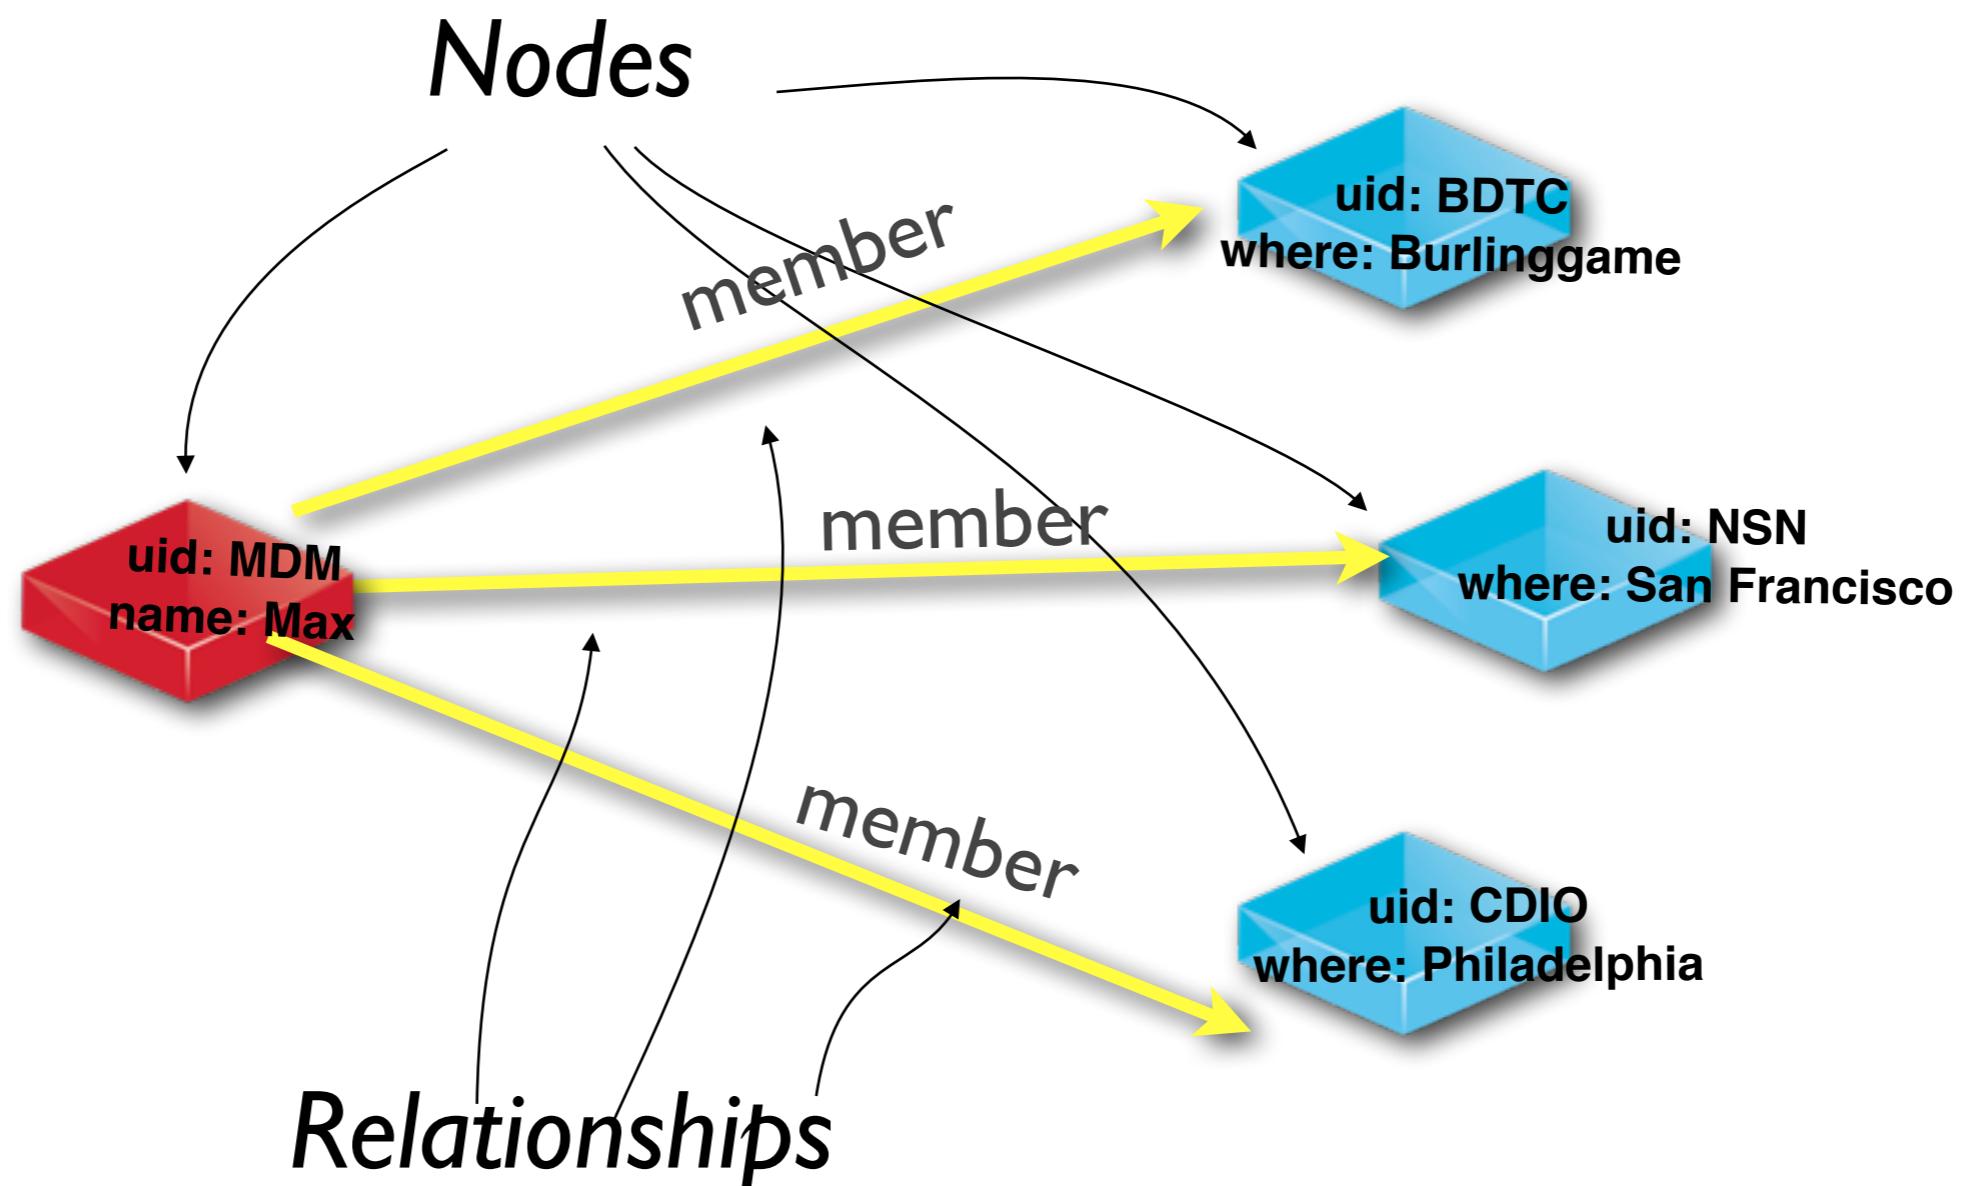
**People**

**Attend**

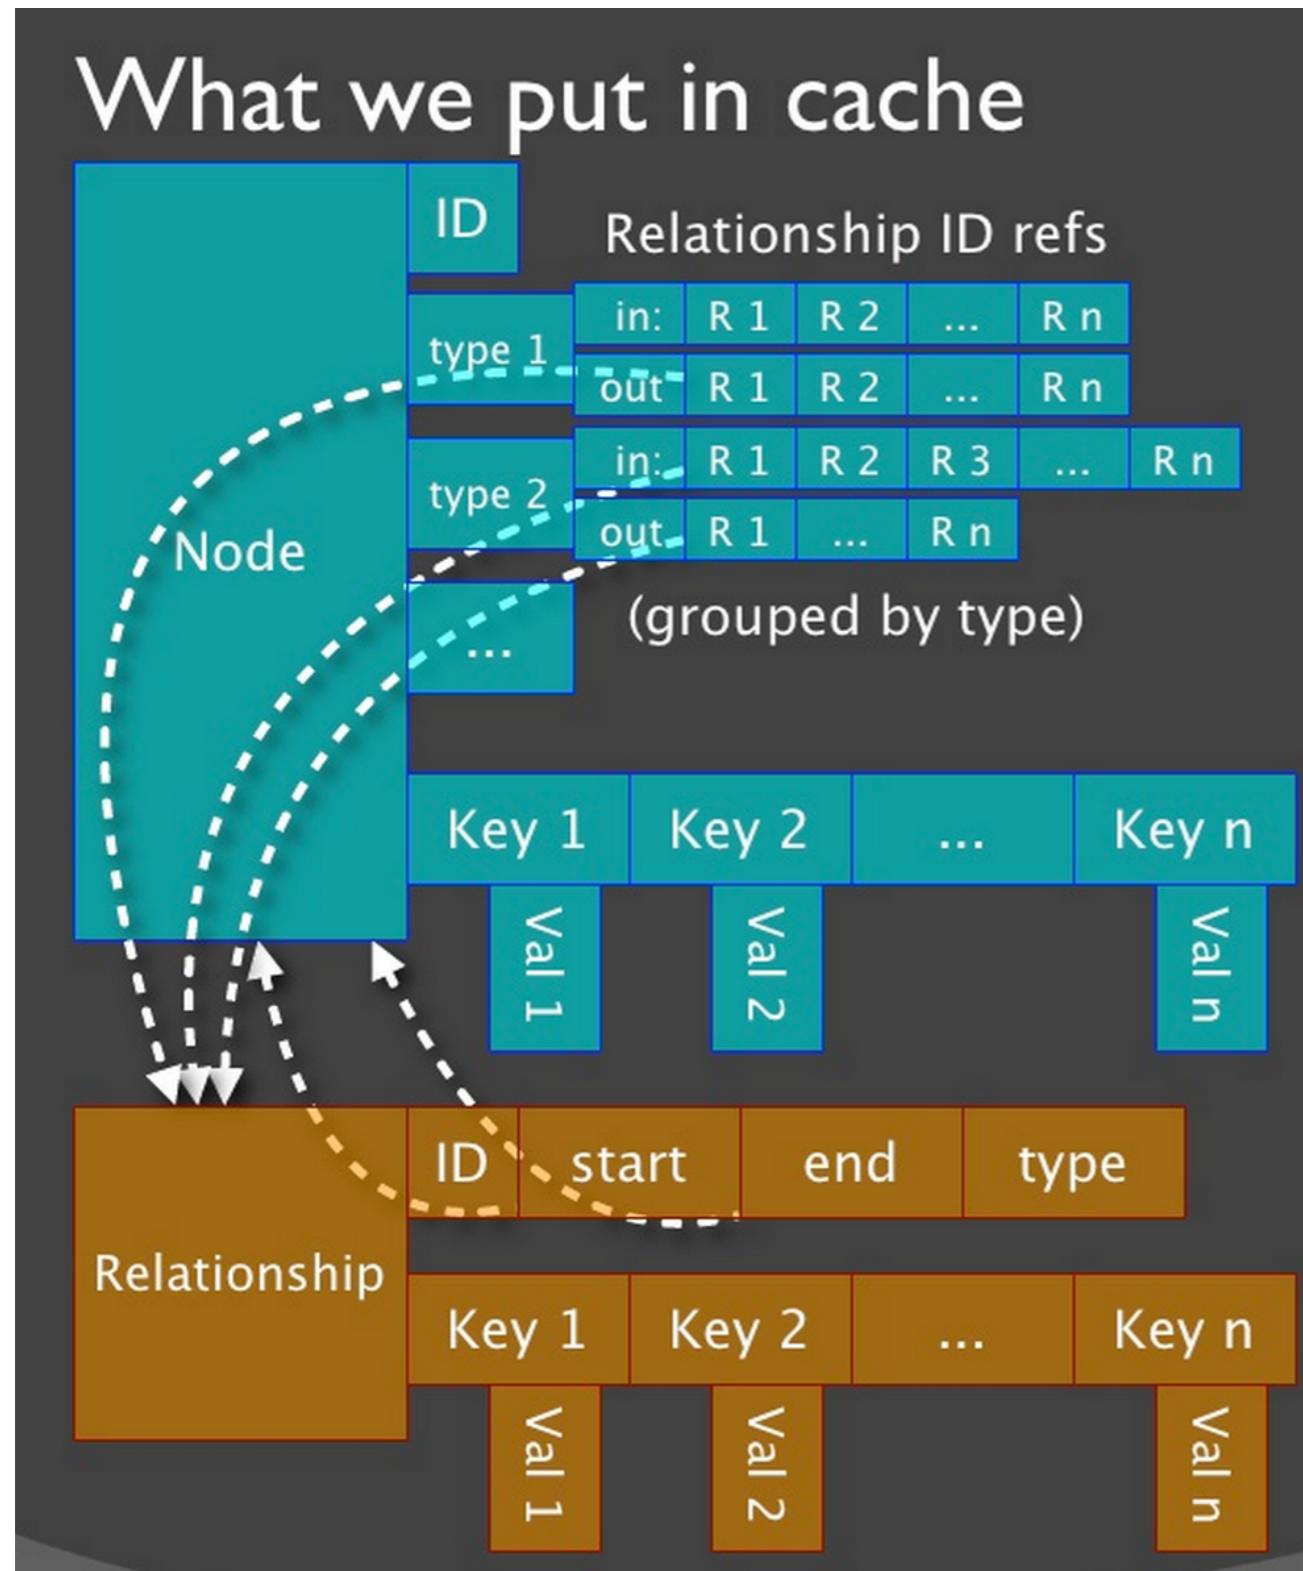
**Conferences**



# A Property Graph



# The Neo4j Secret Sauce



- Pointers instead of look-ups
- Do all your “Joining” on creation
- Spin spin spin through this data structure

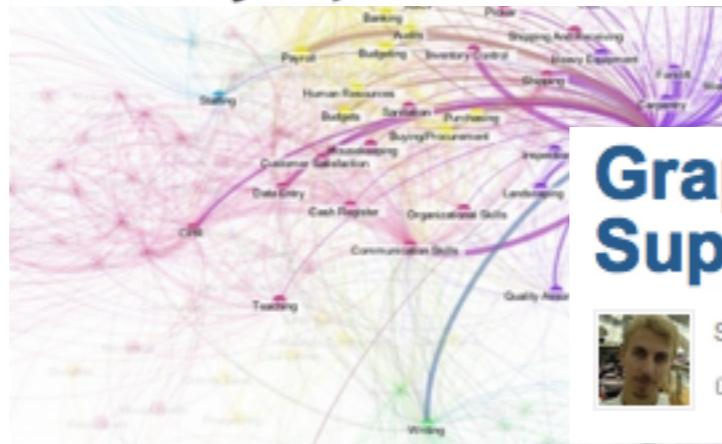
# Graph Buzz!

MacArthur 'Genius Grant' Winner Maria Chudnovsky on Graph Theory

Wednesday, October 03, 2012

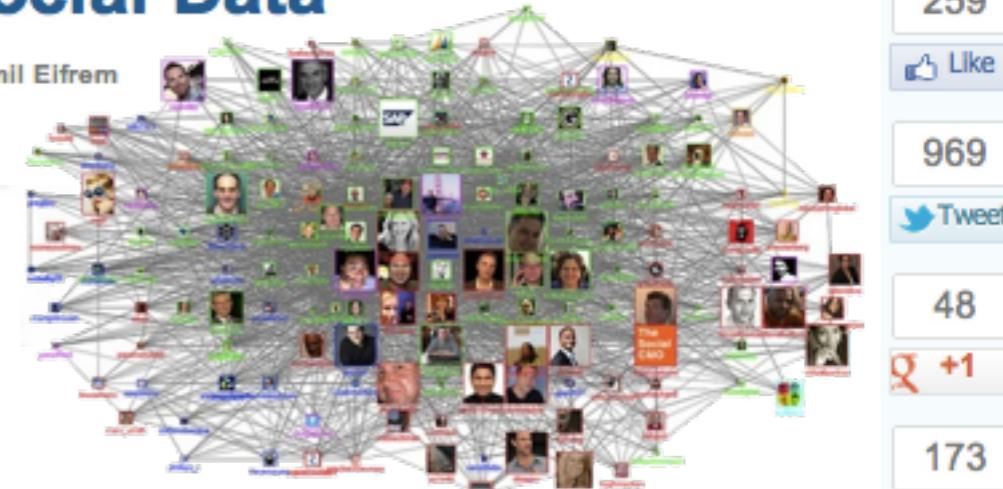
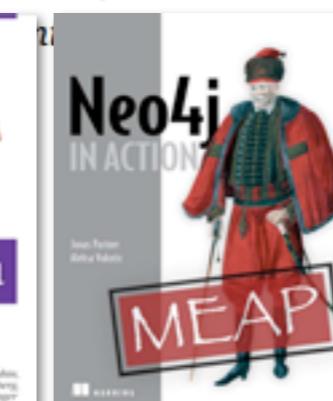
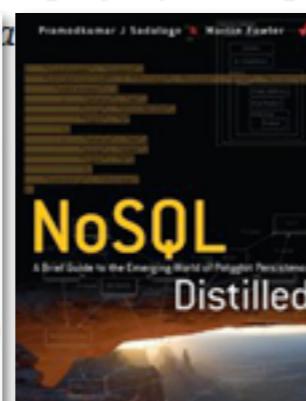
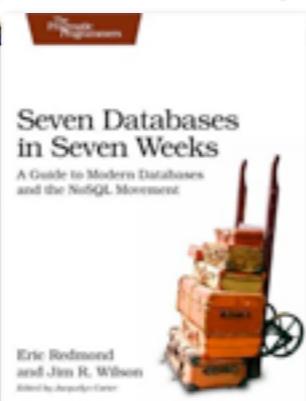
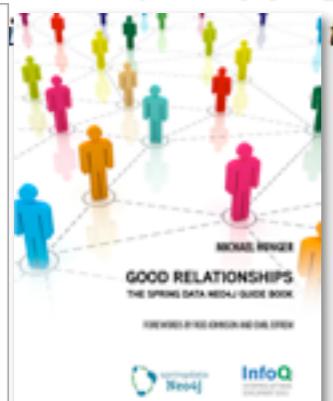
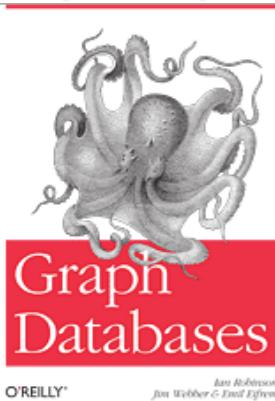


## Bright Launches Bright Packed With Jobs Data Seeking Tips



**Facebook's Social Graph, Neo4j show rising use of graph databases**

**Summary:** Facebook's Social Graph -- the database underlying its Graph Search engine unveiled yesterday-- is just one of many graph databases being employed for complex, connected data. Neo4j



I saw my own Interest Graph and it's scary-accurate. We'd certainly pay for the ability to use the Gravity personalization technology I saw today at TechCrunch to help target content to users.

TC Michael Arrington, TechCrunch

[InfoWorld Home](#) / [InfoWorld Tech Watch](#) / Buzz grows around graph databases



The First Word on Tech  
**INFOWORLD TECH WATCH**

AUGUST 29, 2012

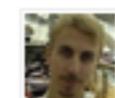
## Buzz grows around graph databases

Interest in graph databases will continue to grow, given its ability to analyze data delivered in a non-relational format, such as social networking data

By Paul Krill | InfoWorld

Follow @pjkrill

## Graph Databases: The New Way to Access Super Fast Social Data



September 26, 2012 by Emil Eifrem

1

259

Like

969

Tweet

48

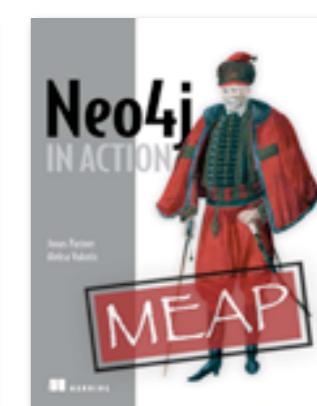
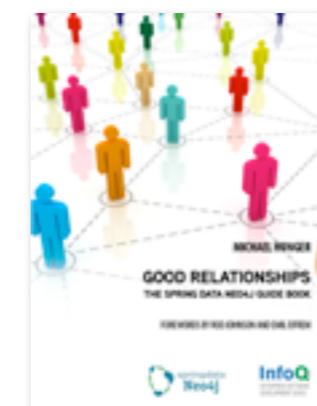
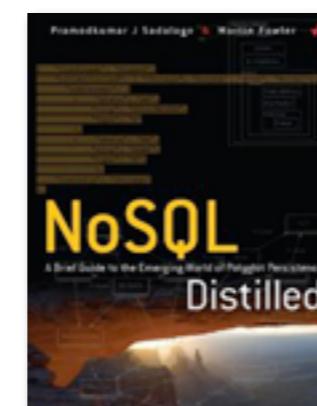
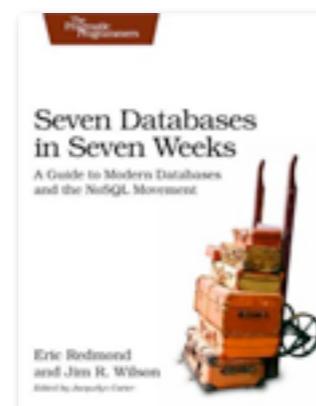
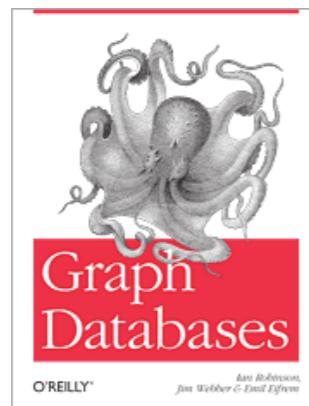
+1

173

Share

# The Neo4j Graph Database

- Neo4j is the **leading graph database** in the world today
  - Most widely deployed: **500,000+** downloads
  - Largest ecosystem: active forums, code contributions, etc
  - Most mature product: in development since 2000, in **24/7 production** since **2003**



neo4j leaves me speechless. Good job at building the best graph database in the world! ★

2:41 AM Mar 24th via Tweetie  
Retweeted by you and 1 other

Reply Retweeted (Undo)

 **fokussiertnet**  
Daniel

If #Cassandra rules its league (#distributed #decentralized #ColumnFamily) then #Neo4j does it in its own, which is #GraphDB. #noSQL ★

3:12 AM Apr 8th via TweetDeck  
Retweeted by you

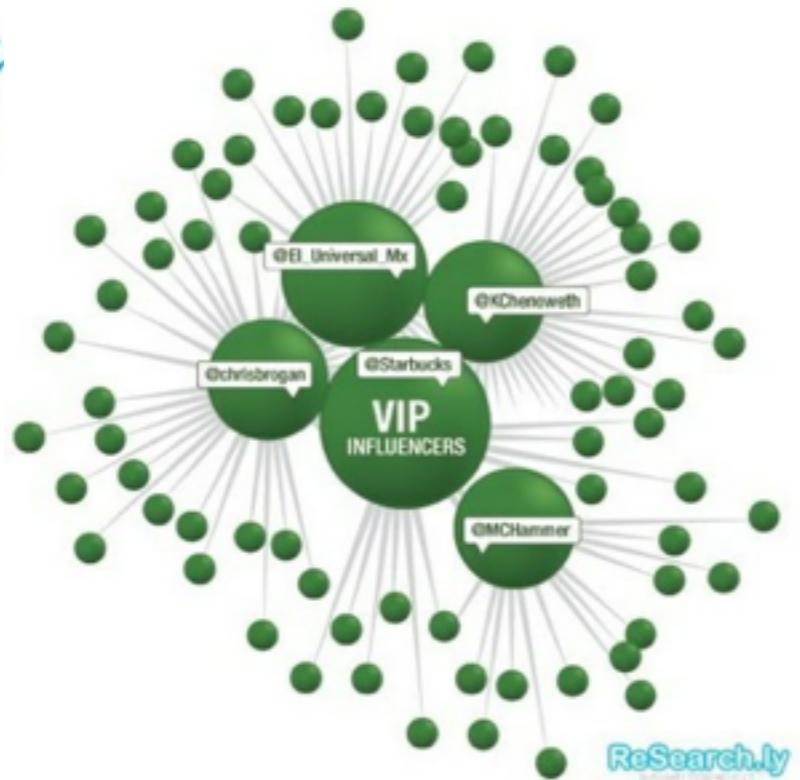
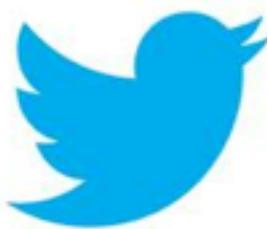
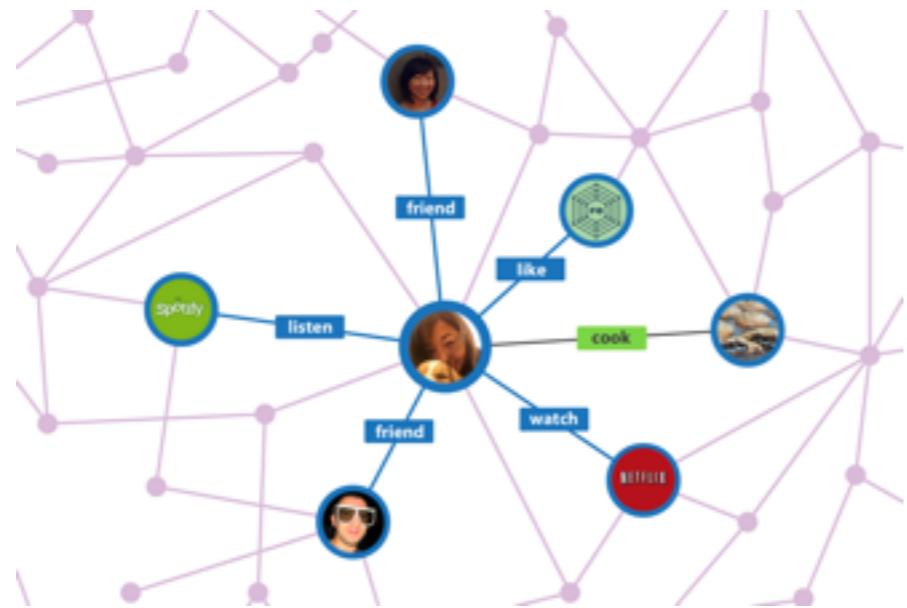
Reply Retweeted (Undo)

 **alisohani**  
Ali Sohani

# Early Adopters of Graph Tech



facebook.



Google™



# Survival of the Fittest

## Evolution of Web Search

### Pre-1999

#### WWW Indexing

**Index**

Page numbers in bold then refer to key term definitions  
 Page numbers in italics refer to images or diagrams  
 Page numbers followed by a “t” indicate a table

**A**

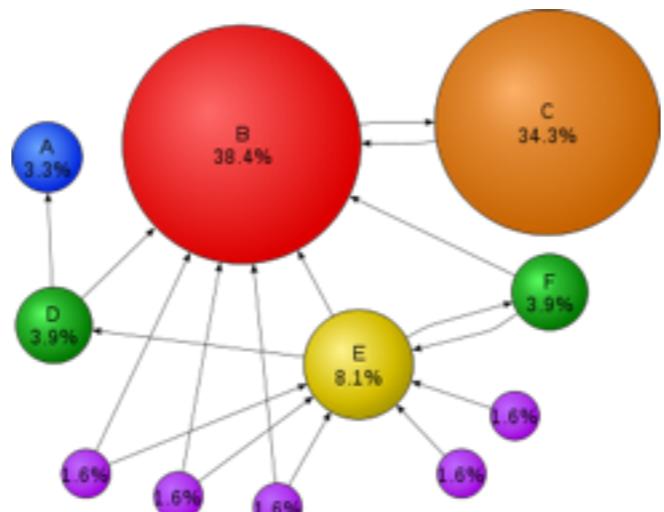
- absolute temperature scale, 380–381
- absolute zero, 381
- acceleration of gravity, A.23t
- accuracy, A.5
- acetic acid ( $\text{CH}_3\text{COOH}$ )

  - buffer, 575–576, 581–582
  - conjugate acid-base pairs, 540
  - ionization constant, 553, 558
  - manufacture of, 451
  - titration, 590–592
  - as weak acid, 148, 149, 510–512
  - acid-base pairs, conjugate, 848–849
  - acid-base reactions, 538
  - ionization of water, 540–542
  - gas-forming exchange, 550–551
  - net ionic equations for, 148–150
  - nernst equation, 146–150, 261–266
  - of salts, 148–151, 561–566
  - acid-base titrations, 597–598
  - amide solutions, 846, 850–852, A.32n–A.33t
  - and ionization constant, 850–851
  - and ionization constant ( $K_{\text{a}}$ ). *See* ionization constants, acids ( $K_{\text{a}}$ )
  - anhydrides, 143. *See also* acid-base reactions, ionization constants, acidic ( $K_{\text{a}}$ ), specific esters, e.g., carboxylic acids, lacto acids
  - Brensted-Lowry concept, 533–544
  - buffer solutions, 575–586
  - conjugate acid-base pairs, 540–544
  - equilibrium constant, 47%
  - ionization constants, 510–511, A.23t
  - Lewis, 546–548
  - organic, 544
  - pH scale, 547–550
  - properties, 143–145
  - solubility of salts, 597–598
  - solutions, 846, 850–852, A.32n–A.33t
  - strengths, 143–146, 533–536
  - titrations, 597–598
  - water's role, 540
  - weak acids, 55, 210–211
  - activated complex, 433
  - activation energies ( $E_a$ ), 434, 438–440, 443, 447, 449–450
  - activation energy, 449
  - activation, 466, 247, 703–706
  - activity series, metals, 159–160
  - actual yields, 123
  - addition, 43, A.6, A.9
  - addition, significant figures in, A.6

Discrete Data 

### 1999 - 2012

#### Google Invents PageRank

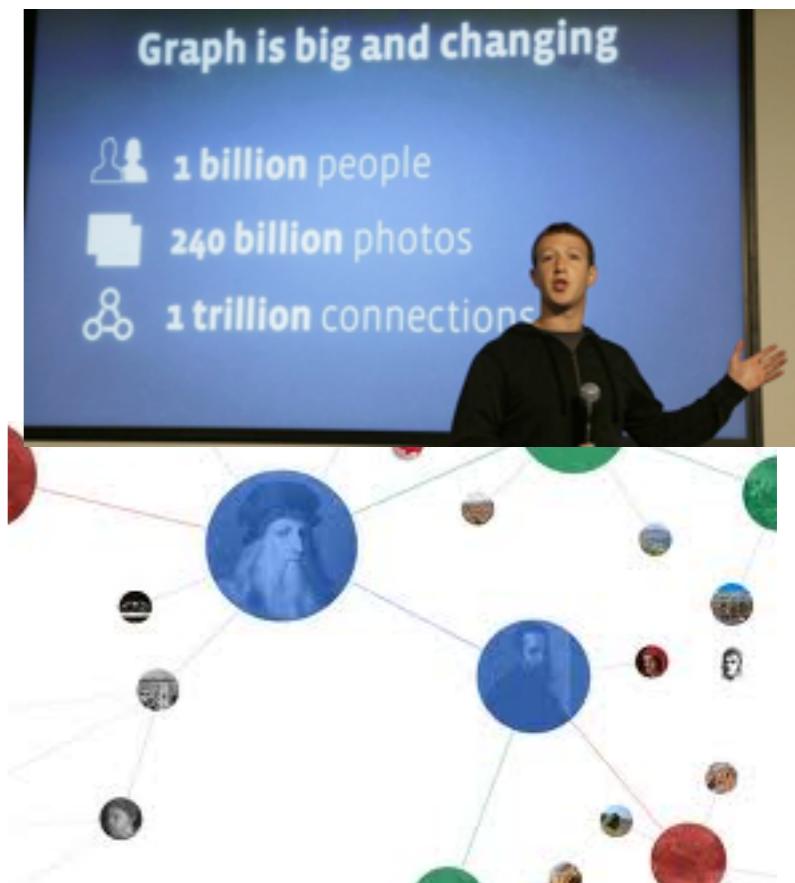


Connected Data  
(Simple)



#### 2012-?

#### Google Knowledge Graph, Facebook Graph Search



Connected Data  
(Rich)

# Open Source Example

Neo Graph Search    Home    Profile    Graph Search    Likes    Friends    Visualization    [Max De Marzi](#)

## Graph Search

friends who like Neo4j    [Search](#)    Try: [friends who like wine](#) , [people who like wine and cheese](#) , [people who like cycling and live in Florida](#)

**Cypher Query:**

```
START me = node({me}), thing = node:things({thing})
MATCH me -[:friends]-> people, people -[:likes]-> thing
RETURN DISTINCT people , people.uid, people.name, people.image_url
LIMIT 100
Parameters: {"me"=>1, "thing"=>"name: Neo4j"}
```

Want your own Graph Search? Contact [me](#) and learn more about [Neo4j](#) and [NeoTechnology](#)

 Peter Neubauer

 Andres Taylor

<http://maxdemarzi.com/2013/01/28/facebook-graph-search-with-cypher-and-neo4j/>

# Survival of the Fittest

## Evolution of Online Recruiting

### 1999

#### Keyword Search

[Home](#) > **Job Search**

Search over **150,000** U.S. jobs.

Perform your search below or get [tips on searching](#).

#### Location Search:

----- Select a location -----

- Alaska-Anchorage
- Alaska-Fairbanks
- Alaska-Juneau
- Alaska-Valdez
- Alabama-Anniston

#### Category Search:

----- Select a category -----

- Accounting/Finance/Banking
- Administrative/Clerical
- Creative Arts/Media
- Education/Training
- Engineering/Architecture/Design

#### Keyword Search:

[Search Jobs](#) [Clear](#)

Discrete Data



### 2011-12

#### Social Discovery



#### Most jobs are found through an inside connection

Each friend that joins Glassdoor allows you to see more connections at more companies

9 friends on glassdoor	3,905 inside connections	3,882 companies
------------------------	--------------------------	-----------------

[Invite more friends](#) — ask them to share their connections

#### Jobs with Connections

##### [Sr. Statistical Analyst, Product Innovation](#)

Netflix — Los Gatos, CA

From: Job.com — 1 days ago

##### [Creative Director](#)

frog design — San Francisco, CA

From: Experiteer — 8 days ago

##### [Java Server Software Engineer](#)

Electronic Arts — Redwood City, CA

From: Experiteer — 3 days ago

##### [EMERGENCY MEDICAL TECH](#)

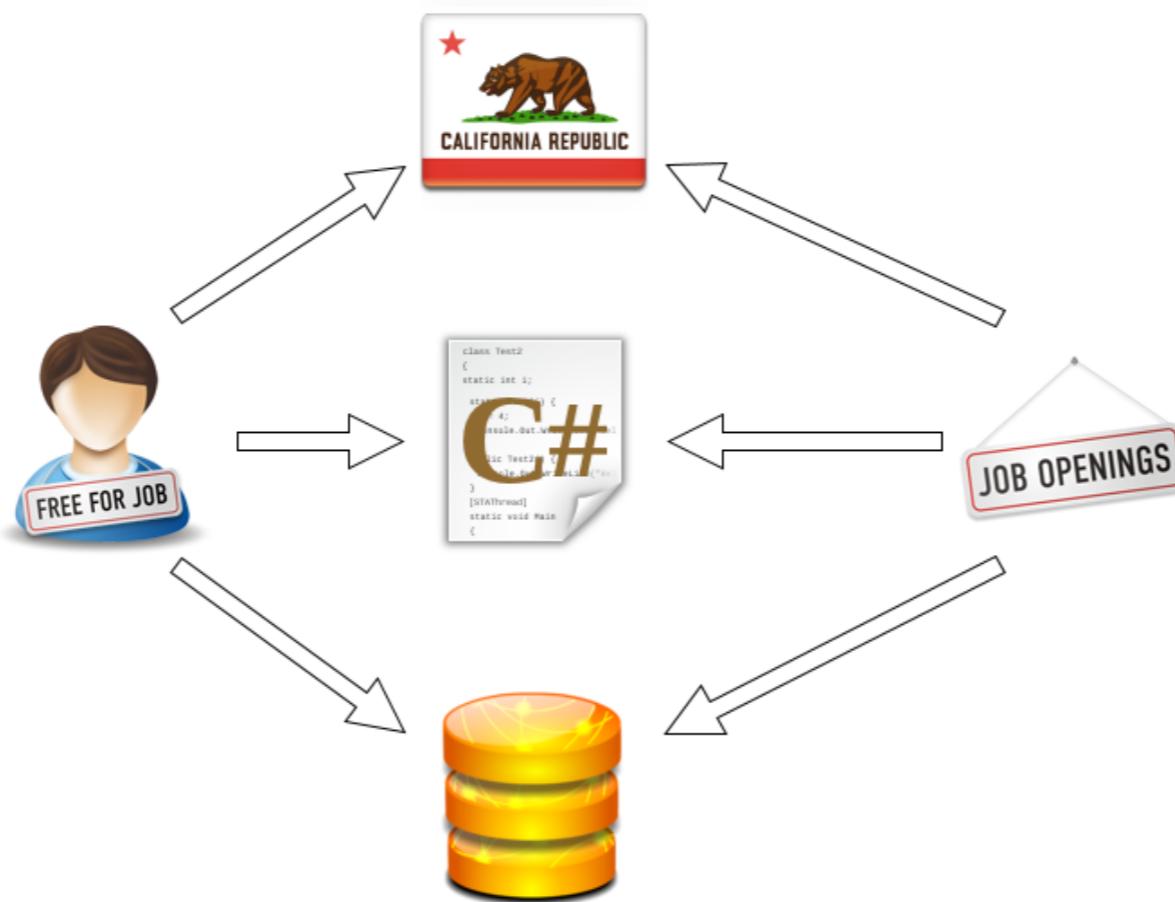
U. S. NAVY — Fremont, CA

From: Monster — 14 days ago

[Want better jobs? Tell us your current job title](#)

Connected Data

# Open Source Example



<http://maxdemarzi.com/2012/10/18/matches-are-the-new-hotness/>

# Open Source Example

```
START me=node:users_index(name={user})
MATCH skills<-[ :has]-me-[ :lives_in]->city<-[ :in_location]-job-[ :requires]->requirements
WHERE me-[ :has]->()-<-[ :requires]-job
WITH DISTINCT city.name AS city_name,
      job.name AS job_name,
      LENGTH(me-[ :has]->()-<-[ :requires]-job) AS matching_skills,
      LENGTH(job-[ :requires]->()) AS job_requires,
      COLLECT(DISTINCT requirements.name) AS req_names,
      COLLECT(DISTINCT skills.name) AS skill_names
RETURN city_name, job_name,
       FILTER(name IN req_names WHERE NOT name IN skill_names) AS missing
ORDER BY matching_skills / job_requires DESC, job_requires
LIMIT 10
```

<http://maxdemarzi.com/2012/10/18/matches-are-the-new-hotness/>

# Open Source Example

## Matching Jobs for Daniel

Reload the page or click [Match Another User](#)

Job	Match	Missing
Spring-C-Django-SQL in <i>Dallas</i>	100% Match	
Neo4j-C-CSS in <i>Dallas</i>	Partial Match	<span>+ CSS</span>
C-Spring-Java in <i>Dallas</i>	Partial Match	<span>+ Java</span>
Java-Redis-C-Neo4j in <i>Dallas</i>	Partial Match	<span>+ Java</span> <span>+ Redis</span>

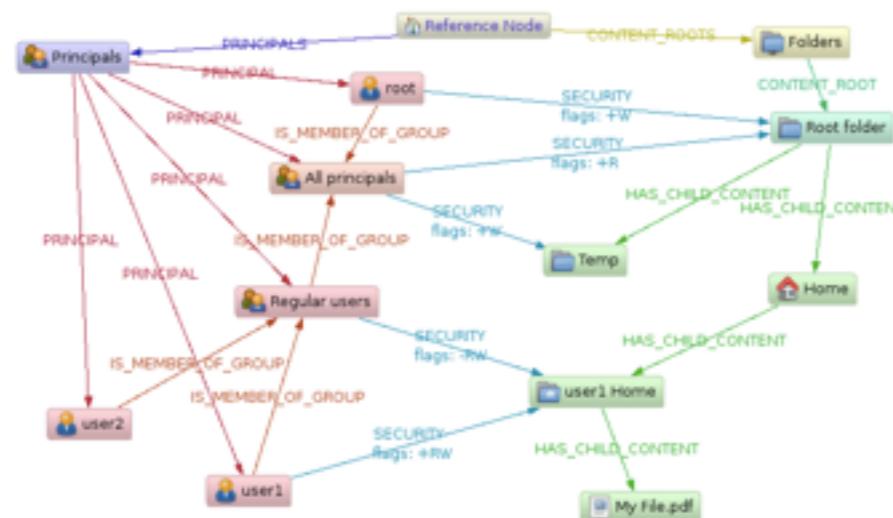
<http://maxdemarzi.com/2012/10/18/matches-are-the-new-hotness/>

# Emergent Graph in Other Industries

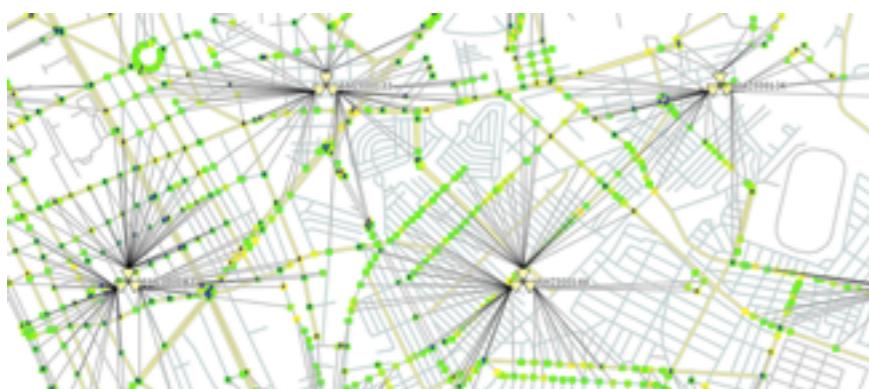


(Actual Neo4j Graphs)

## Content Management & Access Control



## Network Cell Analysis



## Bioinformatics



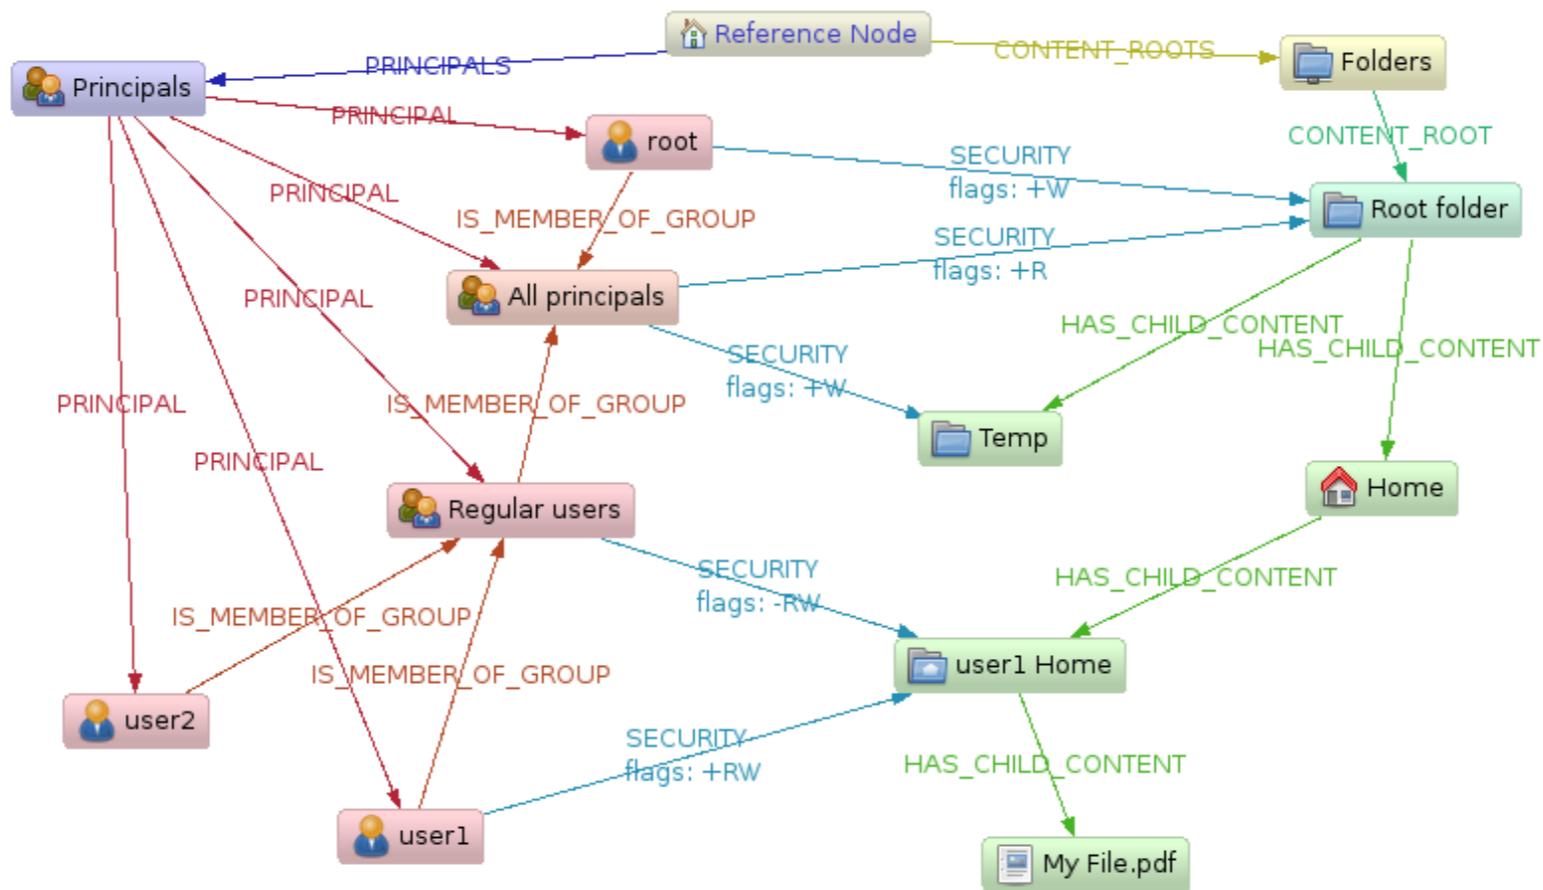
## Insurance Risk Analysis



## Network Asset Management



# Open Source Example



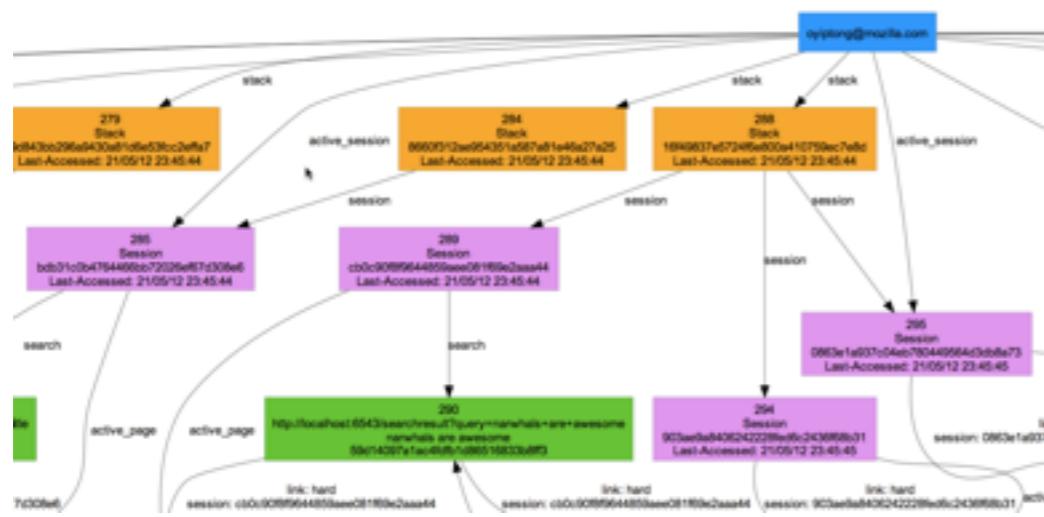
<http://maxdemarzi.com/2013/03/18/permission-resolution-with-neo4j-part-1/>

# Emergent Graph in Other Industries

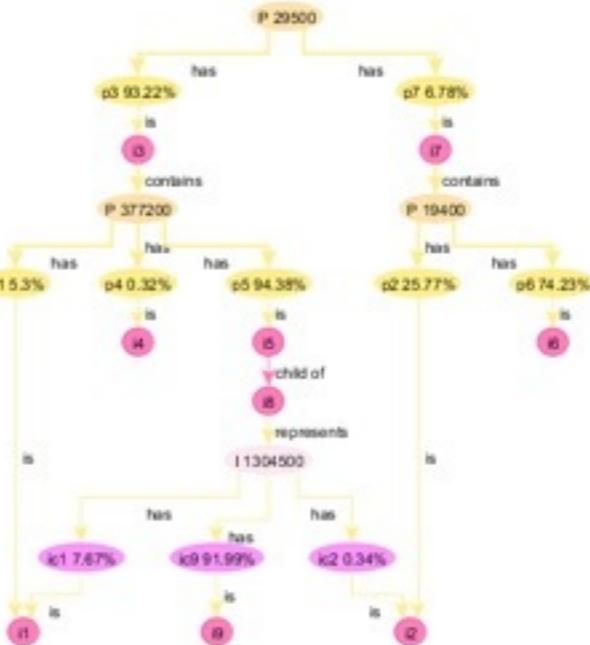
(Actual Neo4j Graphs)



## Web Browsing



## Portfolio Analytics



# Open Source Example

Neo Love Matching

[Match Another User](#)

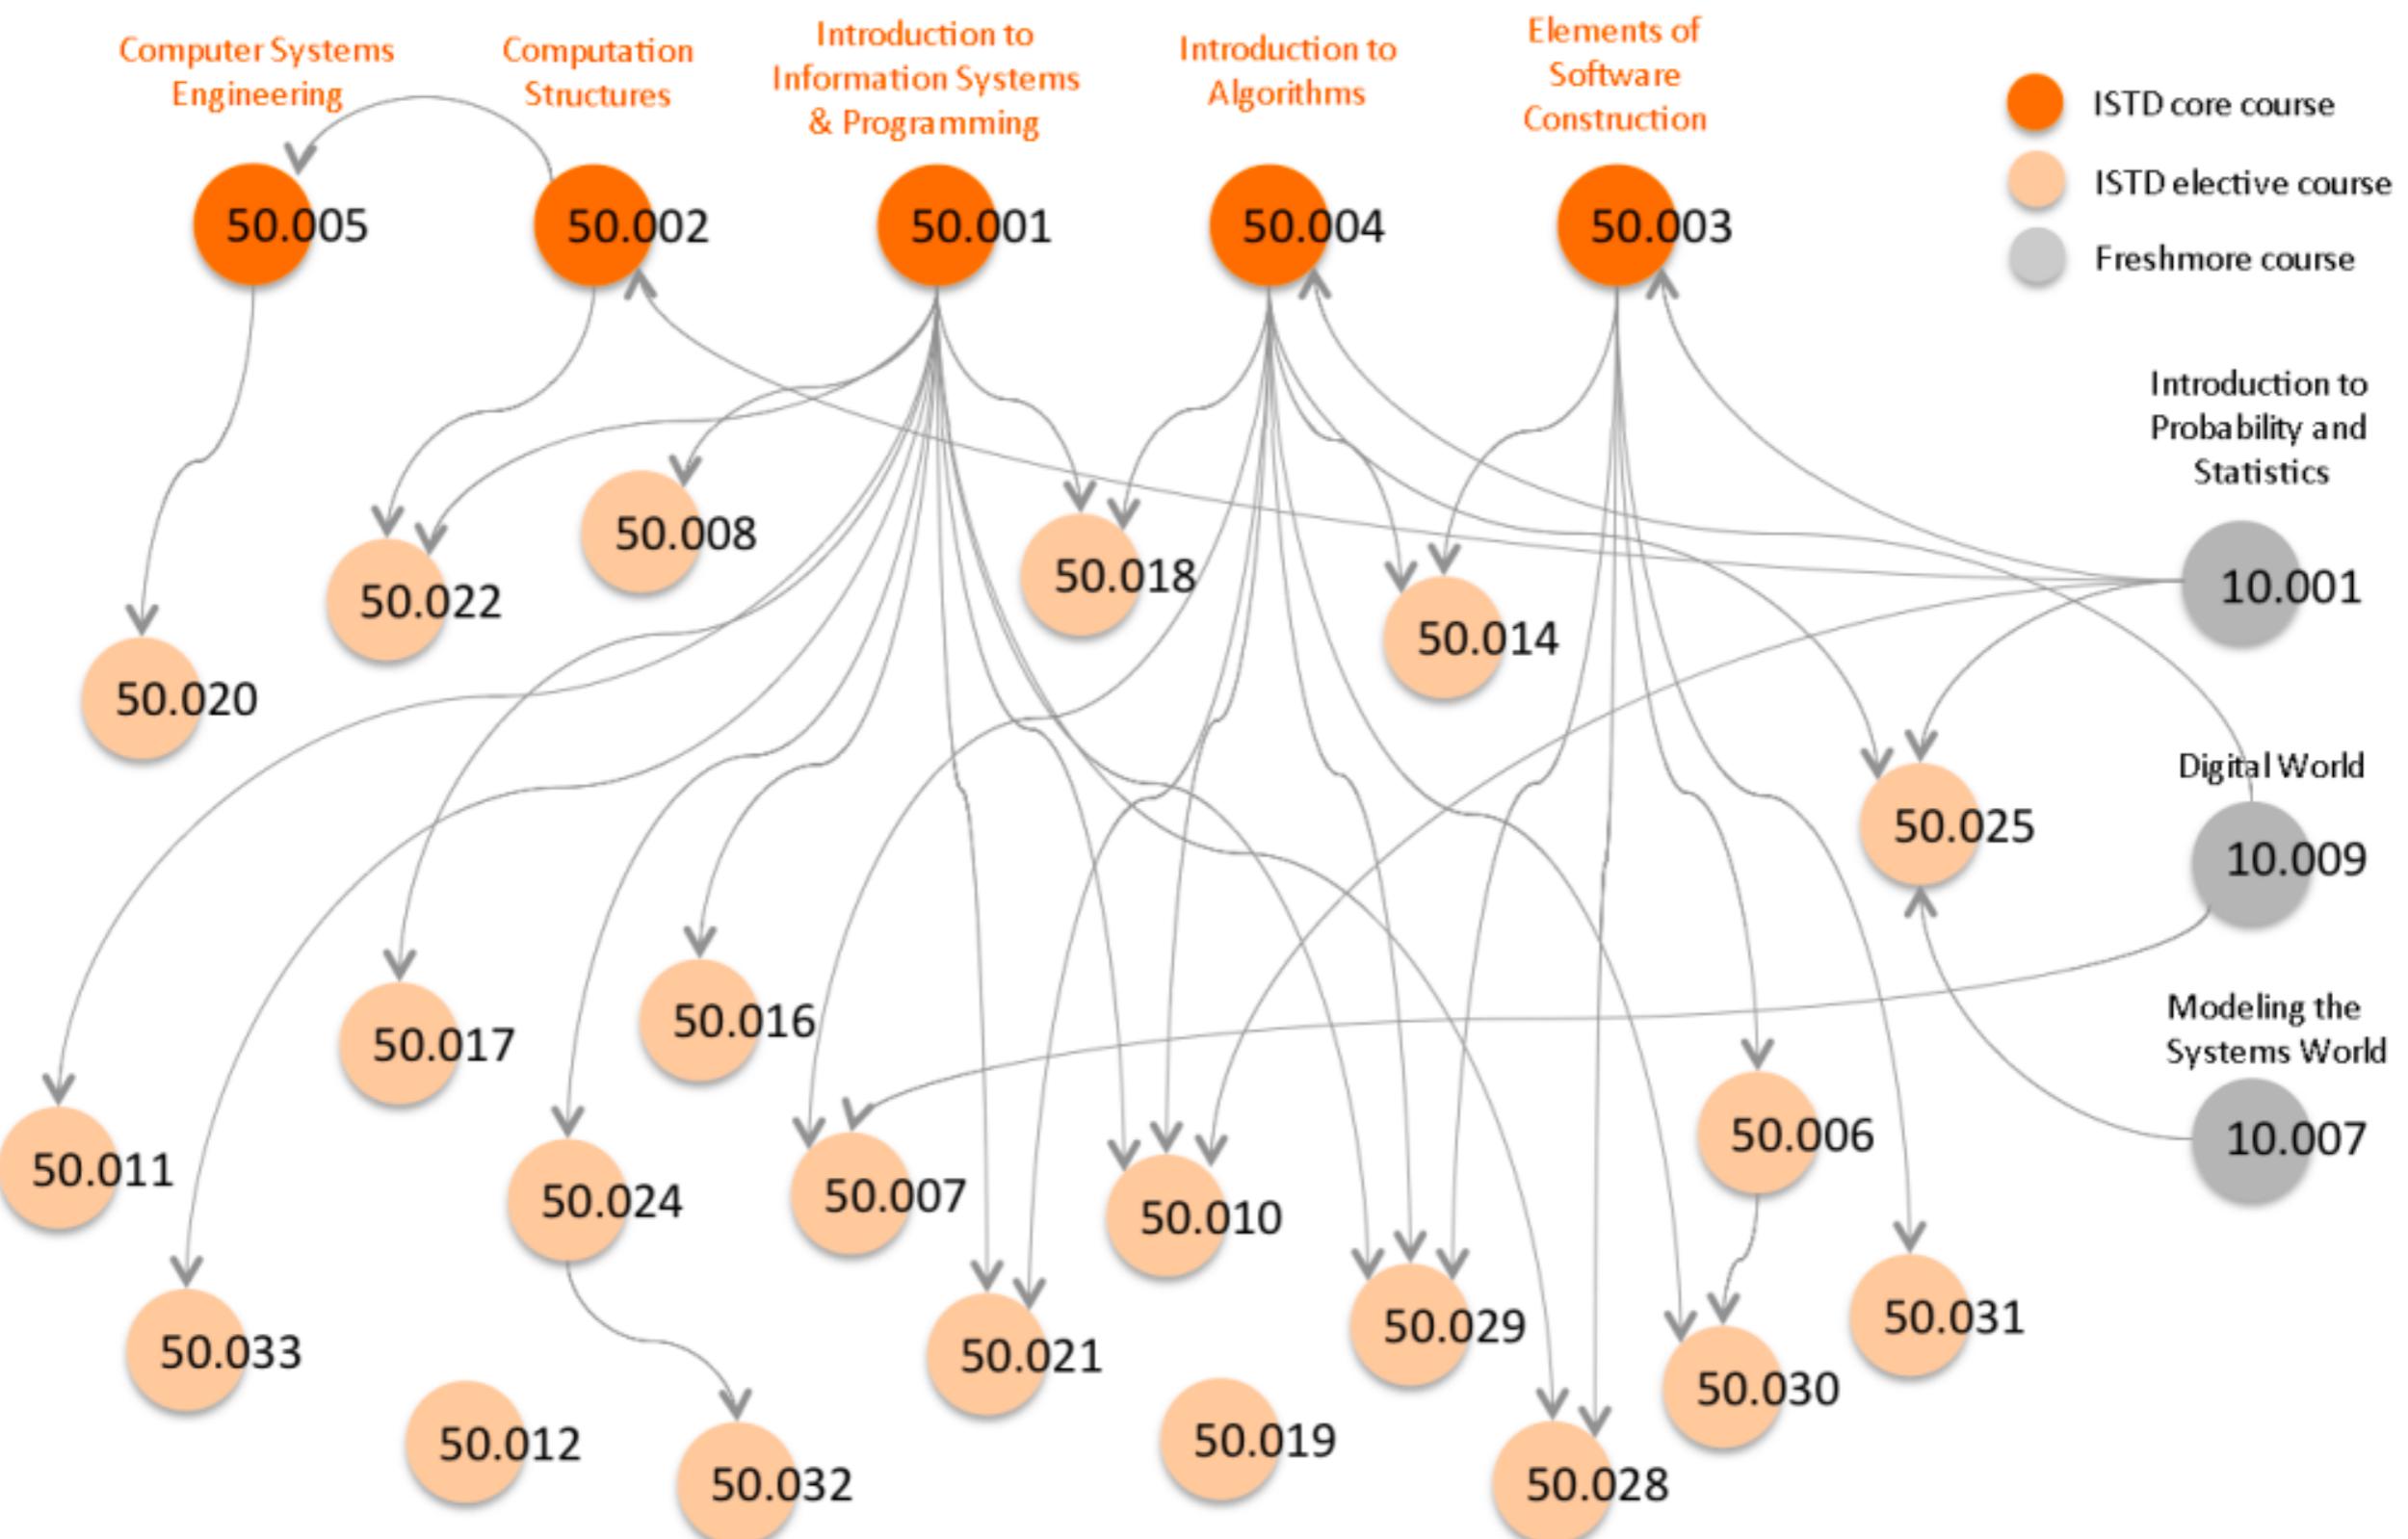
## Matching Loves for Luvenia

*Reload the page or click [Match Another User](#)*

Love	Compatibility	Your Matches	Their Matches
Chet in Philadelphia	2.0	<span style="background-color: #4CAF50; color: white; padding: 2px 10px; border-radius: 5px;">✓ Humorous</span> <span style="background-color: #4CAF50; color: white; padding: 2px 10px; border-radius: 5px;">✓ Thoughtful</span>	<span style="background-color: #4CAF50; color: white; padding: 2px 10px; border-radius: 5px;">✓ Likeable</span>
Mckinley in Philadelphia	1.0	<span style="background-color: #4CAF50; color: white; padding: 2px 10px; border-radius: 5px;">✓ Serious</span>	<span style="background-color: #4CAF50; color: white; padding: 2px 10px; border-radius: 5px;">✓ Creative</span>
Gene in Philadelphia	1.0	<span style="background-color: #4CAF50; color: white; padding: 2px 10px; border-radius: 5px;">✓ Serious</span>	<span style="background-color: #4CAF50; color: white; padding: 2px 10px; border-radius: 5px;">✓ Likeable</span>
Keith in Philadelphia	1.0	<span style="background-color: #4CAF50; color: white; padding: 2px 10px; border-radius: 5px;">✓ Quiet</span>	<span style="background-color: #4CAF50; color: white; padding: 2px 10px; border-radius: 5px;">✓ Inventive</span>

<http://maxdemarzi.com/2013/04/19/match-making-with-neo4j/>

# Curriculum Graph



# Early Adopter Segments

(What we expected to happen - view from several years ago)



Core Industries & Use Cases:	Web / ISV	Finance & Insurance	Datacom / Telecom
Network & Data Center Management			
MDM			
Social			
Geo			

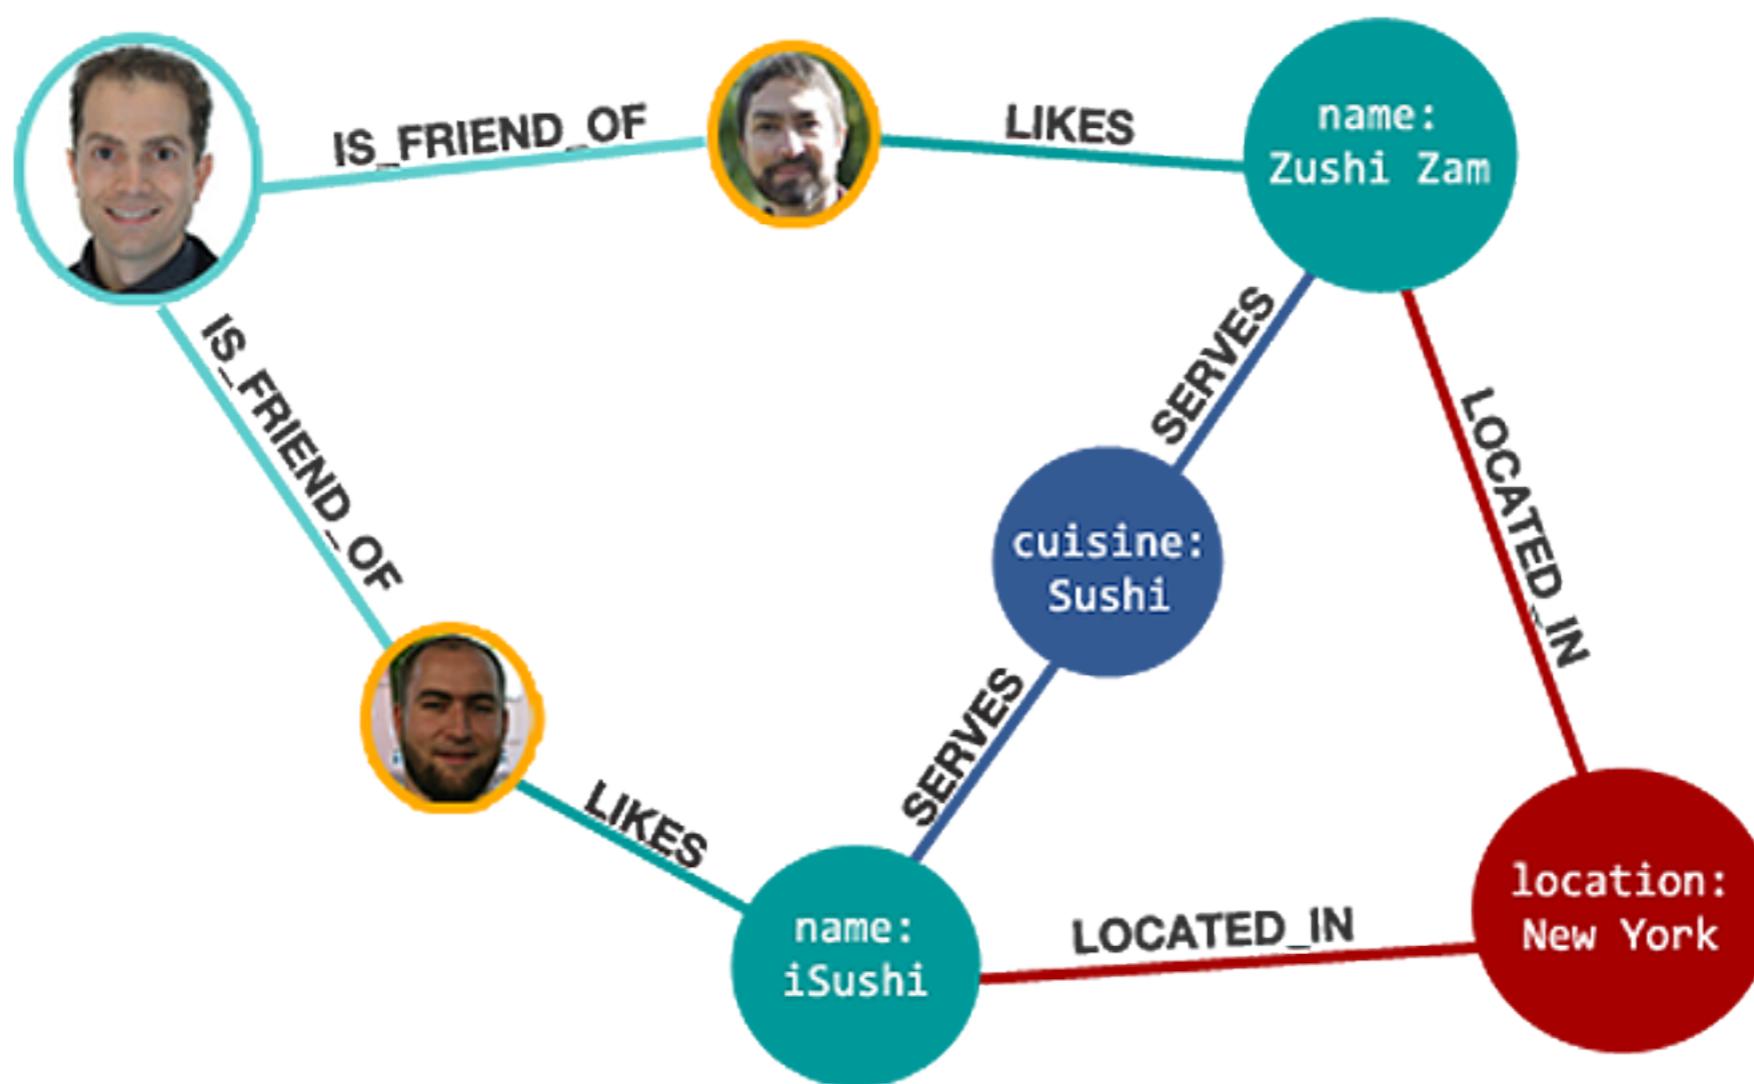
# Neo4j Adoption Snapshot

## Select Commercial Customers (Community Users Not Included)



Core Industries & Use Cases:	Software	Financial Services	Telecommunications	Web Social, HR & Recruiting	Health Care & Life Sciences	Media & Publishing	Energy, Services, Automotive, Gov't, Logistics, Education, Gaming, Other
Network & Data Center Management	Zenoss NetApp SERENA gen VIRTUAL INSTRUMENTS		hp SFR				Juice PLUS+ gonefinestay teachscape
MDM / System of Record			CISCO Deutsche Telekom maaili Let's connect	wooz world EQUILAR viadeo glassdoor	ZEPHYR HEALTH INC HealthUnlocked	LIFECHURCH.TV	
Social	Glowbl	ICE Global Markets in clear view			SharePractice	SQUIDOO indiatimes	gameSYS Accenture Global 500 Logistics shuttle
Geo	DingLicom		Justdial India's No. 1 local search engine	classmates.com			
Identity & Access Mgmt	aikux.com identropy	Global 500 Finance	telenor	Dshini SHRM SOCIETY FOR HUMAN RESOURCE MANAGEMENT	SevenBridges genomics	zeebox LifeWay Biblical Solutions for Life	
Content Management	springcm Adobe			hinge careerbuilder InfoJobs	Curaspan HEALTH GROUP	dedbel <fuseworks/> Perigee CHIP	DOSB NEW MEDIA GMBH DEUTSCHER OLYMPISCHER SPORTBUND
Recommendations	LIQUID COMMON Social Web Mobile	AXON ACTIVE Focusing on your decisions	kitedesk	moviepilot	janssen	DRAKER Impact Technologies A Sikorsky Innovations Company	research now compete
BI, CRM, Impact Analysis, Fraud Detection, Resource Optimization, etc.	SODIFRANCE CONSEIL, TECHNOLOGIES & SERVICES IT idMISSION ENHANCING LEARNING	Global 500 Telcommunication	DRW TRADING GROUP	NexLP	KLM van Kompagnie	LOCKHEED MARTIN	Global 500 Energy Global 500 Aerospace

# What Can You Do With Graphs?

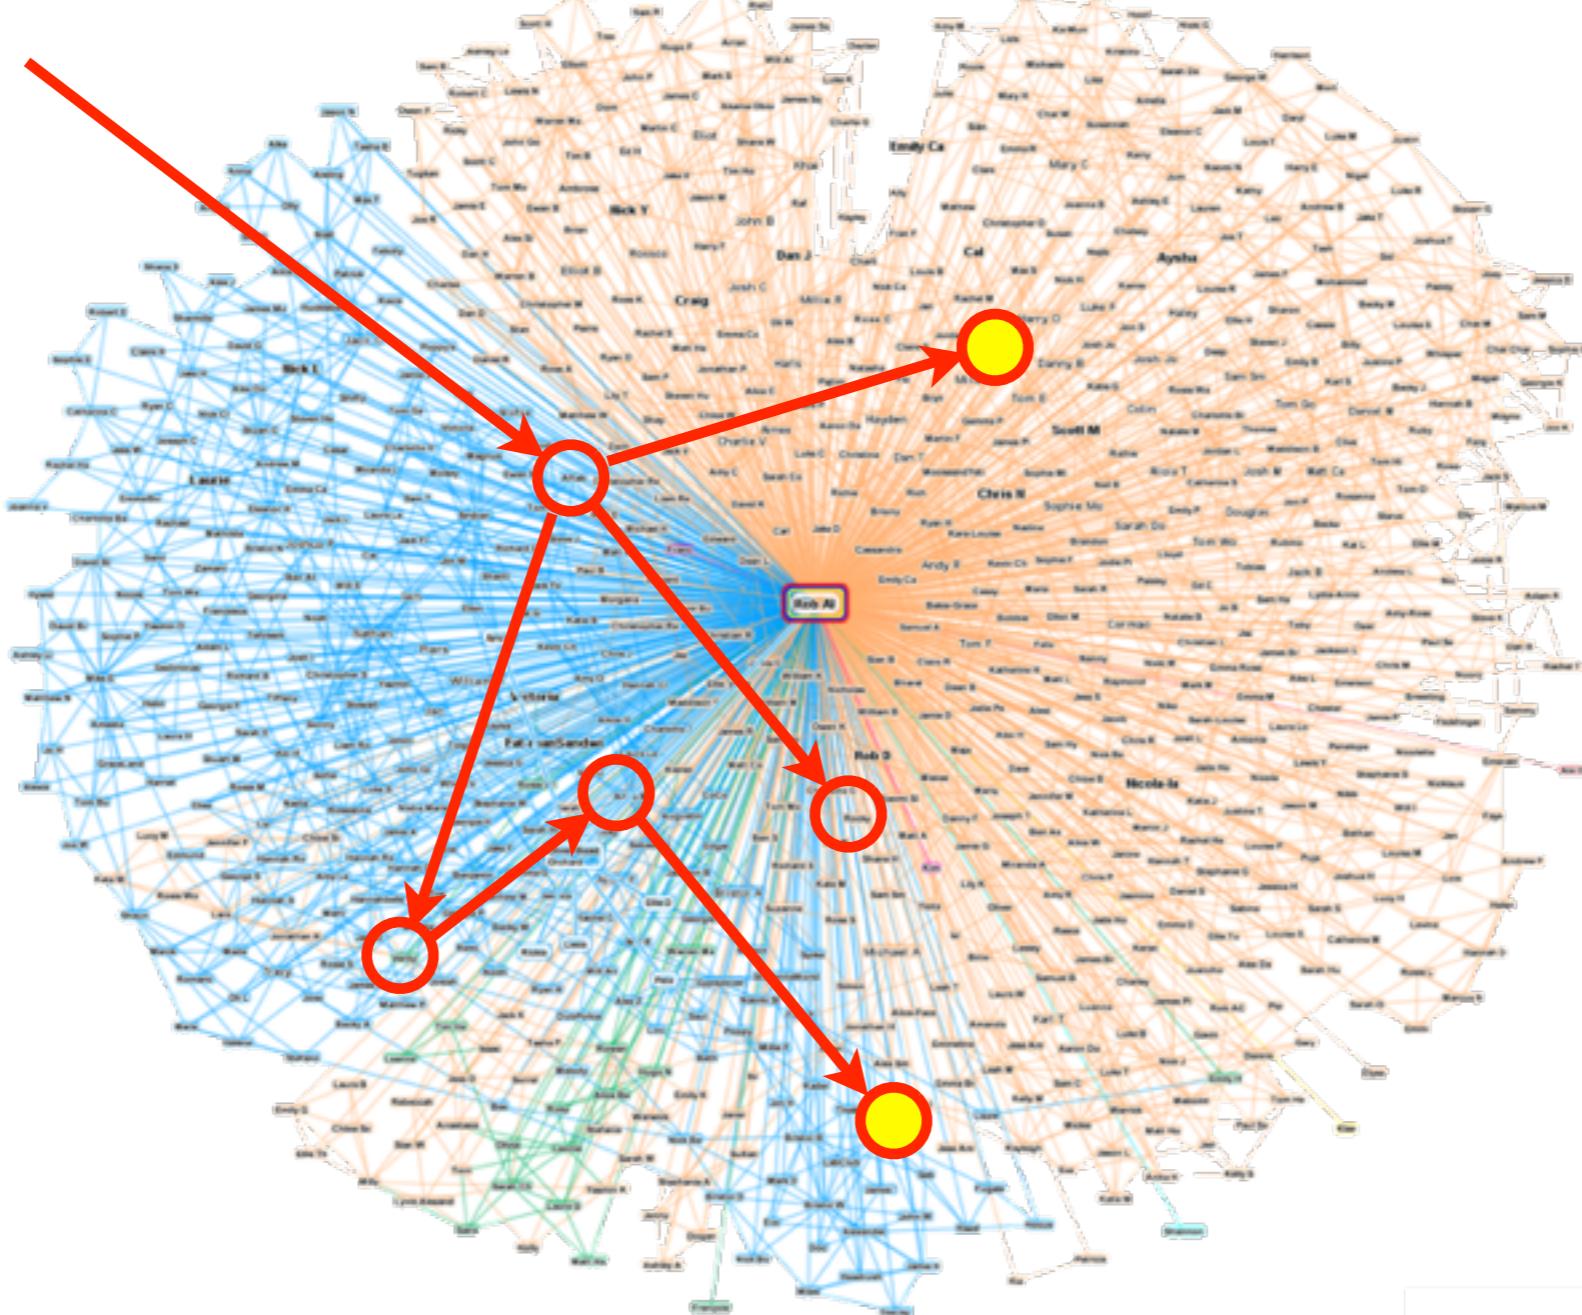




```
MATCH (me:Person) - [:IS_FRIEND_OF] -> (friend),  
(friend) - [:LIKES] -> (restaurant),  
(restaurant) - [:LOCATED_IN] -> (city:Location),  
(restaurant) - [:SERVES] -> (cuisine:Cuisine)
```

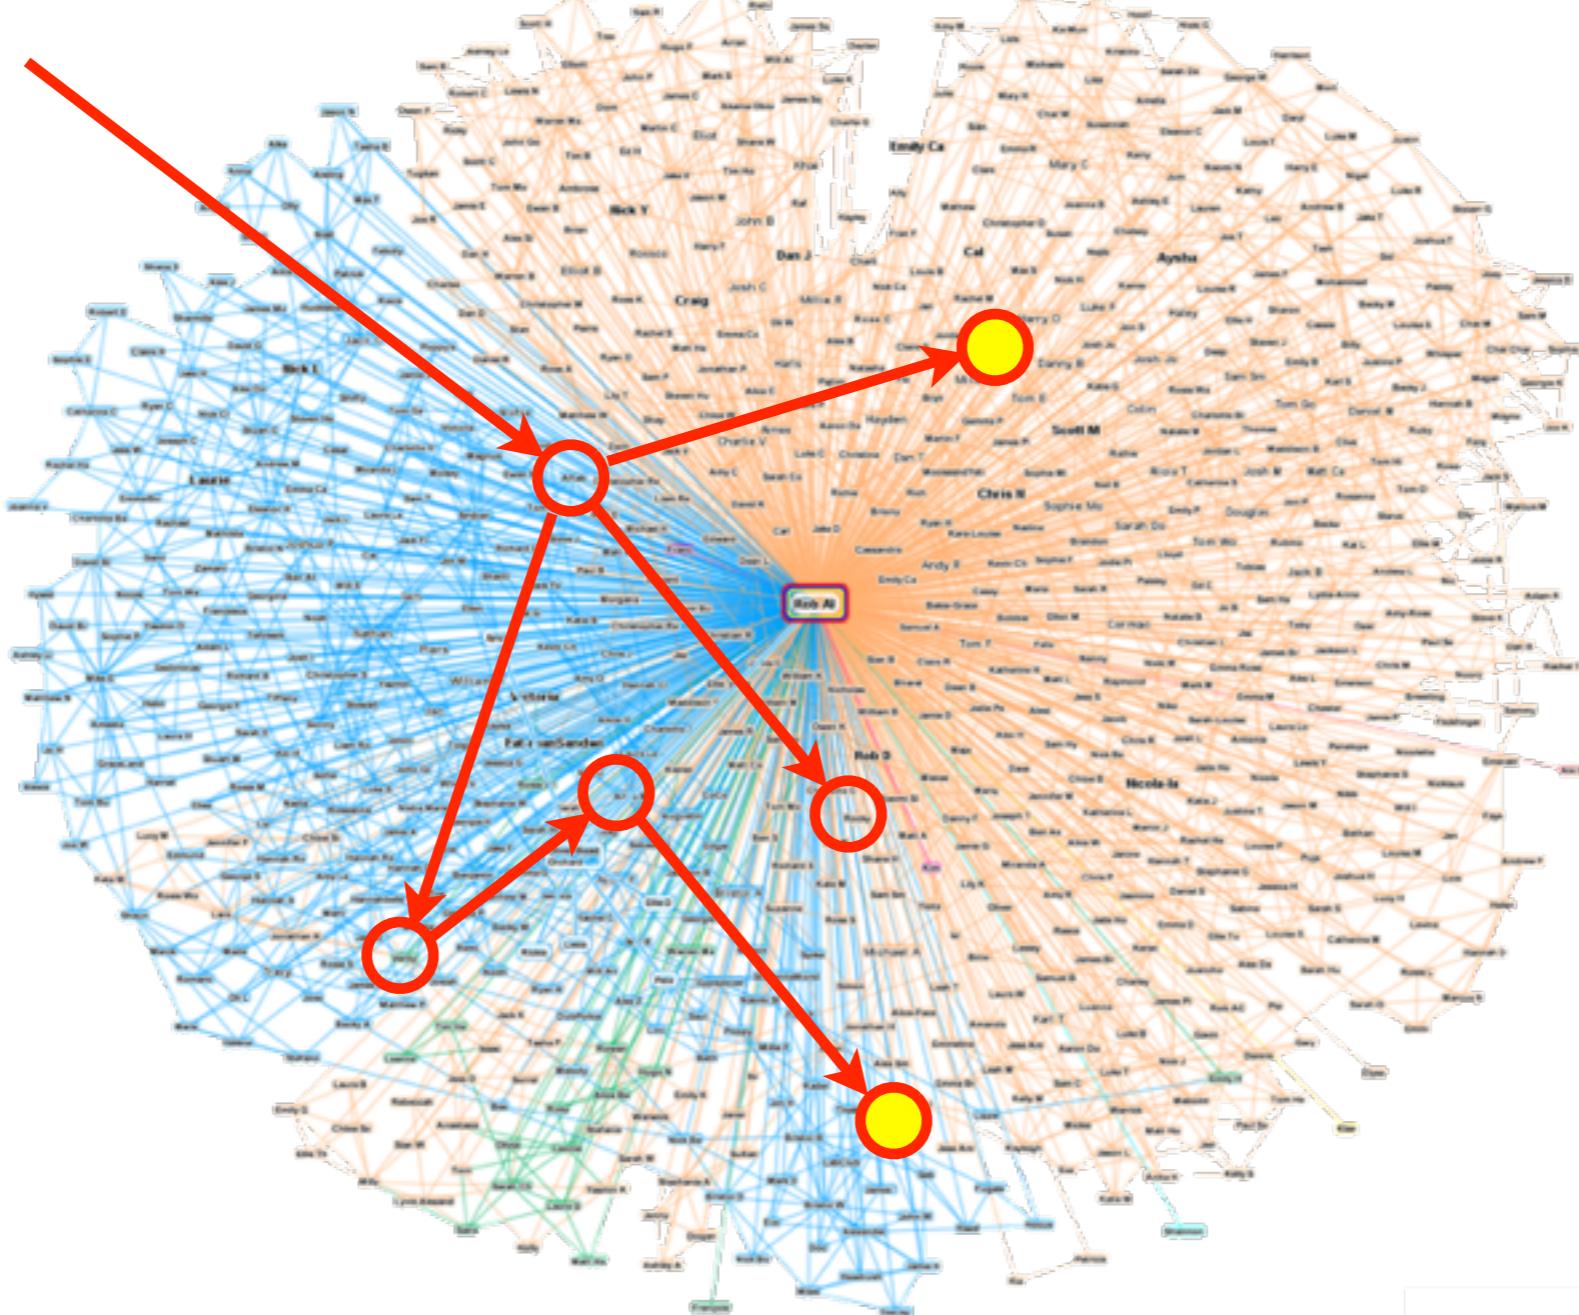
```
WHERE me.name = 'Philip' AND city.location='New York' AND  
cuisine.cuisine='Sushi'
```

```
RETURN restaurant.name
```



# Of course.. a graph is a graph is a graph

*What drugs will bind to protein X and not interact with drug Y?*



# Connected Query Performance

# Connected Query Performance



*Query Response Time\* =  
f(graph density, graph size, query degree)*

- **Graph density** (avg # rel's / node)
- **Graph size** (total # of nodes in the graph)
- **Query degree** (# of hops in one's query)

## RDBMS:

>> exponential slowdown as each factor increases

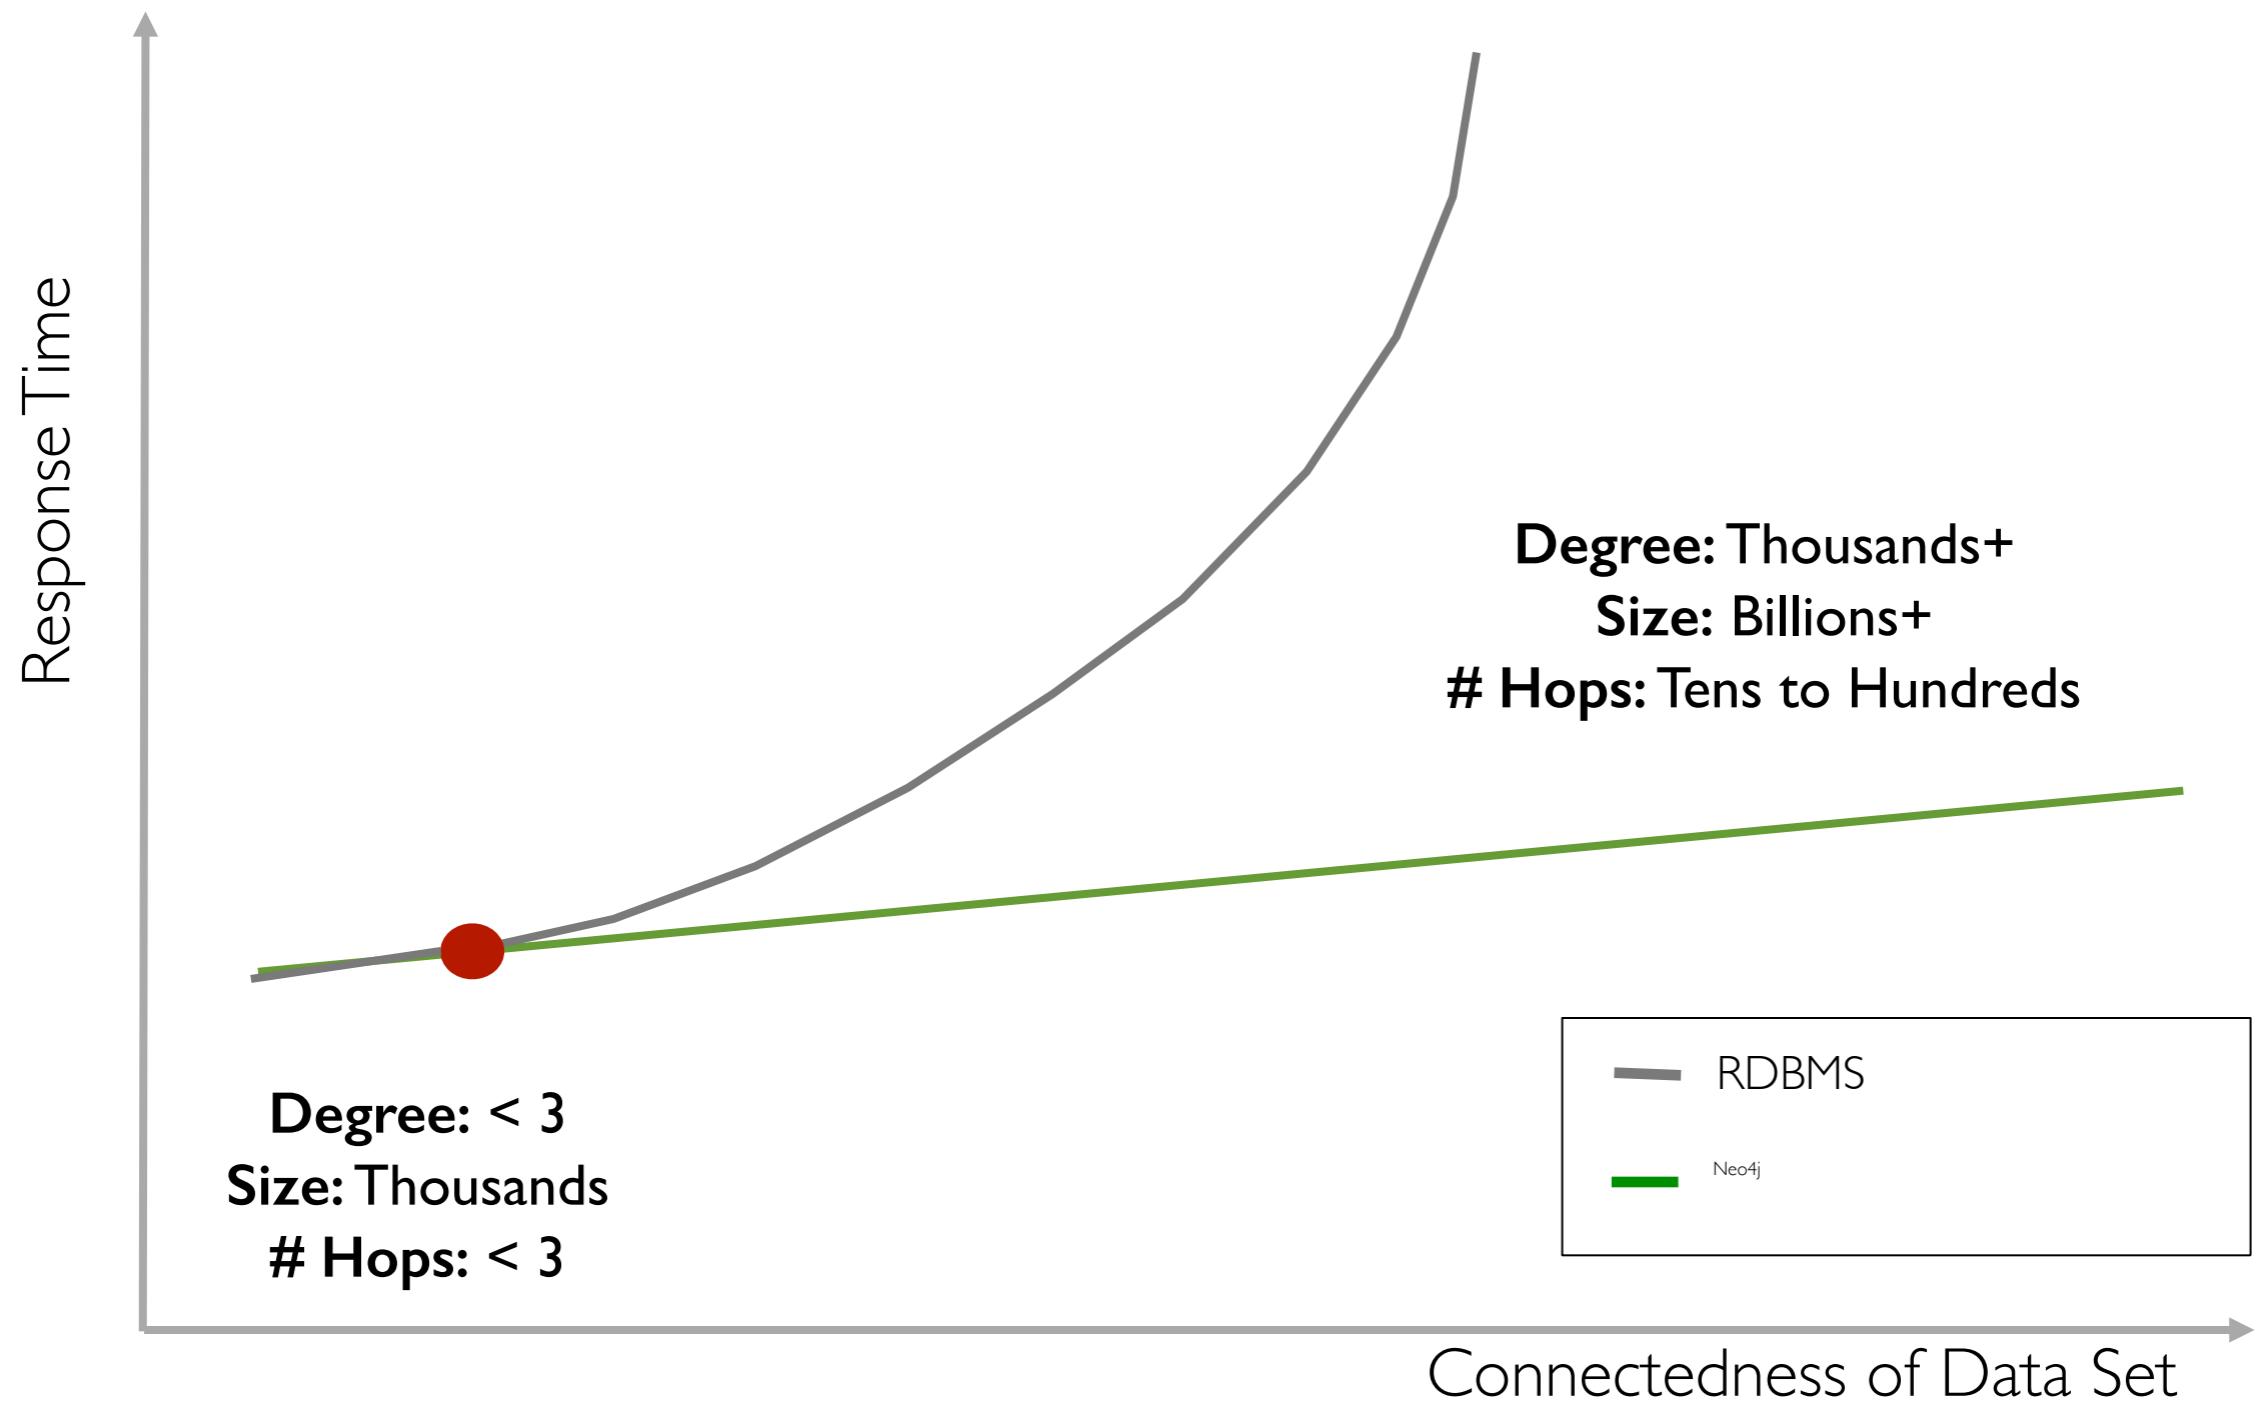
## Neo4j:

>> Performance remains constant as graph size increases

>> Performance slowdown is linear or better as density & degree increase

# Connected Query Performance

## RDBMS vs. Native Graph Database



# Graph db performance

- a sample social graph
  - with ~1,000 persons
- average 50 friends per person
- `pathExists(a,b)` limited to depth 4
- caches warmed up to eliminate disk I/O

Database	# persons	query time
MySQL		
Neo4j		
Neo4j		

# Graph db performance

- a sample social graph
  - with ~1,000 persons
- average 50 friends per person
- `pathExists(a,b)` limited to depth 4
- caches warmed up to eliminate disk I/O

Database	# persons	query time
MySQL	1,000	2,000 ms
Neo4j		
Neo4j		

# Graph db performance

- a sample social graph
  - with ~1,000 persons
- average 50 friends per person
- `pathExists(a,b)` limited to depth 4
- caches warmed up to eliminate disk I/O

Database	# persons	query time
MySQL	1,000	2,000 ms
Neo4j	1,000	2 ms
Neo4j		

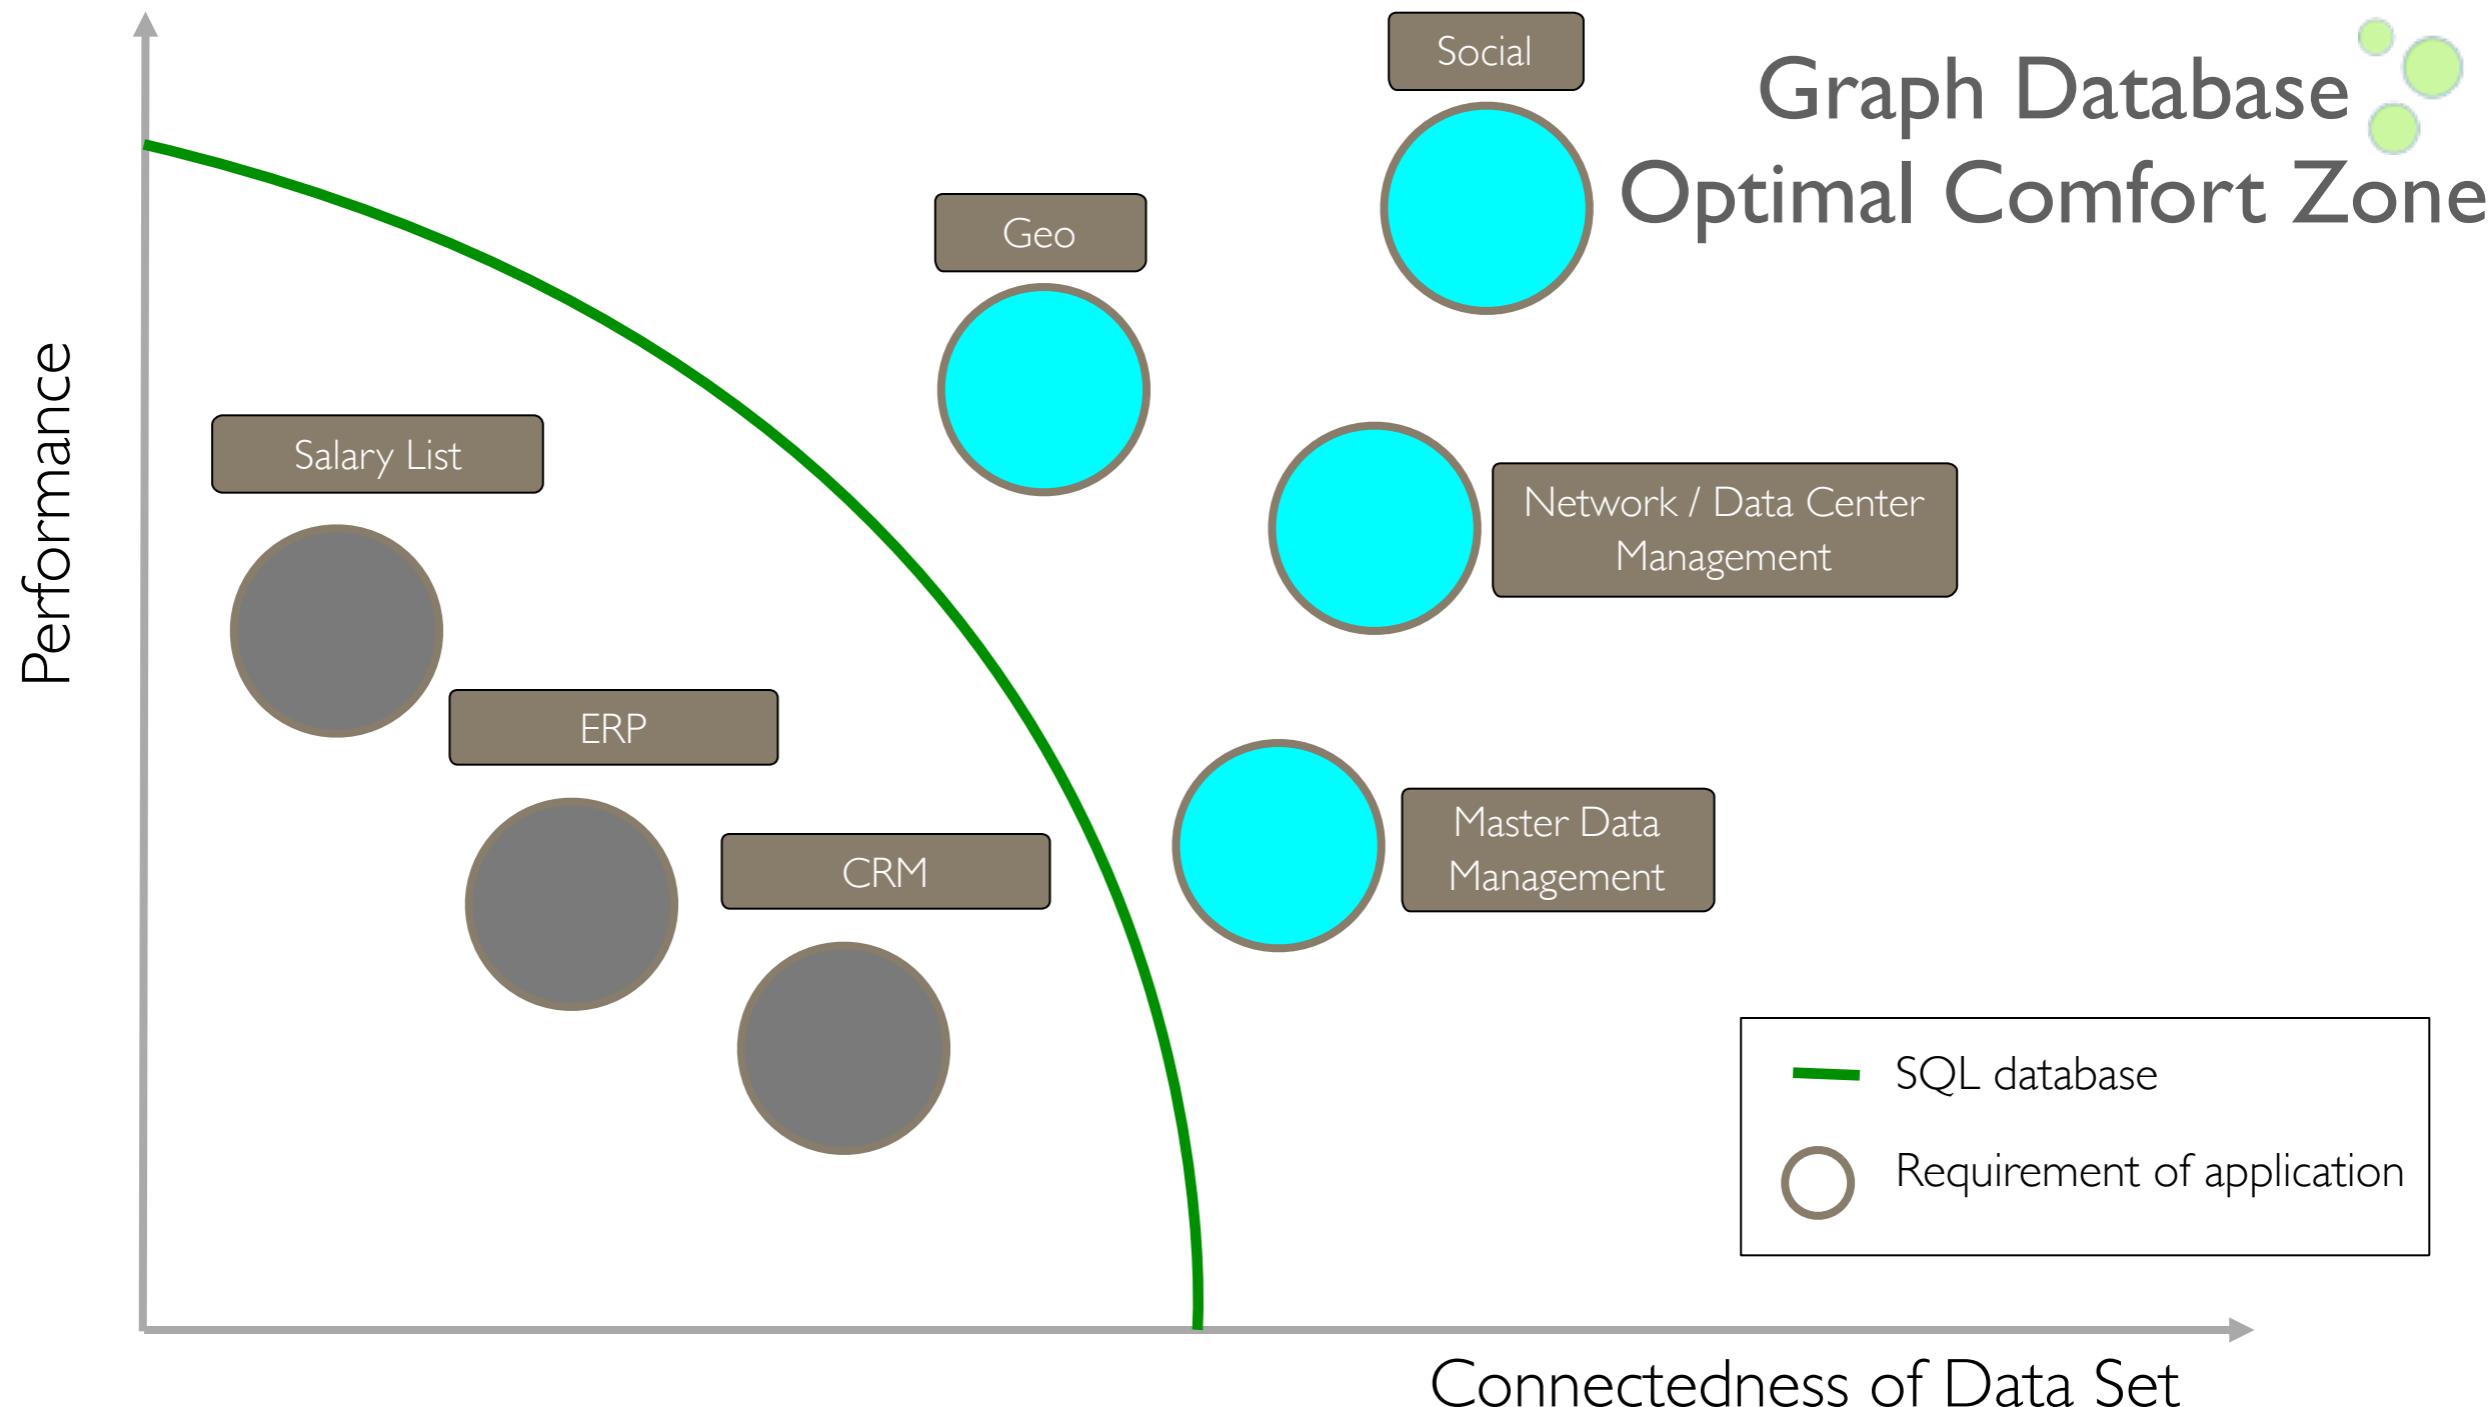
# Graph db performance

- a sample social graph
  - with ~1,000 persons
- average 50 friends per person
- `pathExists(a,b)` limited to depth 4
- caches warmed up to eliminate disk I/O

Database	# persons	query time
MySQL	1,000	2,000 ms
Neo4j	1,000	2 ms
Neo4j	1,000,000	2 ms

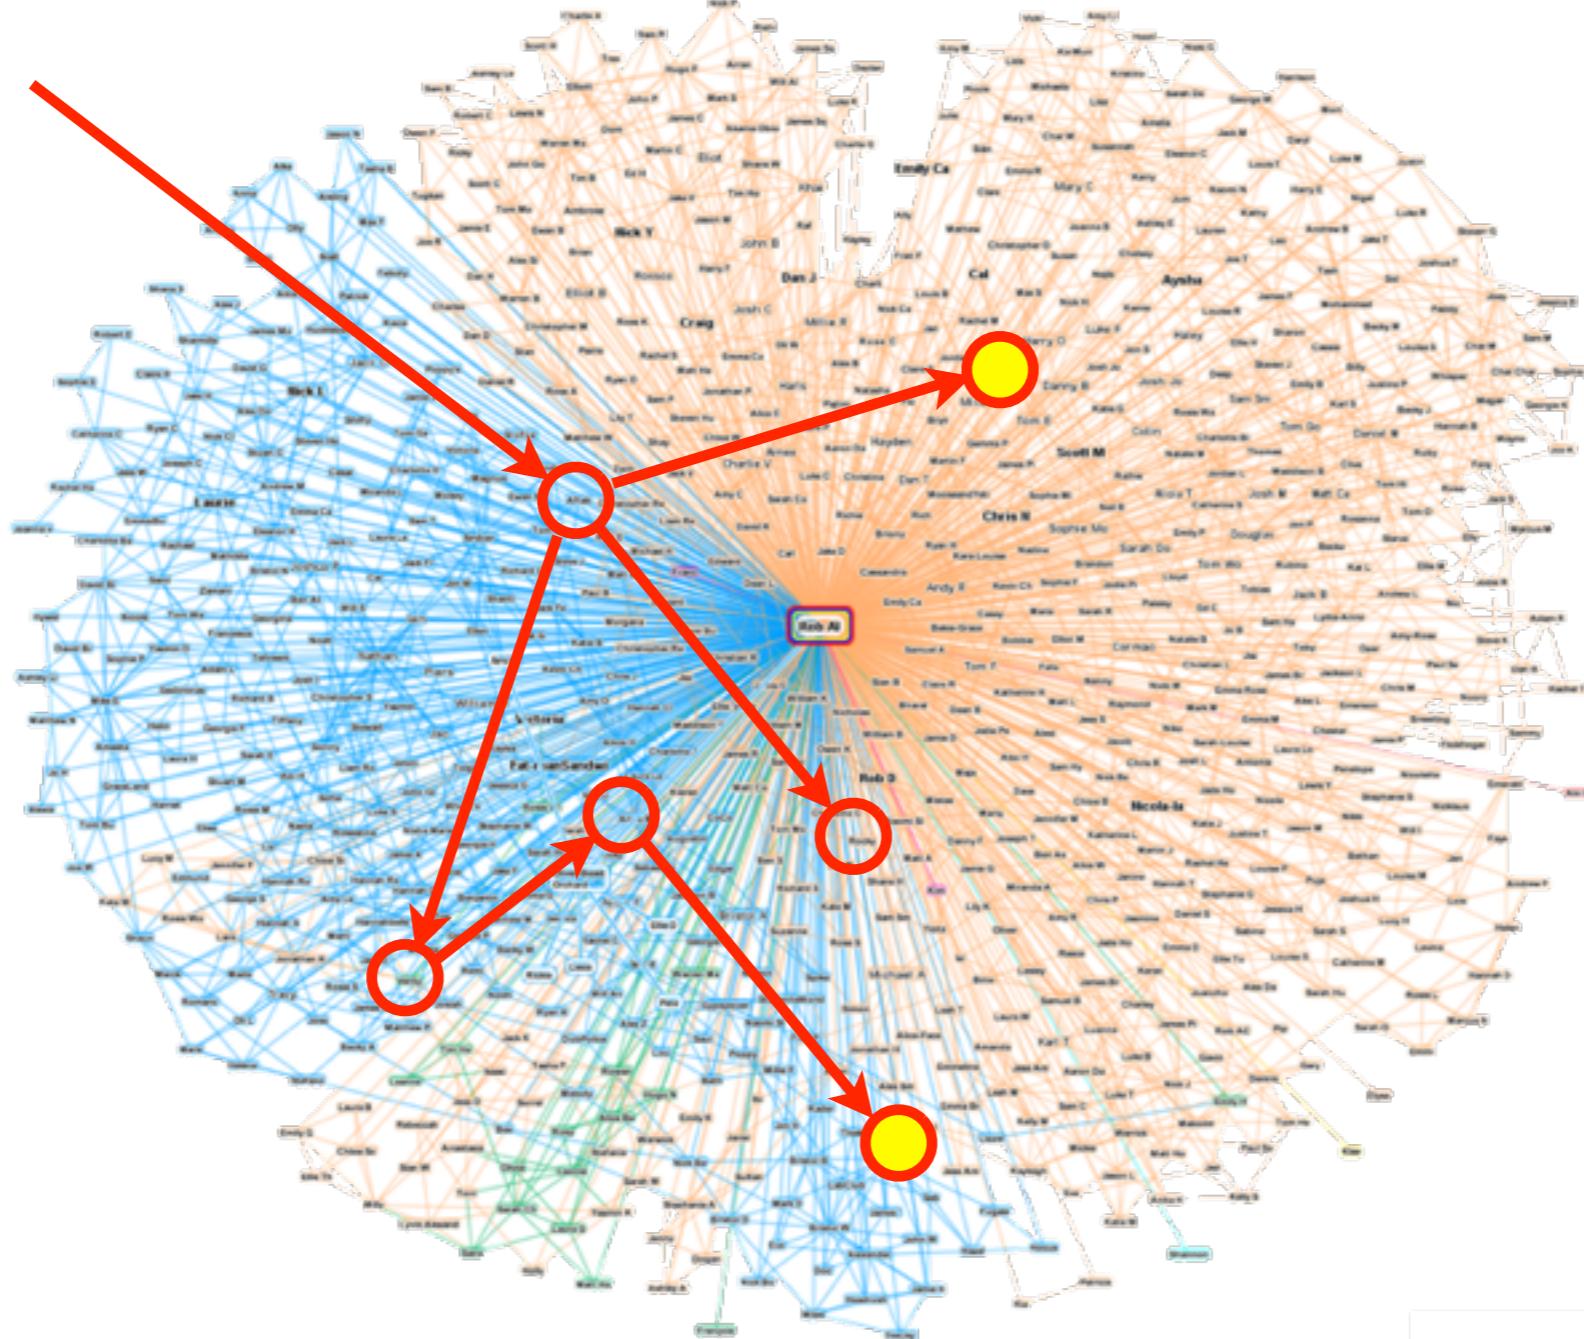
\*Additional Third Party Benchmark Available in *Neo4j in Action*: <http://www.manning.com/partner/>

# The Zone of SQL Adequacy



# Graph Technology Ecosystem

# #1: Graph Local Queries

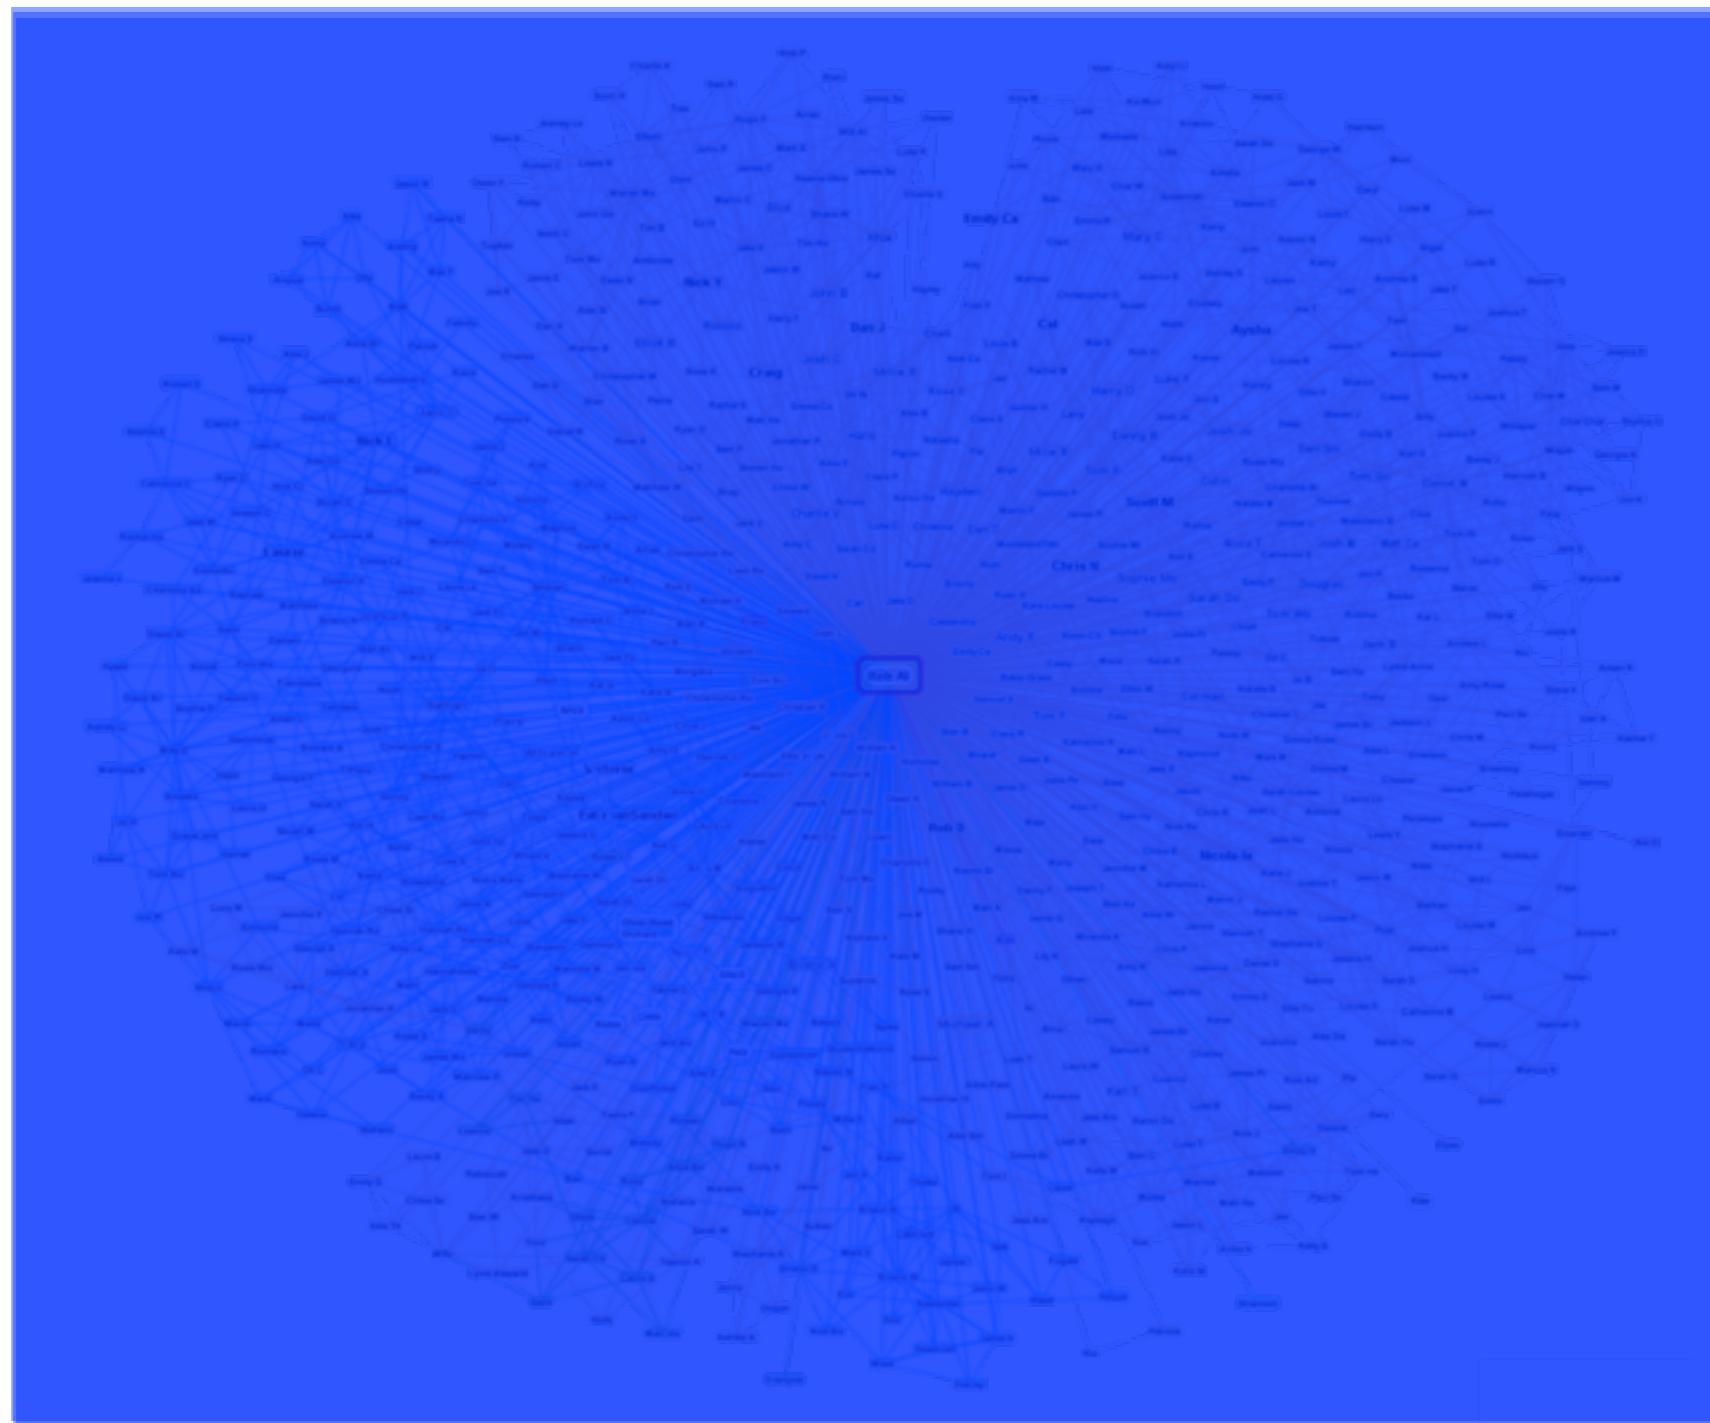


e.g. Recommendations, Friend-of-Friend, Shortest Path

# #2: Graph Global Queries



How many restaurants, on average, has each person liked?



# What is a Graph Database

“A **graph database**... is an online database management system with CRUD methods that expose a graph data model”<sup>1</sup>

- Two important properties:
  - **Native graph storage engine:** *written from the ground up to manage graph data*
  - **Native graph processing,** including *index-free adjacency* to facilitate traversals

# Graph Databases are Designed to:

1. Store inter-connected data
2. Make it easy to make sense of that data
3. Enable extreme-performance operations for:
  - Discovery of connected data patterns
  - Relatedness queries > depth 1
  - Relatedness queries of arbitrary length
4. Make it easy to evolve the database

# Top Reasons People Use Graph Databases

- 1. Problems with Join performance.**
- 2. Continuously evolving data set (often involves wide and sparse tables)**
- 3. The Shape of the Domain is naturally a graph**
- 4. Open-ended business requirements necessitating fast, iterative development.**

# Graph Compute Engine

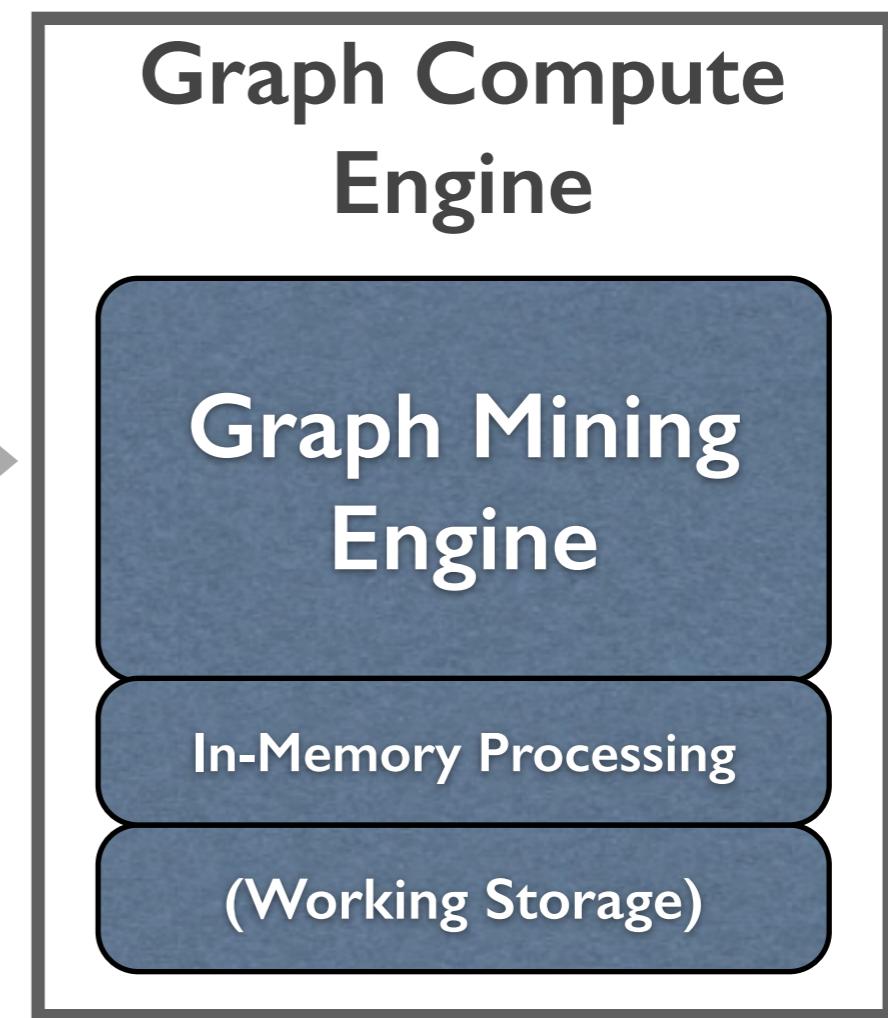
Processing engine that enables graph global computational algorithms to be run against large data sets



System(s)  
of Record



Data extraction,  
transformation,  
and load

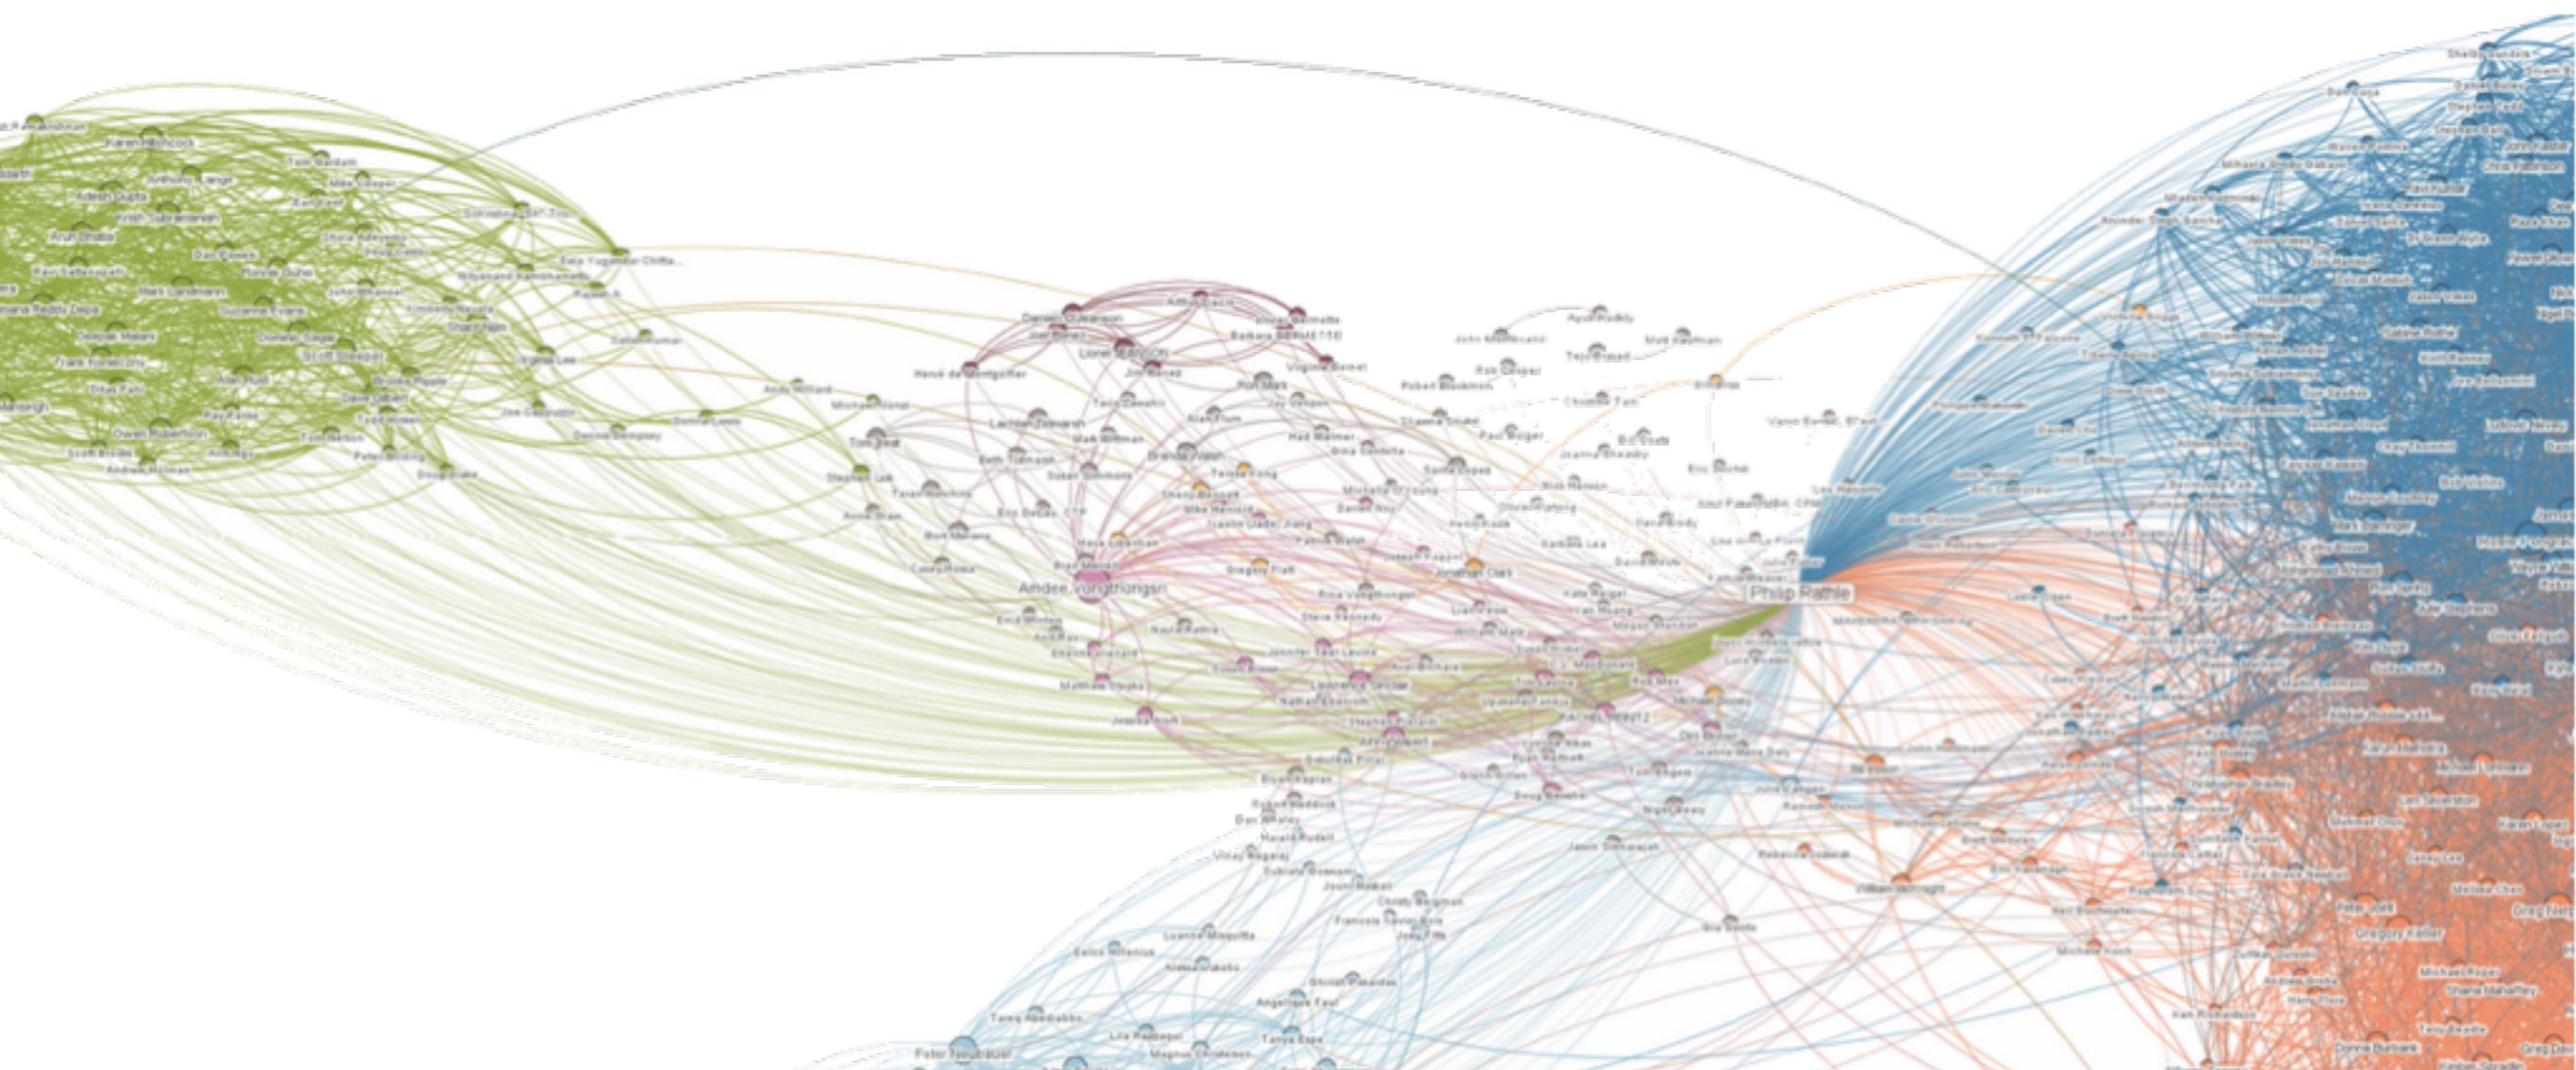


# Connected Data



**Real-Time/  
OLTP**

**Offline/  
Batch**



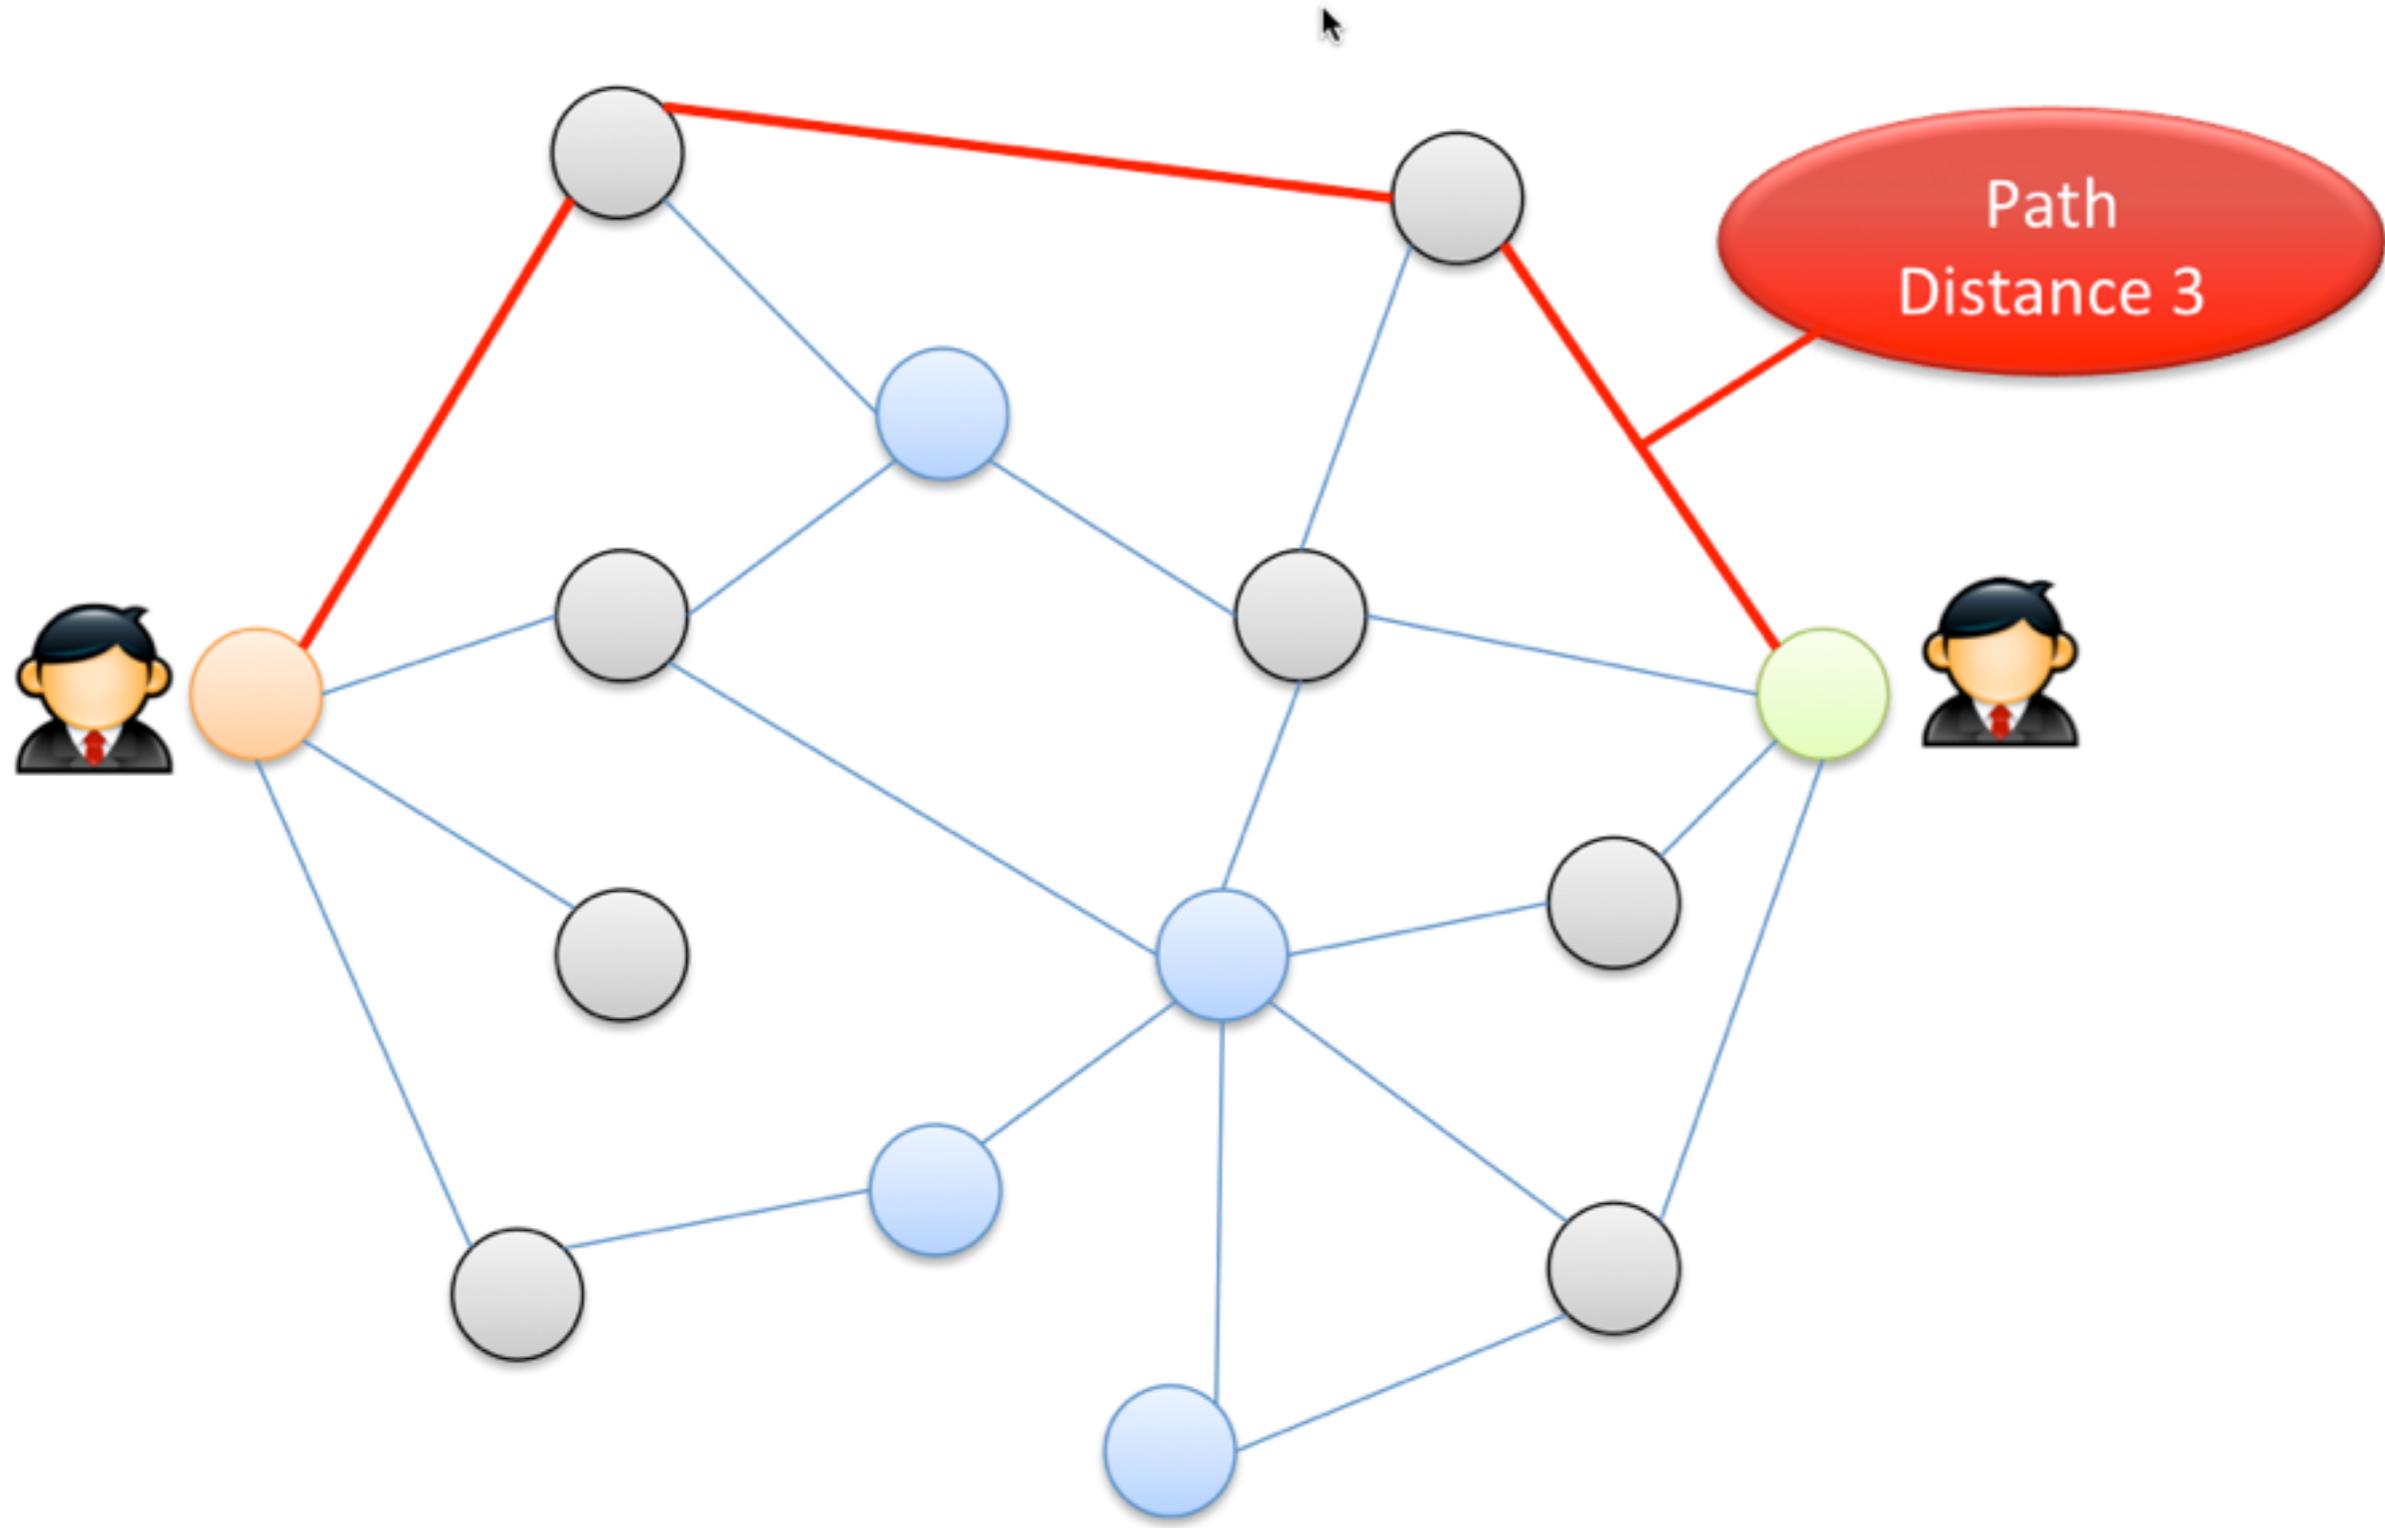
The Viadeo Group forms the world's second largest professional social network with over 50 million members.

- N°1 in France with over 7 million members,
- N°1 in China with over 14 millions members.
- Strong presence in Europe and in emerging markets, strong growth in South America
- Presence in Russia since December 2011 through a joint venture with leading media group : Sanoma Independent Media
- April 2012 : record funding of 24 million euros (FSI, historical shareholders and new entrants)
- Headquarters in Paris and local offices on 5 continents
- 400 employees



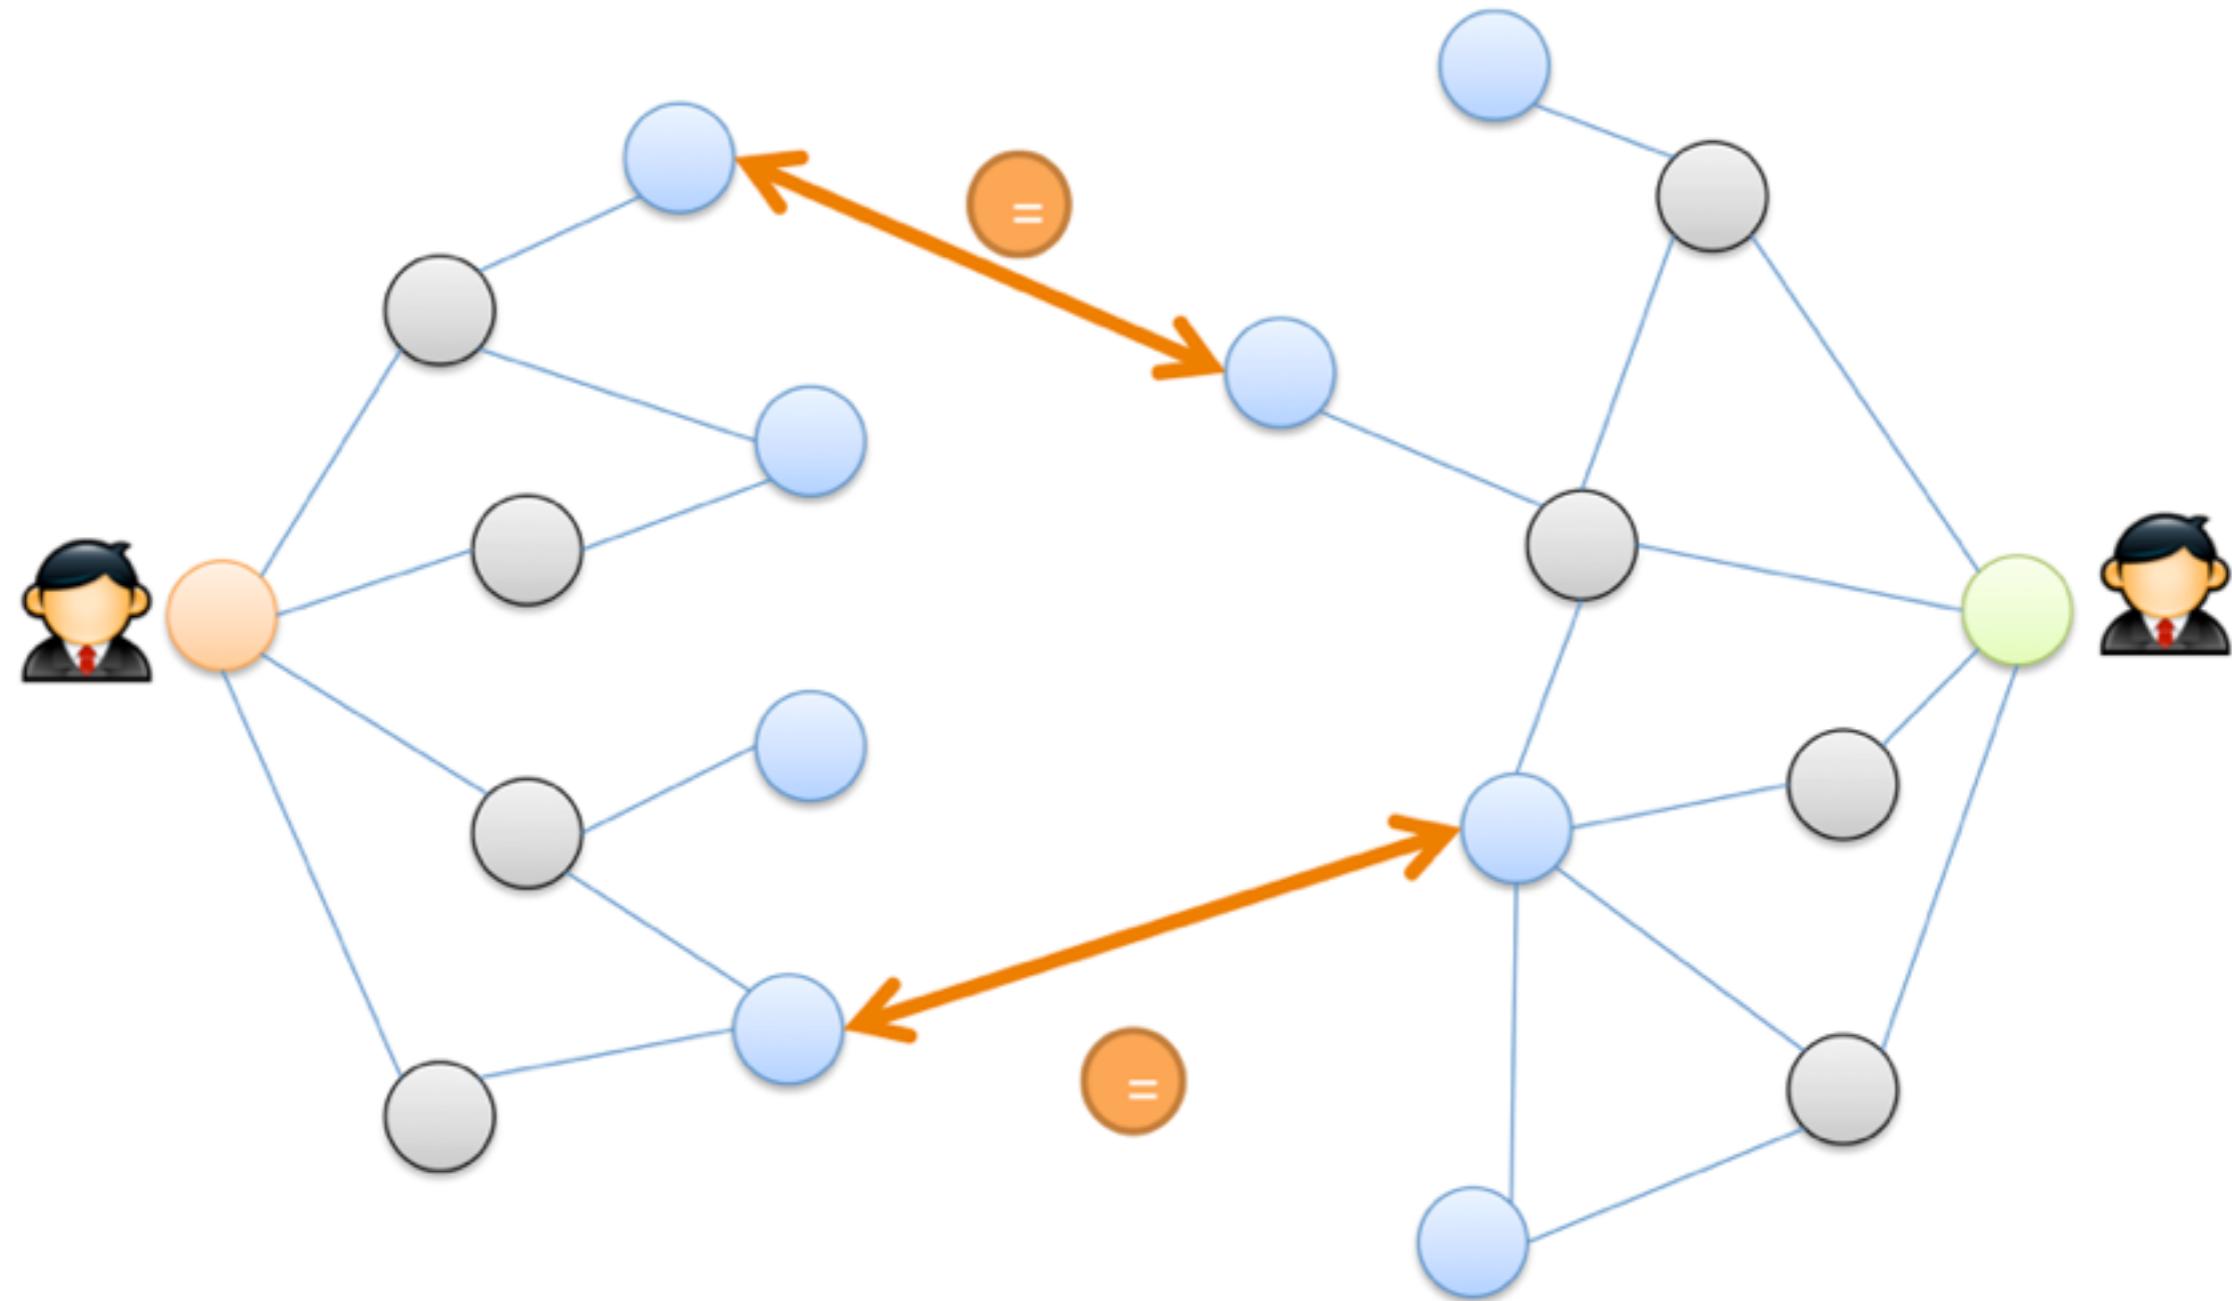
- The Viadeo Group consists of:

[ApnaCircle.com](http://ApnaCircle.com)[天際網  
www.tianji.com](http://www.tianji.com)



- In-house algorithm
- Network storage in **MySQL Database**

```
CREATE TABLE `Network` (
    `memberId` int(11) NOT NULL DEFAULT '0',
    `L1` mediumblob NOT NULL,
    `L2` mediumblob NOT NULL,
    PRIMARY KEY (`memberId`)
) ENGINE=InnoDB;
```



## LIMITATIONS

- 1) Important latency for complete update



- 2) Massive bandwidth impact for internal network

- 3) 48 hours to restart from scratch



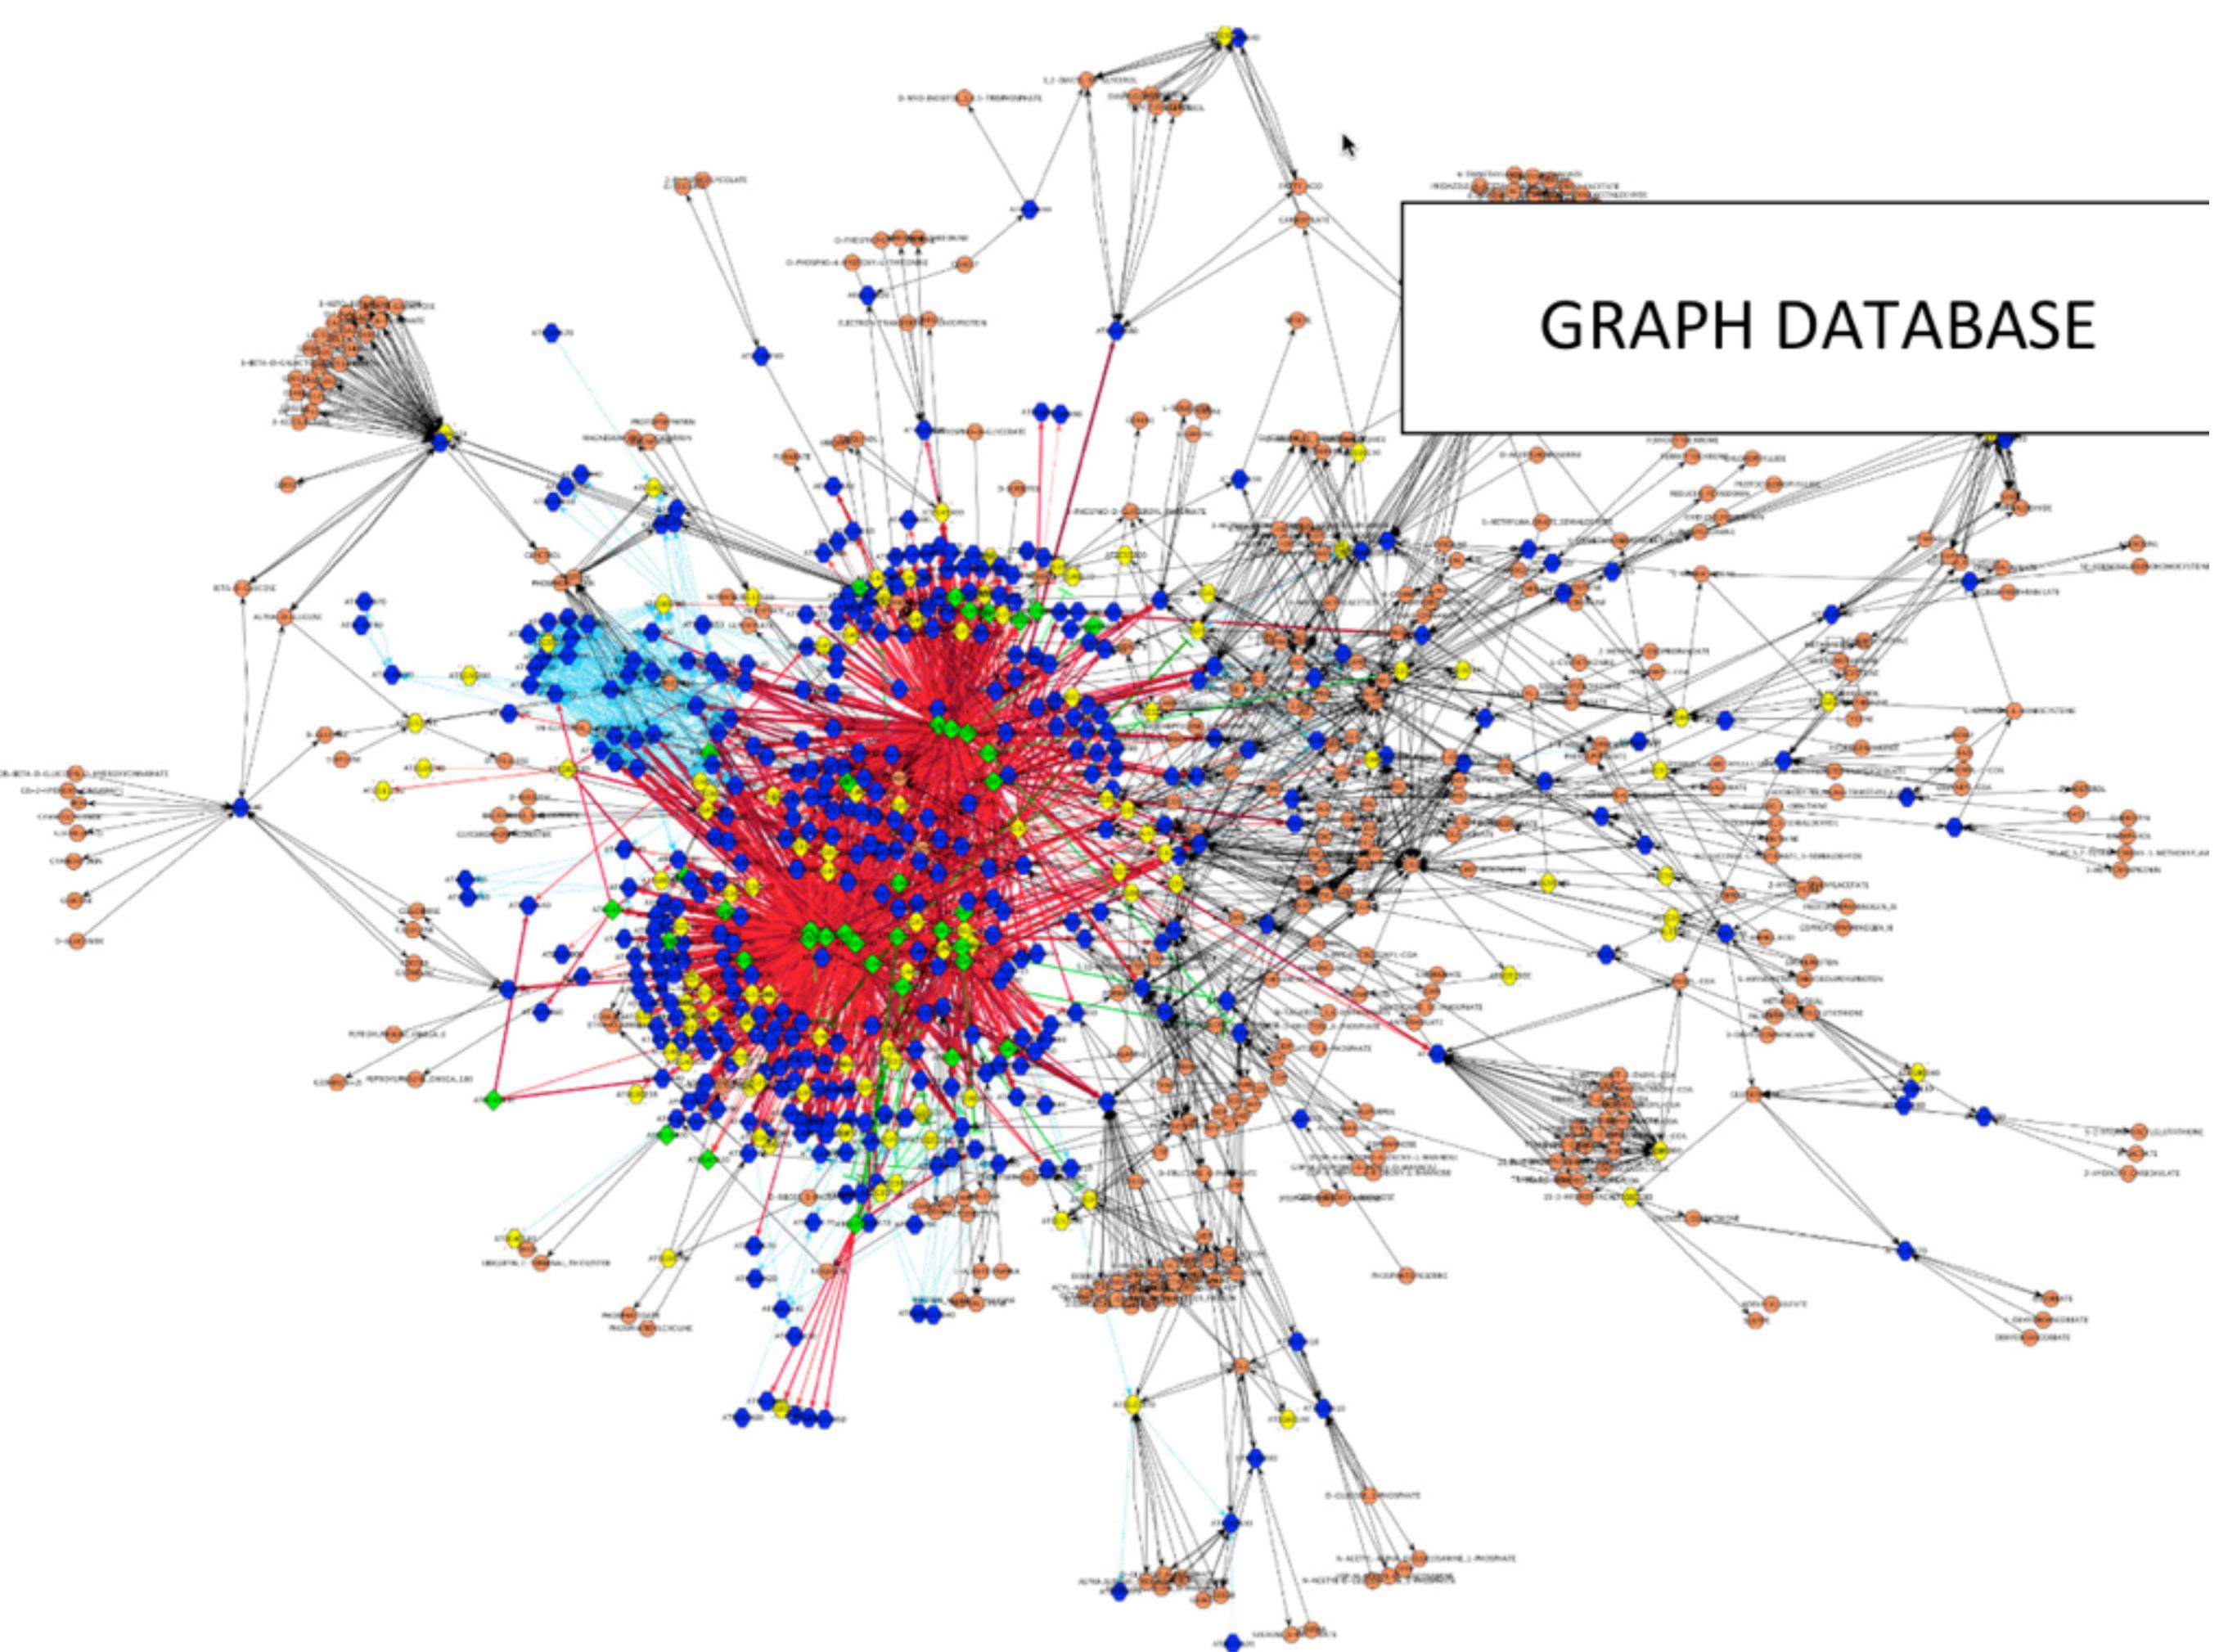
## Update the network (old-fashioned style)

Member A and Member B are now in contacts

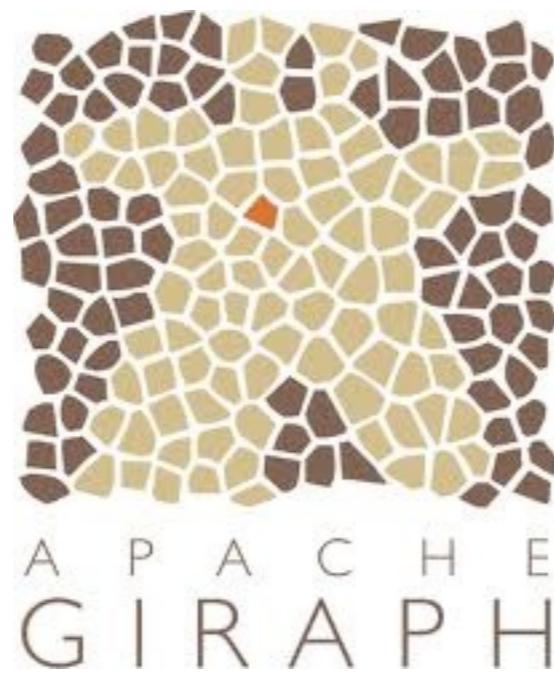
- Update of A.L1 + B.L1 and A.L2 + B.L2
- Retrieving A.L1 + B.L1 and update \*.L2

Example:

- A has 500 contacts
  - B has 150 contacts
- ➔  $500 + 150 + 2 = \textbf{652 updates!}$



## GRAPH DATABASE



Wait what?



New Users?  
Real Time Updates?

## BENEFITS



Very easy to integrate  
(less than 2 months)

Instantaneous  
graph updates



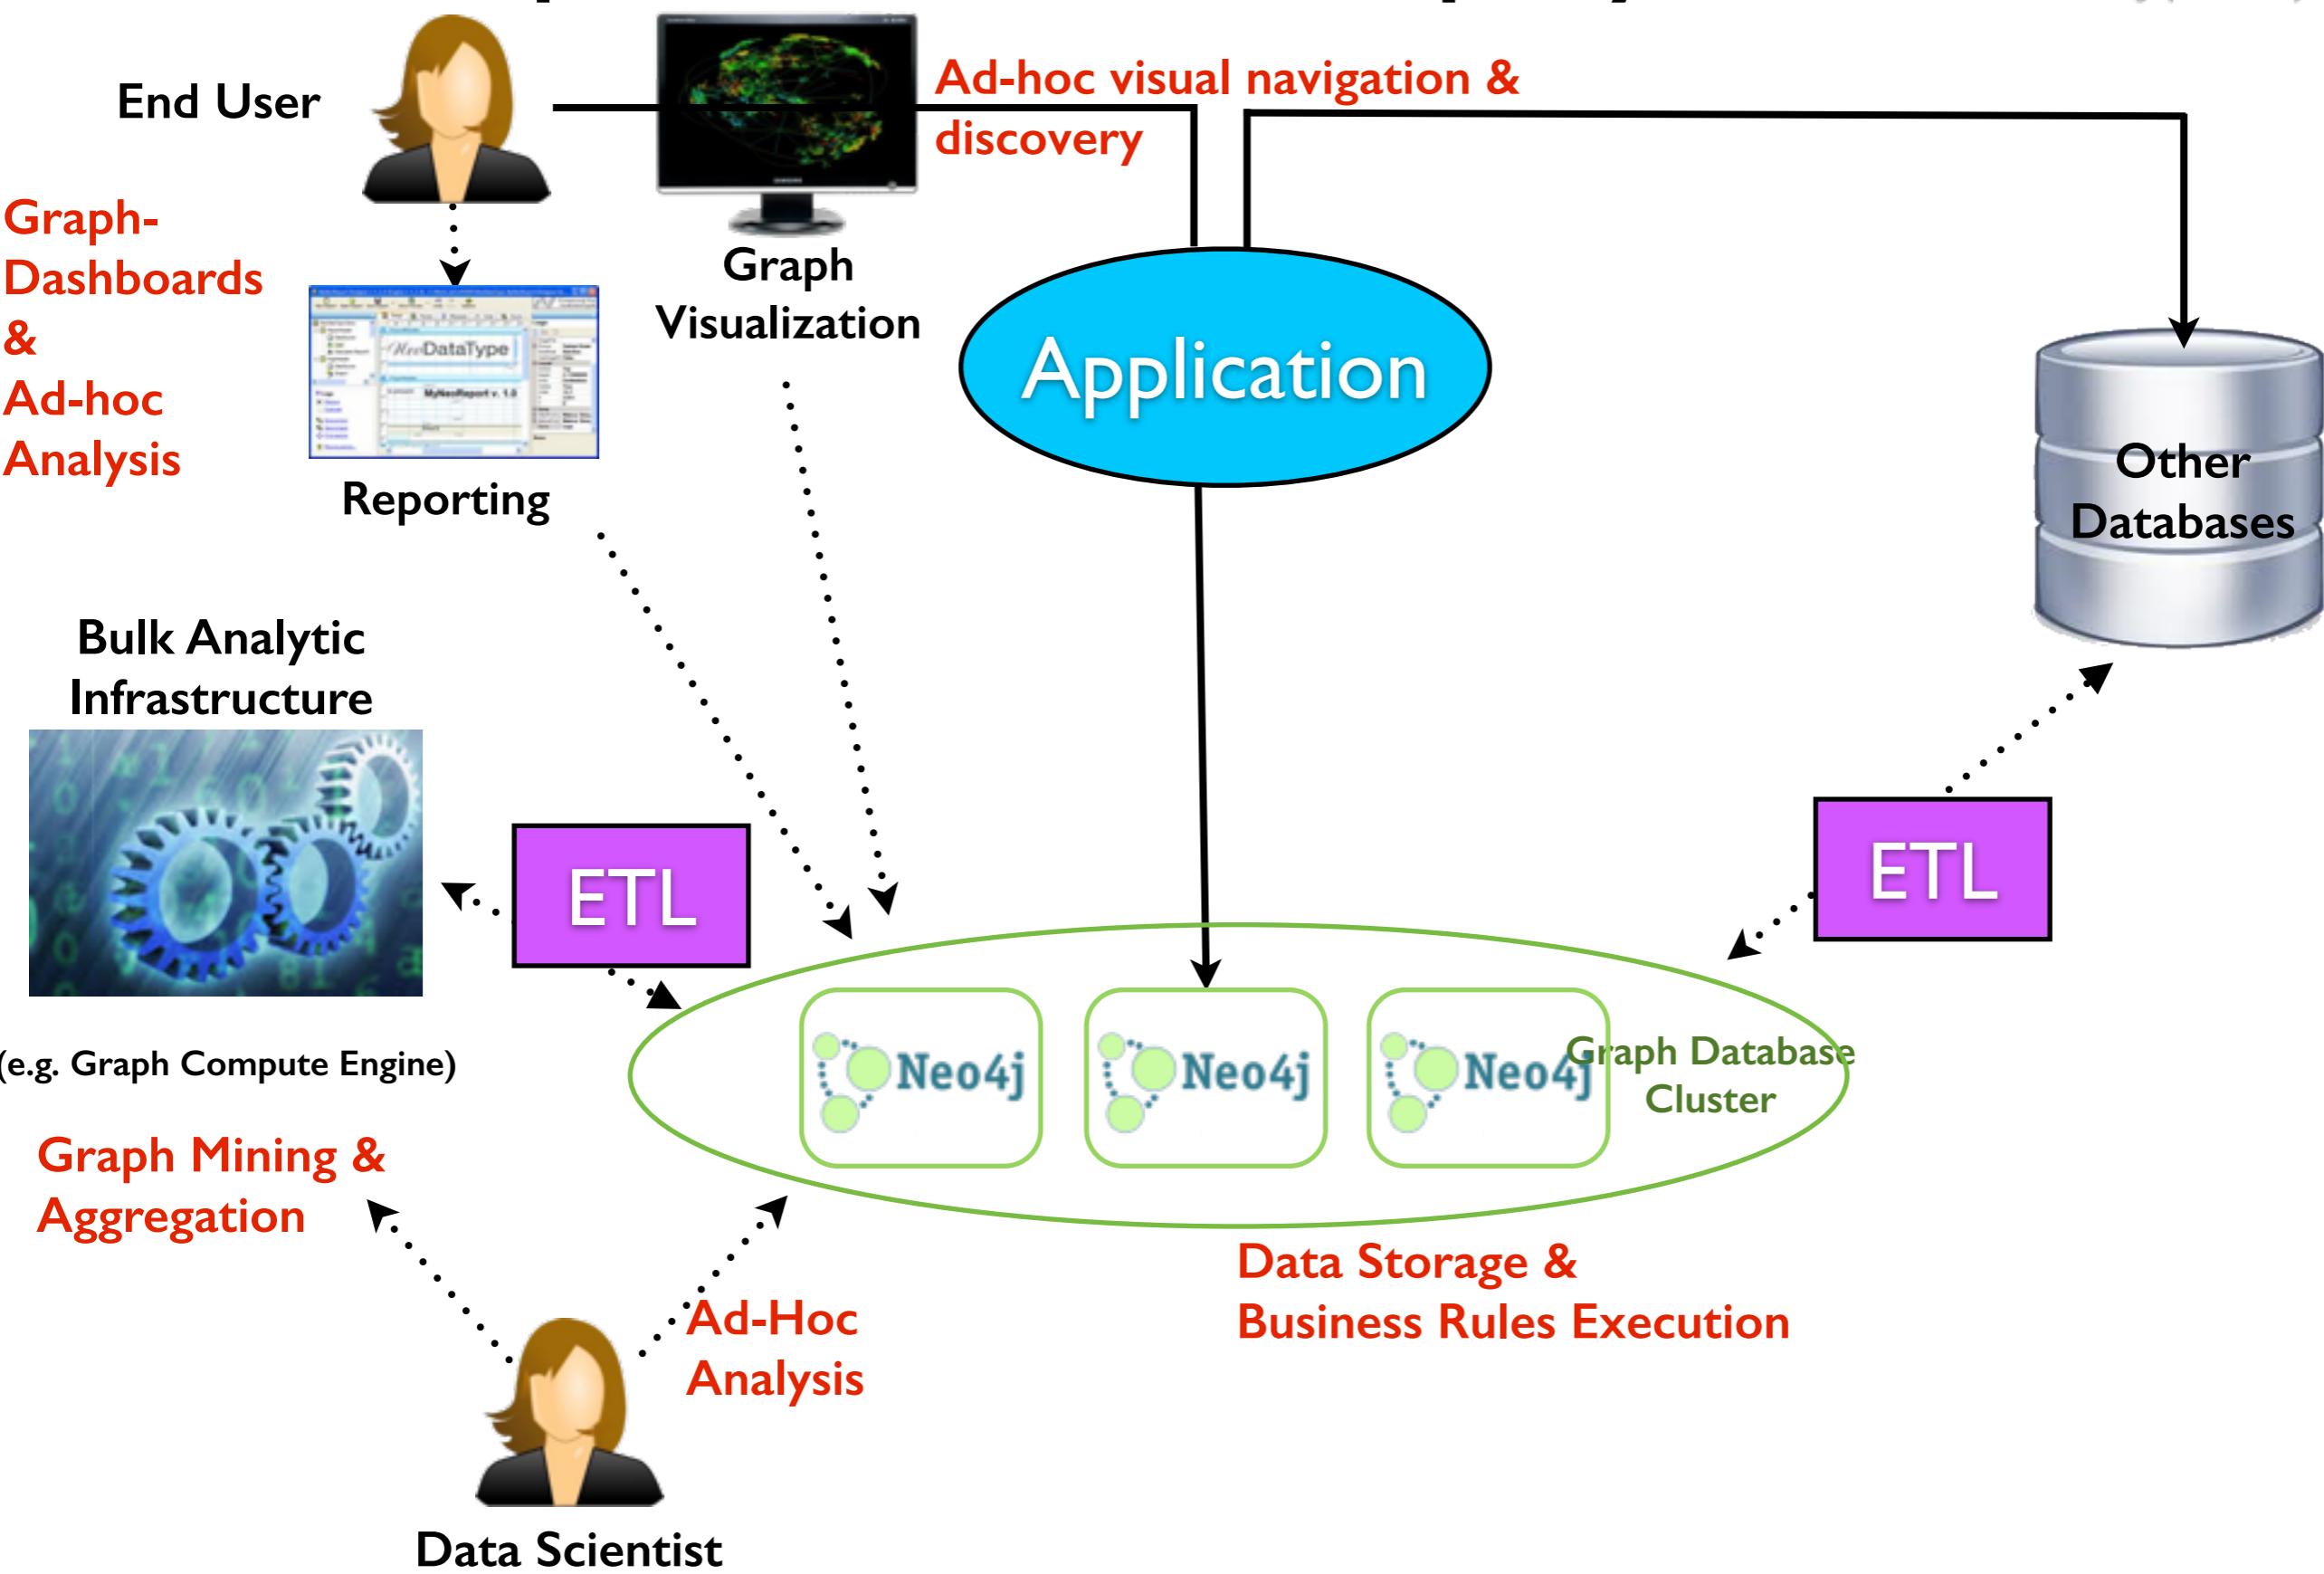
Backup /  
Restore



High Availability



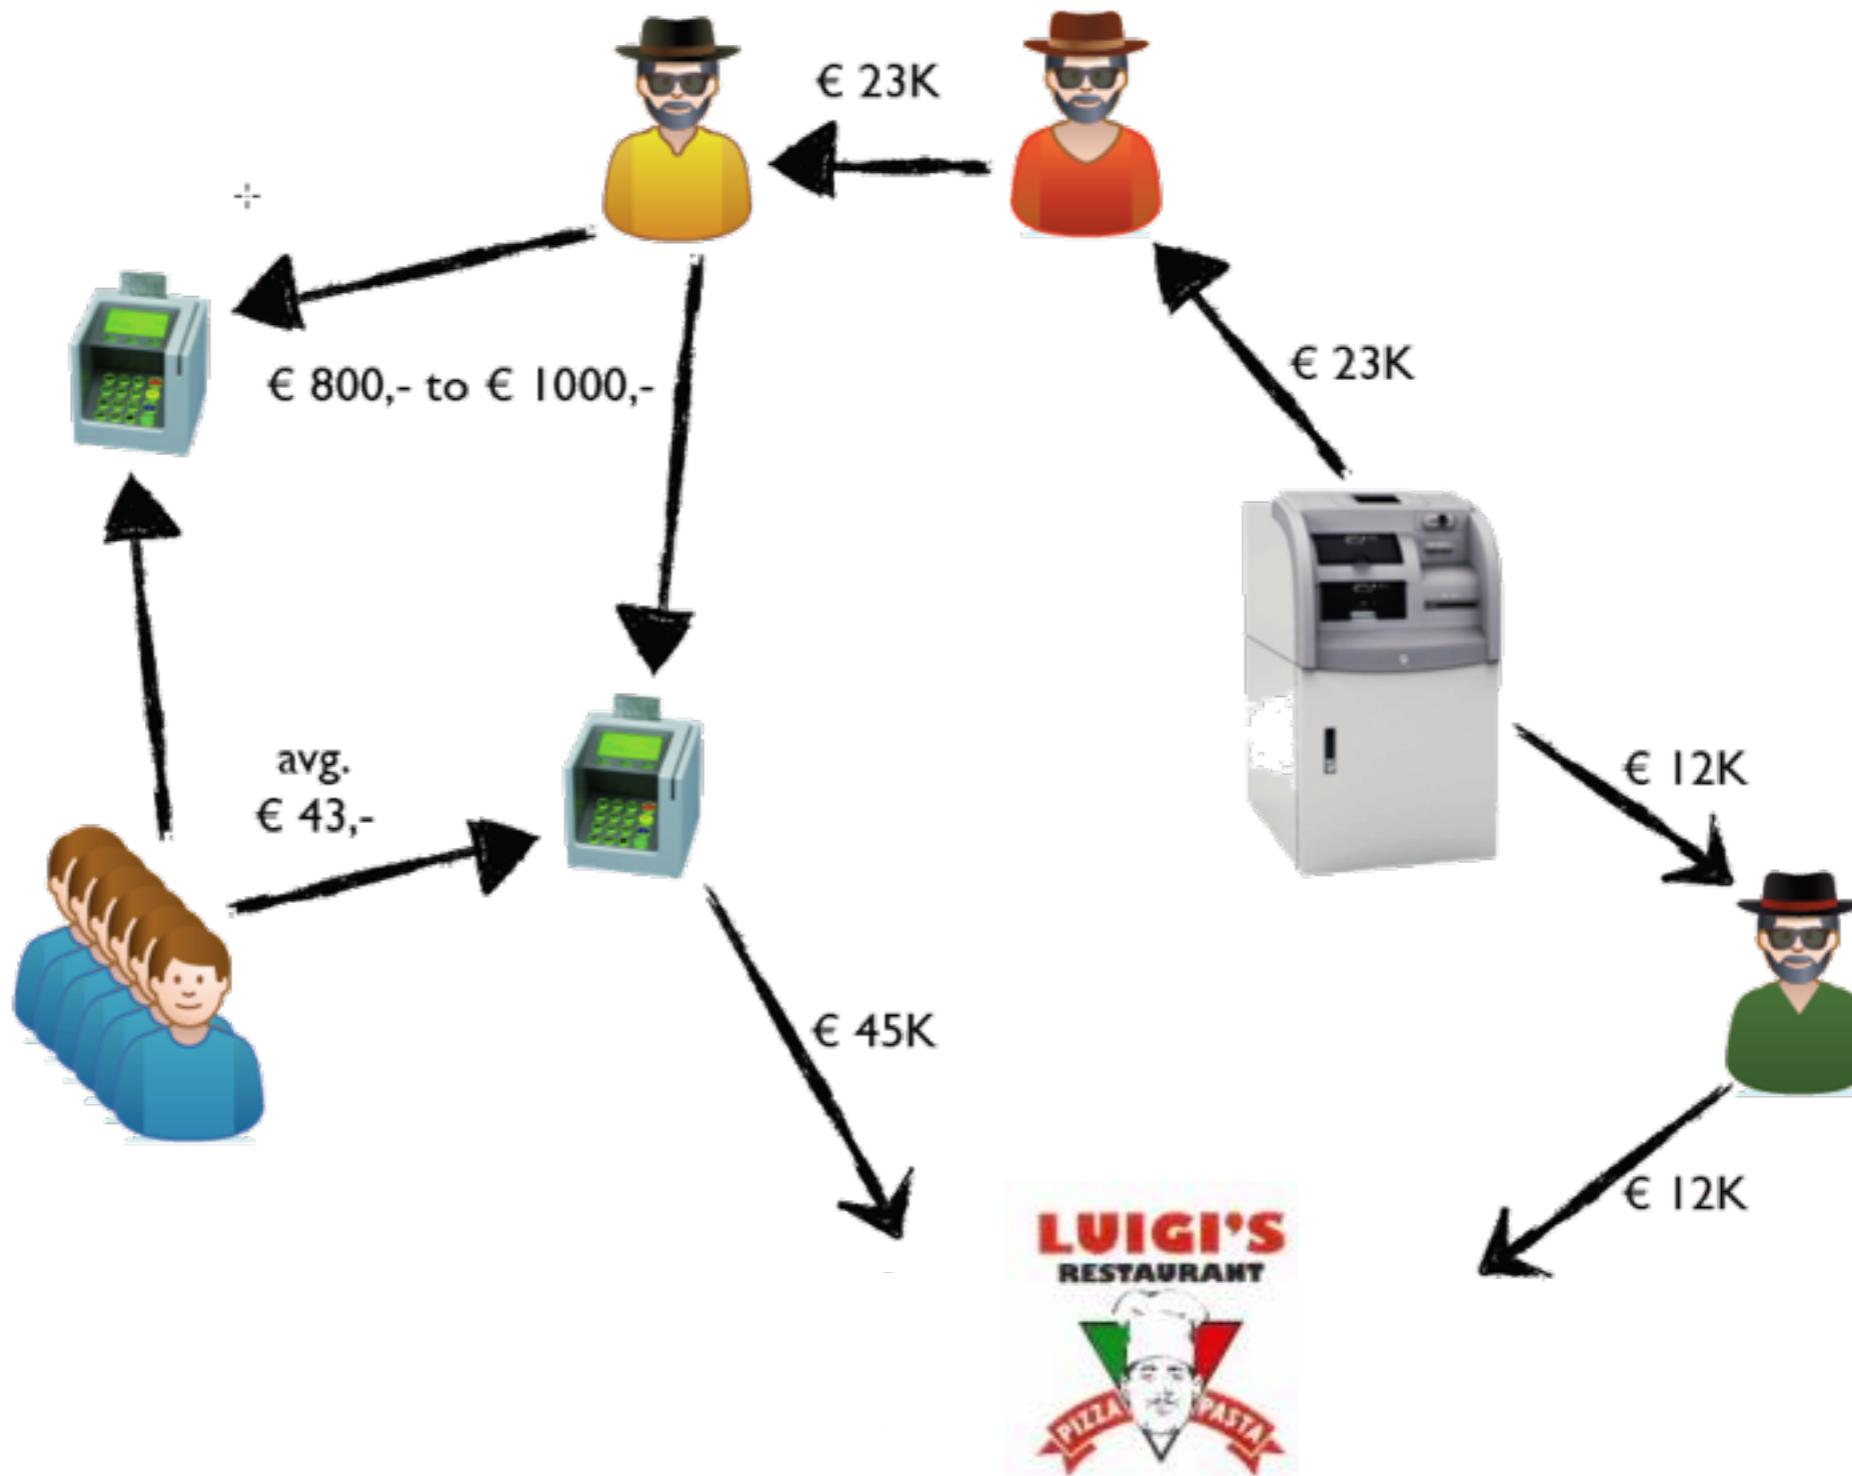
# Graph Database Deployment



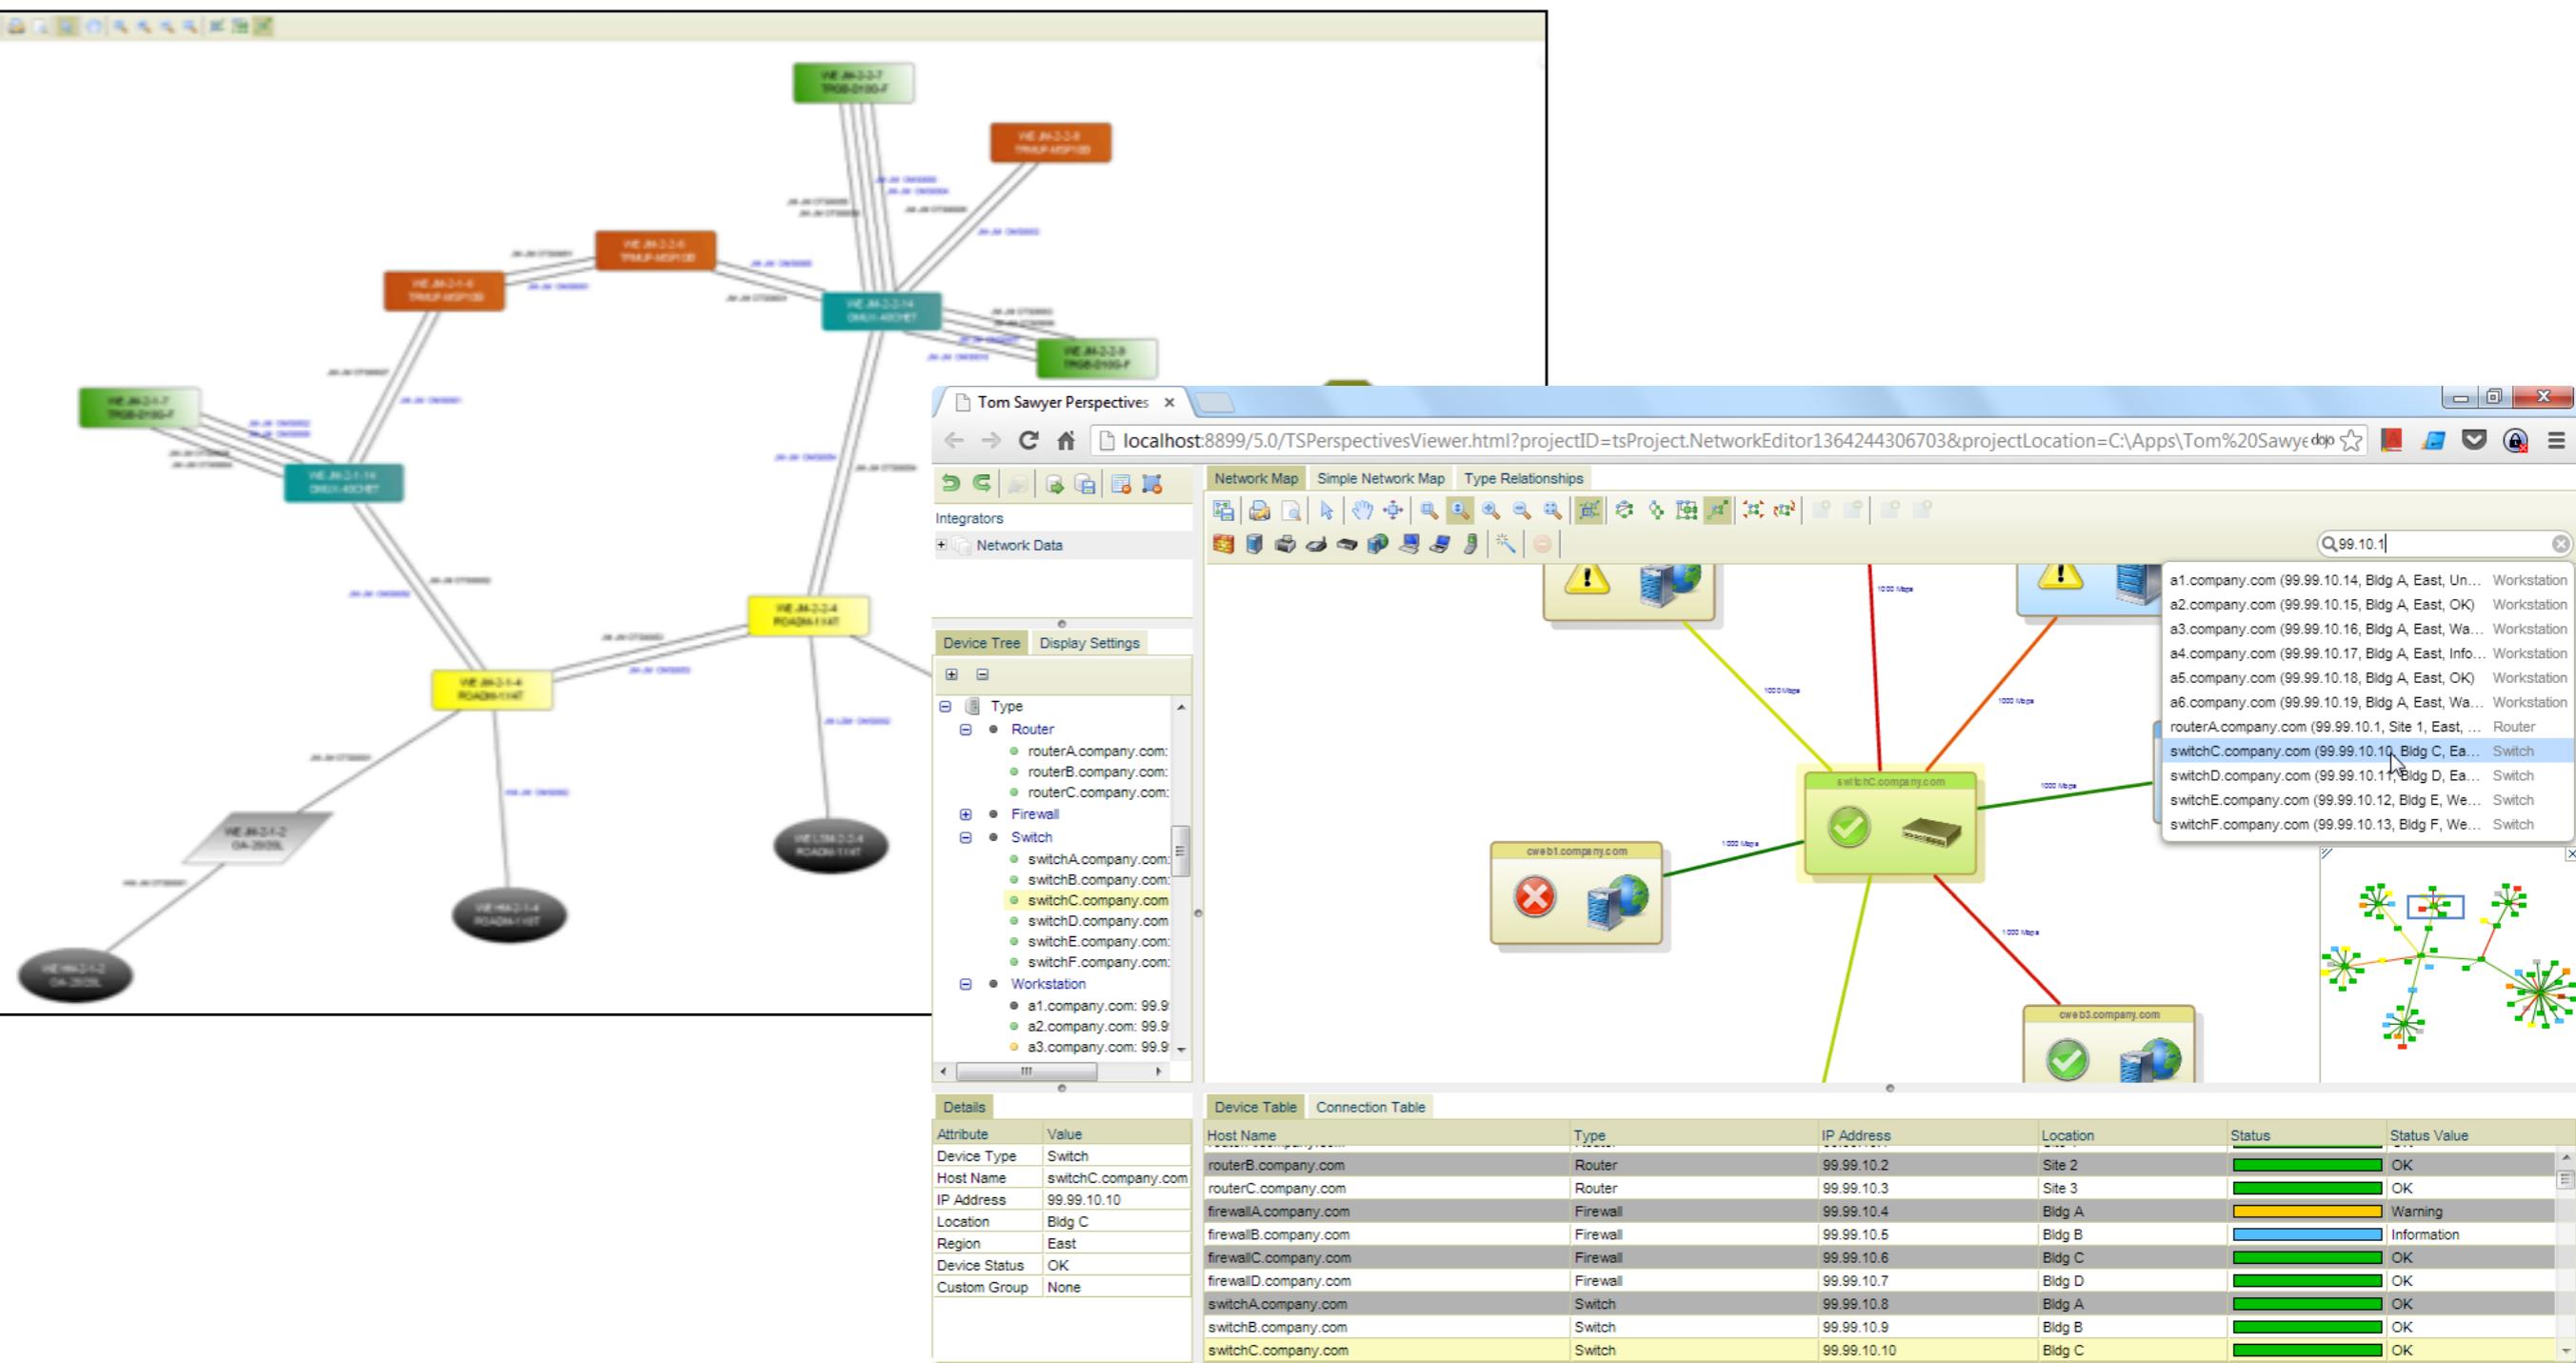
# Graph Dashboards

## The Power of Visualization

# Fraud Detection & Money Laundering

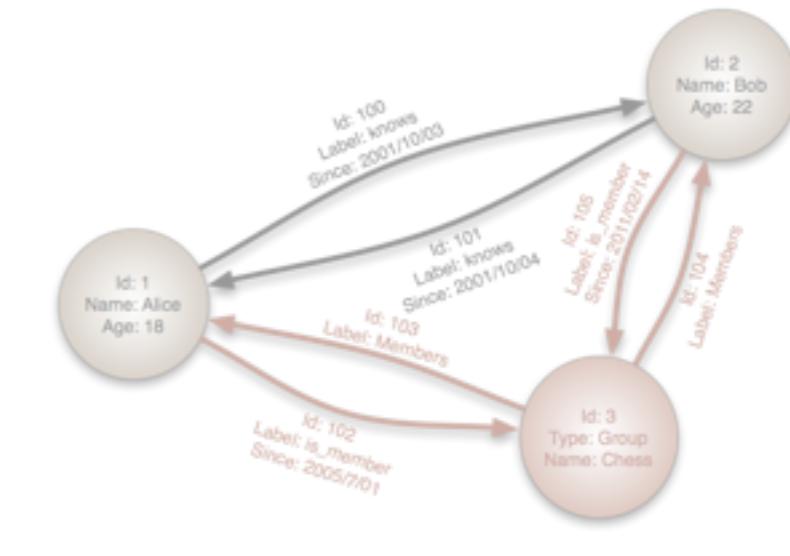


# IT Service Dependencies



# Working with Graphs

## Case Studies & Working Examples



# Cypher

## ASCII art Graph Patterns



```
MATCH (A) -[:LOVES]-> (B)  
WHERE A.name = "A"  
RETURN B as lover
```

# Social Example

# Practical Cypher

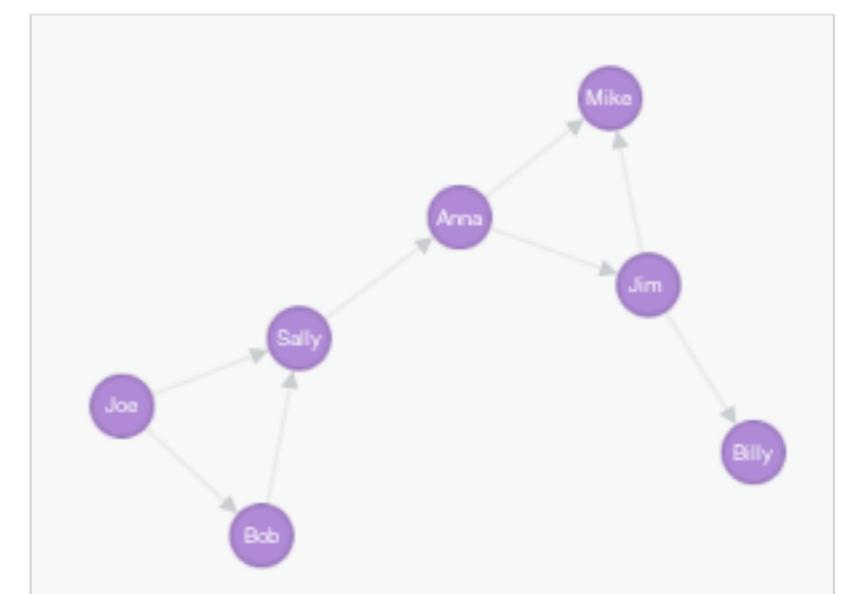
## Social Graph - Create



### CREATE

```
(joe:Person {name:"Joe"}) ,  
(bob:Person {name:"Bob"}) ,  
(sally:Person {name:"Sally"}) ,  
(anna:Person {name:"Anna"}) ,  
(jim:Person {name:"Jim"}) ,  
(mike:Person {name:"Mike"}) ,  
(billy:Person {name:"Billy"}) ,
```

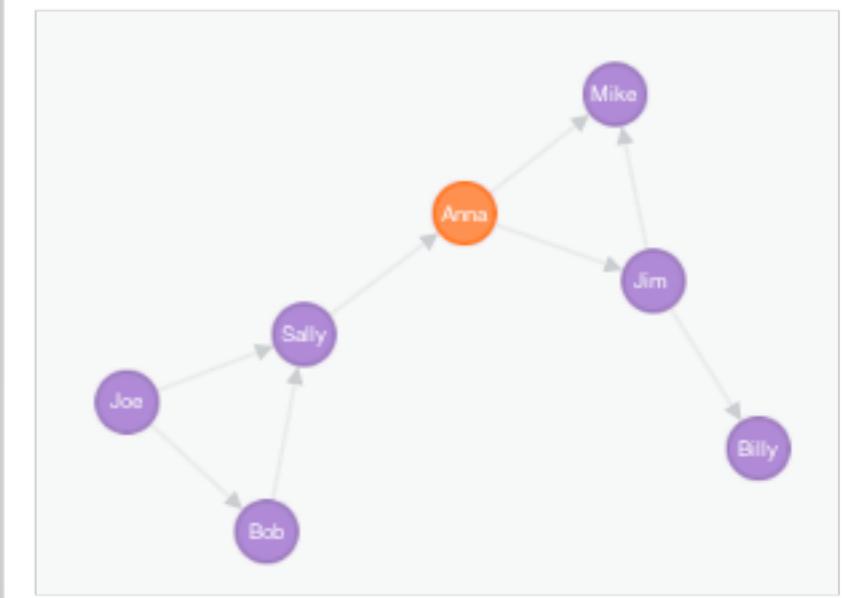
```
(joe)-[:KNOWS]->(bob) ,  
(joe)-[:KNOWS]->(sally) ,  
(bob)-[:KNOWS]->(sally) ,  
(sally)-[:KNOWS]->(anna) ,  
(anna)-[:KNOWS]->(jim) ,  
(anna)-[:KNOWS]->(mike) ,  
(jim)-[:KNOWS]->(mike) ,  
(jim)-[:KNOWS]->(billy)
```



# Practical Cypher

## Social Graph - Friends of Joe's Friends

```
MATCH (person)-[:KNOWS]-(friend),  
      (friend)-[:KNOWS]-(foaf)  
WHERE person.name = "Joe"  
  AND NOT(person-[:KNOWS]-foaf)  
RETURN foaf
```



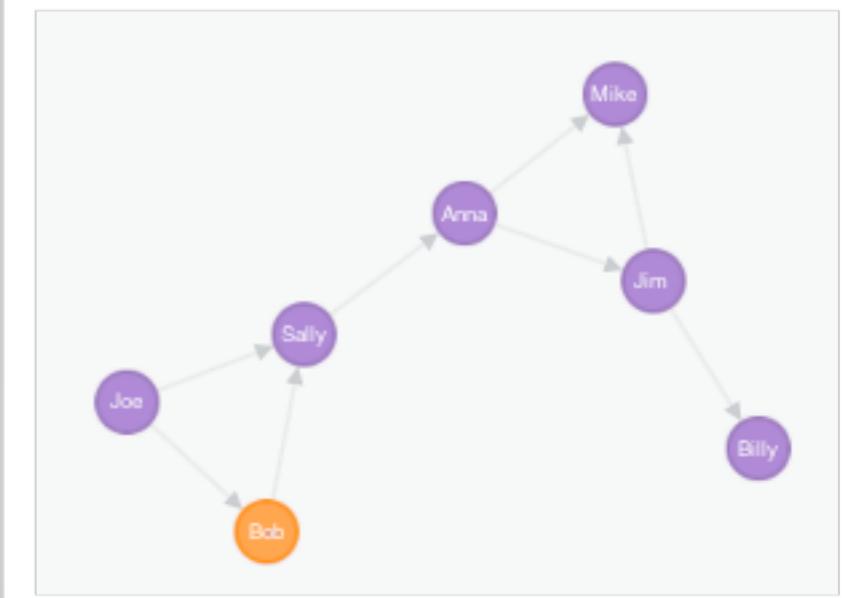
foaf

{name:"Anna"}

# Practical Cypher

## Social Graph - Common Friends

```
MATCH (person1)-[:KNOWS]-(friend),  
      (person2)-[:KNOWS]-(friend)  
WHERE person1.name = "Joe"  
      AND person2.name = "Sally"  
RETURN friend
```



friend

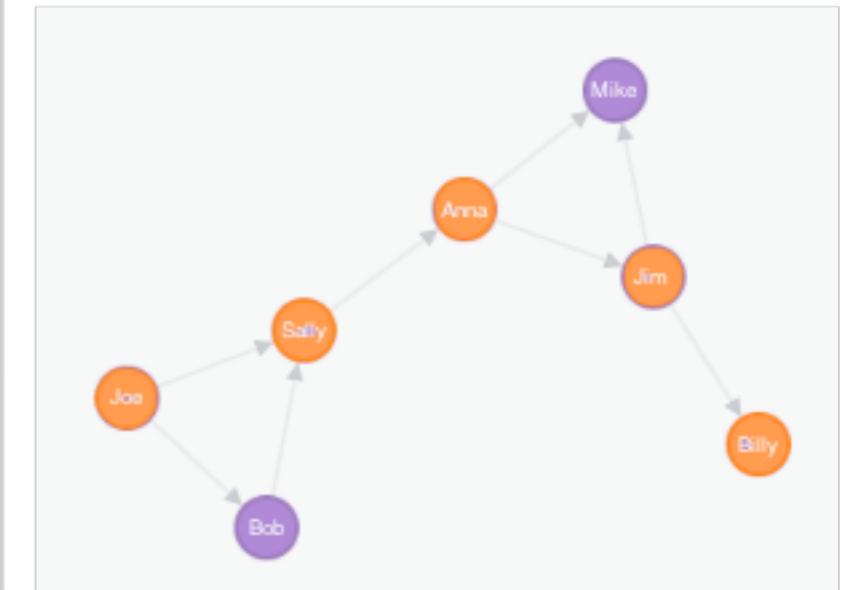
```
{name:"Bob"}
```

# Practical Cypher

## Social Graph - Shortest Path

```

MATCH path = shortestPath(
  (person1)-[:KNOWS*..6]-(person2)
)
WHERE person1.name = "Joe"
  AND person2.name = "Billy"
RETURN path
    
```



**path**

```

{start:"13759",
nodes: [ "13759", "13757", "13756", "13755", "13753" ],
length:4,
relationships:[ "101407", "101409", "101410", "101413" ],
end:"13753"}
```

# Network Management Example

# Practical Cypher

## Network Management - Create



### CREATE

```
(crm {name:"CRM"}) ,  
(dbvm {name:"Database VM"}) ,  
(www {name:"Public Website"}) ,  
(wwwvm {name:"Webserver VM"}) ,  
(srv1 {name:"Server 1"}) ,  
(san {name:"SAN"}) ,  
(srv2 {name:"Server 2"}) ,  
  
(crm)-[:DEPENDS_ON]->(dbvm) ,  
(dbvm)-[:DEPENDS_ON]->(srv2) ,  
(srv2)-[:DEPENDS_ON]->(san) ,  
(www)-[:DEPENDS_ON]->(dbvm) ,  
(www)-[:DEPENDS_ON]->(wwwvm) ,  
(wwwvm)-[:DEPENDS_ON]->(srv1) ,  
(srv1)-[:DEPENDS_ON]->(san)
```



# Practical Cypher

## Network Management - Impact Analysis

```
// Server 1 Outage
MATCH (n)<-[ :DEPENDS_ON* ]-(upstream)
WHERE n.name = "Server 1"
RETURN upstream
```

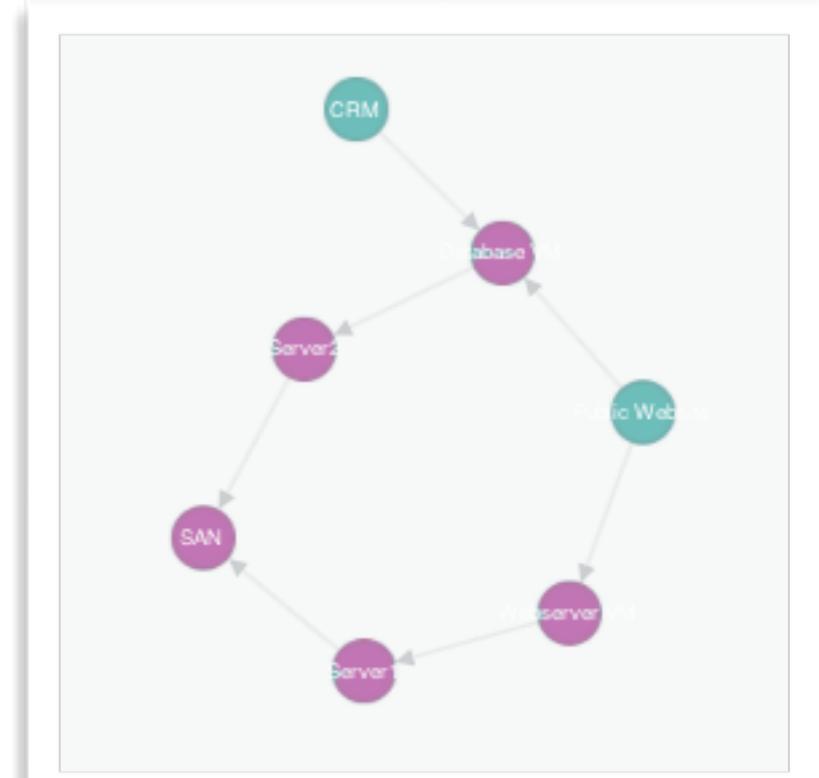


upstream
{name:"Webserver VM"}
{name:"Public Website"}

# Practical Cypher

## Network Management - Dependency Analysis

```
// Public website dependencies
MATCH (n)-[ :DEPENDS_ON* ]->(downstream)
WHERE n.name = "Public Website"
RETURN downstream
```



downstream
{name:"Database VM"}
{name:"Server 2"}
{name:"SAN"}
{name:"Webserver VM"}
{name:"Server 1"}

# Practical Cypher

## Network Management - Statistics

```
// Most depended on component
MATCH (n)<-[ :DEPENDS_ON* ]-(dependent)
RETURN n,
       count(DISTINCT dependent)
           AS dependents
ORDER BY dependents DESC
LIMIT 1
```



n	dependents
{name: "SAN"}	6



# Questions ?

