

HIGH PERFORMANCE COMPUTING IN THE OPEN DATA SCIENCE ERA

Stan Seibert, Continuum Analytics
Sergey Maidanov, Intel



Presenter Bio



Stan Seibert

High Performance Python
Team Lead
Continuum Analytics

Stan Seibert received his Ph.D. in Physics from the University of Texas at Austin.

Prior to joining Continuum, he worked at Los Alamos National Laboratory, the University of Pennsylvania, and Mobi.

Stan has 10+ years of experience in areas including scientific and parallel computing, system administration, open-source software development, and Monte Carlo simulation.

Presenter Bio



Sergey Maidanov
Software Eng. Manager
Intel

Sergey Maidanov has 15+ years of experience in numerical analysis with a range of contributions to Intel software products such as Intel® MKL, Intel® IPP, Intel compilers and others.

He currently leads a team of software engineers working on the optimized Intel® Distribution for Python. Among his recently completed projects was the Intel® Data Analytics Acceleration Library.

Sergey received a master's degree in Mathematics from the State University of Nizhny Novgorod.

Agenda

1. High Performance Python for Data Science
2. Finding Bottlenecks with Profiling
3. Intel Tools and Libraries for High Performance
4. Speeding up Python with Compilation
5. Conclusion

HIGH PERFORMANCE PYTHON FOR DATA SCIENCE



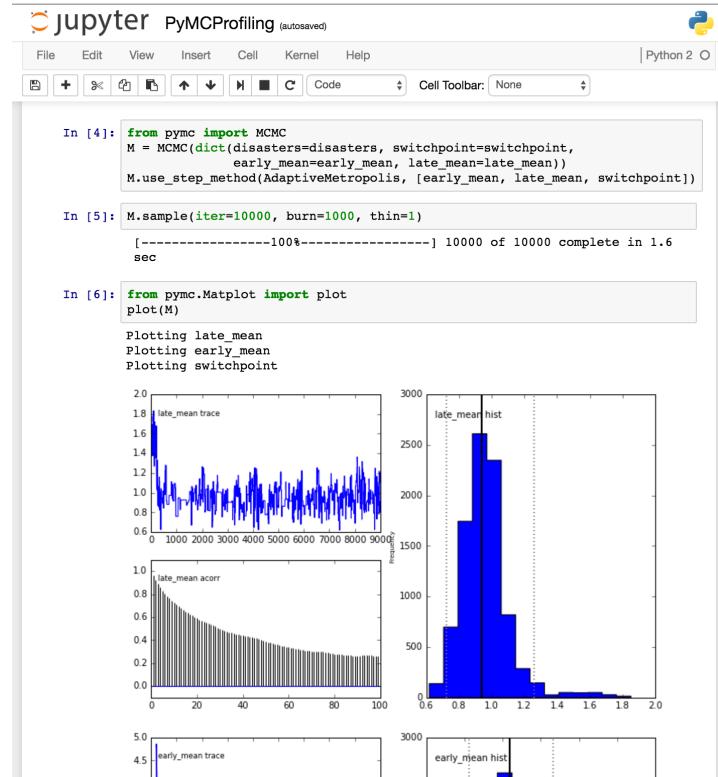
ANACONDA[®]

The Challenge of Data Science

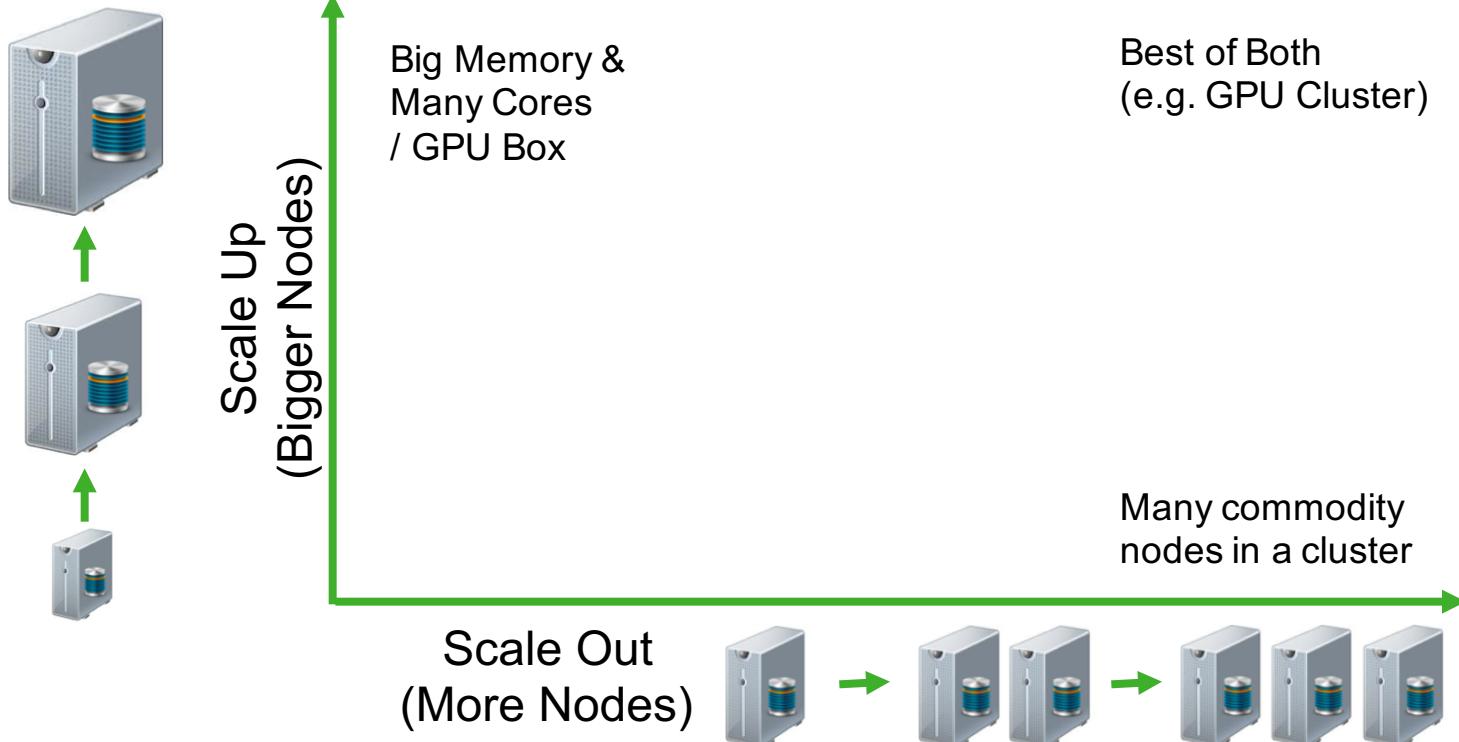
- Data sizes constantly growing
- Need high performance to turn data into insight:
 - *Faster*: React more quickly to information
 - *Cheaper*: Spend less on computing
 - *Better*: Tackle previously impossible tasks

The Transformative Power of Performance

- Performance also enables interactivity
- Interactive analysis allows algorithms to extend your intuition!



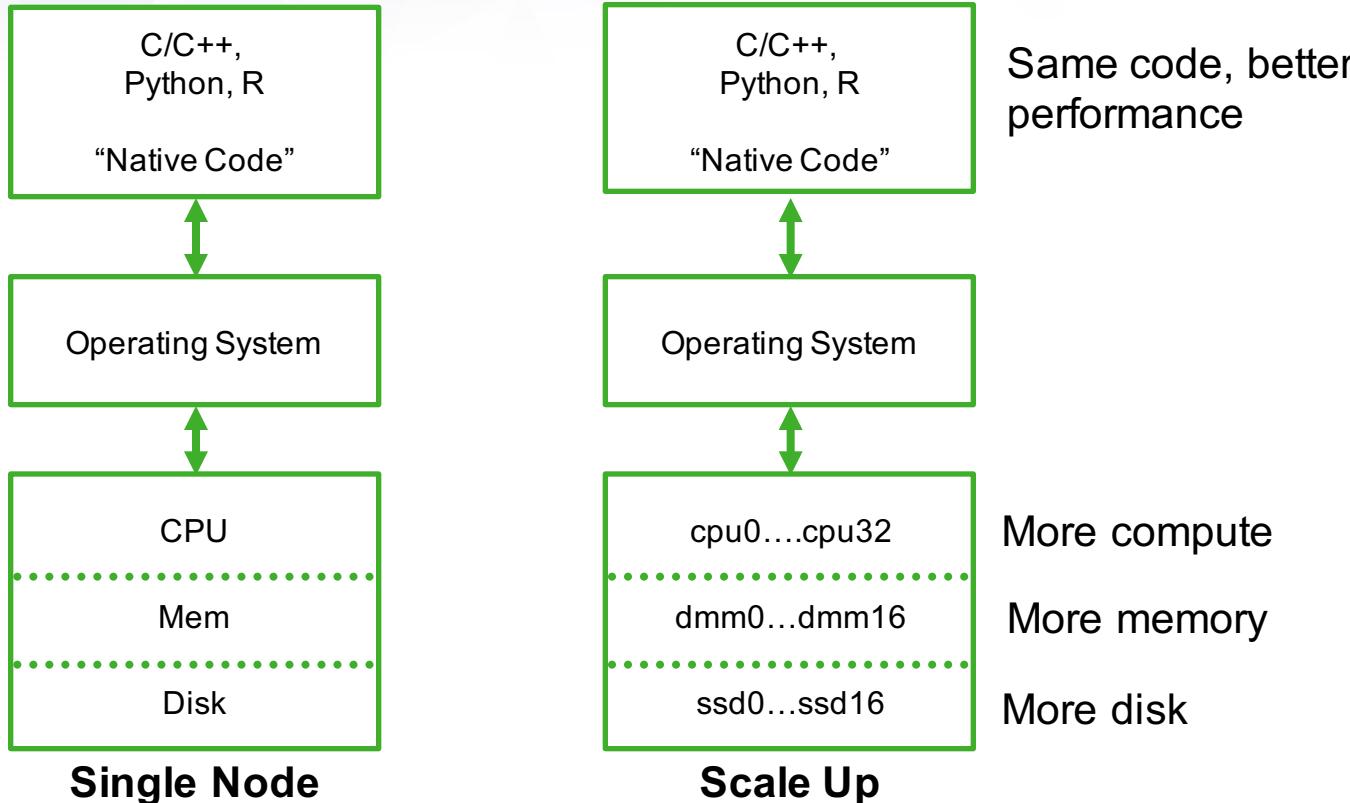
Scale Up vs Scale Out



What to scale?

- Roughly speaking, program execution time is taken up by:
 1. **I/O time**: Time spent reading or writing data (files on disk, databases, cloud storage, etc.)
 2. **Compute time**: Time spent doing calculations on data already in memory
- Everyone focuses on #2, but often #1 is *equally important and frequently overlooked*

Scaling Up



Computers come in many sizes



4 CPU cores,
mobile GPU



8 CPU cores,
midrange GPU



16 CPU cores
per node,
high end GPUs



32 cores,
No GPU

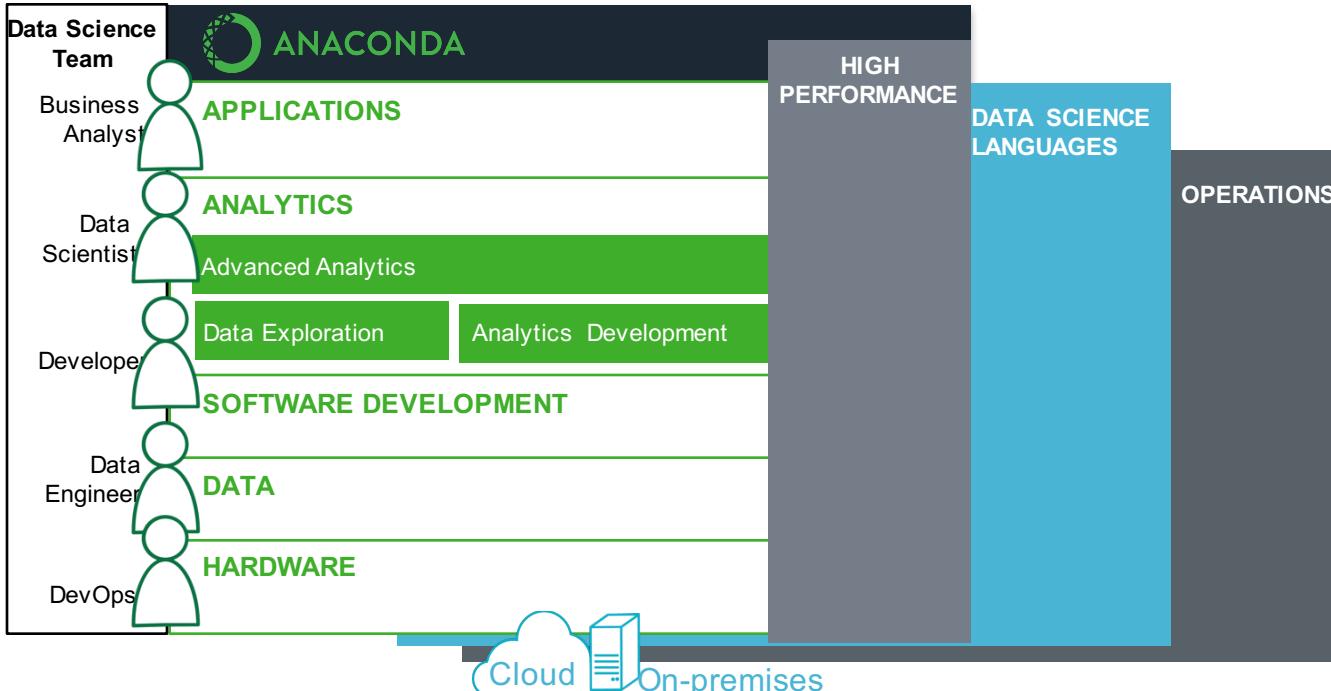
Same code, different devices, maximum performance

Why is Python Great for Data Science?

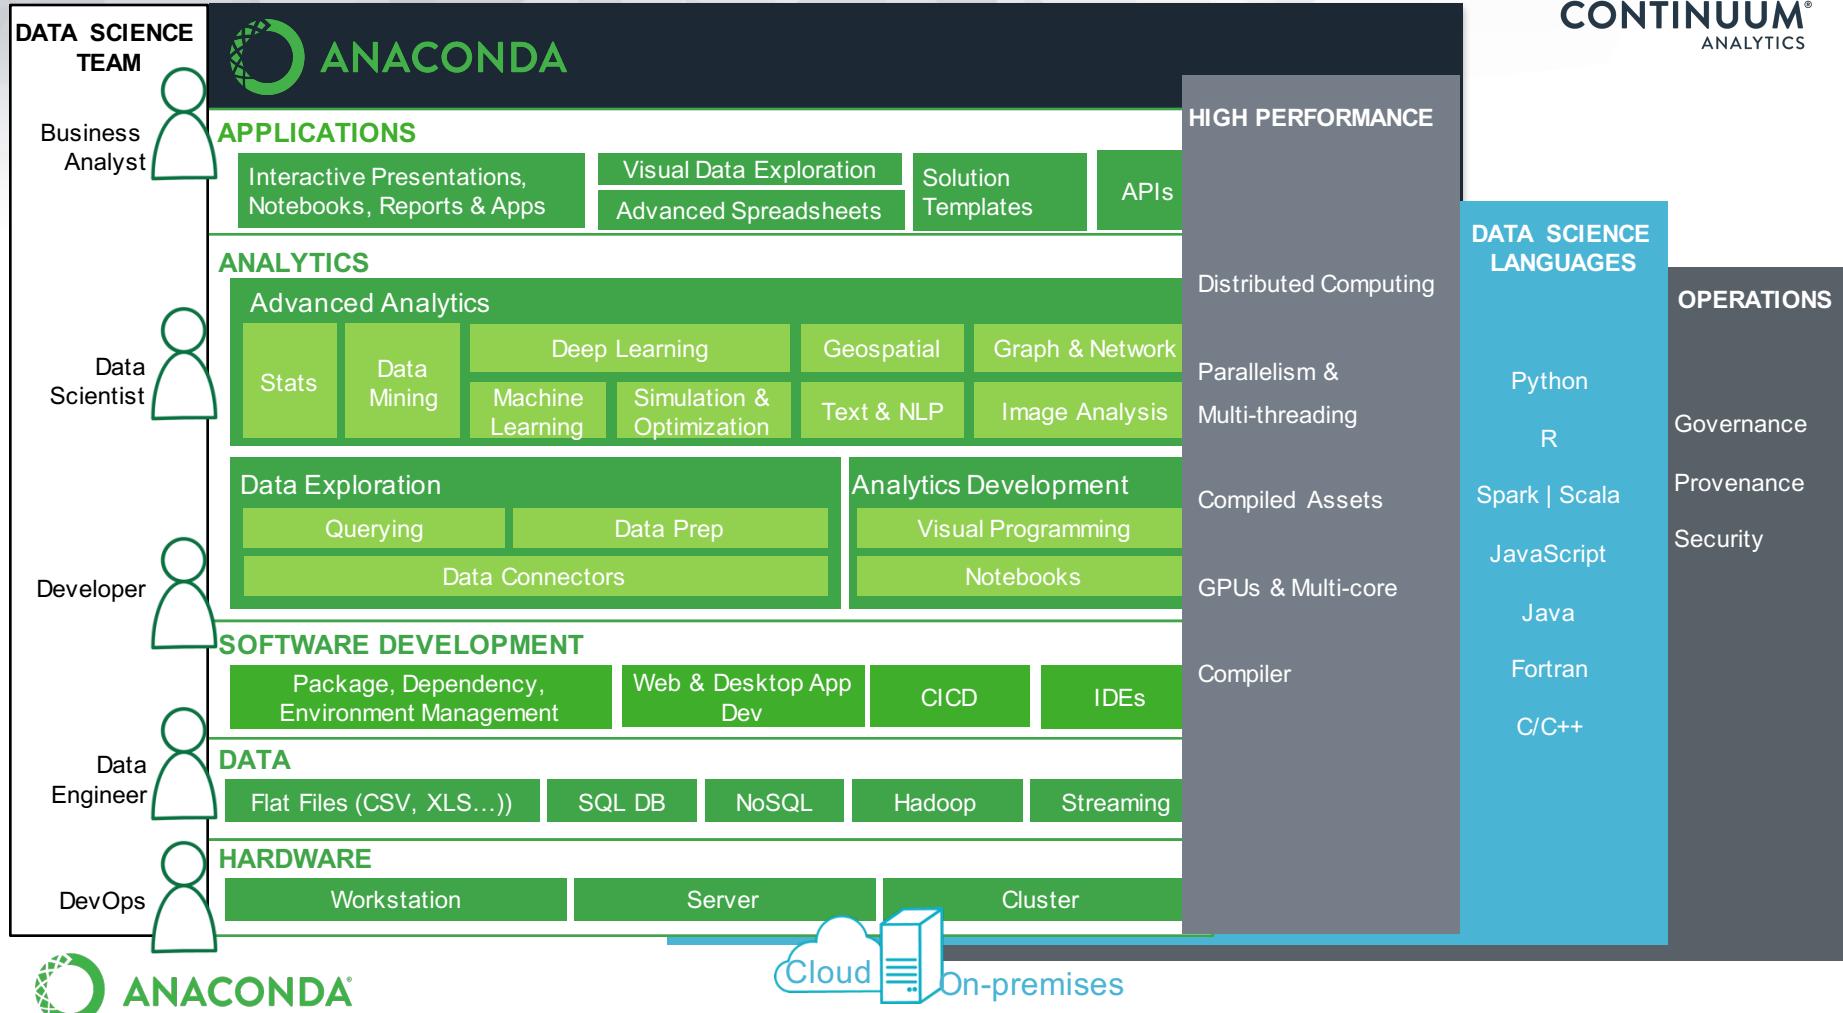
- Easy to use
 - Simple easy to read & write syntax
- “Batteries included”
 - Ships with lots of basic functions
- Innovation from Open Source
 - Open access to a huge variety of existing libraries and algorithms
- Can get high performance when you need it...
 - ... with Anaconda

Introducing Anaconda

The Open Data Science Platform Powered by Python



- Enterprise Ready Platform
 - Simplify administration
 - Use modern data science
 - Collaborate with the entire team
 - Leverage modern architectures
 - Integrate data sources
 - Accelerate performance



Tools for Performance

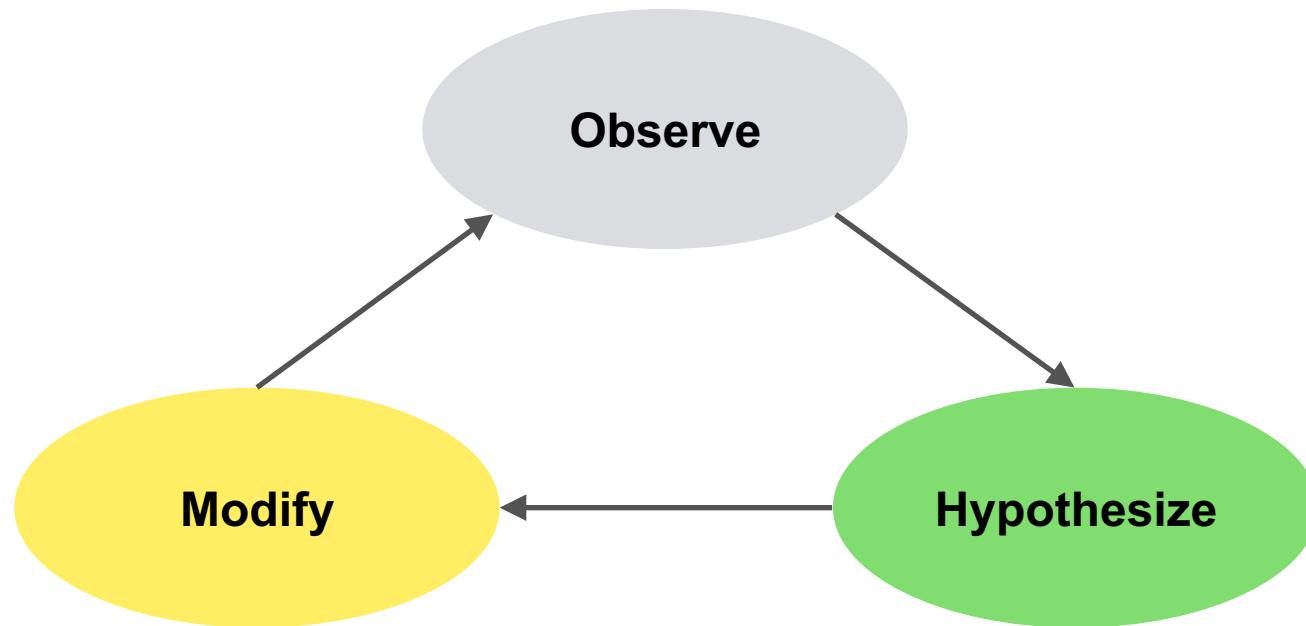
- Profiling: Find the bottlenecks
- Optimized Libraries: Solve the common problems
- Compilation: Accelerate your custom algorithms

FINDING BOTTLENECKS WITH PROFILING



ANACONDA[®]

Optimization is the Scientific Process



Profiling: Observe your Algorithm

- A profiler records the amount of time each part of your program takes to execute.
- Usually record data at the function level, but some go down to the line level.
- We'll talk about both high and low level profilers.

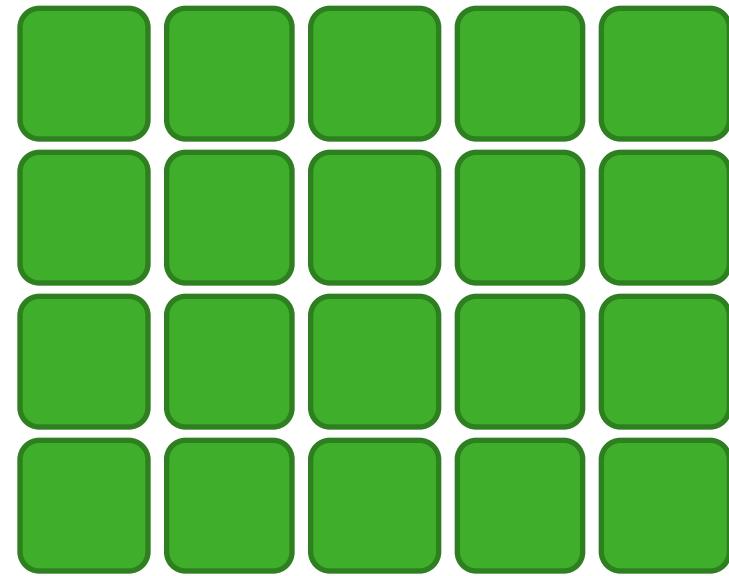
The Anaconda Profiler

- Included with Anaconda Workgroup and Enterprise
- Extends Python profiler with “data profiling”
 - Record the data types and array sizes
 - Data types and sizes are critical for deciding how to optimize a function

Data-Driven Optimization Strategies



Large Data: Look to more threads,
SIMD, memory access patterns



Small Data: Look to batch calls,
group data into larger blocks, caching

Visualization with the Anaconda Profiler

```
In [9]: profiler.plot(p)
```

```
Out[9]:
```

ResetCall Stack

Style: Sunburst

Depth: 5

Cutoff: 1/1000

Name:
Exponential.set_value(value:float64, force:bool)
Cumulative Time:
1.65 s (7.00 %)
File:
PyMCOBJects.py
Line:
837
Directory:
/Users/stan/anaconda/envs/pymc/lib/python2.7/site-packages/pymc/



Visualization with the Anaconda Profiler

Search:

ncalls	tottime	percall	cumtime	percall	filename:lineno(function)
150/50	0.006636	0.0001327	0.01559	0.0003117	linalg.py:532(cholesky(a:ndarray(dtype=float64, shape=(3, 3))))
100/50	0.00222	4.44e-05	0.00376	7.52e-05	linalg.py:139(_commonType())
100/50	0.001763	3.526e-05	0.002728	5.456e-05	linalg.py:106(_makearray(a:ndarray(dtype=float64, shape=(3, 3))))
200/100	0.001399	1.399e-05	0.001399	1.399e-05	linalg.py:111(isComplexType(t:type))
150/50	0.001087	2.174e-05	0.001087	2.174e-05	linalg.py:209(_assertNdSquareness())
100/50	0.000834	1.668e-05	0.000834	1.668e-05	linalg.py:124(_realType(t:type, default:NoneType))
100/50	0.000588	1.176e-05	0.000588	1.176e-05	linalg.py:198(_assertRankAtLeast2())
50	9.5e-05	1.9e-06	9.5e-05	1.9e-06	linalg.py:101(get_linalg_error_extobj(callback:function))

Showing 1 to 8 of 8 entries (filtered from 161 total entries)

INTEL TOOLS AND LIBRARIES FOR HIGH PERFORMANCE



ANACONDA[®]



Boost Python* Performance with Intel® Tools

Sergey Maidanov

Software Engineering Manager for Intel® Distribution for Python*

Performance Profiling with Intel® Vtune™ Amplifier

What makes Intel® VTune™ Amplifier special

Right tool for high performance application profiling at all levels

- Function-level and line-level hotspot analysis, down to disassembly
- Call stack analysis
- Low overhead
- Mixed-language, multi-threaded application analysis
- Advanced hardware event analysis for native codes (Cython, C++, Fortran) for cache misses, branch misprediction, etc.

Feature	cProfile	Line_profiler	Intel® VTune™ Amplifier
Profiling technology	Event	Instrumentation	Sampling, hardware events
Analysis granularity	Function-level	Line-level	Line-level, call stack, time windows, hardware events
Intrusiveness	Medium (1.3-5x)	High (4-10x)	Low (1.05-1.3x)
Mixed language programs	Python	Python	Python, Cython, C++, Fortran

Download Intel® VTune™ Amplifier 2017 Beta for evaluation

Optimization Notice

Copyright © 2016, Intel Corporation. All rights reserved.

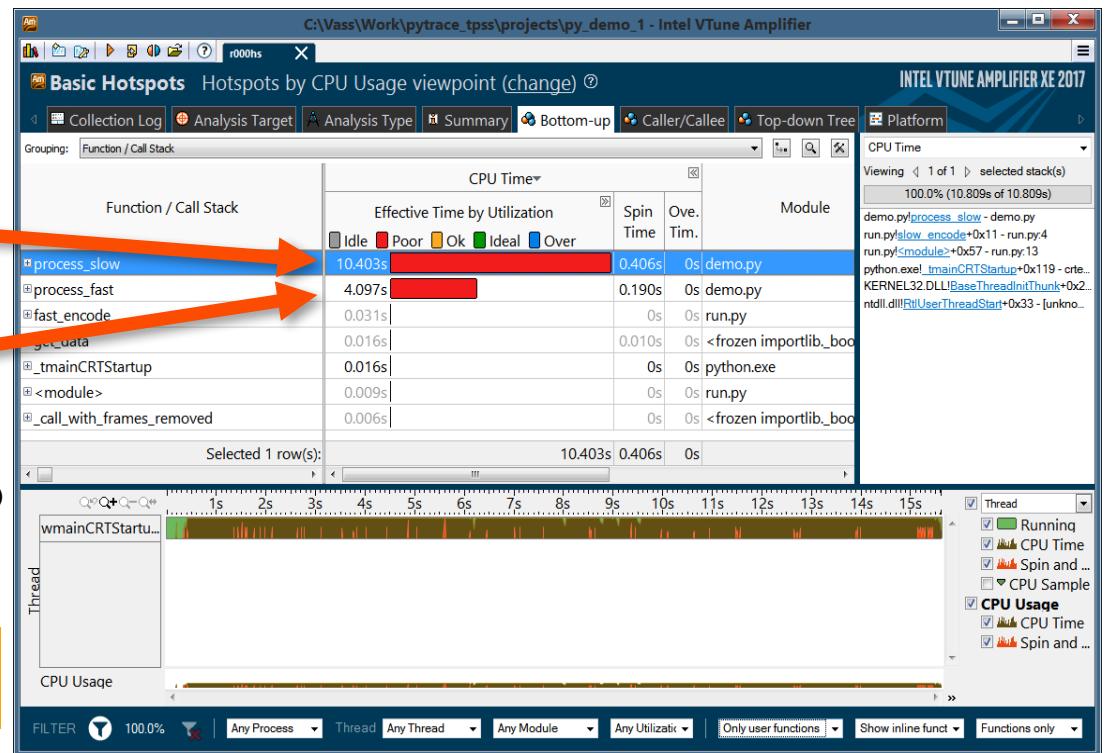
*Other names and brands may be claimed as the property of others.



Encoder demo: Strings vs. Lists

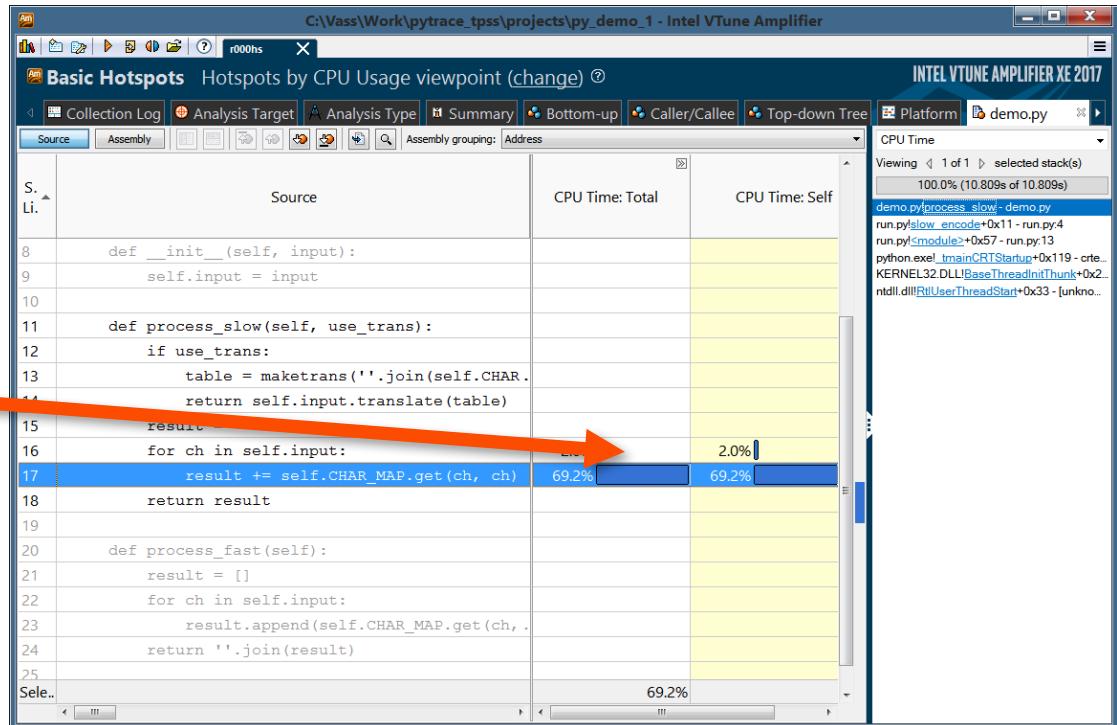
```
class Encoder:  
    CHAR_MAP = { 'a': 'b', 'b': 'c'}  
    def __init__(self, input):  
        self.input = input  
  
    def process_slow(self):  
        result = ''  
        for ch in self.input:  
            result += self.CHAR_MAP.get(ch, ch)  
        return result  
  
    def process_fast(self):  
        result = []  
        for ch in self.input:  
            result.append(self.CHAR_MAP.get(ch, ch))  
        return ''.join(result)
```

Strings concatenation is >2x worse than appending to a list

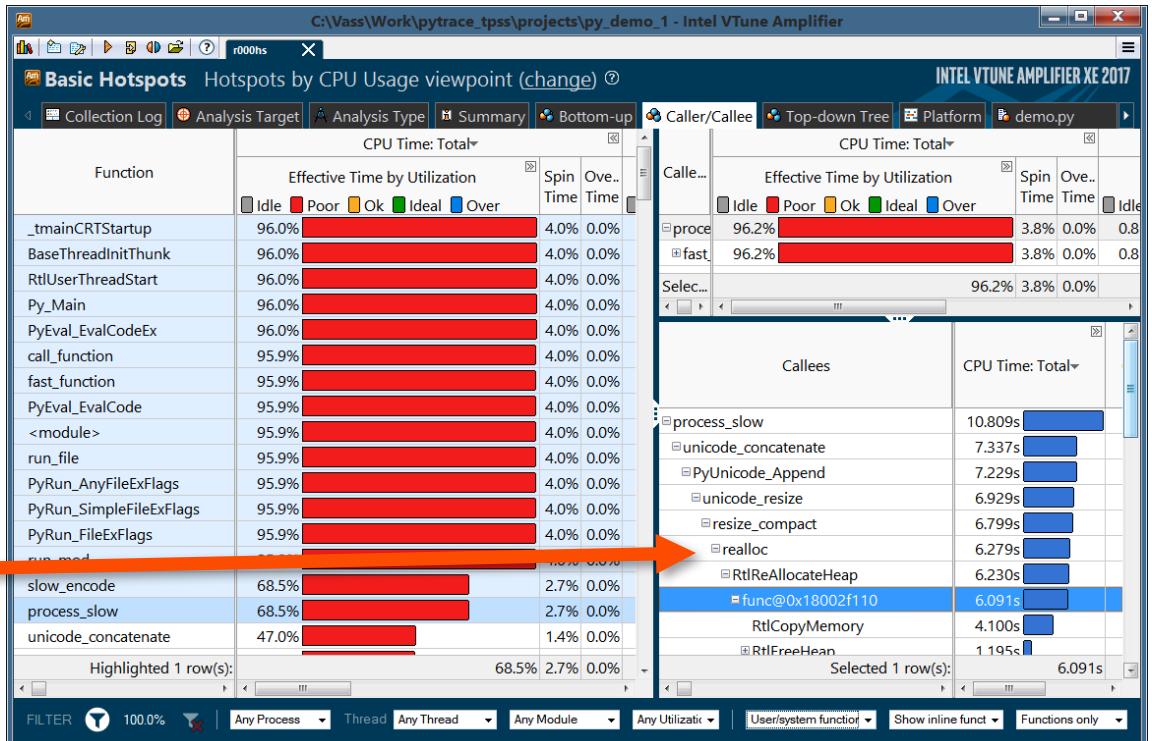


Encoder demo: Line-level statistics

70% time spent to retrieve a value from the dictionary and to concatenate it to the string



Encoder demo: What's in native code (call stack)



Each concatenation results in expensive memory reallocation

Optimization Notice

Copyright © 2016, Intel Corporation. All rights reserved.
*Other names and brands may be claimed as the property of others.



How Intel Enables Python Performance Through Libraries



Energy



Signal Processing



Financial Analytics



Engineering Design



Digital Content Creation



Science & Research

Family of Intel libraries optimized for IA

- Numerical computing, data analytics and data processing
 - Intel® Math Kernel Library
 - Intel® Data Analytics Acceleration Library
 - Intel® Integrated Performance Primitives
- Multi-threading
 - Intel® Treading Building Blocks
- Multi-node parallelism
 - Intel® MPI



Energy



Signal Processing



Financial Analytics



Engineering Design



Digital Content Creation



Science & Research

Optimization Notice

Copyright © 2016, Intel Corporation. All rights reserved.

*Other names and brands may be claimed as the property of others.



Our approach

1. Enable hooks to optimized Intel libraries in the most popular numerical/data processing packages
 - NumPy, SciPy, Scikit-Learn, PyTables, Scikit-Image, MPI4Py, ...
2. Available through Intel® Distribution for Python* and as Conda packages that can be installed in your Anaconda* distribution
 - Most optimizations eventually upstreamed to home open source projects and available through other Python distributions

More cores → More Threads → Wider vectors

The diagram illustrates the evolution of Intel processors from single-core to multi-core. It shows a sequence of images: a single-core processor, followed by two-core processors (5100 and 5500 series), four-core processors (5600 series), six-core processors (E5-2600 v2 series), twelve-core processors (E5-2600 v3 series), and eighteen-core processors (Future Xeon Processor). This leads to the Xeon Phi coprocessor, which features multiple cores and wider SIMD vectors.

Processor	Core(s)	Threads	SIMD Width	Vector ISA
Intel® Xeon® Processor 64-bit	1	2	128	Intel® SSE3
Intel® Xeon® Processor 5100 series	2	8	128	Intel® SSE3
Intel® Xeon® Processor 5500 series	4	12	128	Intel® SSE4.2
Intel® Xeon® Processor 5600 series	6	24	256	Intel® AVX
Intel® Xeon® Processor E5-2600 v2 series	12	36	256	Intel® AVX
Intel® Xeon® Processor E5-2600 v3 series	18	tbd	~ 512	Intel® AVX2
Future Intel® Xeon® Processor ¹	Tbd	tbd	~ 512	Intel® AVX-512
Intel® Xeon Phi™ x100 Coprocessor	57-61	TBD	IMCI 512	IMCI 512
Intel® Xeon Phi™ x200 Processor & Coprocessor	228-244	TBD	512	Intel® AVX-512

Optimization Notice

Copyright © 2016, Intel Corporation. All rights reserved.

*Other names and brands may be claimed as the property of others.



How Intel Enables Python Performance Through Intel® MKL



Energy



Signal Processing



Financial Analytics



Engineering Design



Digital Content Creation



Science & Research

Optimized mathematical building blocks

Intel® Math Kernel Library (Intel MKL)

Linear Algebra

- BLAS
- LAPACK
- ScaLAPACK
- Sparse BLAS
- Sparse Solvers
 - Iterative
- PARDISO* SMP & Cluster

Up to
100x
faster

Fast Fourier Transforms

- Multidimensional
- FFTW interfaces
- Cluster FFT

Up to
10x
faster!

Vector Math

- Trigonometric
- Hyperbolic
- Exponential
- Log
- Power
- Root

Up to
10x
faster!

Vector RNGs

- Multiple BRNG
- Support methods for independent streams creation
- Support all key probability distributions

Up to
60x
faster!

Summary Statistics

- Kurtosis
- Variation coefficient
- Order statistics
- Min/max
- Variance-covariance

And More

- Splines
- Interpolation
- Trust Region
- Fast Poisson Solver

Optimization Notice

Copyright © 2016, Intel Corporation. All rights reserved.
*Other names and brands may be claimed as the property of others.

Functional domain in **this color** accelerate respective NumPy, SciPy, etc. domain

Configuration info: - Versions: Intel® Distribution for Python 2017 Beta, icc 15.0; Hardware: Intel® Xeon® CPU E5-2698 v3 @ 2.30GHz (2 sockets, 16 cores each, HT=OFF), 64 GB of RAM, 8 DIMMS of 8GB@2133MHz; Operating System: Ubuntu 14.04 LTS.



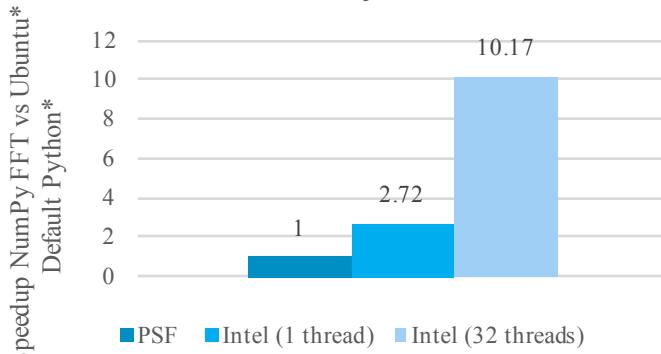
Optimized FFT show case

Intel® Math Kernel Library (Intel MKL)

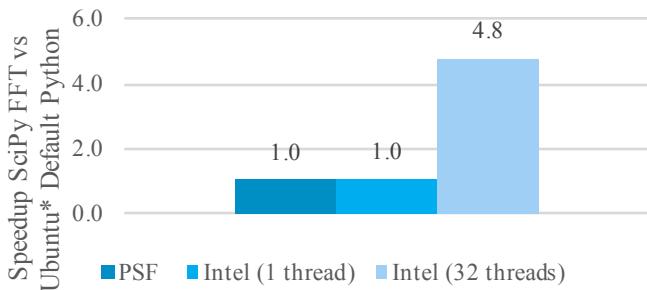
- Original SciPy FFT implementation is about 2x faster than original NumPy FFT
- Intel engineers bridged NumPy and SciPy implementations via common layer and embedded MKL FFT calls, what measurably accelerates both NumPy and SciPy
 - NumPy and SciPy are computationally compatible
 - FFT descriptors caching applied for enhanced performance in repetitive and multidimensional FFT calculations

Available starting Intel® Distribution for Python* 2017 Beta

NumPy FFT NumPy vs Ubuntu*
Default Python*



SciPy FFT Intel SciPy vs Ubuntu*
Vanilla Python*



Optimization Notice

Copyright © 2016, Intel Corporation. All rights reserved.

*Other names and brands may be claimed as the property of others.

Configuration info: - Versions: Intel® Distribution for Python 2017 Beta, icc 15.0; Hardware: Intel® Xeon® CPU E5-2698 v3 @ 2.30GHz (2 sockets, 16 cores each, HT=OFF), 64 GB of RAM, 8 DIMMS of 8GB@2133MHz; Operating System: Ubuntu 14.04 LTS.

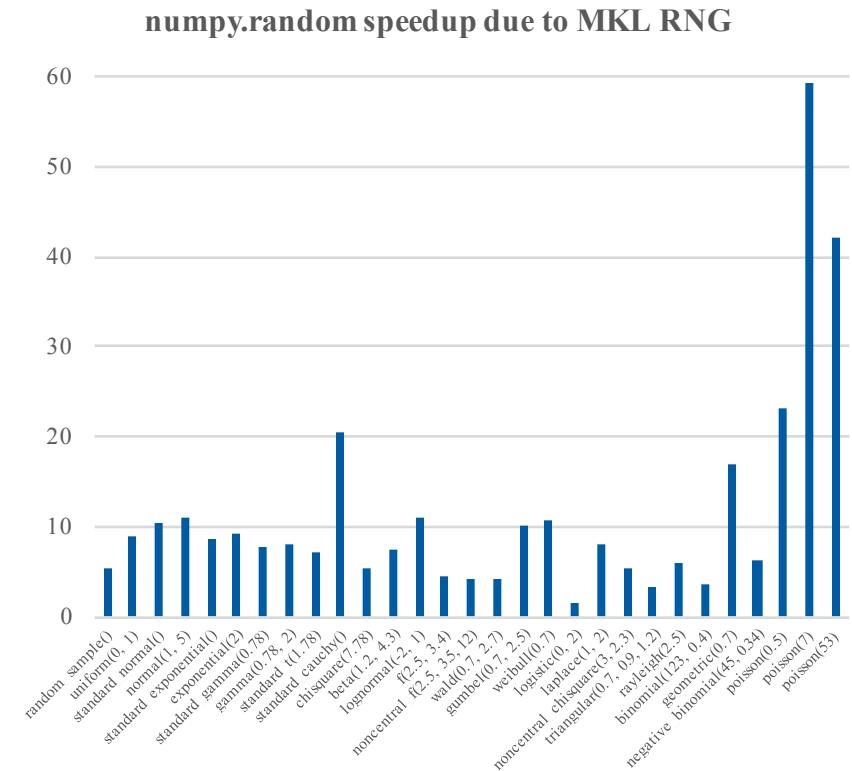


Optimized RNG show case

Intel® Math Kernel Library (Intel MKL)

- Implemented `numpy.random` in vector fashion to enable vector MKL RNG and VML calls
- Enabled multiple BRNG
- Enabled multiple distribution transformation methods

Initial data. Final data to be available in the update for Intel® Distribution for Python* 2017 Beta



Optimization Notice

Copyright © 2016, Intel Corporation. All rights reserved.
*Other names and brands may be claimed as the property of others.

Configuration info: - Versions: Intel® Distribution for Python (internal build 28.42016), ioc 15.0; Hardware: Intel® Xeon® CPU E5-2630 v3 @ 2.40GHz (16 cores), 32 GB; Operating System: Ubuntu 14.04 LTS.



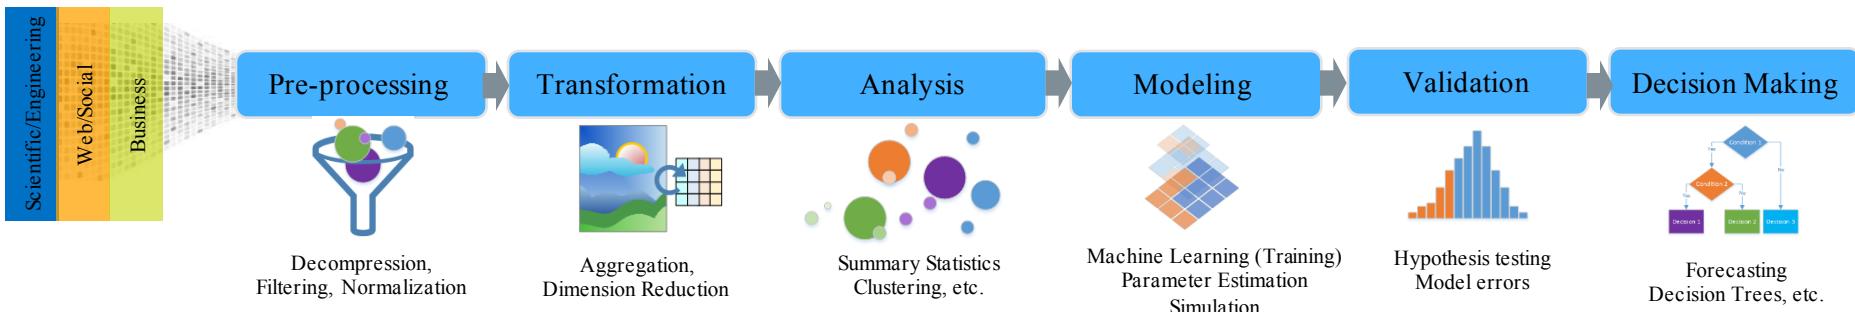
Intel® Data analytics acceleration library



Optimized blocks for data analytics pipelines

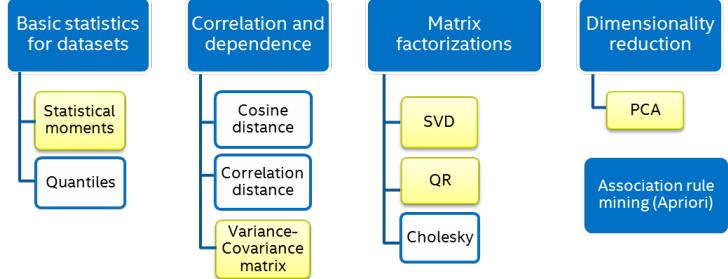
Intel® Data Analytics Acceleration Library (Intel DAAL)

- Optimizes entire analytics pipeline, from data acquisition to training and scoring
 - Allows different stages to be executed on different devices, e.g. server and edge
- Bridges prototyping and production analytics
 - Multiple language interfaces C++, Java, Python
 - Runtime CPU detection, runs best from Intel® Atom to Intel® Xeon and Intel® Xeon Phi processor families
 - Code samples for MPI, Hadoop, Spark

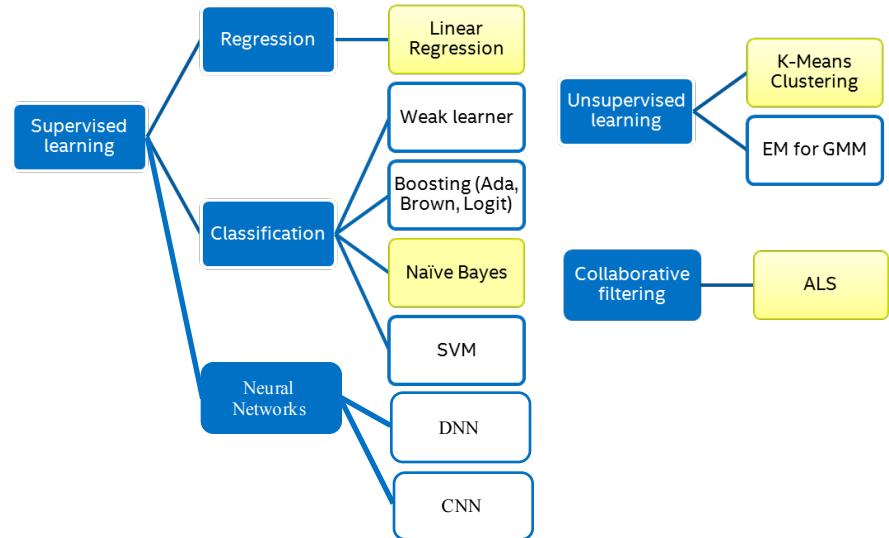


Intel DAAL algorithms at a glance

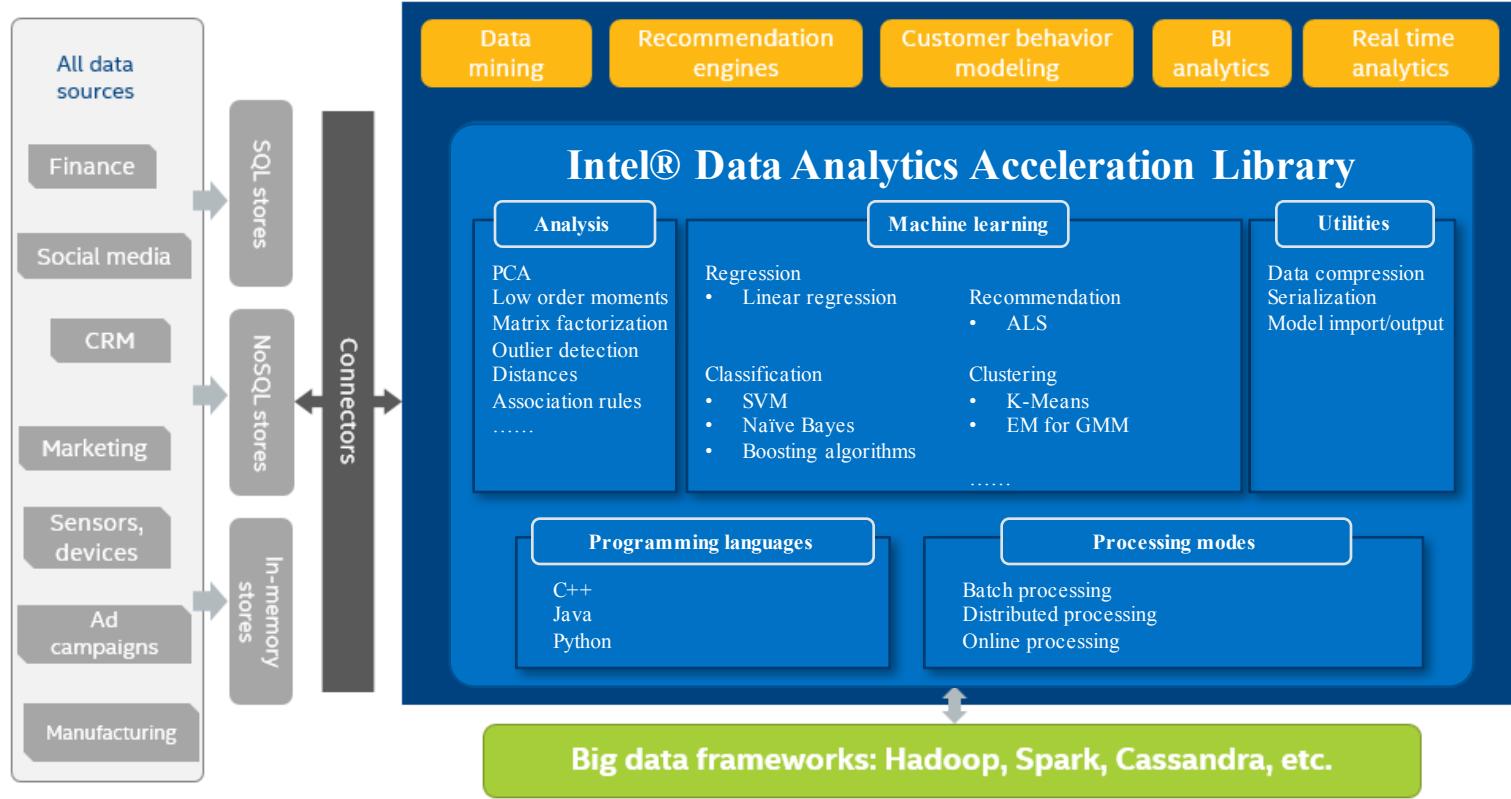
- Classic and modern data mining and machine learning algorithms
 - Supervised and unsupervised learning
 - Neural networks (DNN, CNN)
 - Batch, online, and distributed



 Algorithms support streaming and distributed processing in the current release.



Where DAAL fits



How DAAL accelerates Python

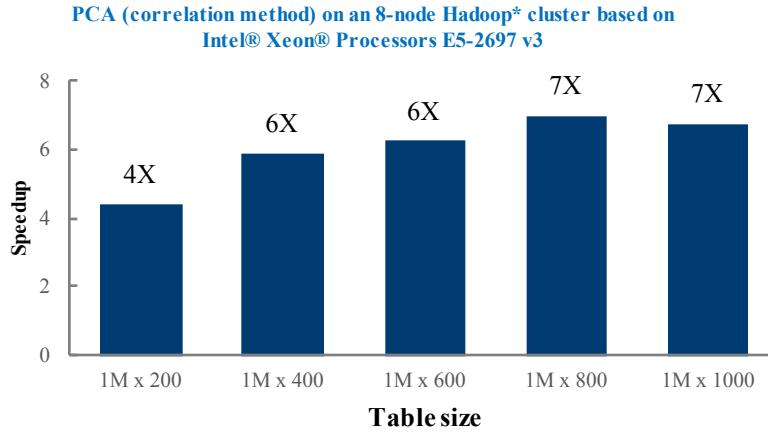
1. Available as PyDAAL package

- Pre-installed with Intel® Distribution for Python*
- Through anaconda.org as Conda* package
- Part of Intel® DAAL and Intel® Parallel Studio XE products

2. Accelerates Scikit-learn

- Work in progress, more news soon

PCA Performance Boosts Using Intel® DAAL vs. Spark* MLLib



Configuration Info - Versions: Intel® Data Analytics Acceleration Library 2016, CDH v5.3.1, Apache Spark* v1.2.0; Hardware: Intel® Xeon® Processor E5-2699 v3, 2 Eighteen-core CPUs (45MB LLC, 2.3GHz), 128GB of RAM per node; Operating System: CentOS 6.6 x86_64.

Software and workloads used in performance tests may have been optimized for performance only on Intel microprocessors. Performance tests, such as SYSmark and MobileMark, are measured using specific computer systems, components, software, operations and functions. Any change to any of those factors may cause the results to vary. You should consult other information and performance tests to assist you in fully evaluating your contemplated purchases, including the performance of that product when combined with other products. * Other brands and names are the property of their respective owners. Benchmark Source: Intel Corporation

Optimization Notice: Intel's compilers may or may not optimize to the same degree for non-Intel microprocessors for optimizations that are not unique to Intel microprocessors. These optimizations include SSE2, SSE3, and SSSE3 instruction sets and other optimizations. Intel does not guarantee the availability, functionality, or effectiveness of any optimization on microprocessors not manufactured by Intel. Microprocessor-dependent optimizations in this product are intended for use with Intel microprocessors. Certain optimizations not specific to Intel microarchitecture are reserved for Intel microprocessors. Please refer to the applicable product User and Reference Guides for more information regarding the specific instruction sets covered by this notice. Notice revision #20110804 .

Optimization Notice

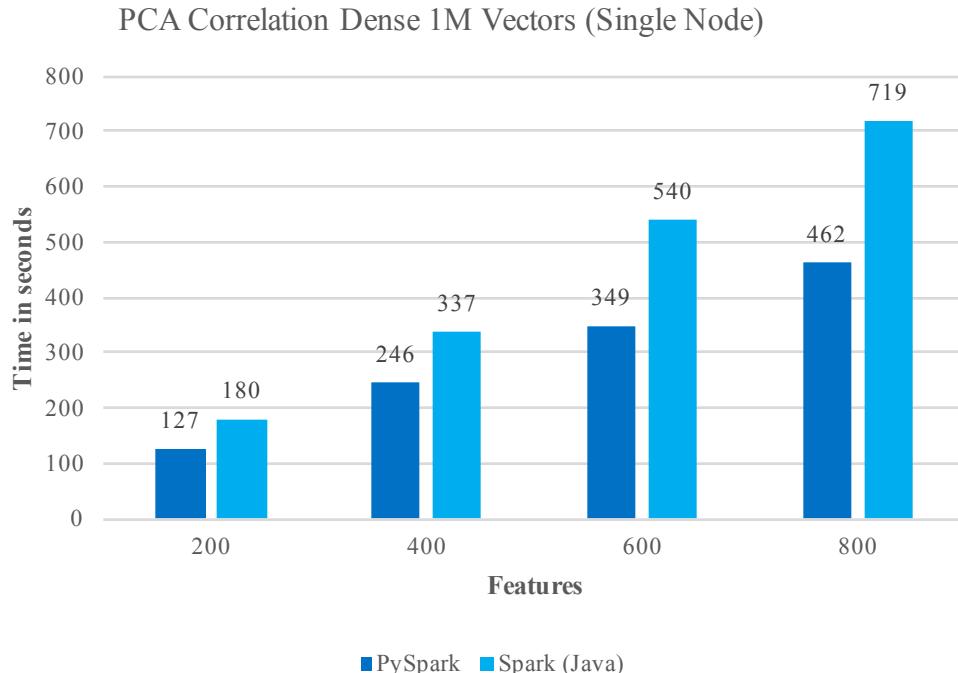
Copyright © 2016, Intel Corporation. All rights reserved.

*Other names and brands may be claimed as the property of others.



PyDAAL and PySpark*

Work in progress to optimize PyDAAL distributed algorithms for PySpark*



Configuration Info - Versions: Intel® Data Analytics Acceleration Library 2016, CDH v5.3.1, Apache Spark* v1.2.0; Hardware: Intel® Xeon® Processor E5-2699 v3, 2 Eighteen-core CPUs (45MB LLC, 2.3GHz), 128GB of RAM per node; Operating System: CentOS 6.6 x86_64.

Software and workloads used in performance tests may have been optimized for performance only on Intel microprocessors. Performance tests, such as SYSmark and MobileMark, are measured using specific computer systems, components, software, operations and functions. Any change to any of those factors may cause the results to vary. You should consult other information and performance tests to assist you in fully evaluating your contemplated purchases, including the performance of that product when combined with other products. * Other brands and names are the property of their respective owners. Benchmark Source: Intel Corporation

Optimization Notice: Intel's compilers may or may not optimize to the same degree for non-Intel microprocessors for optimizations that are not unique to Intel microprocessors. These optimizations include SSE2, SSE3, and SSSE3 instruction sets and other optimizations. Intel does not guarantee the availability, functionality, or effectiveness of any optimization on microprocessors not manufactured by Intel. Microprocessor-dependent optimizations in this product are intended for use with Intel microprocessors. Certain optimizations not specific to Intel microarchitecture are reserved for Intel microprocessors. Please refer to the applicable product User and Reference Guides for more information regarding the specific instruction sets covered by this notice. Notice revision #20110804 .

Optimization Notice

Copyright © 2016, Intel Corporation. All rights reserved.

*Other names and brands may be claimed as the property of others.



SPEEDING UP PYTHON WITH COMPIRATION



ANACONDA[®]

Interpreters vs. Compilers

Compiled Language	Interpreted Language
Must be translated from source code to machine code before execution	Can be executed line-by-line immediately and interactively
Application can be distributed standalone, without the compiler.	Application must be distributed with the interpreter.
Executes much faster	Executes more slowly
Examples: C, C++, FORTRAN	Examples: Python

Just-In-Time Compilation

- JIT compilers combine benefits of interpreted and compiled languages:
 - Allow interactive development
 - Translate algorithms during execution to machine code for speed
 - Can be applied to interpreted languages, like Python and Javascript

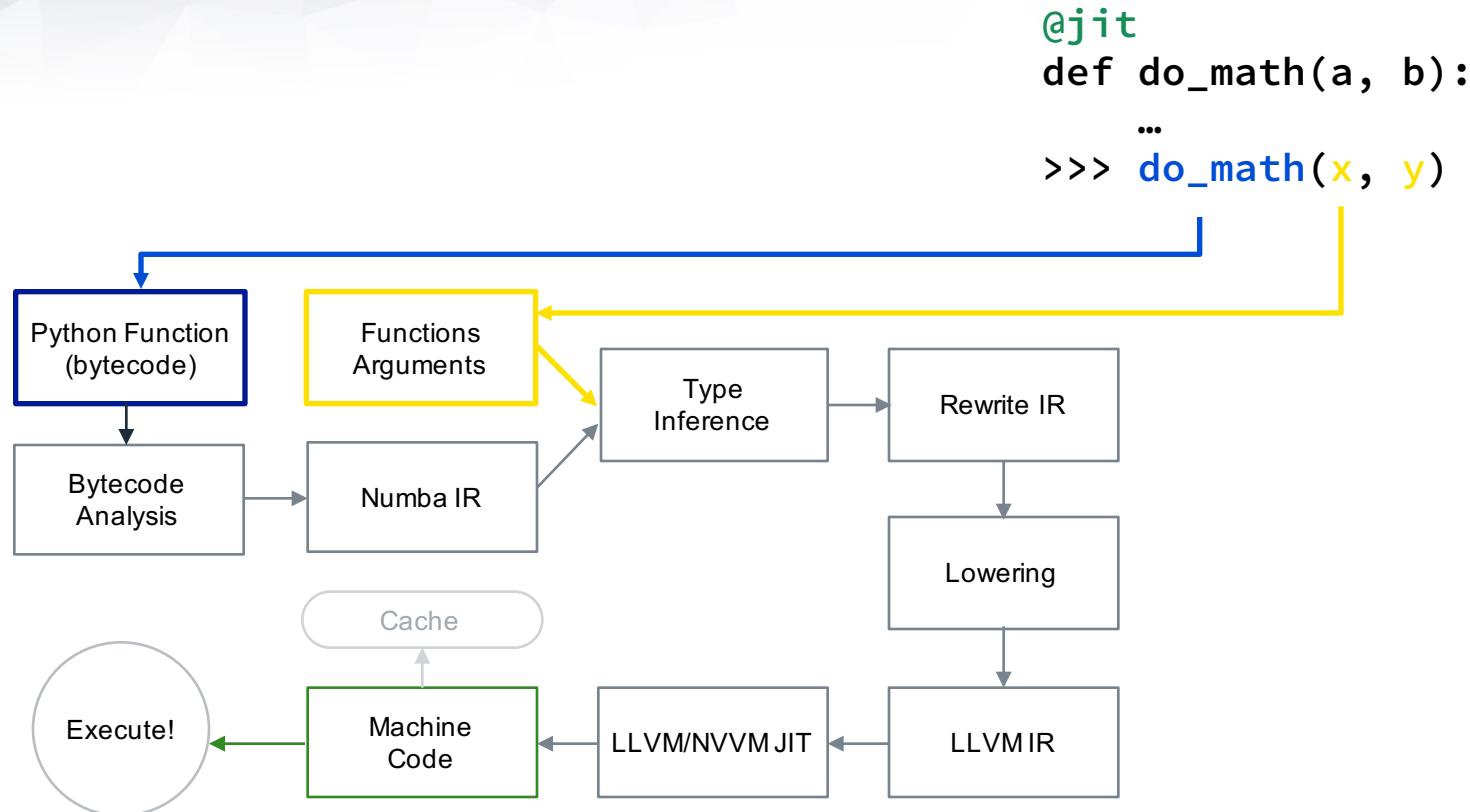
Numba: A Python JIT Compiler

- Compile specific functions, leave the rest of your application unchanged.
- Designed for numerical workloads.
- Does not replace the standard interpreter: Keep using your favorite Python packages!
- Built-in support for NumPy arrays.

Supported Platforms

OS	HW	SW
Windows (7 and later)	32 and 64-bit x86 CPUs	Python 2 and 3
OS X (10.9 and later)	CUDA-capable NVIDIA GPUs	NumPy 1.7 through 1.11
Linux (~RHEL 5 and later)	HSA-capable AMD GPUs	

How Does Numba Work?



Example: Filter an array

```
In [87]: @jit(nopython=True)
def nan_compact(x):
    out = np.empty_like(x)
    out_index = 0
    for element in x:
        if not np.isnan(element):
            out[out_index] = element
            out_index += 1
    return out[:out_index]
```

```
In [88]: a = np.random.uniform(size=10000)
a[a < 0.2] = np.nan
np.testing.assert_equal(nan_compact(a), a[~np.isnan(a)])
```

```
In [89]: %timeit a[~np.isnan(a)]
%timeit nan_compact(a)
```

```
10000 loops, best of 3: 52 µs per loop
100000 loops, best of 3: 19.6 µs per loop
```

Example: Filter an array

```
In [87]: @jit(nopython=True)
def nan_compact(x):
    out = np.empty_like(x)
    out_index = 0
    for element in x:
        if not np.isnan(element):
            out[out_index] = element
            out_index += 1
    return out[:out_index]
```

Annotations for In [87]:

- `@jit(nopython=True)` → Numba decorator (`nopython=True` not required)
- `out = np.empty_like(x)` → Array Allocation
- `for element in x:` → Looping over ndarray `x` as an iterator
- `if not np.isnan(element):` → Using numpy math functions
- `return out[:out_index]` → Returning a slice of the array

```
In [88]: a = np.random.uniform(size=10000)
a[a < 0.2] = np.nan
np.testing.assert_equal(nan_compact(a), a[~np.isnan(a)])
```

```
In [89]: %timeit a[~np.isnan(a)]
%timeit nan_compact(a)
```

10000 loops, best of 3: 52 µs per loop
 100000 loops, best of 3: 19.6 µs per loop

2.7x Speedup

CONCLUSION



ANACONDA[®]

Anaconda

Accelerating Adoption of Python for Enterprise

ANACONDA®

ENTERPRISE DATA INTEGRATION

With optimized connectors & out-of-core processing

NumPy &
Pandas

Numba

PERFORMANCE

With compiled Python for lightning fast execution

COLLABORATIVE NOTEBOOKS

With publication, authentication, & search

Jupyter/
IPython

Bokeh

VISUAL APPS

for interactivity and streaming data

Parallel Computing

Scaling up Python analytics on your cluster

Dask

Conda

SECURE & ROBUST REPOSITORY

for interactivity and streaming data

Anaconda Subscriptions

ANACONDA	ANACONDA PRO	ANACONDA WORKGROUP	ANACONDA ENTERPRISE
Open Source Modern Analytics Platform Powered by Python	Anaconda with Support & Indemnification	Anaconda with High Performance and Team Collaboration	Anaconda with Scalable High Performance and Team Collaboration
FREE FOREVER	Starting at \$10,000 per year up to 10 users +\$1,000 per year for additional user	Starting at \$30,000 per year up to 10 users More than 10 users? Contact us for a quote.	Starting at \$60,000 per year up to 10 users More than 10 users? Contact us for a quote.
Community Support	Enterprise Support	Enterprise Support	Enterprise Support with dedicated Customer Support Rep.
DOWNLOAD	LEARN MORE	CONTACT US	CONTACT US



Intel Optimization Tools and Libraries

- Intel® Distribution for Python Beta:
 - <https://software.intel.com/en-us/python-distribution>
- Intel® VTune™ Amplifier 2017 – Mixed mode profiling Beta:
 - <https://software.intel.com/en-us/python-profiling>
- Intel Software Development tools and libraries:
 - <https://software.intel.com/en-us/tools-by-segment/technical-enterprise>
- Community licensing (free) Libraries:
 - <https://software.intel.com/sites/campaigns/nest/>

Solutions For Python Performance Optimization

- **Download Anaconda** for free at Continuum.io/downloads
- **Engage with us** to help you speed up and optimize your Python at sales@Continuum.io
- **Take Performance Optimization Deep Dive** course at Continuum.io/training/
- **See Anaconda in action** at Continuum.io/webinars
- **Read our whitepapers** at Continuum.io/whitepapers



Legal Disclaimer & Optimization Notice

INFORMATION IN THIS DOCUMENT IS PROVIDED “AS IS”. NO LICENSE, EXPRESS OR IMPLIED, BY ESTOPPEL OR OTHERWISE, TO ANY INTELLECTUAL PROPERTY RIGHTS IS GRANTED BY THIS DOCUMENT. INTEL ASSUMES NO LIABILITY WHATSOEVER AND INTEL DISCLAIMS ANY EXPRESS OR IMPLIED WARRANTY, RELATING TO THIS INFORMATION INCLUDING LIABILITY OR WARRANTIES RELATING TO FITNESS FOR A PARTICULAR PURPOSE, MERCHANTABILITY, OR INFRINGEMENT OF ANY PATENT, COPYRIGHT OR OTHER INTELLECTUAL PROPERTY RIGHT.

Software and workloads used in performance tests may have been optimized for performance only on Intel microprocessors. Performance tests, such as SYSmark and MobileMark, are measured using specific computer systems, components, software, operations and functions. Any change to any of those factors may cause the results to vary. You should consult other information and performance tests to assist you in fully evaluating your contemplated purchases, including the performance of that product when combined with other products.

Copyright © 2016, Intel Corporation. All rights reserved. Intel, Pentium, Xeon, Xeon Phi, Core, VTune, Cilk, and the Intel logo are trademarks of Intel Corporation in the U.S. and other countries.

Optimization Notice

Intel’s compilers may or may not optimize to the same degree for non-Intel microprocessors for optimizations that are not unique to Intel microprocessors. These optimizations include SSE2, SSE3, and SSSE3 instruction sets and other optimizations. Intel does not guarantee the availability, functionality, or effectiveness of any optimization on microprocessors not manufactured by Intel. Microprocessor-dependent optimizations in this product are intended for use with Intel microprocessors. Certain optimizations not specific to Intel microarchitecture are reserved for Intel microprocessors. Please refer to the applicable product User and Reference Guides for more information regarding the specific instruction sets covered by this notice.

Notice revision #20110804