

ONOS Developer Workshop

for ONOS 1.2.0 (Cardinal)

Open Networking Summit
June 15, 2015



Developer Workshop Sessions



- ONOS Fundamentals
 - architecture, checking out, building & running ONOS
- ONOS Distributed Core
 - OSGi, anatomy of a core subsystem, distributed stores & primitives
- ONOS App Development
 - background, Maven archetypes, app deployment on ONOS cluster
- ONOS Interfaces
 - component config service, flow objective service, intent service
 - CLI, REST and GUI overview
- Wrap-up

ONOS Fundamentals

8:30-10:00



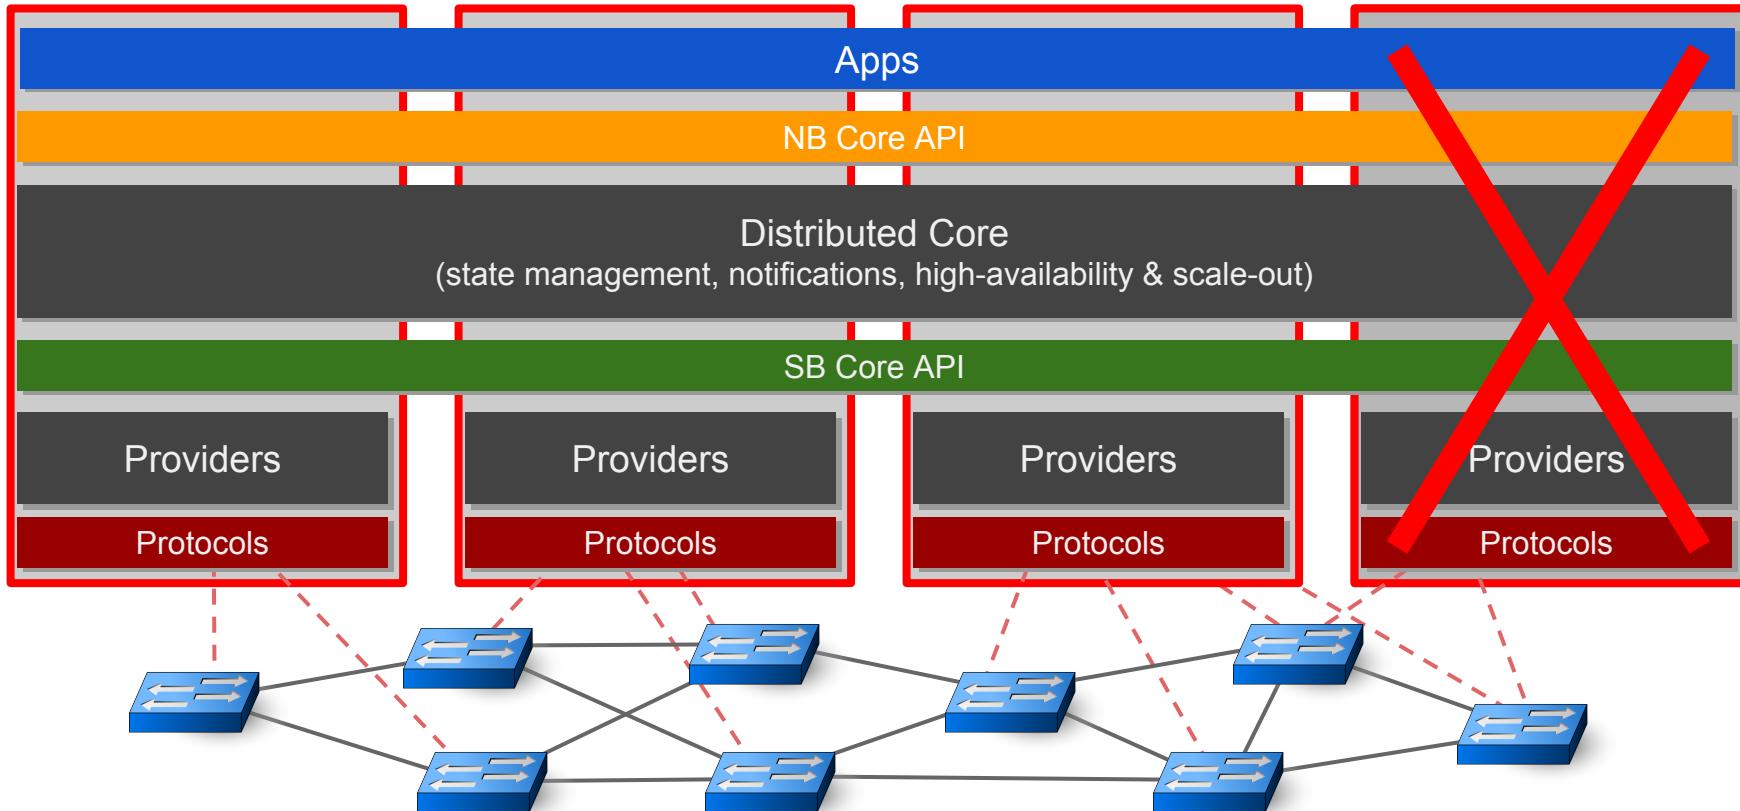
- **ONOS Overview** (~30)
 - tenets, architecture, tiers, subsystems
 - general subsystem structure
- **Getting started with ONOS** (~30)
 - check-out code, environment customization and build
 - describe Maven, hierarchical POM structure
 - run ONOS locally on developer machine
 - CLI & GUI demo with a simple mininet topology (torus,8,8,2)
- **Code organization** (~30)
 - importing ONOS source into IDE (demo with IntelliJ)
 - core (api, net, stores), web (gui, rest), providers (openflow, netconf)

Architectural Tenets

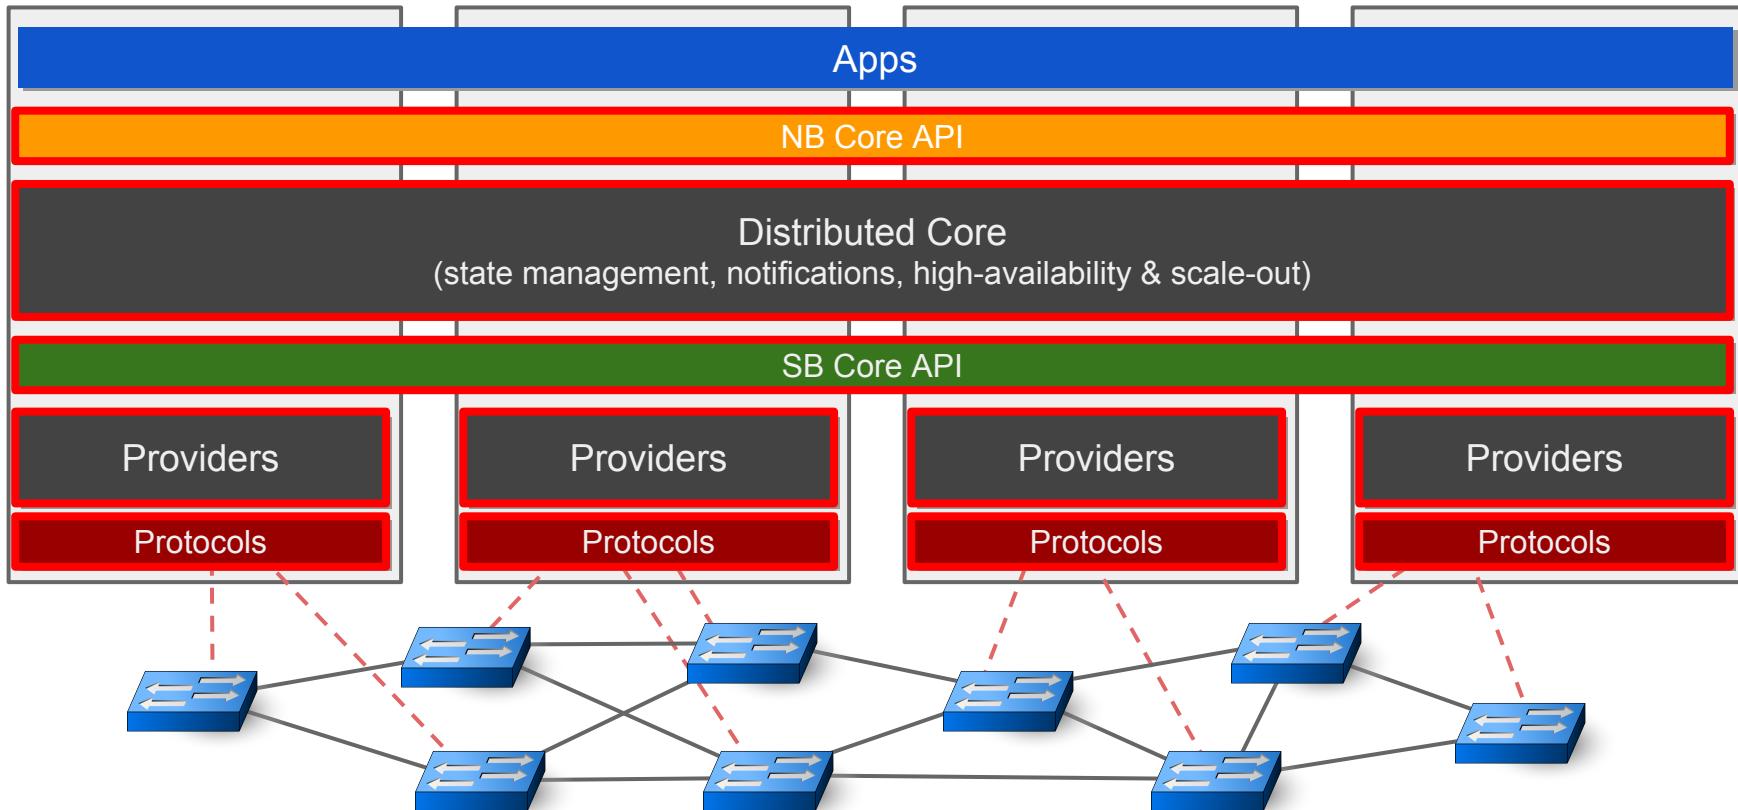


- **High-availability, scalability and performance**
 - required to sustain demands of service provider & enterprise networks
- **Strong abstractions and simplicity**
 - required for development of apps and solutions
- **Protocol and device behaviour independence**
 - avoid contouring and deformation due to protocol specifics
- **Separation of concerns and modularity**
 - allow tailoring and customization without speciating the code-base

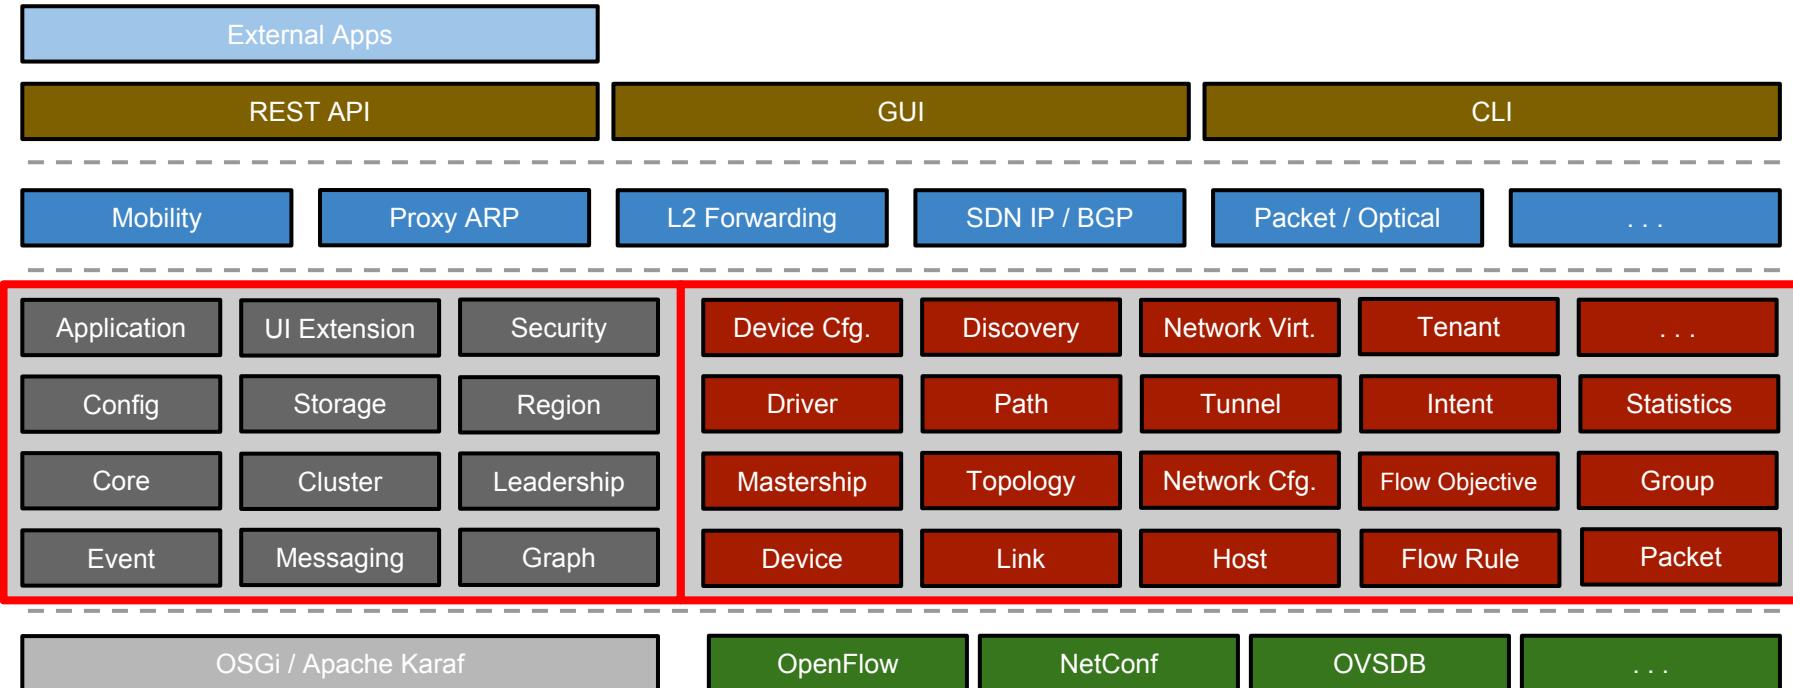
ONOS Distributed Architecture



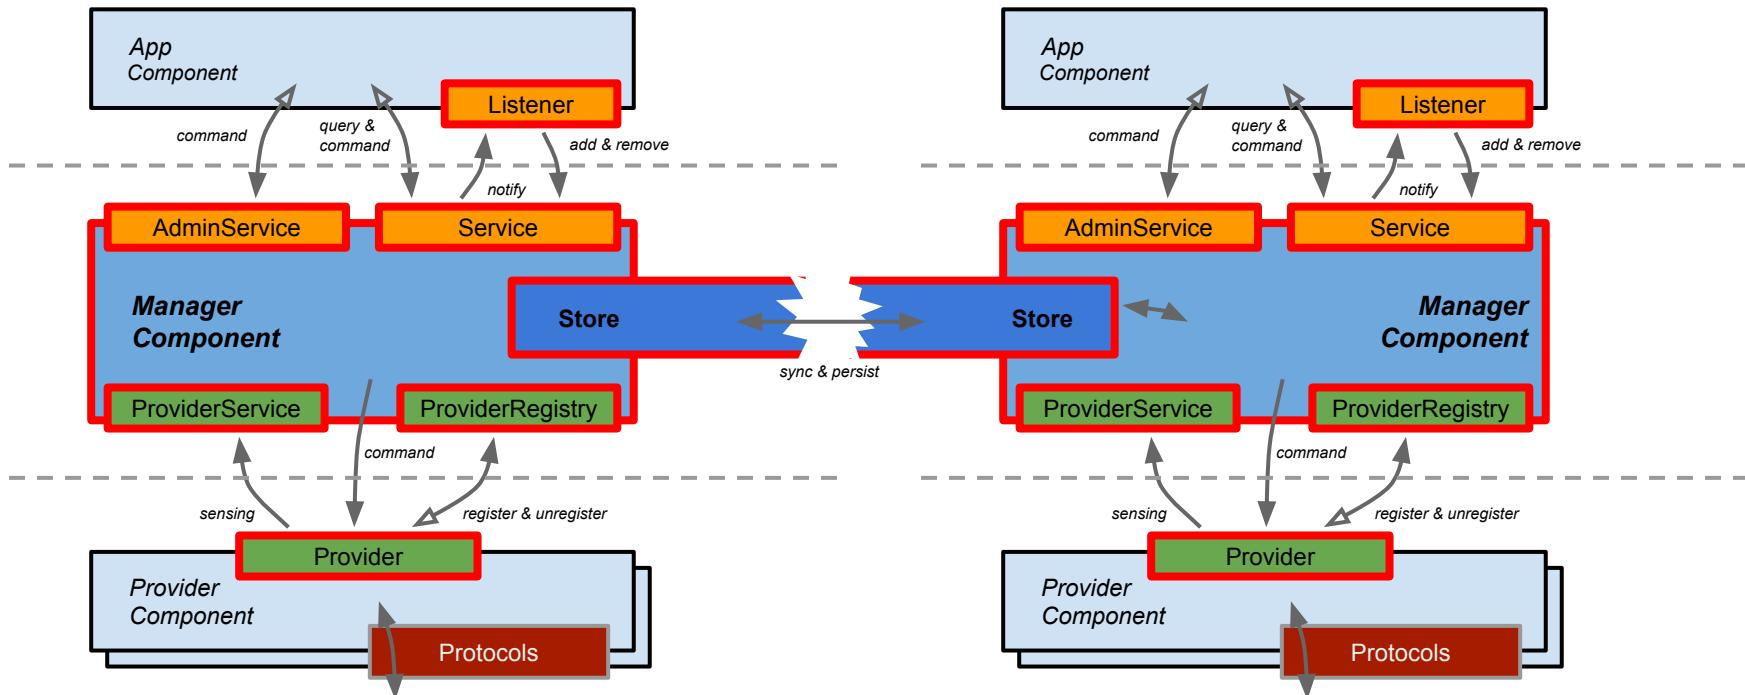
ONOS Distributed Architecture



ONOS Core Subsystems



ONOS Core Subsystem Structure



ONOS Fundamentals

8:30-10:00



- ONOS Overview (~30)
 - tenets, architecture, tiers, subsystems
 - general subsystem structure
- **Getting started with ONOS** (~30)
 - check-out code, environment customization and build
 - describe Maven, hierarchical POM structure
 - run ONOS locally on developer machine
 - CLI & GUI demo with a simple mininet topology (torus,8,8,2)
- Code organization (~30)
 - importing ONOS source into IDE (demo with IntelliJ)
 - core (api, net, stores), web (gui, rest), providers (openflow, netconf)

Development Pre-Requisites



- Bash (or similar shell)
- Oracle Java Development Kit 8+
<http://www.oracle.com/technetwork/java/javase/downloads>
- Apache Maven 3.3.x+
<https://maven.apache.org/download.html>
- Apache Karaf 3.0.3+ (for local development)
<http://karaf.apache.org/index/community/download.html>
- IDE of your choice (IntelliJ, Eclipse, ...)

Getting started with ONOS



- Check-out ONOS code via git
- Tailor bash environment for ONOS development
- Build ONOS via Maven
- Run ONOS locally on developer machine
- Activate pre-installed reactive forwarding application
- Test ONOS via a mininet simulated network
- Explore ONOS GUI & CLI

Check-out ONOS code



- Open a new terminal shell and clone ONOS repo:

```
$ git clone https://gerrit.onosproject.org/onos
```

Tailor shell for ONOS development



- Edit the shell profile:

```
$ vi ~/.bash_profile
```

- Add the following lines at the top:

```
export ONOS_ROOT=~/onos
```

```
source $ONOS_ROOT/tools/dev/bash_profile
```

- Source the modified shell profile

```
$ source ~/.bash_profile
```

Build ONOS via Maven



- Change directory to your ONOS root and invoke Maven

```
$ cd ~/onos
```

```
$ mvn clean install
```

- Alternatively, use **onos-build** command or **ob** alias from anywhere

```
$ onos-build
```

Run ONOS locally



- Use `onos-karaf` command or `ok` alias to setup and run ONOS via Apache Karaf

```
$ onos-karaf clean
```

- Use ONOS shell commands to explore

```
onos> summary
```

```
onos> nodes
```

```
onos> apps -s
```

Activate reactive forwarding app



- Use app command to activate a pre-installed app:

```
onos> app activate org.onosproject.fwd
```

Test ONOS via mininet



- Start ONOS gui via `onos-gui` command or `gui` alias:

```
$ onos-gui
```

- Start 8x8 torus topology with 2 hosts per switch:

```
$ sudo mn --topo torus,8,8,2 \
  --controller remote,ip=192.168.56.1
```

- Ping all hosts in mininet

```
mininet> pingall
```

 ONOS

localhost:8181/onos/ui/index.html#/topo

 Open Network Operating System

ONOS Summary

| | |
|-----------------|-----|
| Devices : | 64 |
| Links : | 256 |
| Hosts : | 512 |
| Topology SCCs : | 1 |

Intents : 0
Tunnels : 0
Flows : 320
Version : 1.3.0*

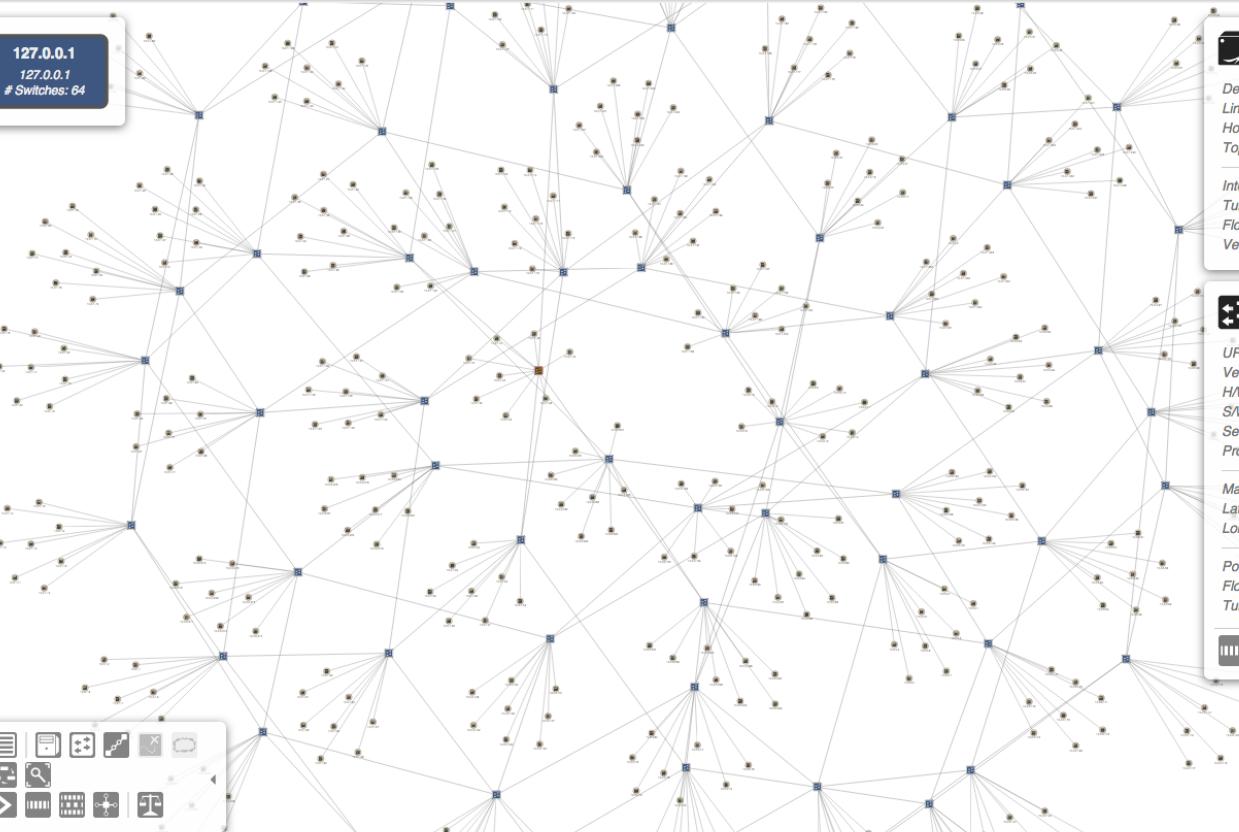
of:0000000000000506

| | |
|-----------------|---------------------|
| URI : | of:0000000000000506 |
| Vendor : | Nicira, Inc. |
| H/W Version : | Open vSwitch |
| S/W Version : | 2.3.1 |
| Serial Number : | None |
| Protocol : | OF_1.3 |

Master : 127.0.0.1
Latitude :
Longitude :

Ports : 13
Flows : 5
Tunnels : 0





Explore ONOS GUI & CLI



- Enable host display via tool icon or **H** key
- Explore Devices, Links and Hosts views, etc.
- Explore ONOS CLI via the following shell commands:

onos> summary

onos> devices

onos> links

onos> hosts

onos> flows

onos> onos:*

ONOS Fundamentals

8:30-10:00



- ONOS Overview (~30)
 - tenets, architecture, tiers, subsystems
 - general subsystem structure
- Getting started with ONOS (~30)
 - check-out code, environment customization and build
 - describe Maven, hierarchical POM structure
 - run ONOS locally on developer machine
 - CLI & GUI demo with a simple mininet topology (torus,8,8,2)
- Code organization (~30)
 - importing ONOS source into IDE (demo with IntelliJ)
 - core (api, net, stores), web (gui, rest), providers (openflow, netconf)

Tour of ONOS code-base



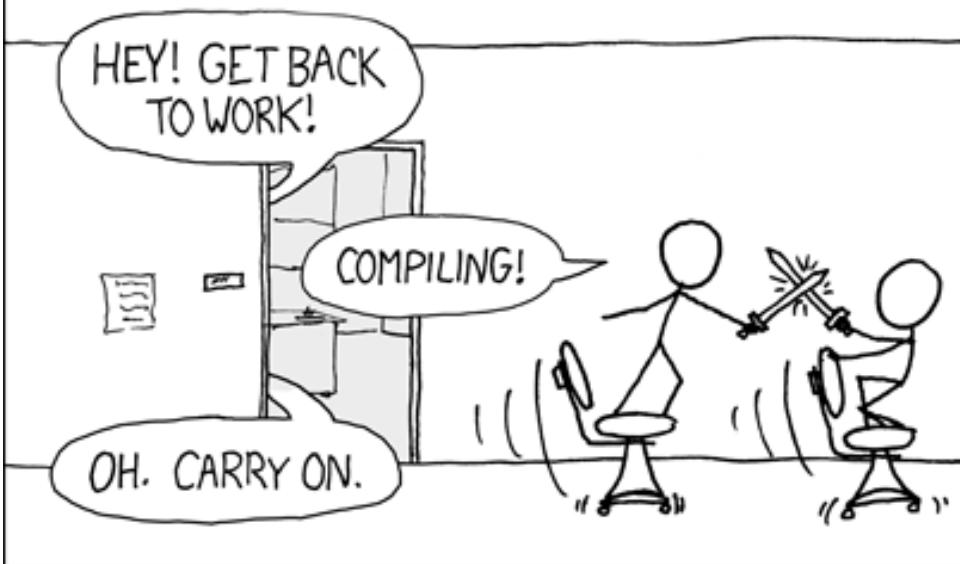
- Import ONOS into IDE
- Tour of ONOS core (api, net, stores, etc.)
- Tour of other areas (apps, providers, web, etc.)

Break



THE #1 PROGRAMMER EXCUSE
FOR LEGITIMATELY SLACKING OFF:

"MY CODE'S COMPILING."



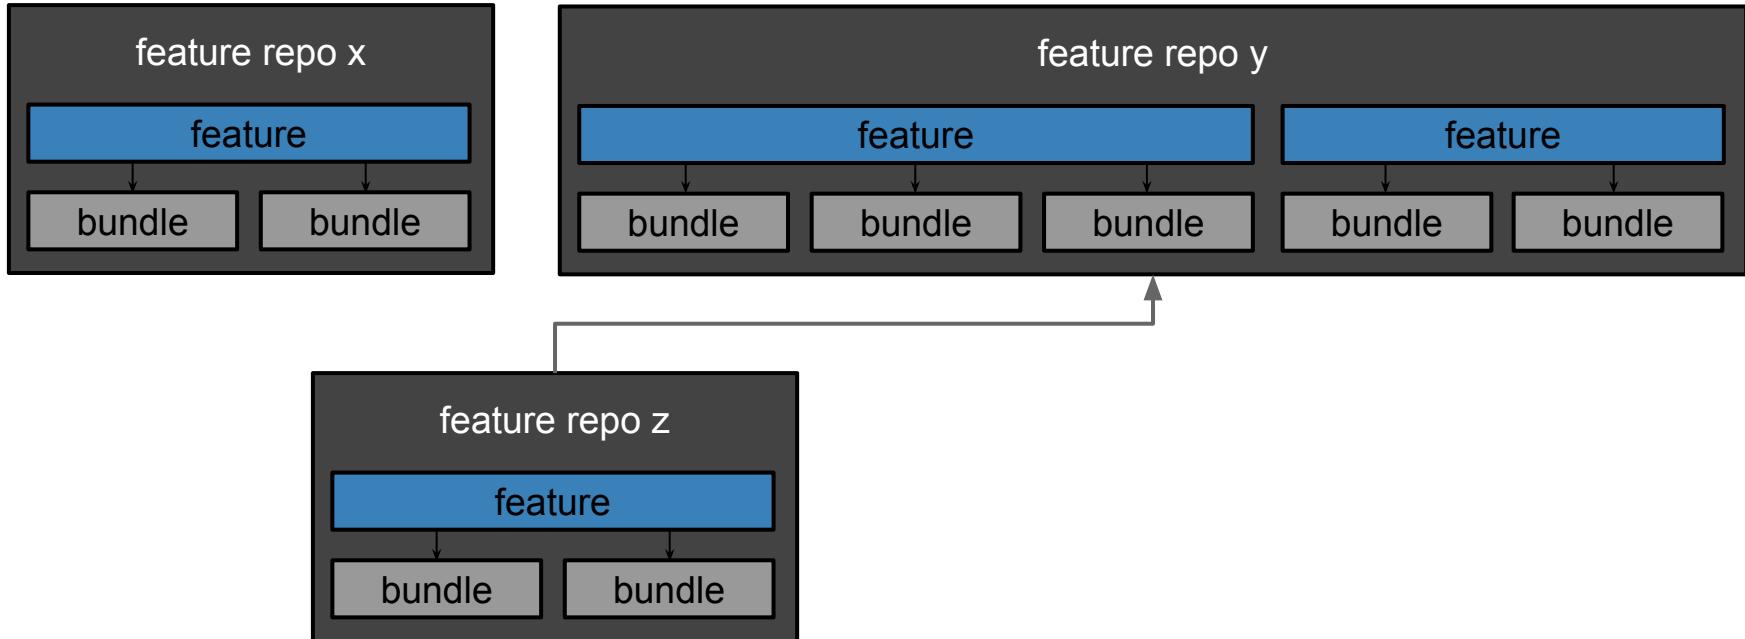
ONOS Distributed Core

10:30-12:00



- **ONOS, OSGi & Apache Karaf** (~15)
 - OSGi bundles, Karaf features & SCR components
 - use `bundle:*`, `feature:*`, `scr:*` commands & webconsole
- Anatomy of an ONOS Core Subsystem (~15)
 - use device subsystem as an example
- Distributed stores & primitives (~60)
 - physically distributed & logically centralized
 - deep dive on distributed topology stores (devices, links, hosts)
 - other types of state treatment for other types of state
 - primitives, e.g. eventually consistent/strongly consistent map, leadership, distributed set, distributed queue, atomic counter

OSGi Bundles & Karaf Features



OSGi Bundles & Karaf Features



- OSGi bundles are Java JAR files with an enhanced Manifest
 - bundles have name and version
 - bundles explicitly require/import other Java packages
 - bundles explicitly provide/export Java packages for others
- Karaf features are means to install or uninstall a set of bundles as a group
 - features are defined via an XML artifact - a feature repository
 - features reference, but do not deliver the bundle JAR artifacts
- Karaf uses Maven repos as OSGi Bundle Repositories for retrieval of feature and bundle artifacts

Service Component Runtime



- Components are effectively stateful singletons whose life-cycle is controlled by the framework
 - components defined by `OSGI-INF/*.xml` files at run-time
 - ONOS uses `maven-scr-plugin` to convert Java annotations to `OSGI-INF/*.xml` files at compile-time
- Components can provide `@Services` to others
- Components can `@Reference` services from others
- `@Activate`, `@Modified` and `@Deactivate` methods serve as component life-cycle hooks

Bundle & Feature Shell Commands



- Bundle related commands

onos> bundle:*

- Feature related commands

onos> feature:*

- Service Component Runtime related commands

onos> scr:*

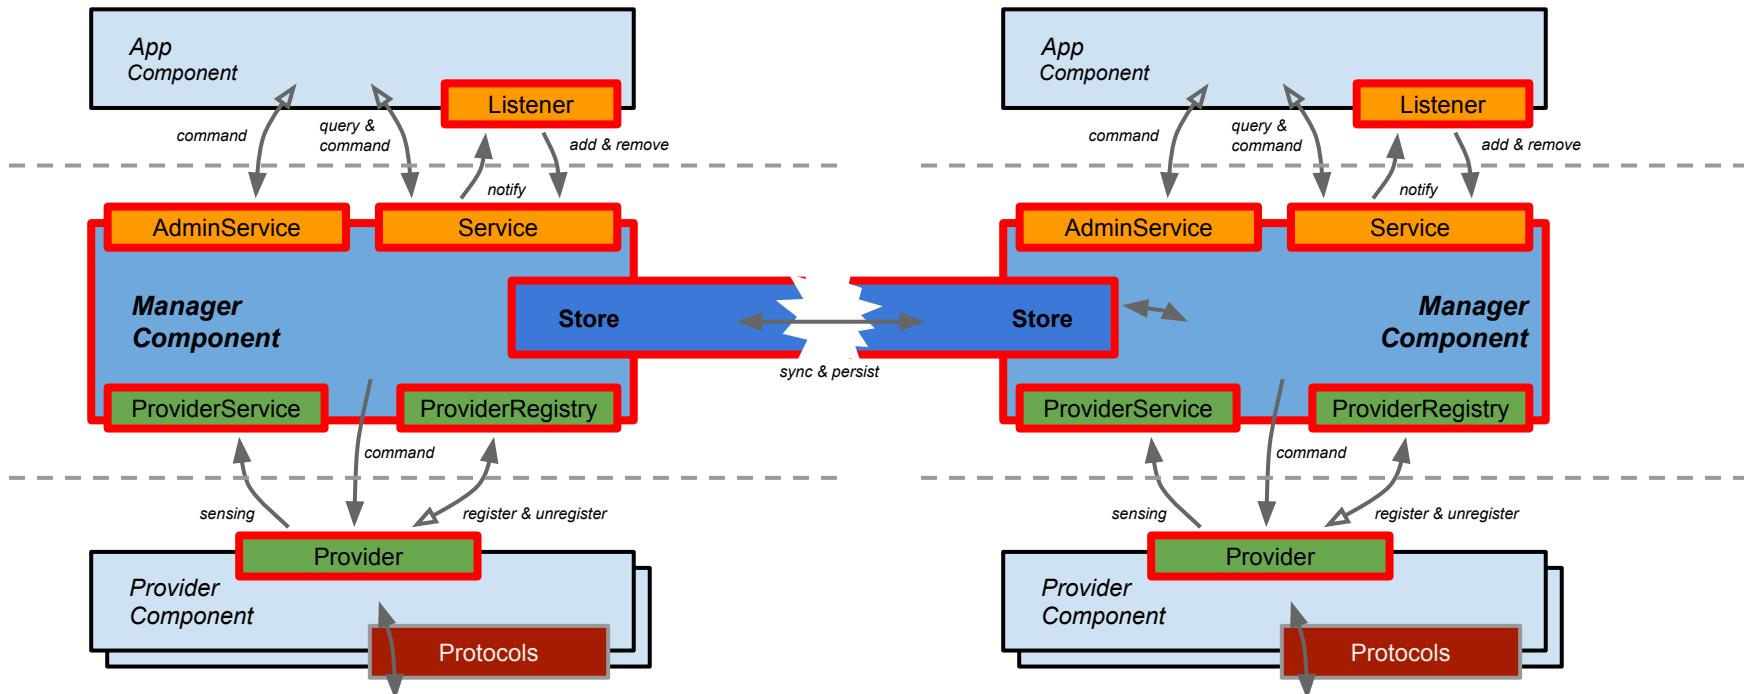
ONOS Distributed Core

10:30-12:00



- ONOS, OSGi & Apache Karaf (~15)
 - OSGi bundles, Karaf features & SCR components
 - use `bundle:*`, `feature:*`, `scr:*` commands & webconsole
- Anatomy of an ONOS Core Subsystem (~15)
 - use device subsystem as an example
- Distributed stores & primitives (~60)
 - physically distributed & logically centralized
 - deep dive on distributed topology stores (devices, links, hosts)
 - other types of state treatment for other types of state
 - primitives, e.g. eventually consistent/strongly consistent map, leadership, distributed set, distributed queue, atomic counter

Anatomy of an ONOS subsystem



Device Subsystem - Northbound



- `DeviceManager` component exposes `DeviceService` to apps and other services
 - note the `@Component` and `@Service` annotations
- Apps can subscribe to asynchronous `DeviceEvent` notifications via `DeviceListener` mechanism
- `DeviceManager` also exposes `DeviceAdminService` to privileged apps for select administrative actions

Device Subsystem - Southbound



- DeviceManager offers DeviceProviderRegistry service
- DeviceProvider entities can register and acquire a reference to DeviceProviderService
- Core uses registered DeviceProvider for sending edicts and queries down to the environment in a protocol-agnostic manner
- Provider uses issued DeviceProviderService for sending device and port sensory data up to the core in a protocol-agnostic manner

Device Subsystem - East / West



- **DeviceManager** delegates to peer **DeviceStore** all cluster-related tasks
 - manager becomes **DeviceStoreDelegate** to the **DeviceStore**
- **DeviceStore** is responsible for indexing, persisting and synchronizing the global device and port state with all cluster peers
- **DeviceStore** delegates incoming **DeviceEvents** to its **DeviceStoreDelegate** for local actions and for local event dispatching

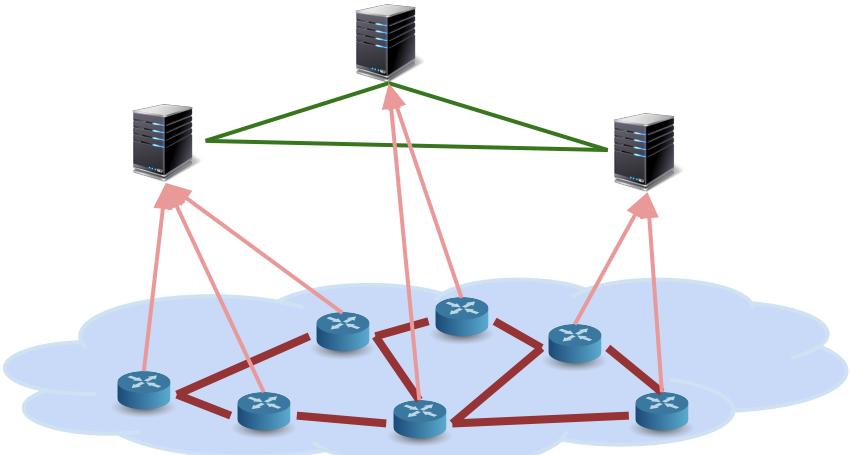
ONOS Distributed Core

10:30-12:00



- ONOS, OSGi & Apache Karaf (~15)
 - OSGi bundles, Karaf features & SCR components
 - use `bundle:*`, `feature:*`, `scr:*` commands & webconsole
- Anatomy of an ONOS Core Subsystem (~30)
 - use device subsystem as an example
- **Distributed stores & primitives** (~60)
 - physically distributed & logically centralized
 - deep dive on distributed topology stores (devices, links, hosts)
 - other types of state treatment for other types of state
 - primitives, e.g. eventually consistent/strongly consistent map, leadership, distributed set, distributed queue, atomic counter

Distributed Stores & Primitives



“The unavoidable price of reliability is simplicity.”

— C.A.R. Hoare

“Simplicity is prerequisite for reliability.”

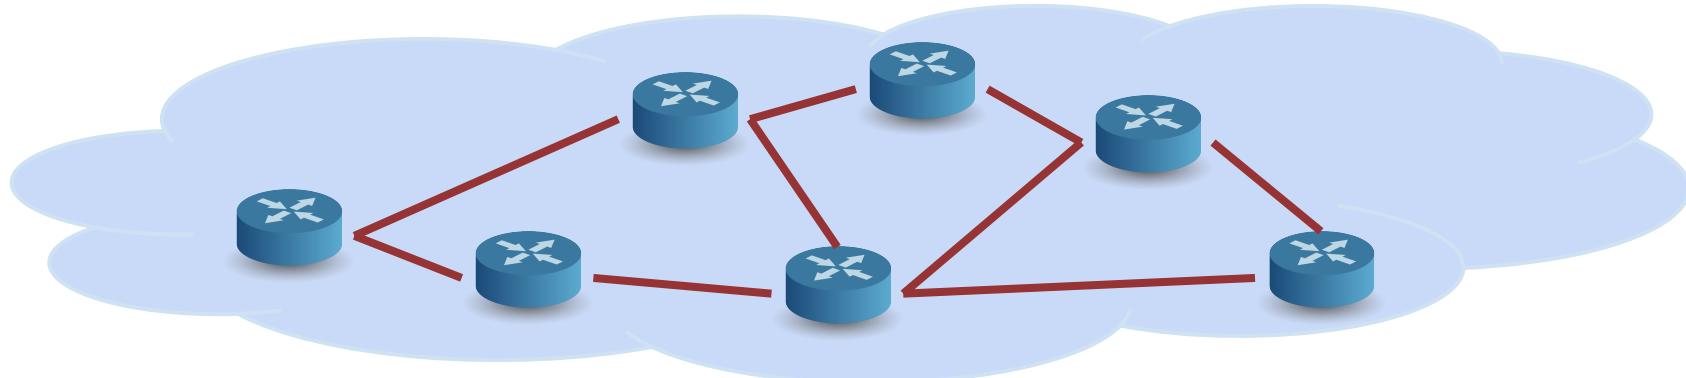
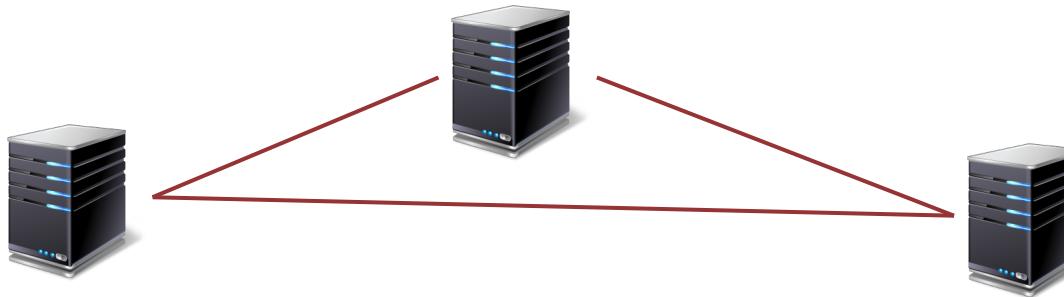
— E. W. Dijkstra

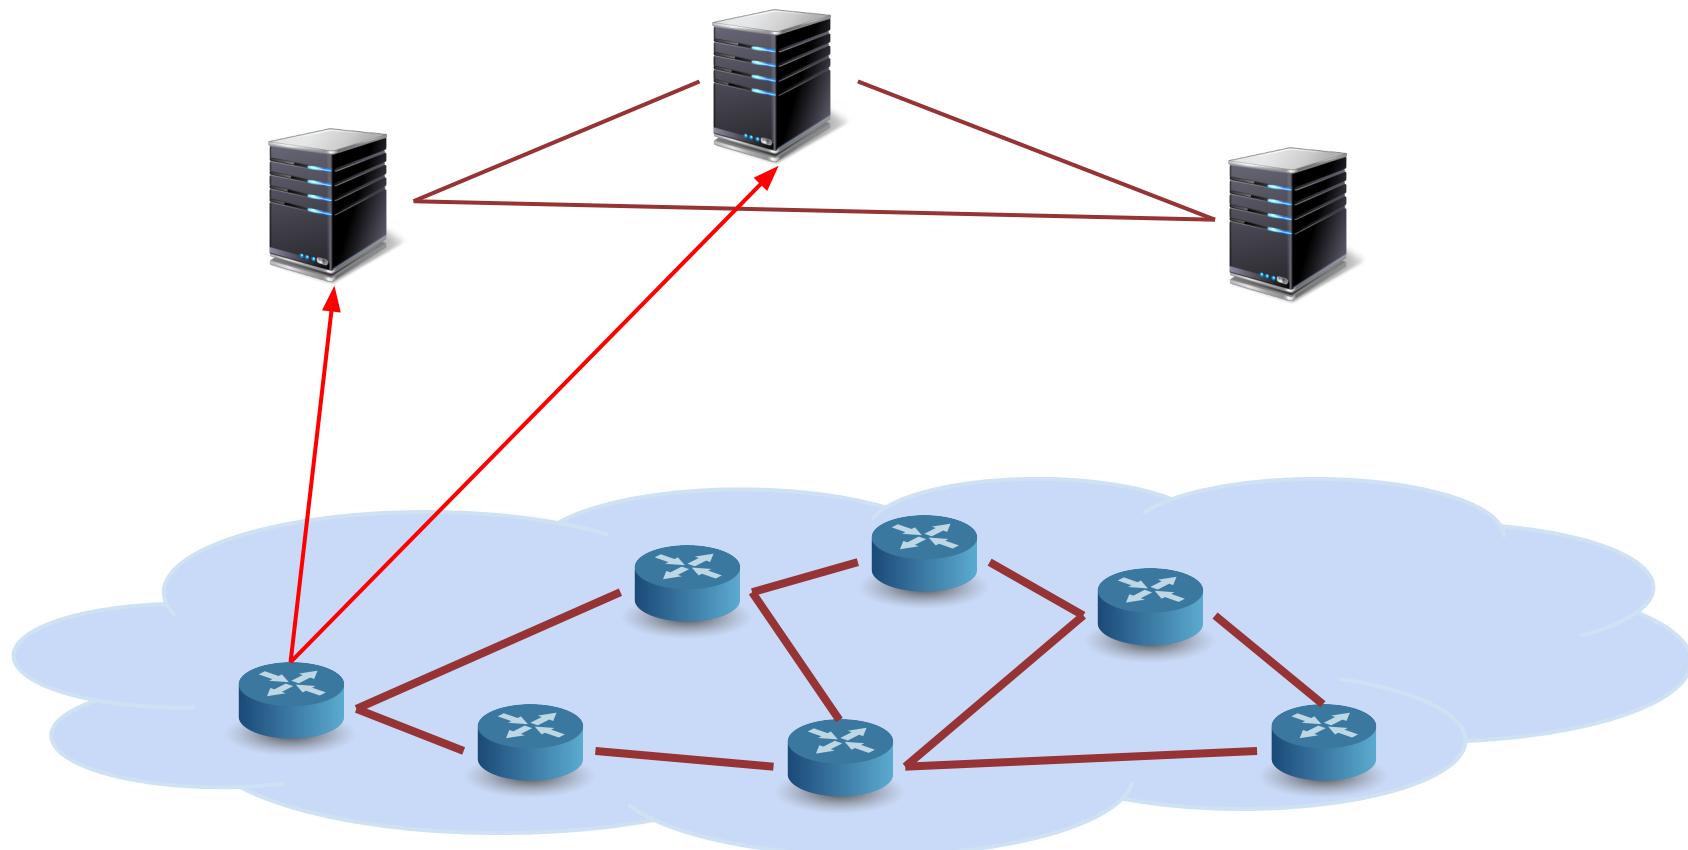
Distributed State Management

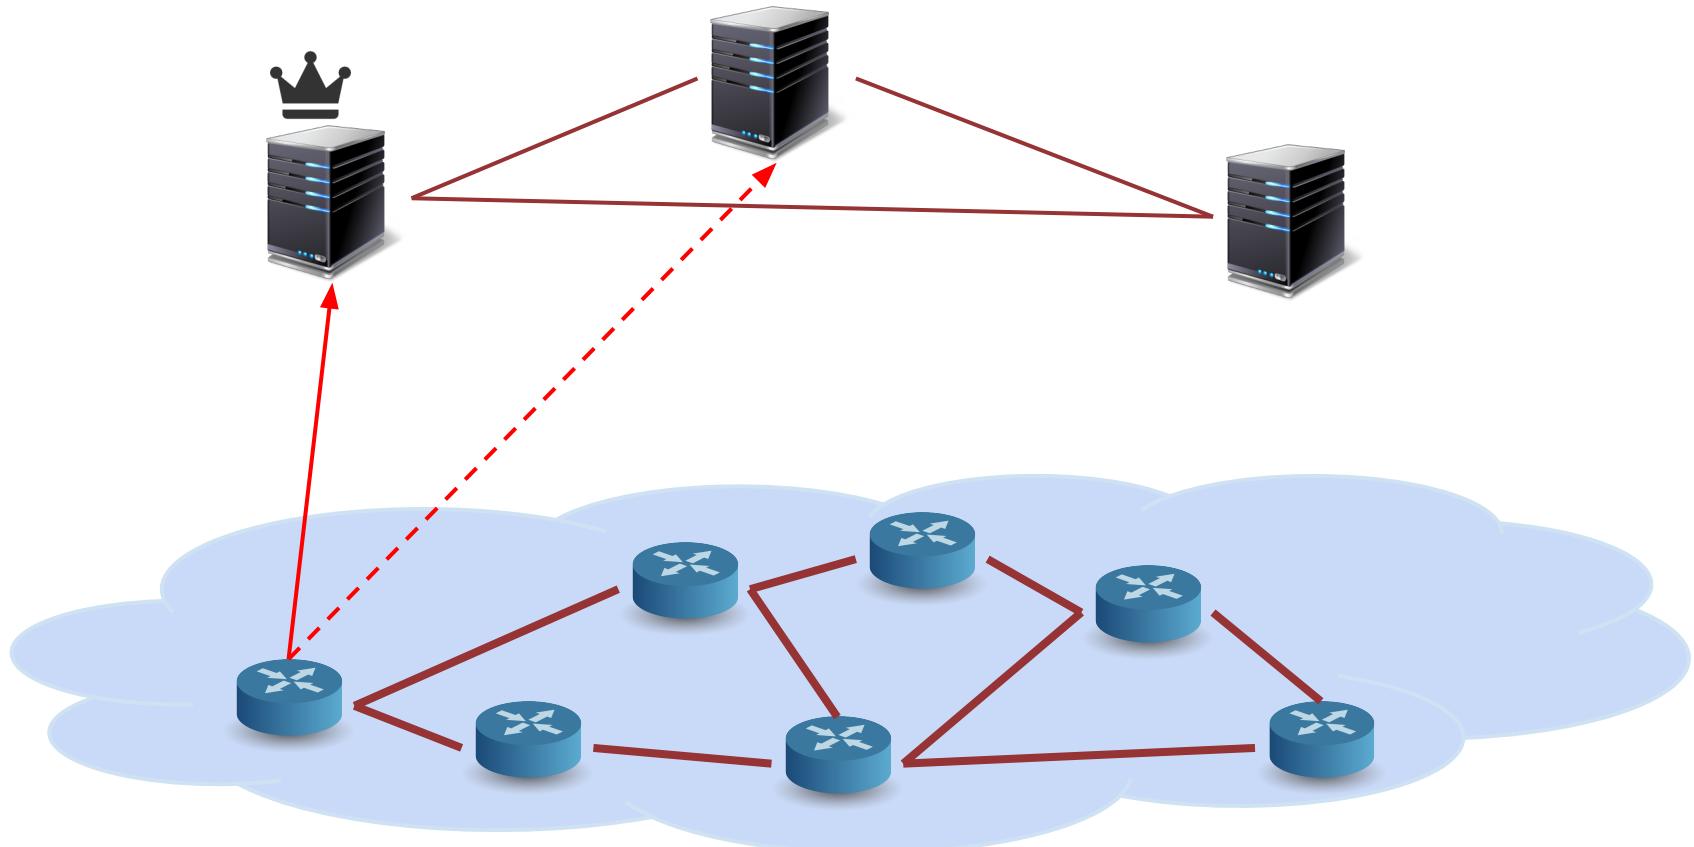


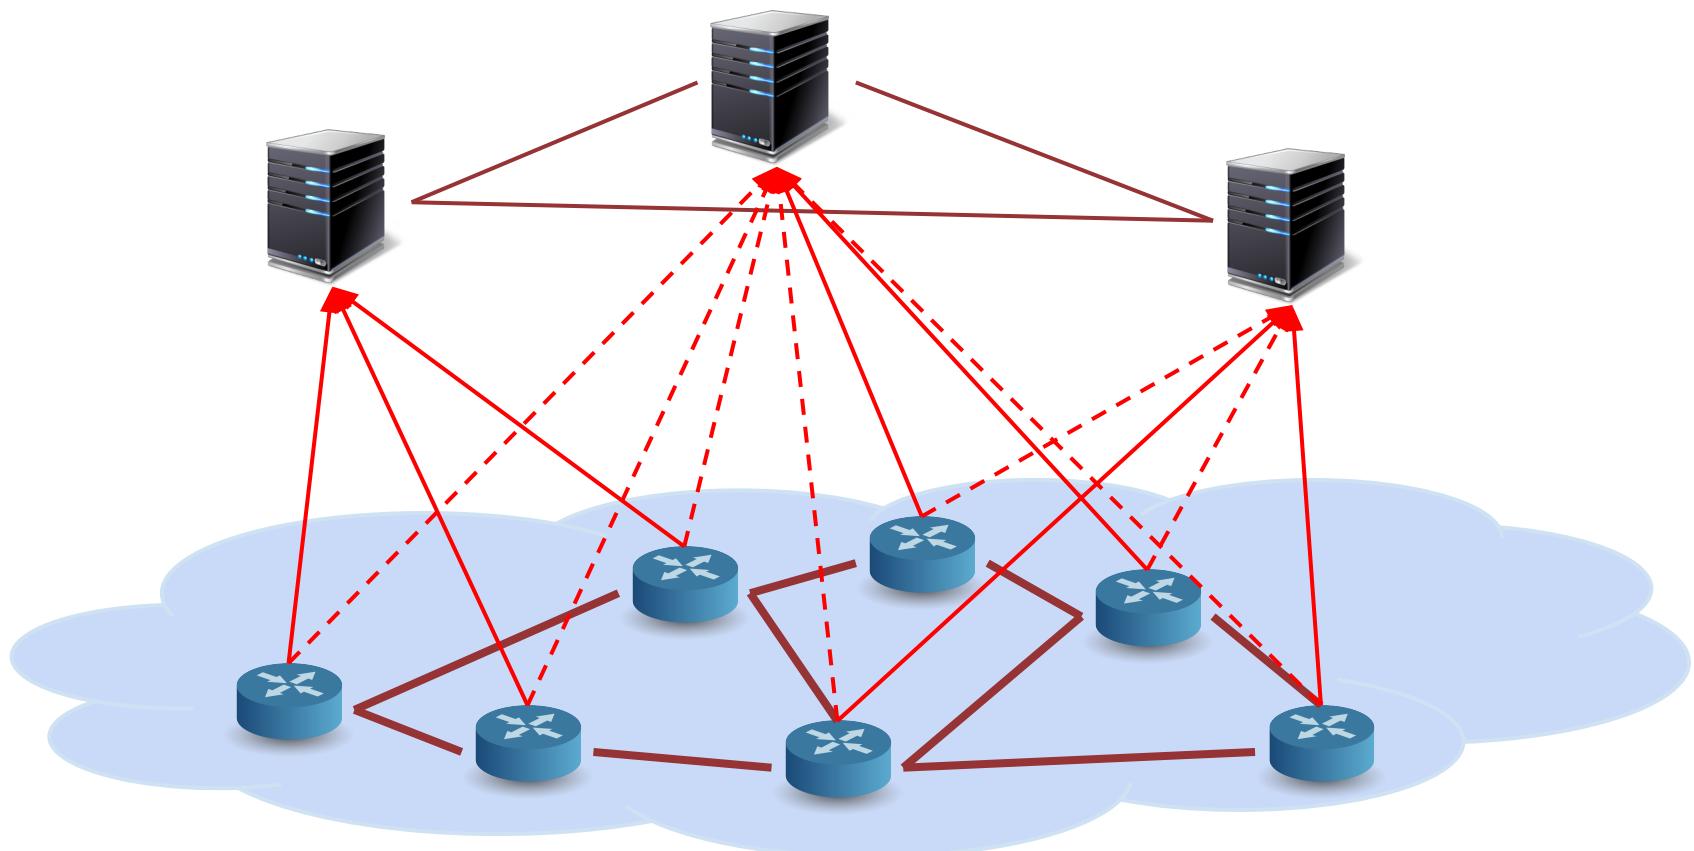
- Using specific examples - introduce various distributed state management primitives ONOS provides

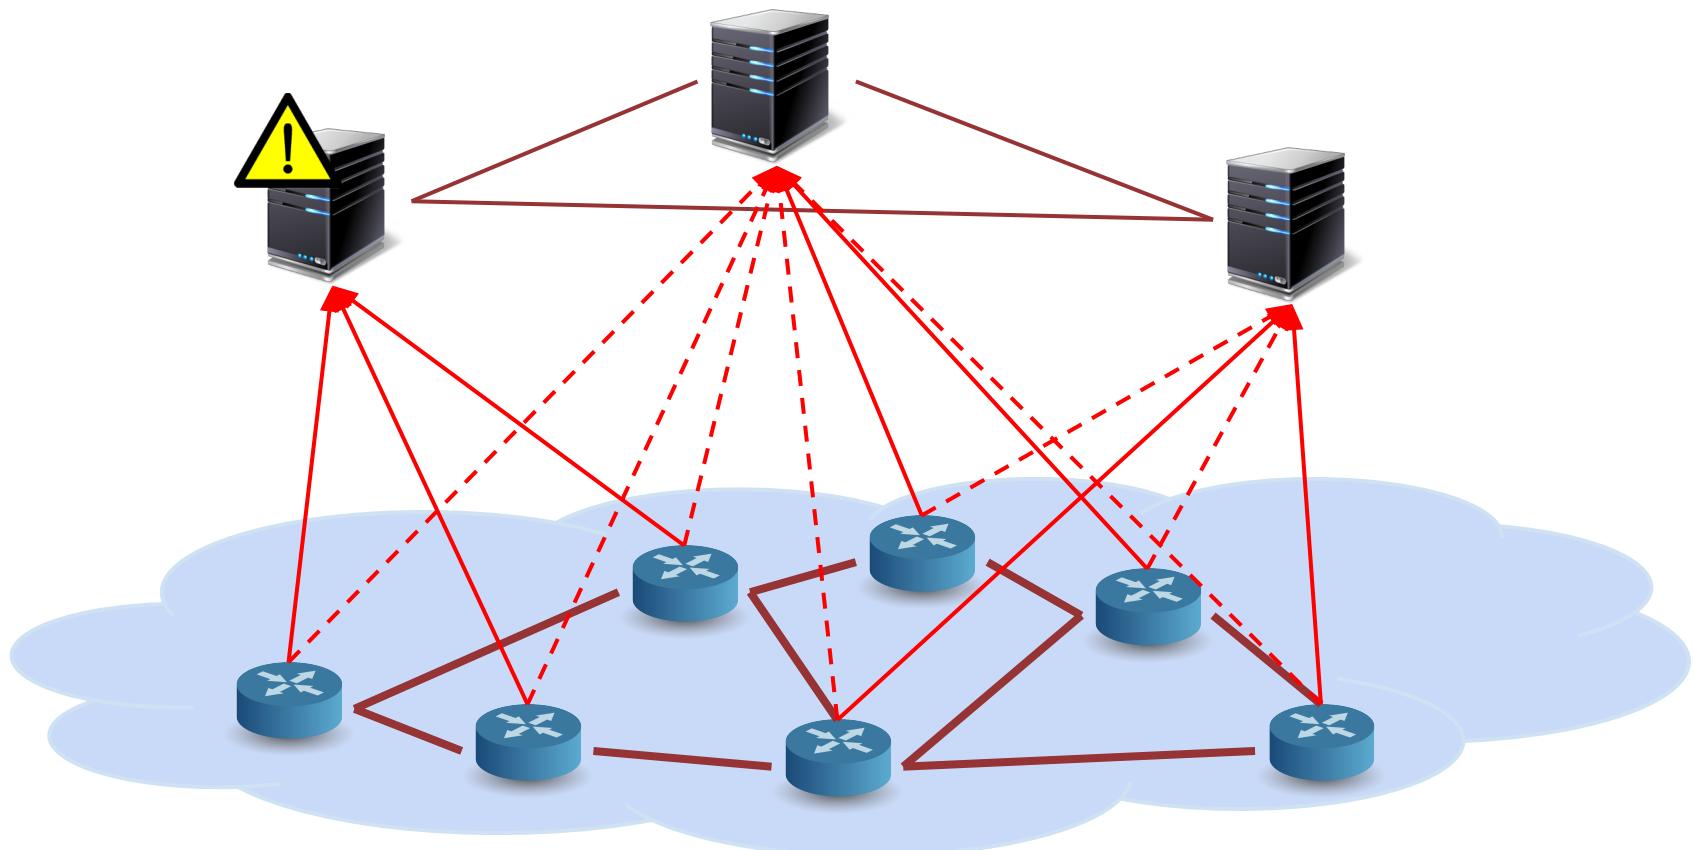
Distributed Control Plane

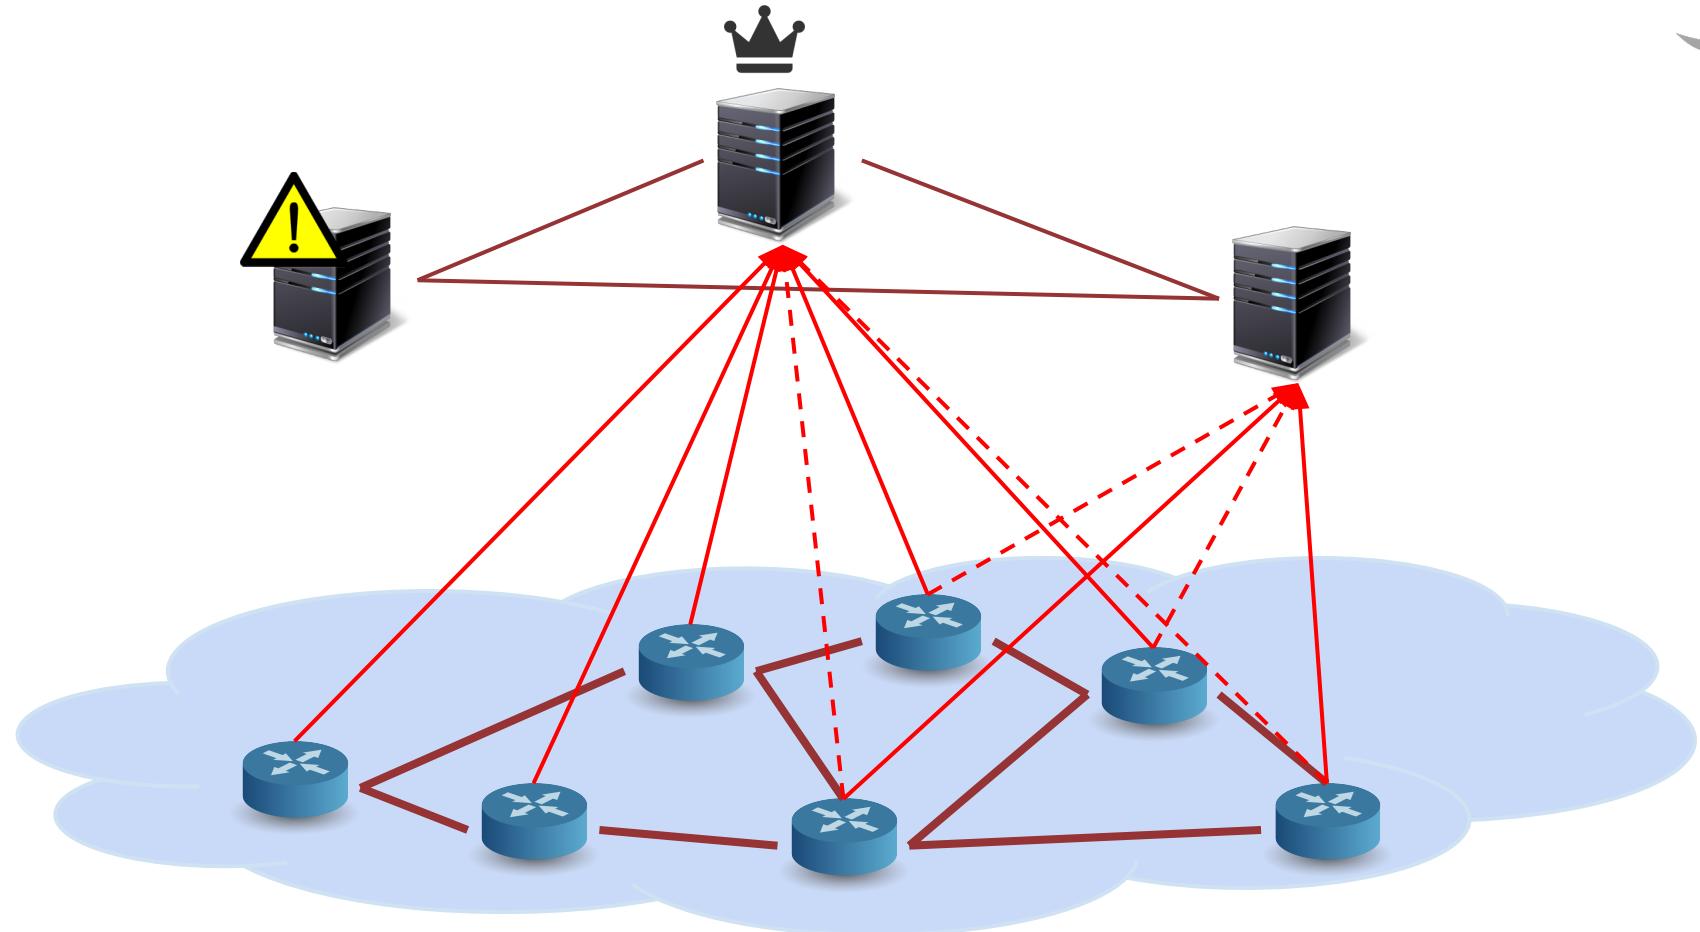












Distributed Primitive: Leadership

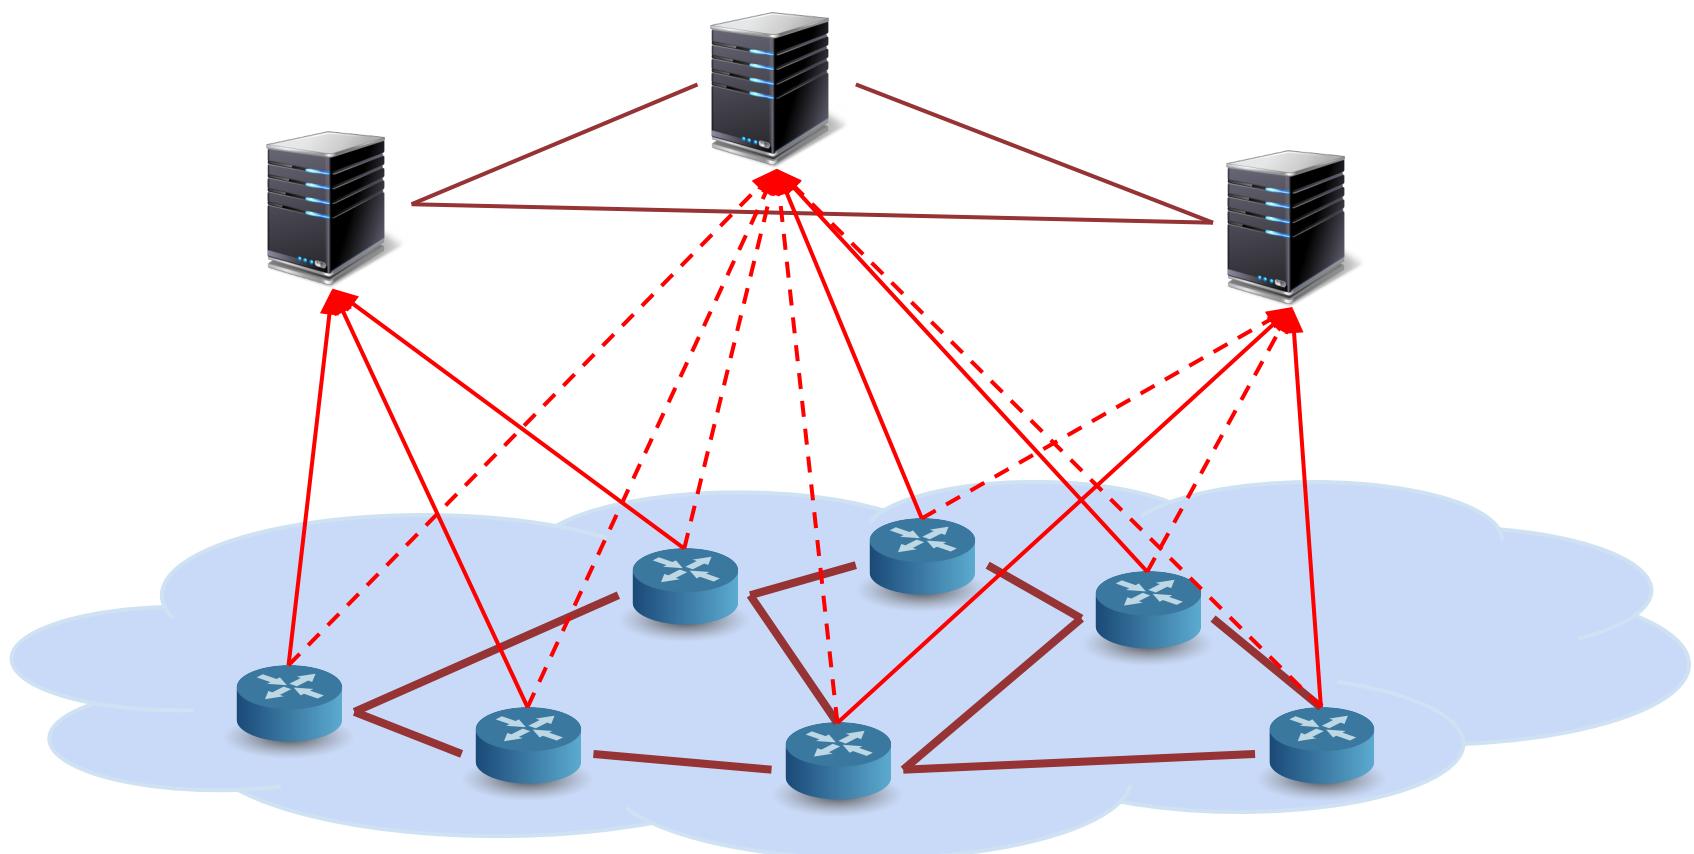


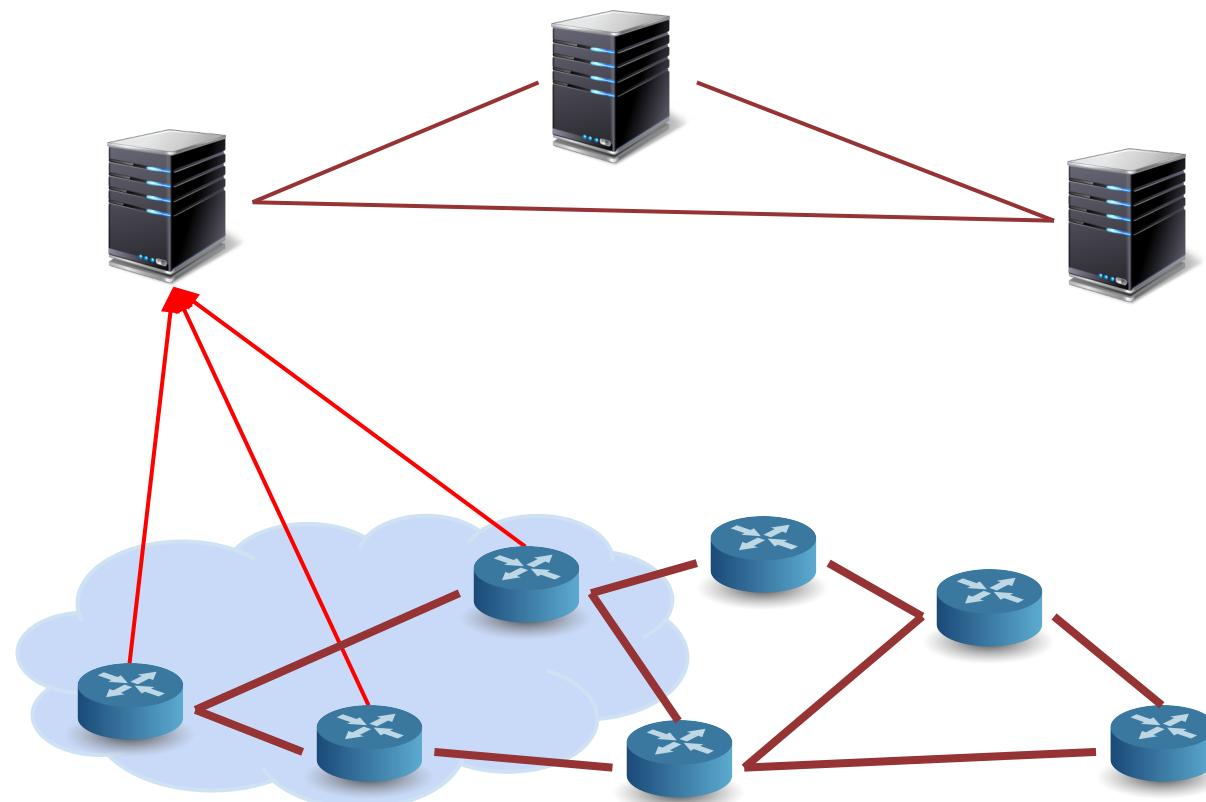
- Support for **mutual exclusion** in a distributed setting
- Available as a primitive for applications to consume
- Implemented via the **Raft consensus** protocol
- Cluster **majority** is needed to reach consensus

Deep dive into Topology state



- Formulate the problem
- Evaluate solution space
- Description of ONOS' approach
- Performance evaluation





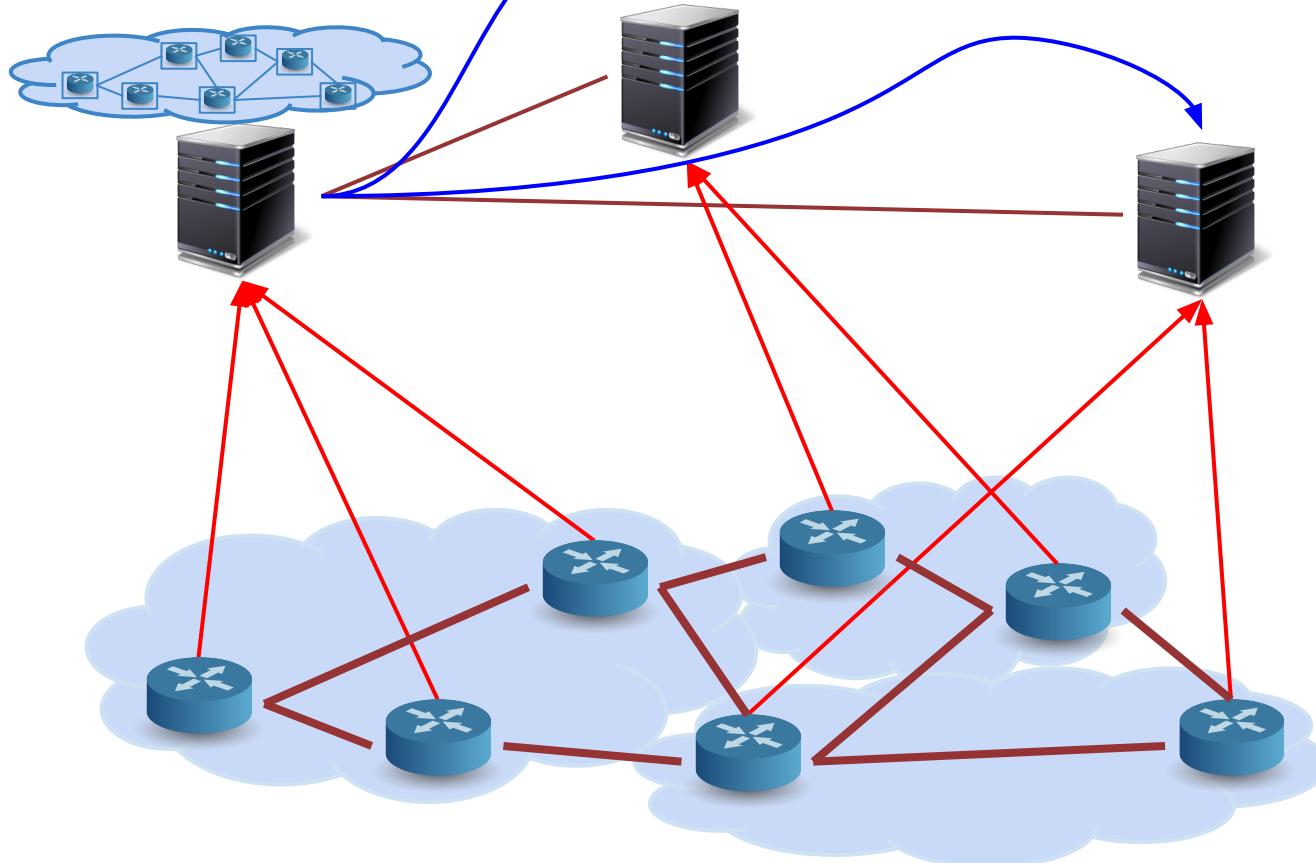
Solution space



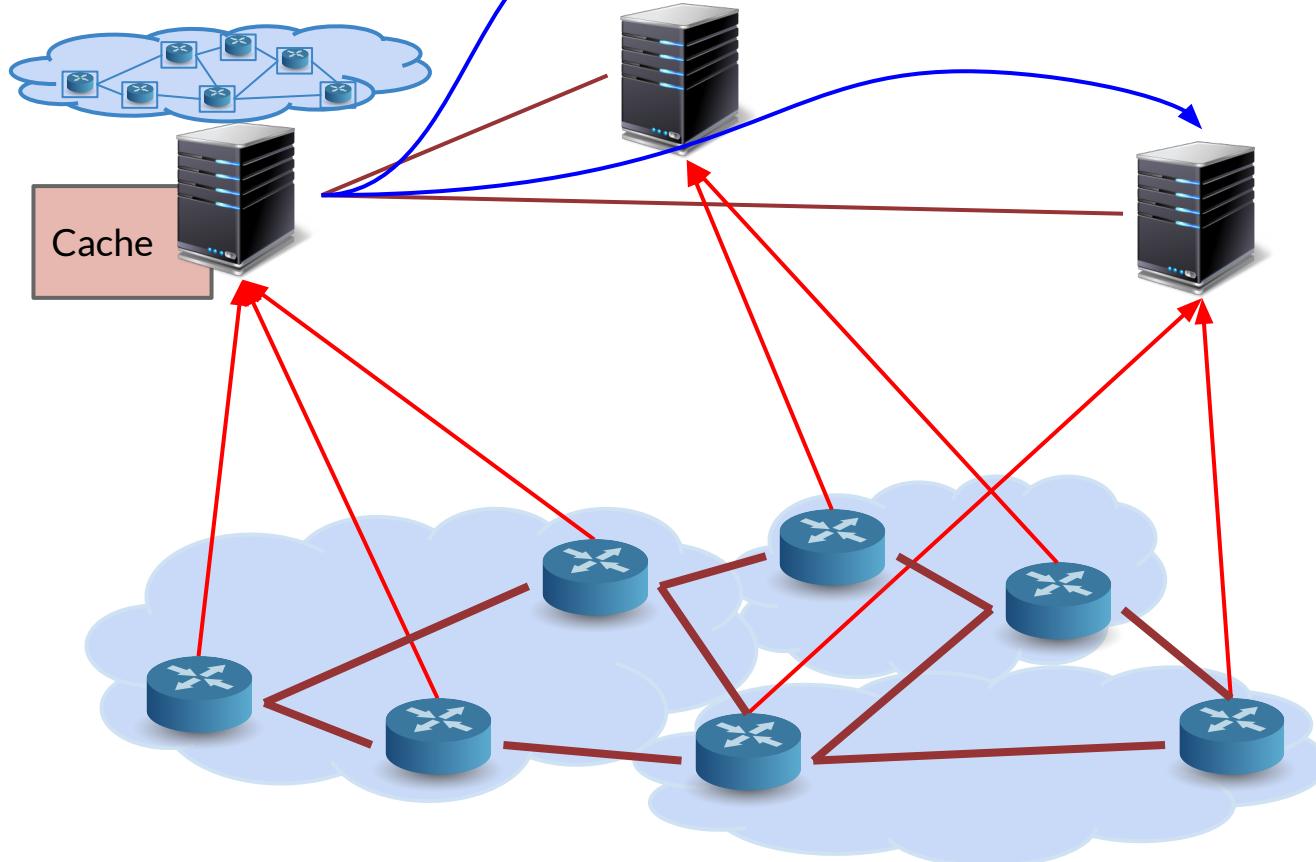
- Partitioned by mastership
- Logically centralized data store
- ONOS approach

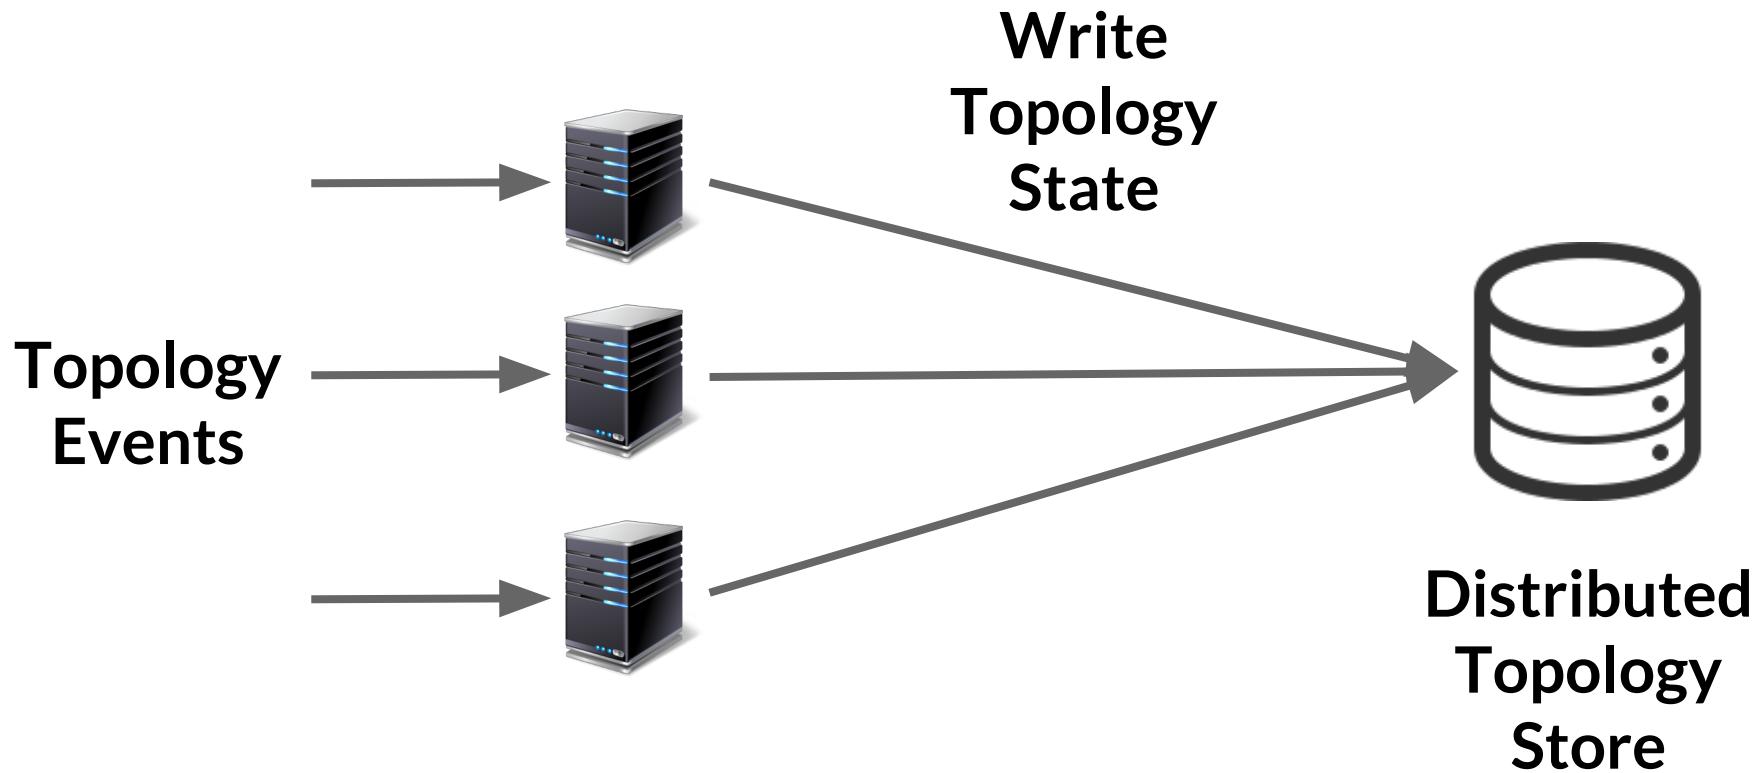


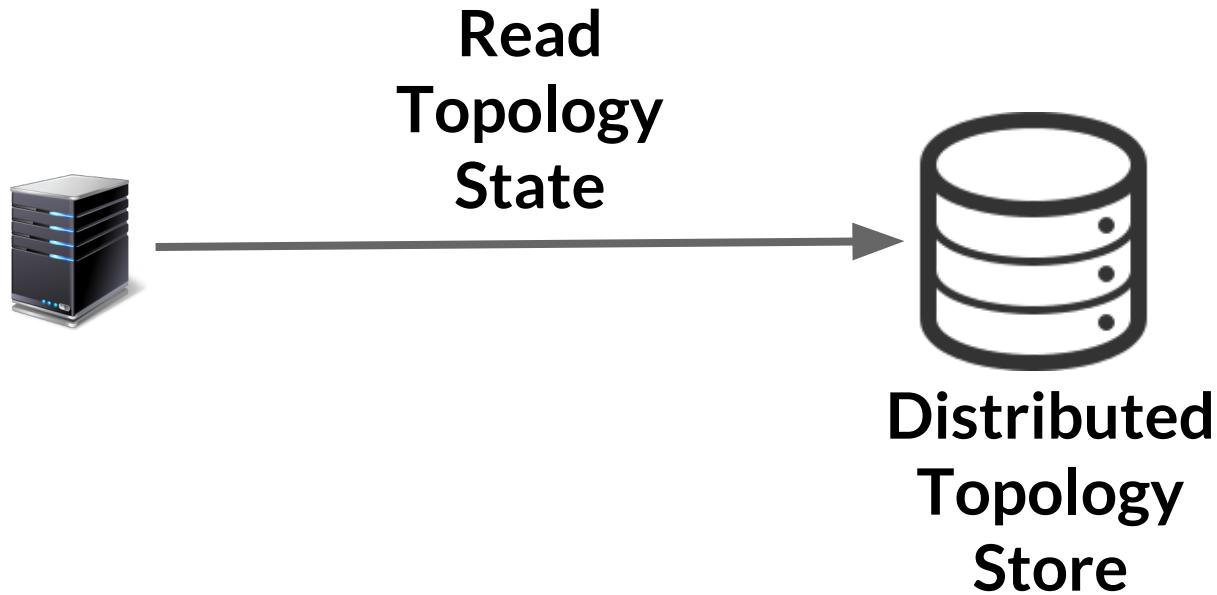
“Tell me about your slice?”

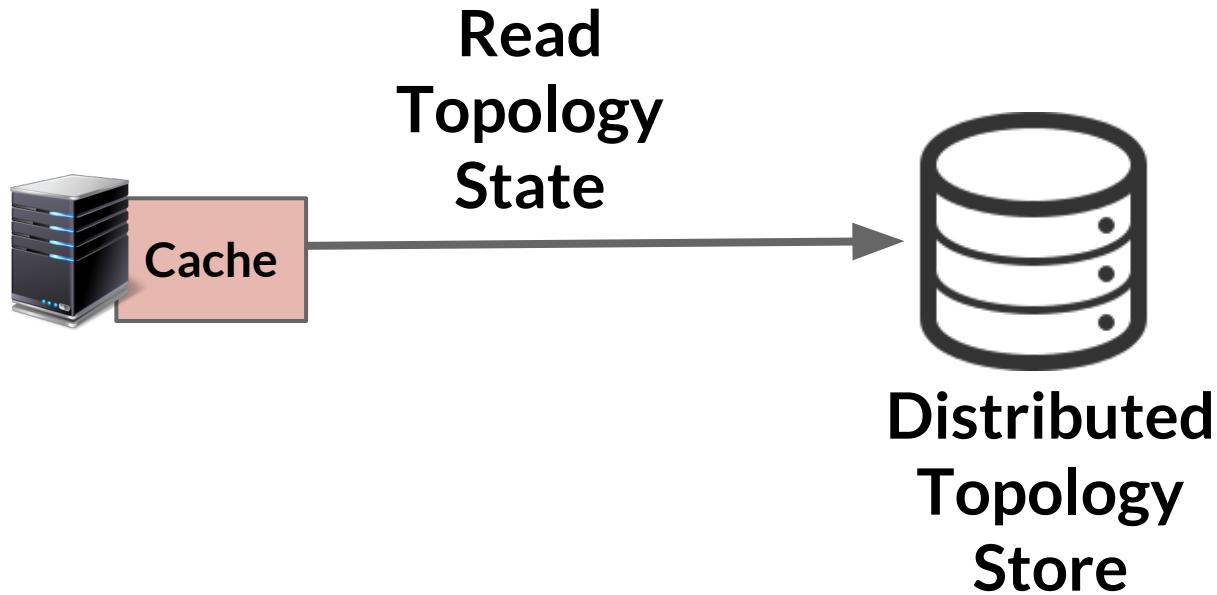


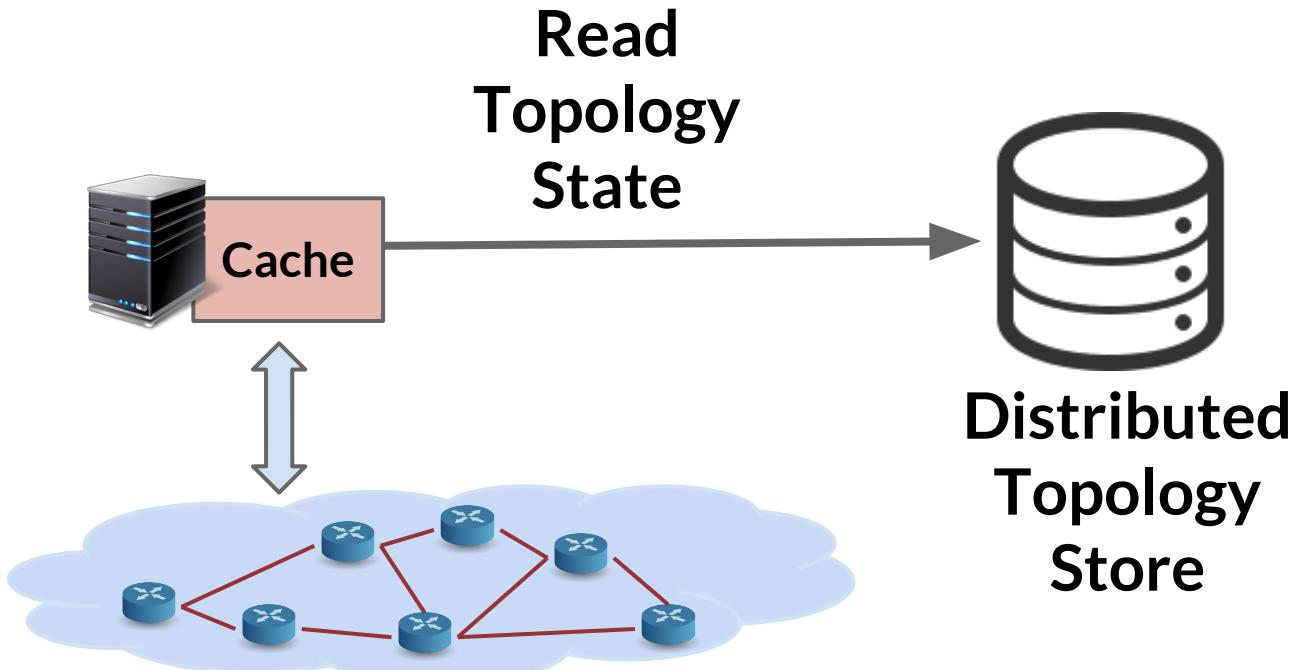
“Tell me about your slice?”







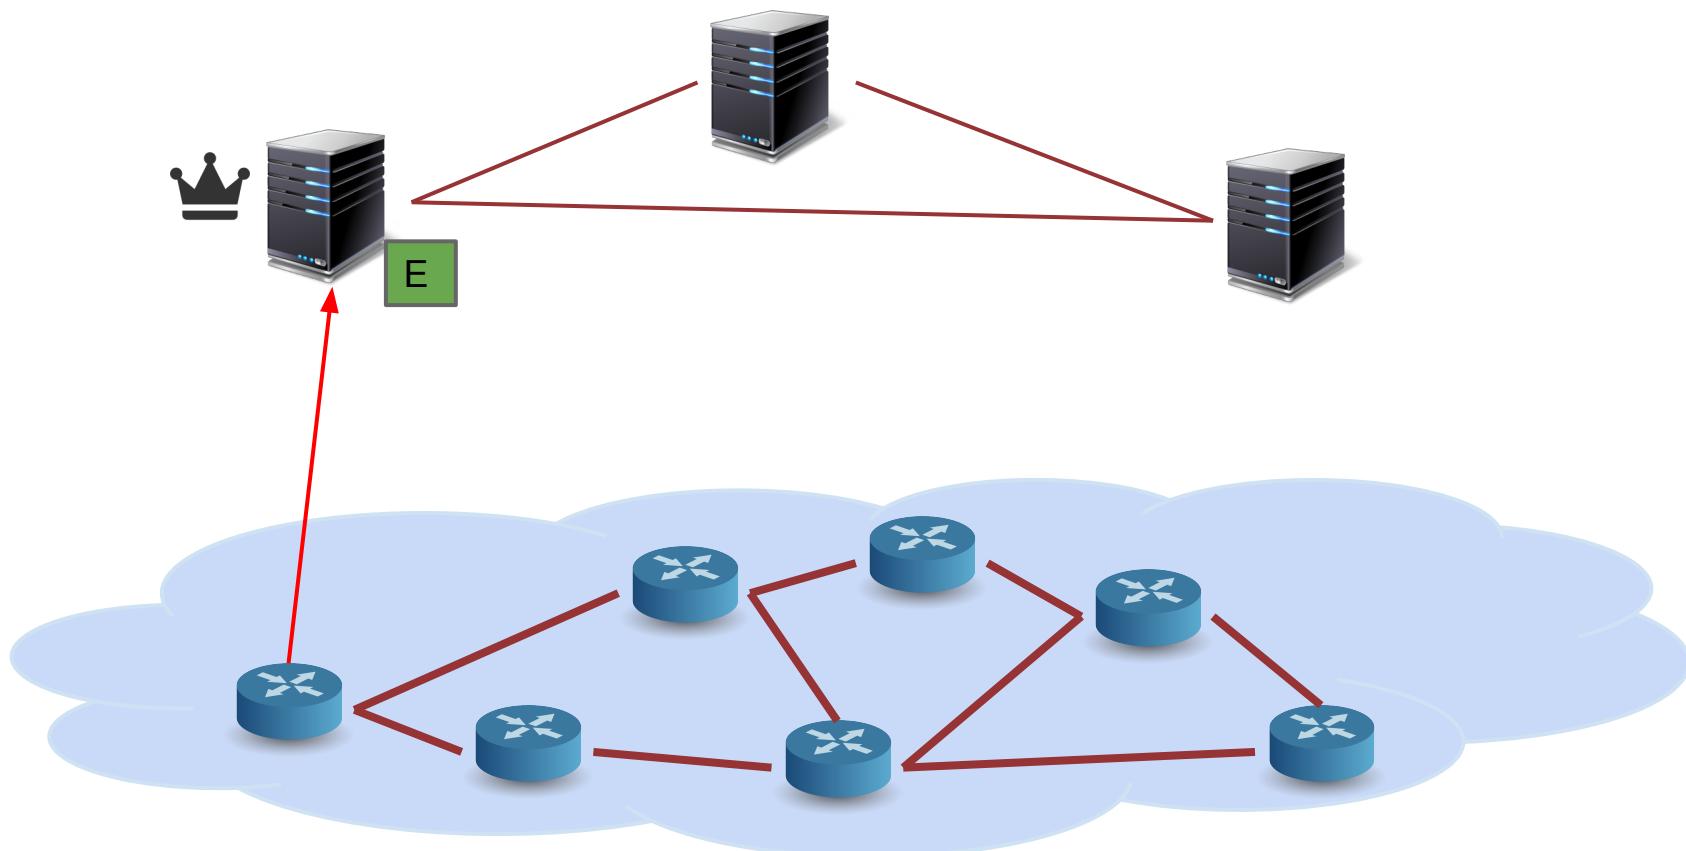


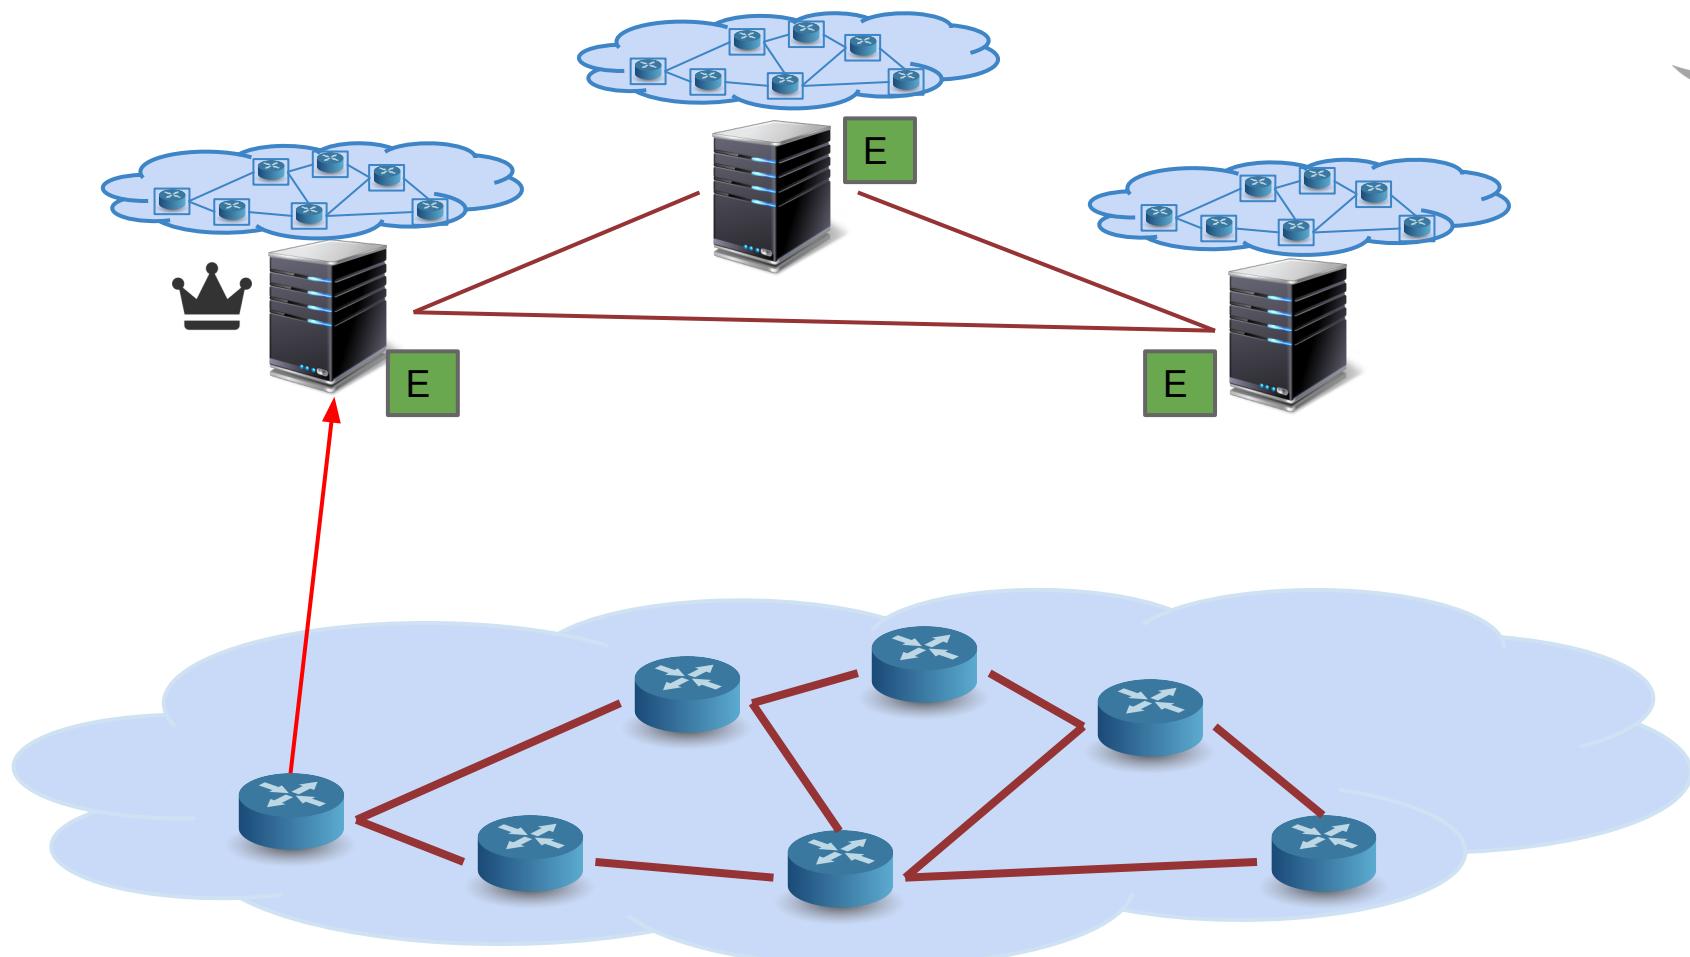


Topology as a State Machine



Events are Switch/Port/Link up/down

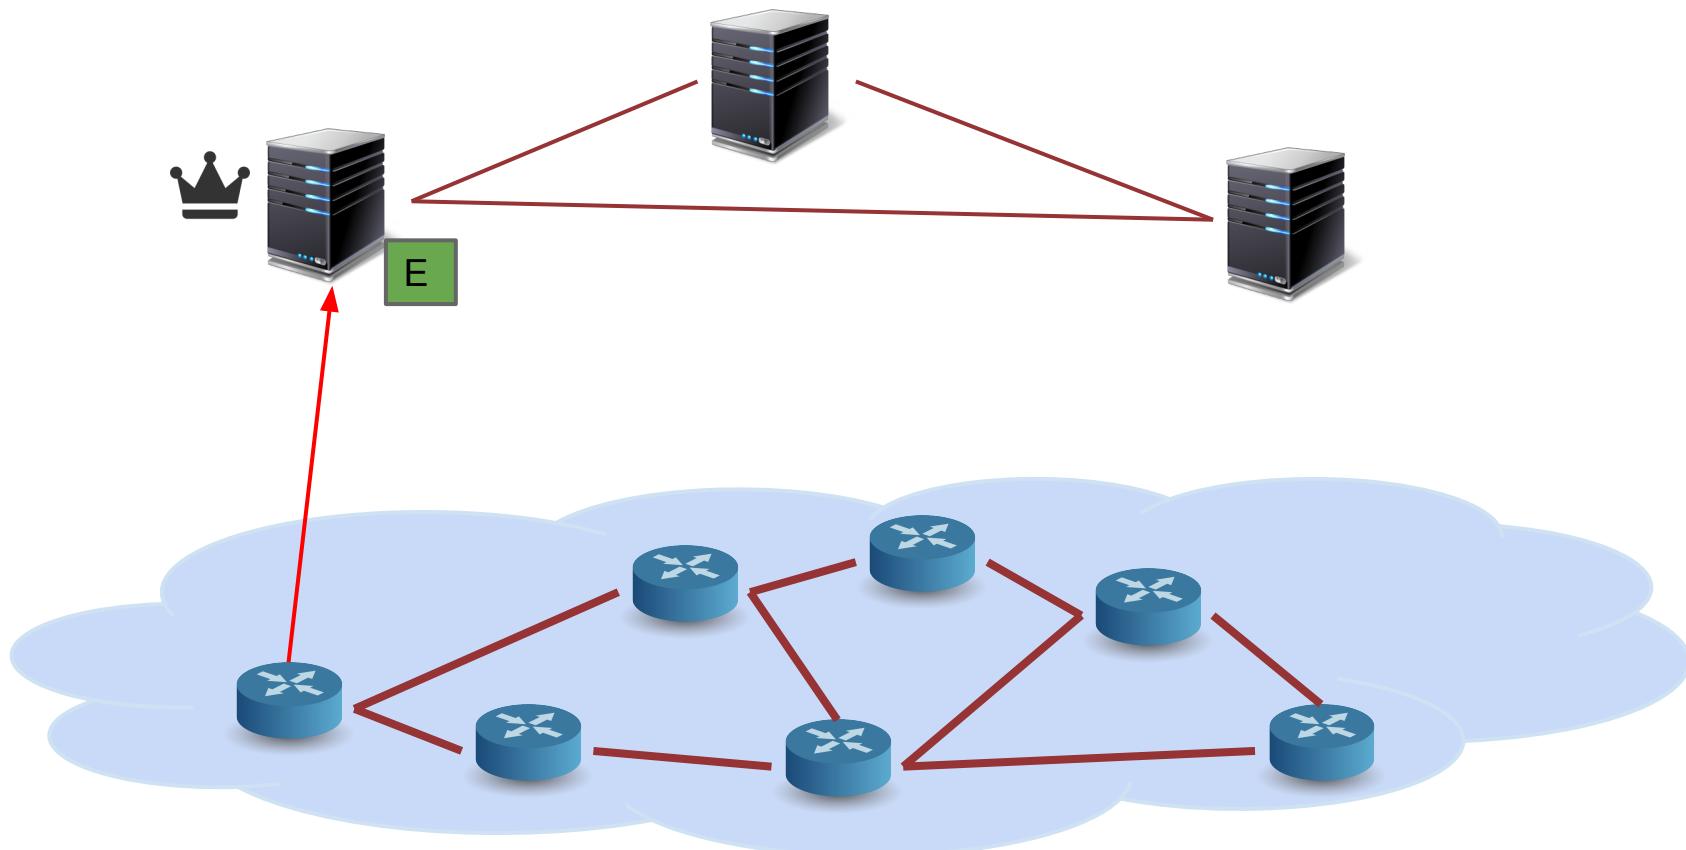


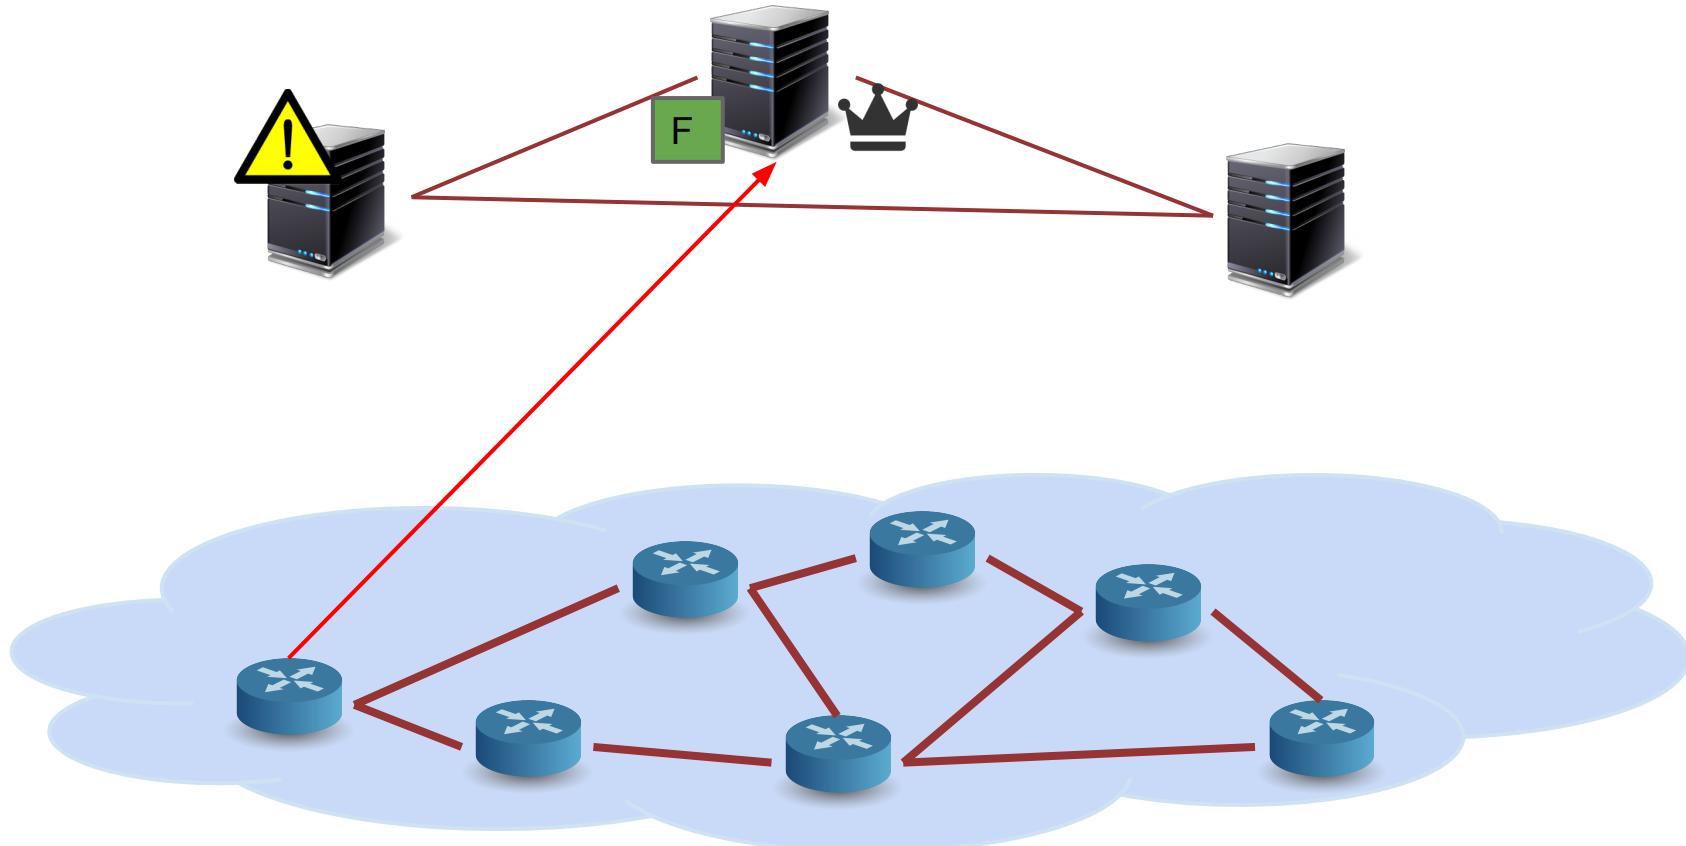


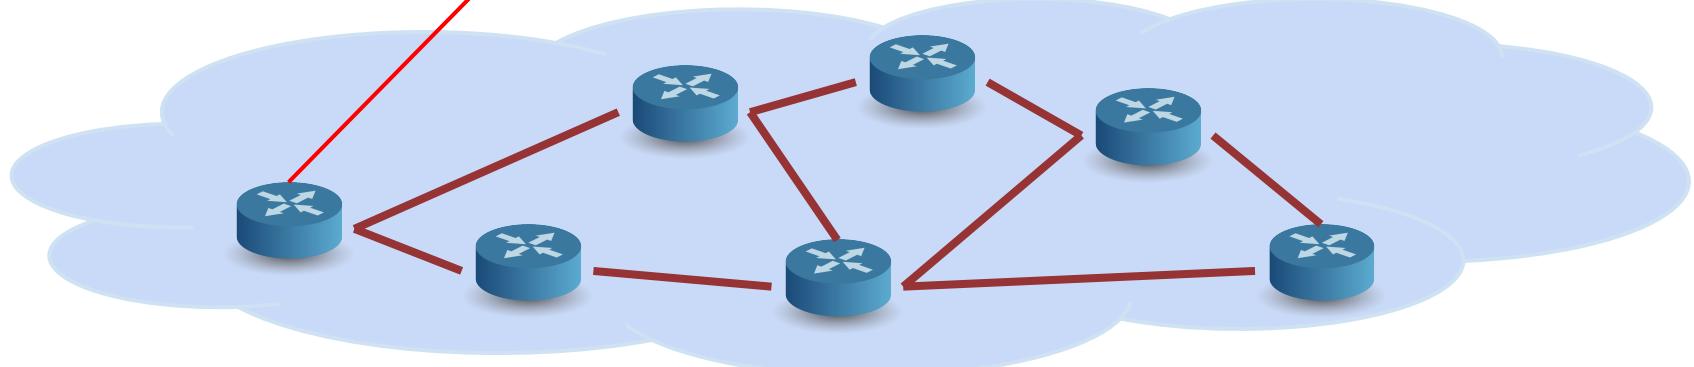
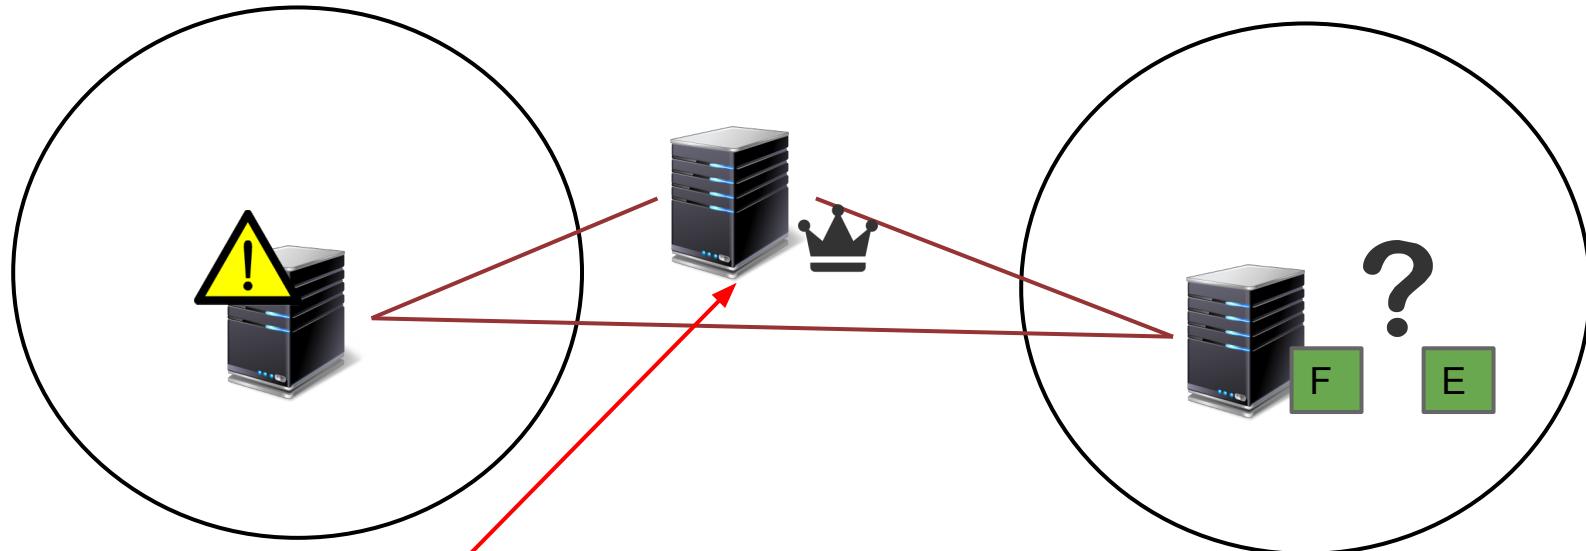


Pros

- Simple
- Fast



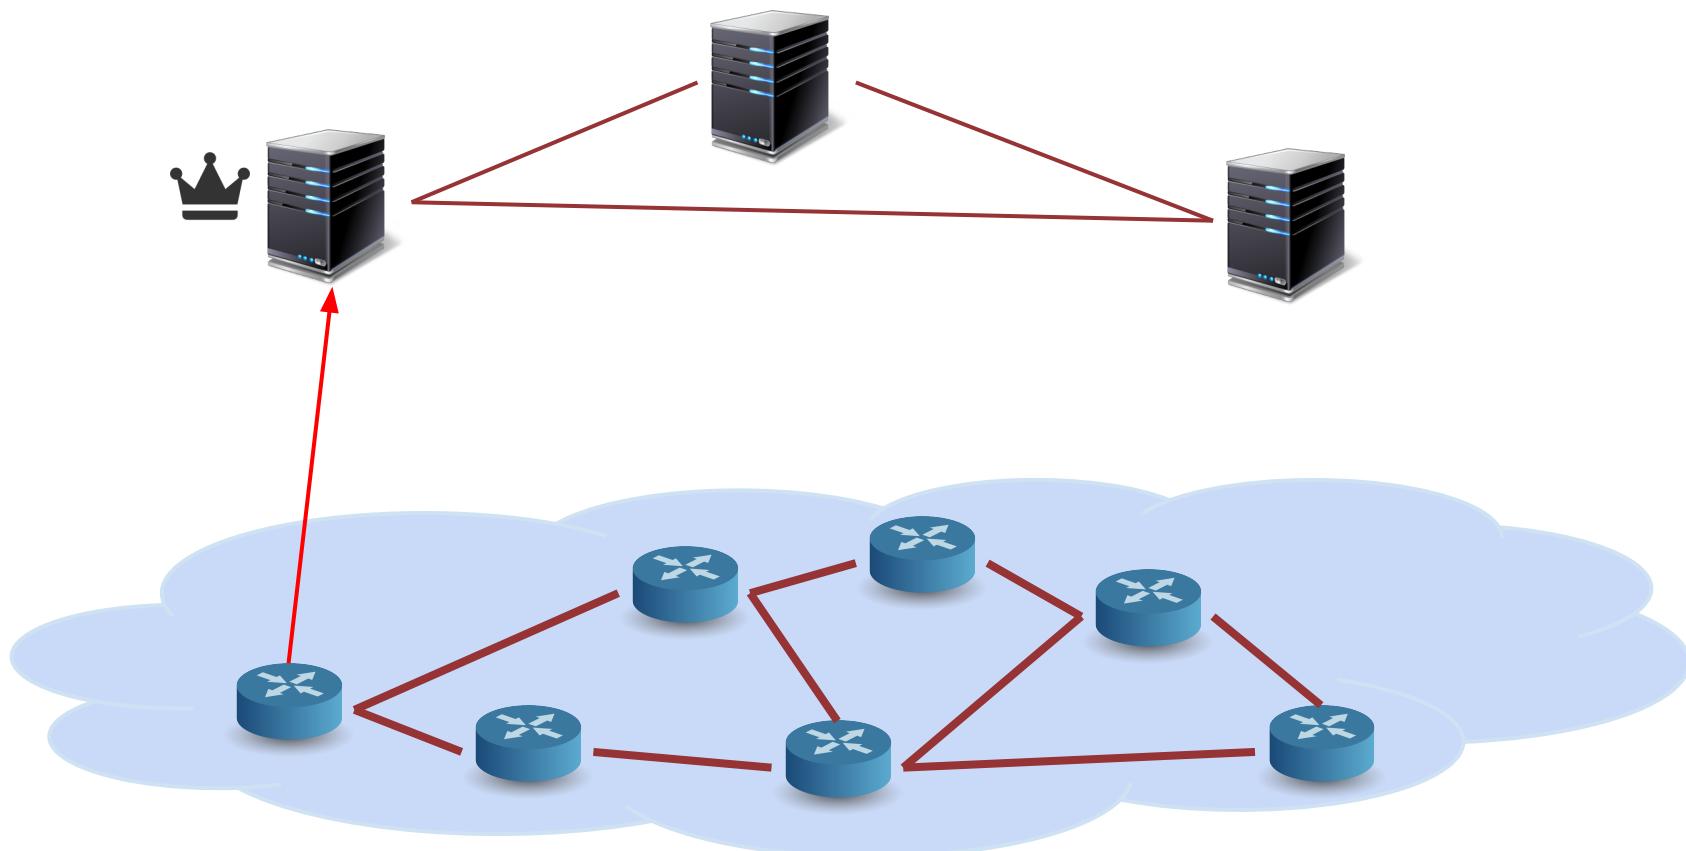


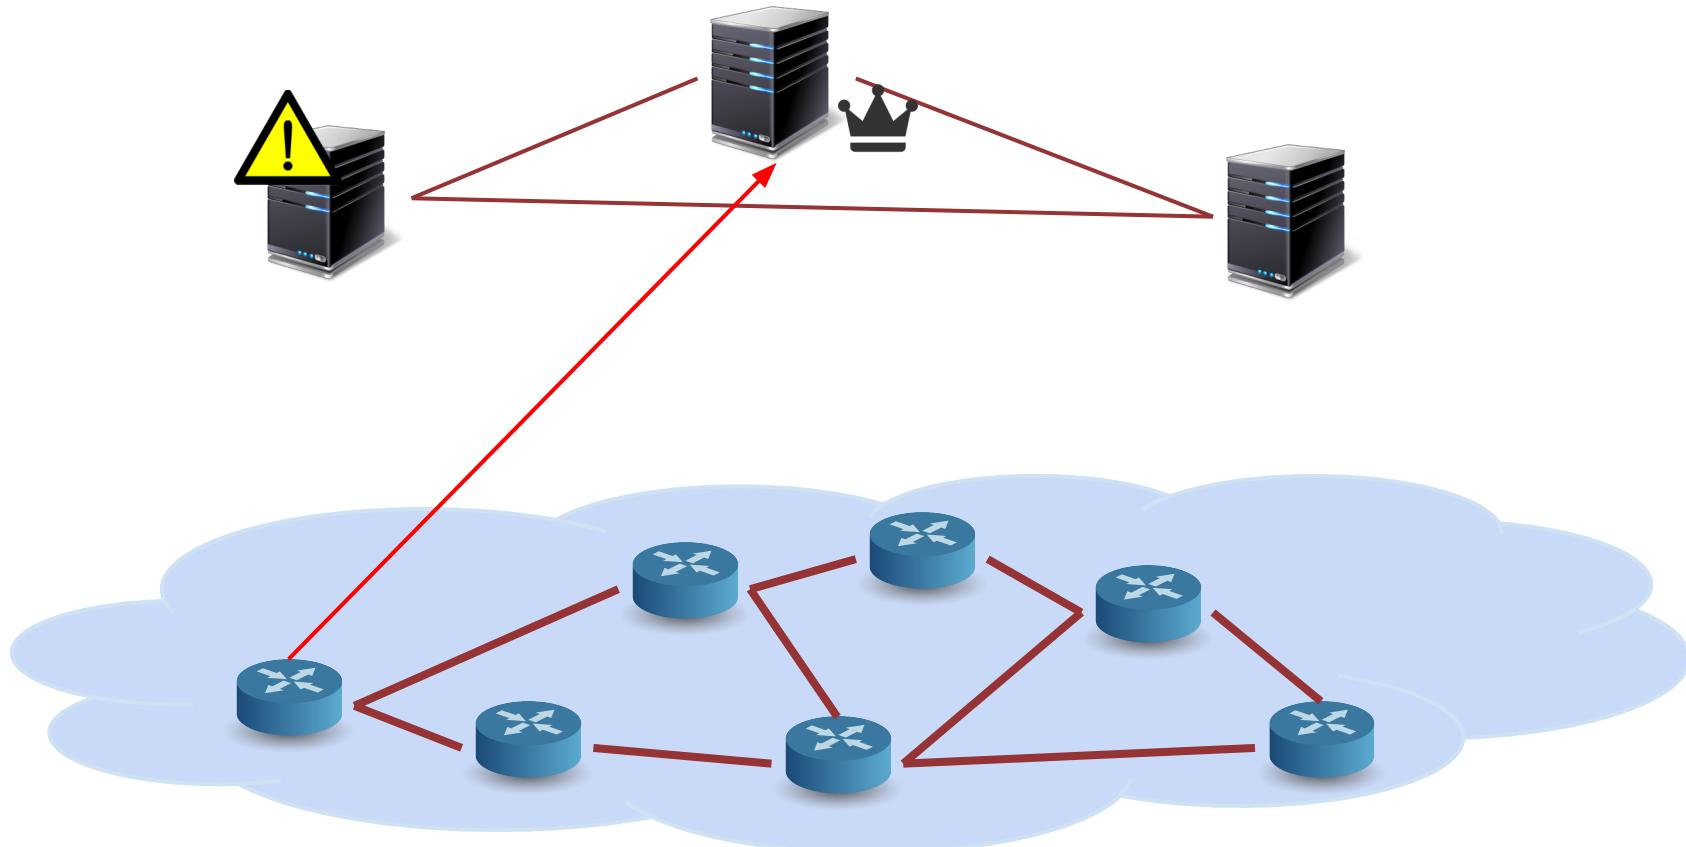




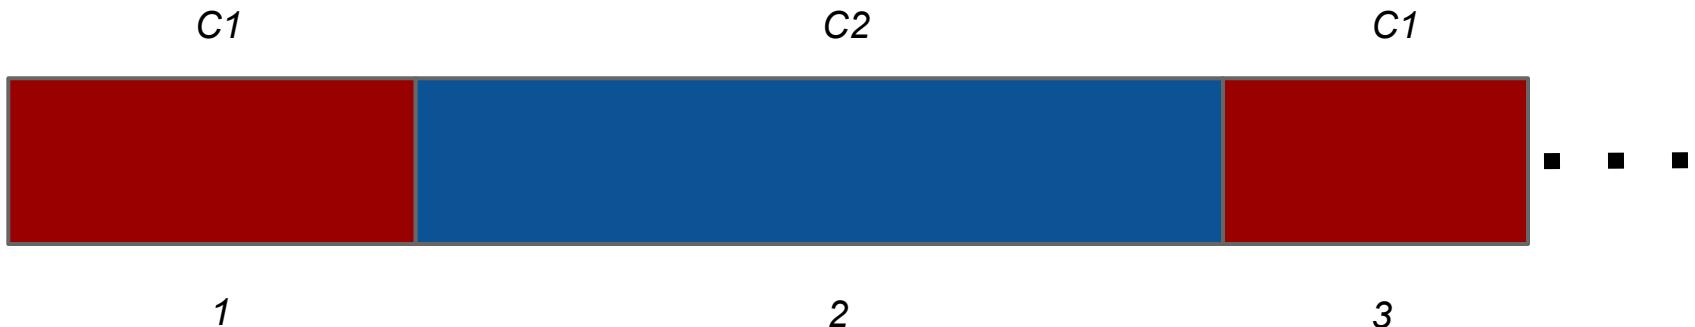
| Pros | Cons |
|---|---|
| <ul style="list-style-type: none">• Simple• Fast | <ul style="list-style-type: none">• Dropped messages• Reordered messages |

How to address the cons without sacrificing the pros?



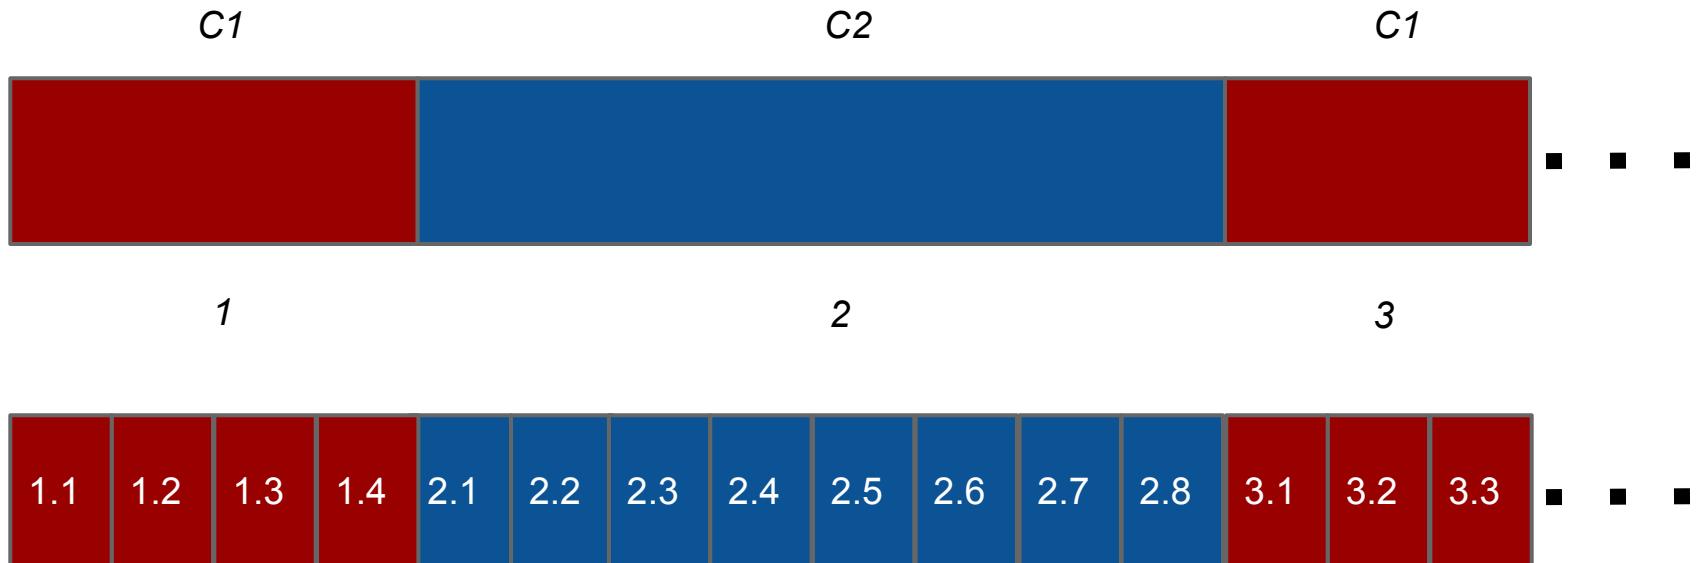


Switch Mastership Terms



We track this in a strongly consistent store

Switch Event Numbers

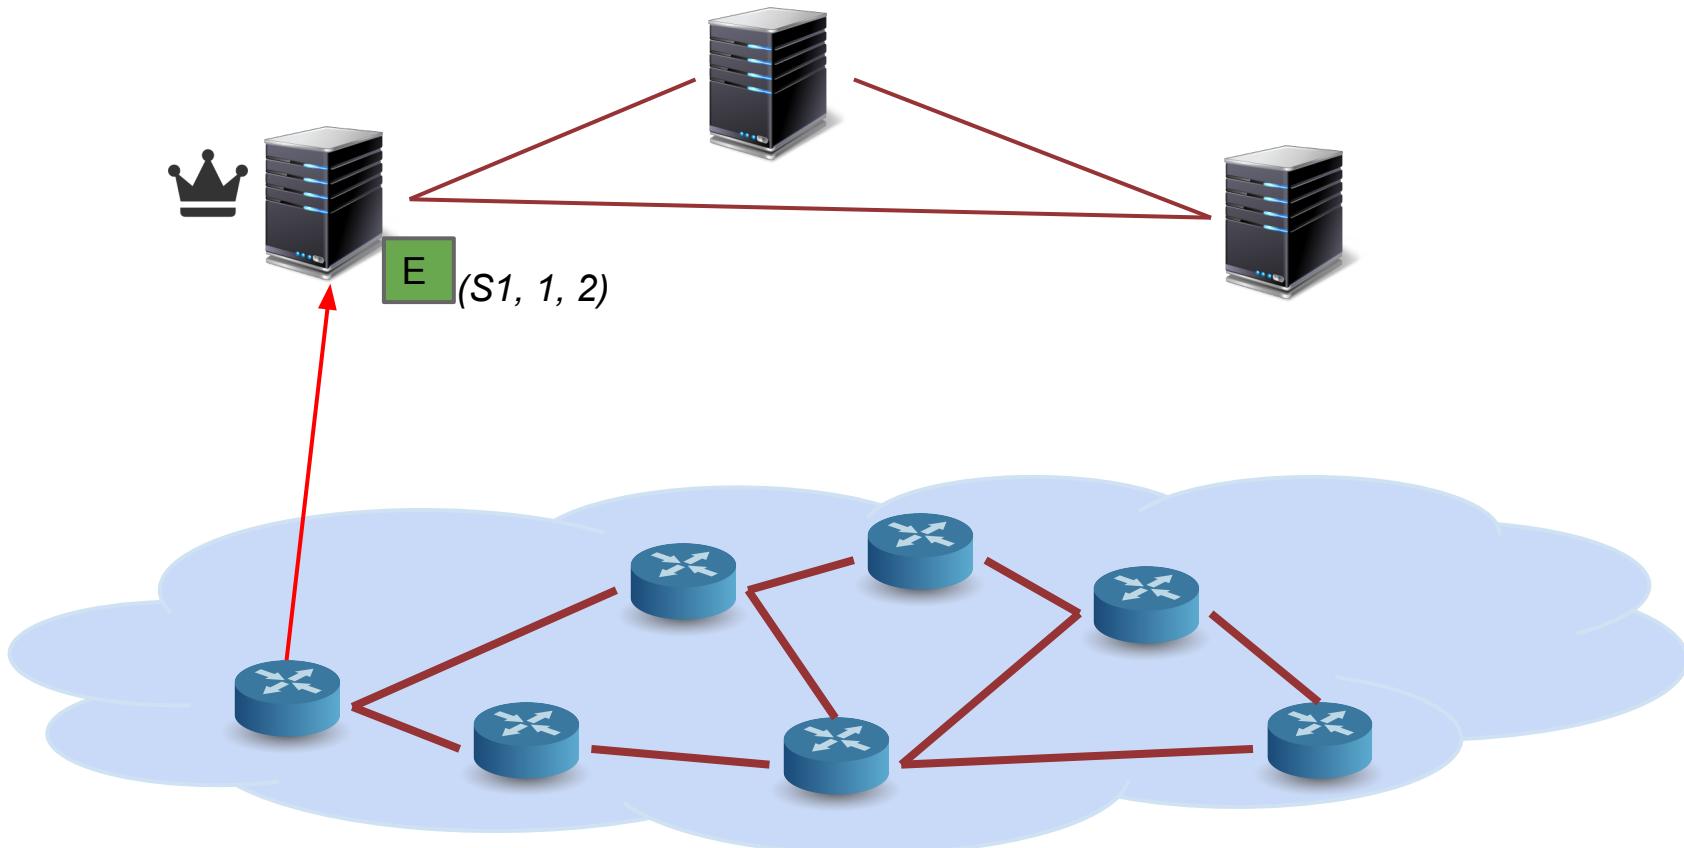


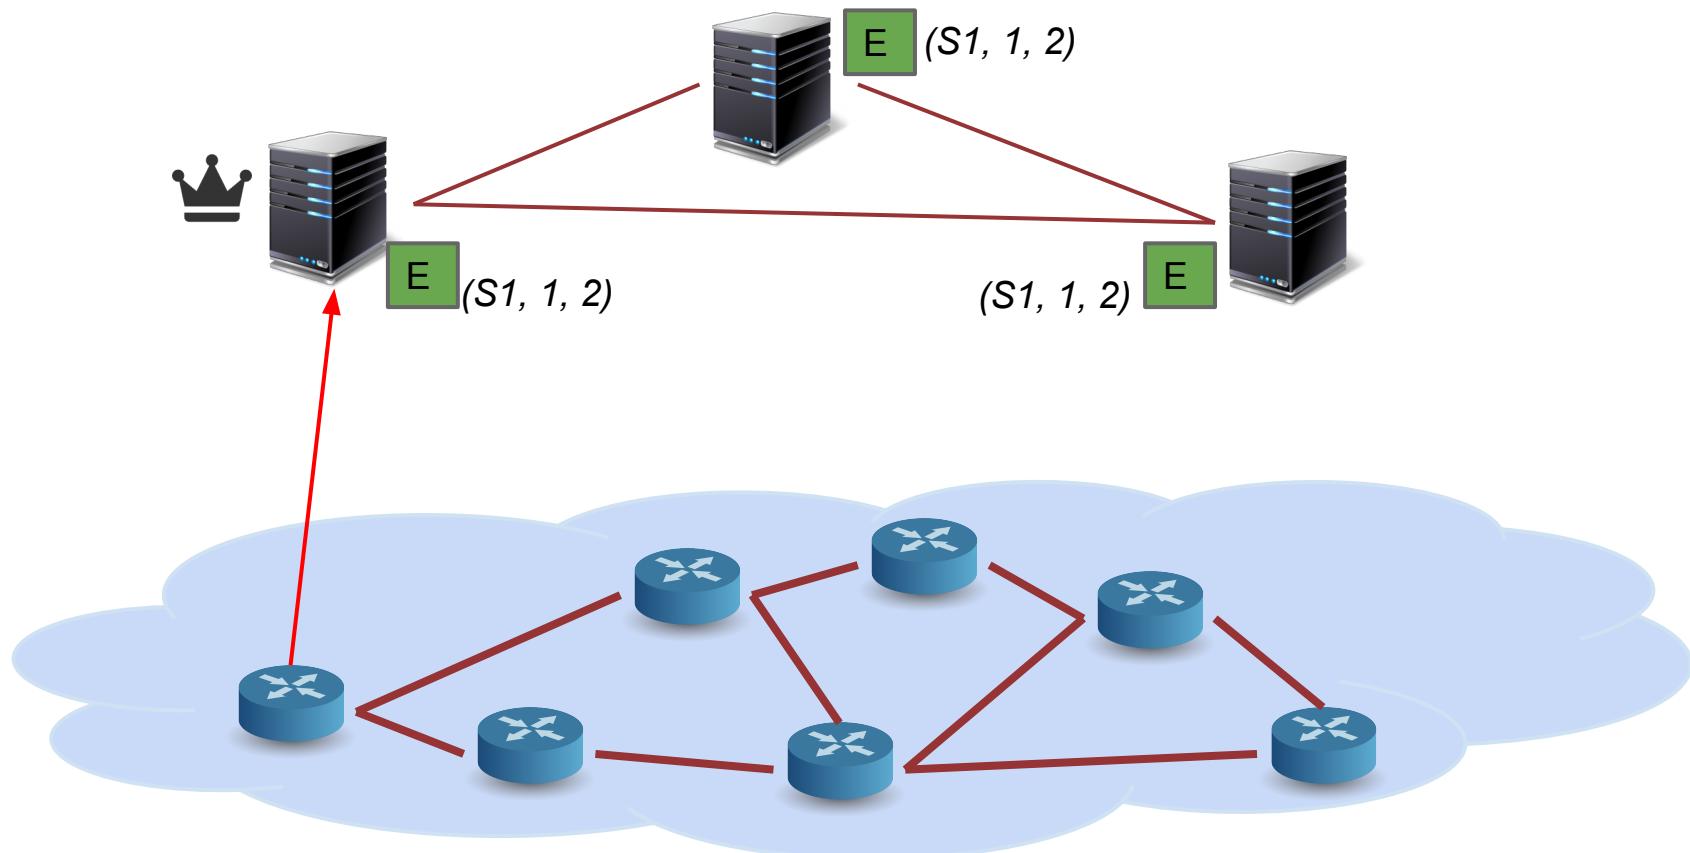
Partial Ordering of Events

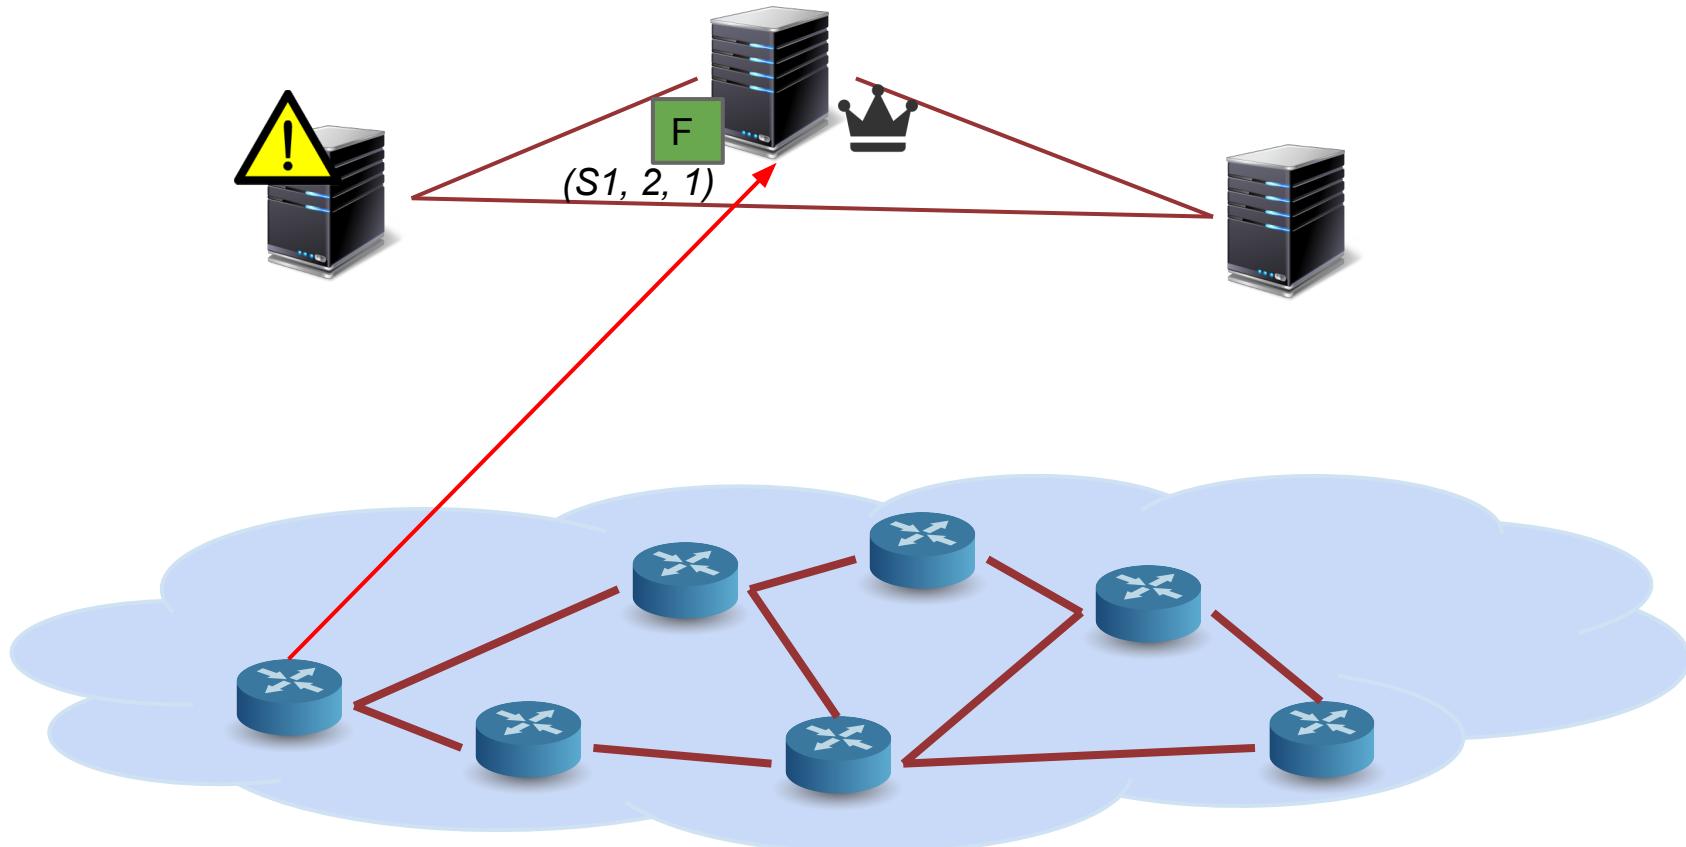


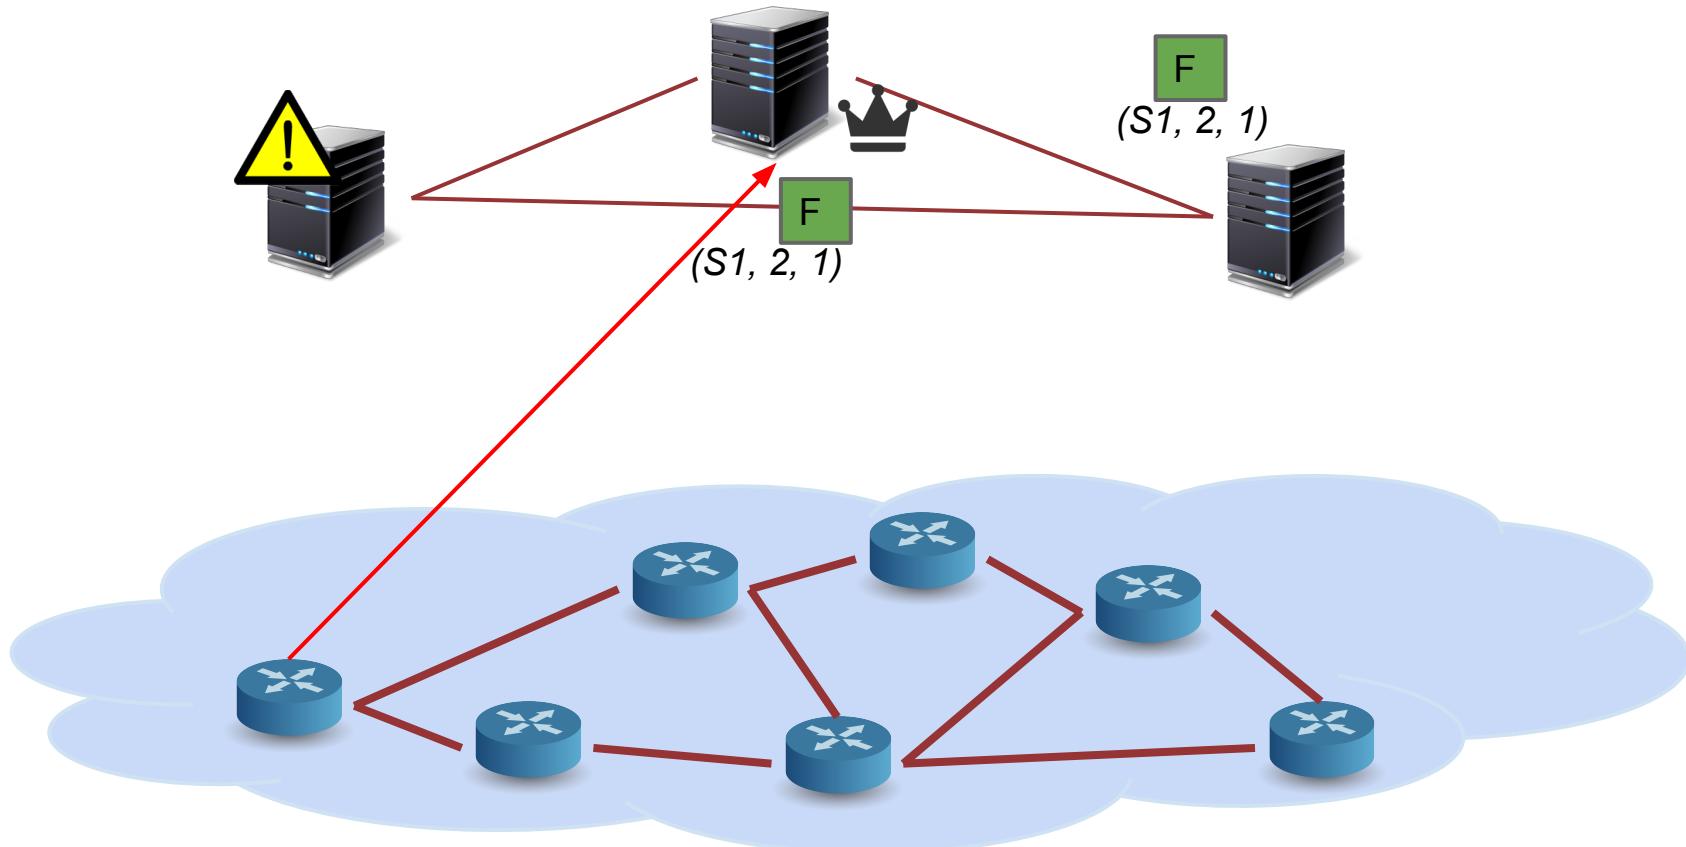
Each event has a unique logical timestamp

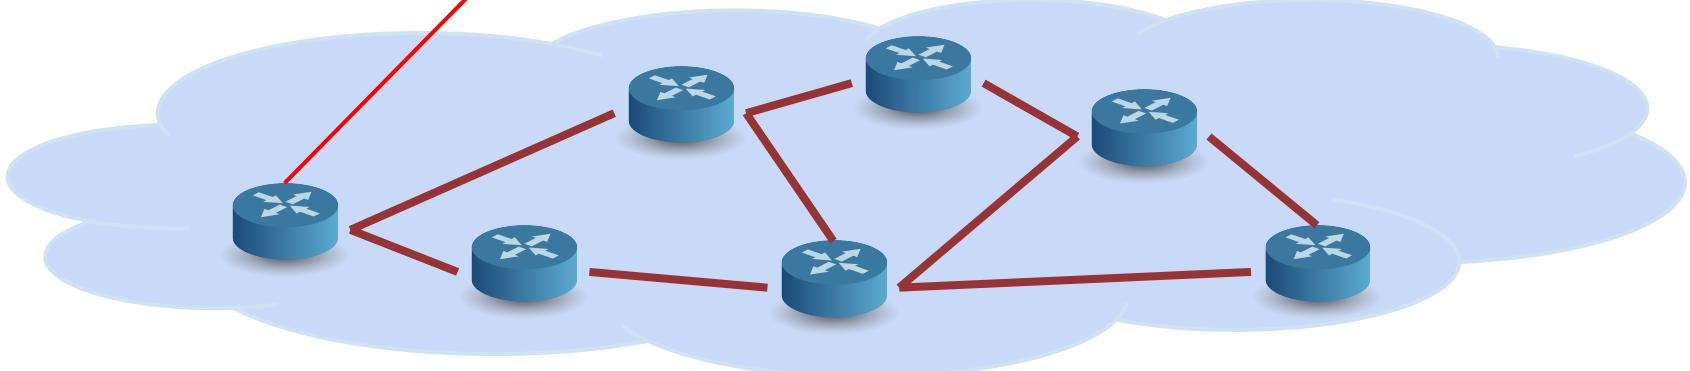
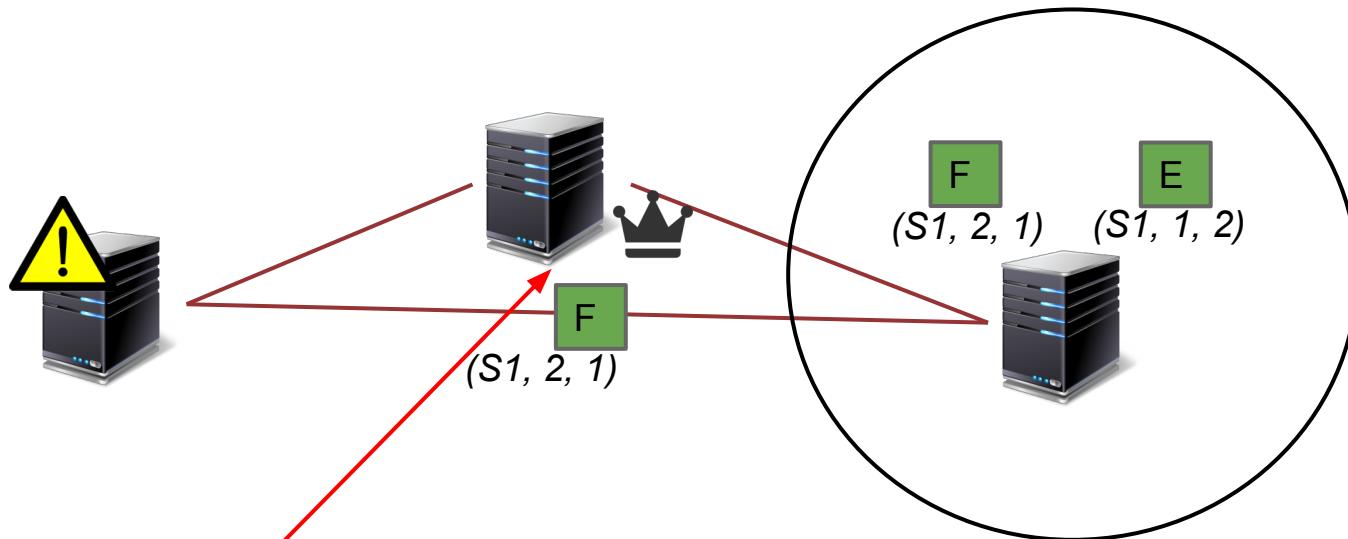
(Switch ID, Term Number, Event Number)

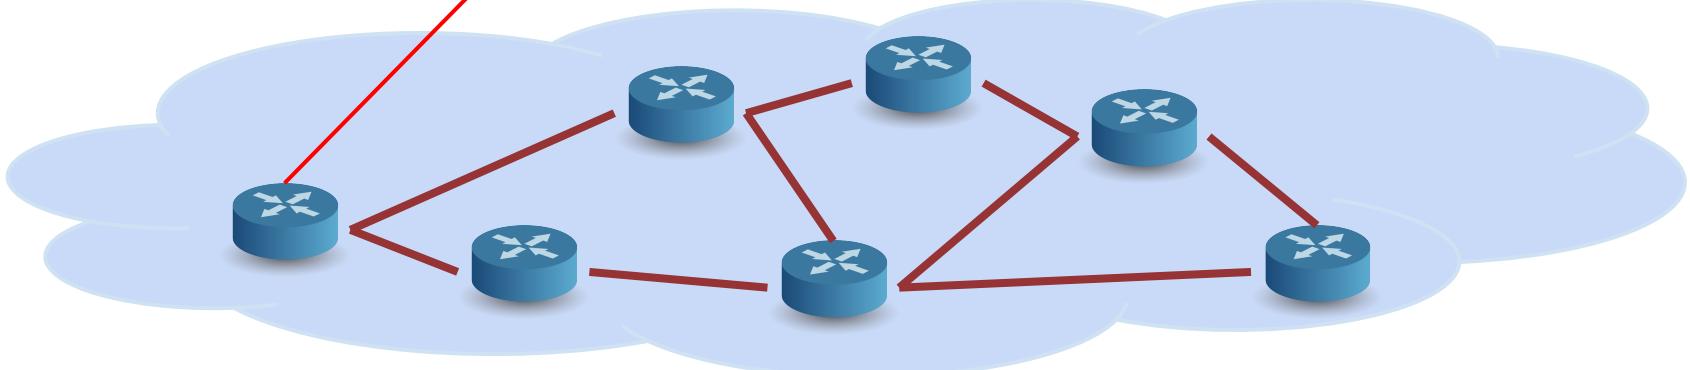
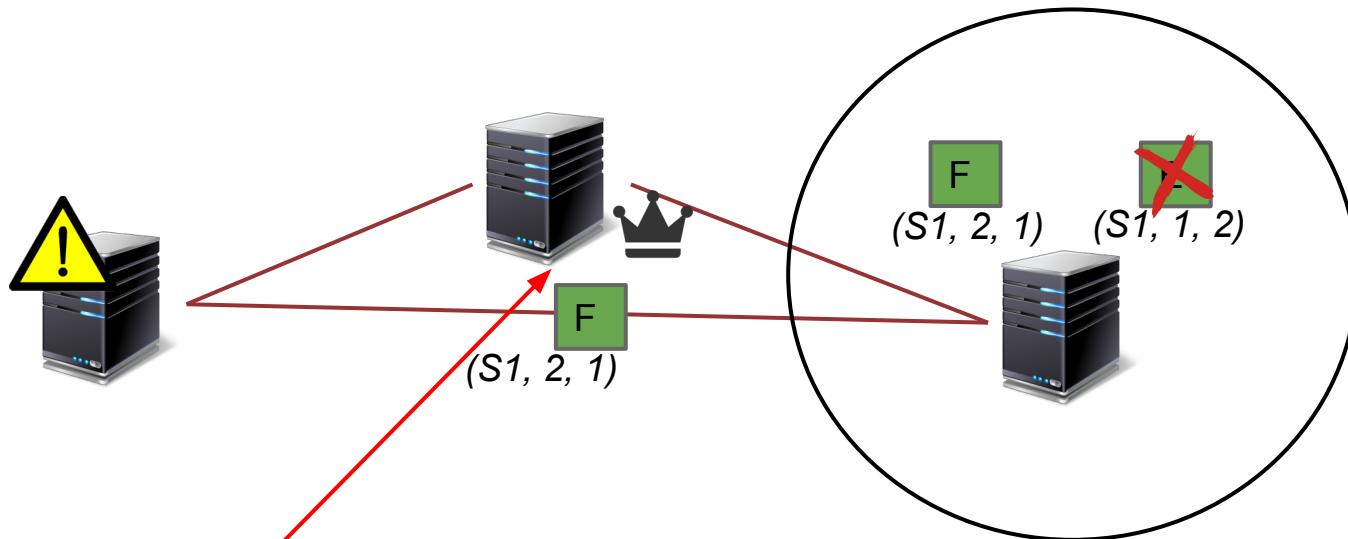












To Summarize



- Each instance has a full copy of network topology
- Events are timestamped on arrival and broadcasted
- Stale events are dropped on receipt

To Summarize



- Each instance has a full copy of network topology
- Events are timestamped on arrival and broadcasted
- Stale events are dropped on receipt

What about dropped messages?



Did you hear
about the port
that went
offline?



© Angie Brown





Anti-Entropy



Gossip-style

Lightweight



Decentralized

A model for state tracking



- If you are the observer, **Eventual Consistency** is the best option
- For topology state, ONOS provides **monotonic read consistency**
- View should be consistent with the network, not with other views

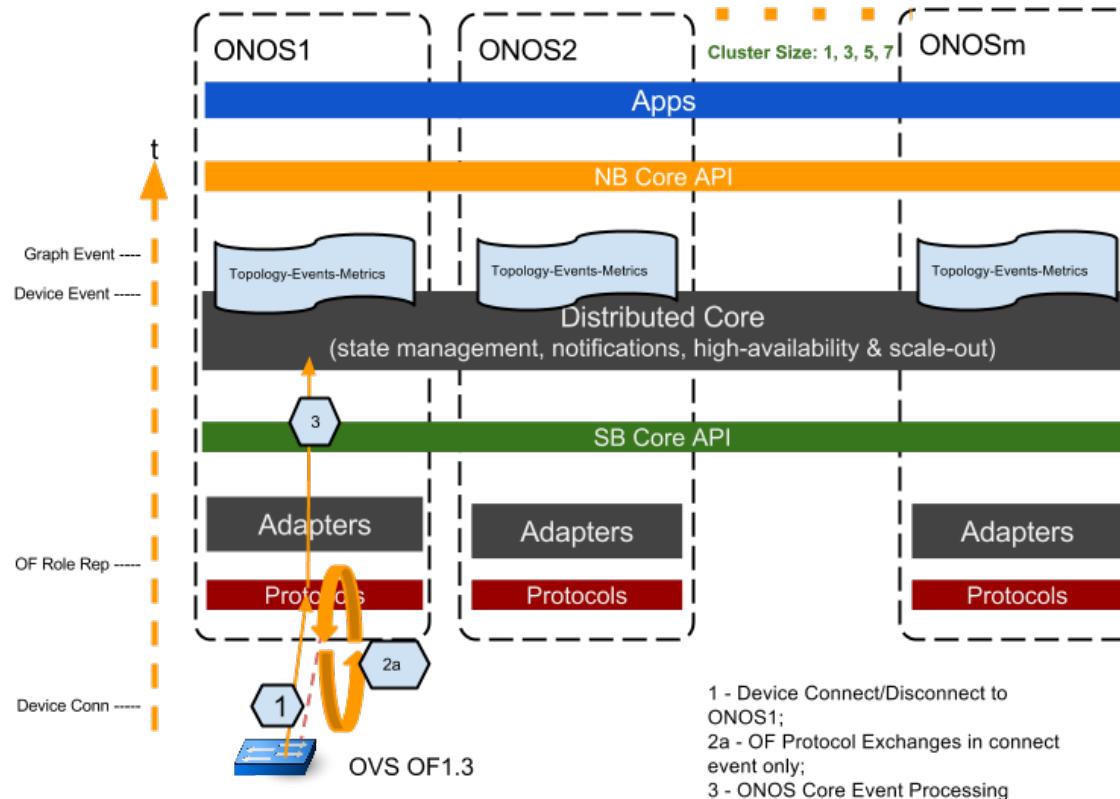
Distributed Primitive



EventuallyConsistentMap<K, V>

- Allows plugging in custom timestamp mechanism

Device & Link Sensing Latency

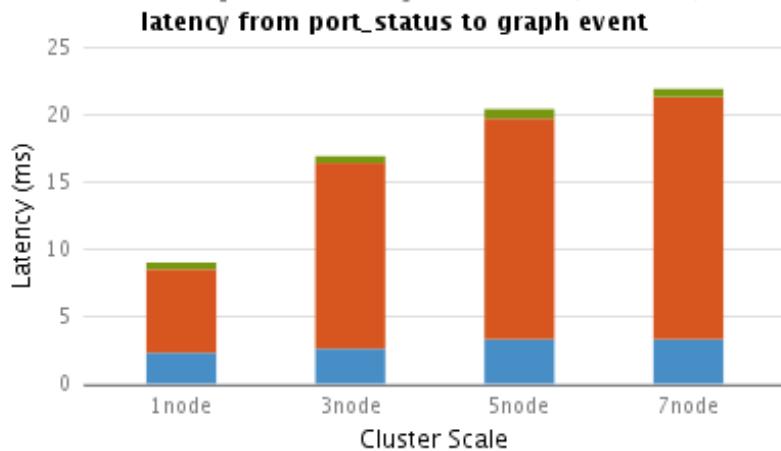


1 - Device Connect/Disconnect to ONOS1;
2a - OF Protocol Exchanges in connect event only;
3 - ONOS Core Event Processing

Link Up/Down Latency

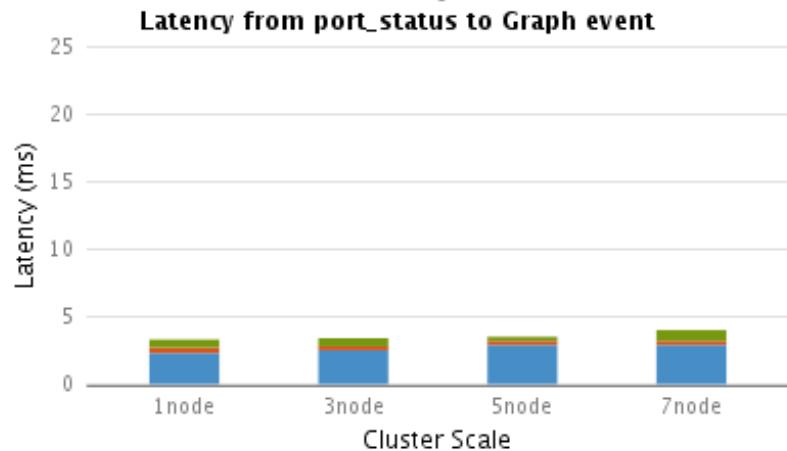


Link Up Latency Tests (Mean)



- OFP Port Status -> Device Event ■ Device Event -> Link Event
- Link Event -> Graph Event

Link Down Latency Tests (Mean)



- OFP Port Status -> Device Event ■ Device Event -> Link Event
- Link Event -> Graph Event

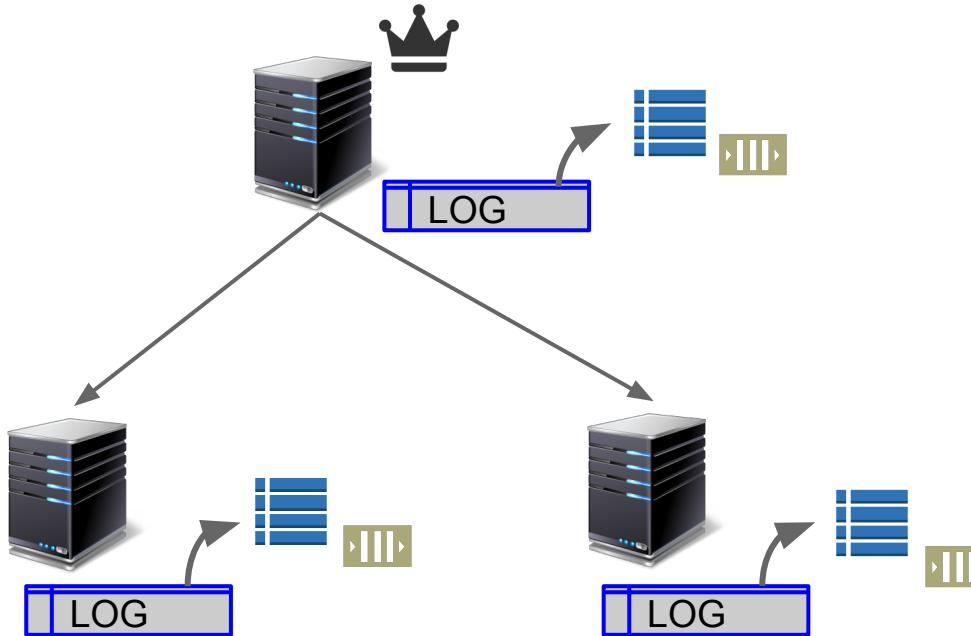
- Since we use LLDP & BDDP to discover links, it takes longer to discover a link coming up than going down
- Port down event trigger immediate teardown of the link.

Next topic: Strong consistency



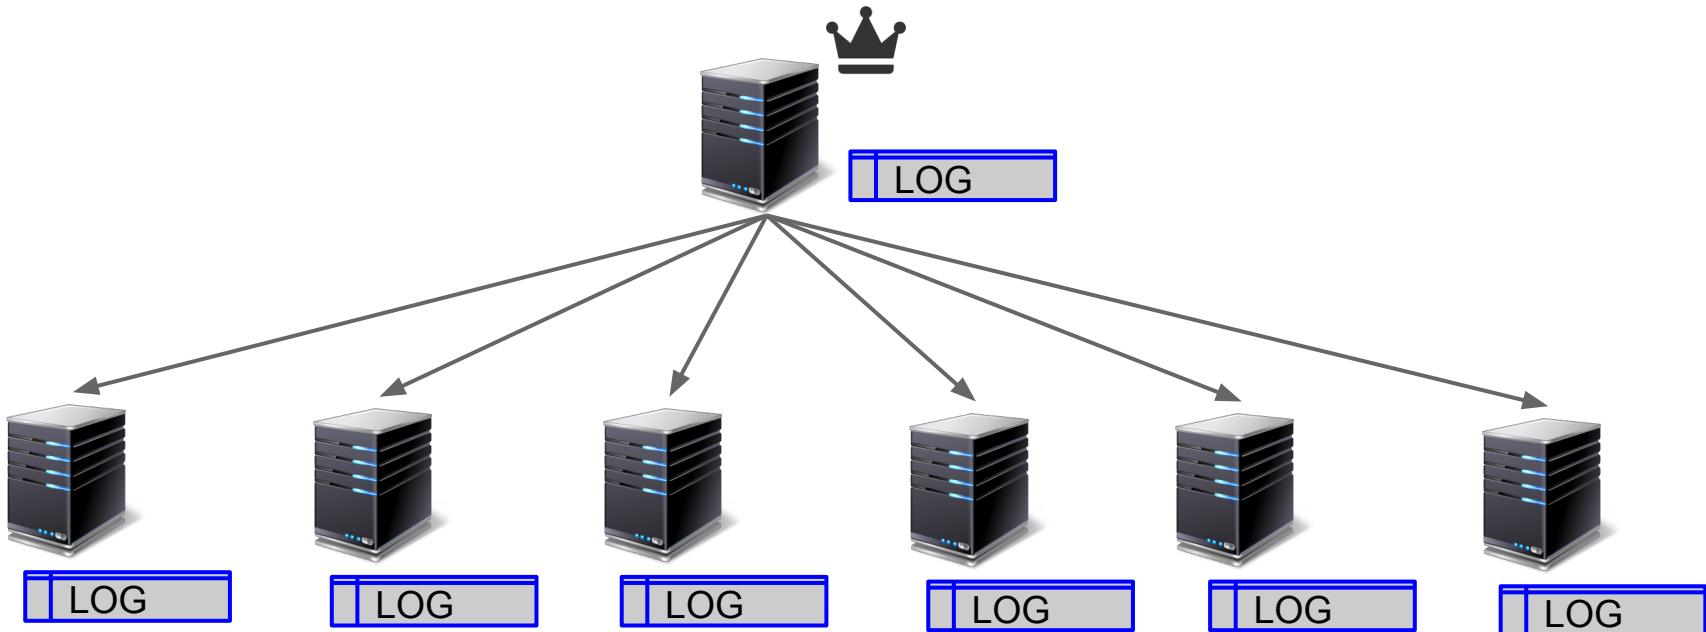
- **EventualConsistency** is necessary but not sufficient
- Some state requires stronger guarantees
 - Reservations is a prime example
- Hard to layer abstractions when you have a store that only provides weak guarantees

Strong Consistency via Consensus

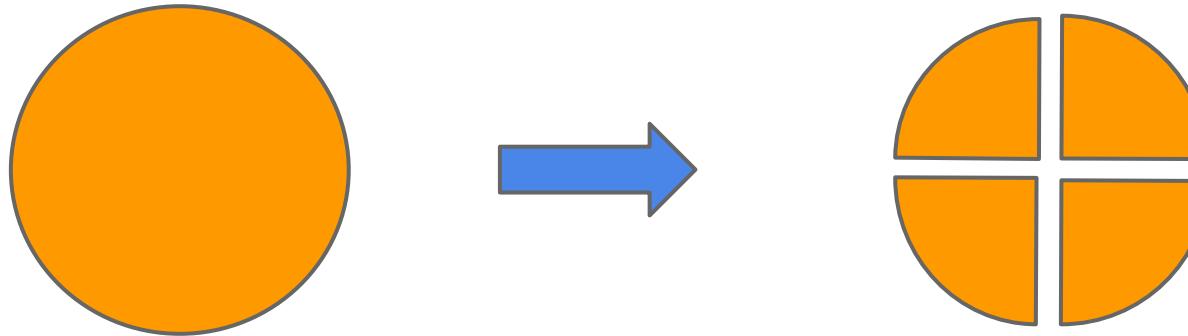


Replicated
State
Machine

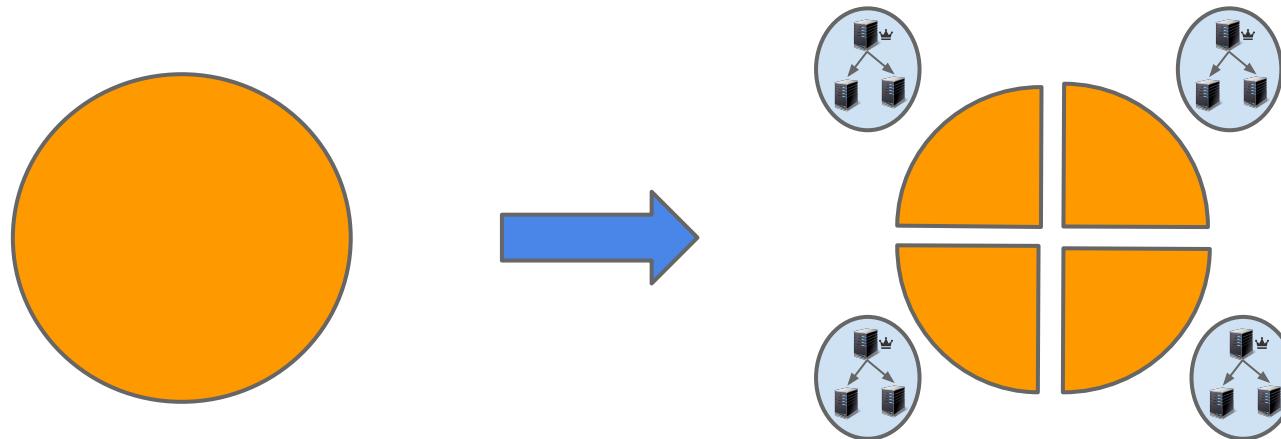
Strong Consistency via Consensus



Scaling Strong Consistency via Partitioning



Scaling Strong Consistency via Partitioning



Distributed Primitives



- ConsistentMap<K, V>
- TransactionalMap<K, V>

Provides transactional semantics for single and group updates

Other Distributed Primitives



- `DistributedSet<E>`
- `DistributedQueue<E>`
- `AtomicCounter`

Final example: Flows

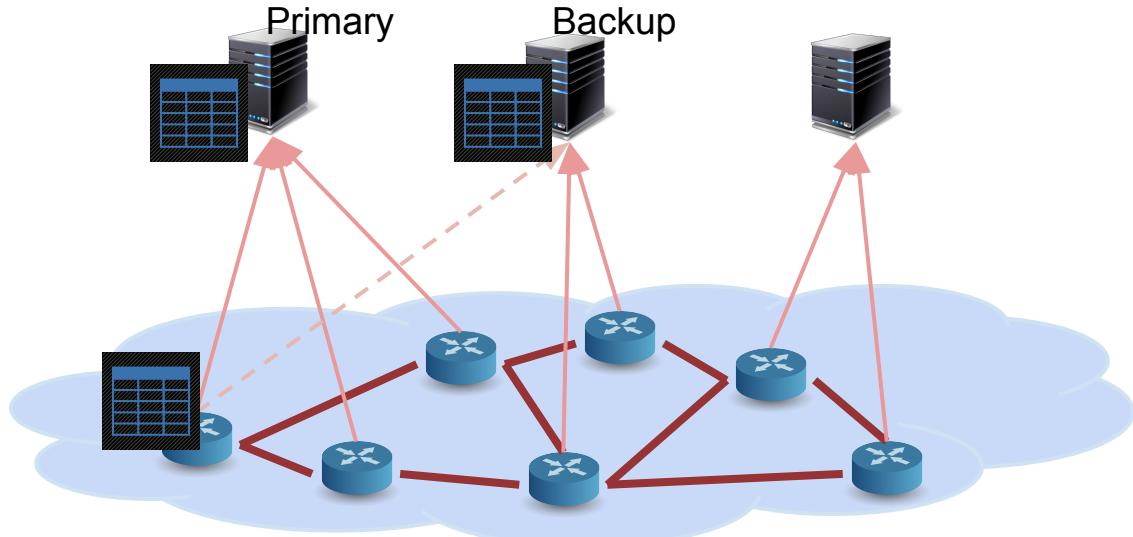


- Forwarding plane rules
- Programmed onto individual switches
- Control plane is the arbiter of truth

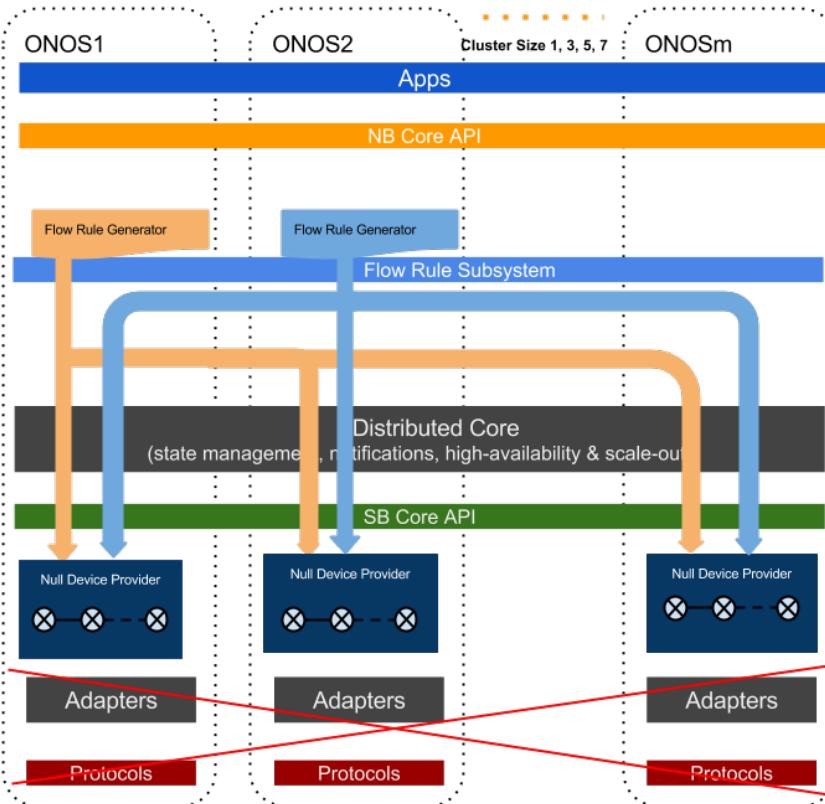
Flows



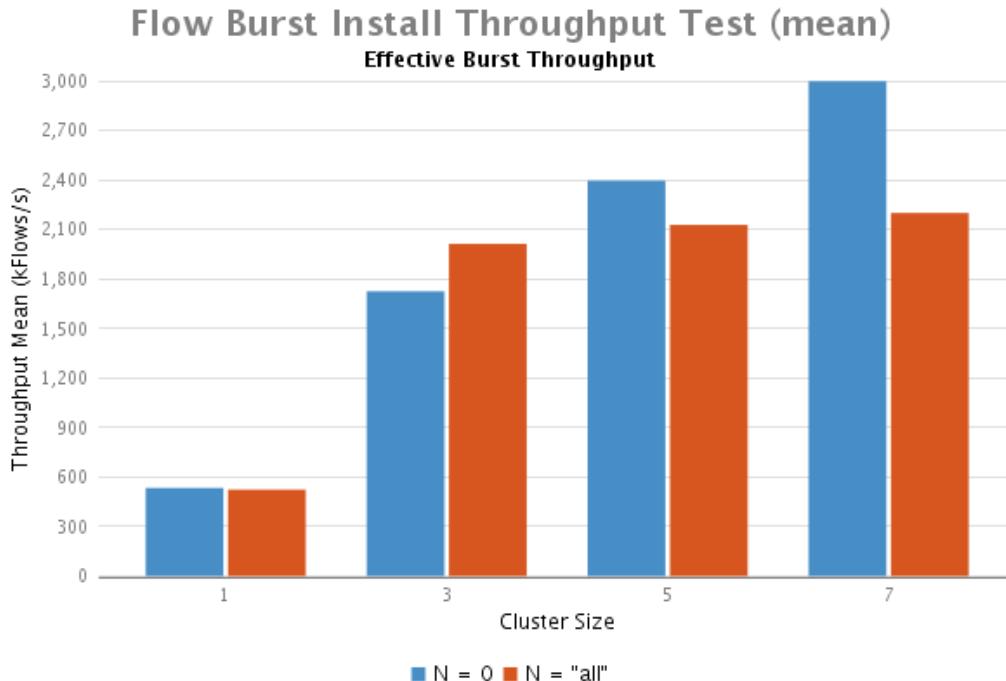
- Primary-backup replication
- Switch master is *primary* and **next** master is *backup*
- Fast read/write access
- Minimal overhead on failover



Flow Rule Operations Throughput

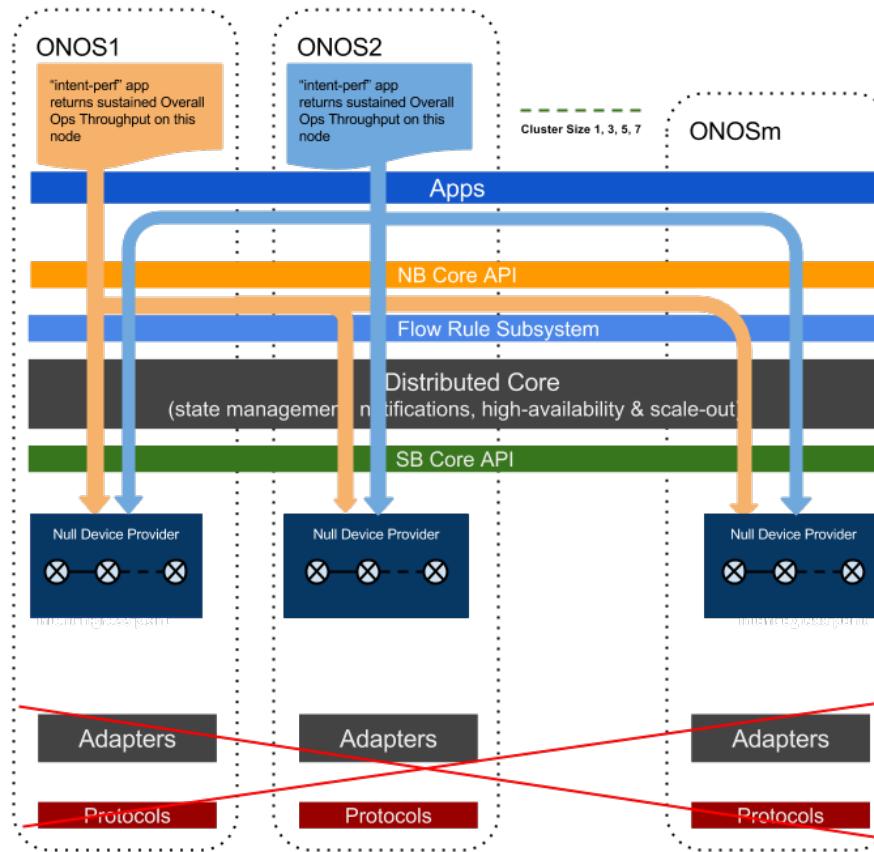


Flow Throughput results

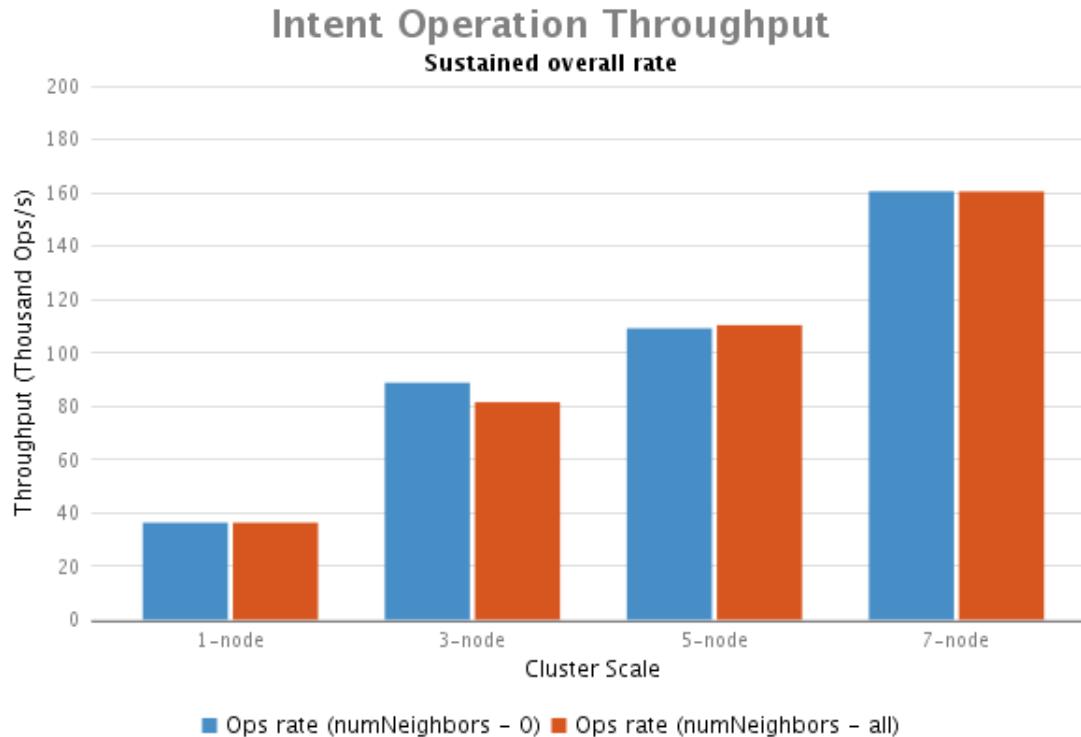


- Single instance can install over 500K flows per second
- ONOS can handle 3M local and 2M non local flow installations
- With 1-3 ONOS instances, the flow setup rate remains constant no matter how many neighbours are involved
- With more than 3 instances injecting load the flow performance drops off due to extra coordination requires.

Intent Throughput Experiment

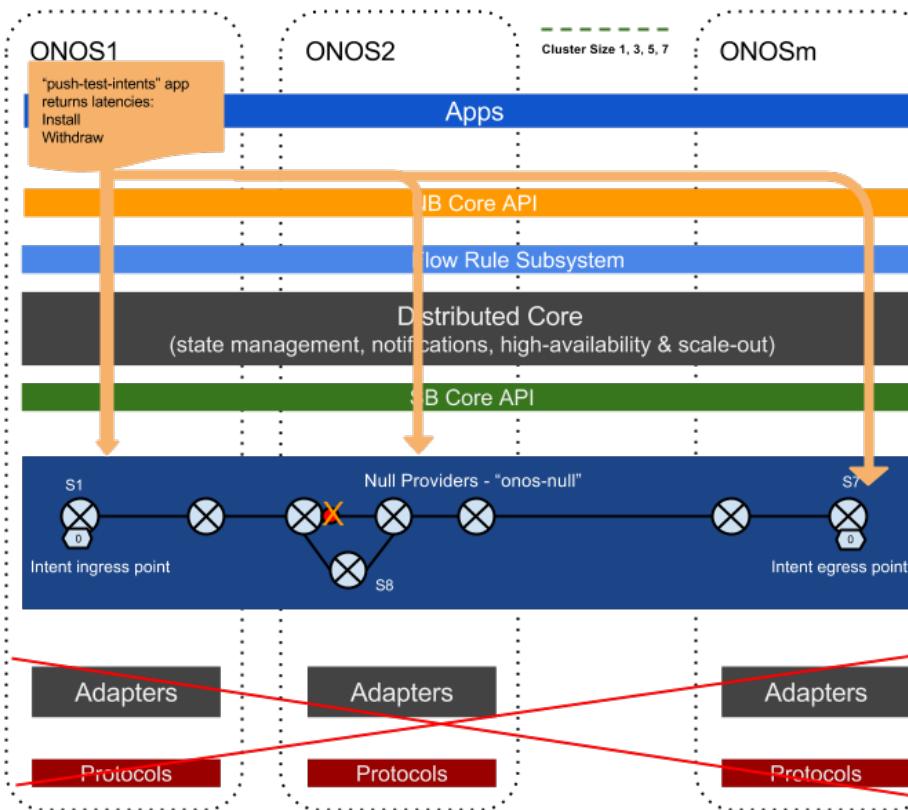


Intent Throughput Results

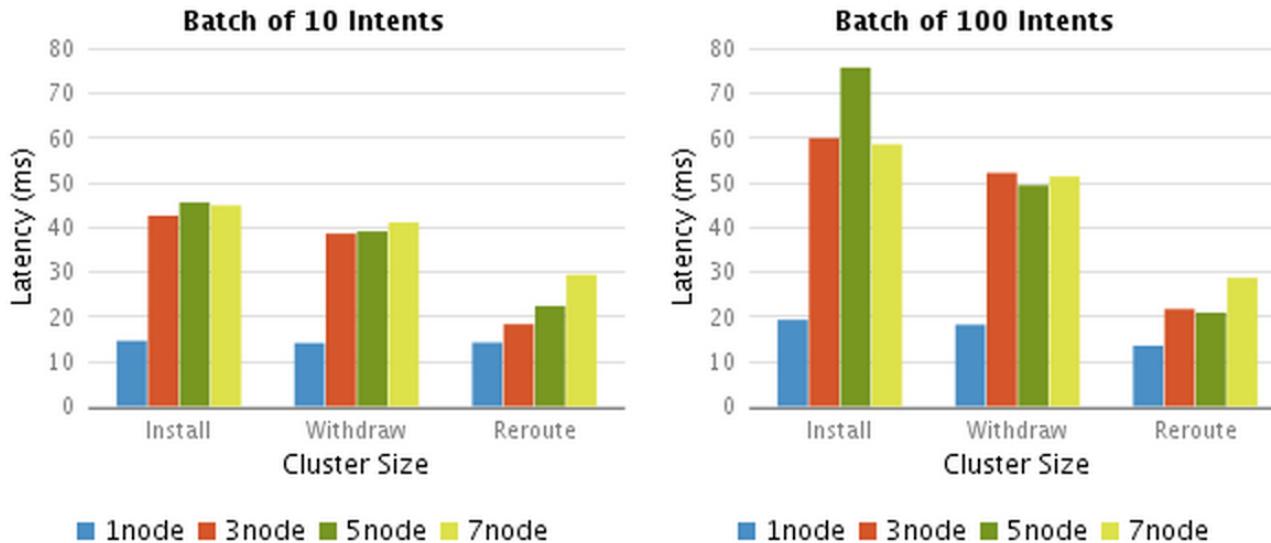


- Processing clearly scales as cluster size increases

Intent Latency Experiment



Intent Latency Results



- Less than 100ms to install or withdraw a batch of intents
- Less than 50ms to process and react to network events
 - Slightly faster because intent objects are already replicated

Lunch Break



YOU KNOW THIS METAL
RECTANGLE FULL OF
LITTLE LIGHTS?

YEAH.



I SPEND MOST OF MY LIFE
PRESSING BUTTONS TO MAKE
THE PATTERN OF LIGHTS
CHANGE HOWEVER I WANT.

SOUNDS
GOOD.



BUT TODAY, THE PATTERN
OF LIGHTS IS ALL WRONG!

OH GOD! TRY
PRESSING MORE
BUTTONS!
IT'S NOT
HELPING!



ONOS Dev Workshop Outline



- ONOS *Fundamentals*
 - *architecture, checking out, building & running ONOS*
- ONOS *Distributed Core*
 - *OSGi, anatomy of a core subsystem, distributed stores & primitives*
- ONOS **App Development**
 - background, Maven archetypes, app deployment on ONOS cluster
- ONOS Interfaces
 - component config service, flow objective service, intent service
 - CLI, REST and GUI overview
- Wrap-up

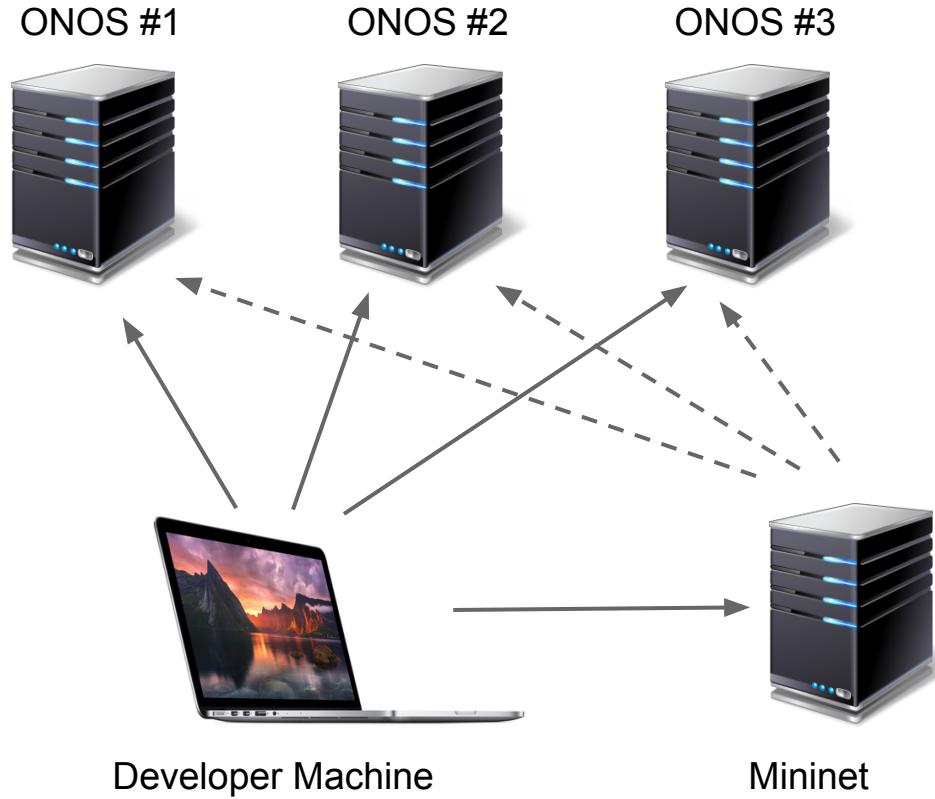
ONOS App Development

13:15-14:45



- Installing ONOS cluster from `onos.tar.gz` (~10)
 - untar, `onos-service`, `onos-form-cluster`
- ONOS Applications (~10)
 - OSGi bundles, Karaf features, ONOS apps & OAR files
 - application lifecycle - install, activate, deactivate, uninstall
 - ONOS CLI & GUI
- Developing ONOS apps (~50)
 - use `mvn archetype:generate` and `onos-create-app`
 - archetype overlays for CLI and UI - generated & pre-canned
 - iterative and demonstrating use of `onos-app` tool or GUI

Working with ONOS cluster



Install ONOS cluster from tar.gz



- Install ONOS on all three instances
 - we will use iTerm to do so simultaneously

```
$ onos-cell demo && onos-iterm-cli
```

- Fetch ONOS bits via wget

```
$ wget http://downloads.onosproject.org/onos-1.3.0.dev.tar.gz
```
- Untar bits

```
$ tar zxf onos-1.3.0.dev.tar.gz
```



Start ONOS instances

- Start ONOS

```
$ cd onos-1.3.0.dev  
$ bin/onos-service server &
```
- Verify install as individual instances

```
$ bin/onos  
onos> nodes
```
- Check logs

```
$ tail -f log/karaf.log
```

Form ONOS cluster



- Form ONOS cluster from individual instances

```
$ bin/onos-form-cluster 10.128.11.1 10.128.11.2 10.128.11.3
```

- Verify install as a cluster

```
$ bin/onos  
onos> nodes
```

- Check logs

```
$ tail -f log/karaf.log
```

Setup test network



- Connect mininet to cluster machines
`$ sudo mn --topo tree,3,4 --controller remote,...`
- Activate proxy ARP and reactive forwarding apps
`onos> app activate org.onosproject.proxyarp`
`onos> app activate org.onosproject.fwd`
- Issue pingall a few times to show unencumbered flow
`mininet> pingall`
`mininet> pingall`

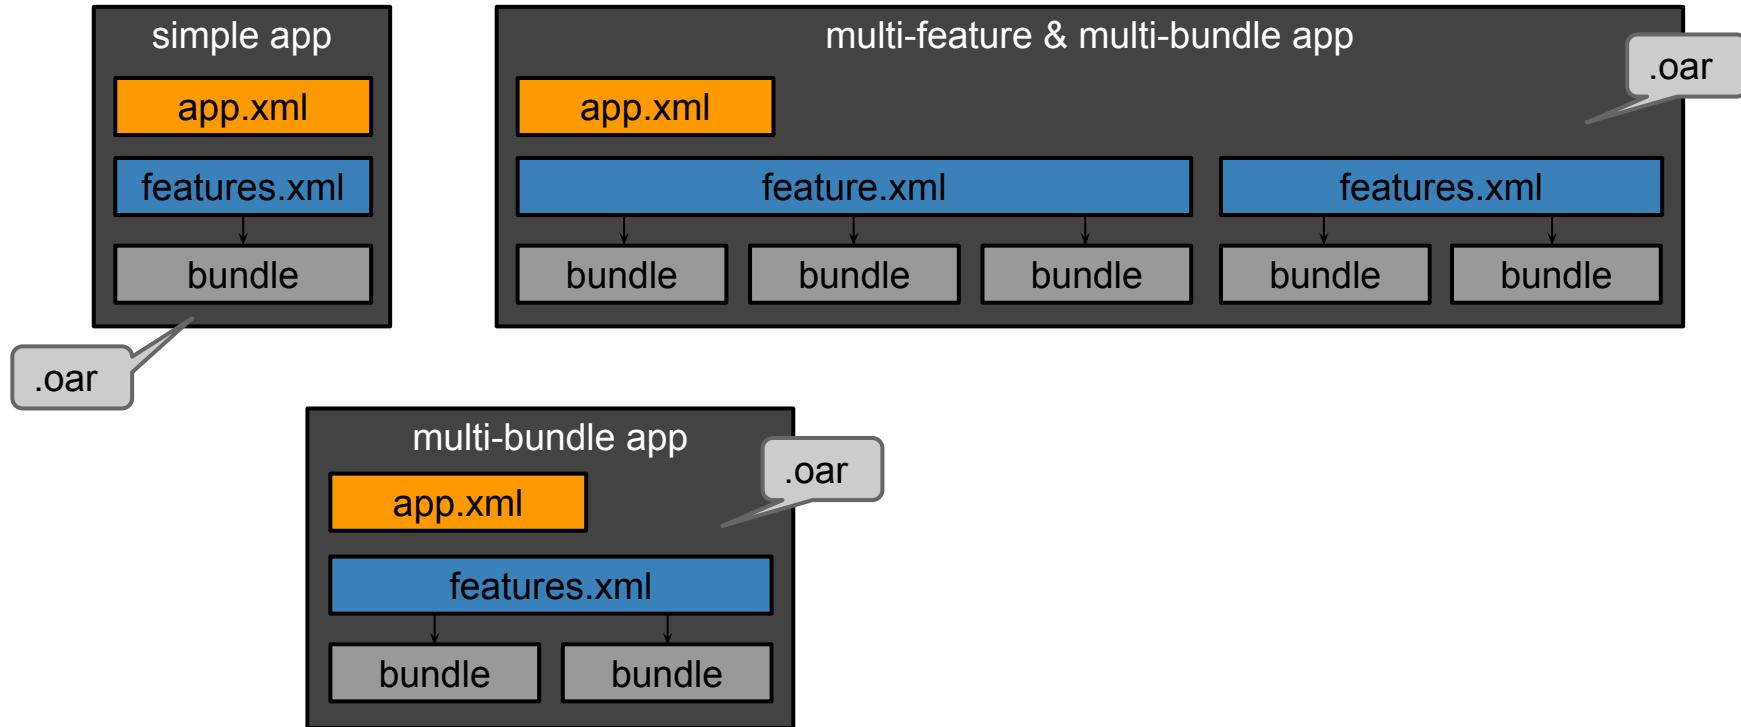
ONOS App Development

13:15-14:45



- Installing ONOS cluster from onos.tar.gz (~10)
 - untar, `onos-service`, `onos-form-cluster`
- ONOS Applications (~10)
 - OSGi bundles, Karaf features, ONOS apps & OAR files
 - application lifecycle - install, activate, deactivate, uninstall
 - ONOS CLI & GUI
- Developing ONOS apps (~50)
 - use `mvn archetype:generate` and `onos-create-app`
 - archetype overlays for CLI and UI - generated & pre-canned
 - iterative and demonstrating use of `onos-app` tool or GUI

Bundles, Features & ONOS Apps



Bundles, Features & ONOS Apps



- Apps are delivered via ONOS App aRchive (**.oar**) files
 - OAR is a JAR with `app.xml`, `features.xml` and bundle artifacts
 - `onos-maven-plugin` generates an ***.oar** file as part of Maven build
- Apps are managed on the entire ONOS cluster
 - via REST API: `GET | POST | DELETE /onos/v1/applications`
 - via shell tool: `onos-app {install|activate|deactivate|uninstall}`
 - via CLI: `onos:app {install|activate|deactivate|uninstall}`
 - via GUI Applications view
- Back-end installation and activation is done via normal feature & bundle services

Developing ONOS apps



- Maven archetypes
 - `onos-bundle-archetype` - basis for an ONOS bundle or an app
 - `onos-cli-archetype` - overlay for apps with CLI extensions
 - `onos-ui-archetype` - overlay for apps with GUI extensions
 - `onos-api-archetype` - basis for an app Java API bundle
- Run `mvn archetype:generate` to create a working minimal project module
- For simpler usage run `onos-create-app` shell tool
 - simply a convenience wrapper for `mvn`

ONOS App Development

13:15-14:45



- Installing ONOS cluster from onos.tar.gz (~10)
 - untar, `onos-service`, `onos-form-cluster`
- ONOS Applications (~10)
 - OSGi bundles, Karaf features, ONOS apps & OAR files
 - application lifecycle - install, activate, deactivate, uninstall
 - ONOS CLI & GUI
- Developing ONOS apps (~50)
 - use `mvn archetype:generate` and `onos-create-app`
 - archetype overlays for CLI and UI - generated & pre-canned
 - iterative and demonstrating use of `onos-app` tool or GUI

One Ping Only ONOS app



One Ping Only ONOS app



Minimal ONOS app project



- Create minimal app project via `onos-create-app` tool
`$ onos-create-app app org.oneping oneping 1.0-SNAPSHOT`
- Build the app via Maven
`$ cd oneping`
`$ mvn clean install`
- Deploy the app to ONOS cluster via `onos-app` tool
`$ onos-app 10.128.11.1 install! target/oneping*.oar`

Code & deploy OnePingOnly app



- “Code” our app as `OnePing.java`
- Rebuild the app via Maven and redeploy it

```
$ mvn clean install
```

```
$ onos-app 10.128.11.1 reinstall! target/oneping*.oar
```

Show one-ping-only enforcement



- Issue one ping between two hosts

```
mininet> h1 ping -c1 h2
```

```
mininet> h1 ping -c1 h3
```

- Issue another ping between the same two hosts

```
mininet> h1 ping -c1 h2
```

- Show failed second ping and messages in the logs

ONOS app overlay archetypes



- Some ONOS archetypes can be used to enhance an existing project with additional capabilities
 - add CLI extensions
 - add GUI extensions
 - others in future
- Simply run **onos-create-app** tool in the same project
- First show that sample command does not exist
onos> sample

Overlay CLI archetype



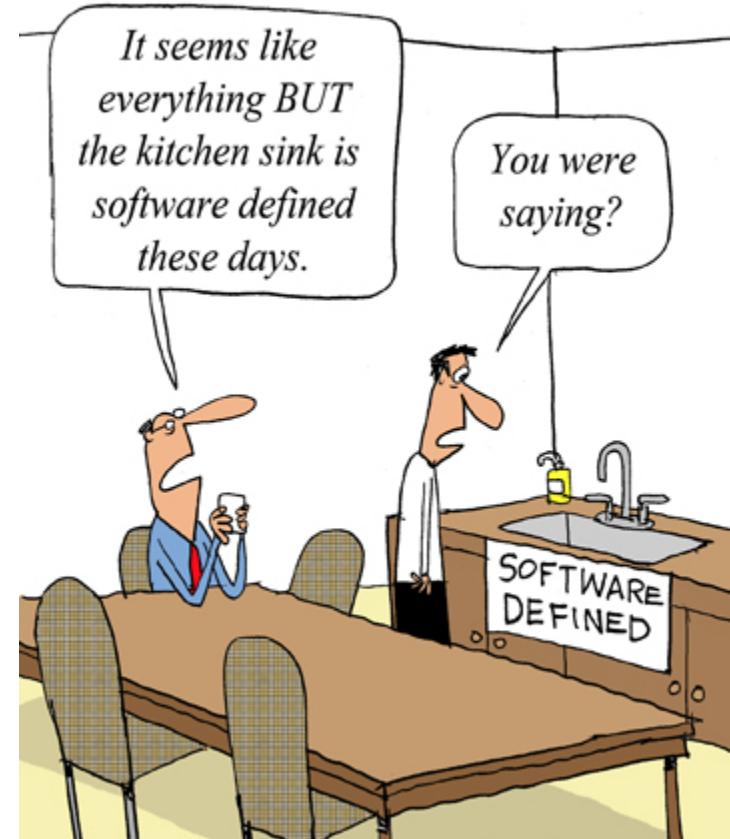
- Overlay CLI additions via `onos-create-app`
\$ `onos-create-app cli org.oneping oneping 1.0-SNAPSHOT`
- Rebuild and redeploy the enhanced app
\$ `mvn clean install`
\$ `onos-app 10.128.11.1 reinstall! target/oneping*.oar`
- Show that sample command now exists
`onos> sample`

Overlay GUI archetype



- Show existing GUI navigation menu
- Overlay GUI additions via `onos-create-app` and rebuild
 - \$ `onos-create-app ui org.oneping oneping 1.0-SNAPSHOT`
- Rebuild and redeploy the enhanced app
 - \$ `mvn clean install`
 - \$ `onos-app 10.128.11.1 reinstall! target/oneping*.oar`
- Refresh GUI and show enhanced GUI navigation menu

Break



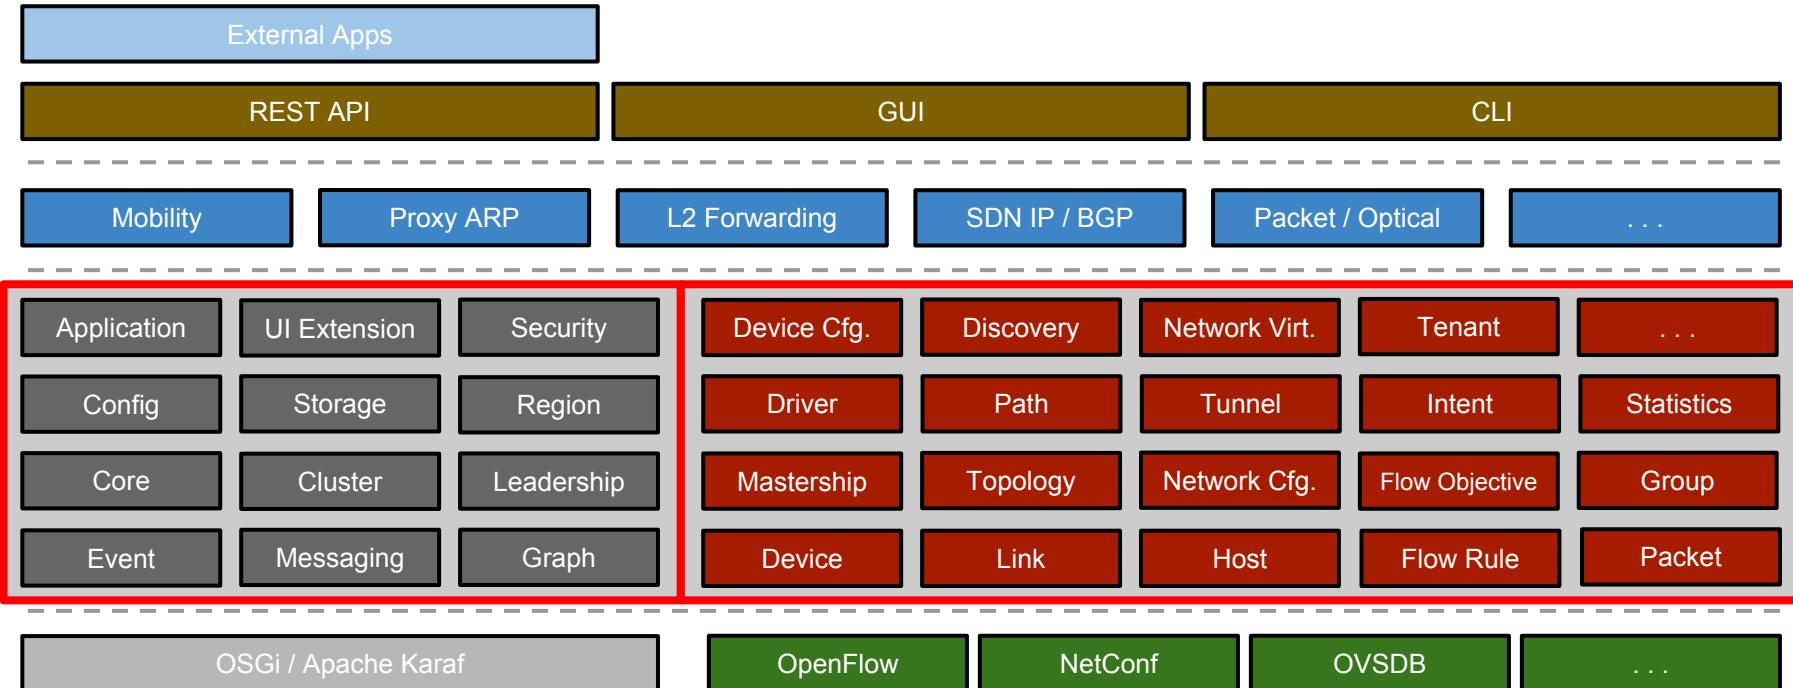
ONOS Interfaces

15:15-17:00



- Component Config Service (~10)
- Device Drivers Service (~10)
- Flow Objective Service (~15)
- Intent Service (~15)
- GUI overview (~15)
- CLI overview (~5)
- REST API overview (~5)
- Wrap-up & Final QA (~15)

ONOS Core Subsystems



ONOS Interfaces

15:15-17:00



- Component Config Service (~10)
- Device Drivers Service (~10)
- Flow Objective Service (~15)
- Intent Service (~15)
- GUI overview (~15)
- CLI overview (~5)
- REST API overview (~5)
- Wrap-up & Final QA (~15)

Component Config Subsystem



- Provide centralized and cluster-wide component configuration
- Avoid drift between configurations on cluster instances
- Built atop the OSGi **ConfigurationAdmin** service
- Components accept configuration via standard OSGi **@Modified** hook
- Catalogue of config properties for each component generated at build-time via **onos-maven-plugin**

Component Config Subsystem



- Component registers its properties on activation

```
void registerProperties(getClass());
```

- Component unregisters its properties on deactivation

```
void unregisterProperties(getClass(), false);
```

- Component provides configuration method hook

```
@Modified
```

```
public void modified(ComponentContext context) { ... }
```

ONOS Interfaces

15:15-17:00



- Component Config Service (~10)
- **Device Drivers Service** (~10)
- Flow Objective Service (~15)
- Intent Service (~15)
- GUI overview (~15)
- CLI overview (~5)
- REST API overview (~5)
- Wrap-up & Final QA (~15)

Device Driver Subsystem



- Isolate device-specific code
- Allow asynchronous delivery & dynamic loading
 - by ON.Lab, vendors or others
- Allow abstracting of how to interact with devices
 - via Java API (**Behaviour** interfaces)
- Avoid a monolithic driver approach
 - different facets of behaviour can come from different sources
- Facilitate specialization via inheritance
 - families of devices share many similar characteristics

Driver



- Representation of a specific family of devices or a specific device
 - has a unique name, e.g. **ciena-6500**
 - supports set of **Behaviour** classes
 - may *inherit* behaviours from another **Driver**
- **DefaultDriver** is a default implementation class

DriverAdminService



- Service for tracking driver providers

`Set<DriverProvider> getProviders()`

`registerProvider(DriverProvider)`

`unregisterProvider(DriverProvider)`

- `DriverProvider` provides a set of driver definitions

`Set<Driver> getDrivers()`

DriverService



Service for locating appropriate drivers for the device

- by driver name
- by manufacturer, H/W version & S/W version patterns
- by supported behaviour
- by device ID

Behaviour & HandlerBehaviour



- **Behaviour** is an interface for talking about a device
- **HandlerBehaviour** is an interface for interacting with a device
 - **OpenFlowController**
 - **OpenFlowPipeliner**
 - ...

DriverData & DriverHandler



- **DriverData** is a container for data about a device
 - provides behaviours for talking about a device
 - has parent driver
- **DriverHandler** is a context for interacting with a device
 - provides behaviours for interacting with a device
 - has **DriverData**
 - has parent driver

ONOS Interfaces

15:15-17:00

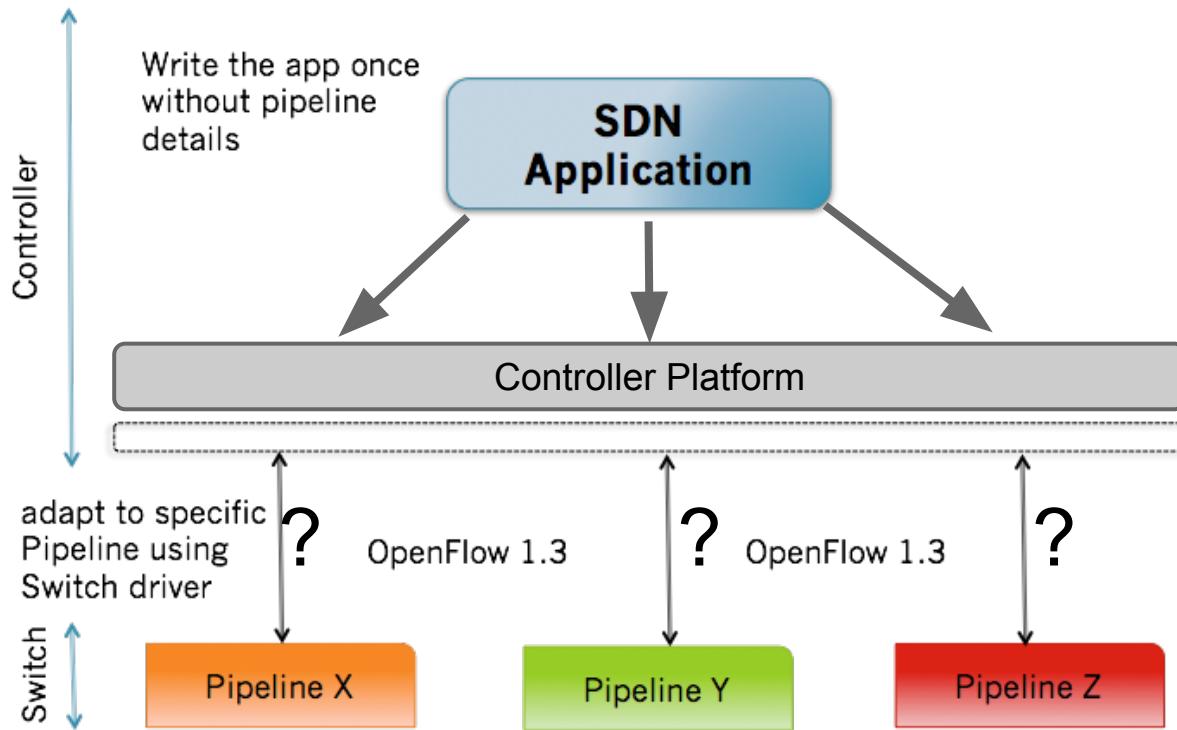


- Component Config Service (~10)
- Device Drivers Service (~10)
- Flow Objective Service (~15)
- Intent Service (~15)
- GUI overview (~15)
- CLI overview (~5)
- REST API overview (~5)
- Wrap-up & Final QA (~15)

Flow Objective Subsystem



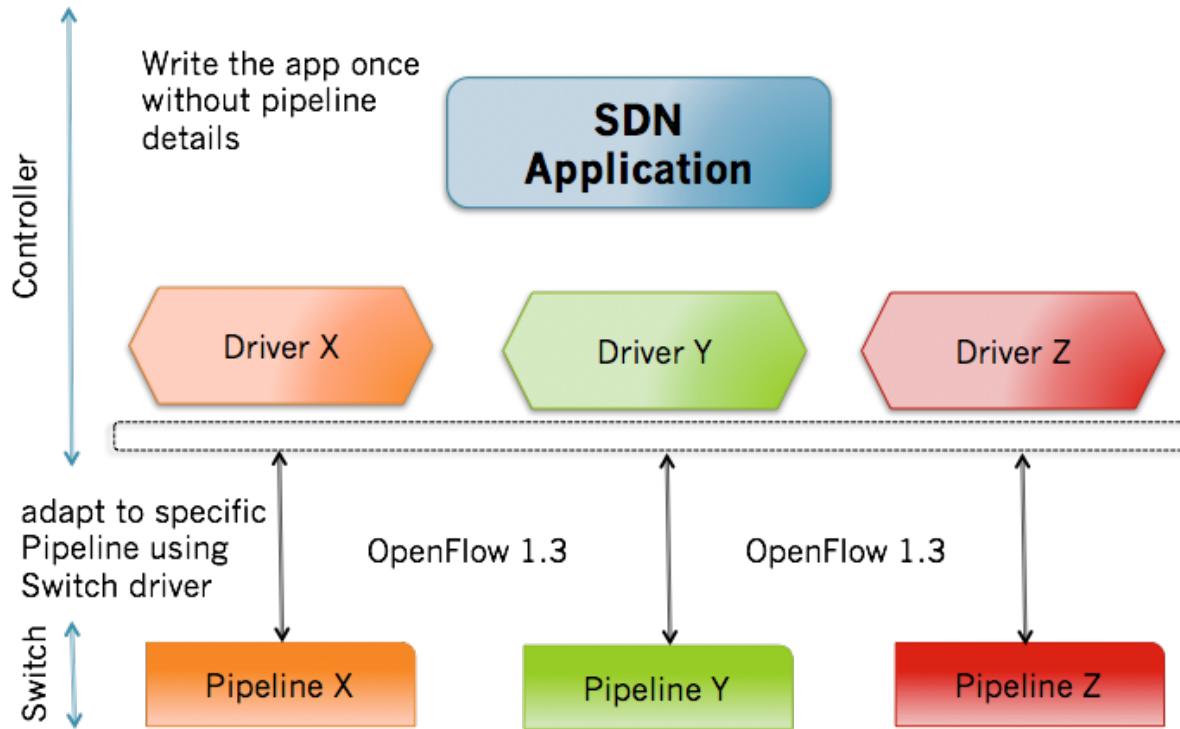
- **Problem:** Applications today must be pipeline aware, effectively making them applicable to specific HW.



Flow Objective Subsystem



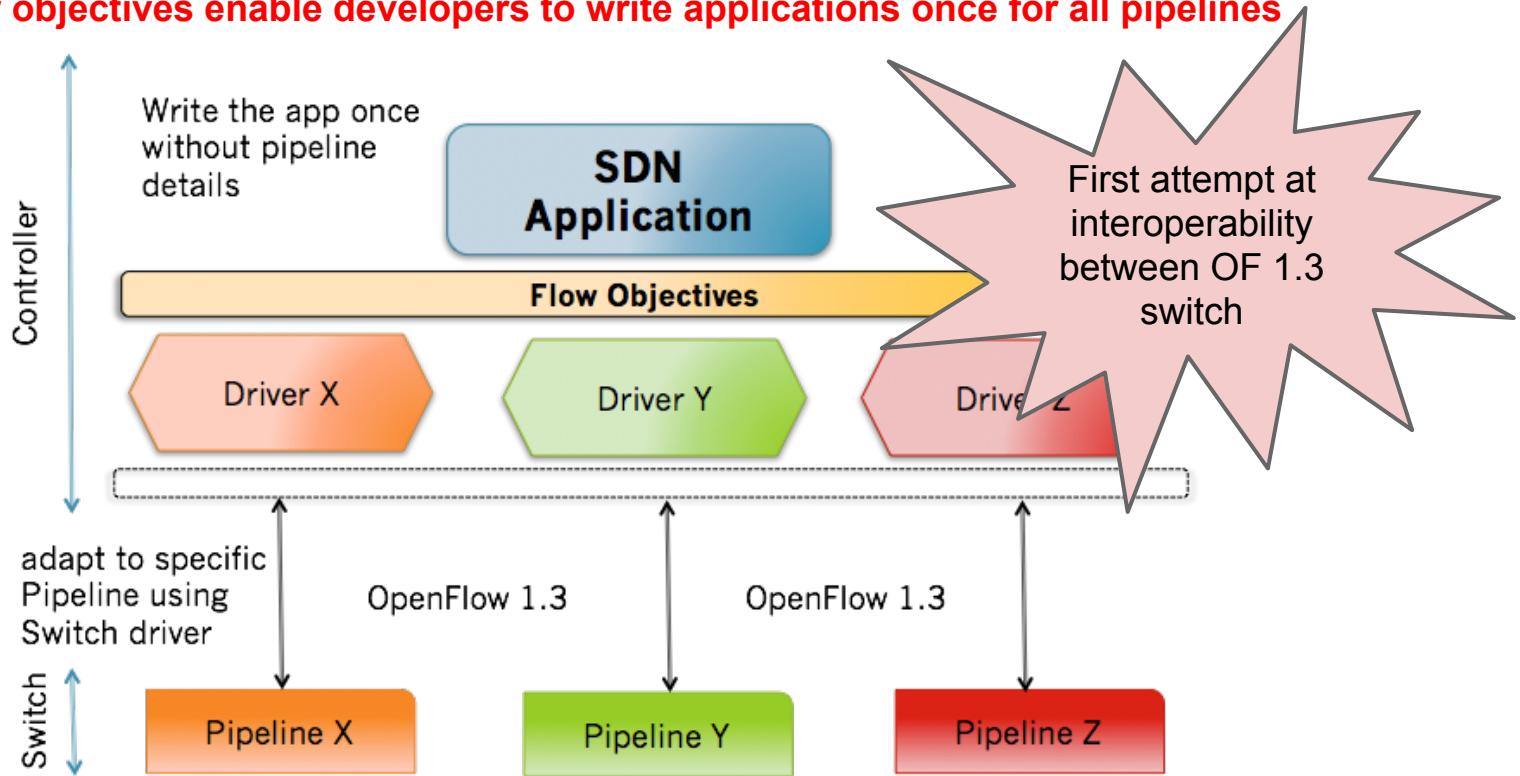
- **Problem:** Applications today must be pipeline aware, effectively making them applicable to specific HW.



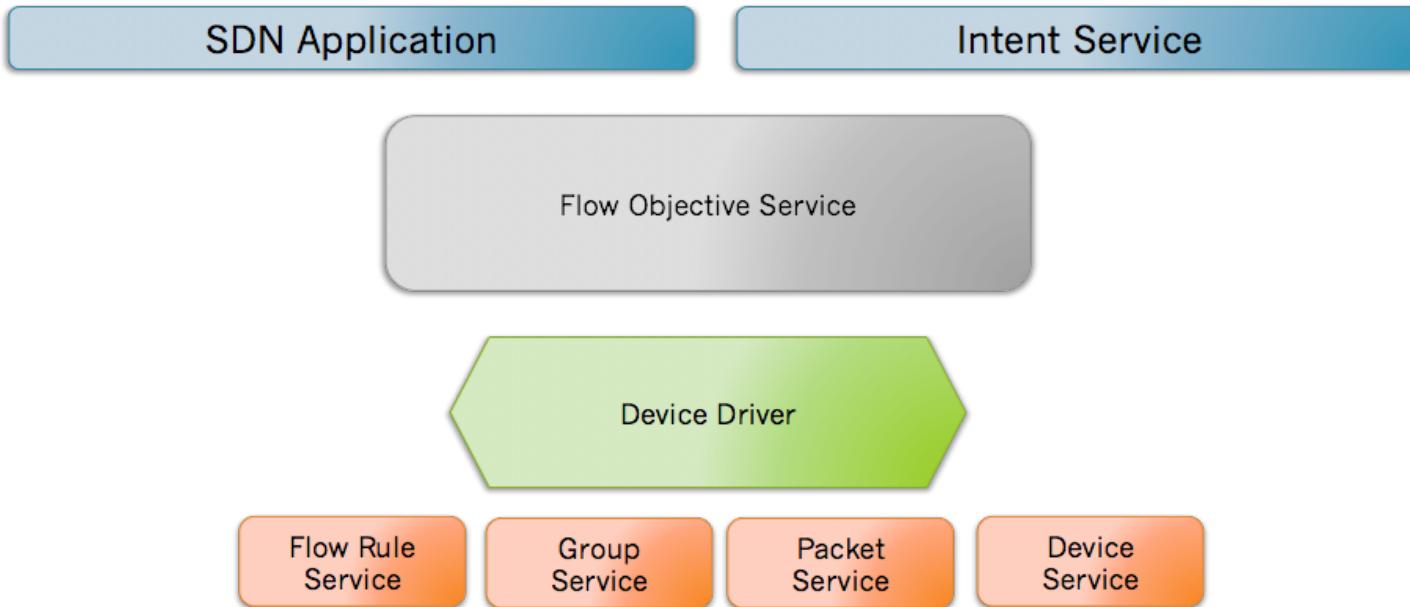
Flow Objective Abstraction



- **Problem:** Applications currently must be pipeline aware, effectively making applicable on specific HW.
Flow objectives enable developers to write applications once for all pipelines



Flow Objective Service



- Applications use Objective to take advantage multi-table architectures
- Other services also make use of the Objective service (eg. Intent Service)
- Device driver translates objectives to the specific flow rules for a given device

Flow Objectives



- Flow Objectives describe a SDN application's objective behind a **flow** it is sending to a **device**
- We currently only have three types of objectives:
 1. Filtering Objective
 2. Forwarding Objective
 3. Next Objective

Filtering Objective



- Filter -> only Permit or Deny options
- On criteria (match fields)

Example:

Peering Router

Switch Port : X

Permit: MAC 1, VLAN 1, IP 1, 2, 3

Permit: MAC 1, VLAN 2, IP 4, 5

Flow Objectives

Corsa Driver

T0
mac

T2
port-vlan

T6
ip

OF-DPA Driver

T0
port

T1
Port
-vlan

T2
mac

T4
ip

Filtering Objective



- ☐ Filter -> only Permit or Deny options
- ☐ On criteria (match fields)

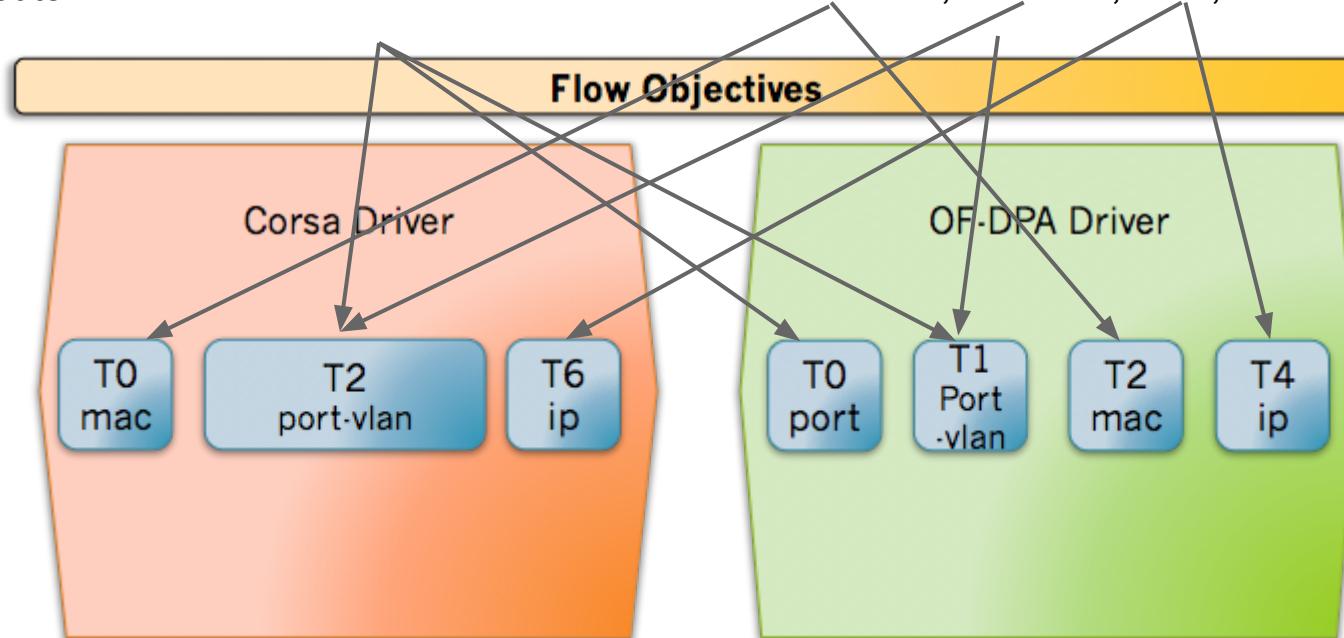
Example:

Peering Router

Switch Port : X

Permit: MAC 1, VLAN 1, IP 1, 2, 3

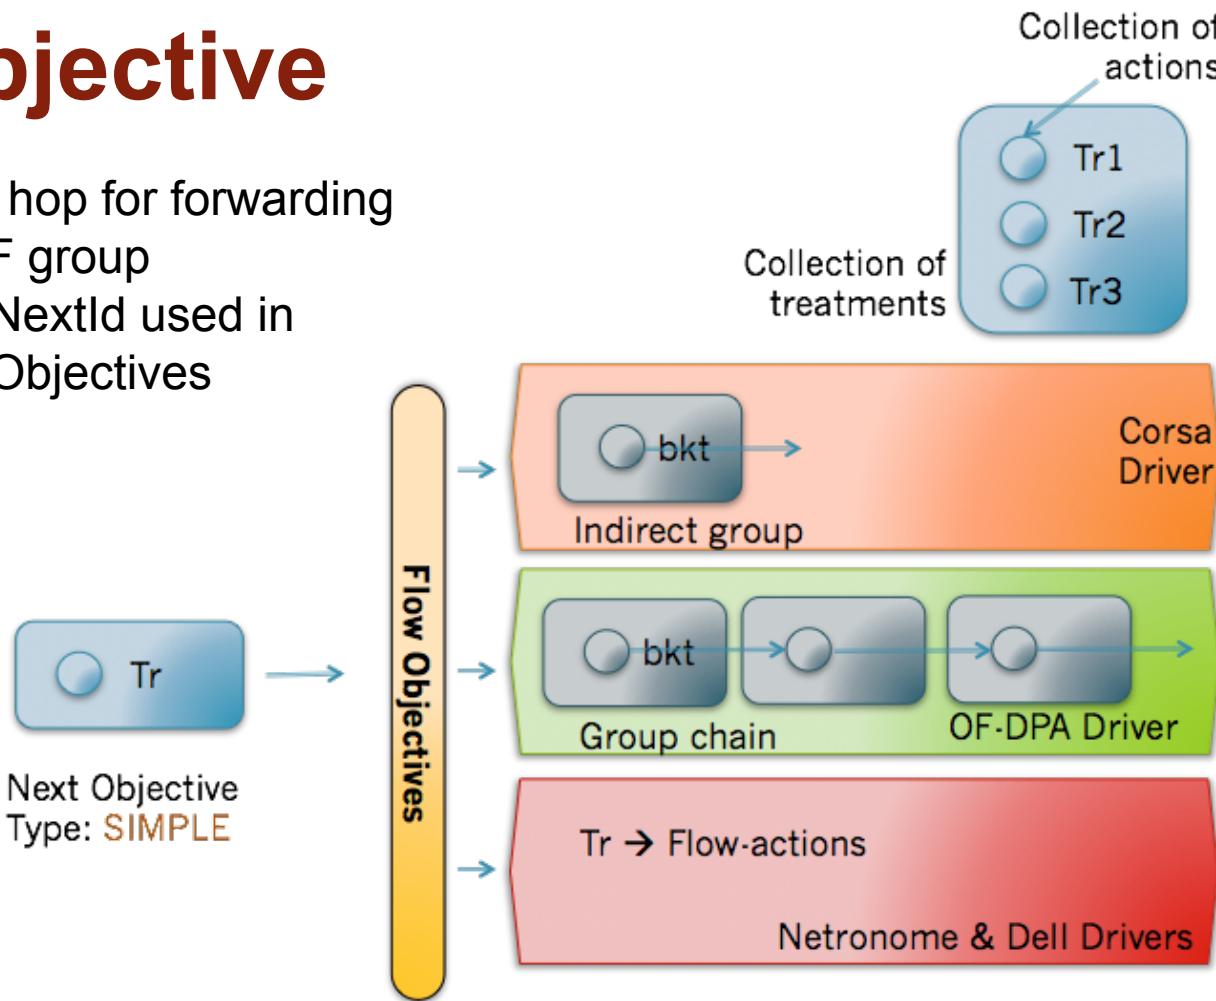
Permit: MAC 1, VLAN 2, IP 4, 5



Next Objective



- Next -> next hop for forwarding
- Similar to OF group
- Keyed by a NextId used in Forwarding Objectives



Forwarding Objective



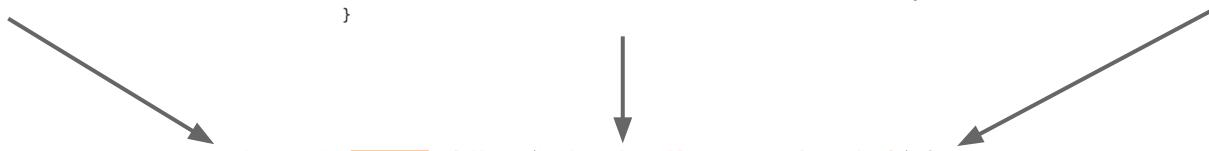
- Forwarding: { Selector -> Next Id }
- Forwarding Types: Specific or Versatile
 - Specific -> MAC, IP, MPLS forwarding tables
 - Versatile -> ACL table
- NextId is resolved to whatever the driver previously built for the corresponding Next Objective

Objectives - Simpler applications

```
for (InterfaceIpAddress ipAddr : intfIps) {  
    log.debug("adding rule for IPs: {}", ipAddr.ipAddress());  
    selector = DefaultTrafficSelector.builder();  
    treatment = DefaultTrafficTreatment.builder();  
  
    selector.matchEthType(Ethernet.TYPE_IPV4);  
    selector.matchIPDst(ipPrefix.valueOf(ipAddr.ipAddress(), 32));  
    treatment.transition(Type.ACL);  
  
    rule = new DefaultFlowRule(deviceId, selector.build(),  
        treatment.build(), HIGHEST_PRIORITY, appId,  
        0, true, FlowRule.Type.IP);  
  
    ops = install ? ops.add(rule) : ops.remove(rule);  
}
```

```
for (MacAddress mac : intfMacs) {  
    log.debug("adding rule for MAC: {}", mac);  
    selector = DefaultTrafficSelector.builder();  
    treatment = DefaultTrafficTreatment.builder();  
  
    selector.matchEthDst(mac);  
    treatment.transition(FlowRule.Type.VLAN_MPLS);  
  
    rule = new DefaultFlowRule(deviceId, selector.build(),  
        treatment.build(),  
        CONTROLLER_PRIORITY, appId, 0,  
        true, FlowRule.Type.FIRST);  
  
    ops = install ? ops.add(rule) : ops.remove(rule);  
}
```

```
for (Map.Entry<PortNumber, VlanId> portVlan : portVlanPair.entrySet()) {  
    log.debug("adding rule for VLAN: {}", portVlan);  
    selector = DefaultTrafficSelector.builder();  
    treatment = DefaultTrafficTreatment.builder();  
  
    selector.matchVlanId(portVlan.getValue());  
    selector.matchInPort(portVlan.getKey());  
    treatment.transition(Type.ETHER);  
    treatment.deferred().popVlan();  
  
    rule = new DefaultFlowRule(deviceId, selector.build(),  
        treatment.build(), CONTROLLER_PRIORITY, appId,  
        0, true, FlowRule.Type.VLAN);  
  
    ops = install ? ops.add(rule) : ops.remove(rule);  
}
```



```
private void processIntfFilters(boolean install, Set<Interface> intfs) {  
    log.info("Processing {} router interfaces", intfs.size());  
    for (Interface intf : intfs) {  
        FilteringObjective.Builder fob = DefaultFilteringObjective.builder();  
        fob.withKey(Criteria.matchInPort(intf.connectPoint().port()))  
            .addCondition(Criteria.matchEthDst(intf.mac()))  
            .addCondition(Criteria.matchVlanId(intf.vlan()));  
        intf.ipAddresses().stream()  
            .forEach(ipaddr -> fob.addCondition(  
                Criteria.matchIPDst(ipaddr.subnetAddress())));  
        fob.permit().fromApp(appId);  
        flowObjectiveService.filter(deviceId,  
            Collections.singletonList(fob.add()));  
    }  
}
```

Flow Objective Summary



- *Flow Objective Service*: **Abstraction** for applications to be **pipeline unaware** while **benefiting** from scalable, multi-table architectures
- Aims to make it **simple** to write apps
- First attempt at achieving **interoperability** between OF 1.3 implementations

ONOS Interfaces

15:15-17:00



- Component Config Service (~10)
- Device Drivers Service (~10)
- Flow Objective Service (~15)
- Intent Service (~15)
- GUI overview (~15)
- CLI overview (~5)
- REST API overview (~5)
- Wrap-up & Final QA (~15)

Building Network Applications



Objective: Connect Host 1 and Host 2



Host 1



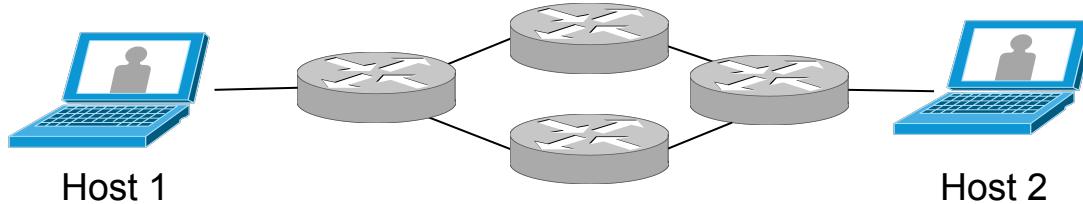
Host 2

Building Network Applications



Objective: Connect Host 1 and Host 2

1. Read/discover the topology

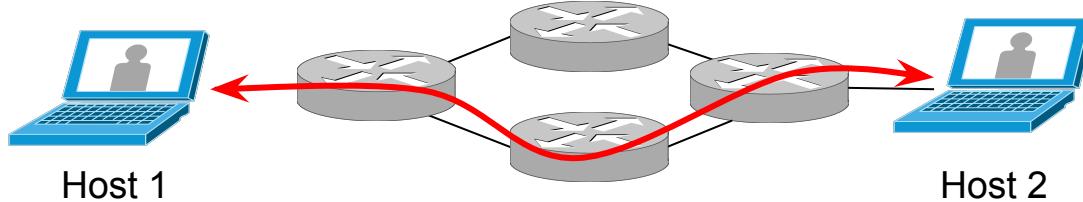


Building Network Applications



Objective: Connect Host 1 and Host 2

1. Read/discover the topology
2. Compute a path

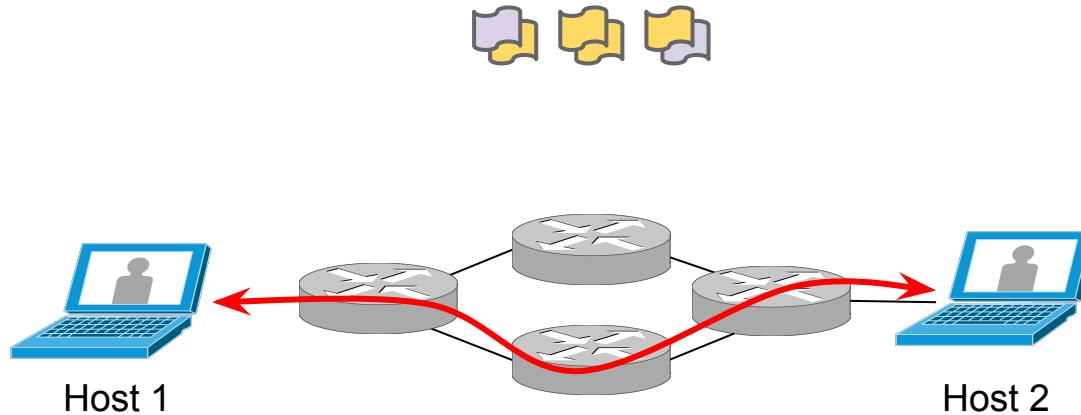


Building Network Applications



Objective: Connect Host 1 and Host 2

1. Read/discover the topology
2. Compute a path
3. Build flow objectives for each device

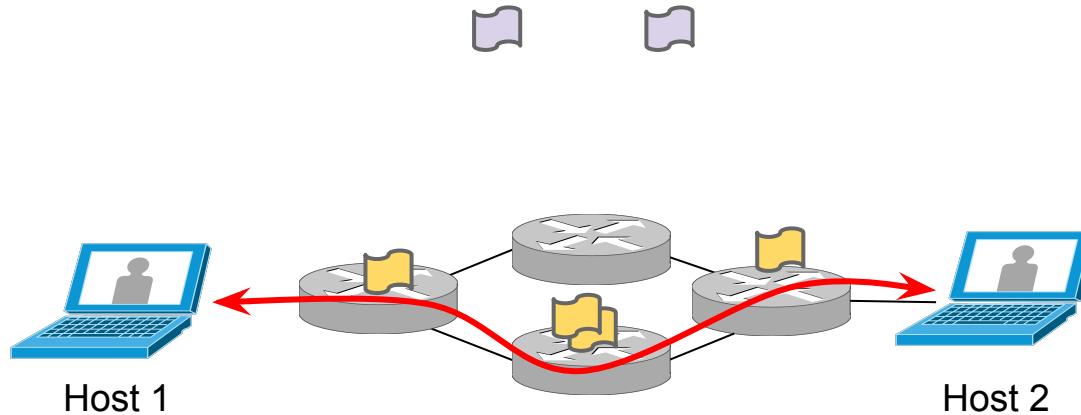


Building Network Applications



Objective: Connect Host 1 and Host 2

1. Read/discover the topology
2. Compute a path
3. Build flow objectives for each device
4. Install rules in consistent way

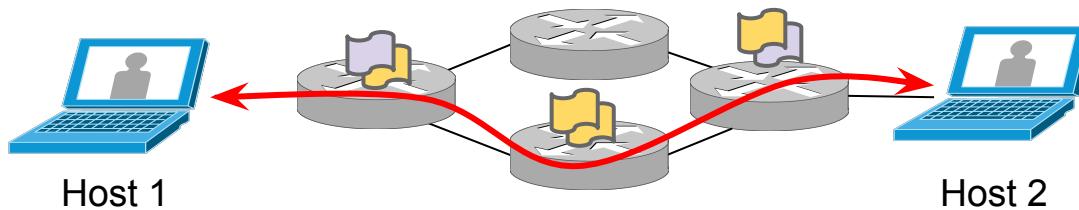


Building Network Applications



Objective: Connect Host 1 and Host 2

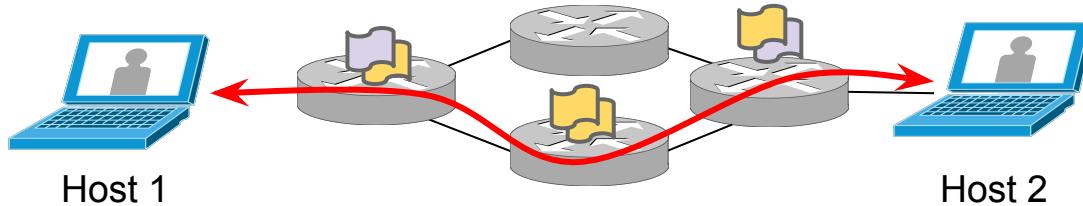
1. Read/discover the topology
2. Compute a path
3. Build flow objectives for each device
4. Install rules in consistent way



Building Network Applications



What can go wrong?

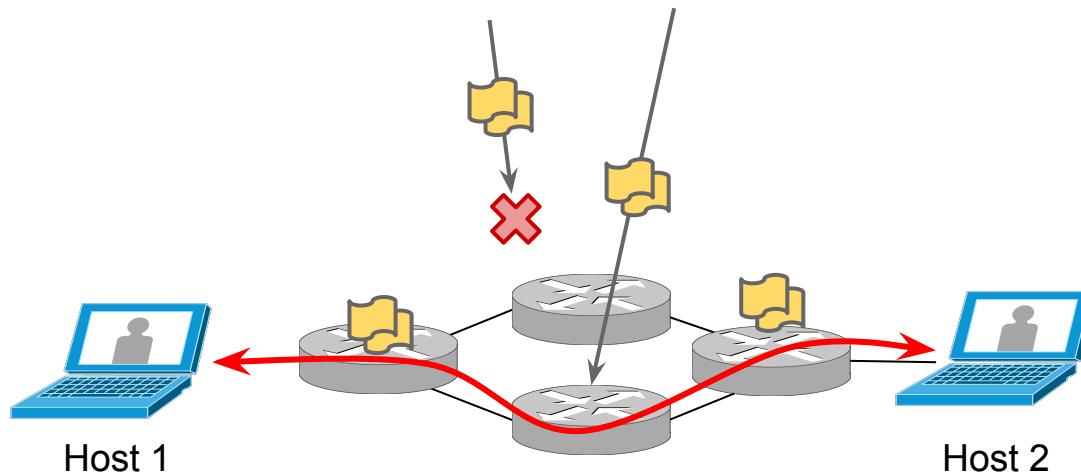


Building Network Applications



What can go wrong?

- Missing / rejected / dropped rules
 - Monitor devices connections
 - Send barriers between rule updates
 - Poll flow state

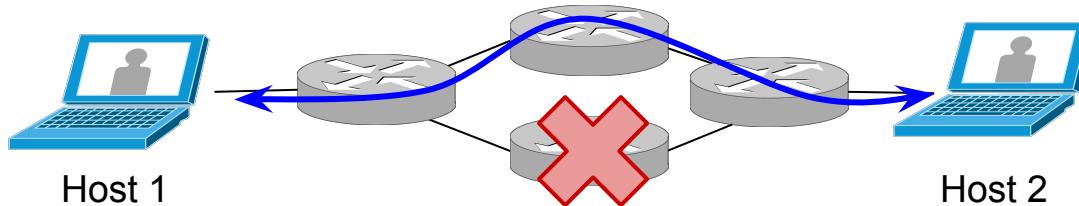


Building Network Applications



What can go wrong?

- Missing / rejected / dropped rules
 - Monitor devices connections
 - Send barriers between rule updates
 - Poll flow state
- Topology changes
 - Listen to switch, port, link and host events
 - Compute new path that leverage or remove old flows



Building Network Applications



- Each application requires complex path computation and rule installation engines and state machines
- Inconsistent behavior in the face of failures
 - Failures may be handled in different ways (or not at all)
- Bugs need to be fixed in multiple places (applications)
- Expensive to upgrade/refactor behavior across all applications; e.g.
 - Improve performance
 - Support new types of devices
 - Implement better algorithms
- Difficult or impossible to resolve conflicts with other applications

Intent Subsystem

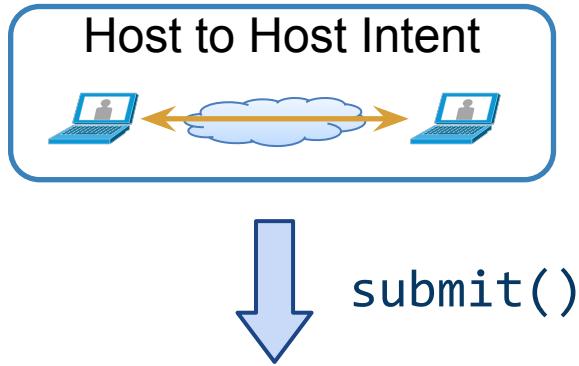


- Provides high-level, network-centric interface that focuses on what should be done rather than how it is specifically programmed
- Abstracts unnecessary network complexity from applications
- Maintains requested semantics as network changes
- High availability, scalability and high performance

Intent Example



Intent Example

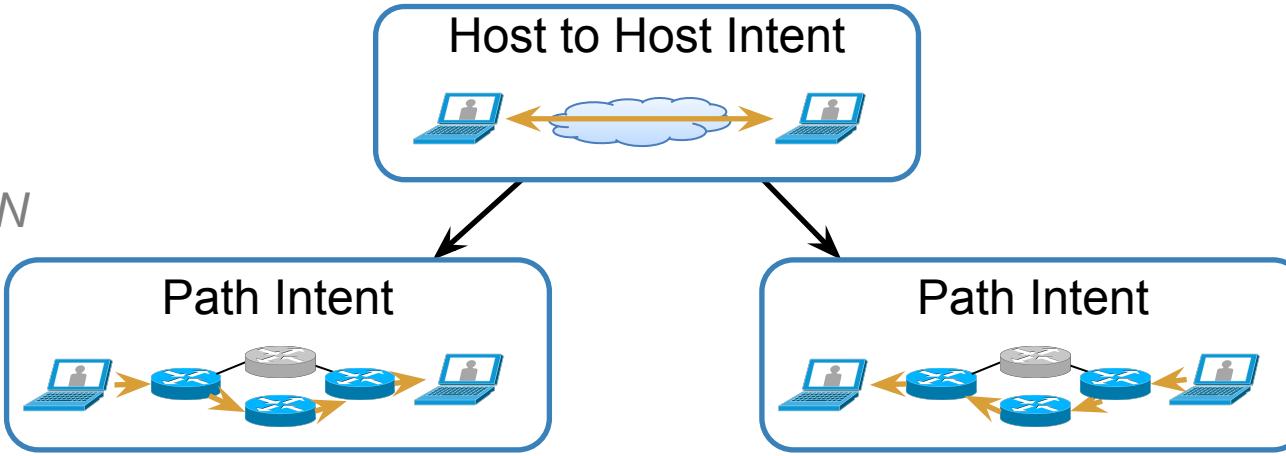


Intent Service API

Intent Example



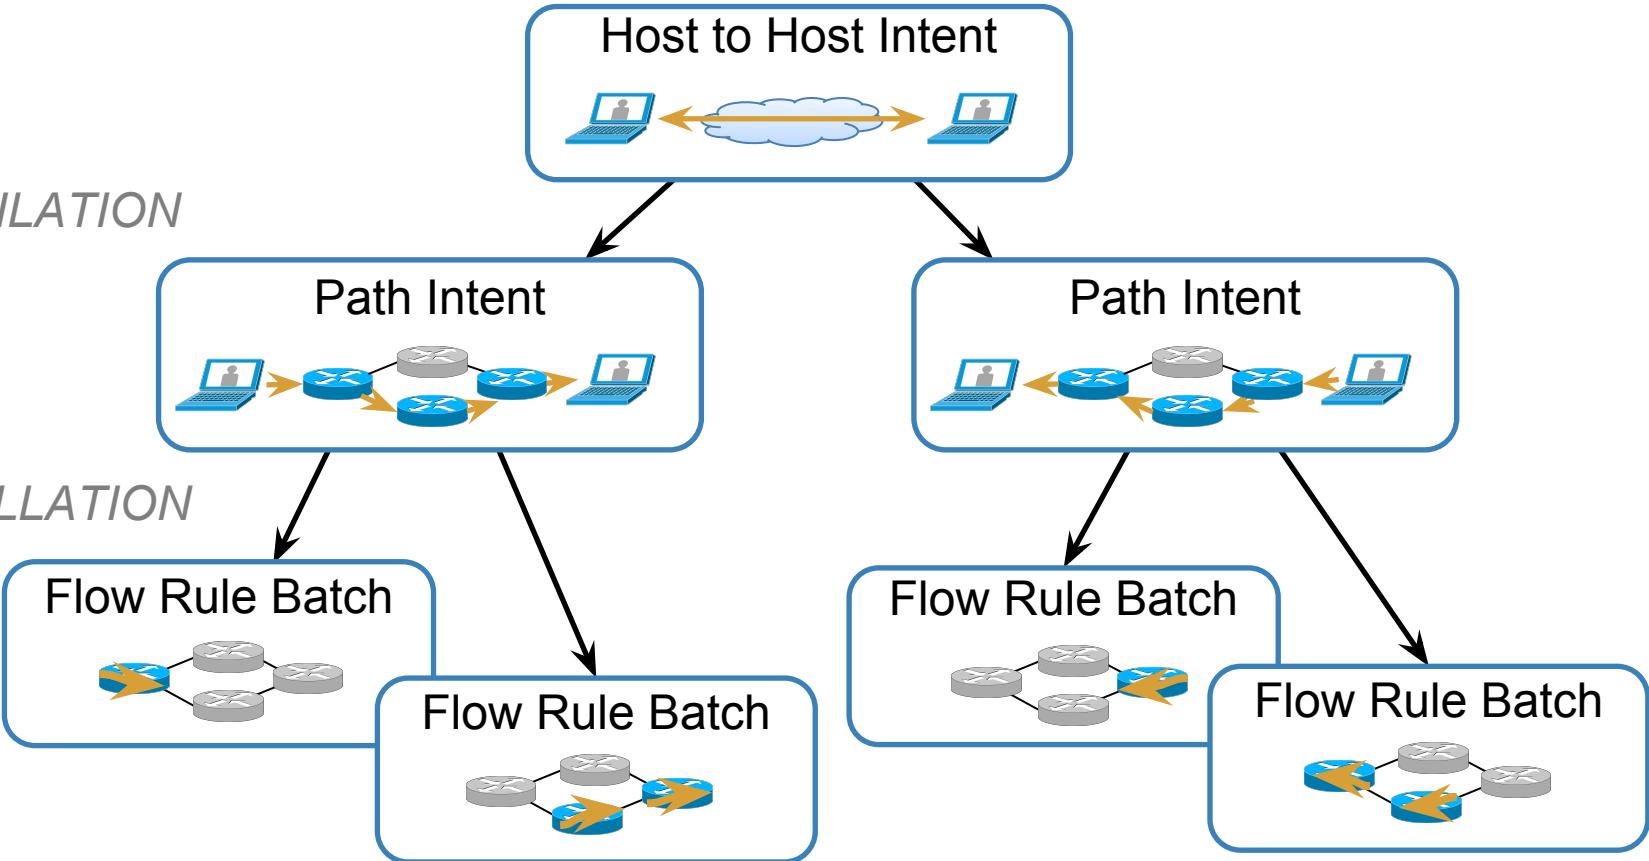
COMPILED



Intent Example



COMPILEMENT



Intent Compilers



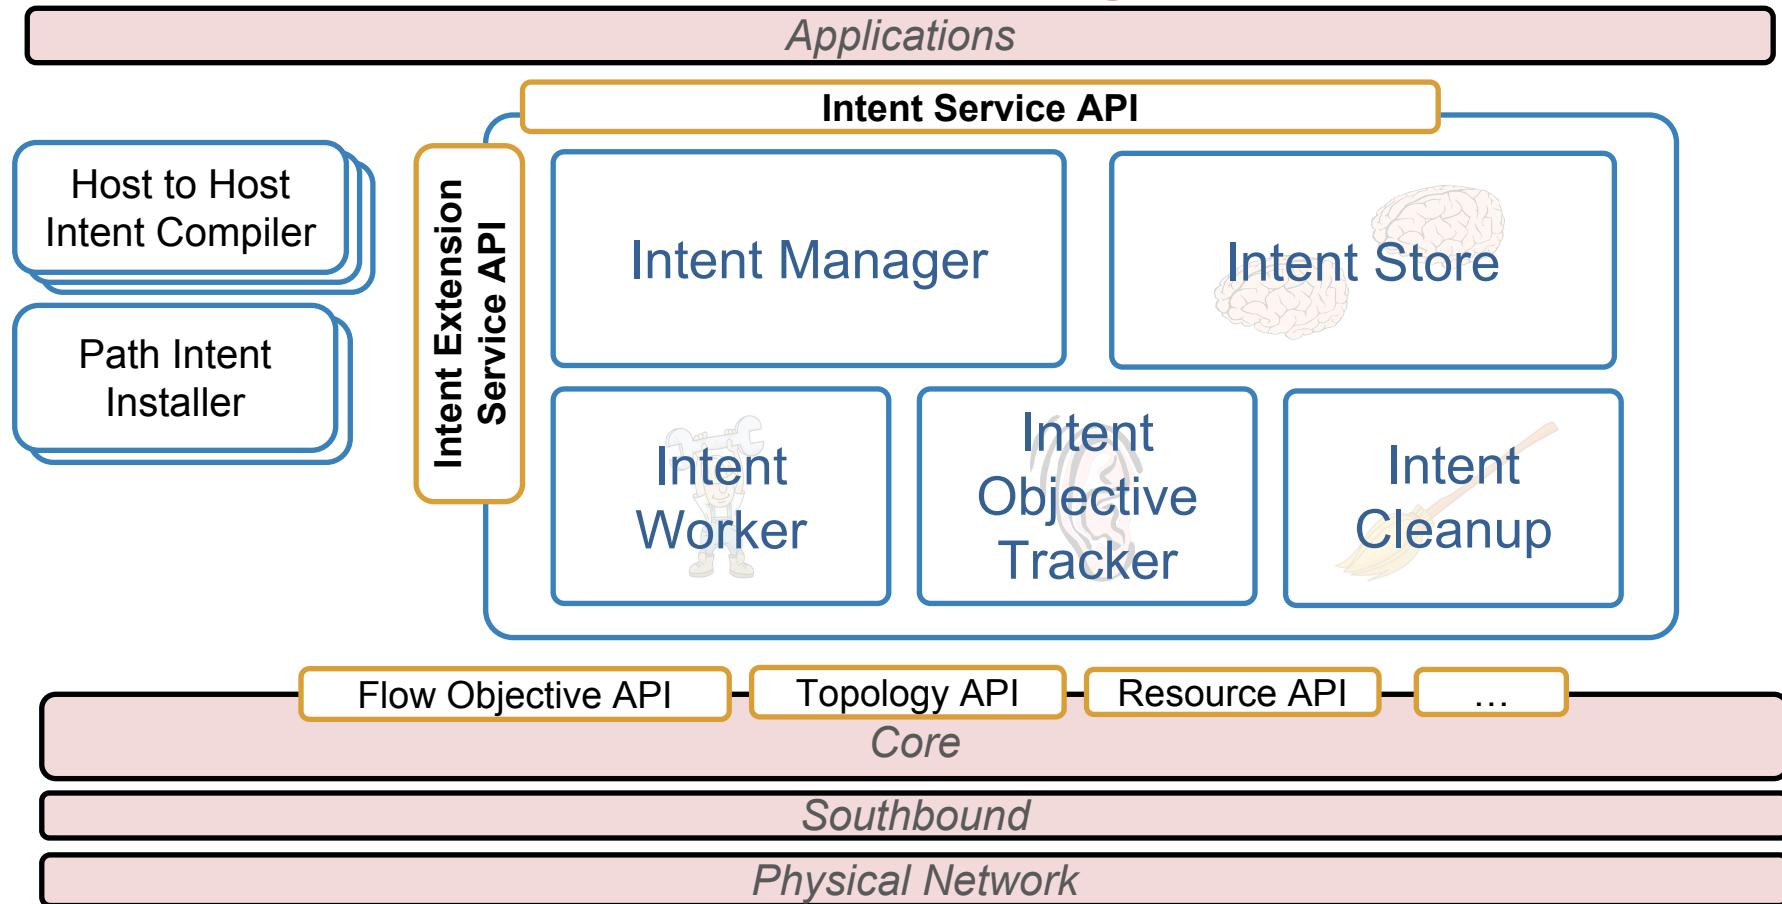
- Produce more specific Intents given the environment
- Works on high-level network objects, rather than device specifics
- “Stateless” – free from HA / distribution concerns

Intent Installers



- Transform Intents into device commands
- Allocate / deallocate network resources
- Define scheduling dependencies for workers
- “Stateless”

Intent Framework Design



ONOS Interfaces

15:15-17:00



- Component Config Service (~10)
- Device Drivers Service (~10)
- Flow Objective Service (~15)
- Intent Service (~15)
- **GUI overview** (~15)
- CLI overview (~5)
- REST API overview (~5)
- Wrap-up & Final QA (~15)

GUI Overview



- Single page application at `/onos/ui`
 - AngularJS
 - d3.js (Data-Driven Documents)
 - SVG
 - Web-Socket
- Shared web-socket for asynchronous & bidirectional exchange
- REST also used for select few interactions
- Dynamically extensible navigation & views
 - apps can contribute new views at run-time

ONOS Interfaces

15:15-17:00



- Component Config Service (~10)
- Device Drivers Service (~10)
- Flow Objective Service (~15)
- Intent Service (~15)
- GUI overview (~15)
- **CLI overview** (~5)
- **REST API overview** (~5)
- Wrap-up & Final QA (~15)

CLI Overview



- Built using Apache Karaf shell mechanism
 - `ssh -p8101 ...`
 - `onos` shell tool provided for convenience
- All commands in `onos`: namespace
- Most arguments support TAB-completion
- Plain-text output format for human readability
- JSON output format for easy scripting via `--json` option
- Noun oriented
 - a few verb-oriented exceptions will be addressed

REST API Overview



- Core REST APIs available under path `/onos/v1`
 - all resources support read (get)
 - not all resources support write (post/put/delete) yet
- Resources available today:

`/topology`

`/cluster`

`/devices`

`/applications`

`/links`

`/configuration`

`/hosts`

`/config`

`/paths`

`/flows`

`/intents`

ONOS Dev Workshop Outline



- **ONOS Fundamentals**
 - *architecture, checking out, building & running ONOS*
- **ONOS Distributed Core**
 - *OSGi, anatomy of a core subsystem, distributed stores & primitives*
- **ONOS App Development**
 - *background, Maven archetypes, app deployment on ONOS cluster*
- **ONOS Interfaces**
 - *component config service, flow objective service, intent service*
 - *CLI, REST and GUI overview*
- **Wrap-up**

Wrap-Up



- Browse ONOS Wiki
<https://wiki.onosproject.org/>
- Watch ONOS how-to screencasts on YouTube
<https://goo.gl/8Druv0>
- Browse ONOS Java API
<http://api.onosproject.org/>
- Join ONOS onos-dev@onosproject.org mailing list
<https://wiki.onosproject.org/display/ONOS/ONOS+Mailing+Lists>
- Engage with ONOS developers & community