

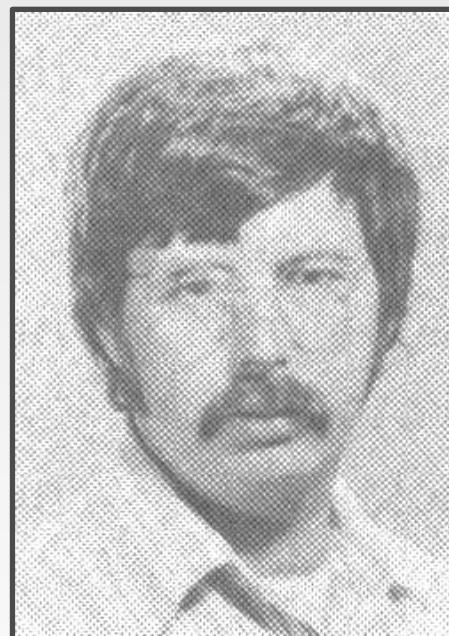
MY DATABASE SYSTEM IS
THE ONLY THING I CAN



TRUST

@ANDY_PAVLO

Thirty Years Ago...





INTERACTIVE TRANSACTIONS



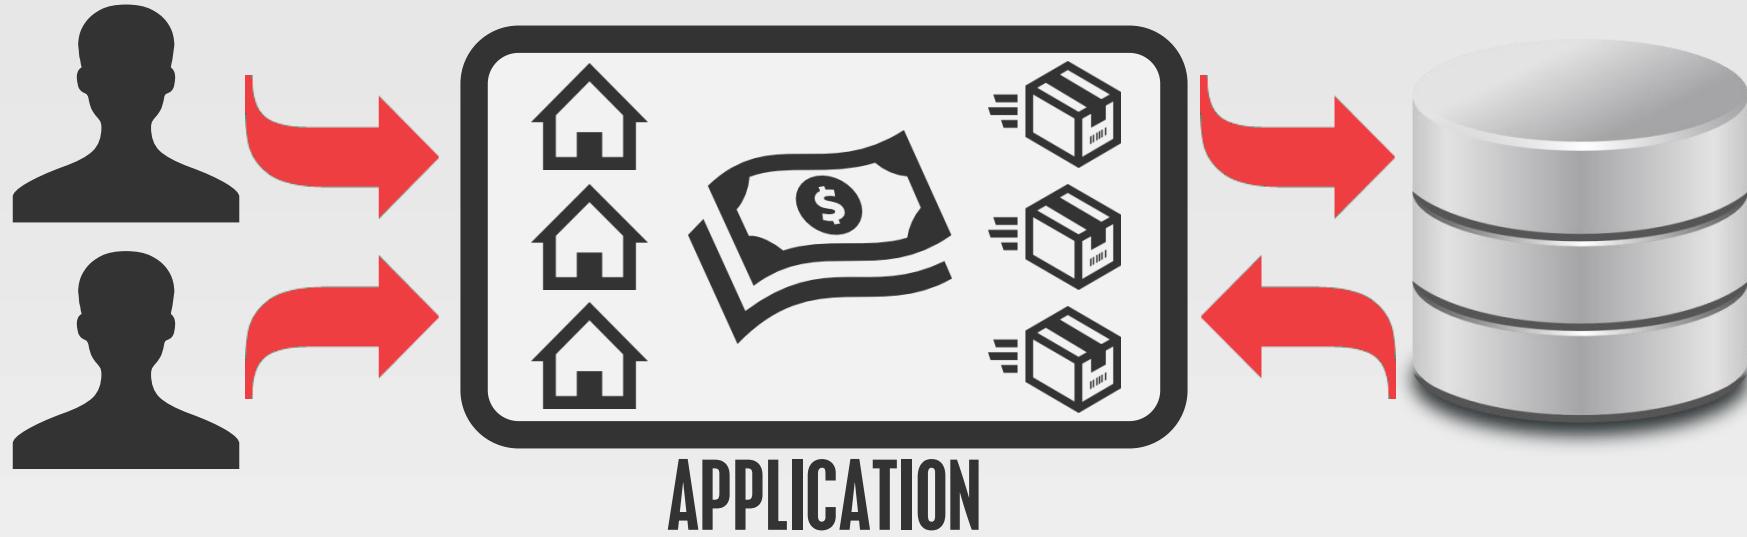
SMALL # OF CPU CORES



SMALL MEMORY SIZES

TPC-C BENCHMARK

Warehouse Order Processing

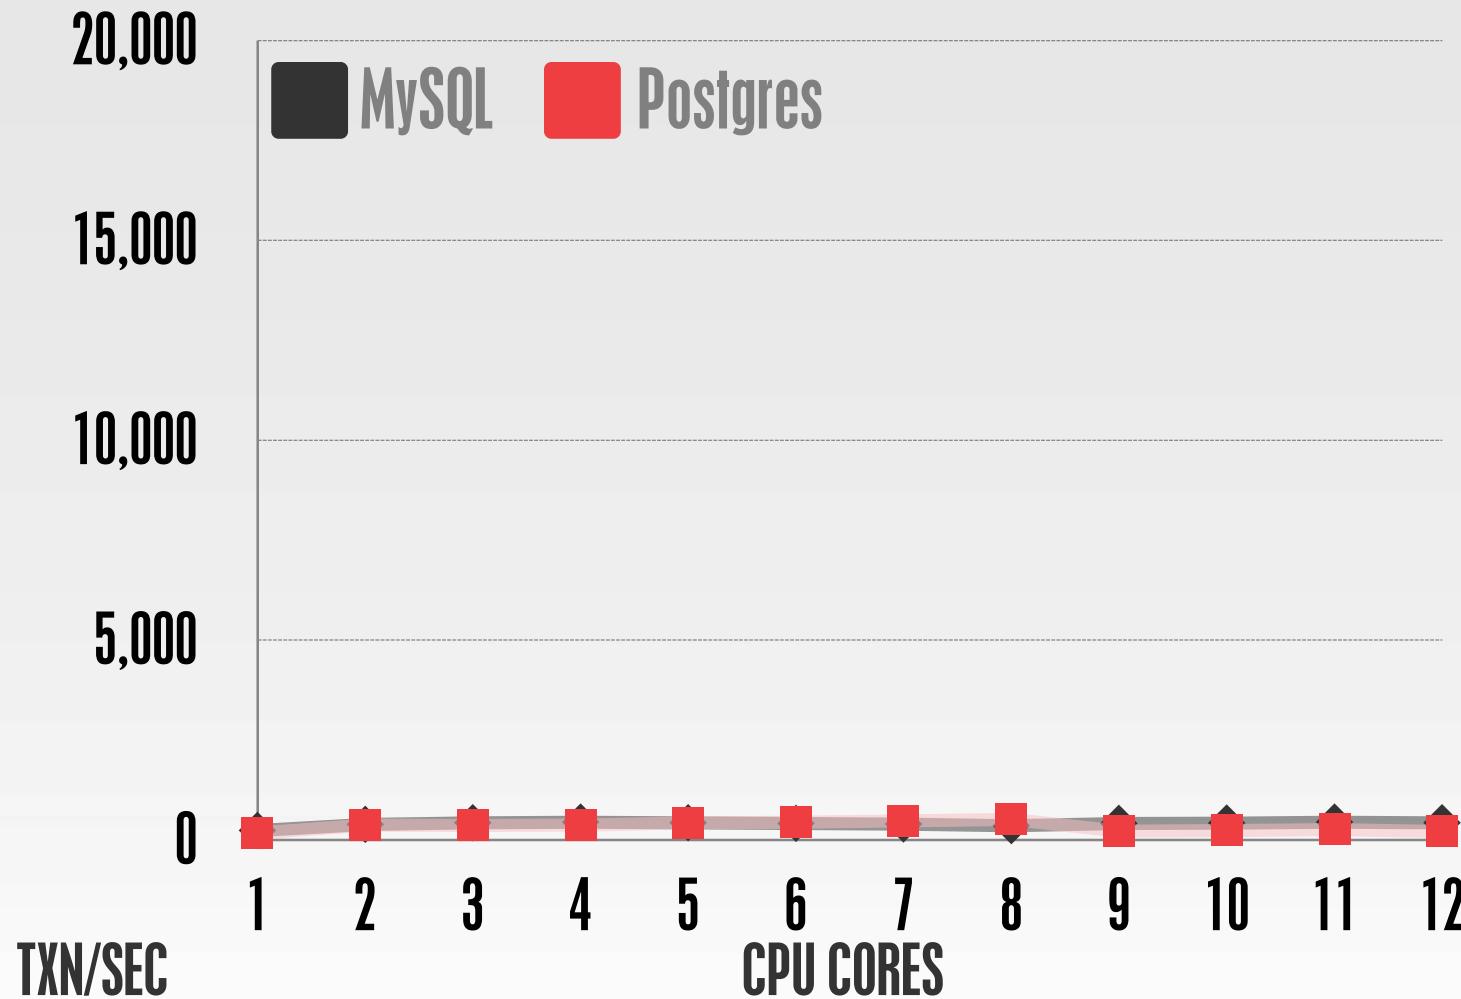


NewOrder Transaction

1. Check item stock level.
2. Create new order information.
3. Update item stock levels.

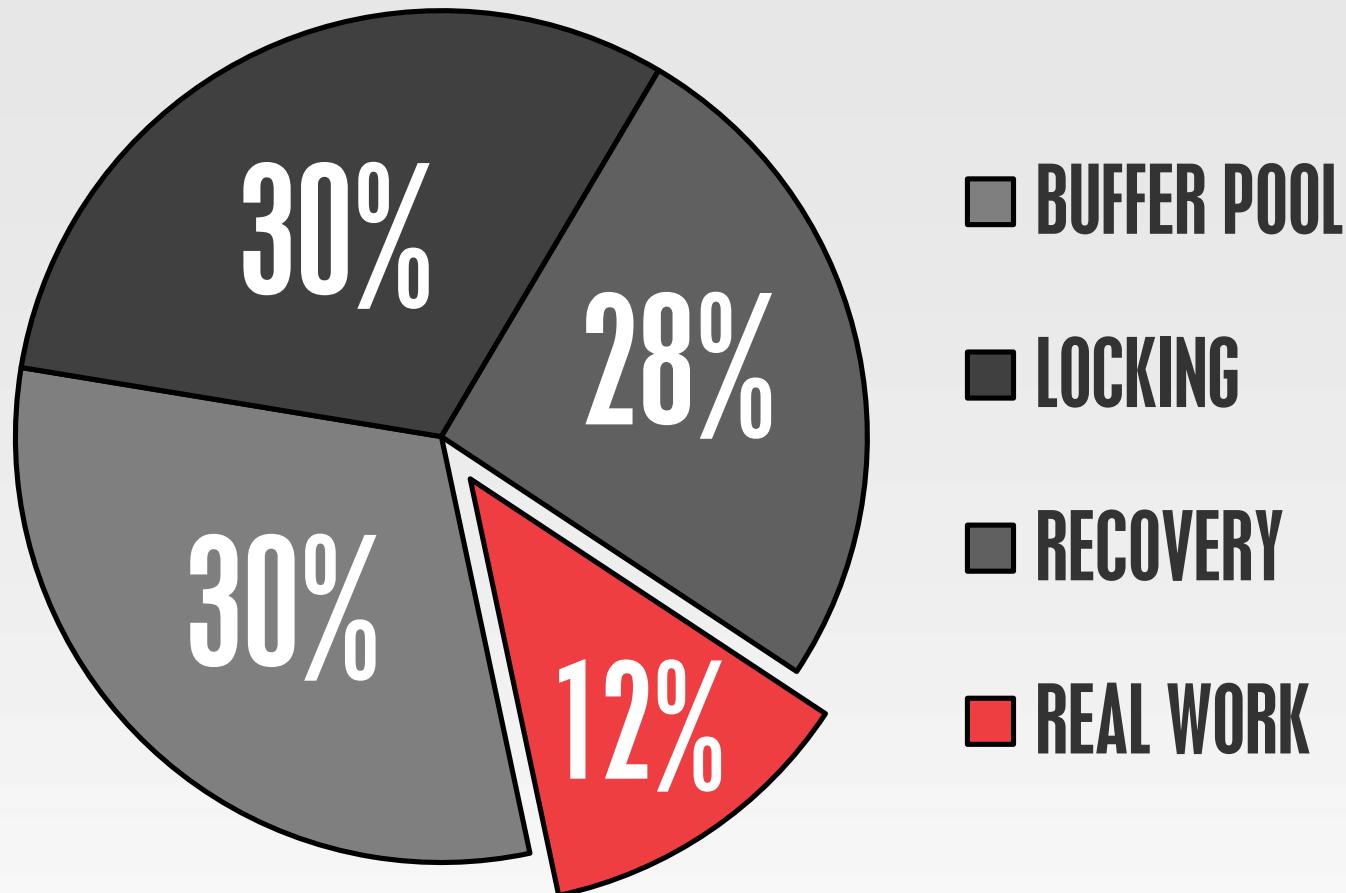
TPC-C BENCHMARK

Warehouse Order Processing



TRADITIONAL DBMS

Measured CPU Cycles



OLTP THROUGH THE LOOKING GLASS,
AND WHAT WE FOUND THERE
SIGMOD, pp. 981-992, 2008.





HARDWARE UPGRADE



REPLICATION



DISTRIBUTED CACHE



SHARDING MIDDLEWARE



NOSQL

**HOW TO SCALE UP
WITHOUT GIVING UP
TRANSACTIONS?**



Distributed Main Memory
Transaction Processing System

 H-STORE: A HIGH-PERFORMANCE, DISTRIBUTED
MAIN MEMORY TRANSACTION PROCESSING SYSTEM
Proc. VLDB Endow., vol. 1, iss. 2, pp. 1496-1499, 2008.



DISK ORIENTED
MAIN MEMORY STORAGE



CONCURRENT EXECUTION
SERIAL EXECUTION



HEAVYWEIGHT RECOVERY
COMPACT LOGGING

PARTITIONS



SINGLE-THREADED EXECUTION ENGINES

**Procedure Name
Input Parameters**



Application



STORED PROCEDURE

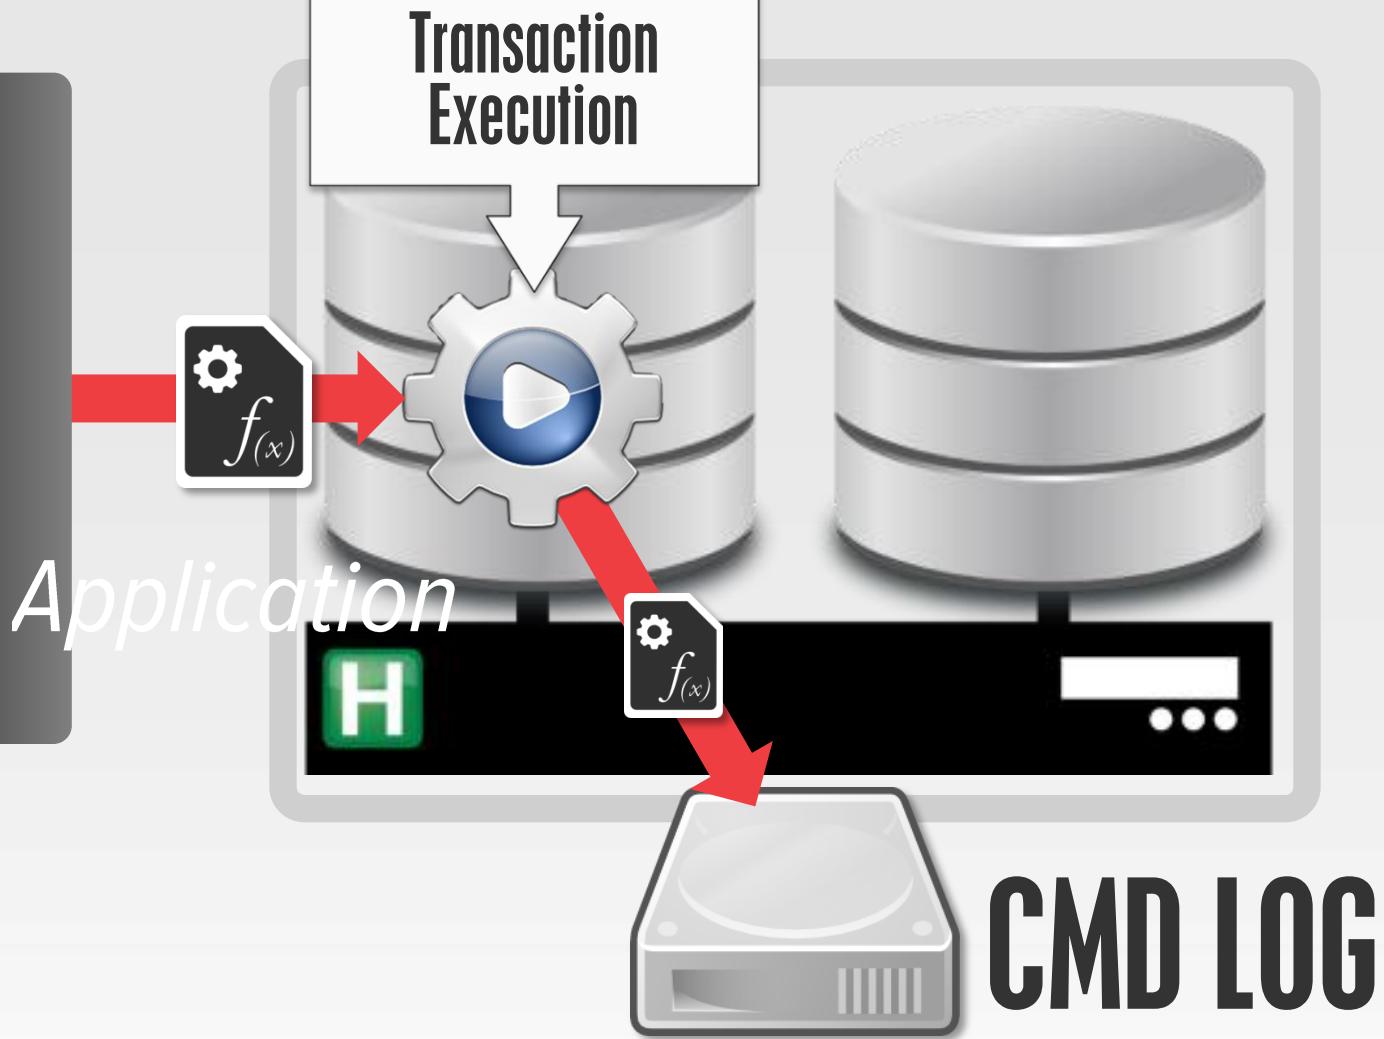
VoteCount:

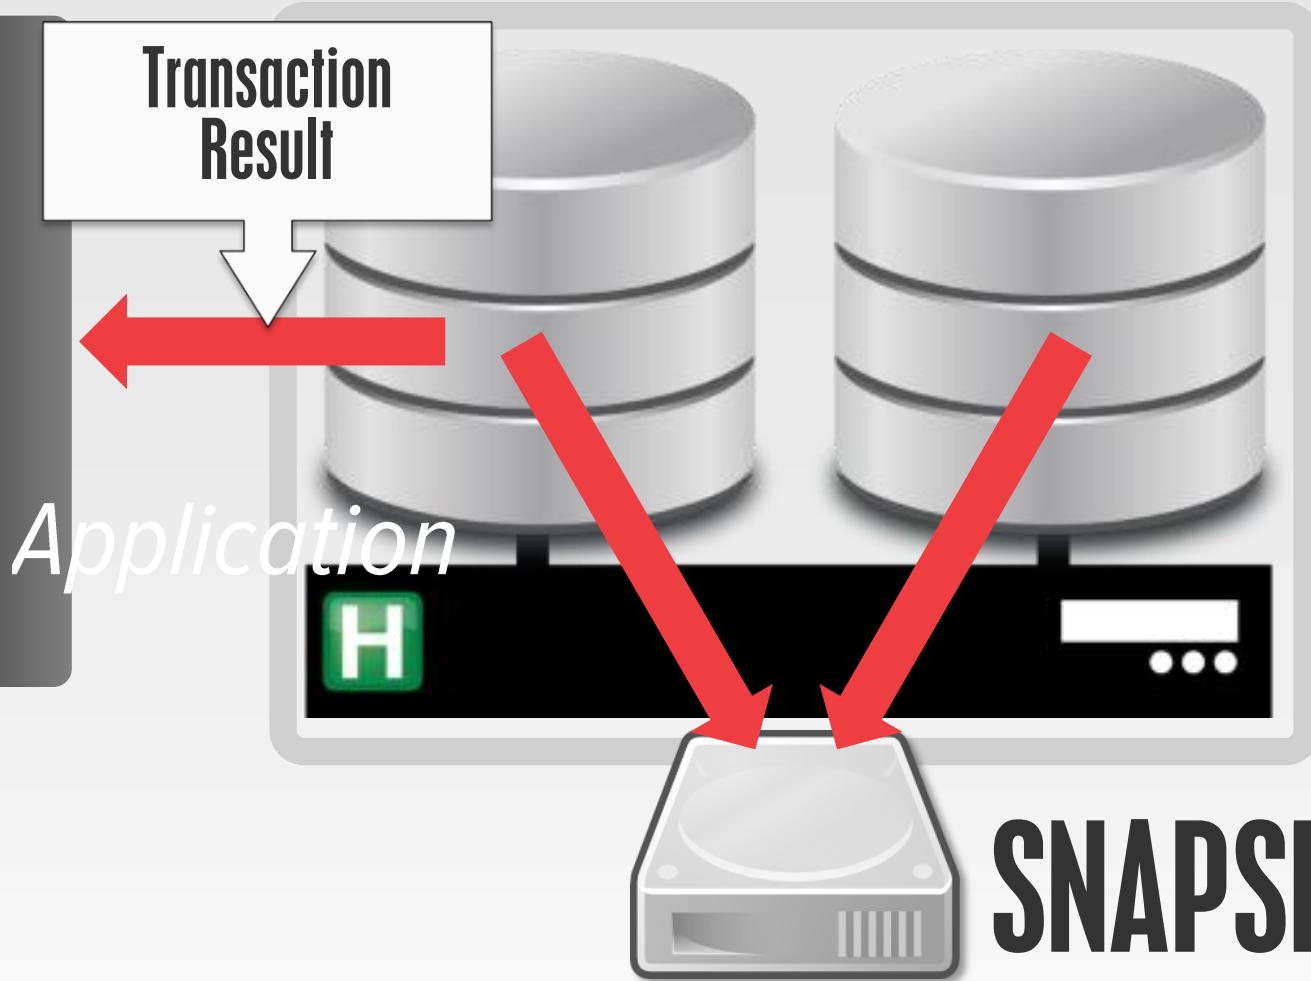
```
SELECT COUNT(*)  
      FROM votes  
     WHERE phone_num = ?;
```

InsertVote:

```
INSERT INTO votes  
VALUES (?, ?, ?);
```

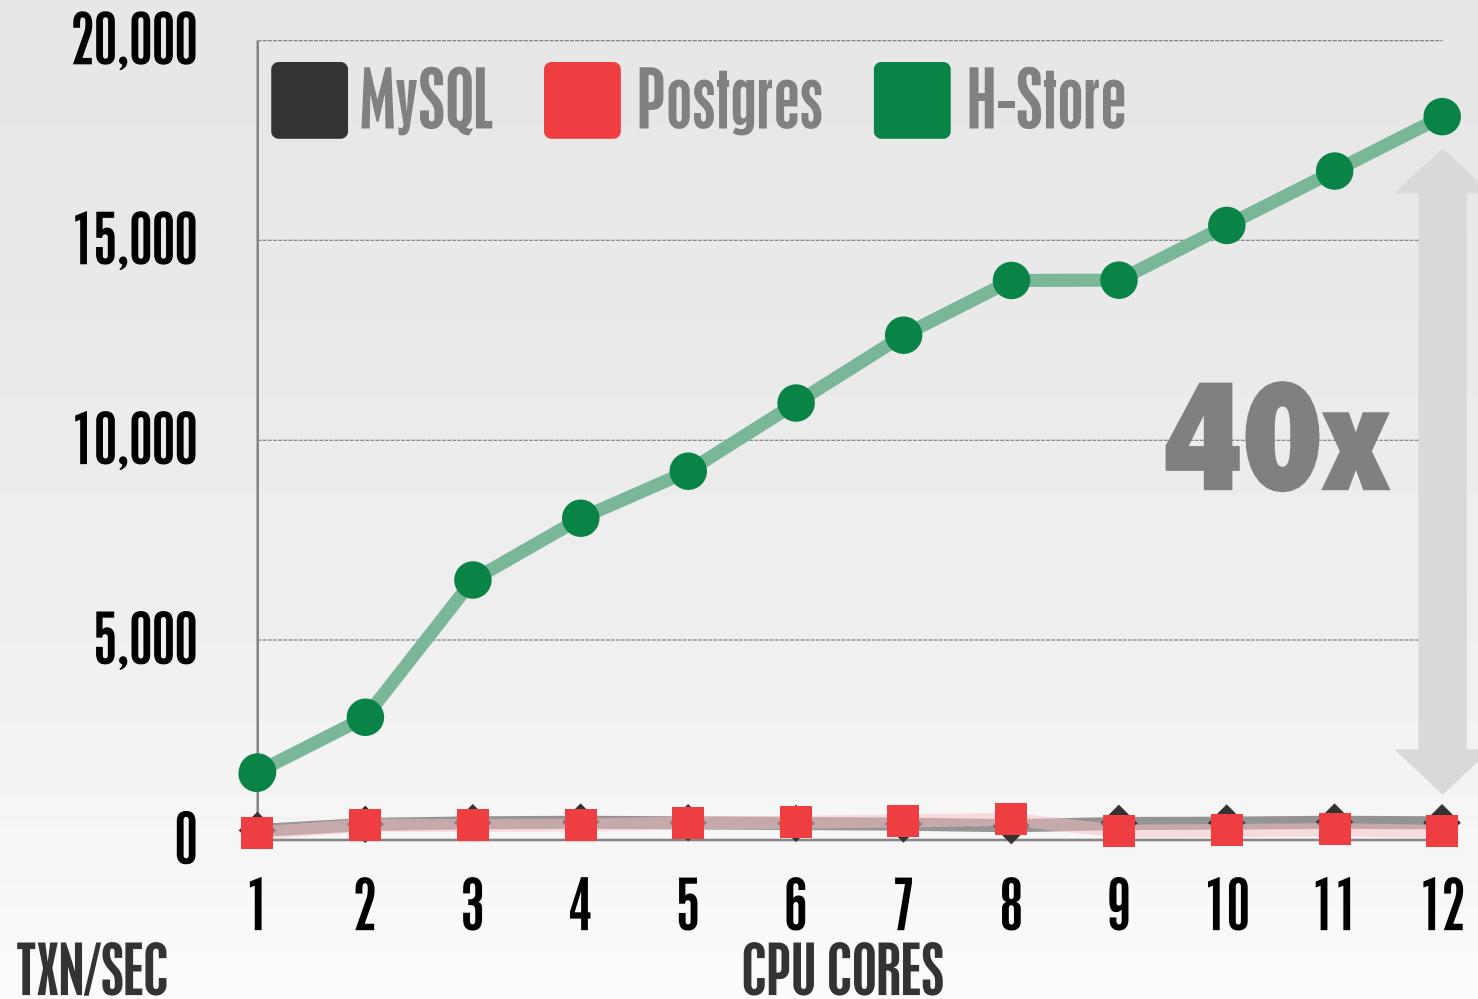
```
run(phoneNum, contestantId, currentTime) {  
    result = execute(VoteCount, phoneNum);  
    if (result > MAX_VOTES) {  
        return (ERROR);  
    }  
    execute(InsertVote, phoneNum,  
           contestantId,  
           currentTime);  
    return (SUCCESS);  
}
```





TPC-C BENCHMARK

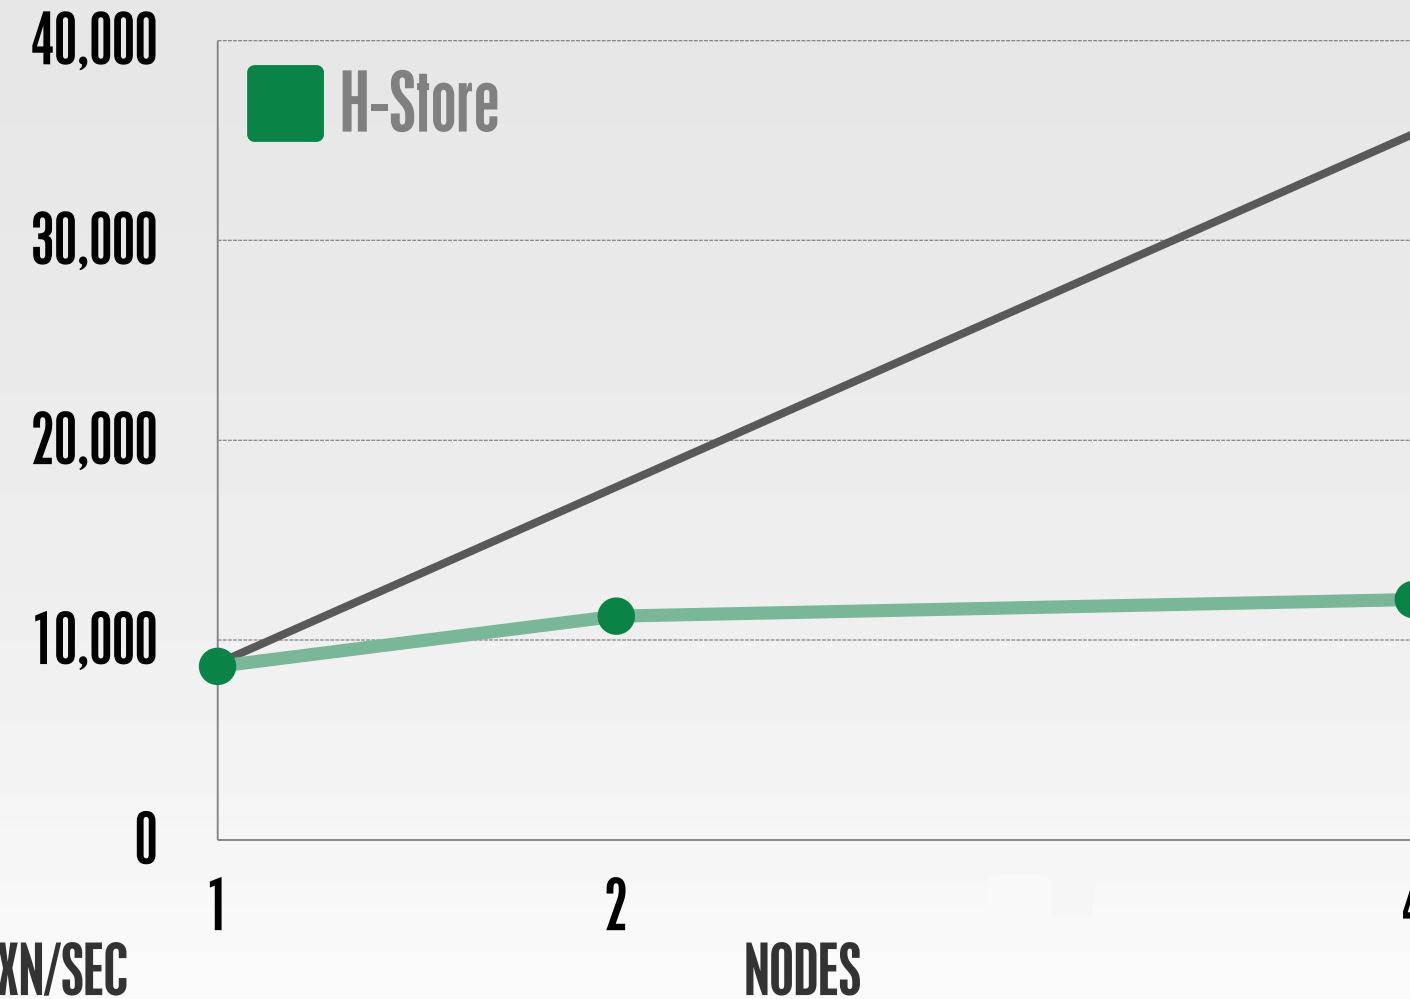
Warehouse Order Processing



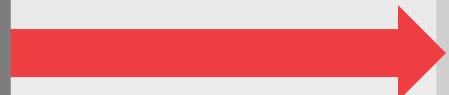
DISTRIBUTED TRANSACTIONS

TPC-C BENCHMARK

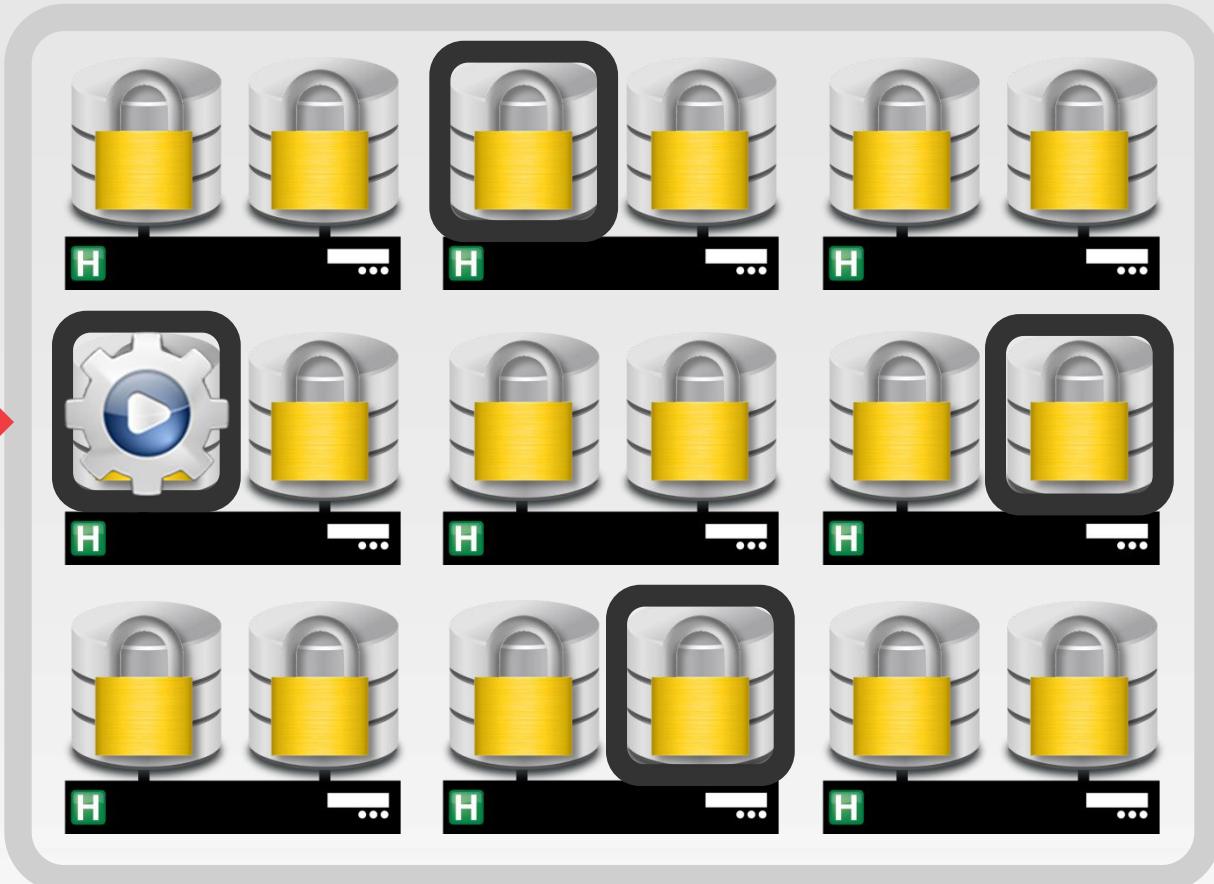
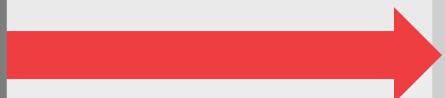
8 Cores per Node
10% Distributed Transactions



DISTRIBUTED TRANSACTIONS



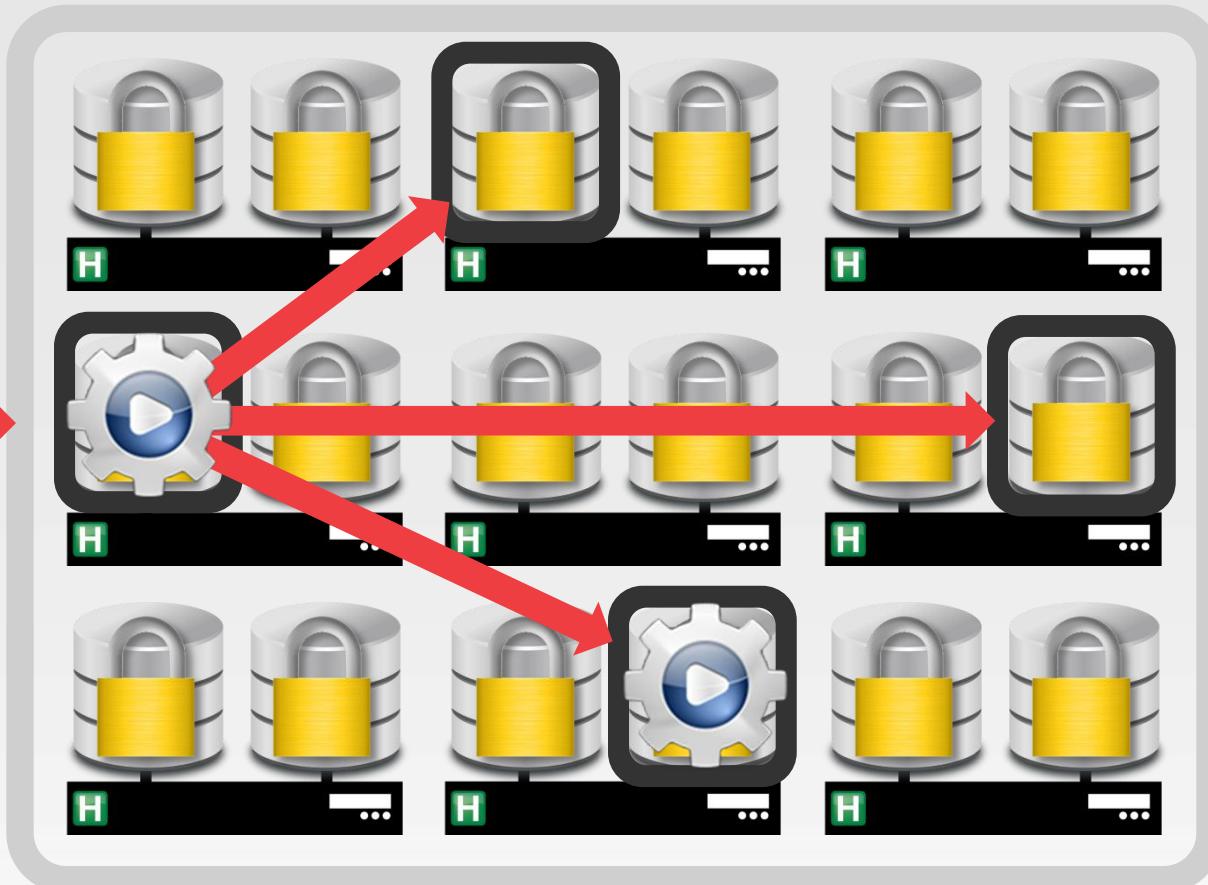
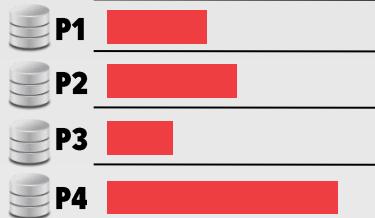
DISTRIBUTED TRANSACTIONS



DISTRIBUTED TRANSACTIONS



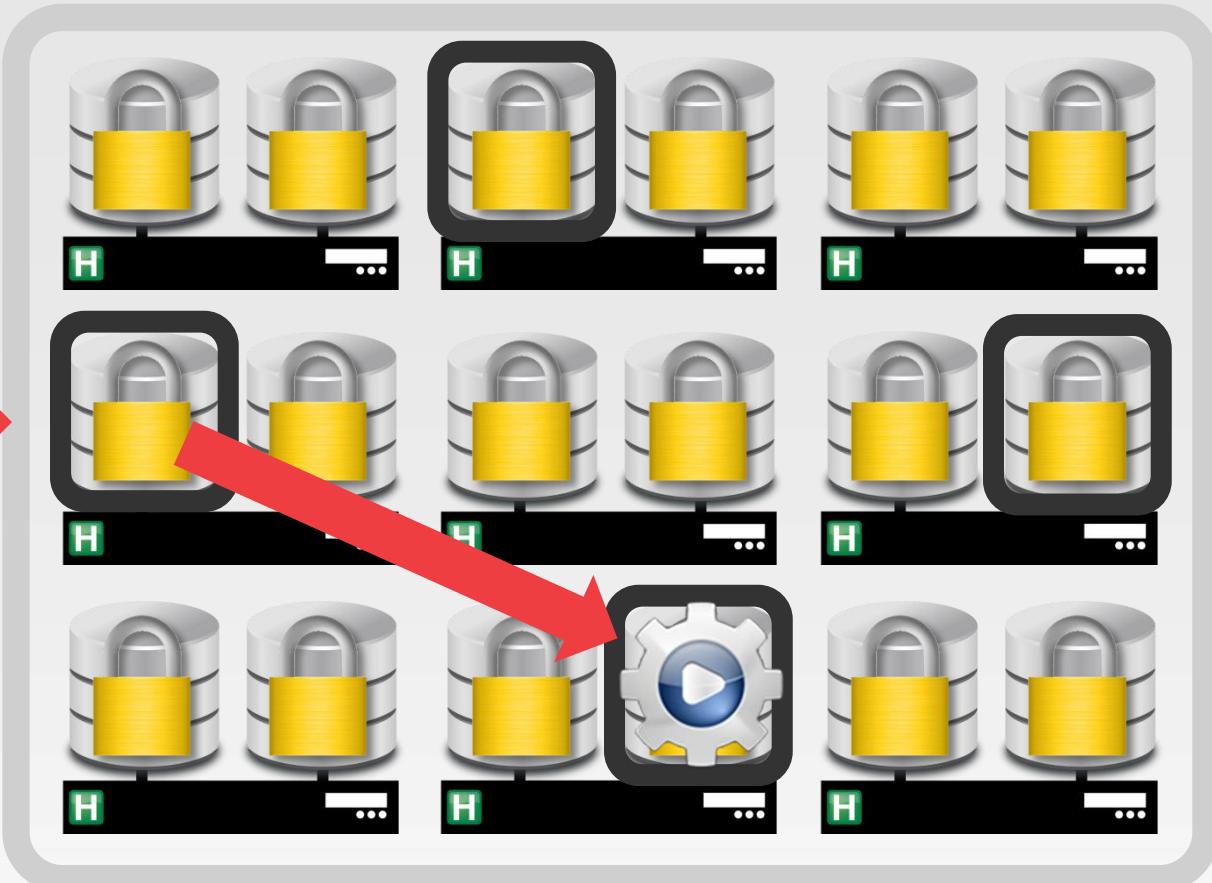
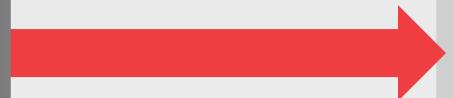
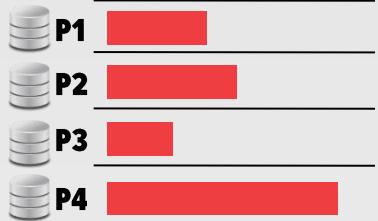
Query Count



DISTRIBUTED TRANSACTIONS



Query Count



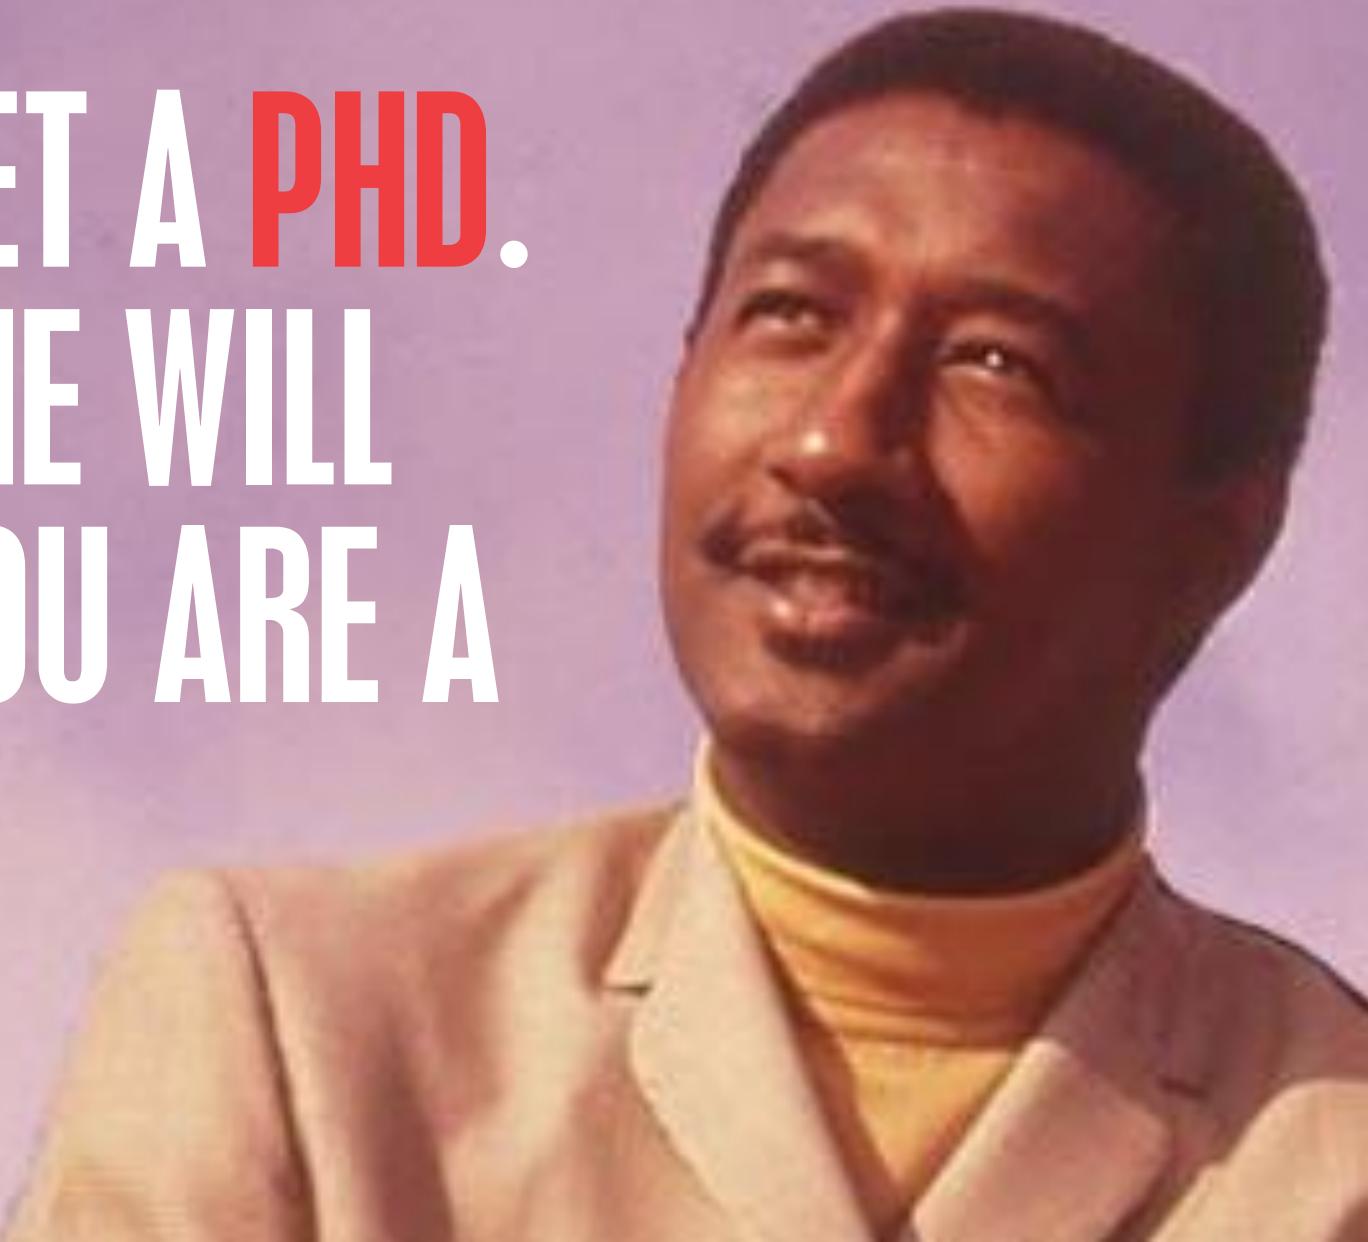
**KNOW WHAT TRANSACTIONS
WILL DO BEFORE THEY START**

BUT PEOPLE ALWAYS
GIVE ME BAD ADVICE

**DON'T GET
INVOLVED WITH
COMPUTERS.
YOU'LL NEVER
MAKE ANY
MONEY.**



DON'T GET A PHD.
EVERYONE WILL
THINK YOU ARE A
JERK.



THE DATABASE SYSTEM
ALWAYS HAS MORE



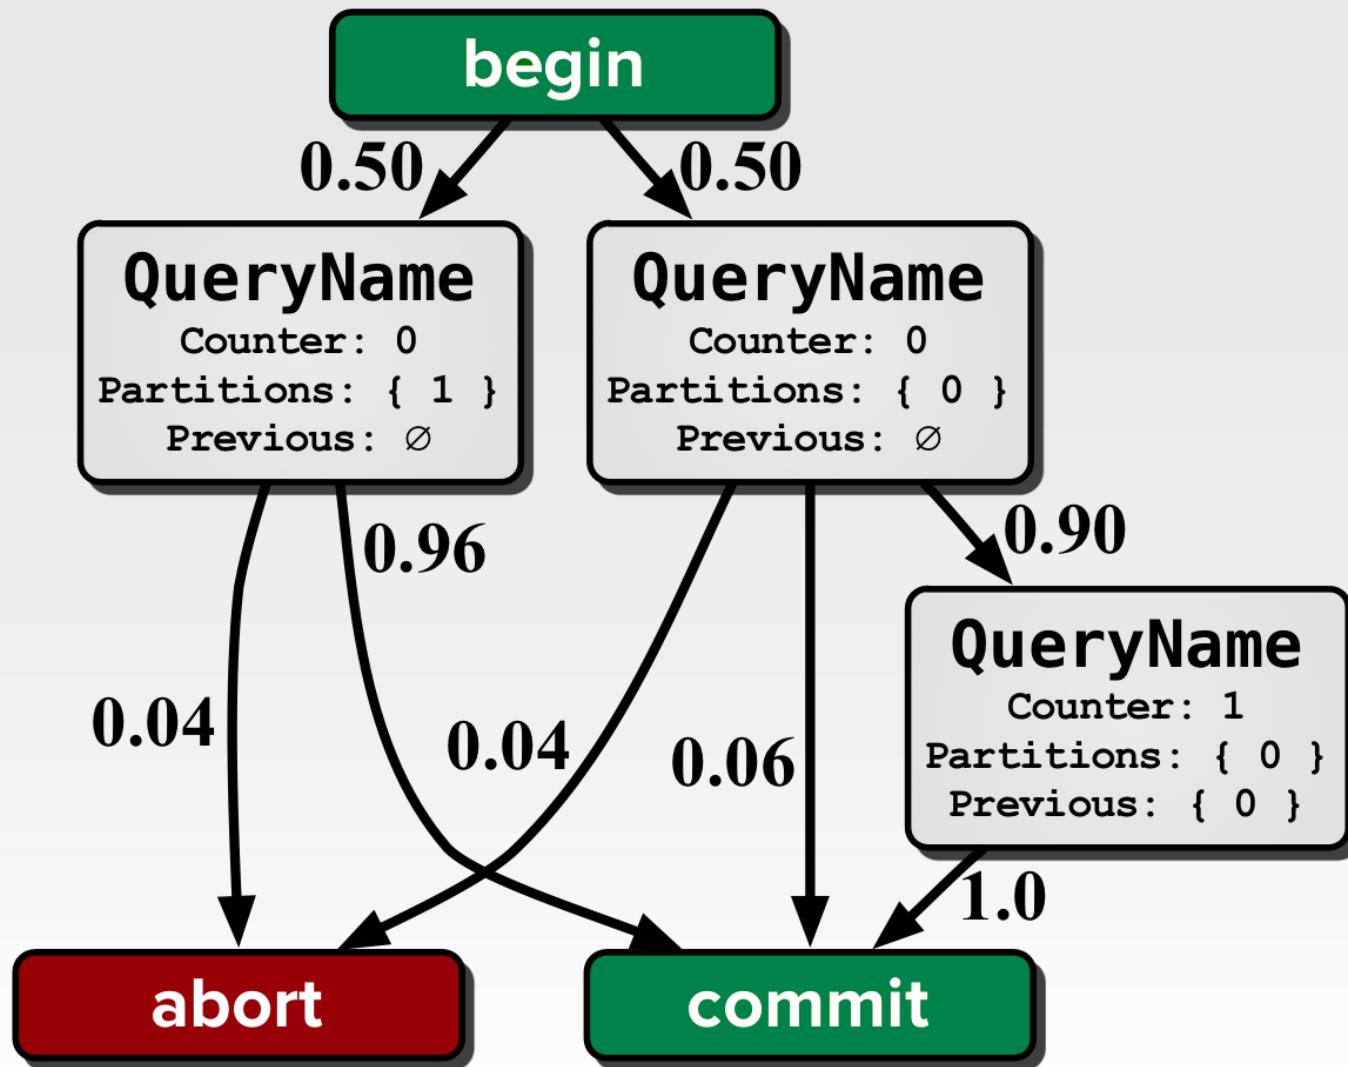
INFORMATION

DO USE MACHINE LEARNING TO PREDICT TRANSACTION BEHAVIOR.

ON PREDICTIVE MODELING FOR OPTIMIZING
TRANSACTION EXECUTION IN PARALLEL OLTP SYSTEMS
Proc. VLDB Endow., Vol 5, Iss. 2, pp. 85-96, 2011



PREDICTIVE MODELS



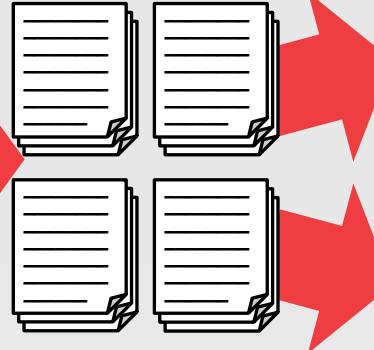
WORKLOAD

```
SELECT * FROM WAREHOUSE  
WHERE W_ID = 10;
```

```
SELECT * FROM DISTRICT WHERE  
D_W_ID = 10 AND D_ID = 9;
```

```
INSERT INTO ORDERS (O_W_ID,  
O_D_ID, O_C_ID,...) VALUES  
(10, 9, 12345,...);  
...
```

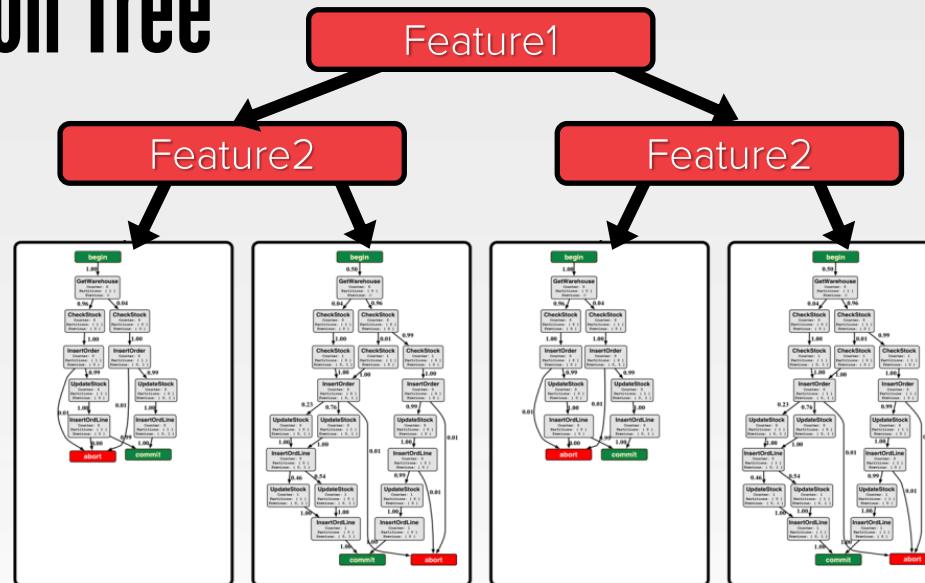
Feature Clusterer



Model Generator

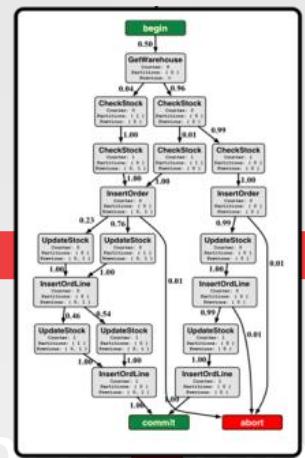
Classifier

Decision Tree

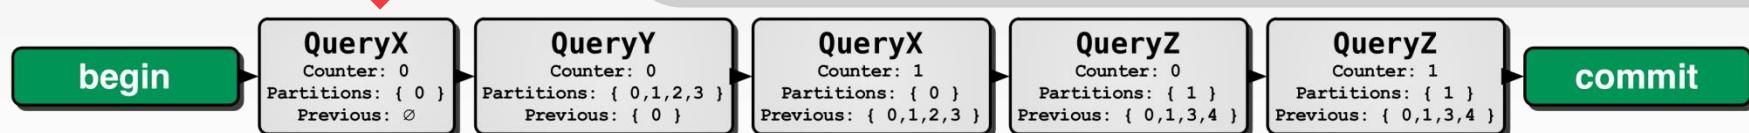


Markov Models

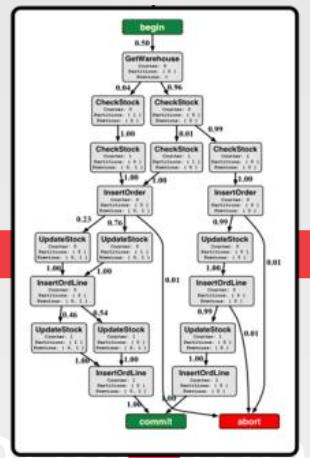
DISTRIBUTED TRANSACTIONS



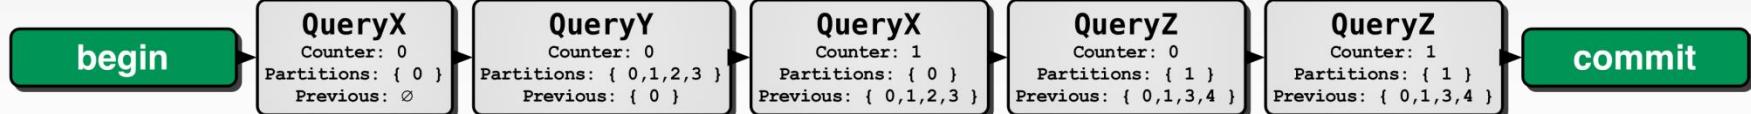
Application



DISTRIBUTED TRANSACTIONS

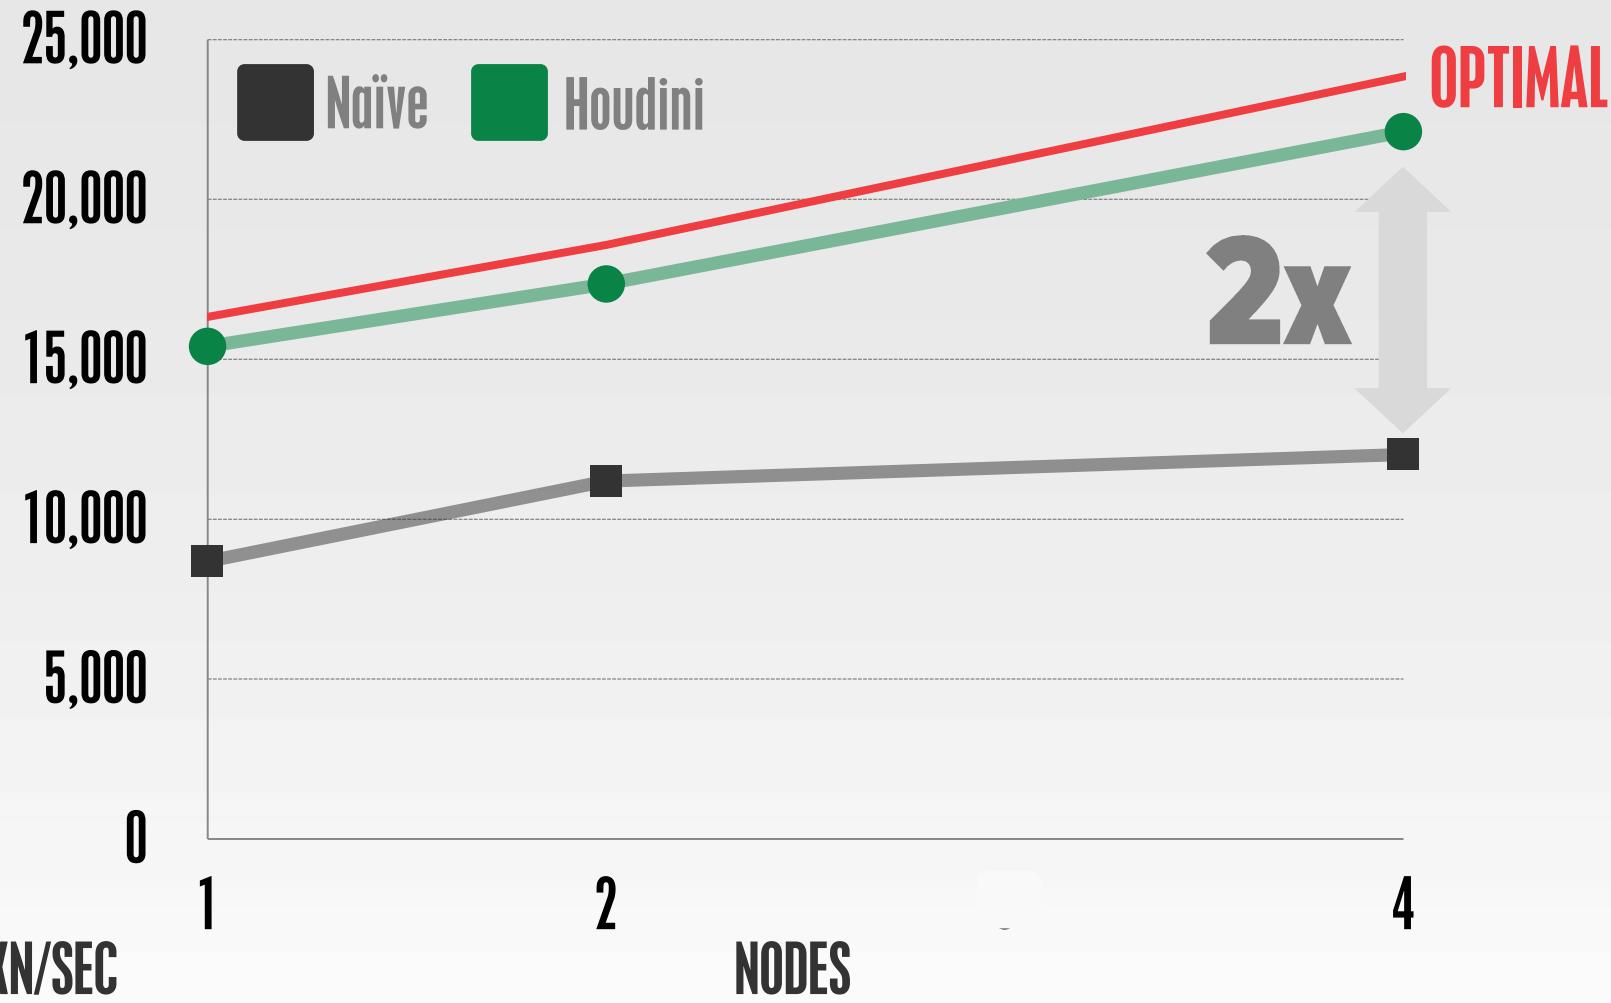


Application



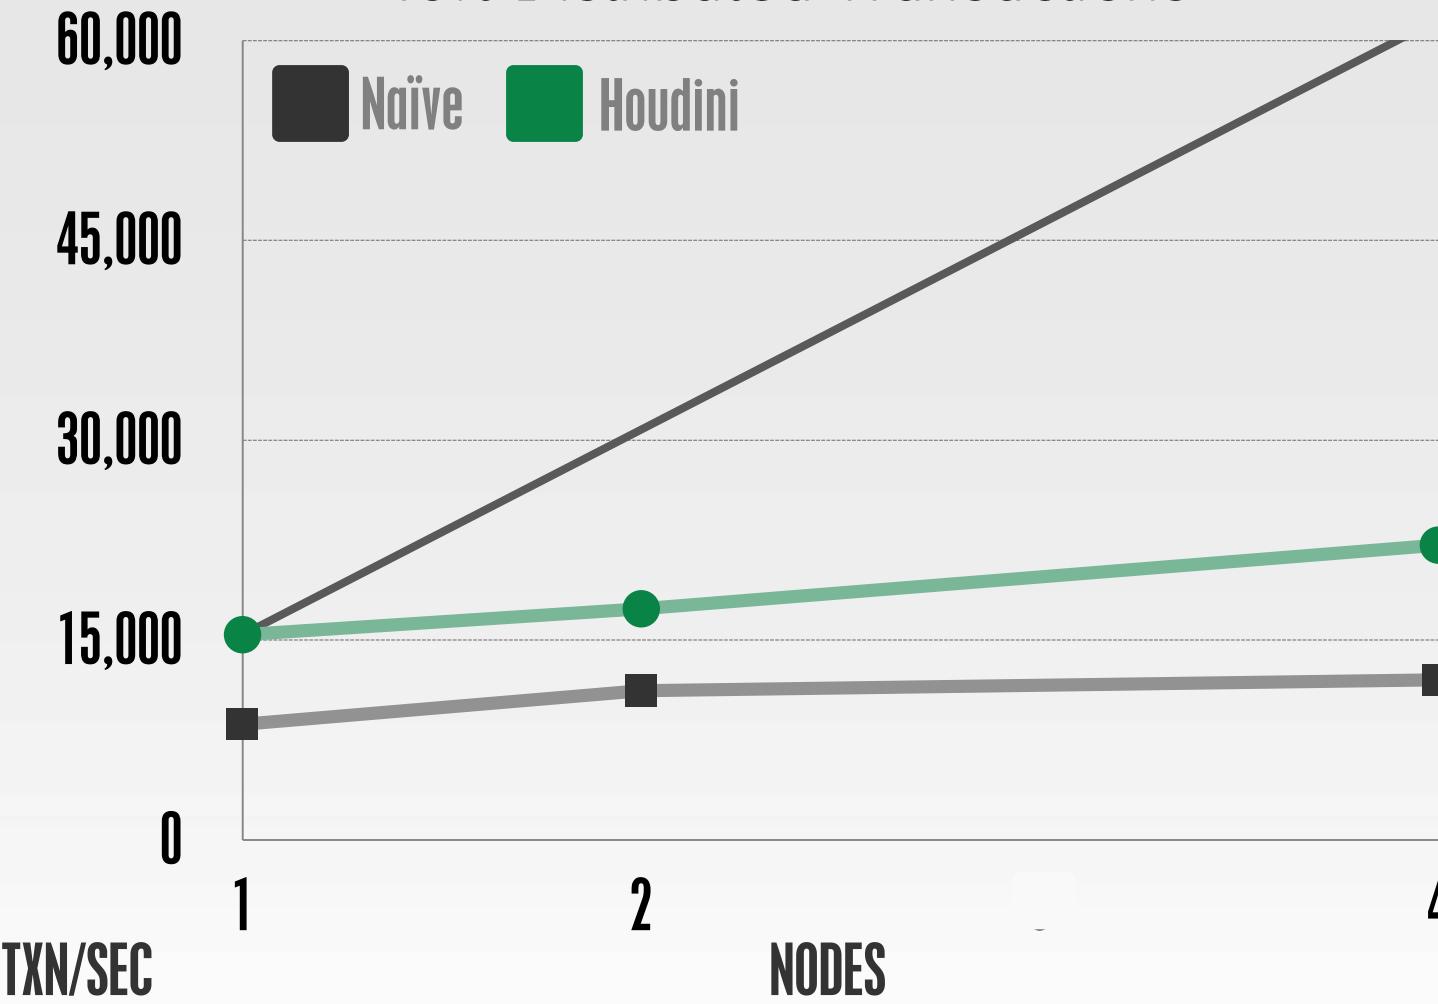
TPC-C BENCHMARK

8 Cores per Node
10% Distributed Transactions



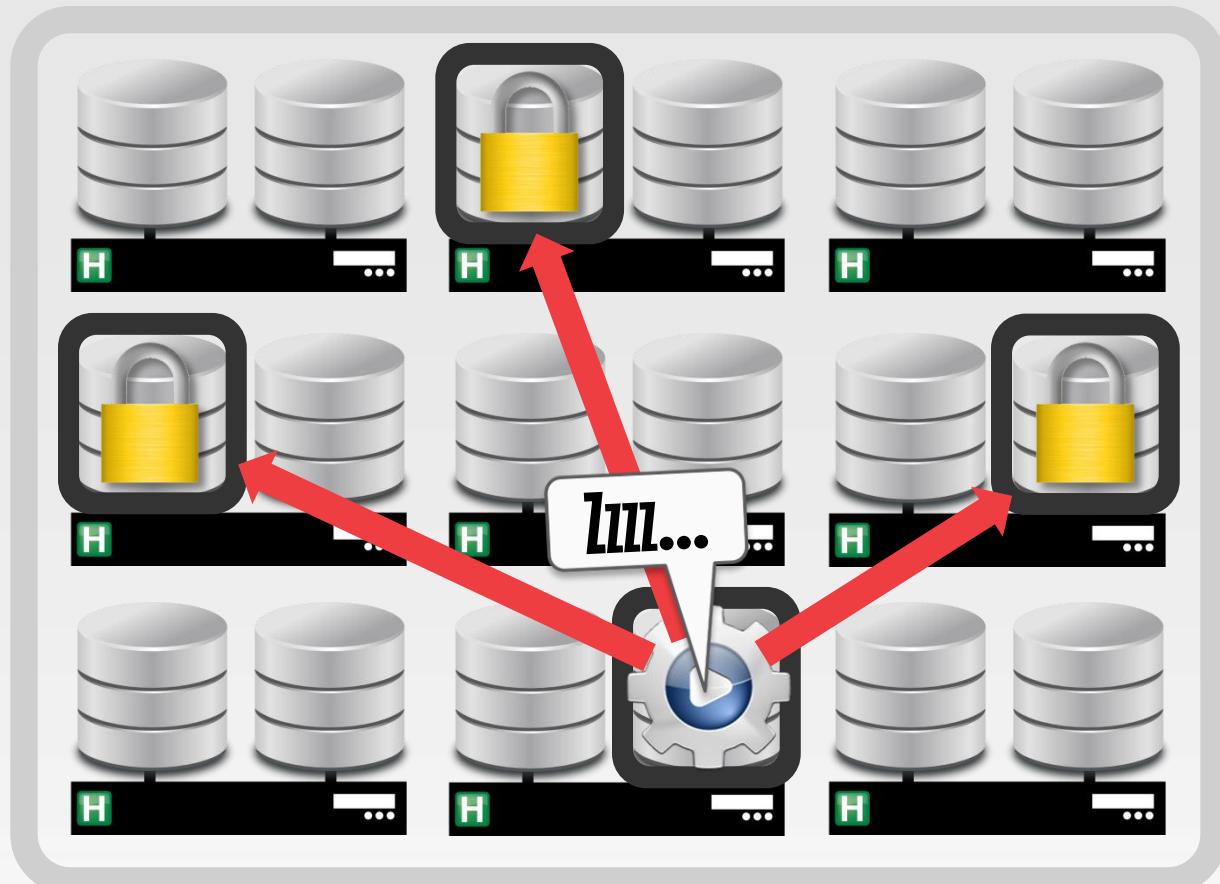
TPC-C BENCHMARK

8 Cores per Node
10% Distributed Transactions



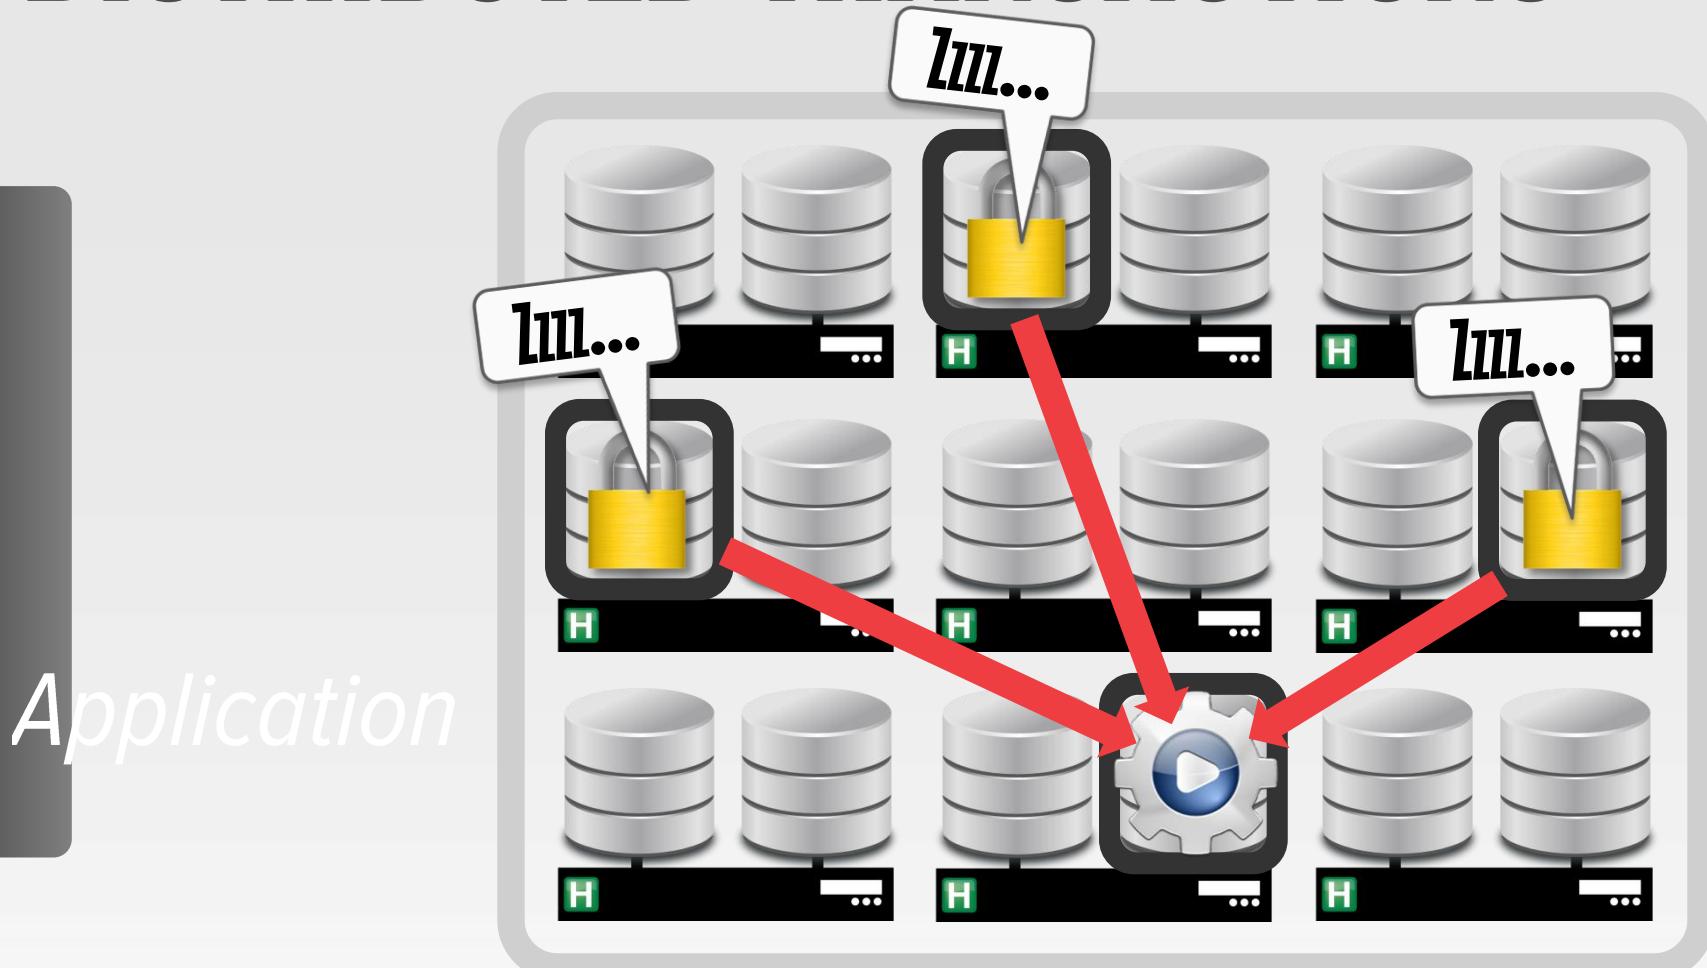
DISTRIBUTED TRANSACTIONS

Application



SP1 - Waiting for Query Result

DISTRIBUTED TRANSACTIONS



SP1 - Waiting for Query Result

SP2 - Waiting for Query Request

DISTRIBUTED TRANSACTIONS



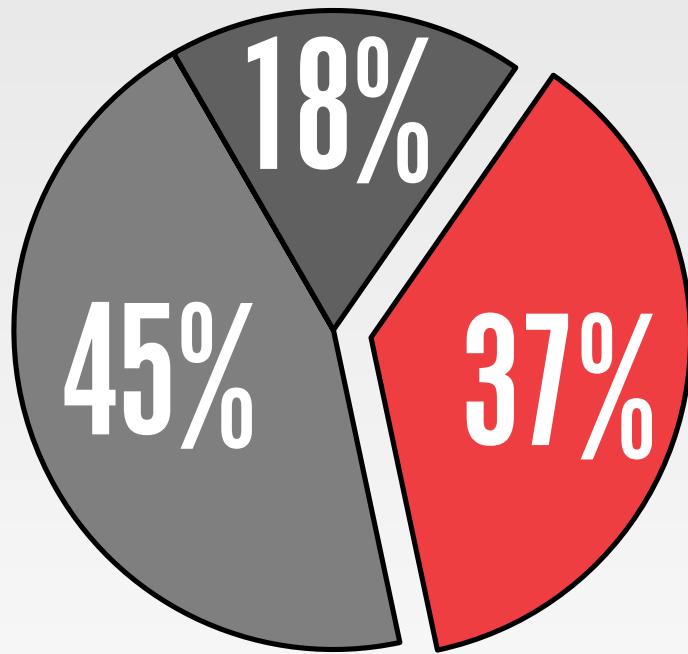
SP1 - Waiting for Query Result

SP2 - Waiting for Query Request

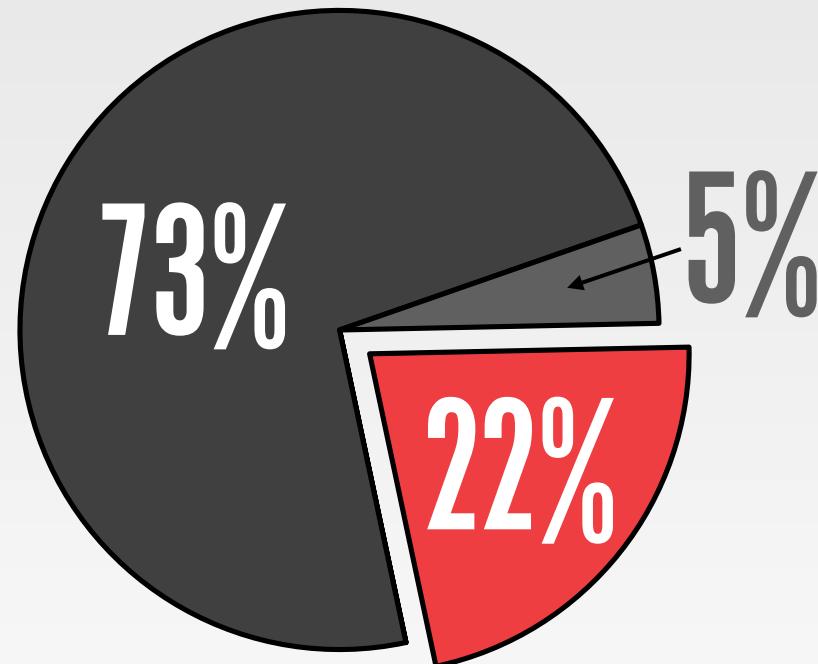
SP3 - Two-Phase Commit

TRANSACTION STALL POINTS

BASE PARTITION



REMOTE PARTITION



- SP1 - Waiting for Query Result
- SP3 - Two-Phase Commit

- SP2 - Waiting for Query Request
- Real Work

**DO SOMETHING USEFUL
WHEN STALLED**

**DON'T BE SURPRISED
IF YOU & KB DON'T
LAST THROUGH
GRAD SCHOOL.**



DON'T BE STAN'S
STUDENT
IF YOU GO TO
BROWN.



DO USE MACHINE LEARNING
TO SCHEDULE
SPECULATIVE TASKS.

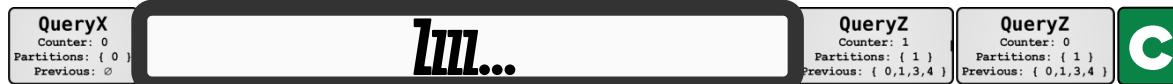
THE ART OF SPECULATIVE EXECUTION
In Progress (August 2013)



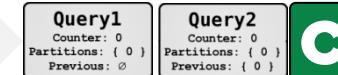
SERIALIZABLE SCHEDULE



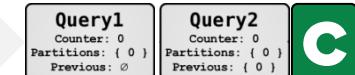
Distributed Transaction



Single-Partition Transaction



Single-Partition Transaction



SERIALIZABLE SCHEDULE



Distributed Transaction



VERIFY



Speculative Transaction



Speculative Transaction

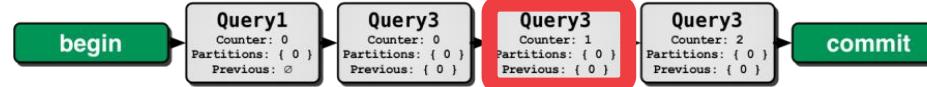


SPECULATIVE TRANSACTIONS

zzz...

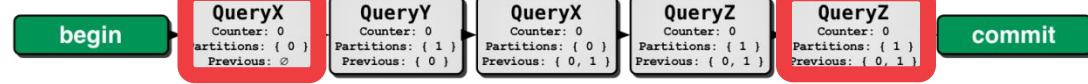


Speculation Candidate:



WRITE X

Distributed Transaction:

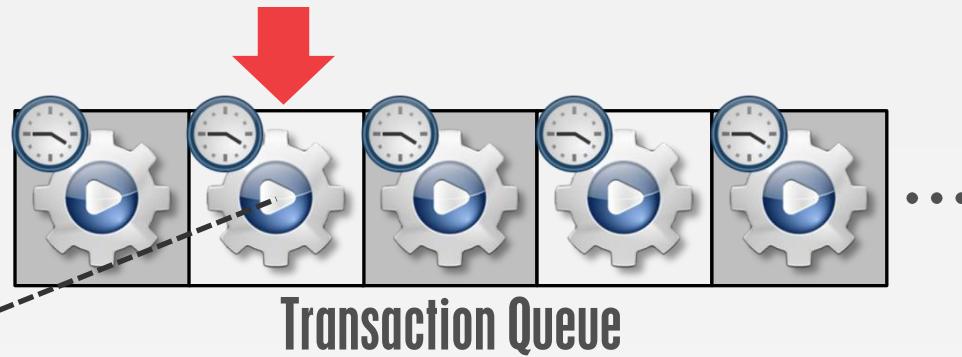


READ X

READ X

SPECULATIVE TRANSACTIONS

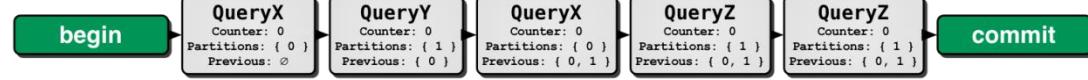
zzz...



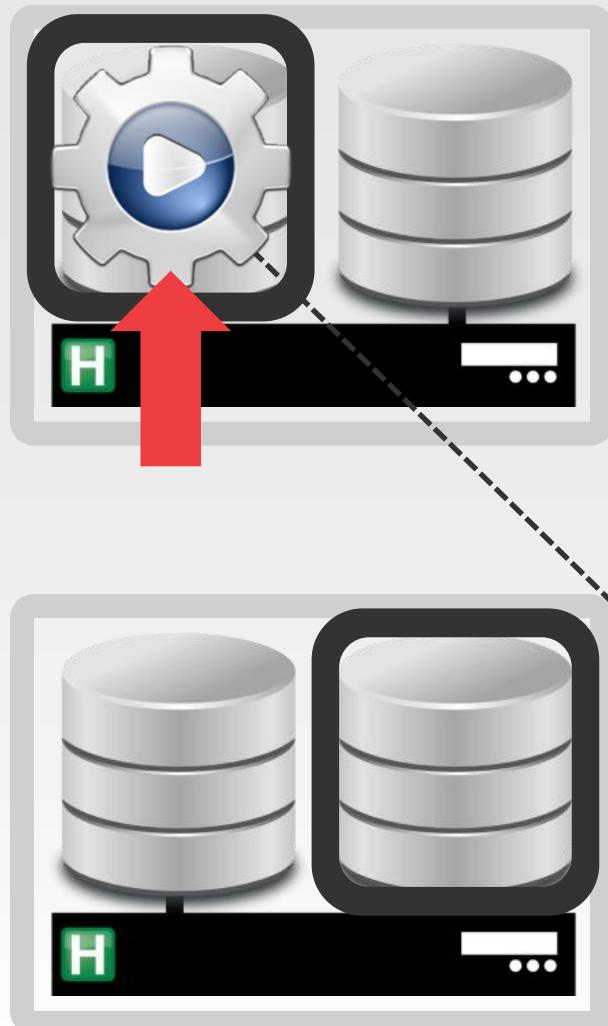
Speculation Candidate:



Distributed Transaction:



SPECULATIVE QUERIES



Distributed Transaction:

begin

QueryX
Counter: 0
Partitions: { 0 }
Previous: \emptyset

QueryY
Counter: 0
Partitions: { 1 }
Previous: { 0 }

QueryX
Counter: 0
Partitions: { 0 }
Previous: { 0, 1 }

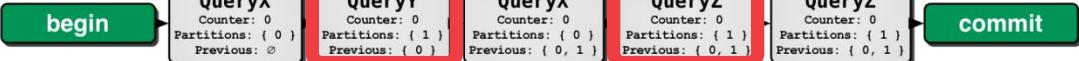
QueryZ
Counter: 0
Partitions: { 1 }
Previous: { 0, 1 }

commit

SPECULATIVE QUERIES



Distributed Transaction:



SPECULATIVE QUERIES



QueryY:

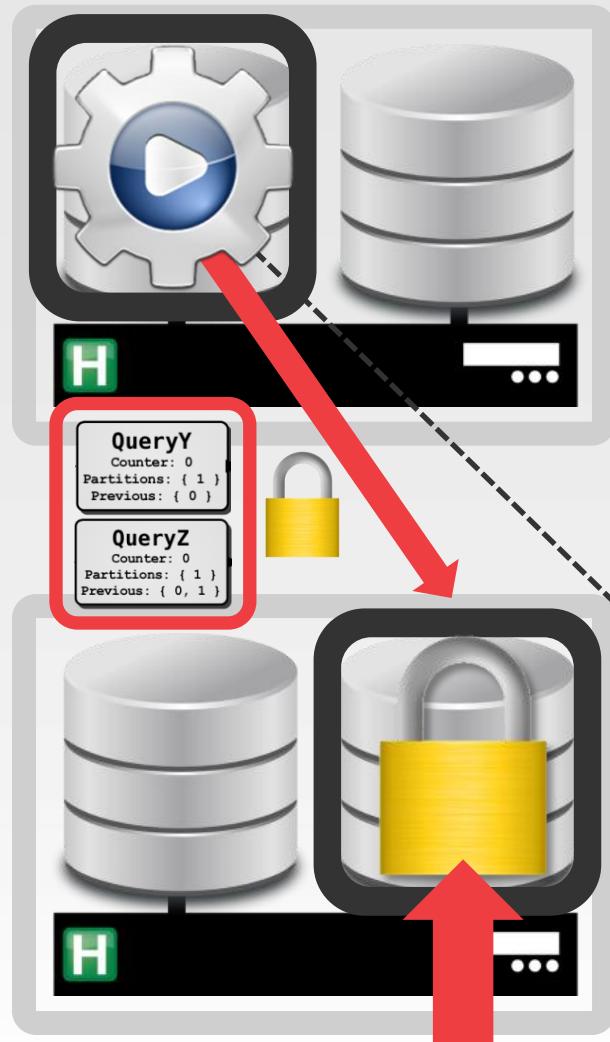
**SELECT S_QTY FROM STOCK
WHERE S_W_ID = ?
AND S_I_ID = ?**



Distributed Transaction:



SPECULATIVE QUERIES



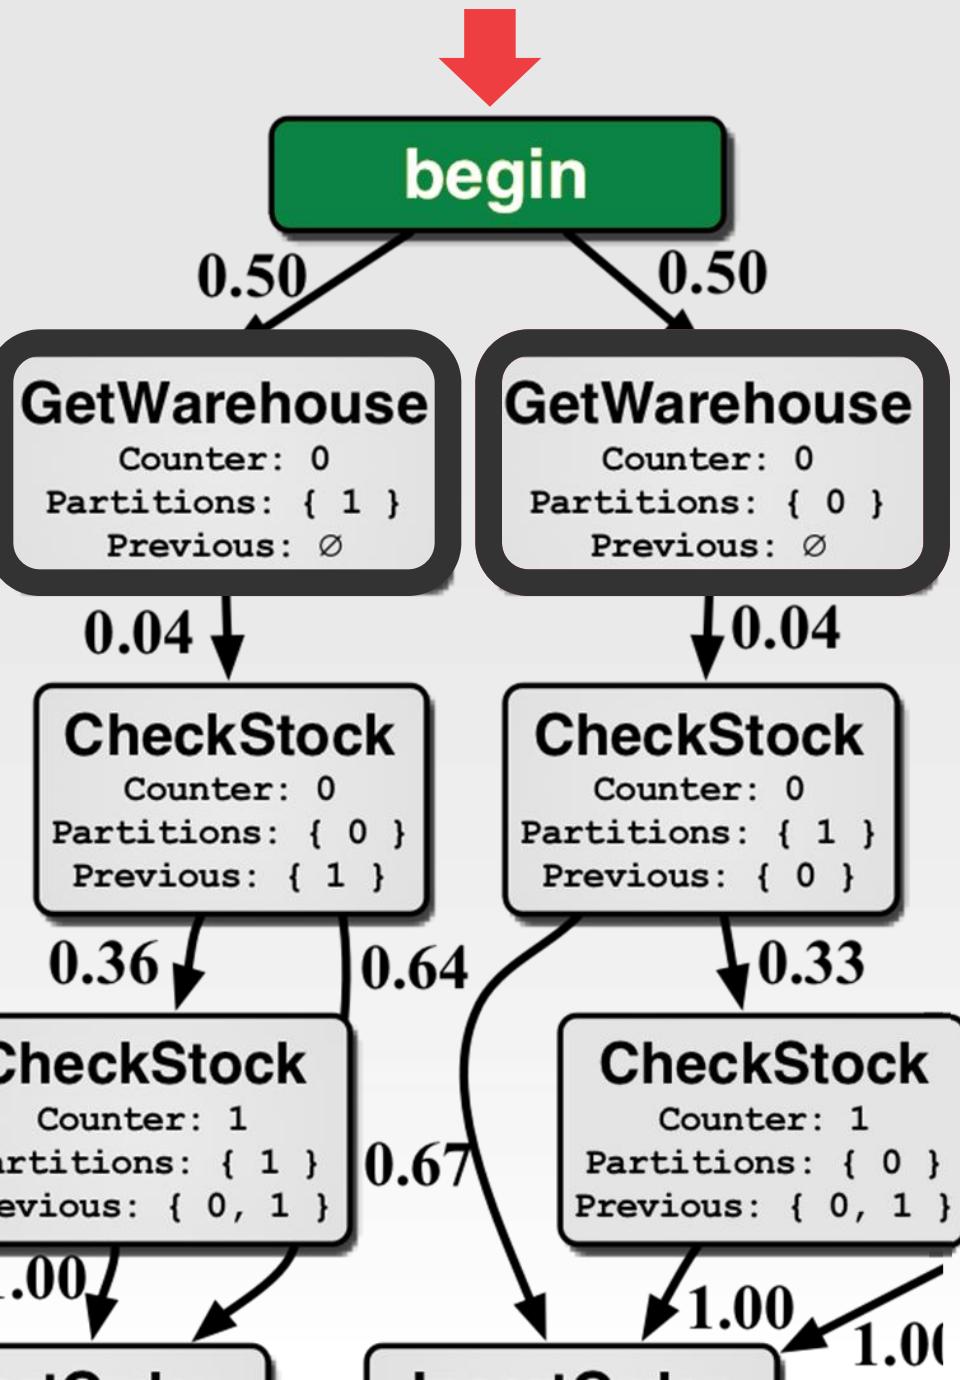
QueryY:

**SELECT S_QTY FROM STOCK
WHERE S_W_ID = ?
AND S_I_ID = ?**



Distributed Transaction:

begin QueryX Counter: 0 Partitions: { } Previous: { } QueryY Counter: 0 Partitions: {1} Previous: { } QueryX Counter: 0 Partitions: { } Previous: {0, 1} QueryZ Counter: 0 Partitions: {1} Previous: {0, 1} QueryZ Counter: 0 Partitions: {1} Previous: {0, 1} commit

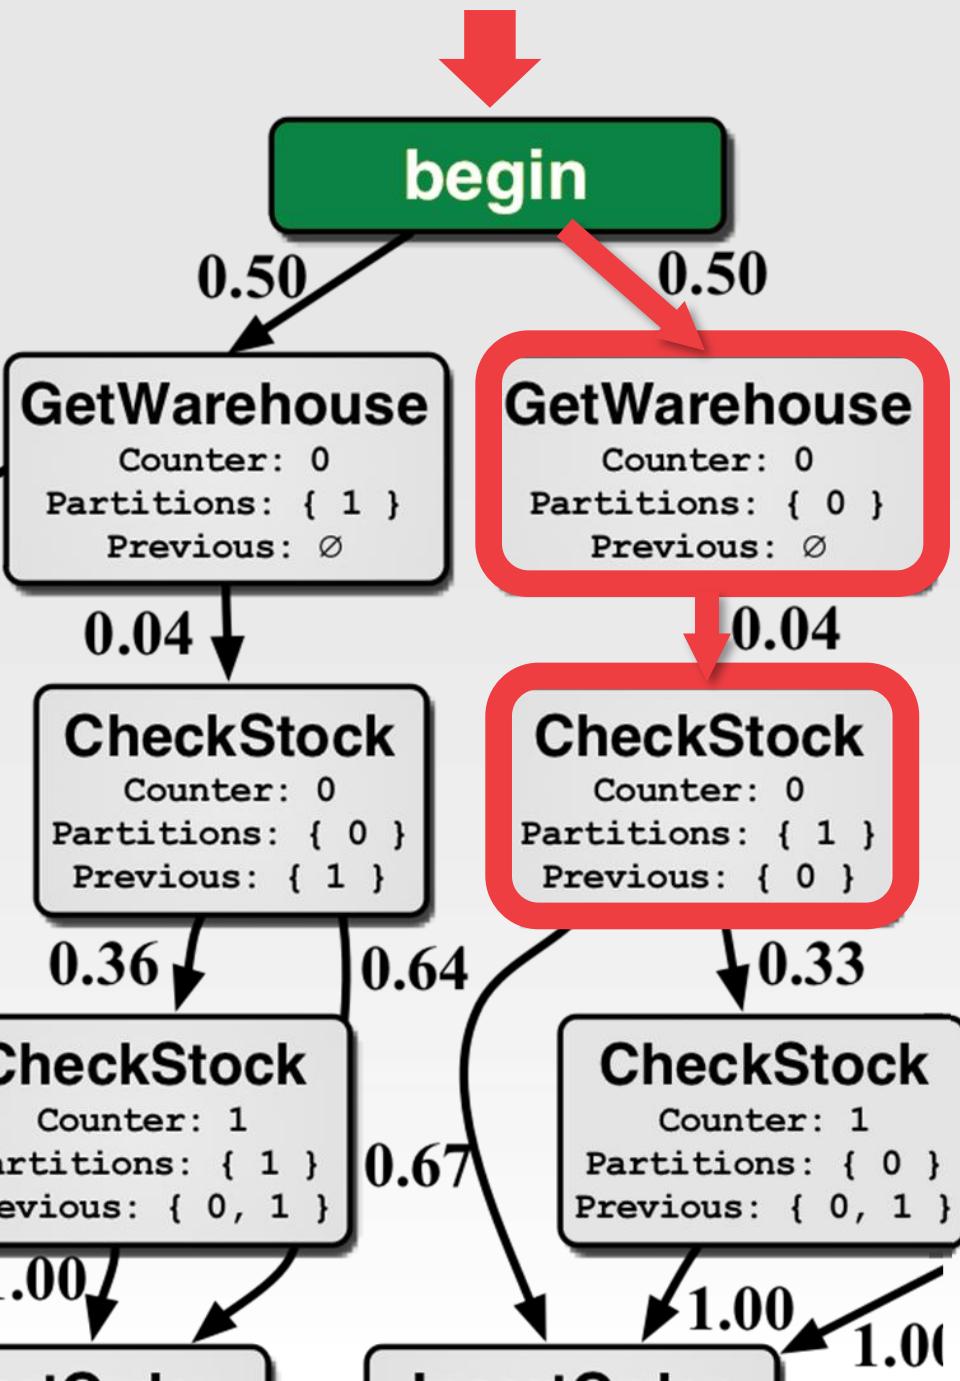


Transaction Parameters:

w_id=0
i_w_ids=[1,0]
i_ids=[1001,1002]

GetWarehouse:

SELECT * FROM WAREHOUSE
WHERE W_ID = ?

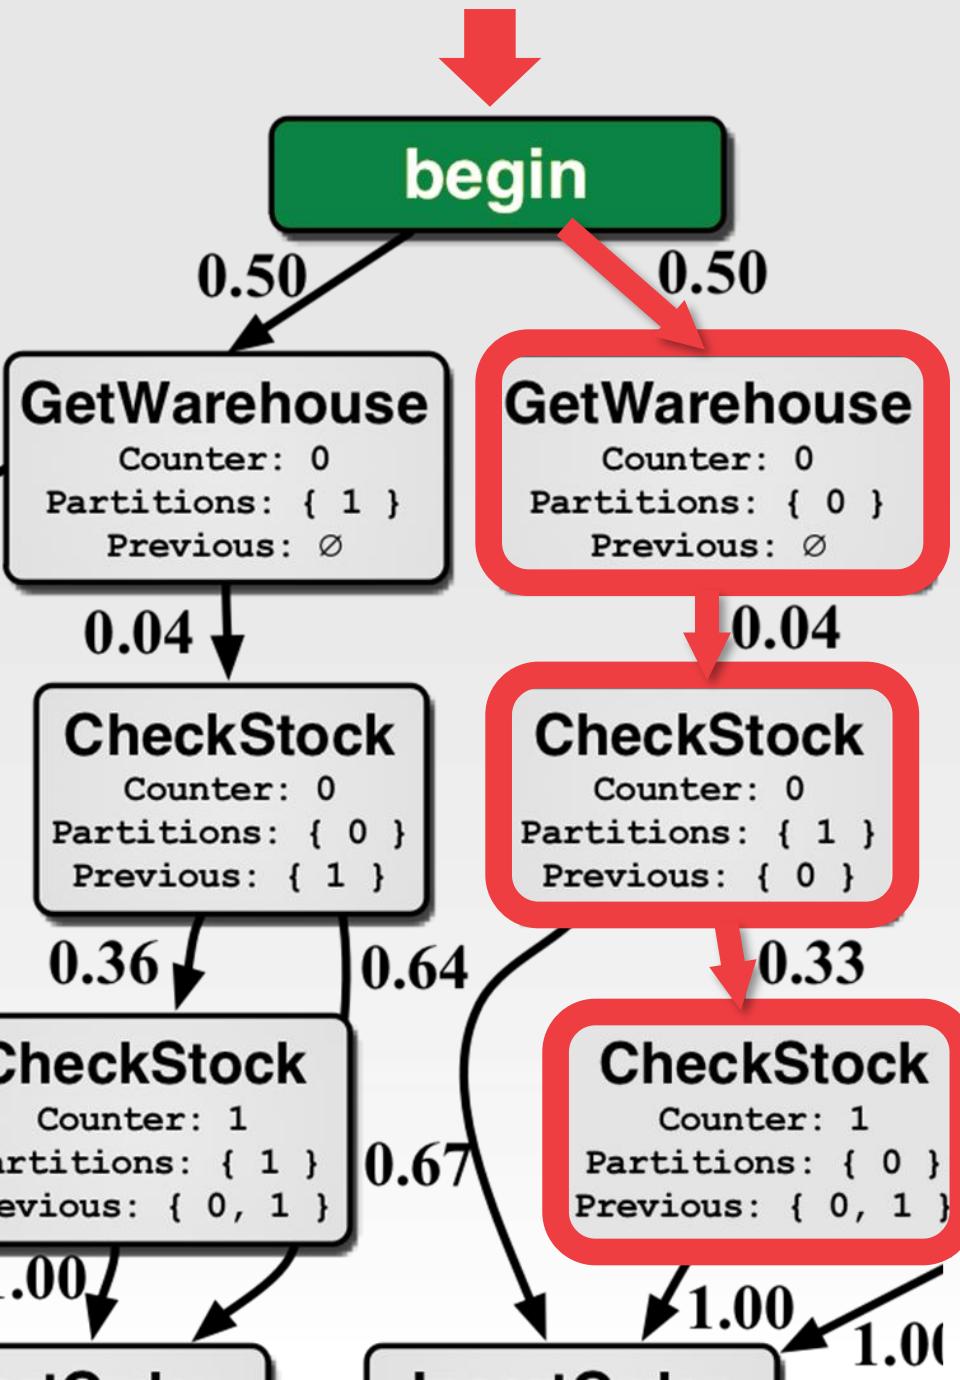


Transaction Parameters:

w_id=0
 i_w_ids=[1, 0]
 i_ids=[1001, 1002]

CheckStock:

SELECT S_QTY FROM STOCK
 WHERE S_W_ID = ?
 AND S_I_ID = ?



CheckStock:

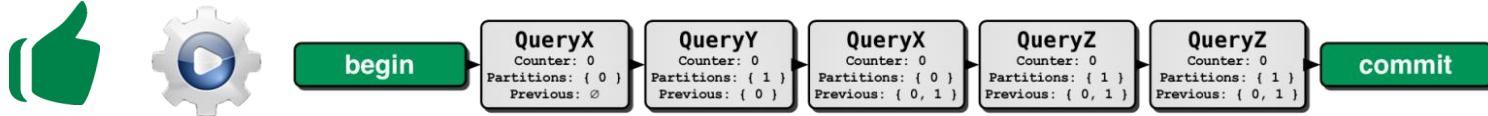
```

SELECT S_QTY FROM STOCK
WHERE S_W_ID = ?
AND S_I_ID = ?;

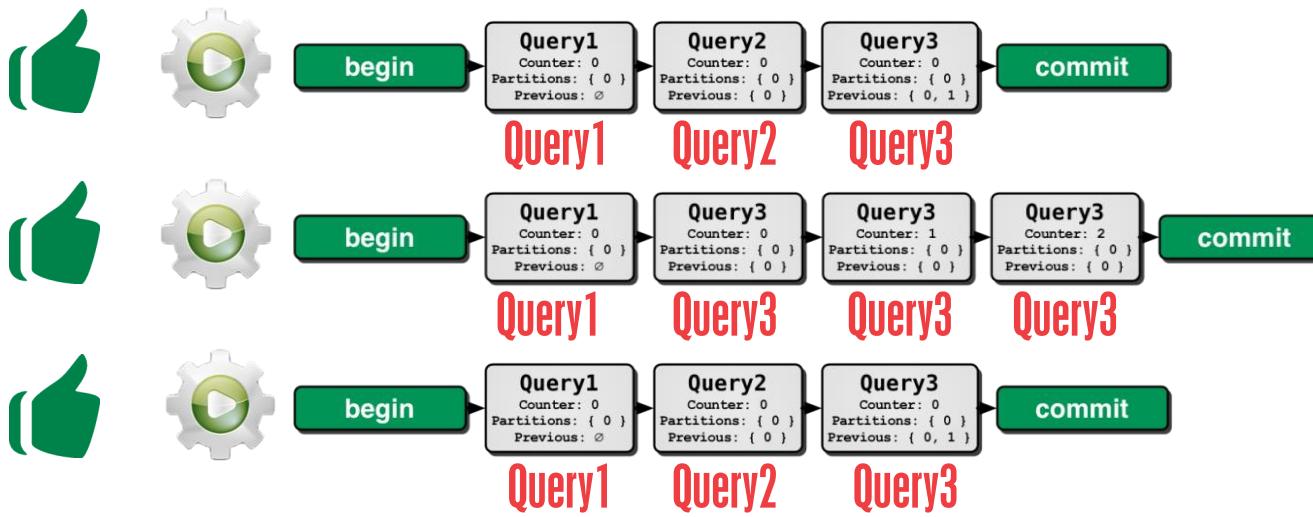
```

VERIFICATION

Distributed Transaction

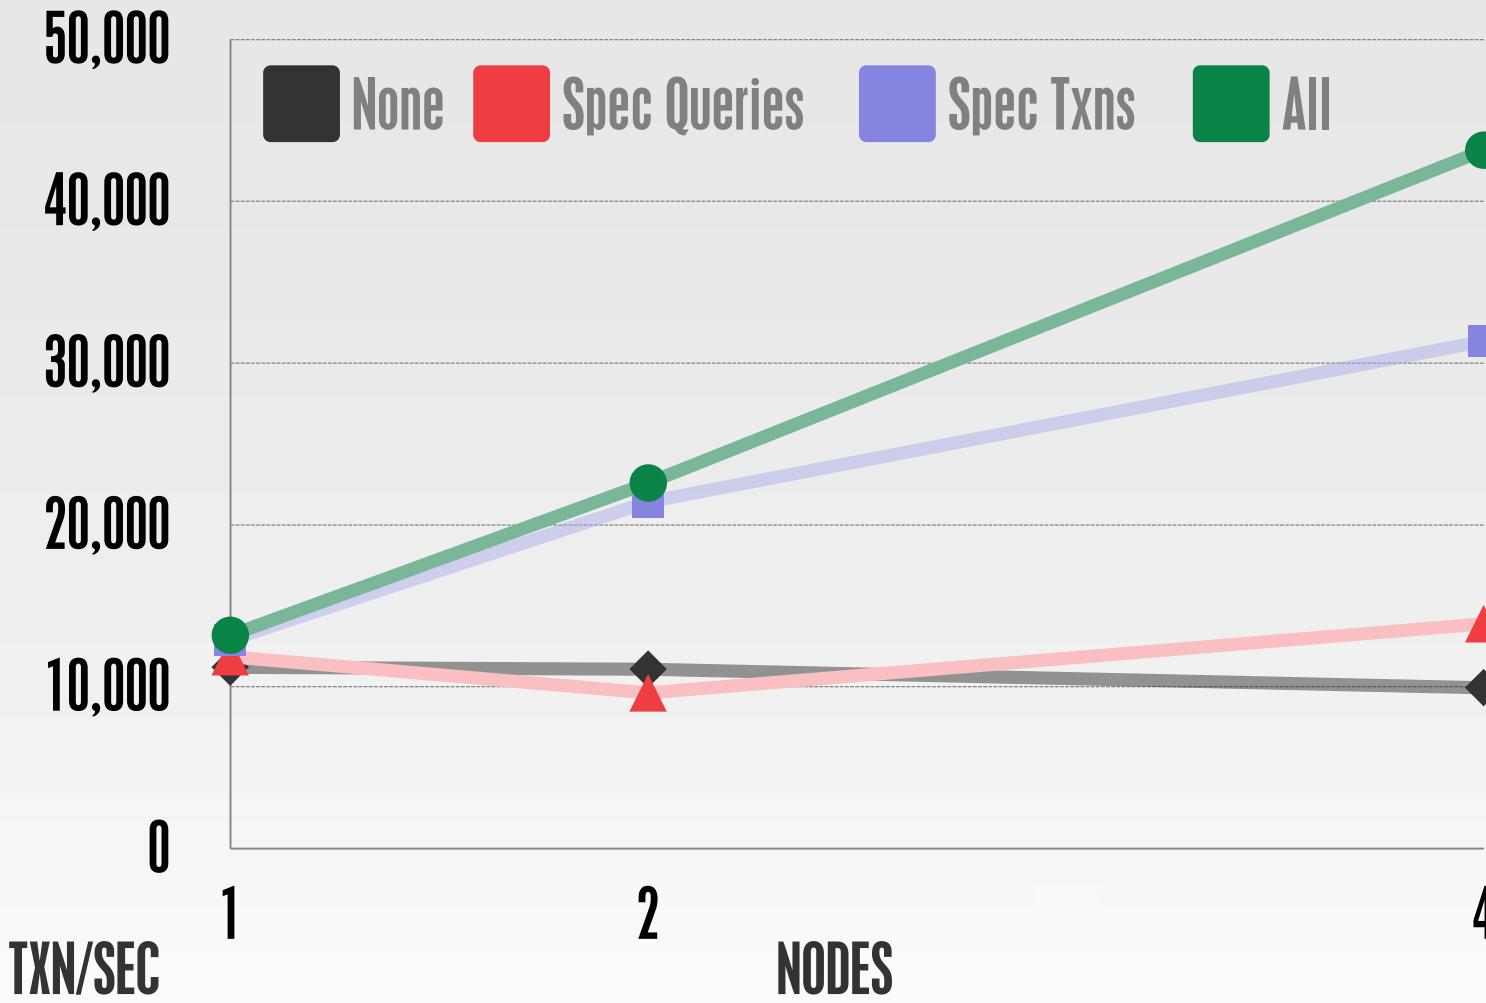


Speculative Transactions



TPC-C BENCHMARK

8 Cores per Node
10% Distributed Transactions





Optimize Single-Partition Execution

H-STORE: A HIGH-PERFORMANCE, DISTRIBUTED MAIN
MEMORY TRANSACTION PROCESSING SYSTEM
Proc. VLDB Endow., vol. 1, iss. 2, pp. 1496-1499, 2008.



Minimize Distributed Transactions

SKEW-AWARE AUTOMATIC DATABASE PARTITIONING IN
SHARED-NOTHING, PARALLEL OLTP SYSTEMS
Proceedings of SIGMOD, 2012.



Identify Distributed Transactions

ON PREDICTIVE MODELING FOR OPTIMIZING TRANSACTION
EXECUTION IN PARALLEL OLTP SYSTEMS
Proc. VLDB Endow., vol. 5, pp. 85-96, 2011.



Utilize Transaction Stalls

THE ART OF SPECULATIVE EXECUTION
In Progress (August 2013)

**FUTURE
WORK**

One Size
Almost
Fits AllTM



H-STORE



S-STORE



N-STORE

Escape From Planet Zdonik

(i.e., Andy Needs to Get Tenure)

Beyond the 'Stores

- Non-Partitionable Workloads.
- The Poor Man's Spanner.
- Scientific Databases.

DON'T MESS IT UP
WITH KB.





**Stan
Zdonik**



**"The Thrill"
Stonebraker**



**Sam
Madden**



**Ugur
Cetintemel**



**David
DeWitt**



**Dan
Abadi**



**Evan
Jones**



**Saurya
Velagapudi**



**Xin
Jia**



**Carlo
Curino**



**Justin
DeBrabant**



**Yang
Zou**



**Visawee
Angkana.**



**Ning
Shi**



**John
Meehan**