

# Example

Customers		
<u>name</u>	street	cityID
peter	minstreet	0
steve	macstreet	1
mike	longstreet	9
tim	unistreet	9
hans	msstreet	5
jens	shortstreet	1
frank	minstreet	0
olaf	macstreet	9
stefan	unistreet	0
alekh	unistreet	7
felix	macstreet	5
jorge	minstreet	9

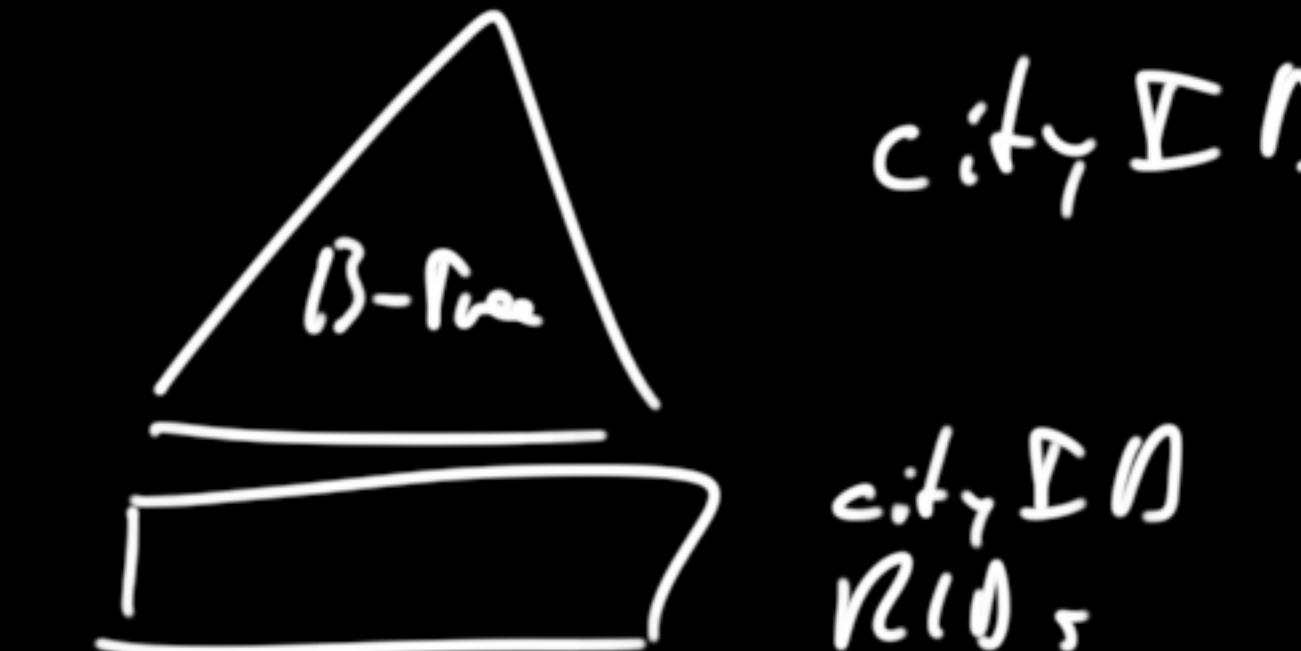
Cities_Dictionary	
<u>cityID</u>	city
5	berlin
1	cuppertino
9	saarbruecken
3	paris
0	new york
7	london

# Sort Both Tables on Join Attributes

Customers		
<u>name</u>	street	cityID
frank	minstreet	0
peter	minstreet	0
stefan	unistreet	0
jens	shortstreet	1
steve	macstreet	1
felix	macstreet	5
hans	msstreet	5
alekh	unistreet	7
jorge	minstreet	9
mike	longstreet	9
olaf	macstreet	9
tim	unistreet	9

Cities_Dictionary	
<u>cityID</u>	city
0	new york
1	cuppertino
3	paris
5	berlin
7	london
9	saarbruecken

① Index on one of the join attributes



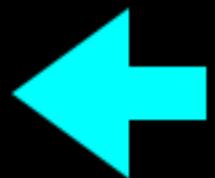
② "Interesting order"  
one or both already sorted

# Values Pointed to are Equal!

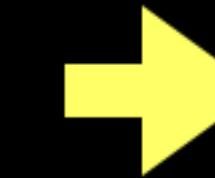
R

Customers		
<u>name</u>	street	cityID
frank	minstreet	0
peter	minstreet	0
stefan	unistreet	0
jens	shortstreet	1
steve	macstreet	1
felix	macstreet	5
hans	msstreet	5
alekh	unistreet	7
jorge	minstreet	9
mike	longstreet	9
olaf	macstreet	9
tim	unistreet	9

PR



PS



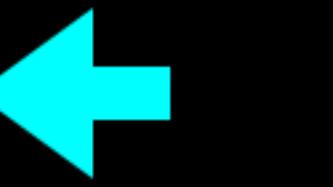
2

Cities_Dictionary	
cityID	city
0	new york
1	cuppertino
3	paris
5	berlin
7	london
9	saarbruecken

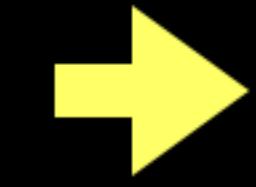
# Found a Join Result

have left pointer

Customers		
<u>name</u>	street	cityID
frank	minstreet	0
peter	minstreet	0
stefan	unistreet	0
jens	shortstreet	1
steve	macstreet	1
felix	macstreet	5
hans	msstreet	5
alekh	unistreet	7
jorge	minstreet	9
mike	longstreet	9
olaf	macstreet	9
tim	unistreet	9

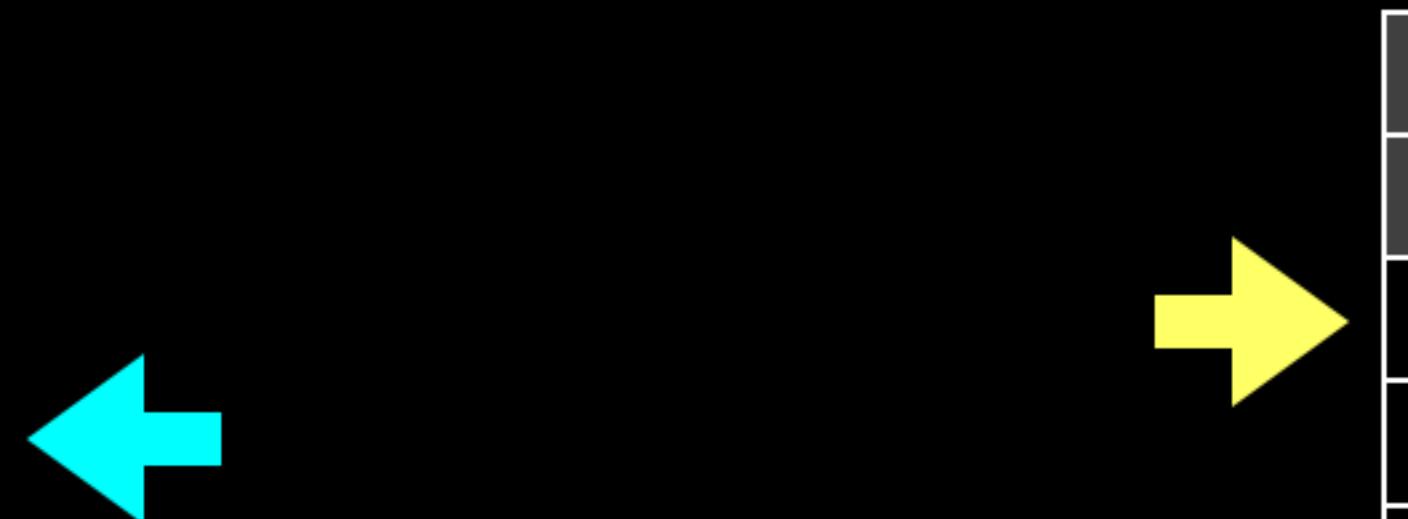


Cities_Dictionary	
<u>cityID</u>	city
0	new york
1	cuppertino
3	paris
5	berlin
7	london
9	saarbruecken



# Values Pointed to are Equal!

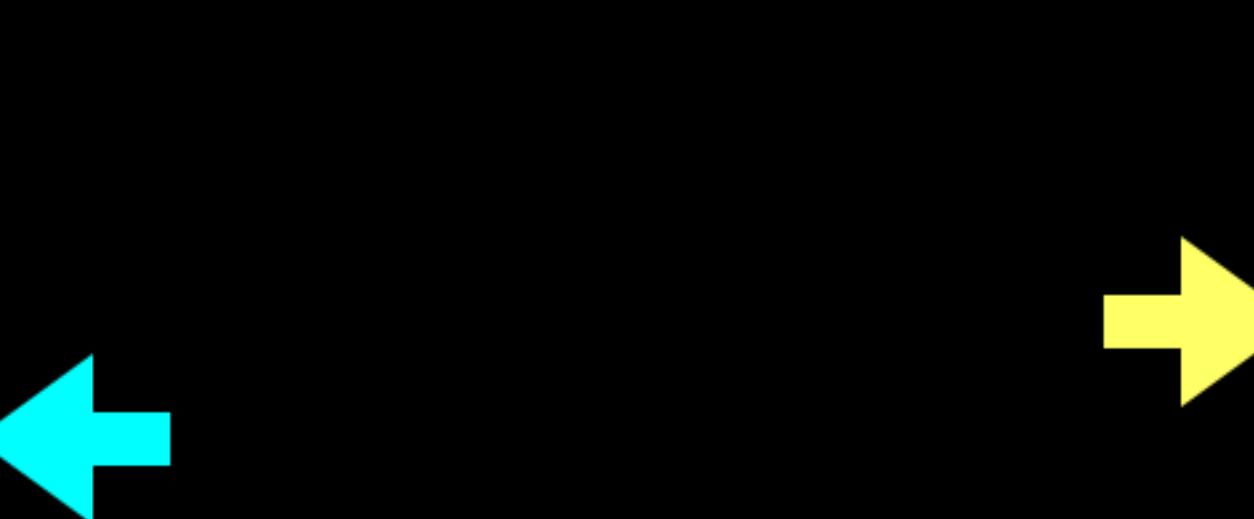
Customers		
<u>name</u>	street	cityID
frank	minstreet	0
peter	minstreet	0
stefan	unistreet	0
jens	shortstreet	1
steve	macstreet	1
felix	macstreet	5
hans	msstreet	5
alekh	unistreet	7
jorge	minstreet	9
mike	longstreet	9
olaf	macstreet	9
tim	unistreet	9



Cities_Dictionary	
<u>cityID</u>	city
0	new york
1	cuppertino
3	paris
5	berlin
7	london
9	saarbruecken

# Found a Join Result

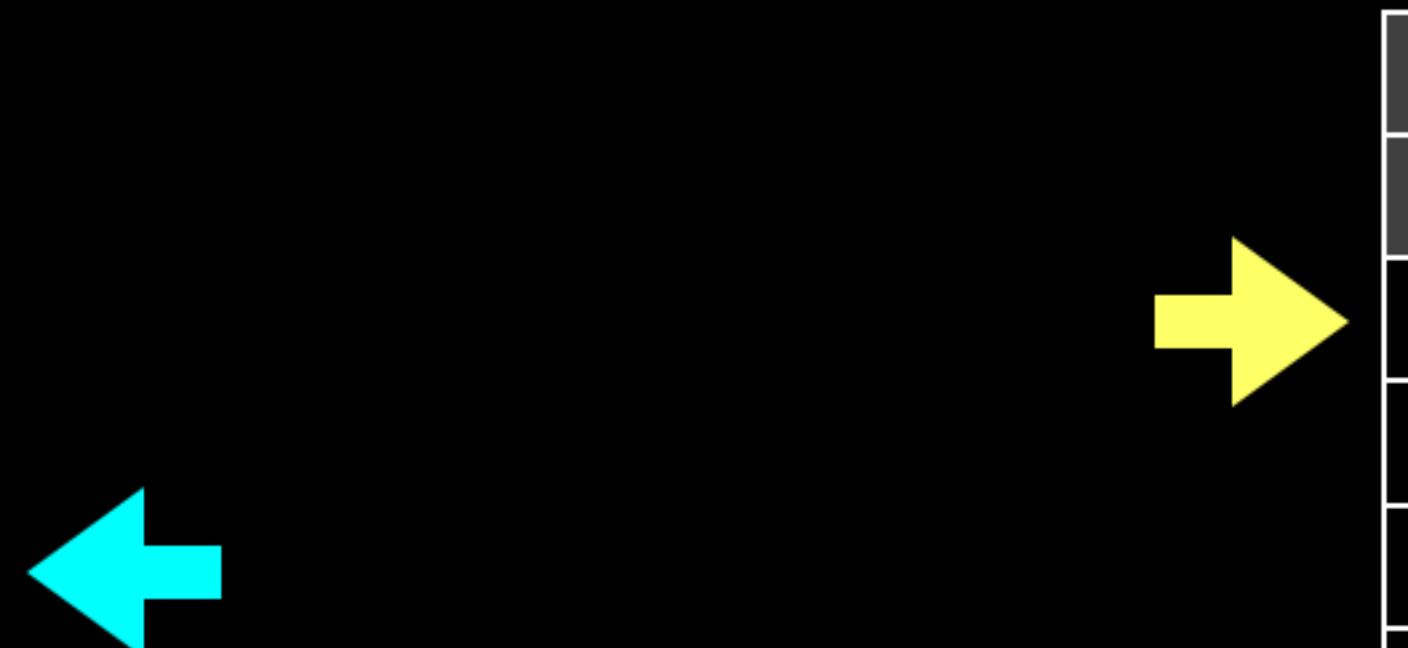
Customers		
<u>name</u>	street	cityID
frank	minstreet	0
peter	minstreet	0
stefan	unistreet	0
jens	shortstreet	1
steve	macstreet	1
felix	macstreet	5
hans	msstreet	5
alekh	unistreet	7
jorge	minstreet	9
mike	longstreet	9
olaf	macstreet	9
tim	unistreet	9



Cities_Dictionary	
<u>cityID</u>	city
0	new york
1	cuppertino
3	paris
5	berlin
7	london
9	saarbruecken

# Values Pointed to are Equal!

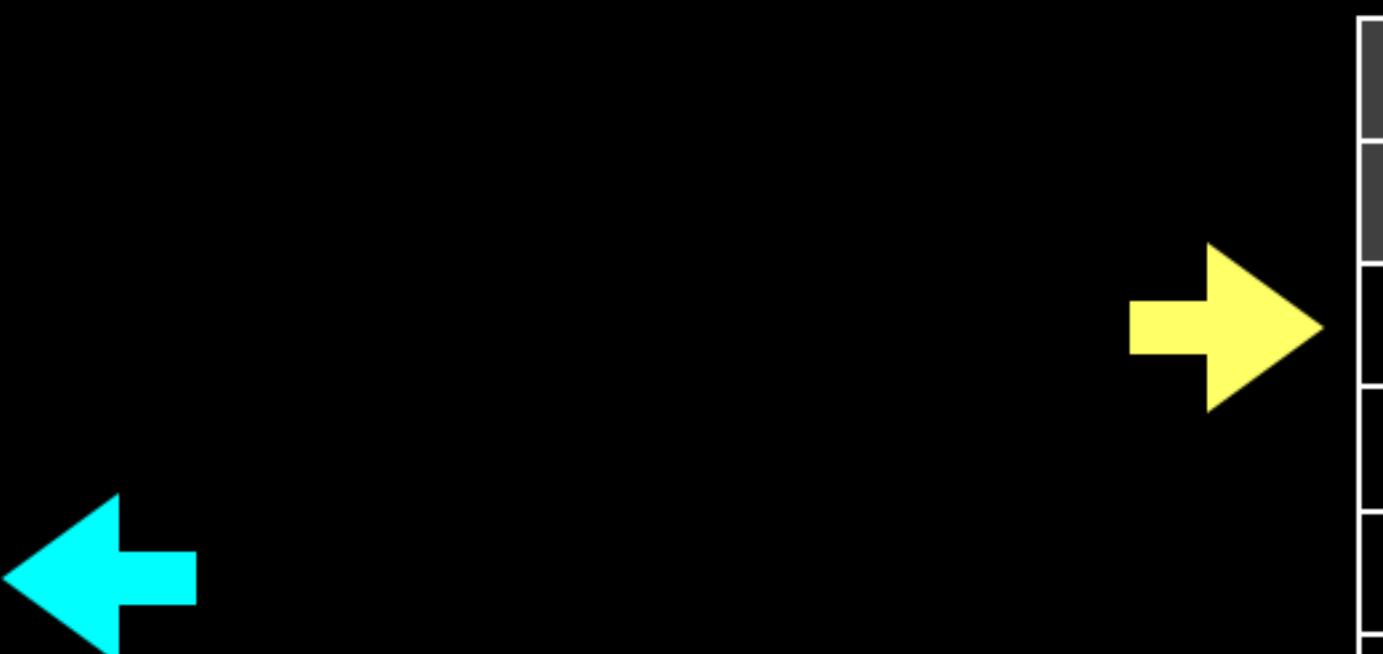
Customers		
<u>name</u>	street	cityID
frank	minstreet	0
peter	minstreet	0
stefan	unistreet	0
jens	shortstreet	1
steve	macstreet	1
felix	macstreet	5
hans	msstreet	5
alekh	unistreet	7
jorge	minstreet	9
mike	longstreet	9
olaf	macstreet	9
tim	unistreet	9



Cities_Dictionary	
<u>cityID</u>	city
0	new york
1	cuppertino
3	paris
5	berlin
7	london
9	saarbruecken

# Found a Join Result

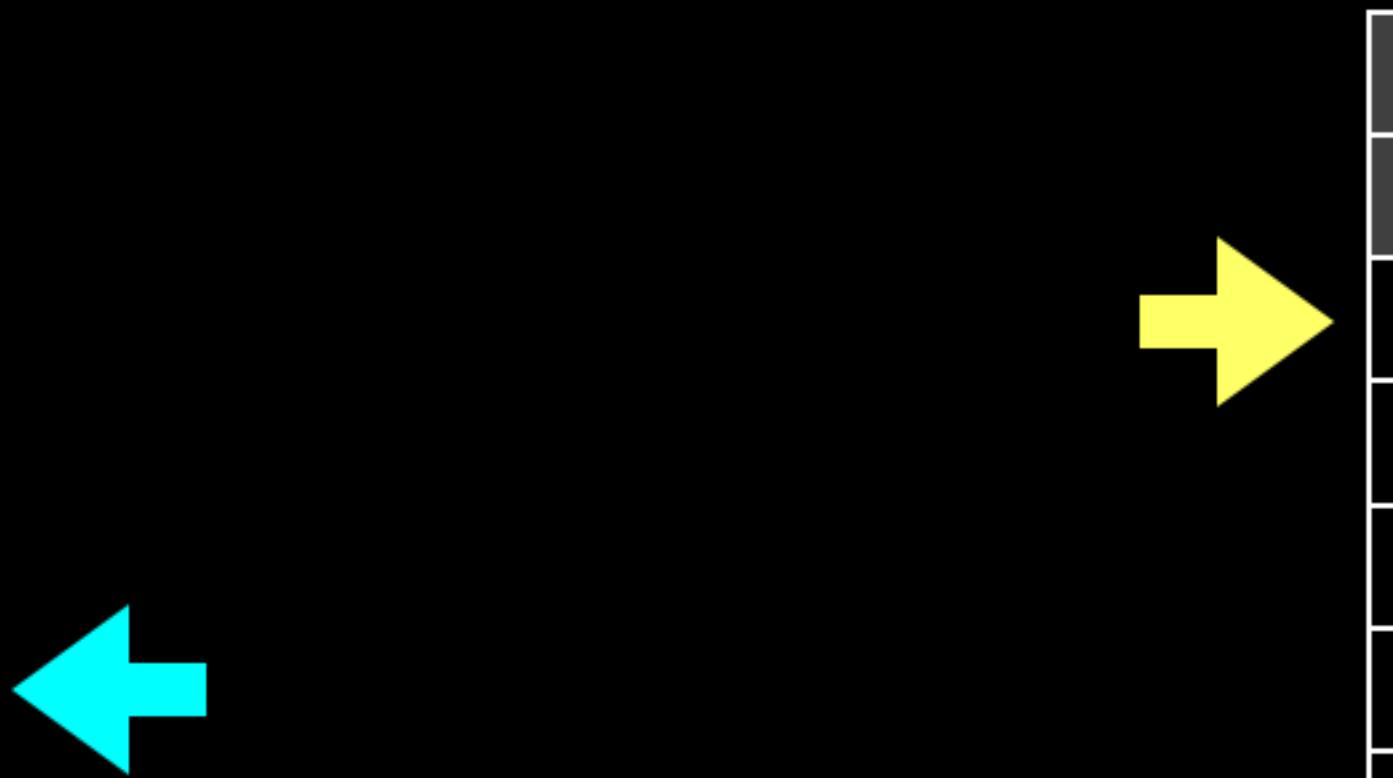
Customers		
<u>name</u>	street	cityID
frank	minstreet	0
peter	minstreet	0
stefan	unistreet	0
jens	shortstreet	1
steve	macstreet	1
felix	macstreet	5
hans	msstreet	5
alekh	unistreet	7
jorge	minstreet	9
mike	longstreet	9
olaf	macstreet	9
tim	unistreet	9



Cities_Dictionary	
<u>cityID</u>	city
0	new york
1	cuppertino
3	paris
5	berlin
7	london
9	saarbruecken

# Values Pointed to are NOT Equal!

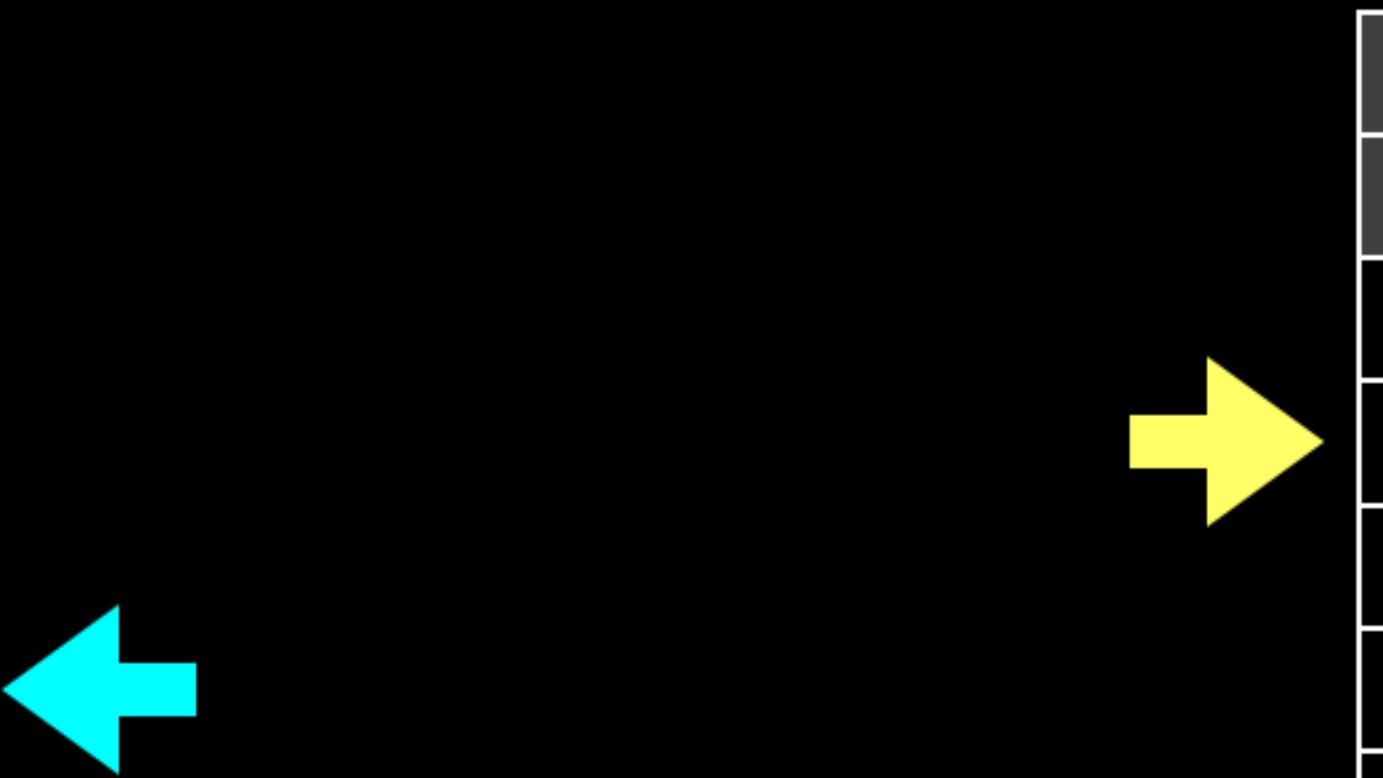
Customers		
<u>name</u>	street	cityID
frank	minstreet	0
peter	minstreet	0
stefan	unistreet	0
jens	shortstreet	1
steve	macstreet	1
felix	macstreet	5
hans	msstreet	5
alekh	unistreet	7
jorge	minstreet	9
mike	longstreet	9
olaf	macstreet	9
tim	unistreet	9



Cities_Dictionary	
<u>cityID</u>	city
0	new york
1	cuppertino
3	paris
5	berlin
7	london
9	saarbruecken

# Values Pointed to are Equal!

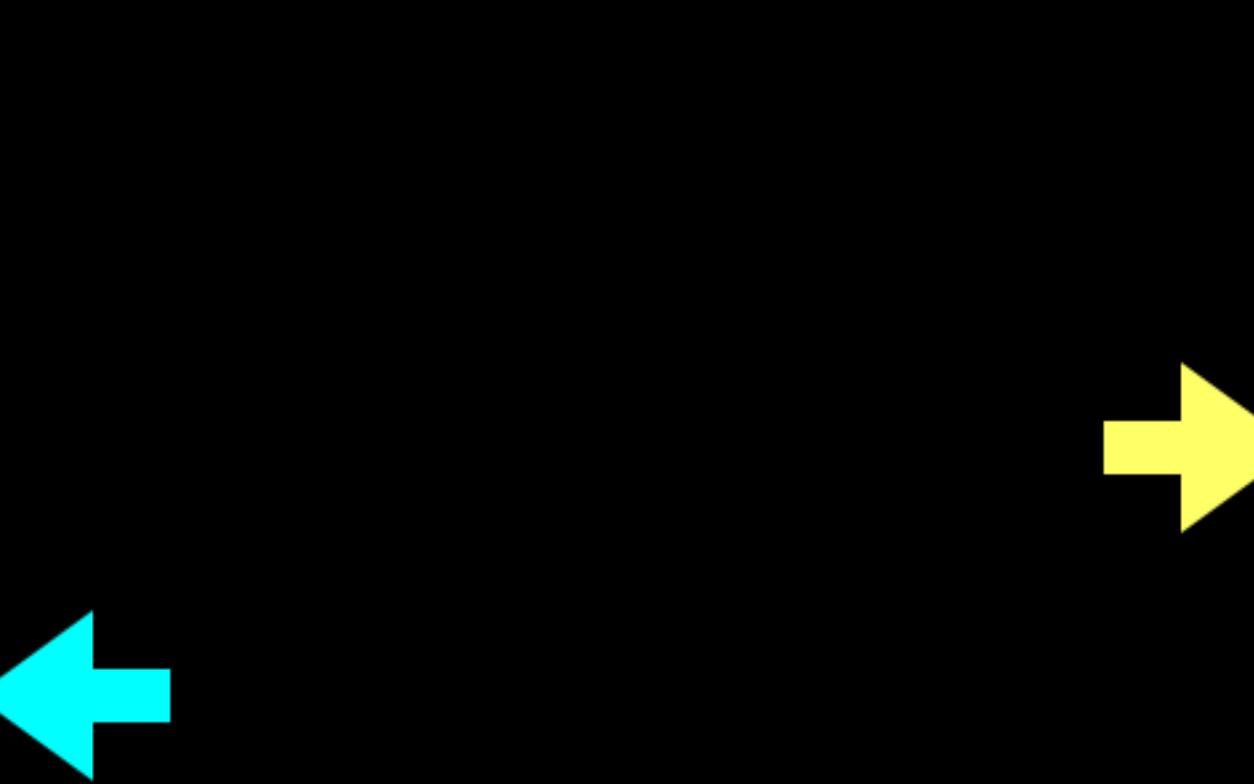
Customers		
<u>name</u>	street	cityID
frank	minstreet	0
peter	minstreet	0
stefan	unistreet	0
jens	shortstreet	1
steve	macstreet	1
felix	macstreet	5
hans	msstreet	5
alekh	unistreet	7
jorge	minstreet	9
mike	longstreet	9
olaf	macstreet	9
tim	unistreet	9



Cities_Dictionary	
<u>cityID</u>	city
0	new york
1	cuppertino
3	paris
5	berlin
7	london
9	saarbruecken

# Found a Join Result

Customers		
<u>name</u>	street	cityID
frank	minstreet	0
peter	minstreet	0
stefan	unistreet	0
jens	shortstreet	1
steve	macstreet	1
felix	macstreet	5
hans	msstreet	5
alekh	unistreet	7
jorge	minstreet	9
mike	longstreet	9
olaf	macstreet	9
tim	unistreet	9

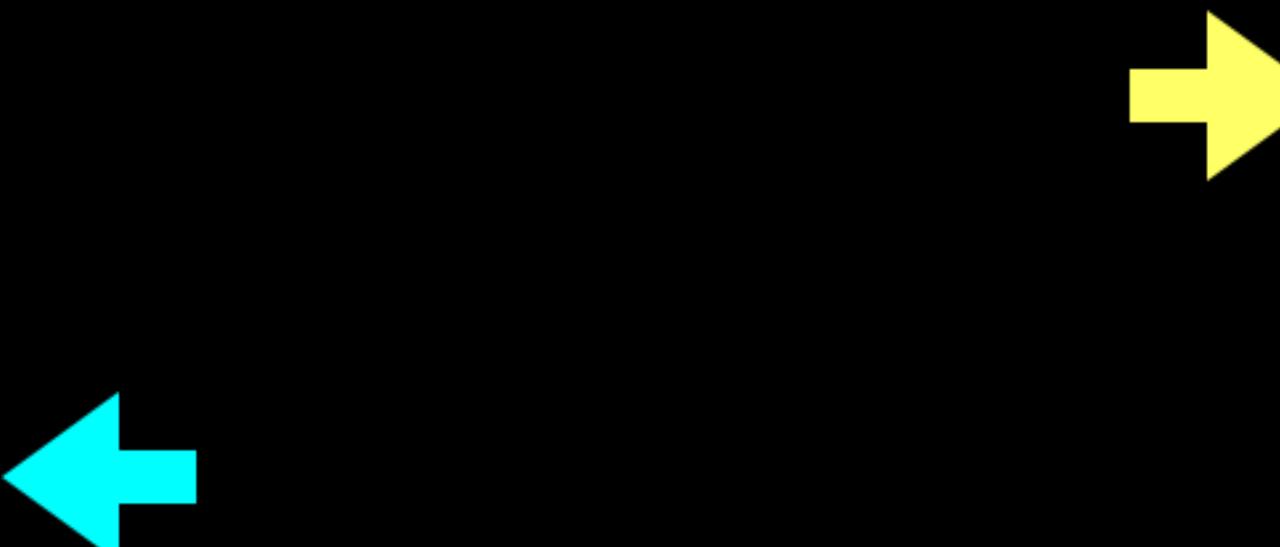


Cities_Dictionary	
<u>cityID</u>	city
0	new york
1	cuppertino
3	paris
5	berlin
7	london
9	saarbruecken

# Values Pointed to are Equal!

Customers		
<u>name</u>	street	cityID
frank	minstreet	0
peter	minstreet	0
stefan	unistreet	0
jens	shortstreet	1
steve	macstreet	1
felix	macstreet	5
hans	msstreet	5
alekh	unistreet	7
jorge	minstreet	9
mike	longstreet	9
olaf	macstreet	9
tim	unistreet	9

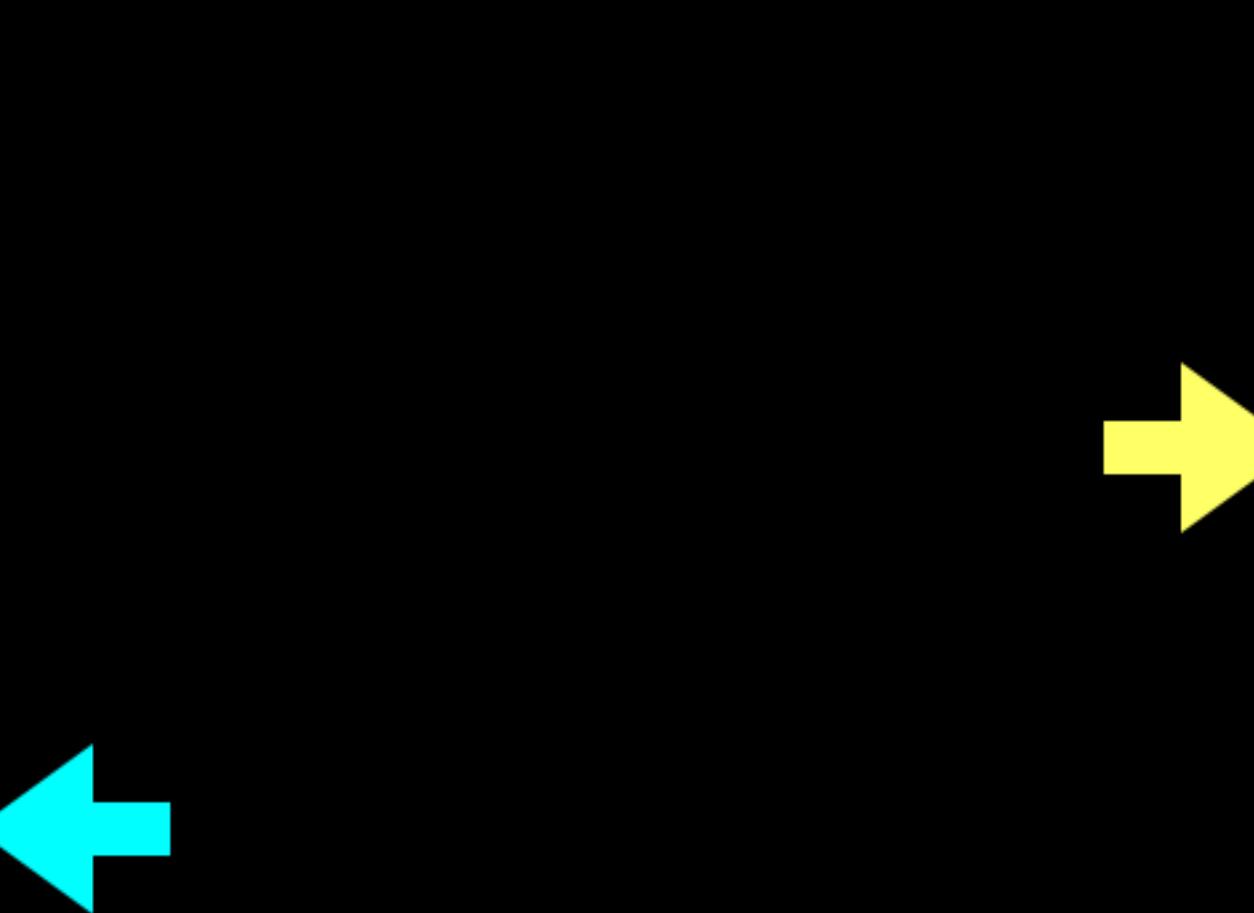
Cities_Dictionary	
<u>cityID</u>	city
0	new york
1	cuppertino
3	paris
5	berlin
7	london
9	saarbruecken



Customers NATURAL JOIN Cities_Dictionary			
<u>name</u>	street	cityID	city
frank	minstreet	0	new york
peter	minstreet	0	new york
stefan	unistreet	0	new york
jens	shortstreet	1	cuppertino

# Found a Join Result

Customers		
<u>name</u>	street	cityID
frank	minstreet	0
peter	minstreet	0
stefan	unistreet	0
jens	shortstreet	1
steve	macstreet	1
felix	macstreet	5
hans	msstreet	5
alekh	unistreet	7
jorge	minstreet	9
mike	longstreet	9
olaf	macstreet	9
tim	unistreet	9

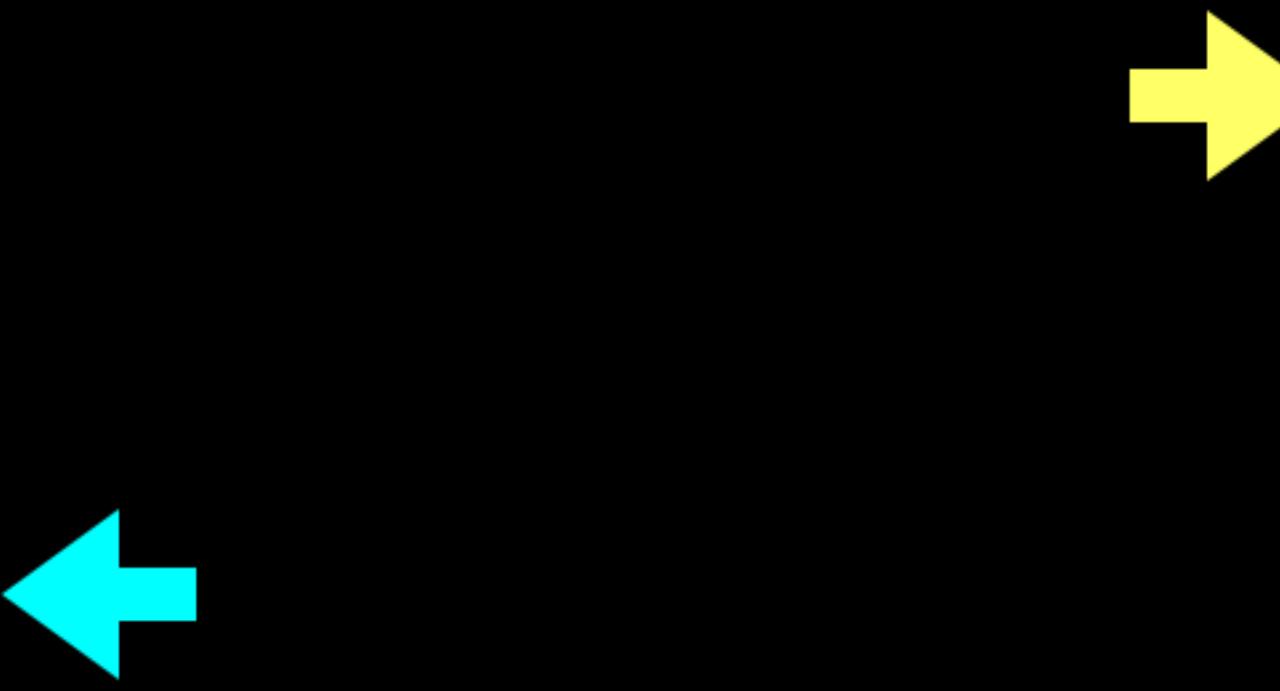


Cities_Dictionary	
<u>cityID</u>	city
0	new york
1	cuppertino
3	paris
5	berlin
7	london
9	saarbruecken

Customers NATURAL JOIN Cities_Dictionary			
<u>name</u>	street	cityID	city
frank	minstreet	0	new york
peter	minstreet	0	new york
stefan	unistreet	0	new york
jens	shortstreet	1	cuppertino
steve	macstreet	1	cuppertino

# Values Pointed to are NOT Equal!

Customers		
<u>name</u>	street	cityID
frank	minstreet	0
peter	minstreet	0
stefan	unistreet	0
jens	shortstreet	1
steve	macstreet	1
felix	macstreet	5
hans	msstreet	5
alekh	unistreet	7
jorge	minstreet	9
mike	longstreet	9
olaf	macstreet	9
tim	unistreet	9



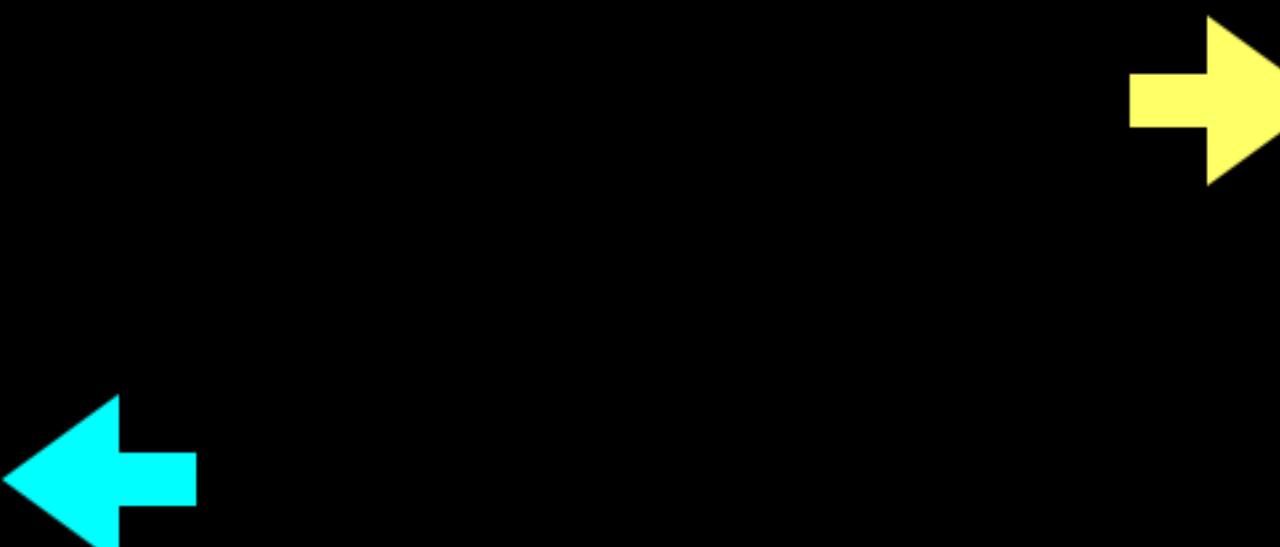
Cities_Dictionary	
<u>cityID</u>	city
0	new york
1	cuppertino
3	paris
5	berlin
7	london
9	saarbruecken

Customers NATURAL JOIN Cities_Dictionary			
<u>name</u>	street	cityID	city
frank	minstreet	0	new york
peter	minstreet	0	new york
stefan	unistreet	0	new york
jens	shortstreet	1	cuppertino
steve	macstreet	1	cuppertino

# Values Pointed to are NOT Equal!

Customers		
<u>name</u>	street	cityID
frank	minstreet	0
peter	minstreet	0
stefan	unistreet	0
jens	shortstreet	1
steve	macstreet	1
felix	macstreet	5
hans	msstreet	5
alekh	unistreet	7
jorge	minstreet	9
mike	longstreet	9
olaf	macstreet	9
tim	unistreet	9

Cities_Dictionary	
<u>cityID</u>	city
0	new york
1	cuppertino
3	paris
5	berlin
7	london
9	saarbruecken

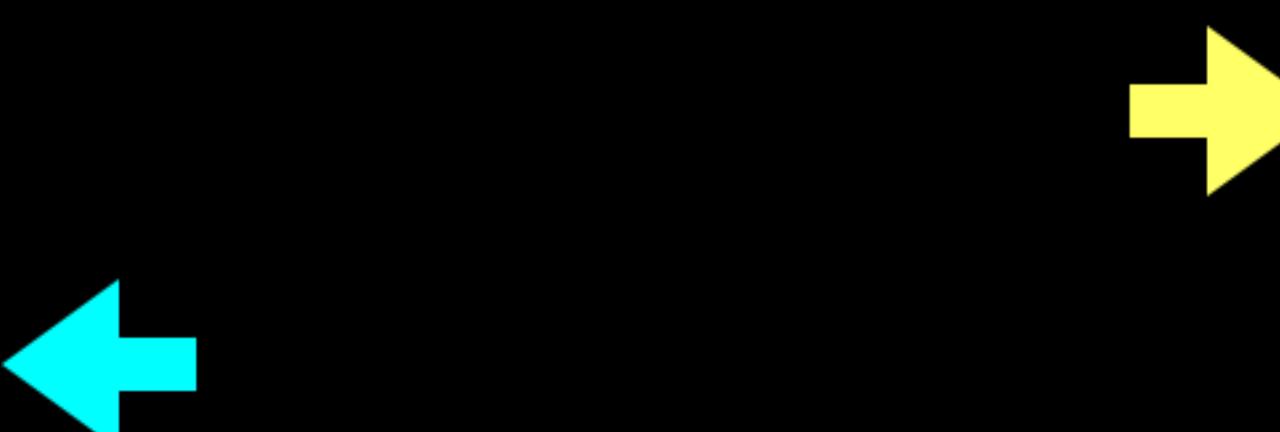


Customers NATURAL JOIN Cities_Dictionary			
<u>name</u>	street	cityID	city
frank	minstreet	0	new york
peter	minstreet	0	new york
stefan	unistreet	0	new york
jens	shortstreet	1	cuppertino
steve	macstreet	1	cuppertino

# Values Pointed to are Equal!

Customers		
<u>name</u>	street	cityID
frank	minstreet	0
peter	minstreet	0
stefan	unistreet	0
jens	shortstreet	1
steve	macstreet	1
felix	macstreet	5
hans	msstreet	5
alekh	unistreet	7
jorge	minstreet	9
mike	longstreet	9
olaf	macstreet	9
tim	unistreet	9

Cities_Dictionary	
<u>cityID</u>	city
0	new york
1	cuppertino
3	paris
5	berlin
7	london
9	saarbruecken

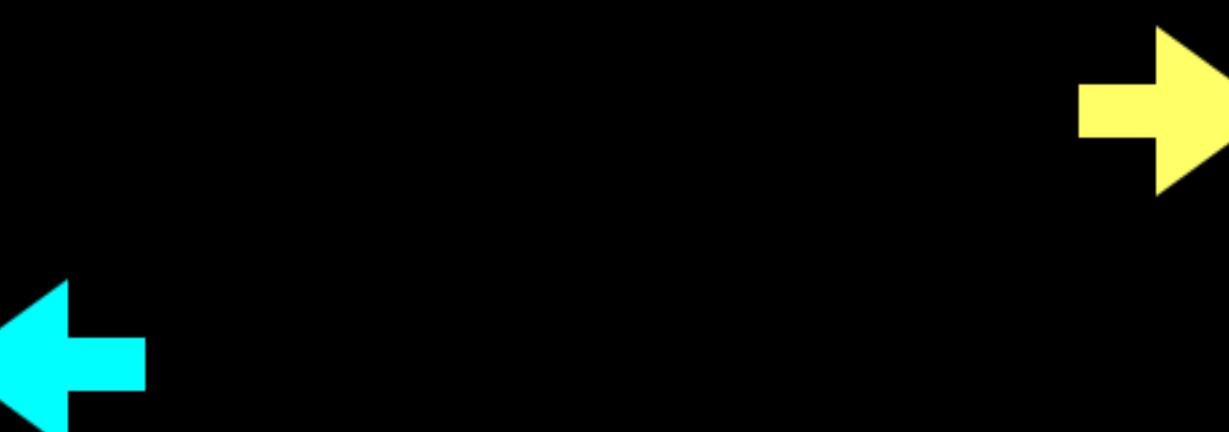


Customers NATURAL JOIN Cities_Dictionary			
<u>name</u>	street	cityID	city
frank	minstreet	0	new york
peter	minstreet	0	new york
stefan	unistreet	0	new york
jens	shortstreet	1	cuppertino
steve	macstreet	1	cuppertino

# Found a Join Result

Customers		
<u>name</u>	street	cityID
frank	minstreet	0
peter	minstreet	0
stefan	unistreet	0
jens	shortstreet	1
steve	macstreet	1
felix	macstreet	5
hans	msstreet	5
alekh	unistreet	7
jorge	minstreet	9
mike	longstreet	9
olaf	macstreet	9
tim	unistreet	9

Cities_Dictionary	
<u>cityID</u>	city
0	new york
1	cuppertino
3	paris
5	berlin
7	london
9	saarbruecken

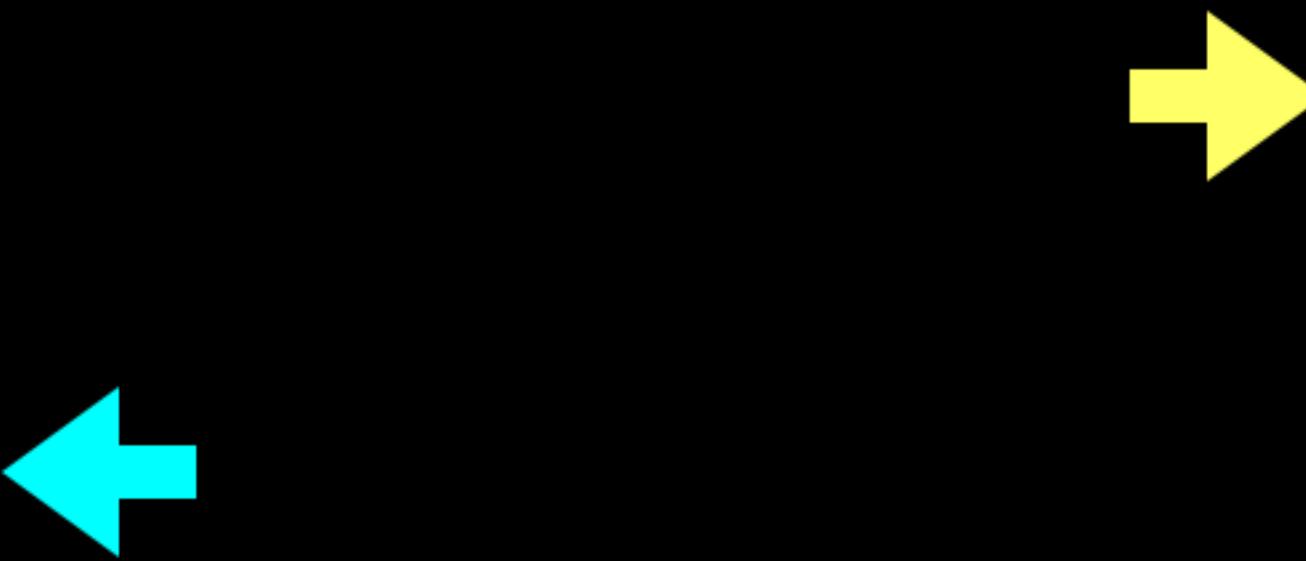


Customers NATURAL JOIN Cities_Dictionary			
<u>name</u>	street	cityID	city
frank	minstreet	0	new york
peter	minstreet	0	new york
stefan	unistreet	0	new york
jens	shortstreet	1	cuppertino
steve	macstreet	1	cuppertino
felix	macstreet	5	berlin

# Values Pointed to are Equal!

Customers		
<u>name</u>	street	cityID
frank	minstreet	0
peter	minstreet	0
stefan	unistreet	0
jens	shortstreet	1
steve	macstreet	1
felix	macstreet	5
hans	msstreet	5
alekh	unistreet	7
jorge	minstreet	9
mike	longstreet	9
olaf	macstreet	9
tim	unistreet	9

Cities_Dictionary	
<u>cityID</u>	city
0	new york
1	cuppertino
3	paris
5	berlin
7	london
9	saarbruecken

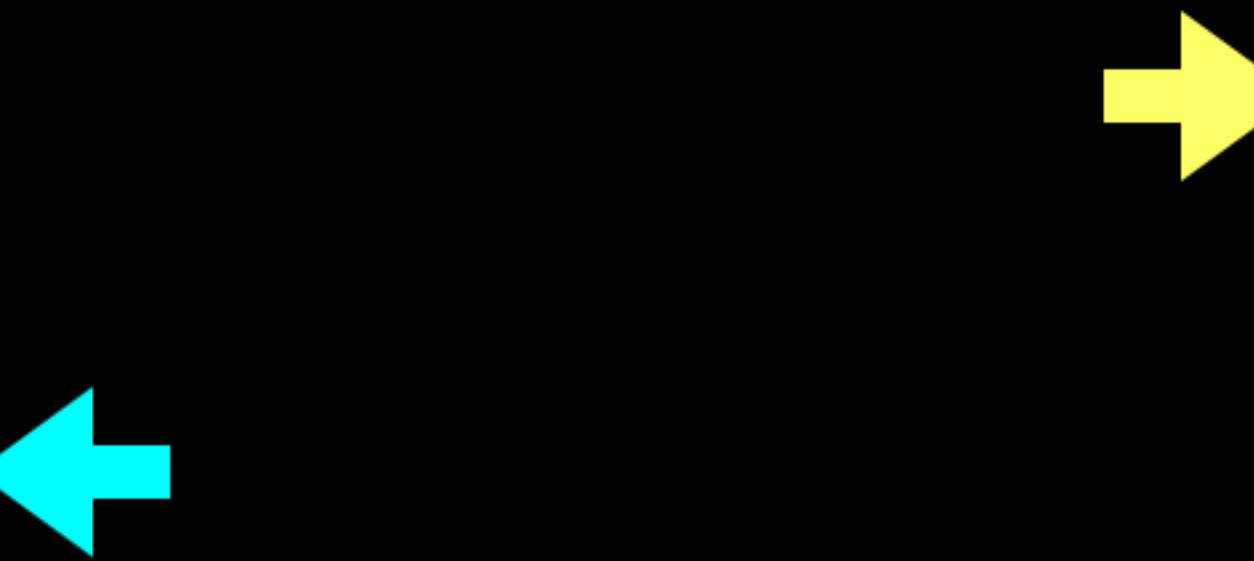


Customers NATURAL JOIN Cities_Dictionary			
<u>name</u>	street	cityID	city
frank	minstreet	0	new york
peter	minstreet	0	new york
stefan	unistreet	0	new york
jens	shortstreet	1	cuppertino
steve	macstreet	1	cuppertino
felix	macstreet	5	berlin

# Found a Join Result

Customers		
<u>name</u>	street	cityID
frank	minstreet	0
peter	minstreet	0
stefan	unistreet	0
jens	shortstreet	1
steve	macstreet	1
felix	macstreet	5
hans	msstreet	5
alekh	unistreet	7
jorge	minstreet	9
mike	longstreet	9
olaf	macstreet	9
tim	unistreet	9

Cities_Dictionary	
<u>cityID</u>	city
0	new york
1	cuppertino
3	paris
5	berlin
7	london
9	saarbruecken



Customers NATURAL JOIN Cities_Dictionary			
<u>name</u>	street	cityID	city
frank	minstreet	0	new york
peter	minstreet	0	new york
stefan	unistreet	0	new york
jens	shortstreet	1	cuppertino
steve	macstreet	1	cuppertino
felix	macstreet	5	berlin
hans	msstreet	5	berlin

# Sort-Merge Join

extensions: intervals, spatial, temporal

R

S

$JP(r,s) := r.x == s.x$

//definition of the join predicate

# Sort-Merge Join

R

S

**JP(r,s)** := r.x == s.x

//definition of the join predicate

**SortMergeJoin( R, S, JP(r,s) ):**

```
sort(R on R.x);           //sort R on join attribute x
sort(S on S.x);           //sort S on join attribute x
Pointer PR = R[0];         //initialize pointer to R
Pointer PS = S[0];         //initialize pointer to S
do:
  if PR.x == PS.x:         //start loop
    output( PR, PS );       //if values match
    PR++;                  //output join result
    PS++;
```

if PR.x == PS.x:

  output( PR, PS );

How to project?

# Sort-Merge Join

R

S

**JP(r,s)** := r.x == s.x

//definition of the join predicate

**SortMergeJoin( R, S, JP(r,s) ):**

```
sort(R on R.x);           //sort R on join attribute x
sort(S on S.x);           //sort S on join attribute x
Pointer PR = R[0];         //initialize pointer to R
Pointer PS = S[0];         //initialize pointer to S
do:
    if PR.x == PS.x:       //start loop
        output( (PR, PS) ); //if values match
        //output join result
    if PR.x <= PS.x:       //find smaller join key (move left first!)
        PR++;               //move pointer to R forward
    else: if PR.x > PS.x   //move pointer to S forward
        PS++;
```

# Sort-Merge Join

R

S

**JP(r,s)** := r.x == s.x

//definition of the join predicate

**SortMergeJoin( R, S, JP(r,s) ):**

```
sort(R on R.x);           //sort R on join attribute x
sort(S on S.x);           //sort S on join attribute x
Pointer PR = R[0];         //initialize pointer to R
Pointer PS = S[0];         //initialize pointer to S
do:
    if PR.x == PS.x:       //if values match
        output( (PR, PS) ); //output join result
    if PR.x <= PS.x:       //find smaller join key (move left first!)
        PR++;               //move pointer to R forward
    else:
        PS++;               //move pointer to S forward
while PR != R.end && PS != S.end //both tables unprocessed
```

→ LA B

# Sort-Merge Join

R

S

**JP(r,s)** := r.x == s.x

//definition of the join predicate

**SortMergeJoin( R, S, JP(r,s) ):**

```
sort(R on R.x);                                //sort R on join attribute x
sort(S on S.x);                                //sort S on join attribute x
Pointer PR = R[0];                             //initialize pointer to R
Pointer PS = S[0];                             //initialize pointer to S
do:
    if PR.x == PS.x:                           //if values match
        output( (PR, PS) );                   //output join result
    if PR.x <= PS.x:                          //find smaller join key (move left first!)
        PR++;                                 //move pointer to R forward
    else:
        PS++;                                 //move pointer to S forward
while PR != R.end && PS != S.end           //both tables unprocessed
...
...
...
```

# Natural Join

*FK  
J*

Customers		
<u>name</u>	street	cityID
frank	minstreet	0
peter	minstreet	0
stefan	unistreet	0
jens	shortstreet	1
steve	macstreet	1
felix	macstreet	5
hans	msstreet	5
alekh	unistreet	7
jorge	minstreet	9
mike	longstreet	9
olaf	macstreet	9
tim	unistreet	9

*↑  
duplicates*

*J PK*

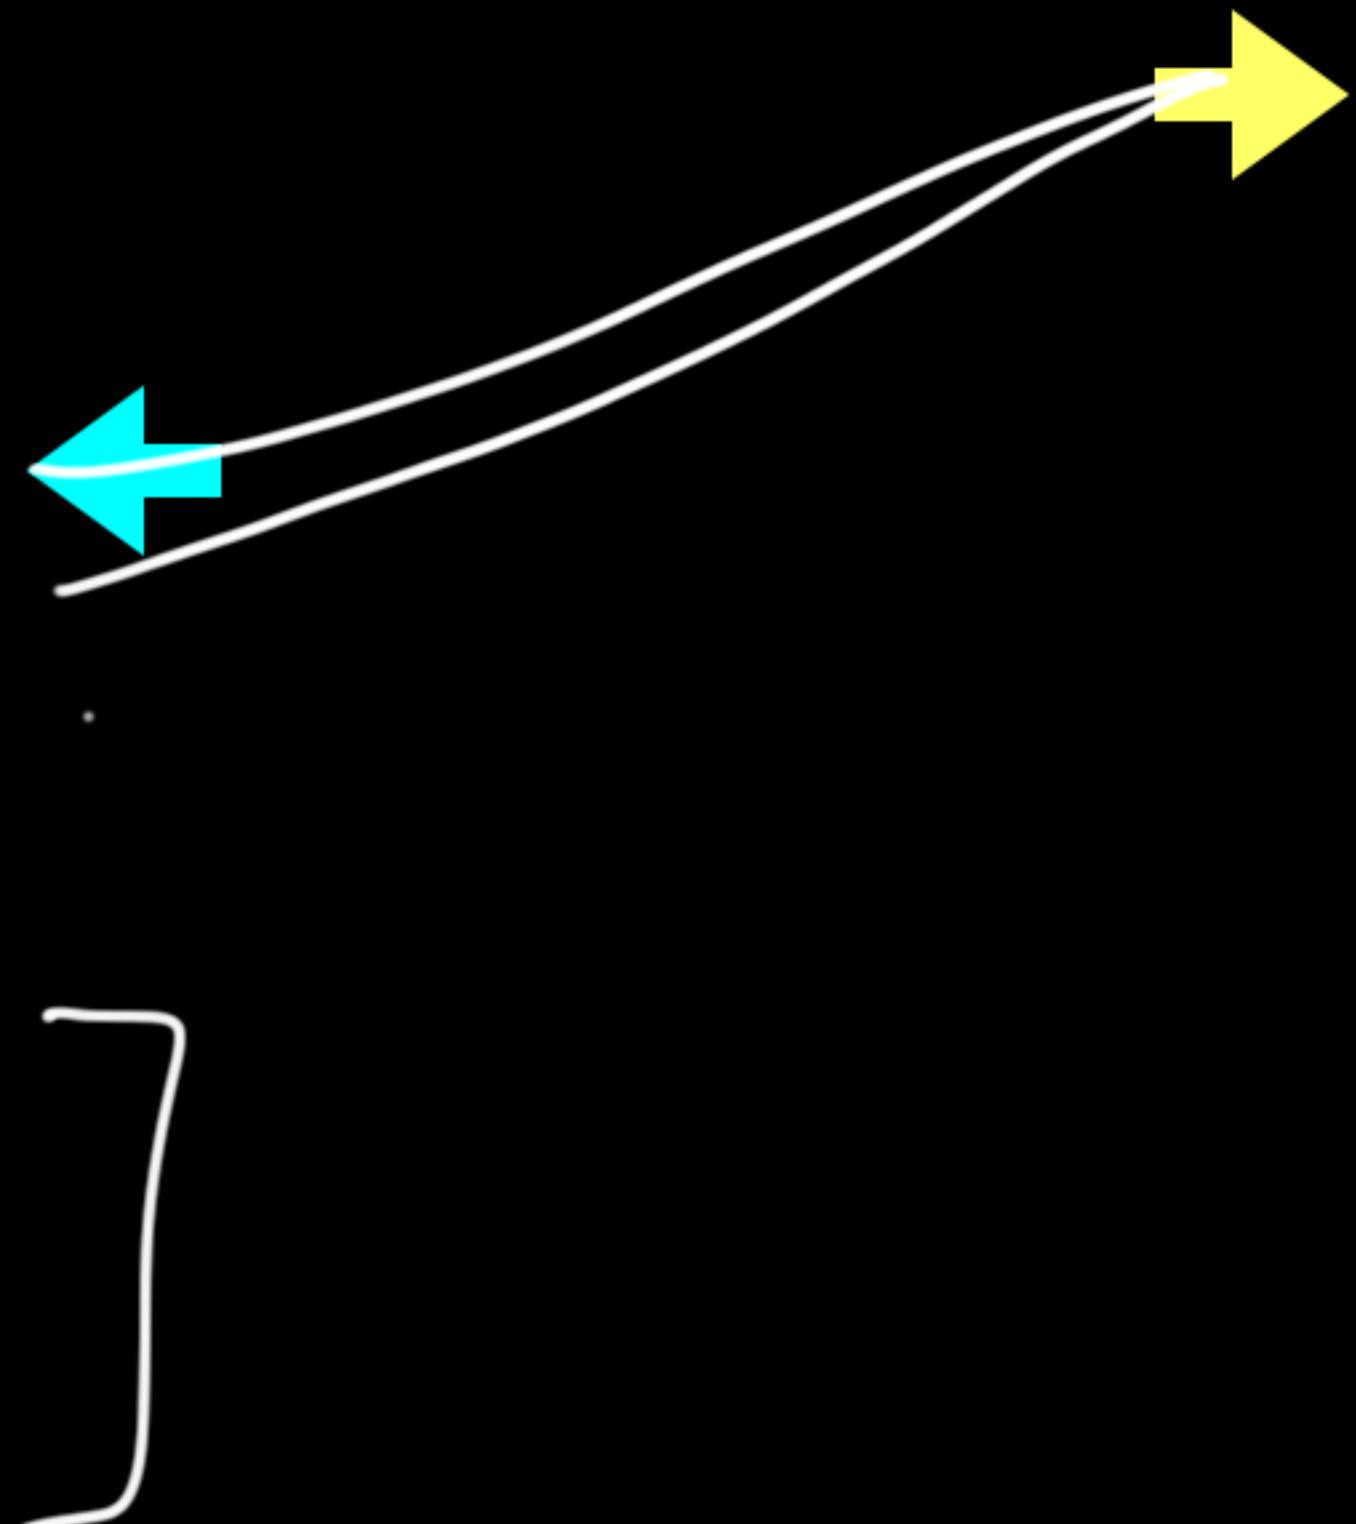
Cities_Dictionary	
<u>cityID</u>	city
0	new york
1	cuppertino
3	paris
5	berlin
7	london
9	saarbruecken

*↑  
no duplicates*

# Duplicates in both Columns!

Customers		
<u>name</u>	street	cityID
frank	minstreet	0
peter	minstreet	0
stefan	unistreet	0
jens	shortstreet	1
steve	macstreet	1
felix	macstreet	5
hans	msstreet	5
alekh	unistreet	7
jorge	minstreet	9
mike	longstreet	9
olaf	macstreet	9
tim	unistreet	9

Cities_Dictionary	
cityNo	city
1	new york
1	cuppertino
1	paris
5	berlin
5	london
9	saarbruecken



# Joins? CoGroups!

Customers		
<u>name</u>	street	cityID
frank	minstreet	0
peter	minstreet	
stefan	unistreet	
jens	shortstreet	1
steve	macstreet	
felix	macstreet	5
hans	msstreet	
alekh	unistreet	7
jorge	minstreet	
mike	longstreet	9
olaf	macstreet	
tim	unistreet	

Cities_Dictionary	
cityNo	city
	new york
1	cuppertino
	paris
5	berlin
	london
9	saarbruecken

SELECT cityID  
 FROM Customers  
 GROUP BY cityID;

$$\left. \begin{array}{l} 1 : 2 \times 3 = 6 \\ 5 : 2 \times 2 = 4 \\ 9 : 4 \times 1 = 4 \end{array} \right\} \rightarrow 14 \text{ tuples}$$

SELECT cityNo  
 FROM Cities\_Dictionary  
 GROUP BY cityNo