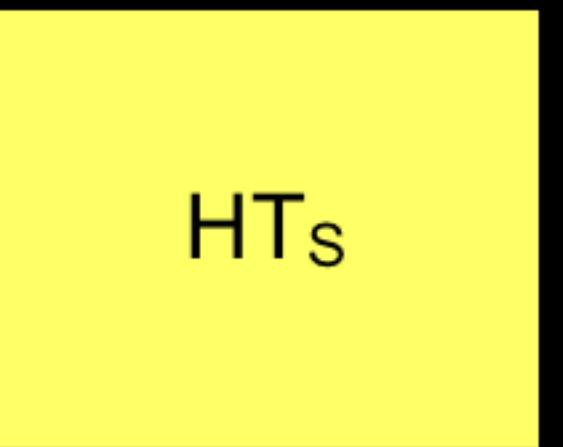


# Example: Double-pipelined Hash Join

hash-table for R



hash-table for S



$r_1$

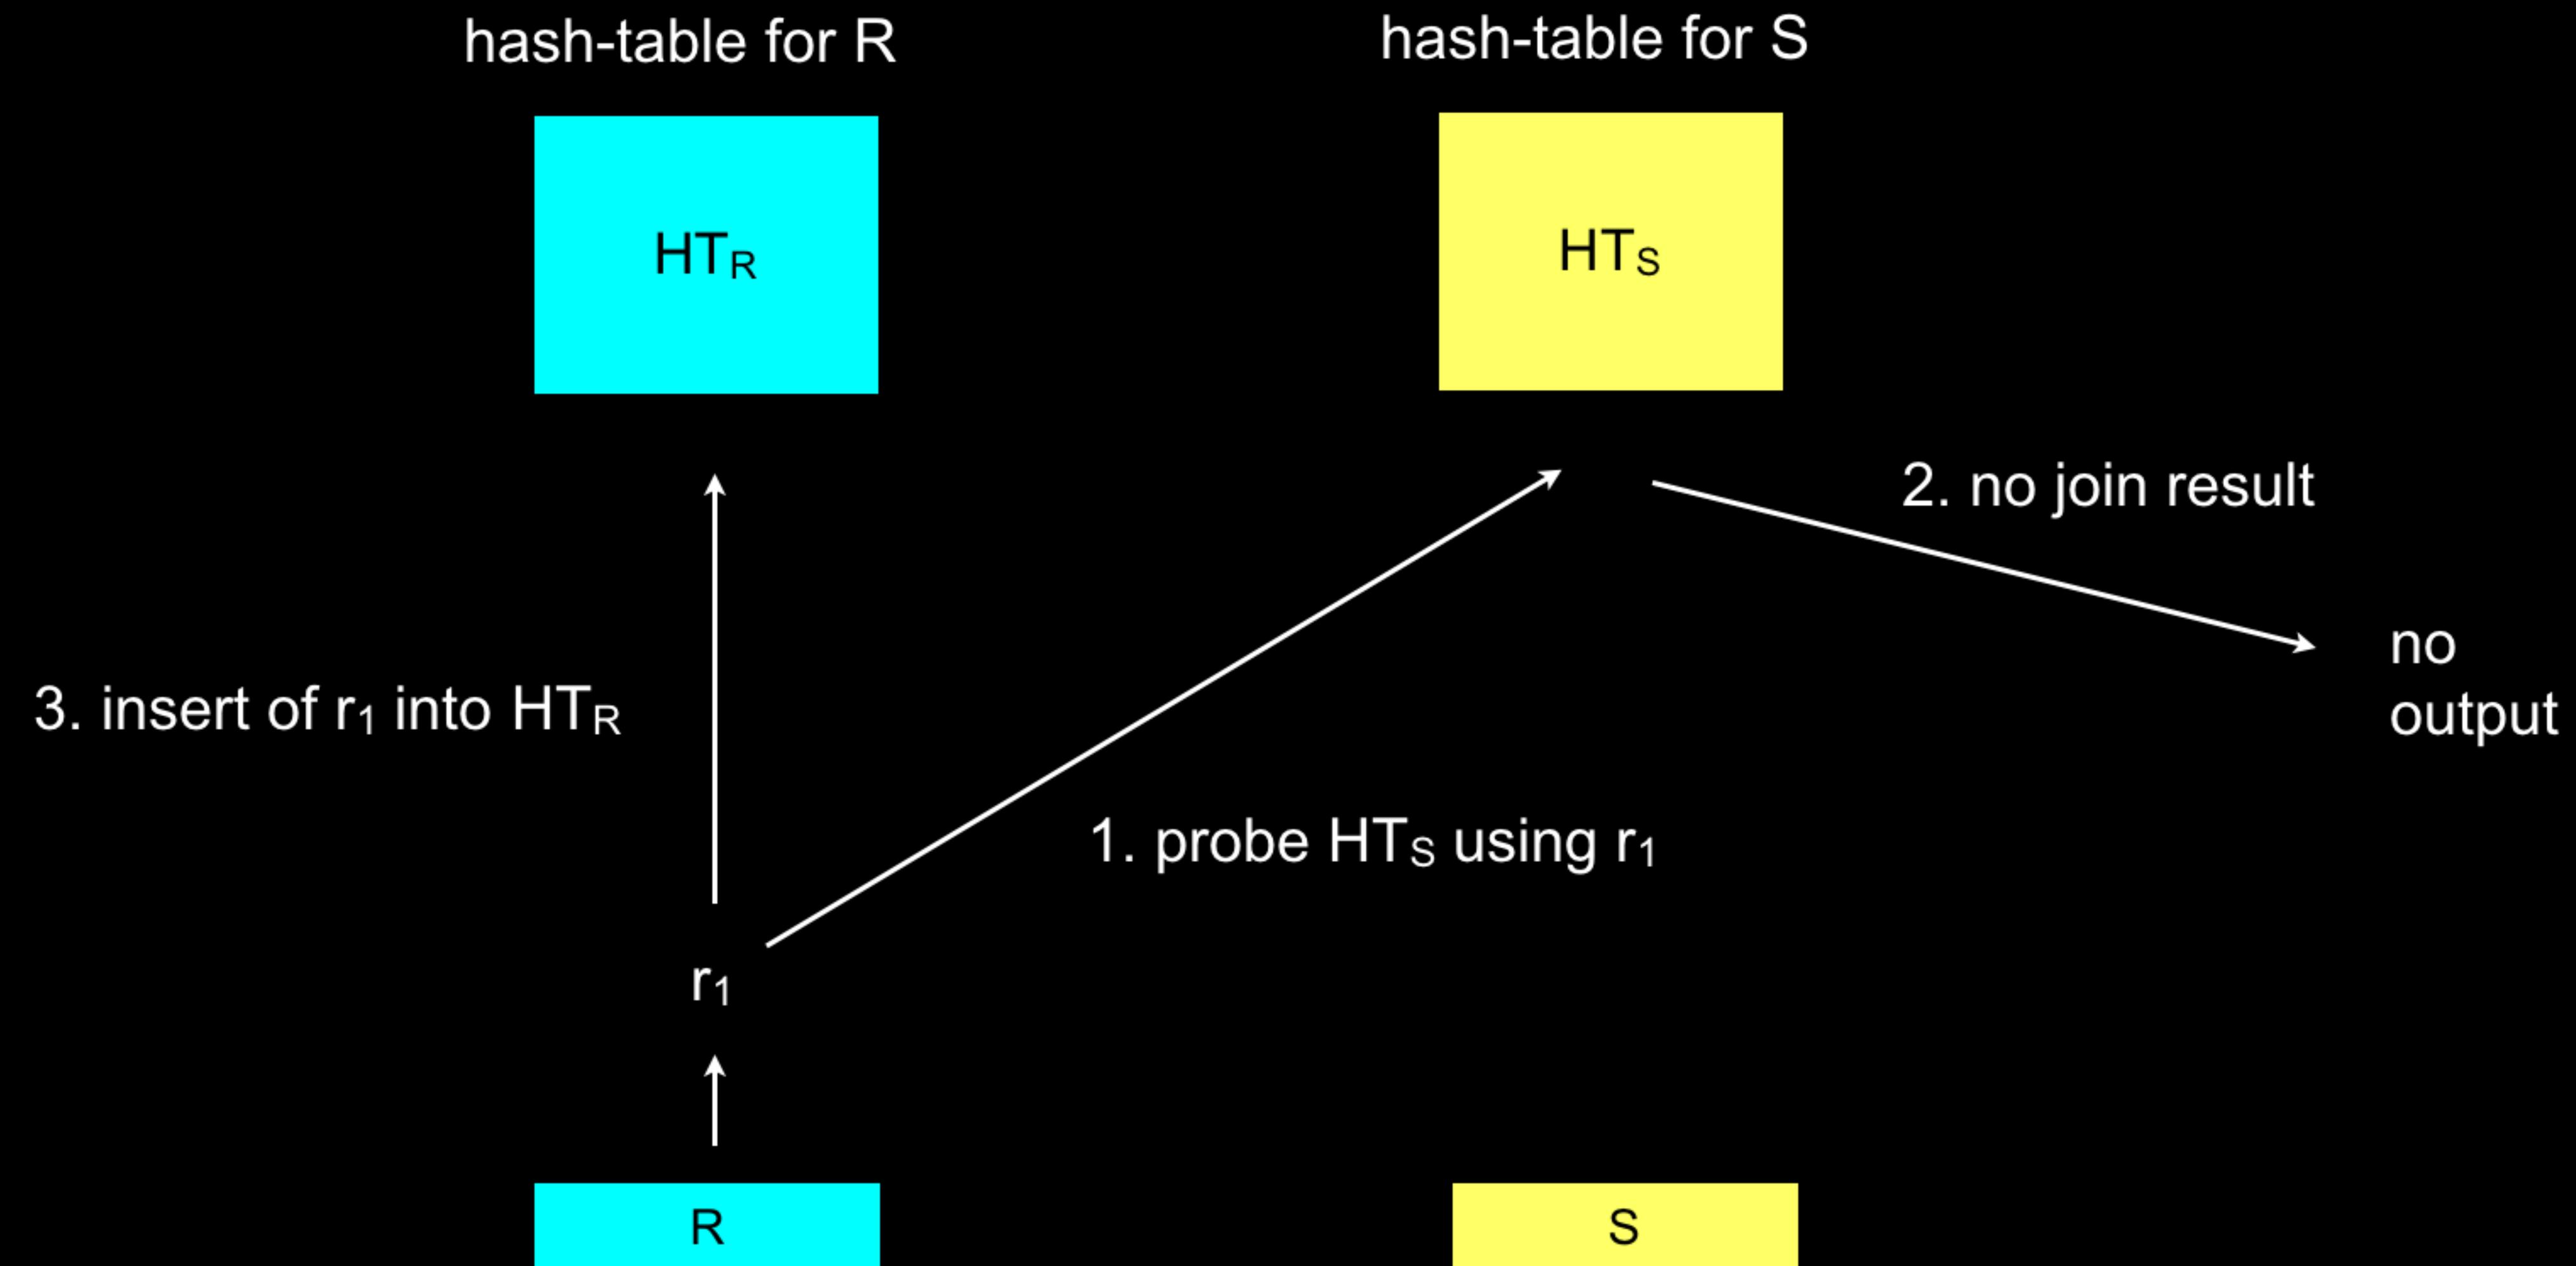


R

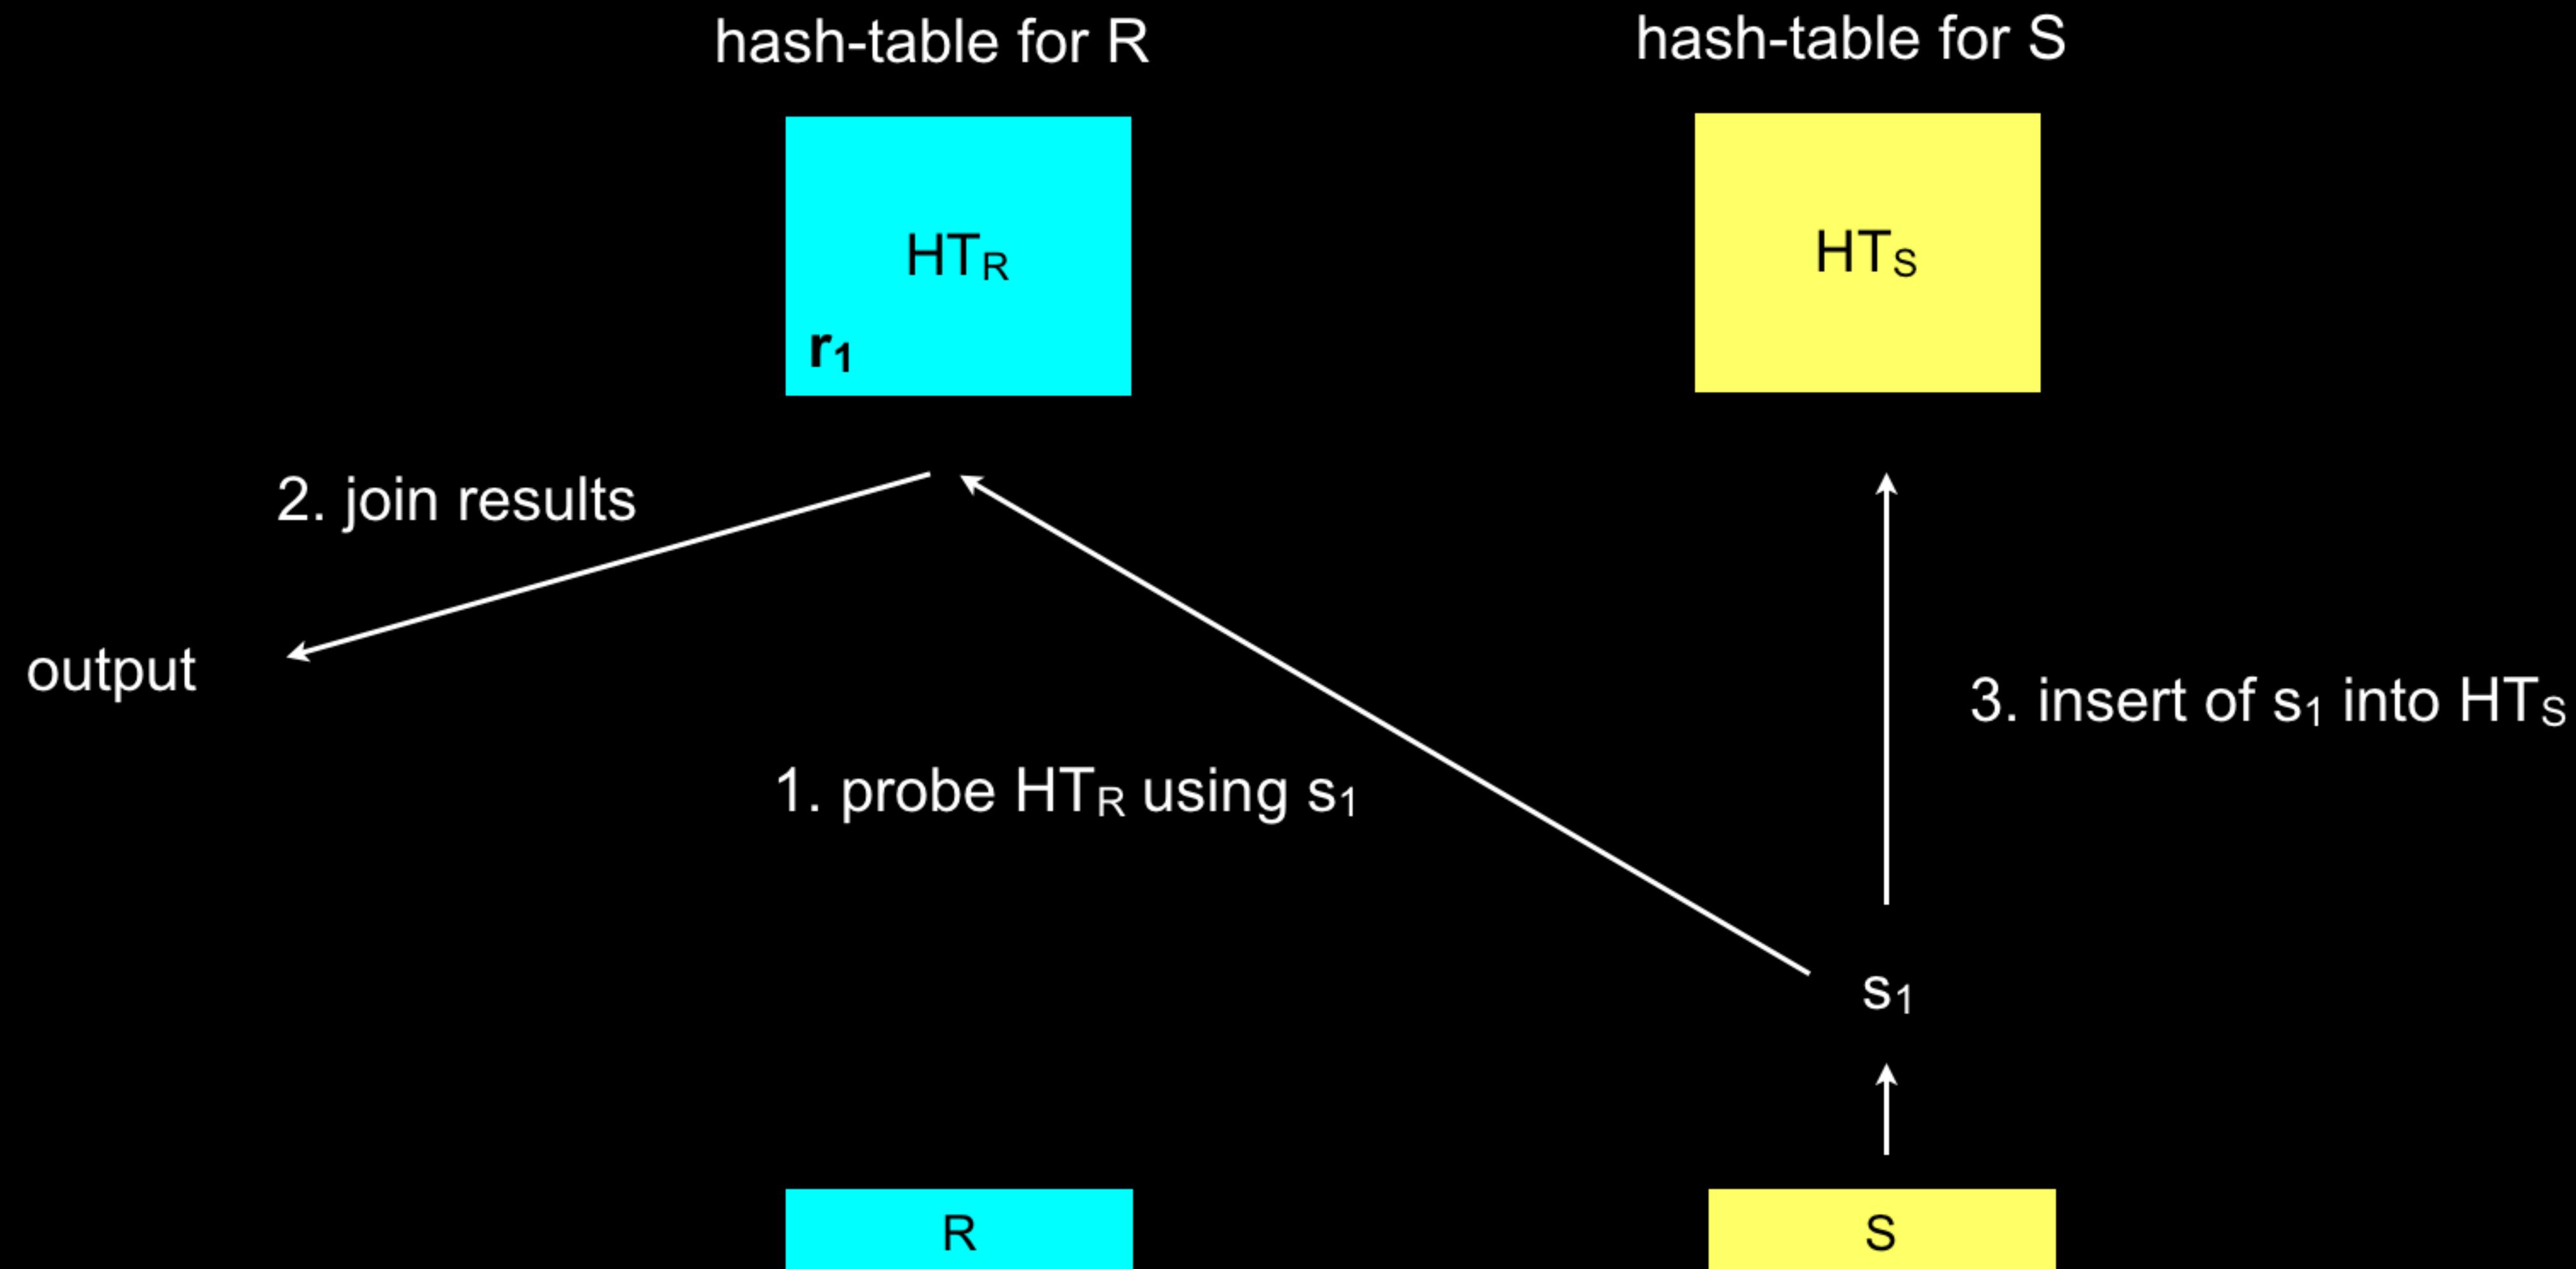
S

streaming environments  
SELECT

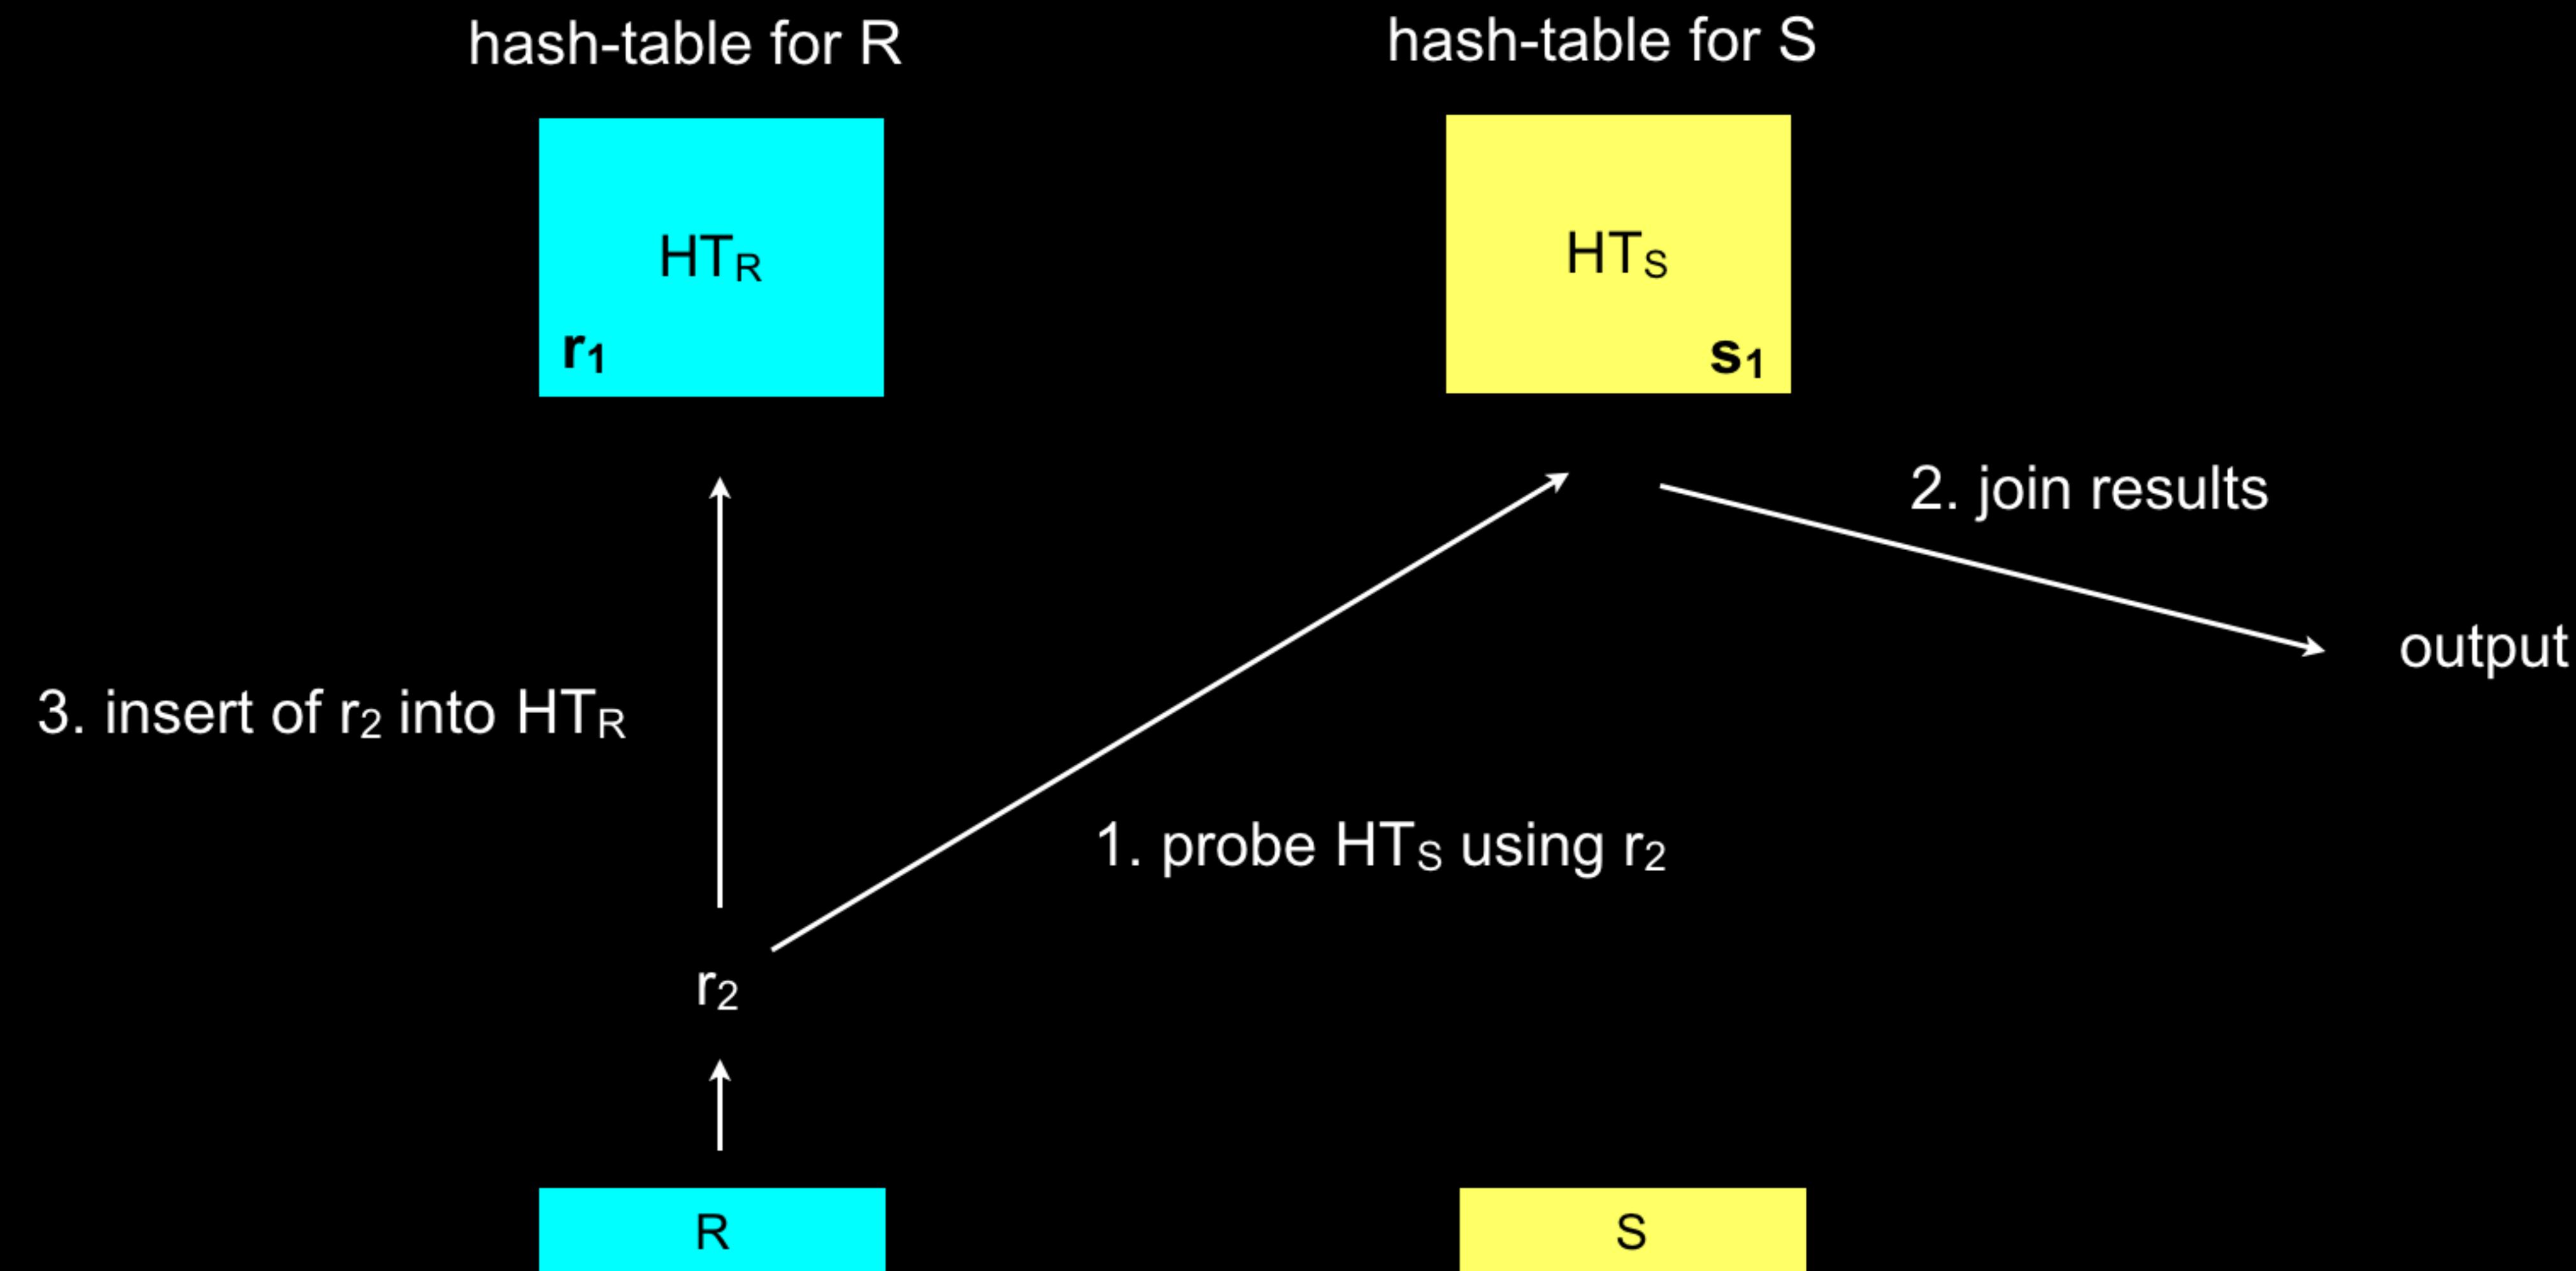
# Example: Double-pipelined Hash Join



# Example: Double-pipelined Hash Join

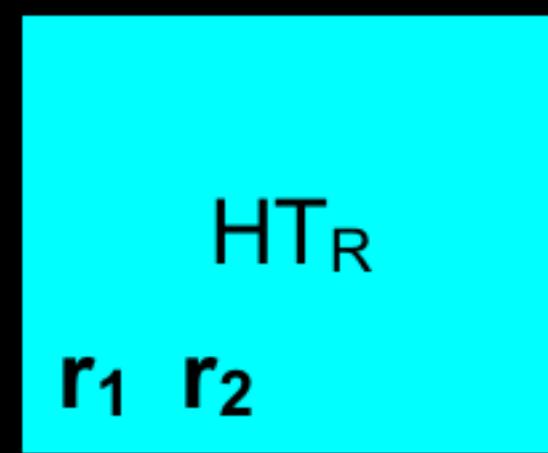


# Example: Double-pipelined Hash Join

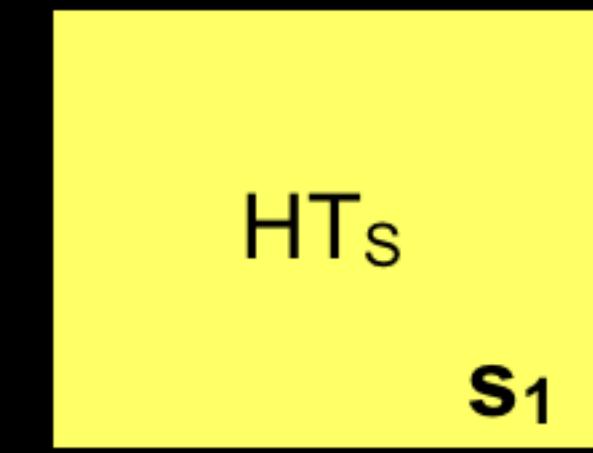


# Example: Double-pipelined Hash Join

hash-table for R



hash-table for S



# Double-pipelined Hash Join

R

S

**JP(r,s) := r.x == s.x**

//definition of the join predicate

**probeAndInsert( tuple, indexToInsert, indexToProbe ):**

$\epsilon R \Rightarrow$  insert into R  
↑  
1 probe S

$\epsilon S \Rightarrow$  insert into S  
1 probe R

# Double-pipelined Hash Join

R

S

**JP(r,s) := r.x == s.x**

//definition of the join predicate

**probeAndInsert( tuple, indexToInsert, indexToProbe ):**

queryResultSet = indexToProbe.query( tuple.x );

//query index for this tuple.x (probe the index)

If NOT queryResultSet.isEmpty():

//did the query return results?

joinResultSet = { tuple } × queryResultSet;

//cross product tuple with the query result

indexToInsert.insert( tuple.x );

//insert the tuple into the corresponding index

return joinResultSet;

//return the actual join result for this tuple

**DoublePipelinedHashJoin( R, S, JP(r,s) ):**

indexOnRX = new HashTable(); indexOnSX = new HashTable();

//initialize empty hash tables

bool readFromR = TRUE;

//flag to determine relation to read from

While R.hasNext() AND S.hasNext():

//if both relations have items to read

If readFromR:

//if we have to read from R in this round

    output( probeAndInsert( R.next(), indexOnRX, indexOnSX ) );

//join the R tuple with the probed tuples of S

Else:

//if we have to read from S in this round

    output( probeAndInsert( S.next(), indexOnSX, indexOnRX ) );

//join the S tuple with the probed tuples of R

    readFromR = NOT readFromR;

//switch to other relation for next round

While R.hasNext():

//if only R has entries left

    output( probeAndInsert( R.next(), indexOnRX, indexOnSX ) );

//probe them against S and output

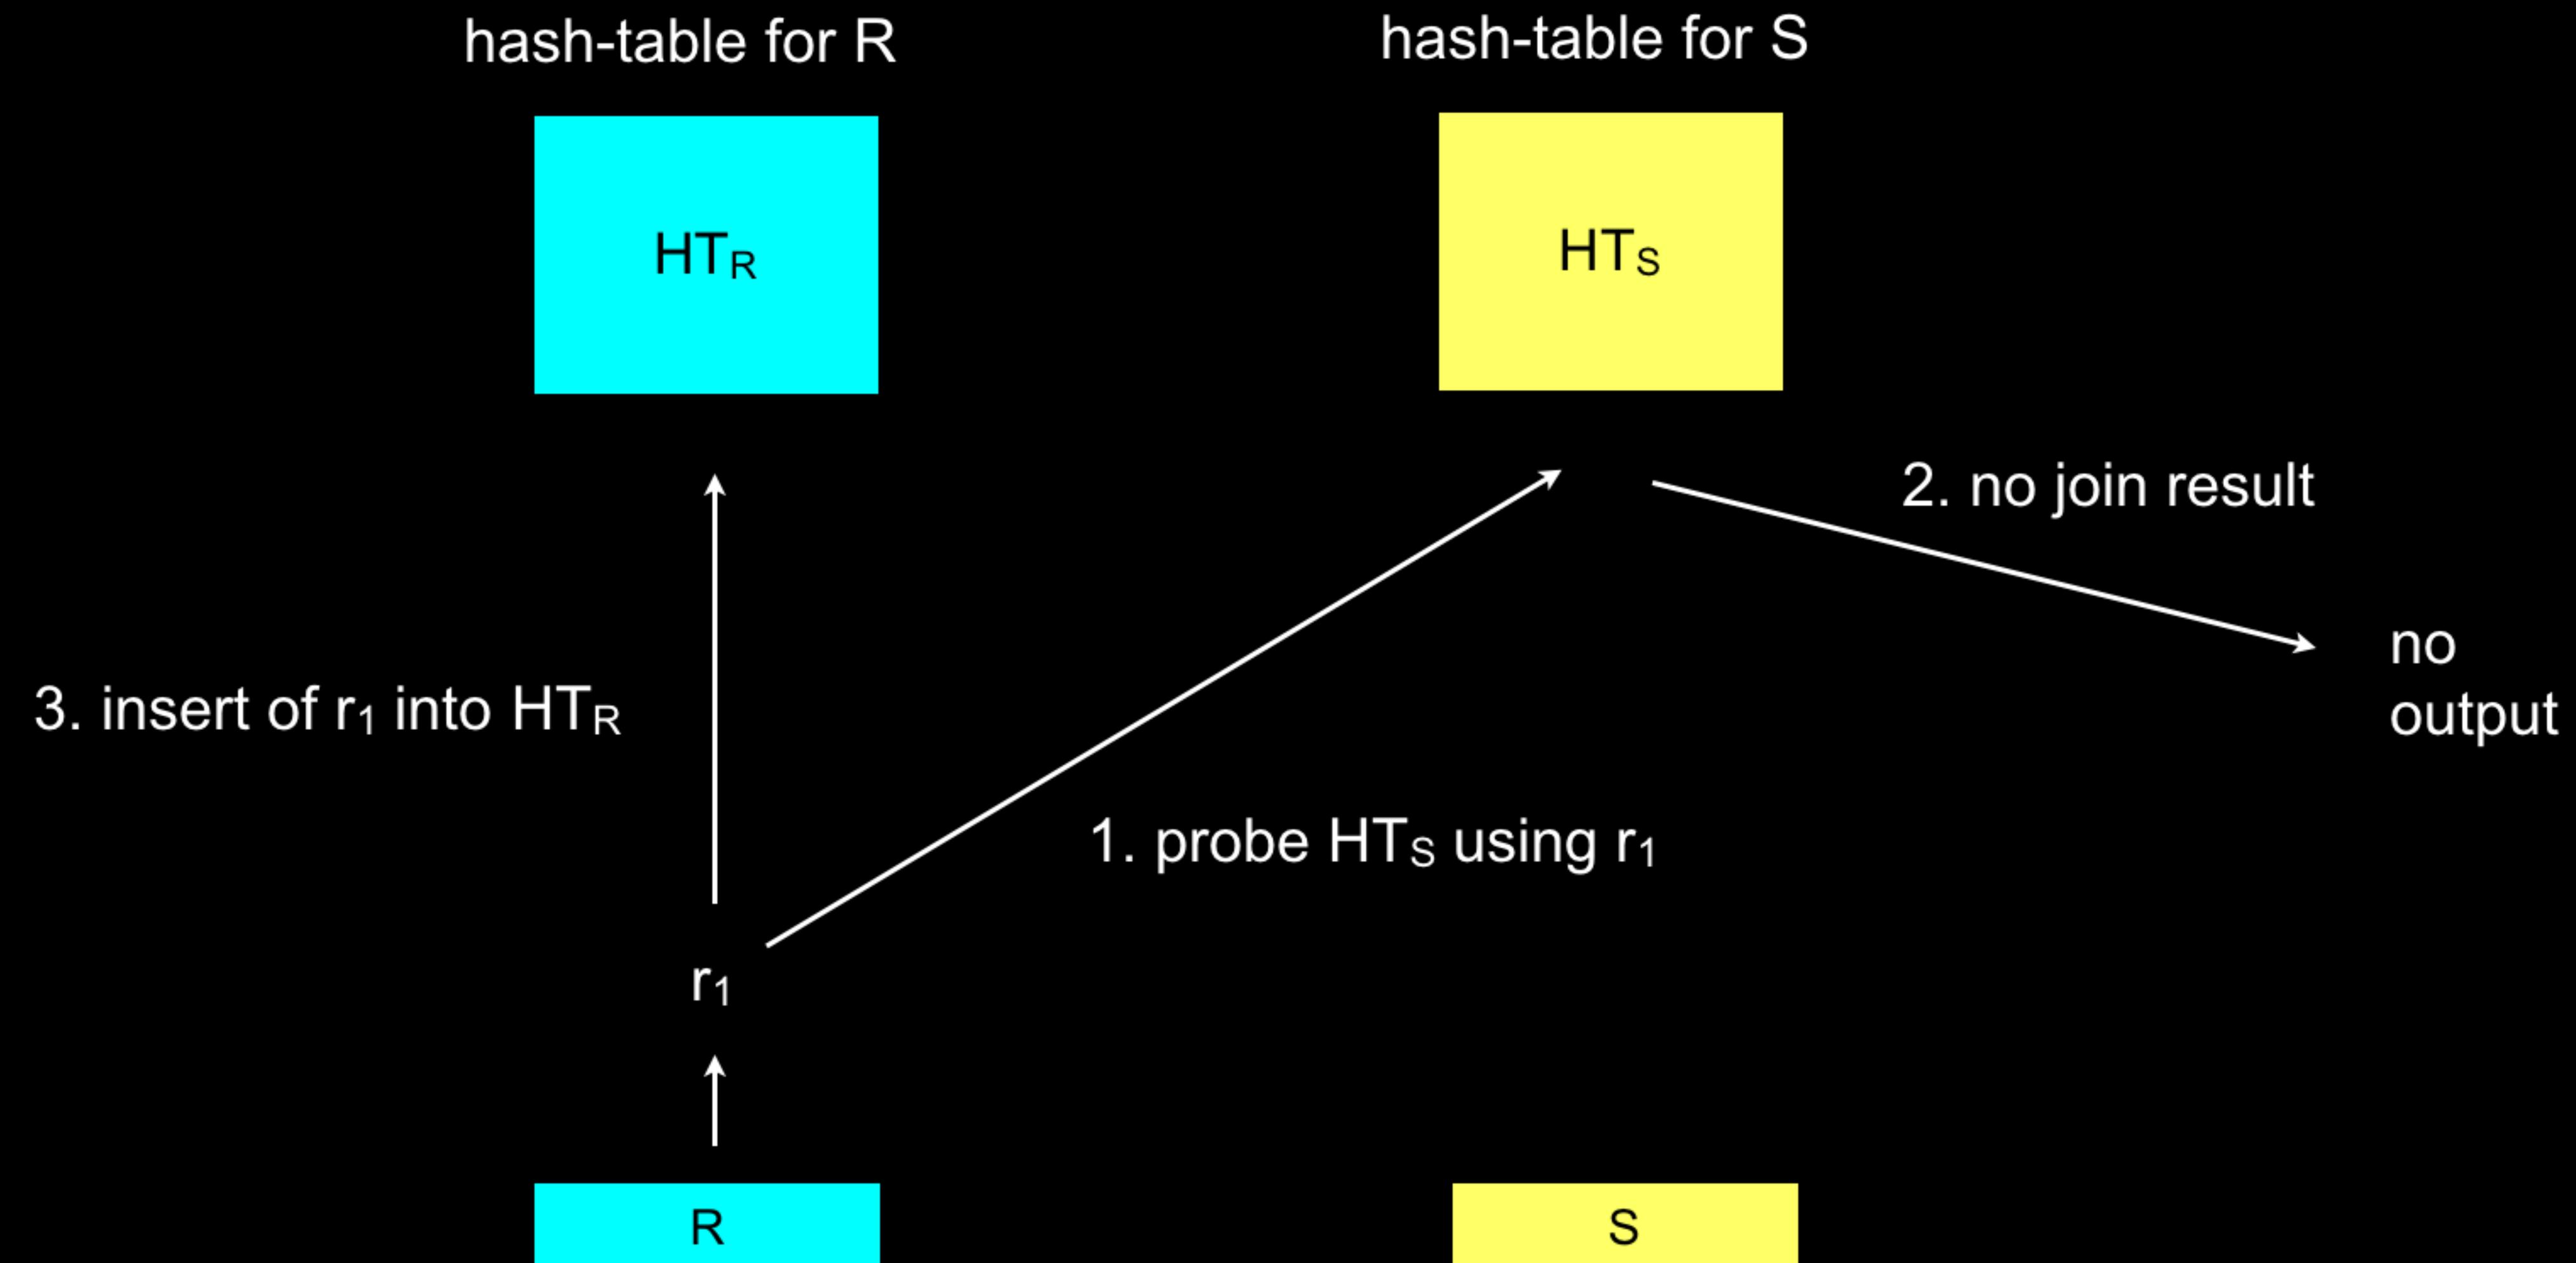
While S.hasNext():

//if only S has entries left

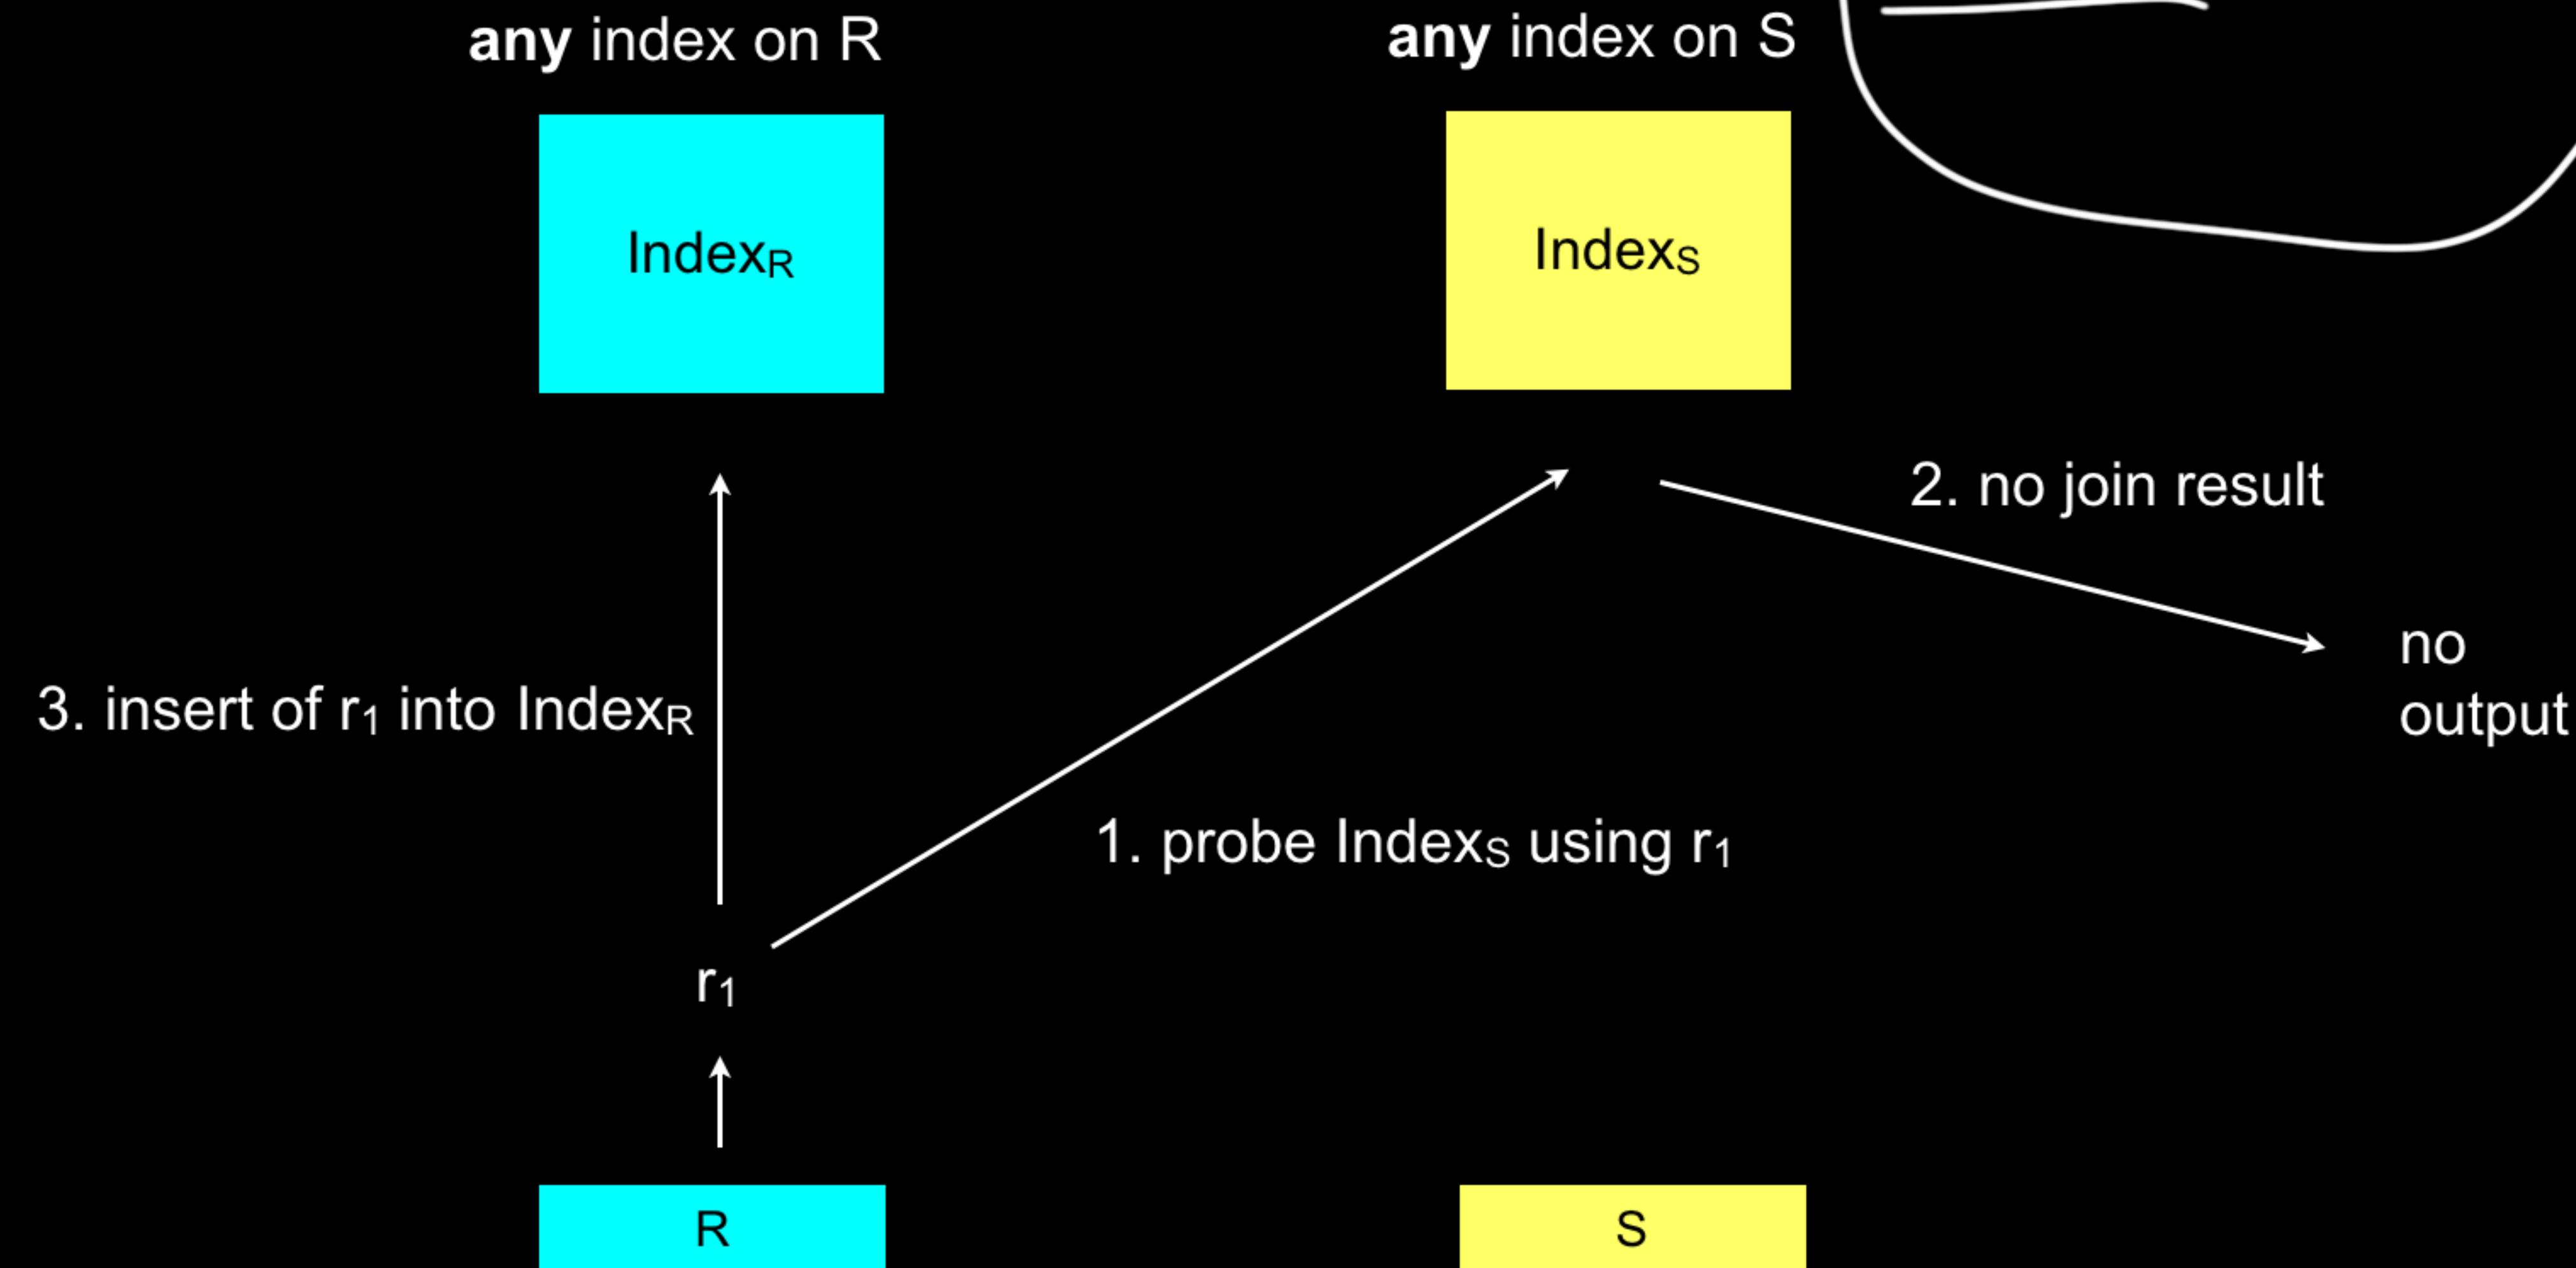
    output( probeAndInsert( S.next(), indexOnSX, indexOnRX ) );

//probe them against R and output

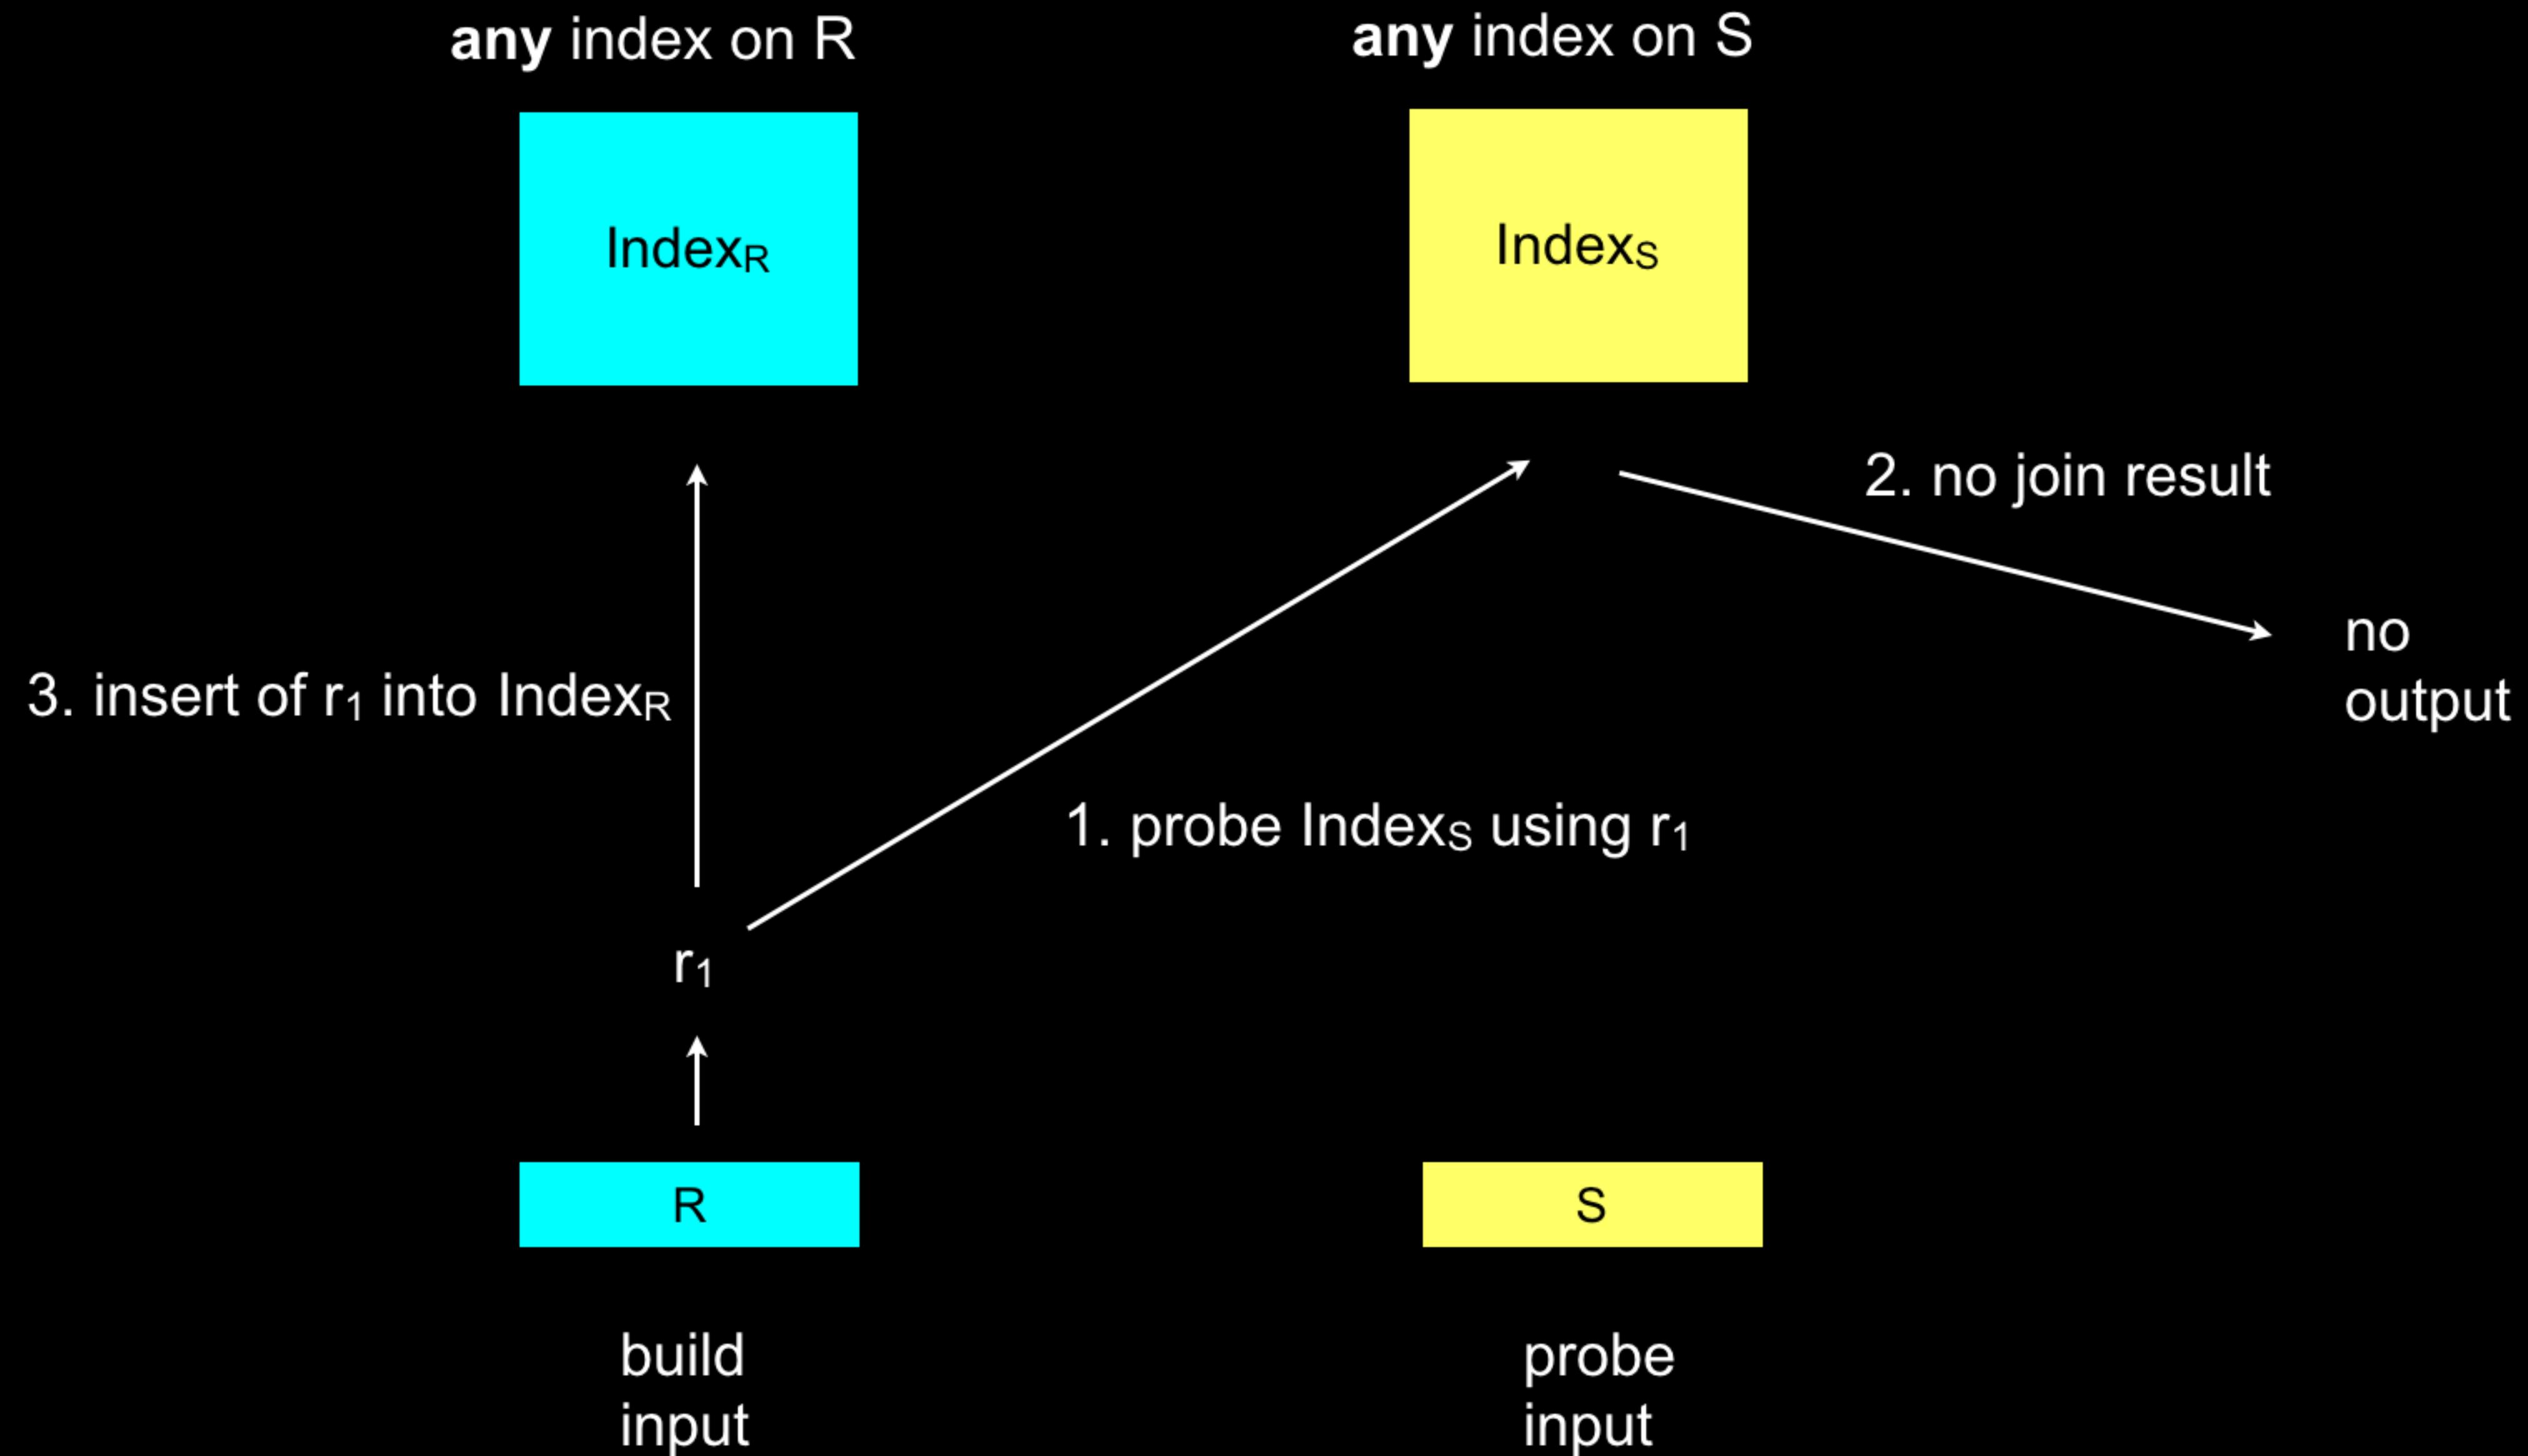
# Example: Double-pipelined Hash Join



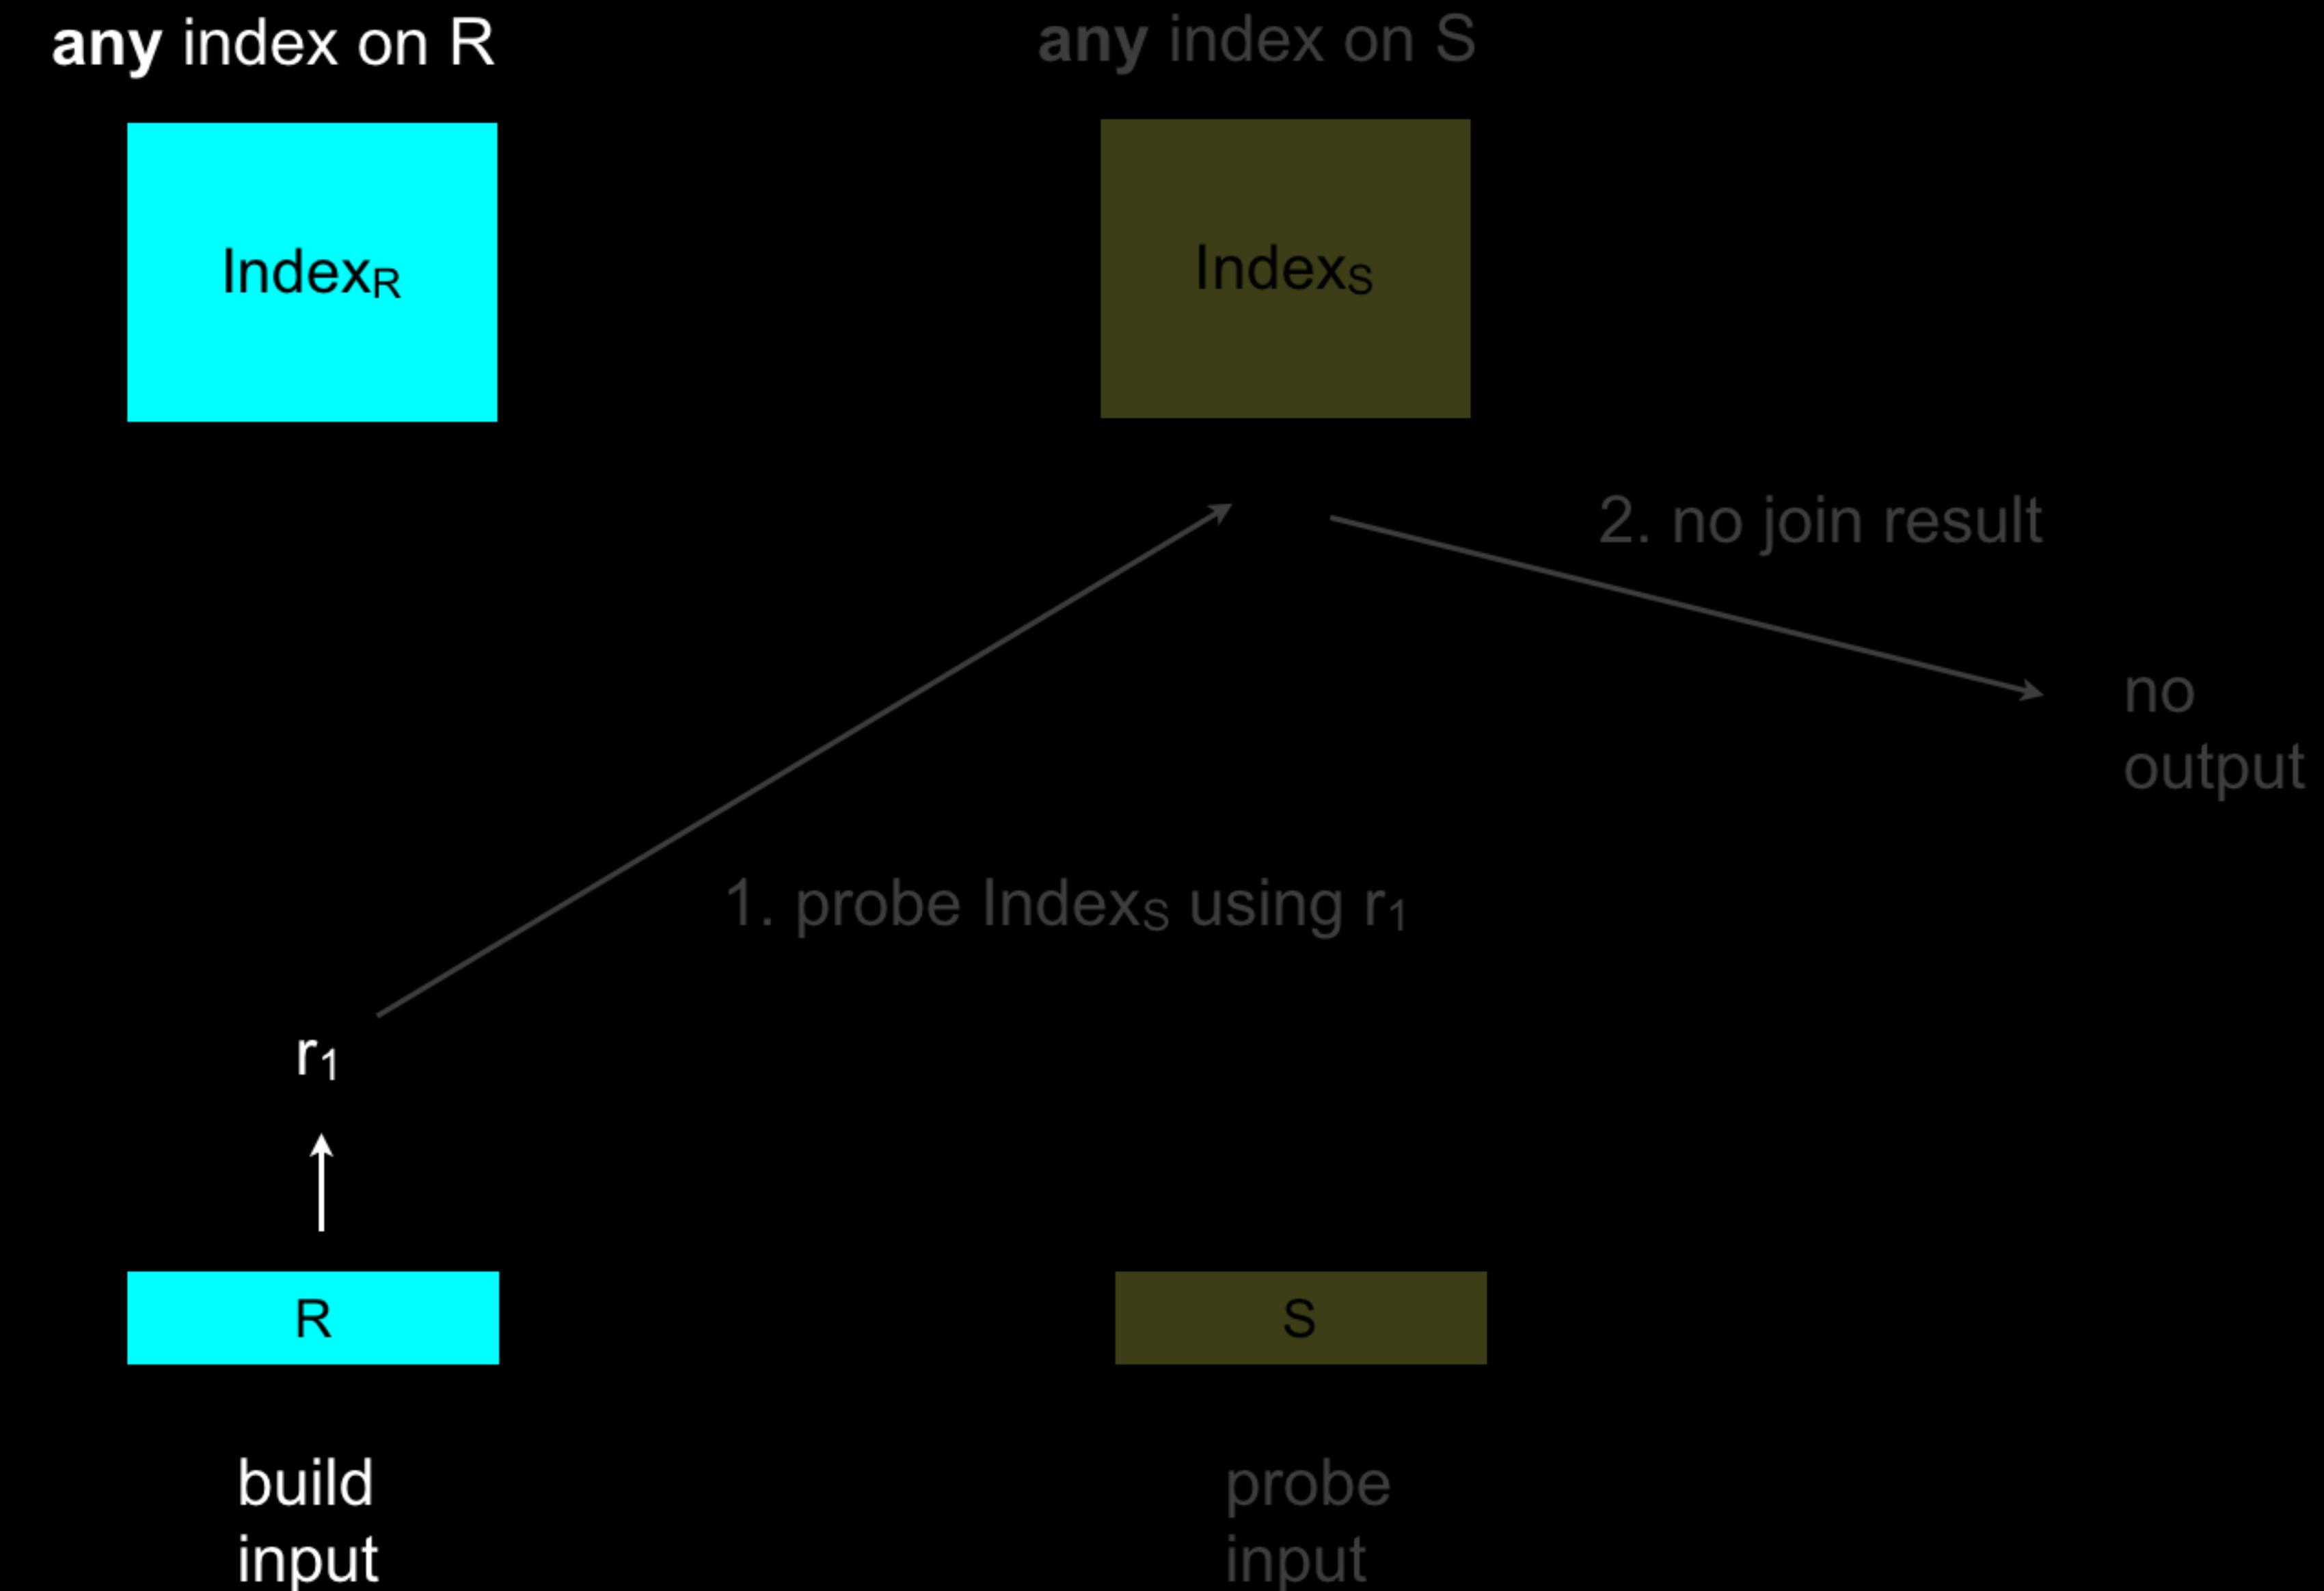
# Example: Double-pipelined Index Join



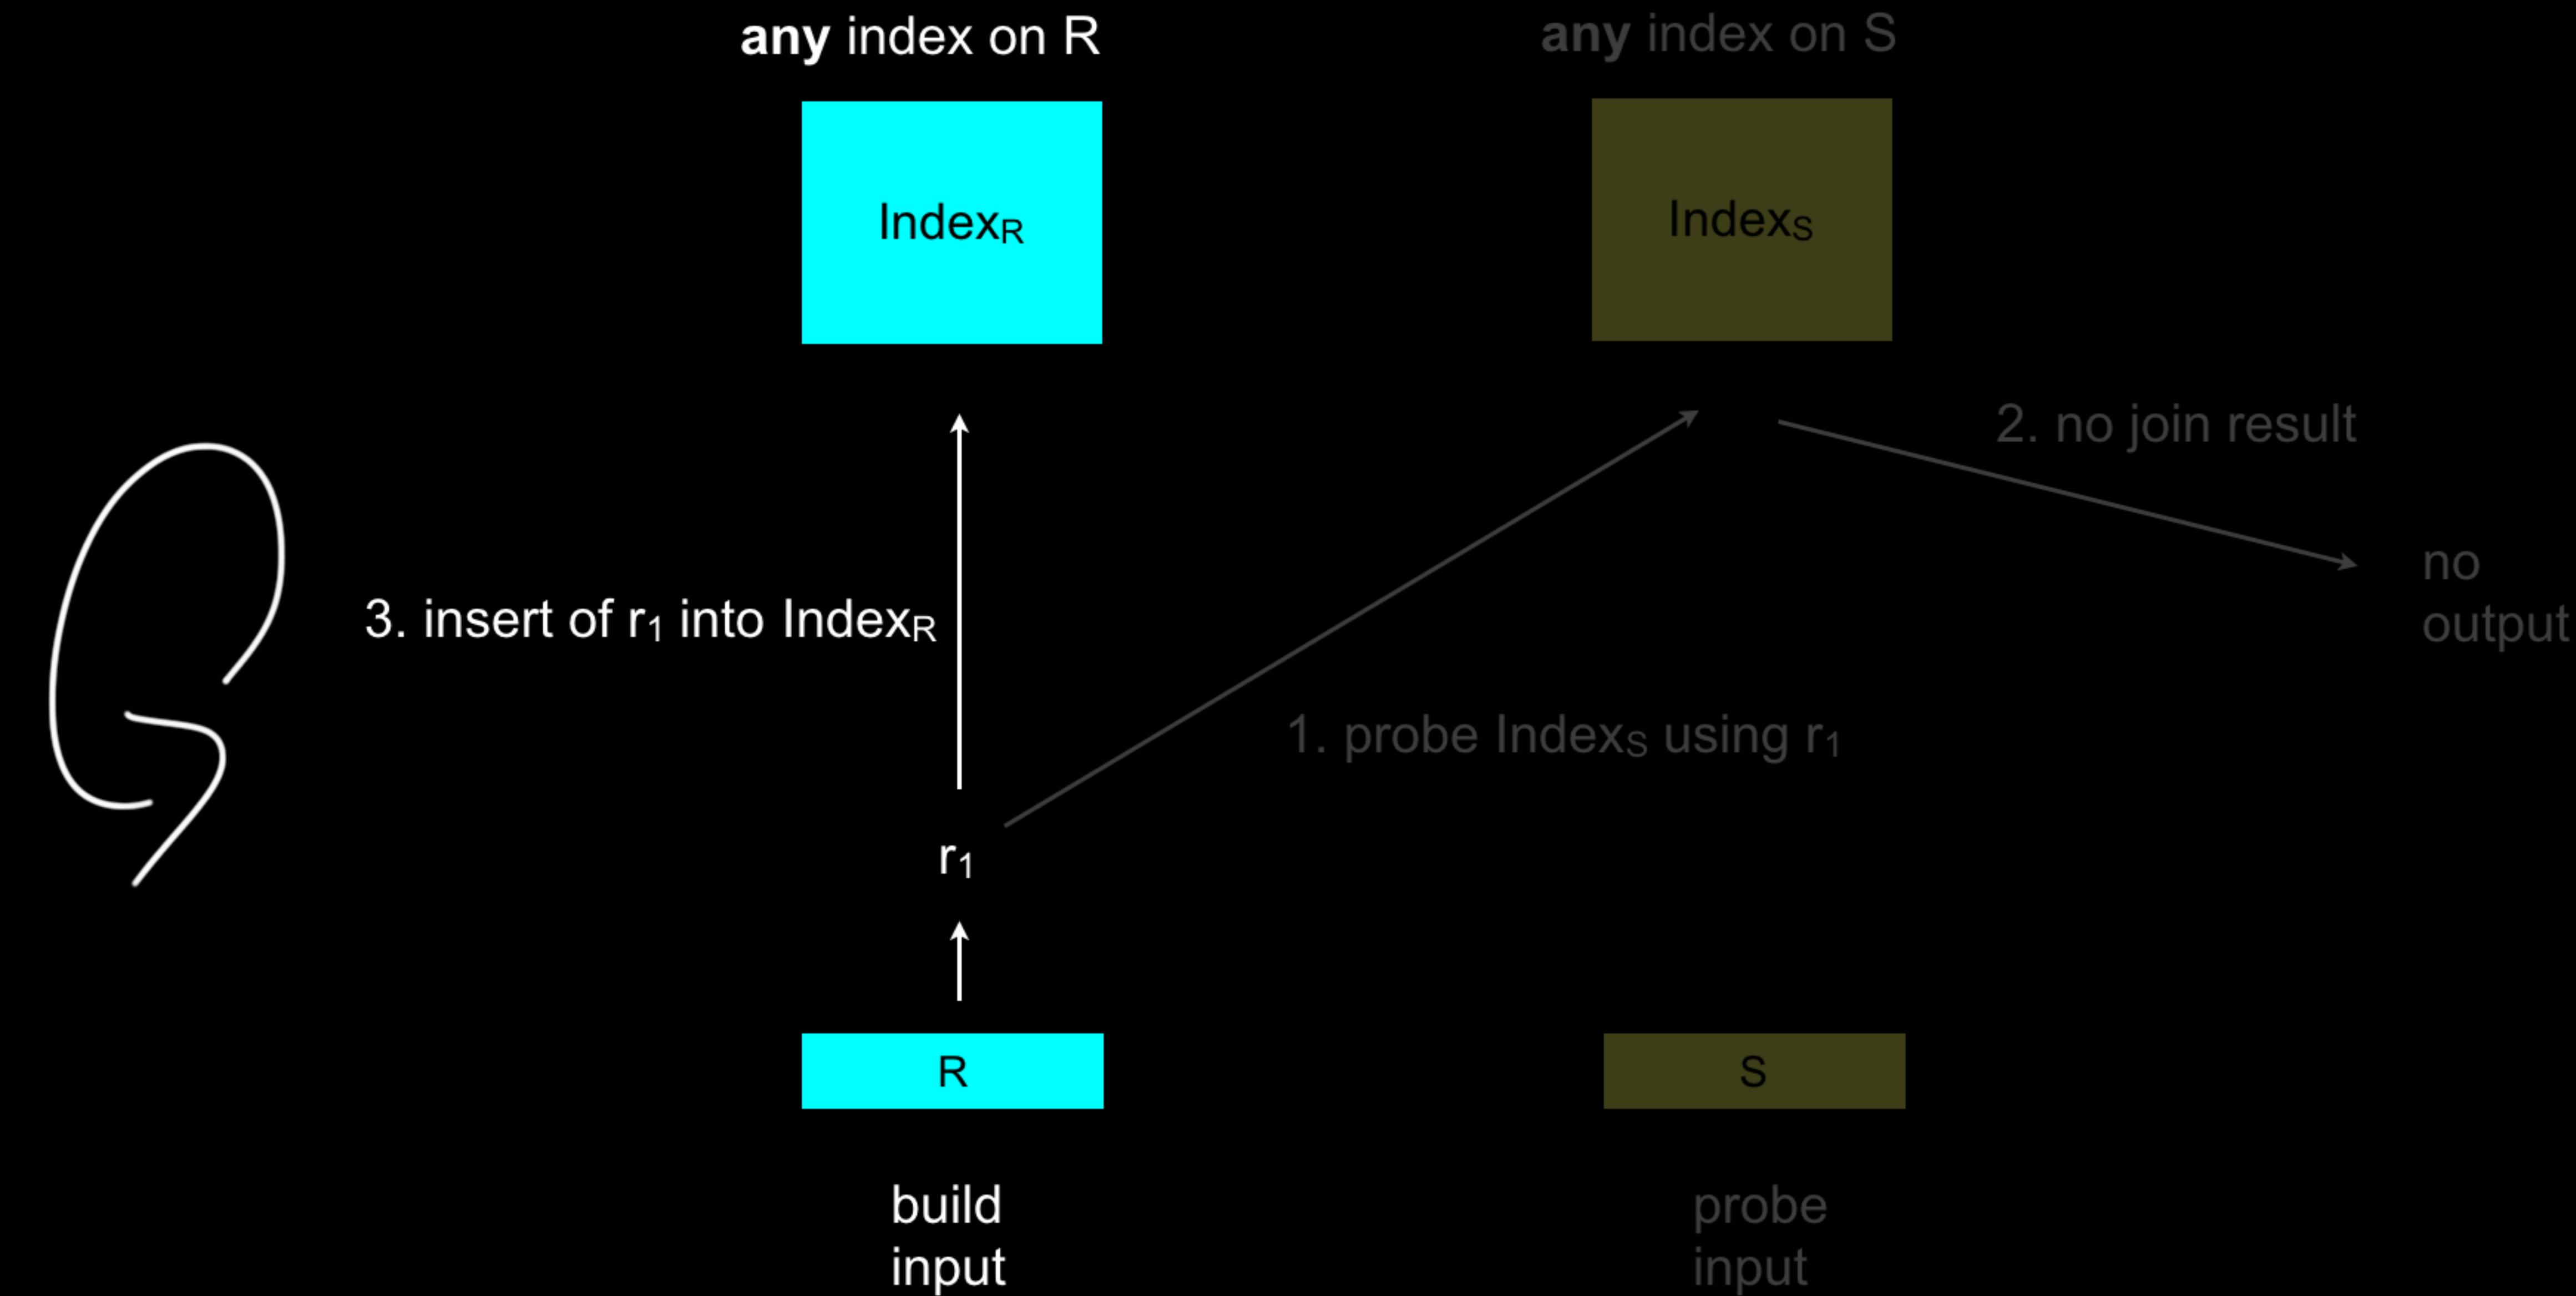
# Example: Index Nested-Loop Join (Viewed Differently)



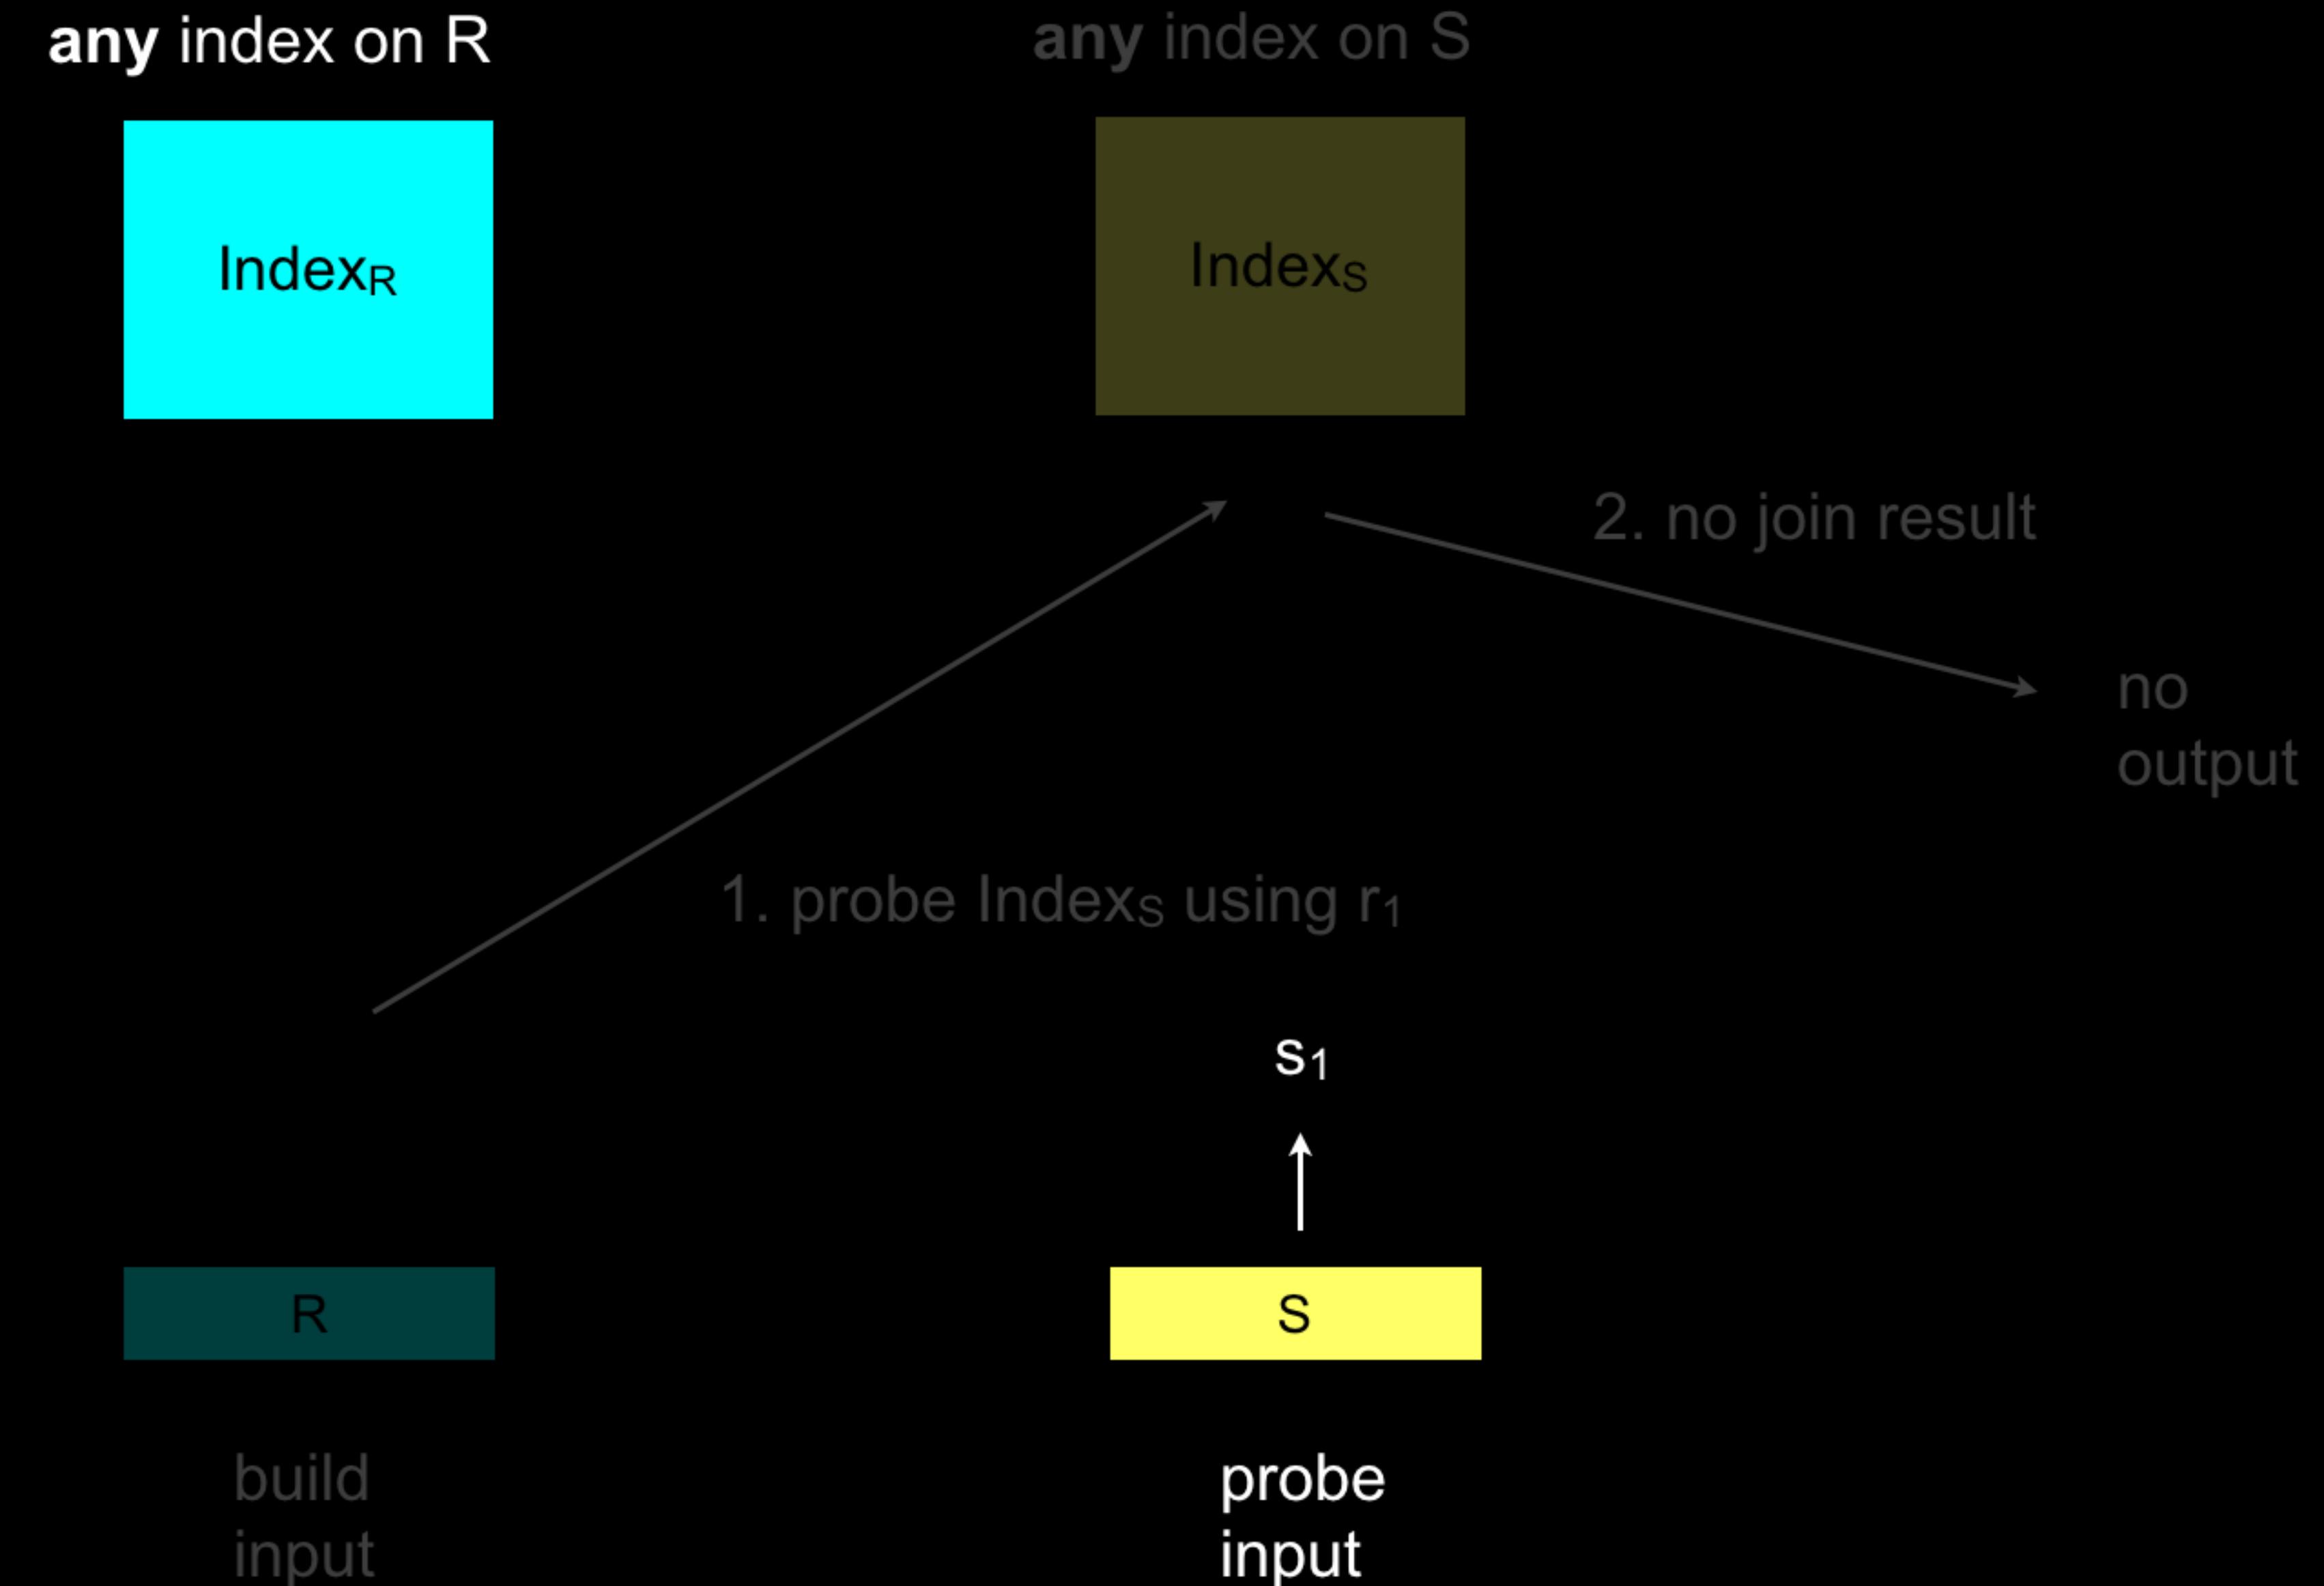
# Build Phase of Index Nested-Loop Join



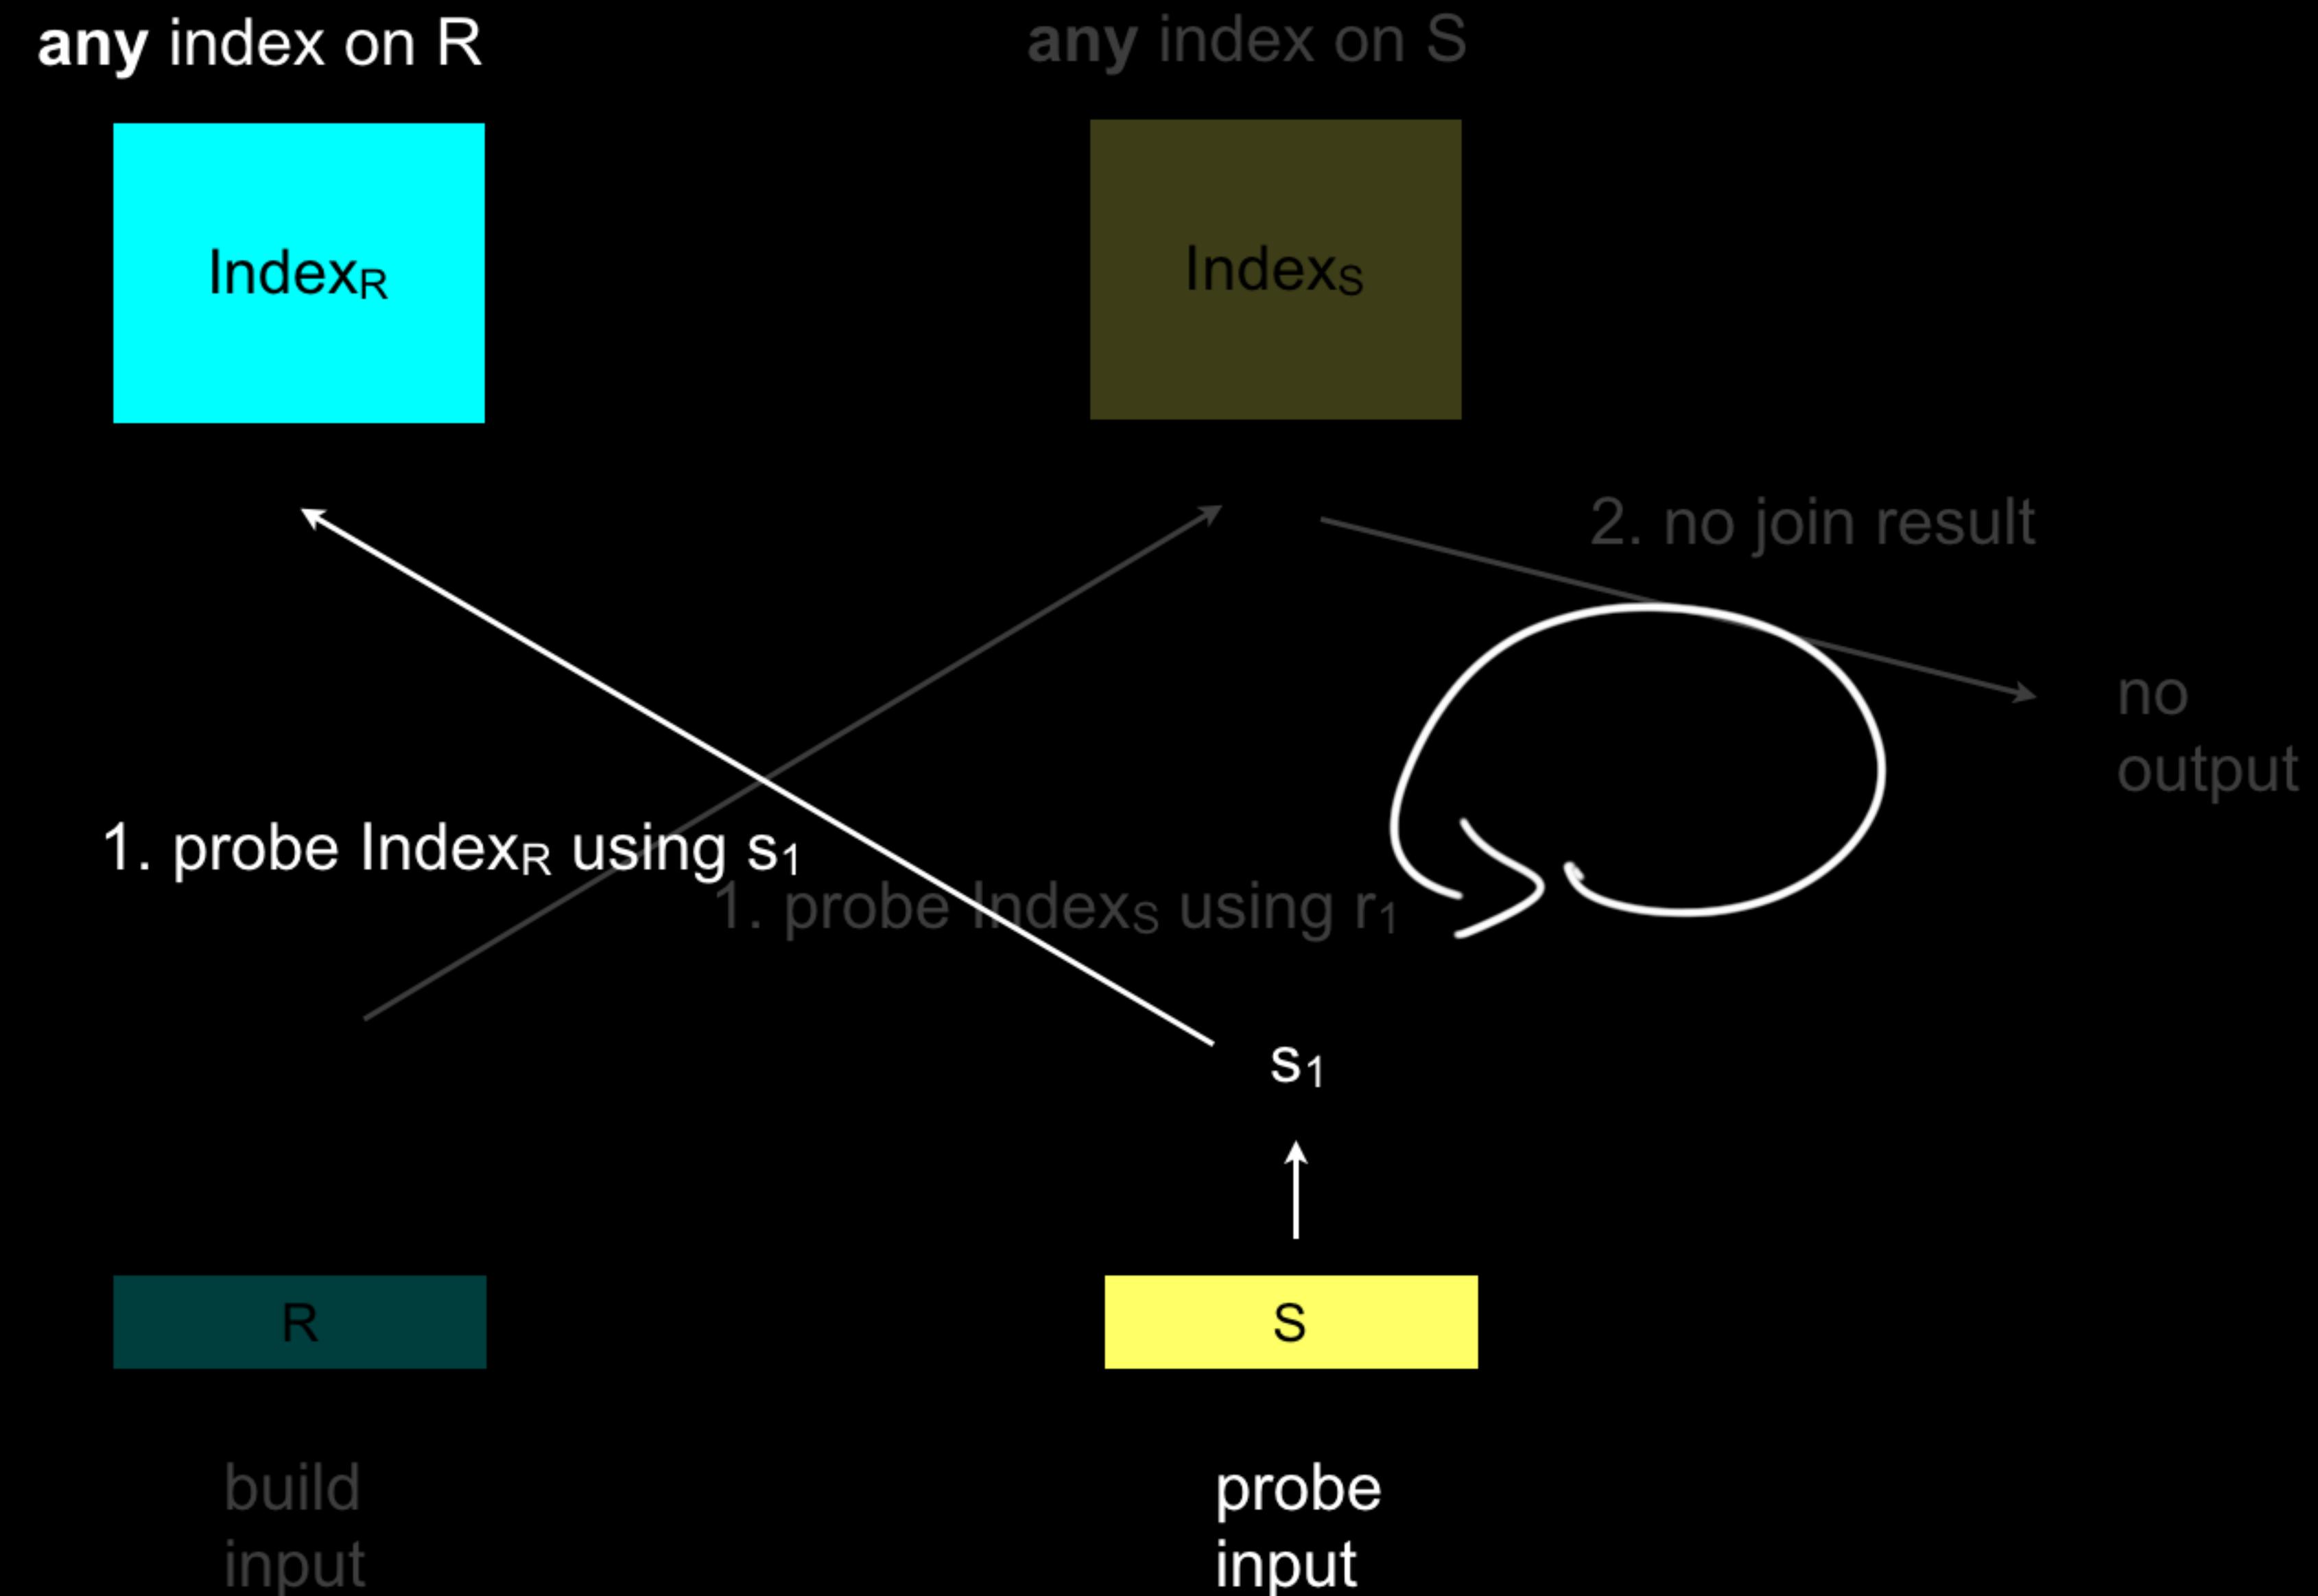
# Build Phase of Index Nested-Loop Join



# Probe Phase of Index Nested-Loop Join



# Probe Phase of Index Nested-Loop Join



# Difference of DPHJ over INLJ: Interleaving of Build and Probe Phases

