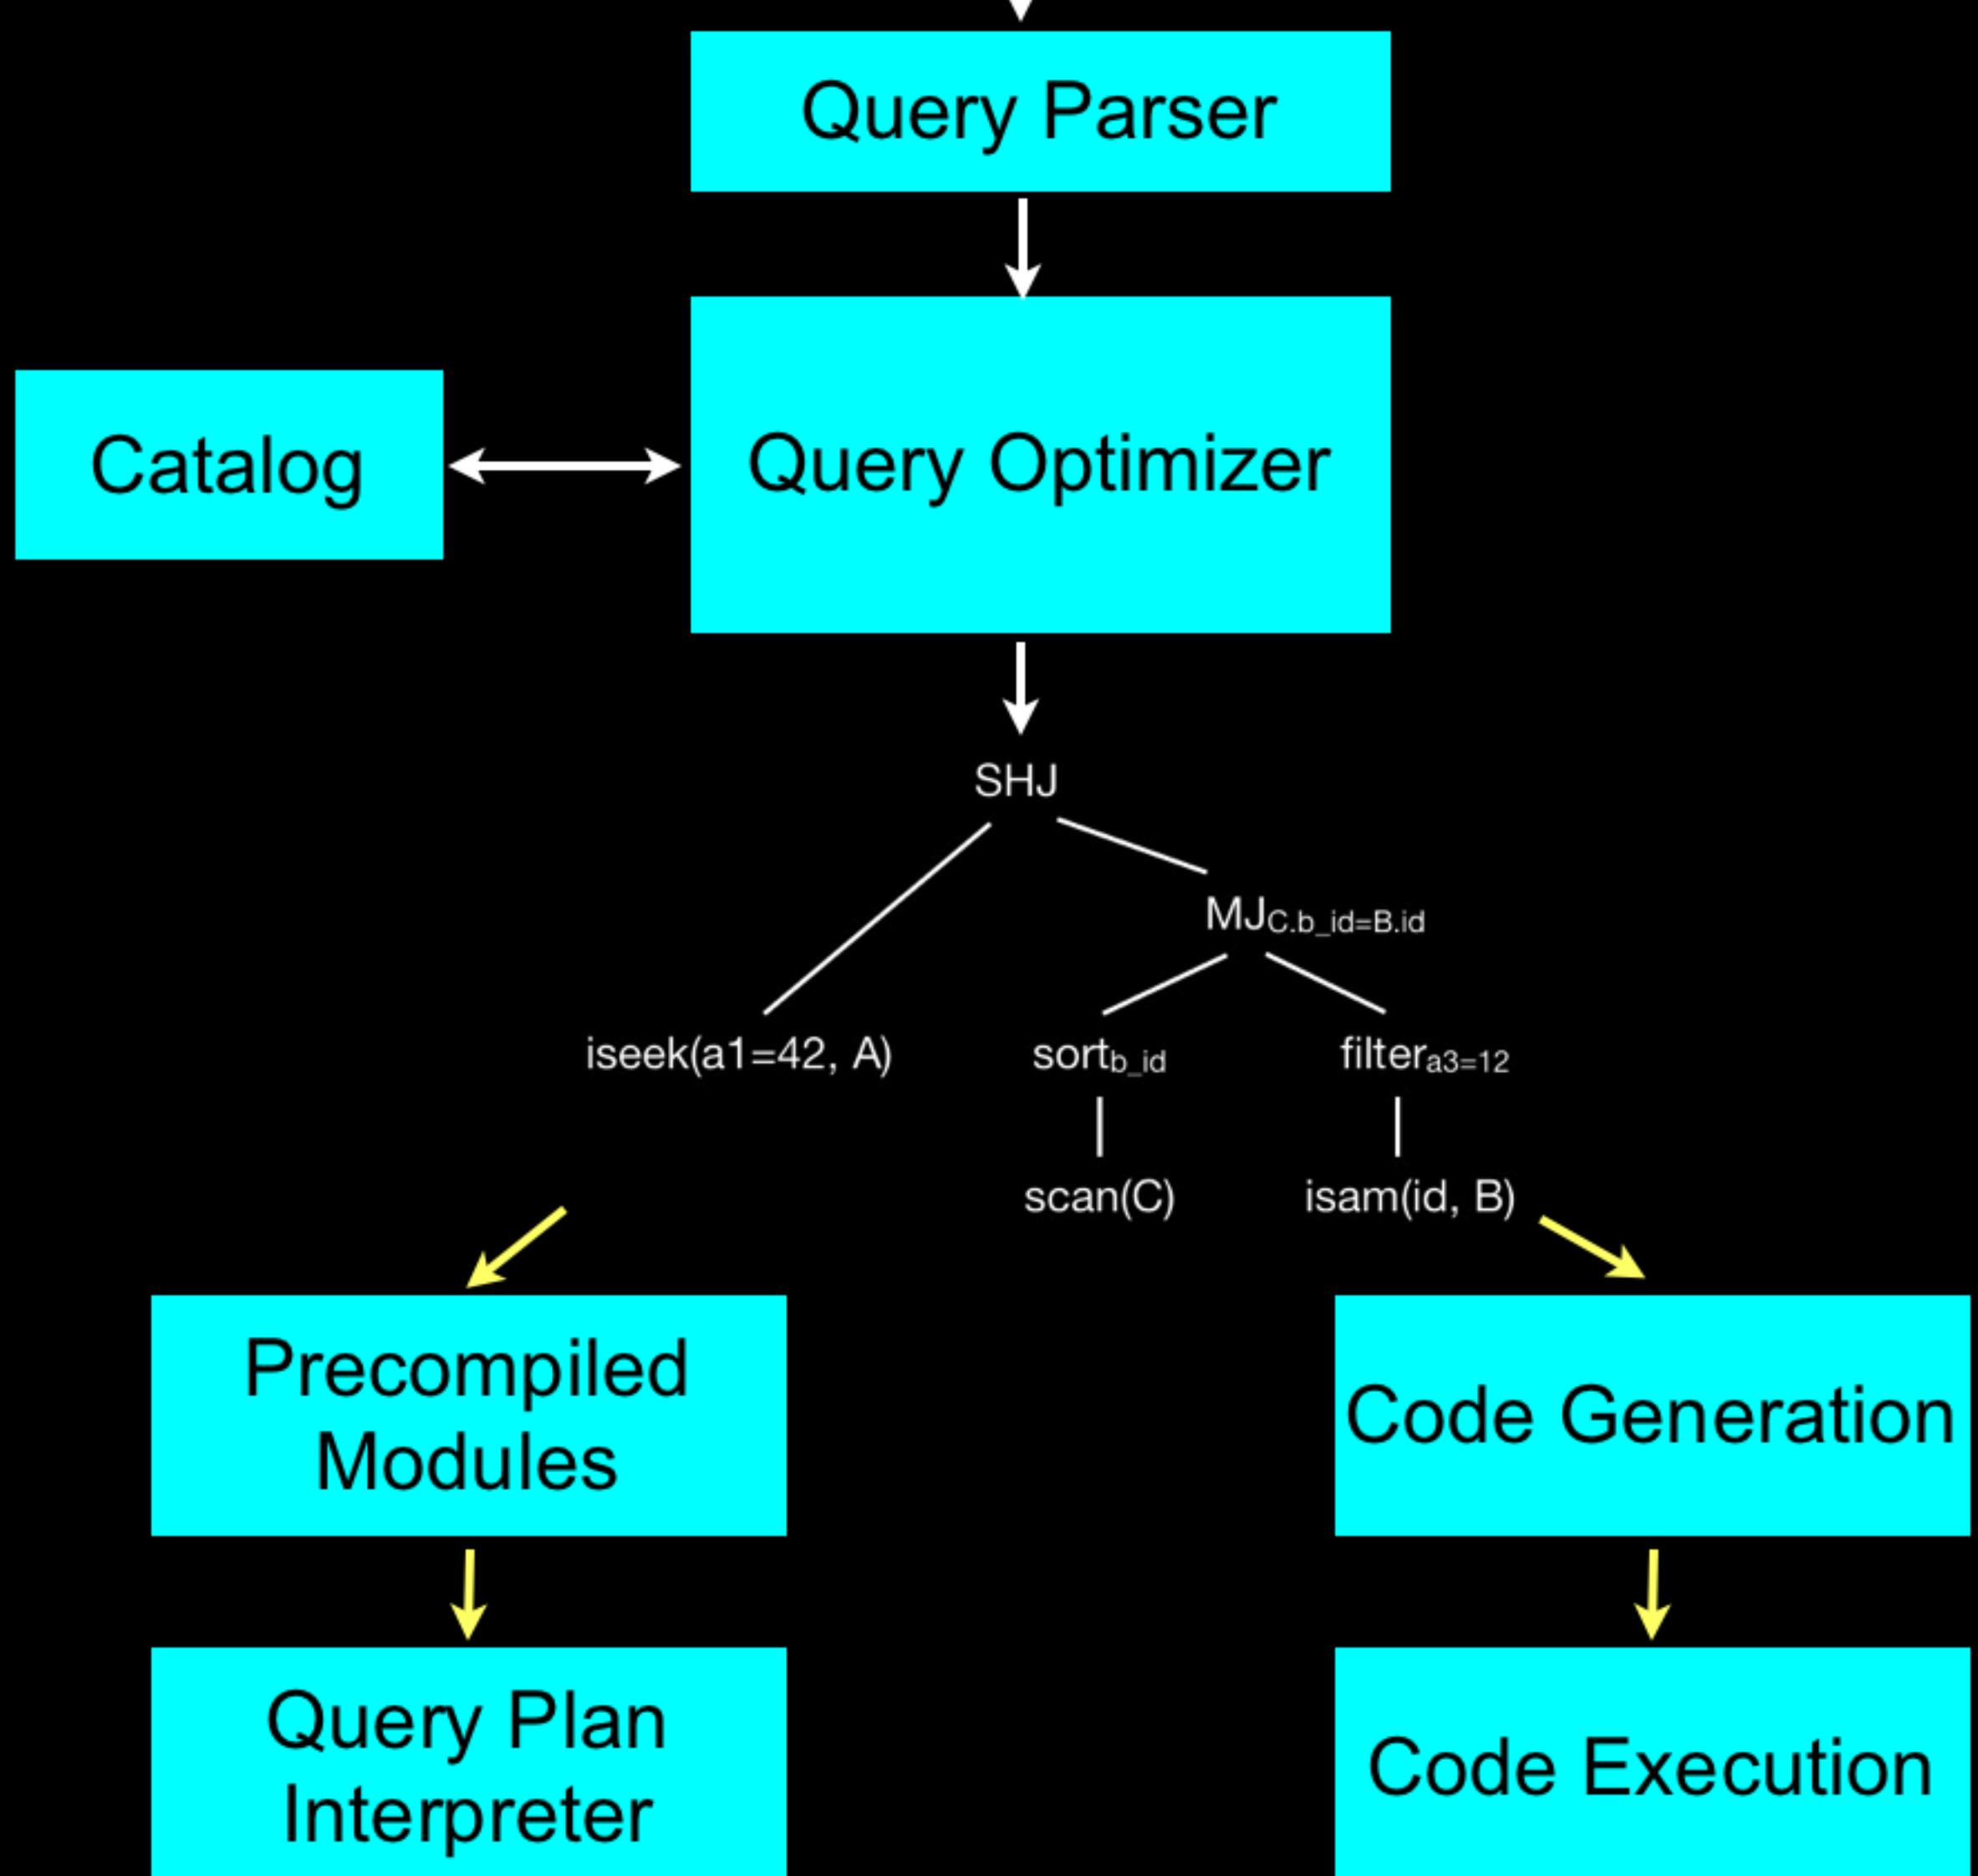


Overview

```

SELECT A.a1, B.a2, B.a3, D.a4
FROM   A     JOIN B ON A.id=B.a_id
        JOIN C ON B.id=C.b_id
WHERE  A.a1=42
       AND B.a3=12
  
```

declarative query



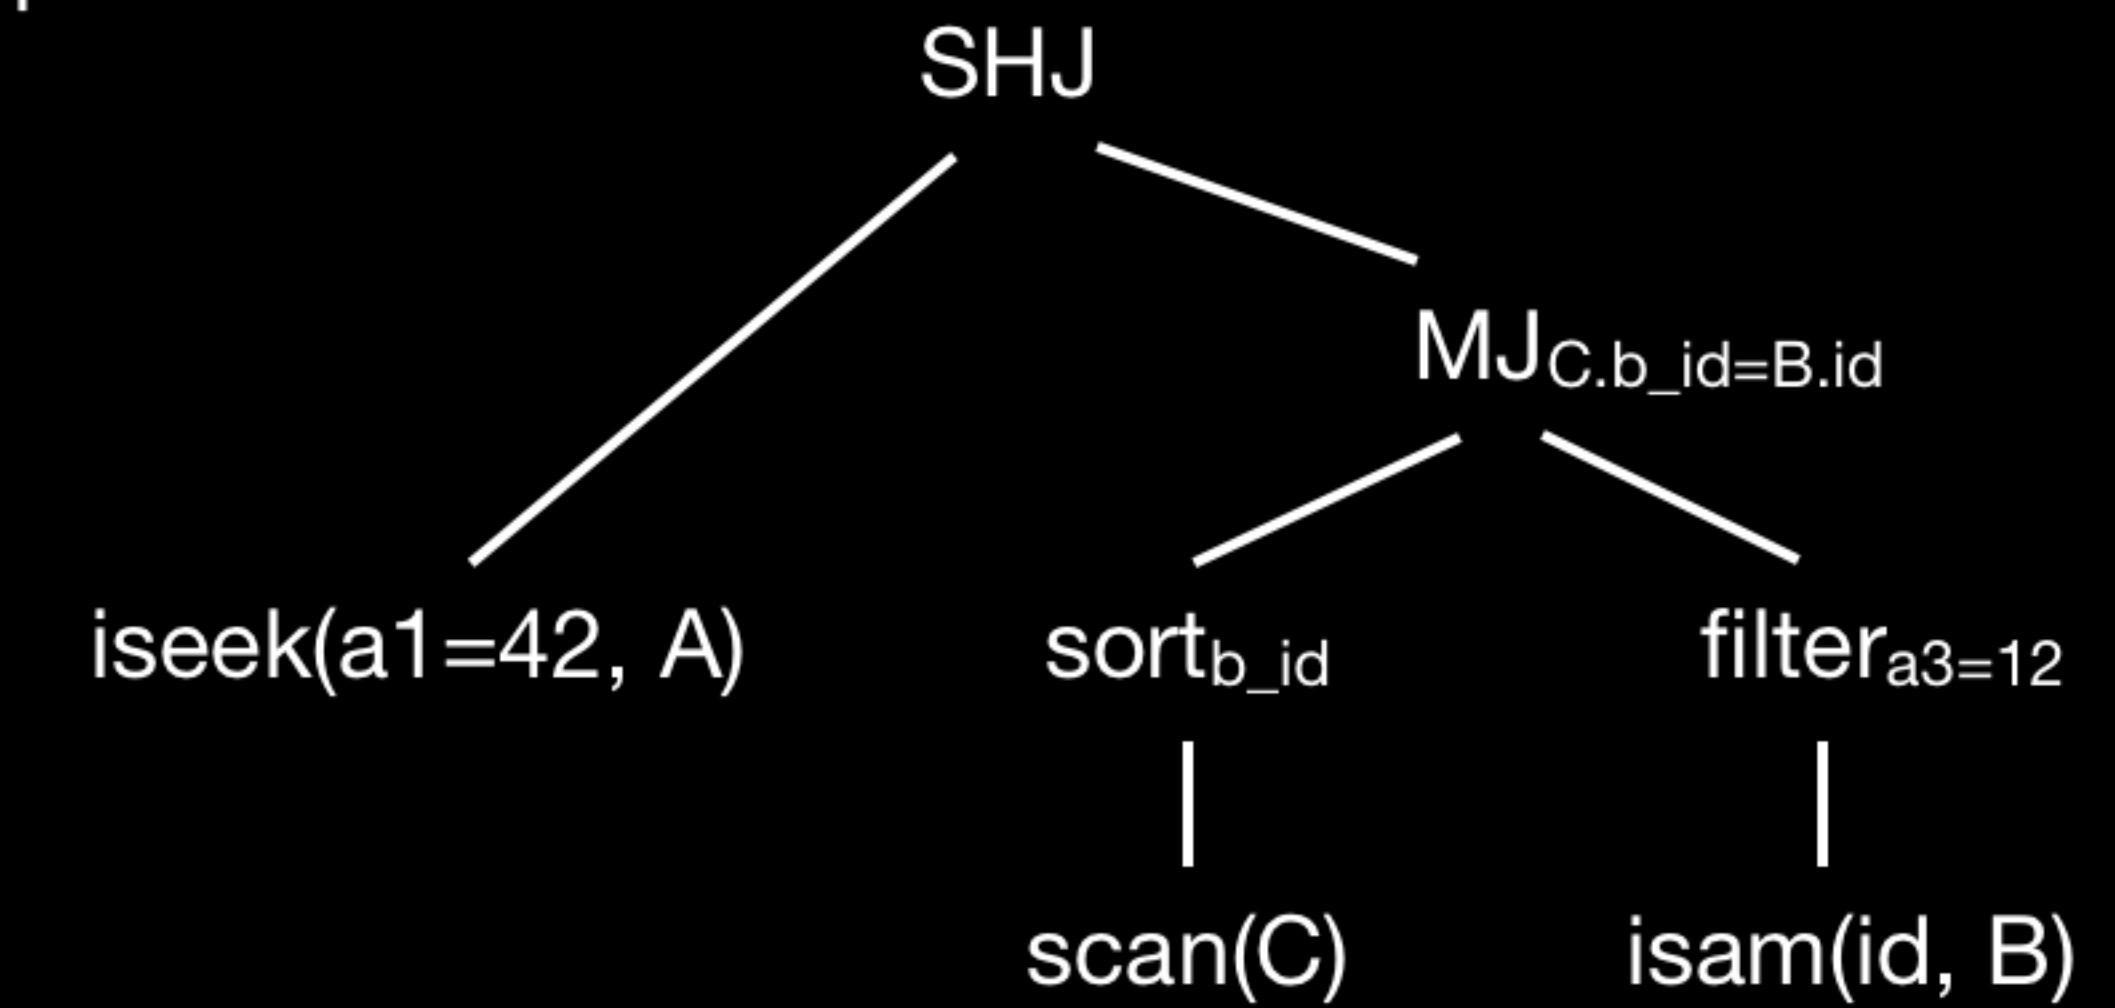
physical query
execution plan
(QEP)

how to execute?

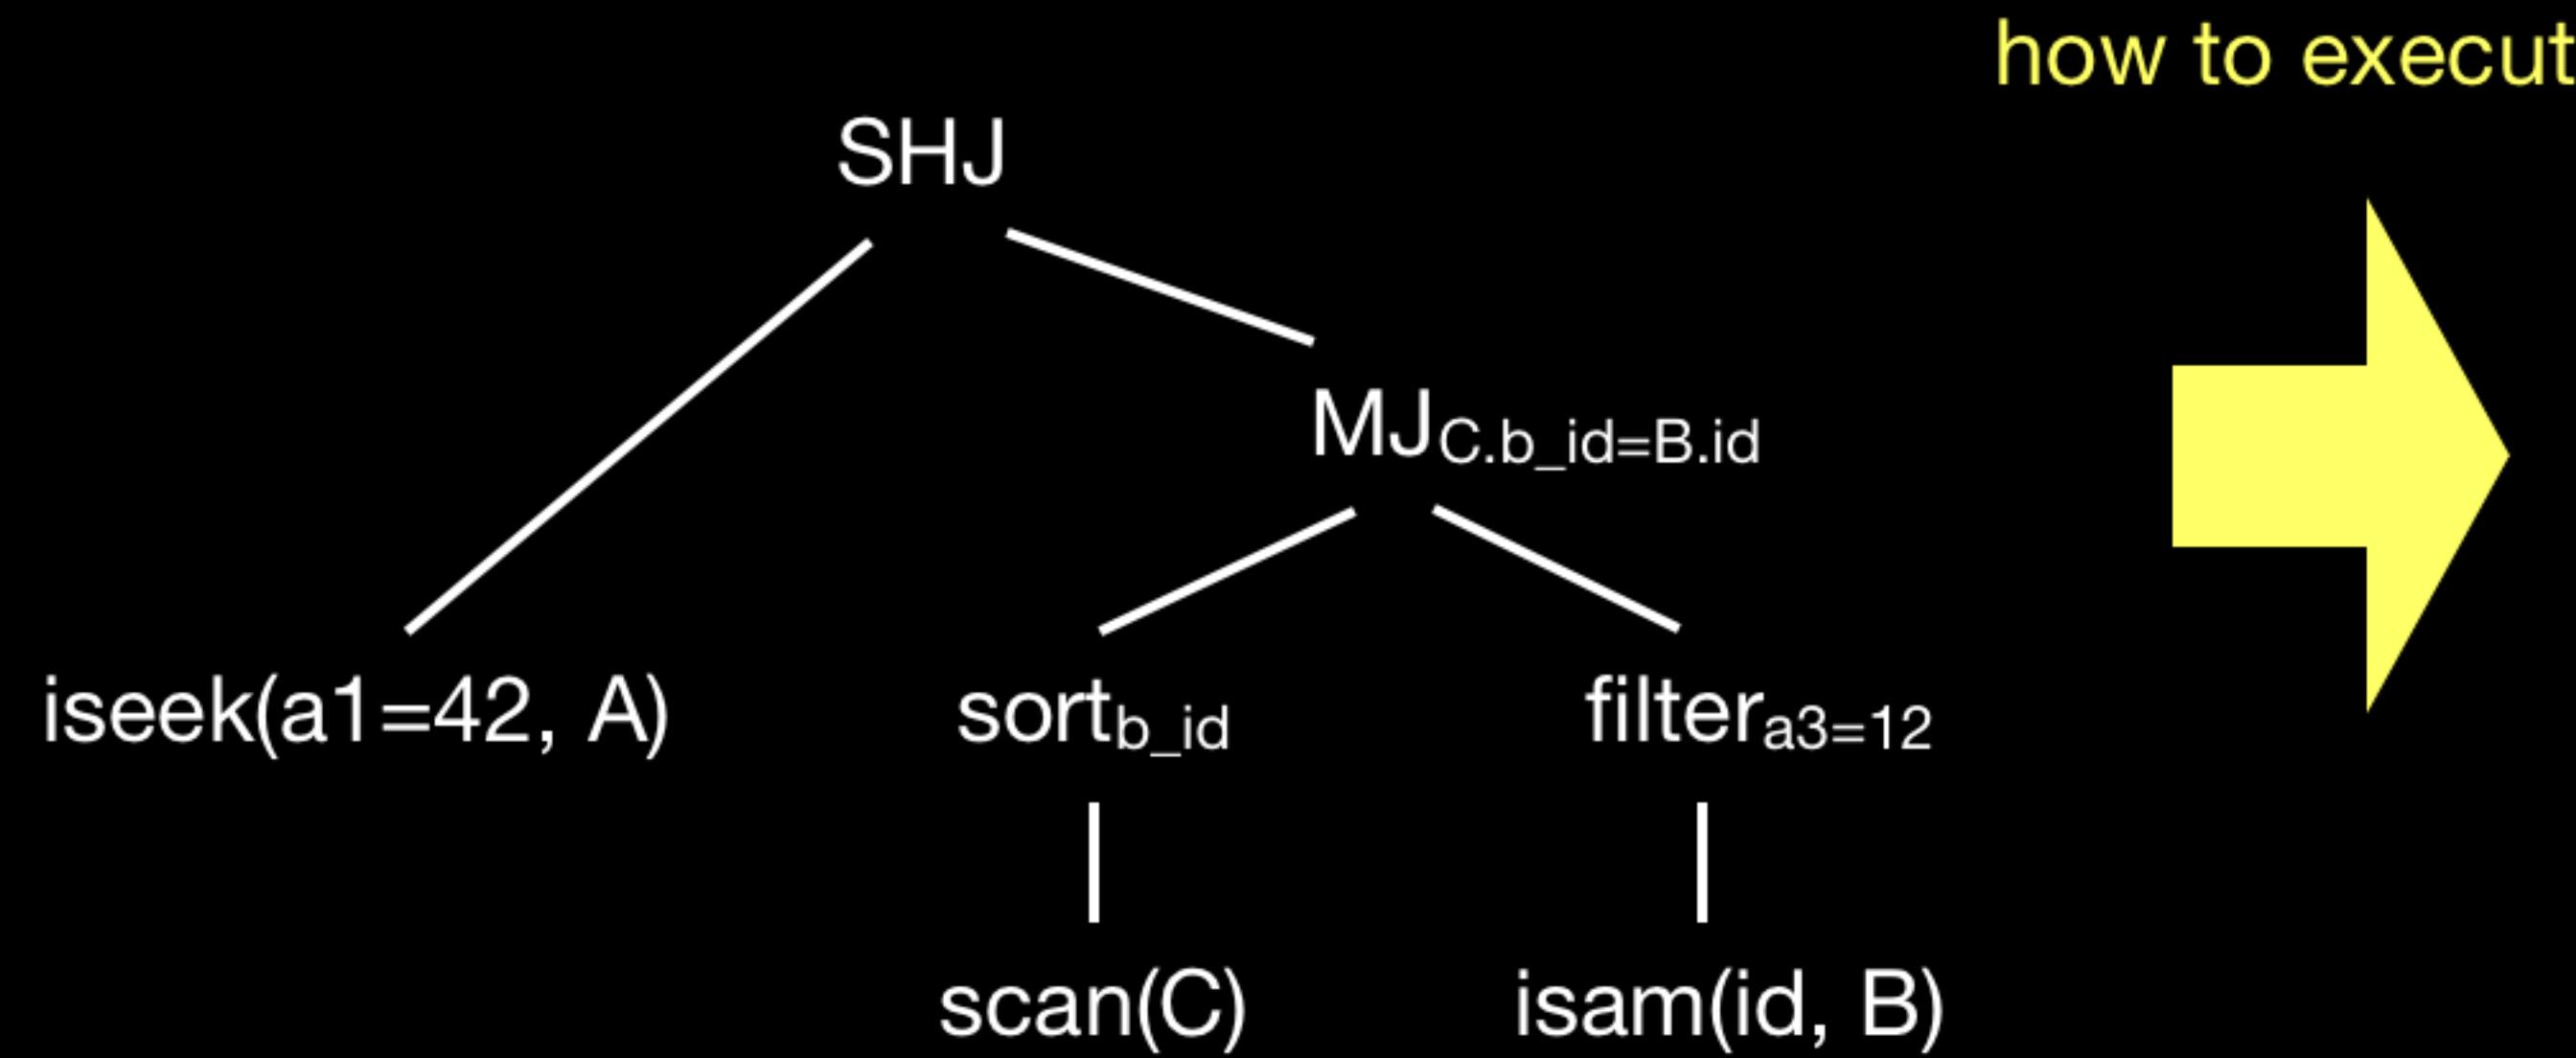
The Story so far:

```
SELECT A.a1, B.a2, B.a3, D.a4  
FROM   A   JOIN B ON A.id=B.a_id  
        JOIN C ON B.id=C.b_id  
WHERE  A.a1=42  
       AND B.a3=12
```

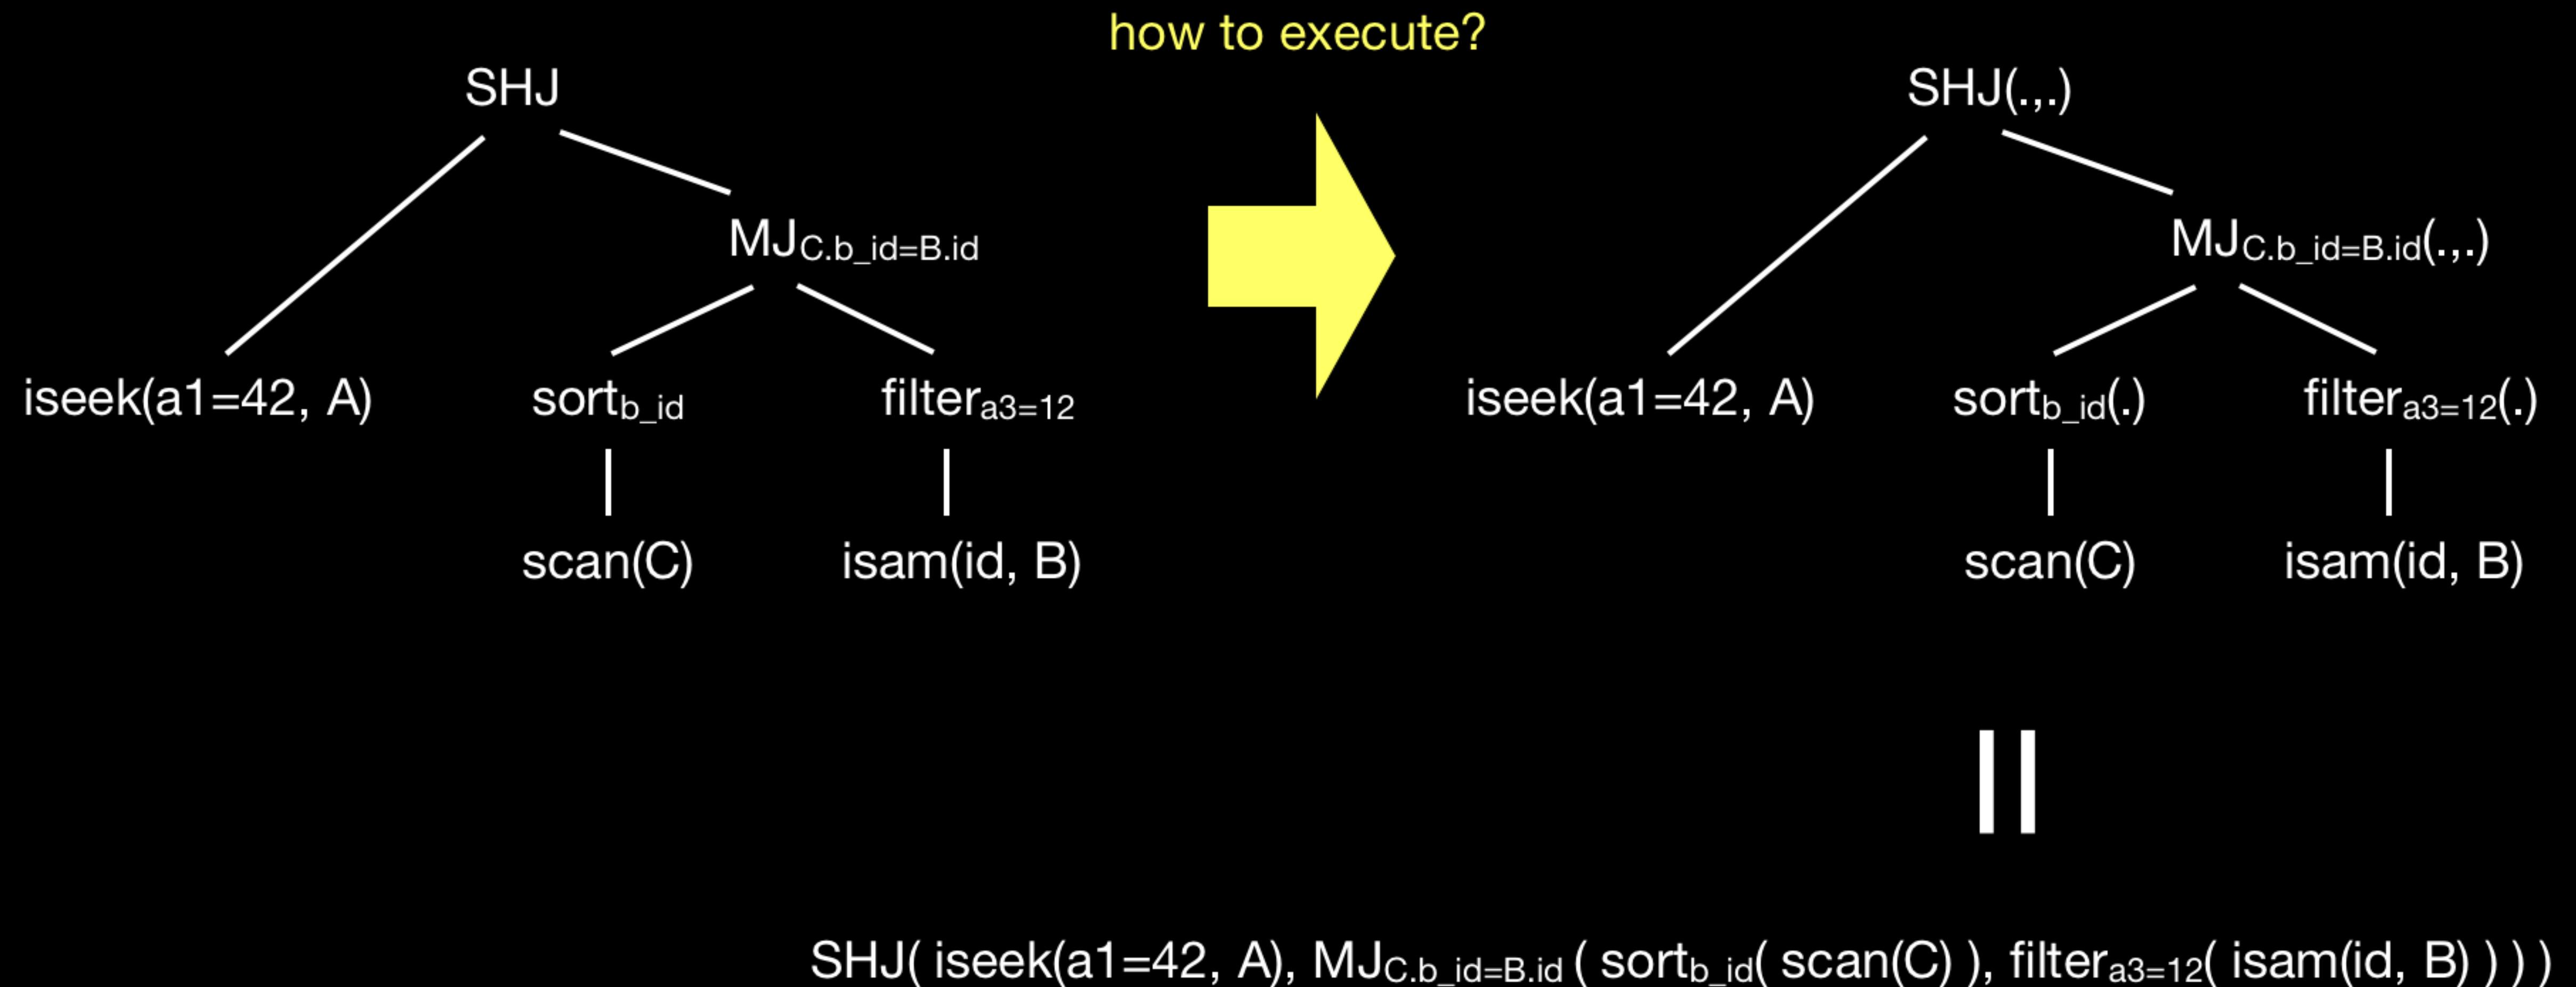
Query Parser/Optimizer



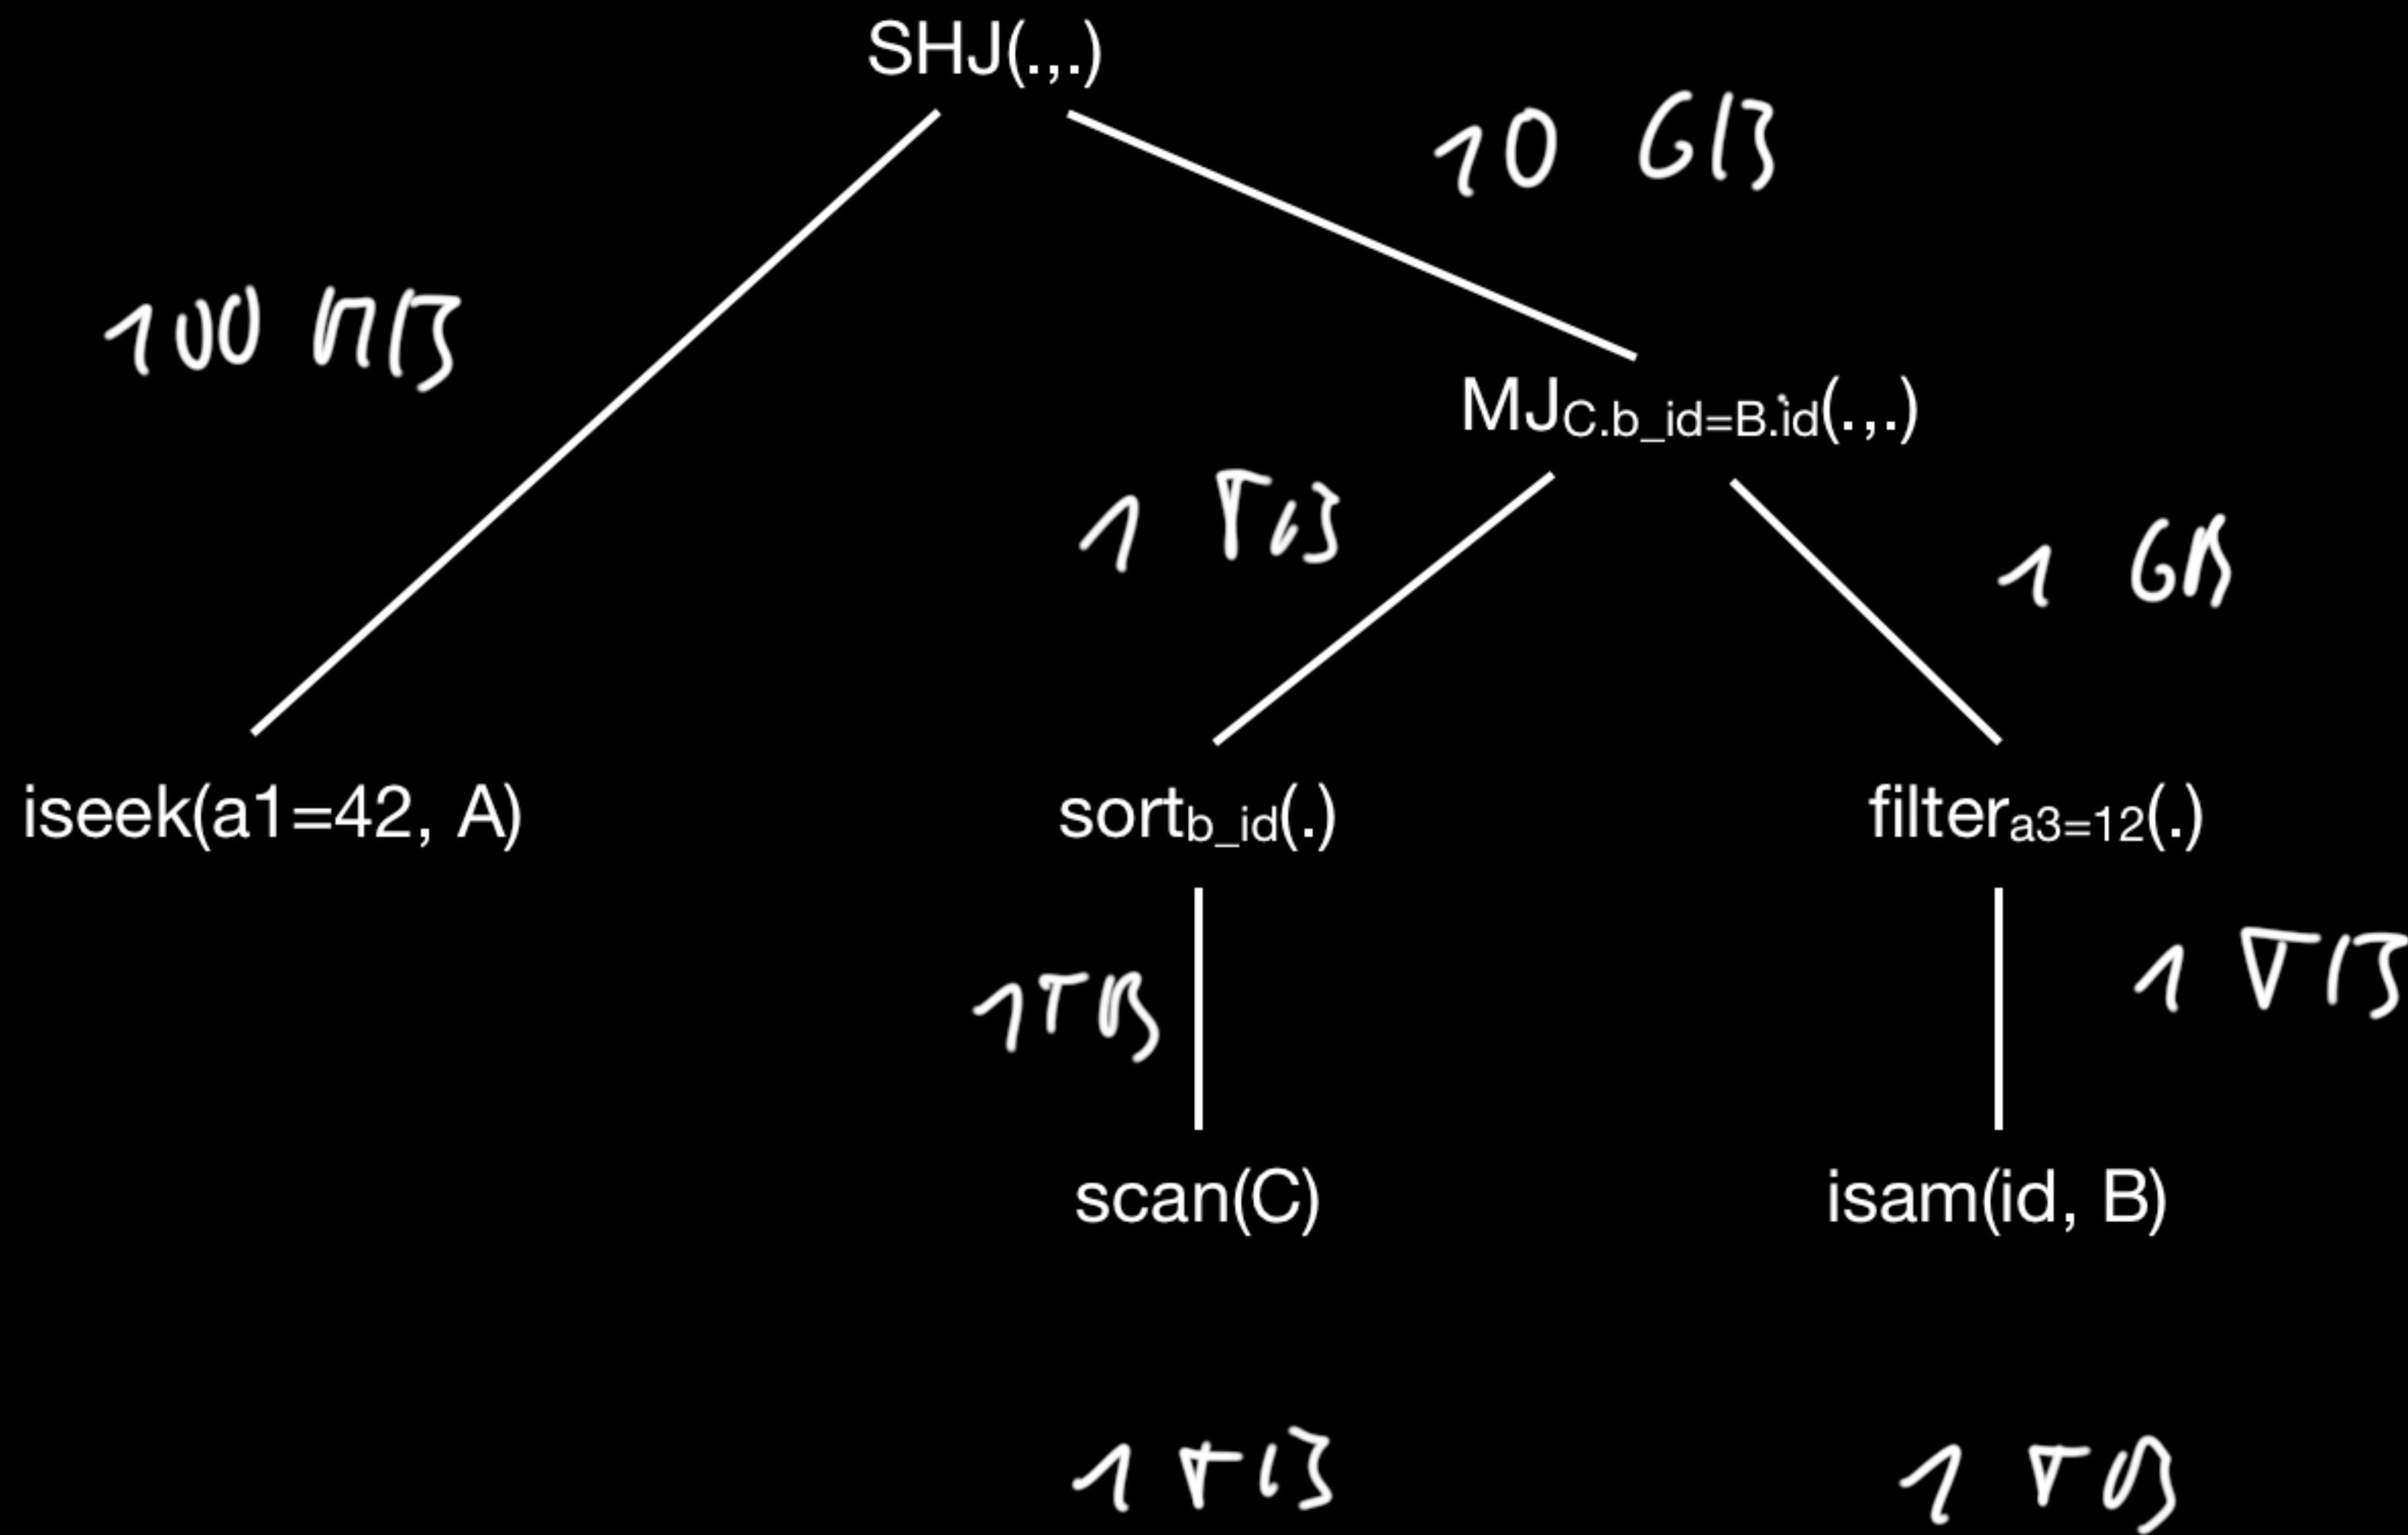
How to Proceed?



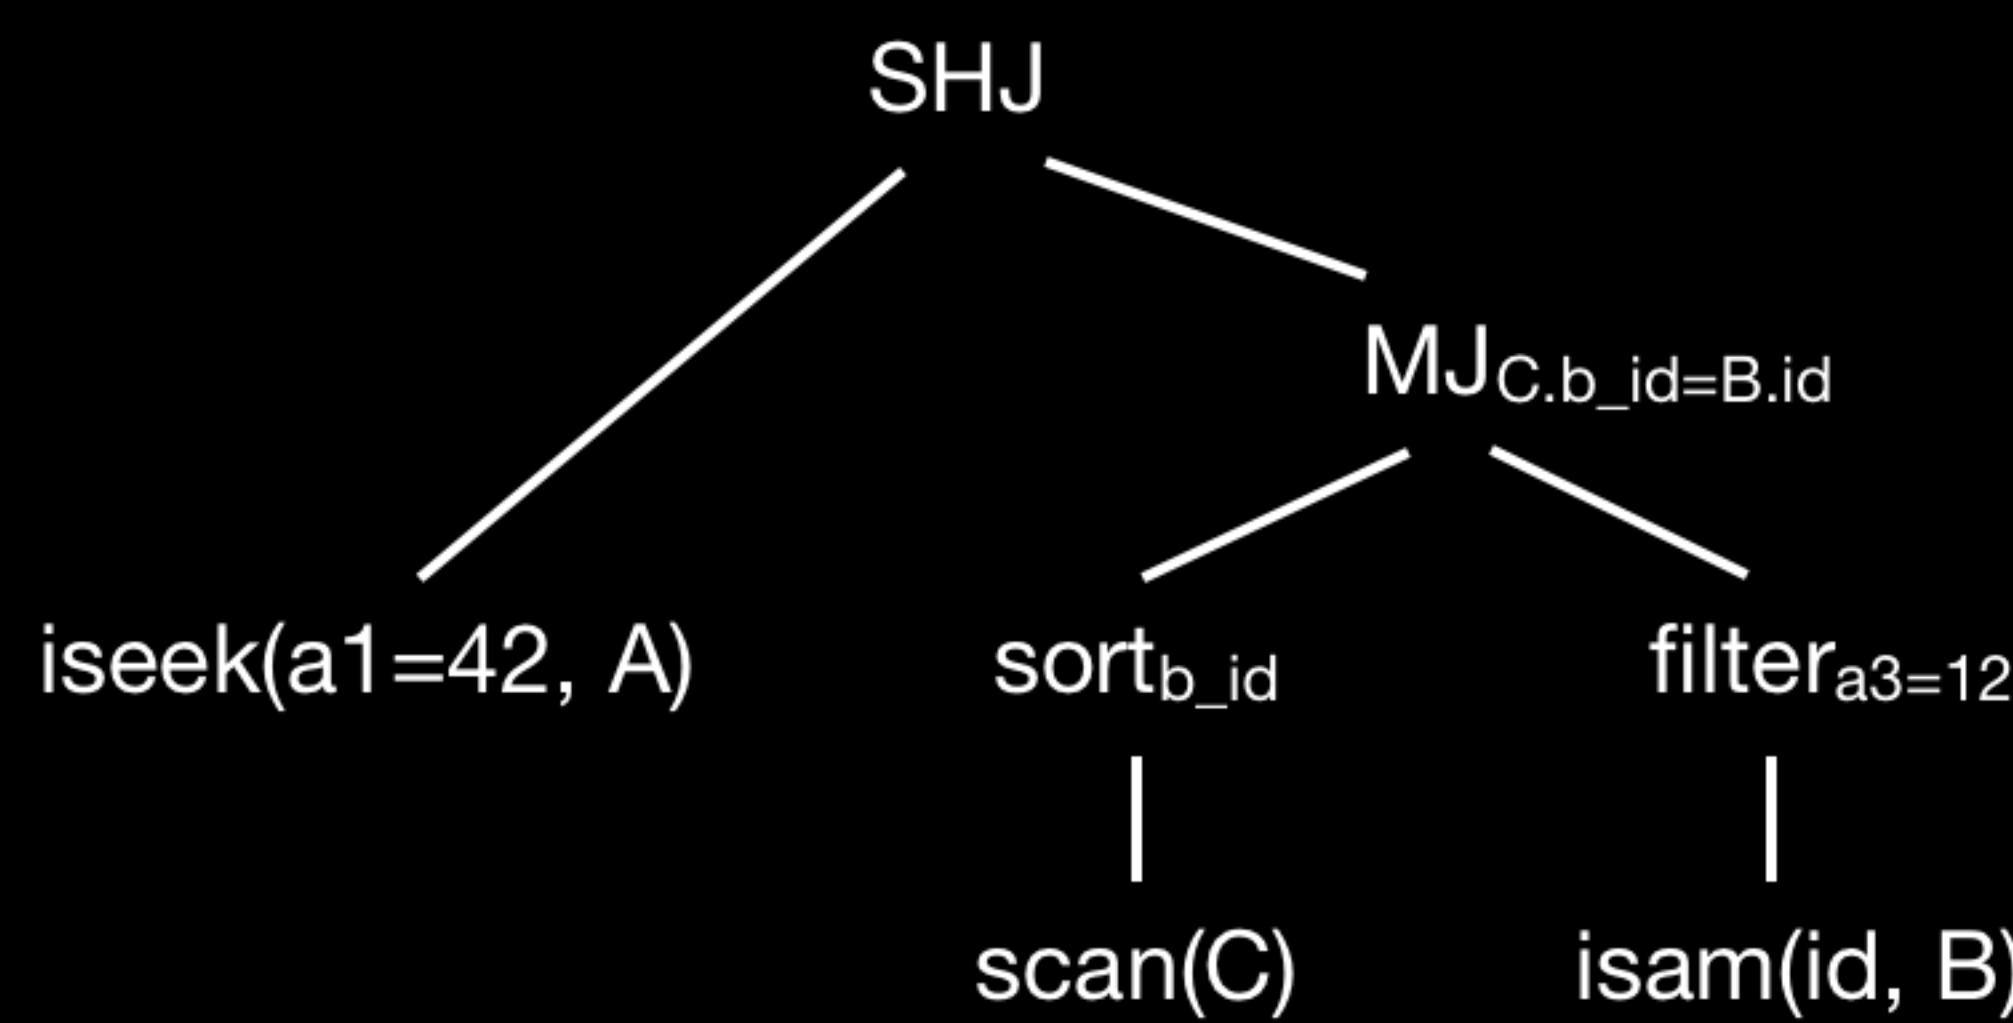
Approach 1: Function Library



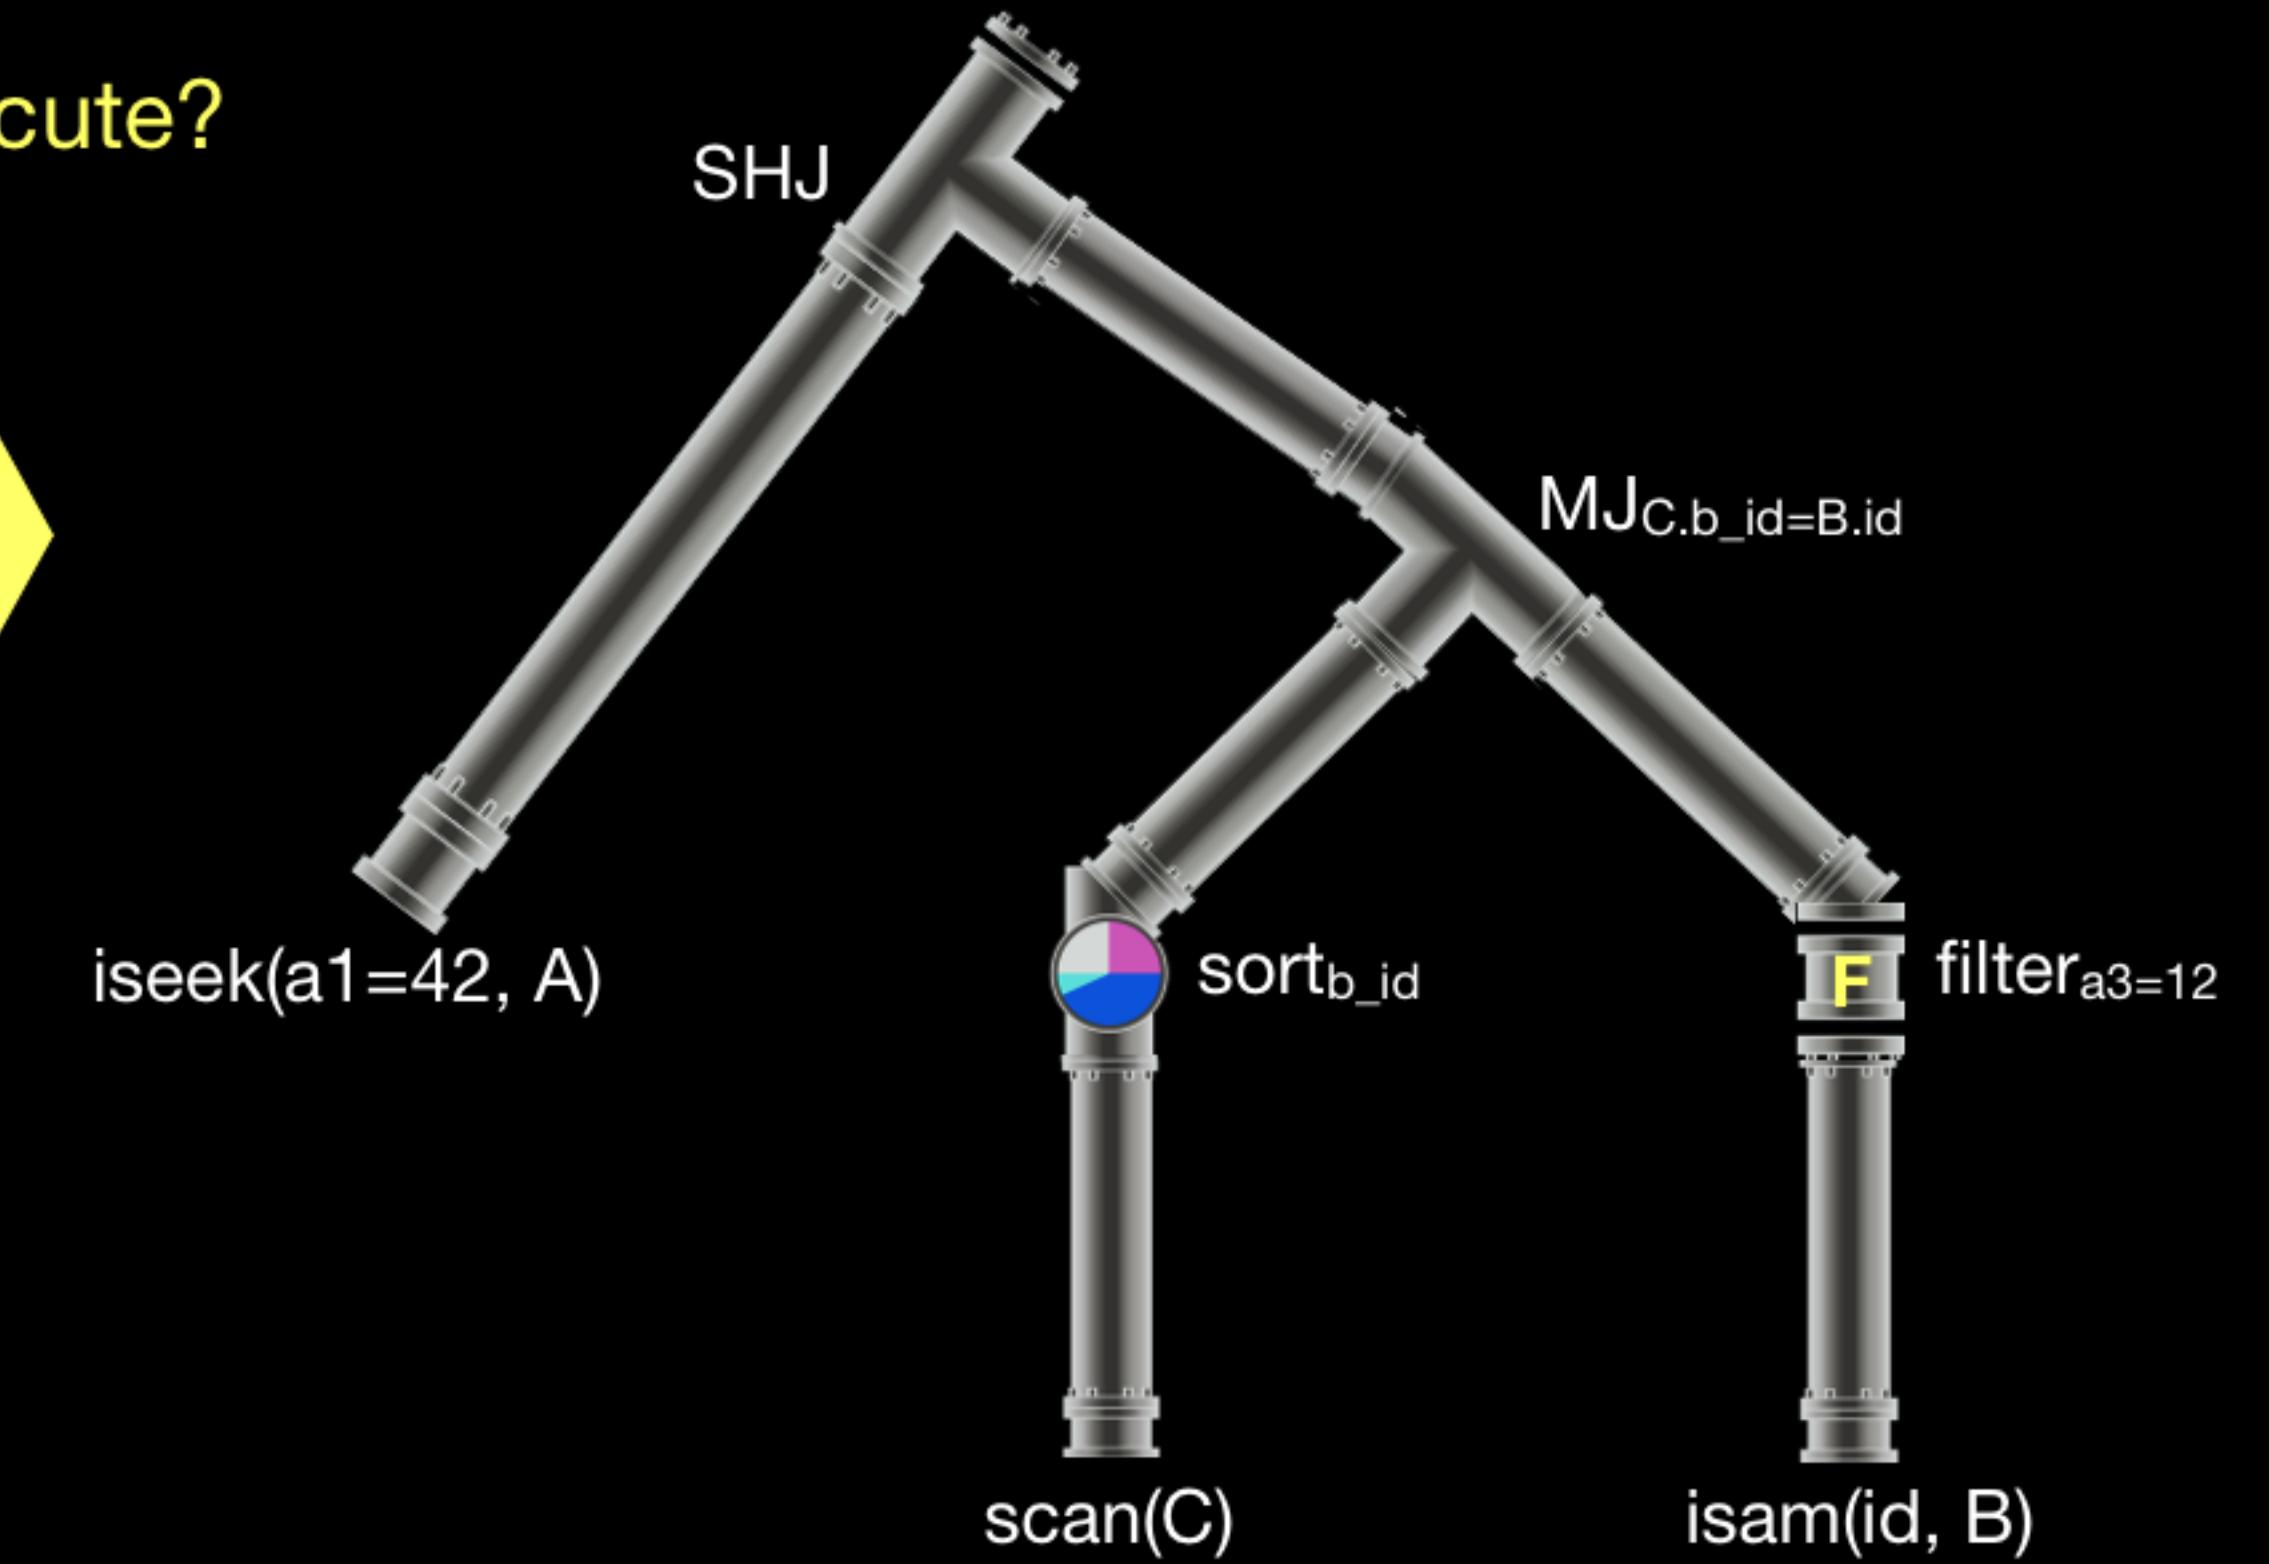
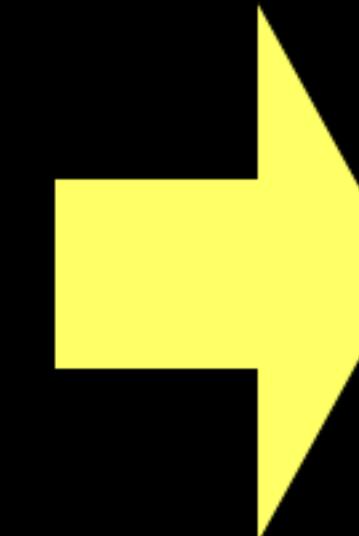
Problem: Intermediate Results



Approach 2: Pipelining



how to execute?

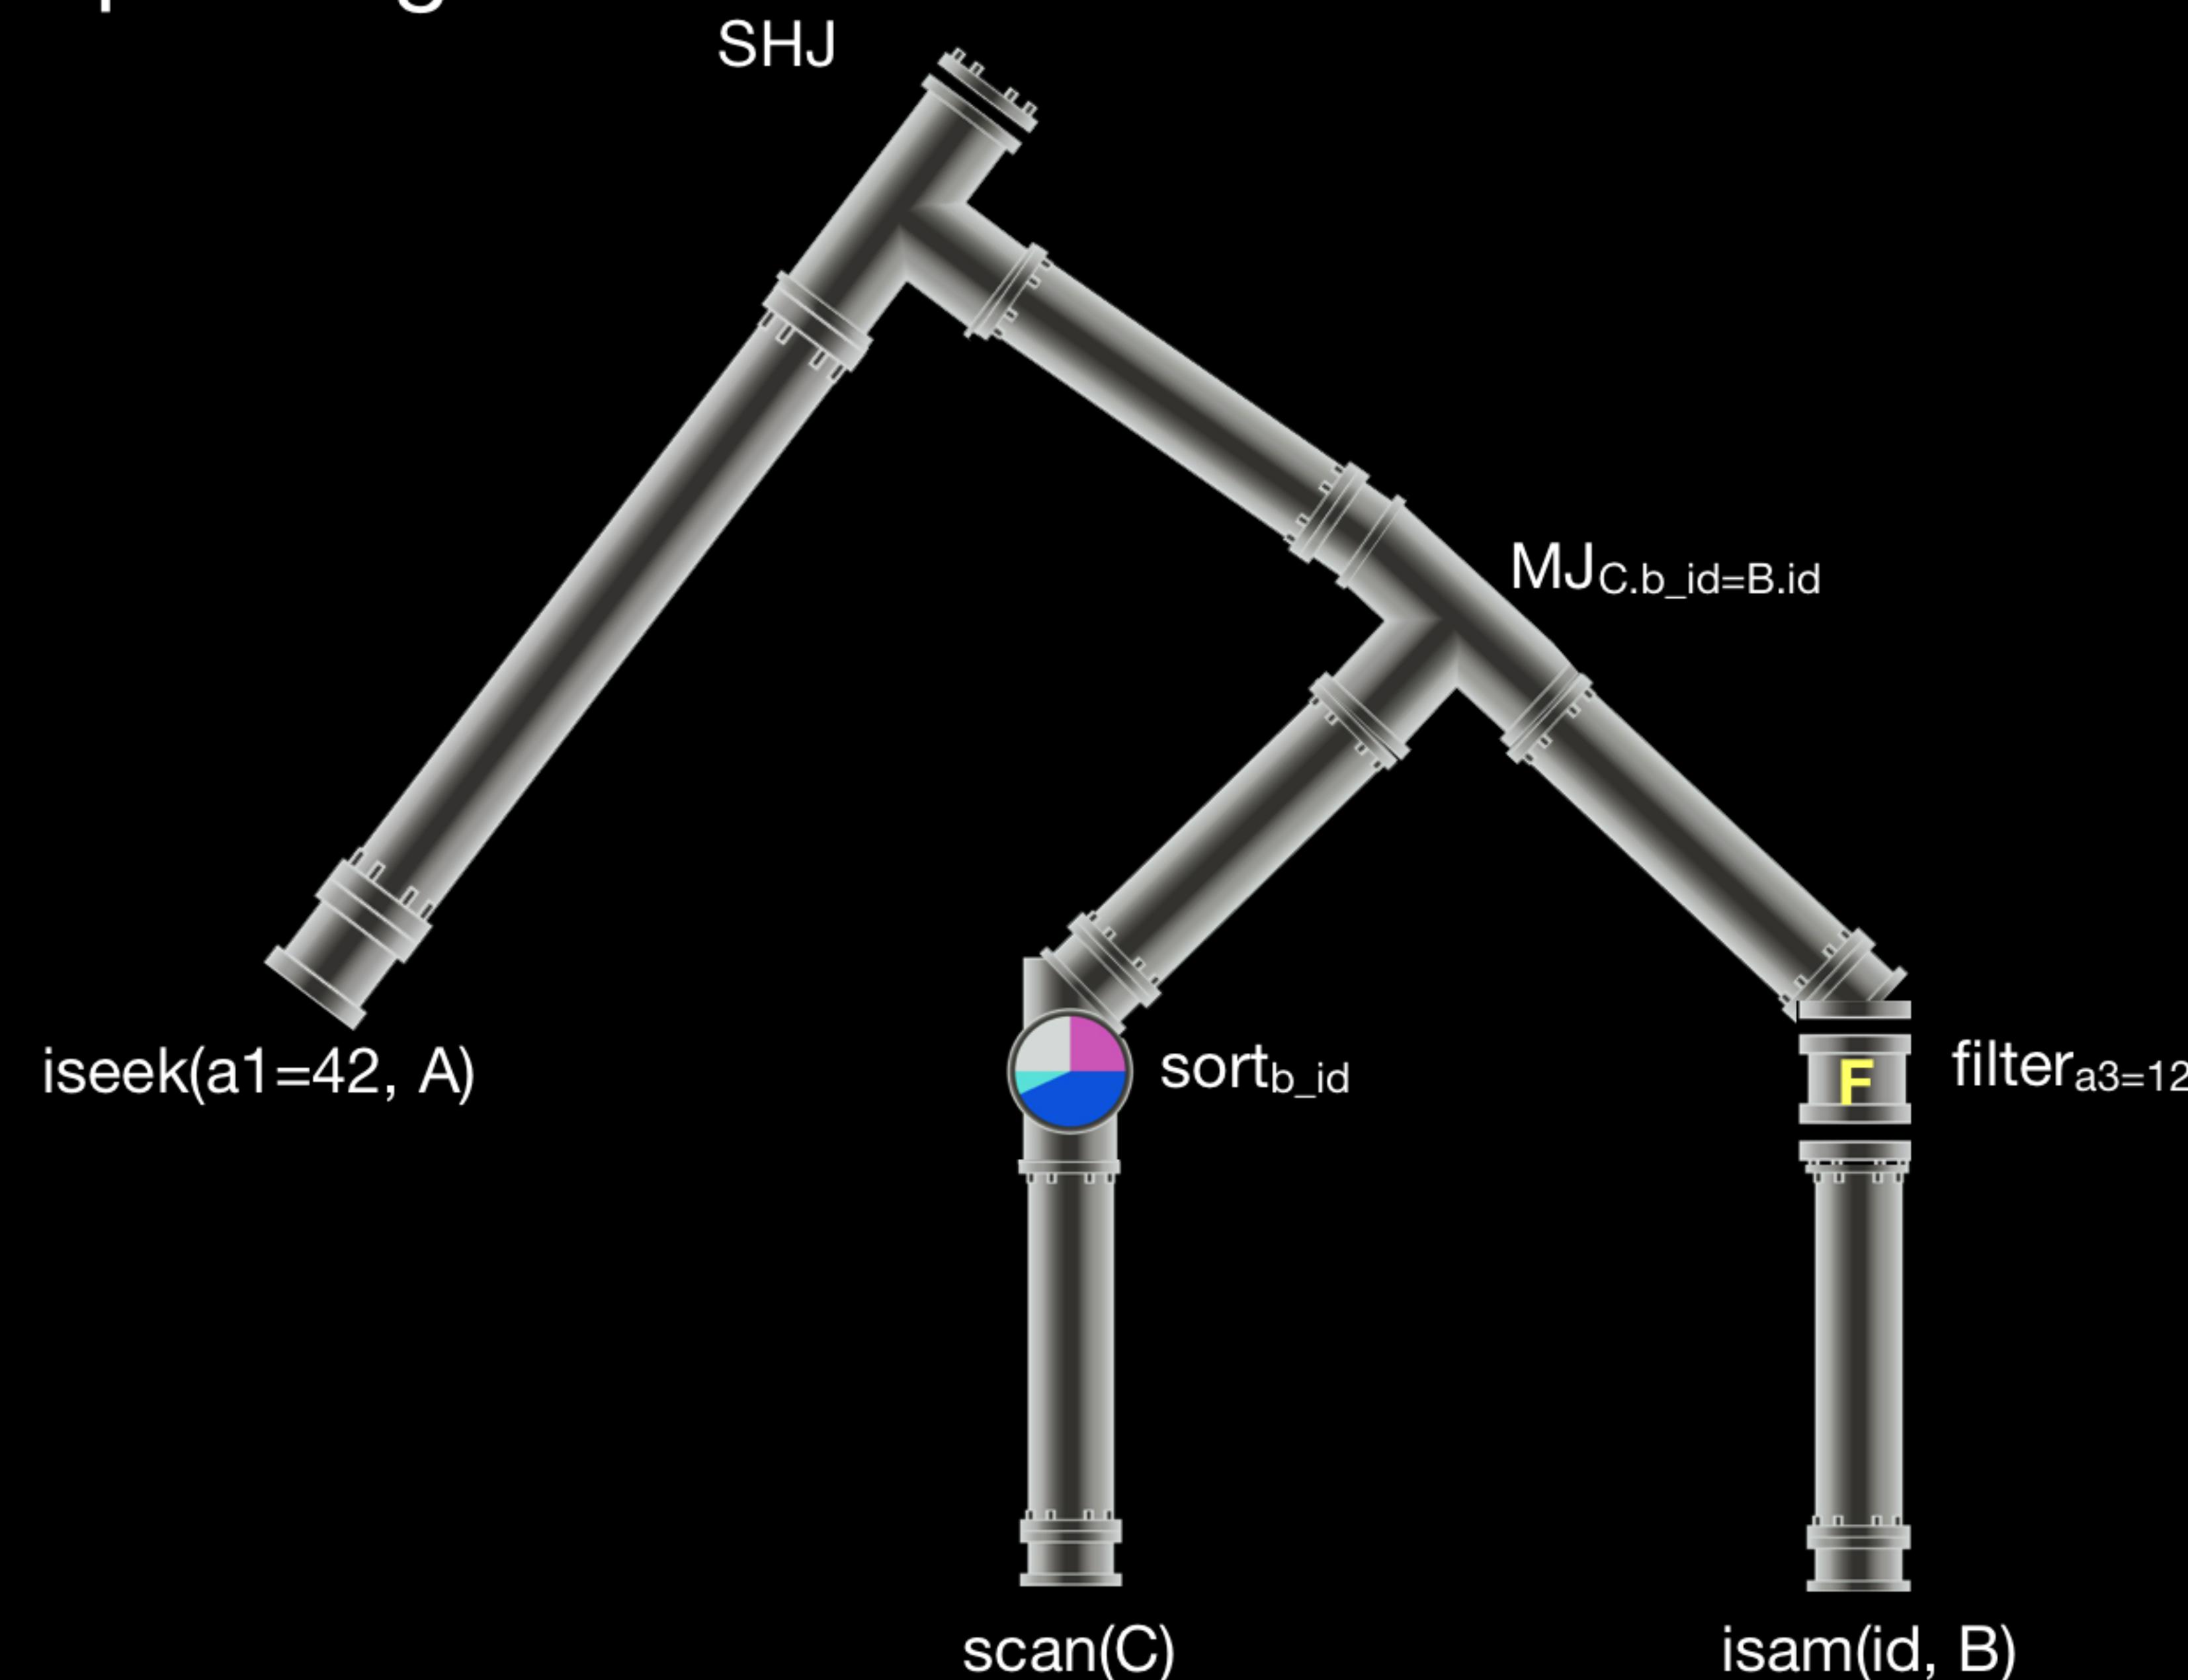


vh:x / l:lux

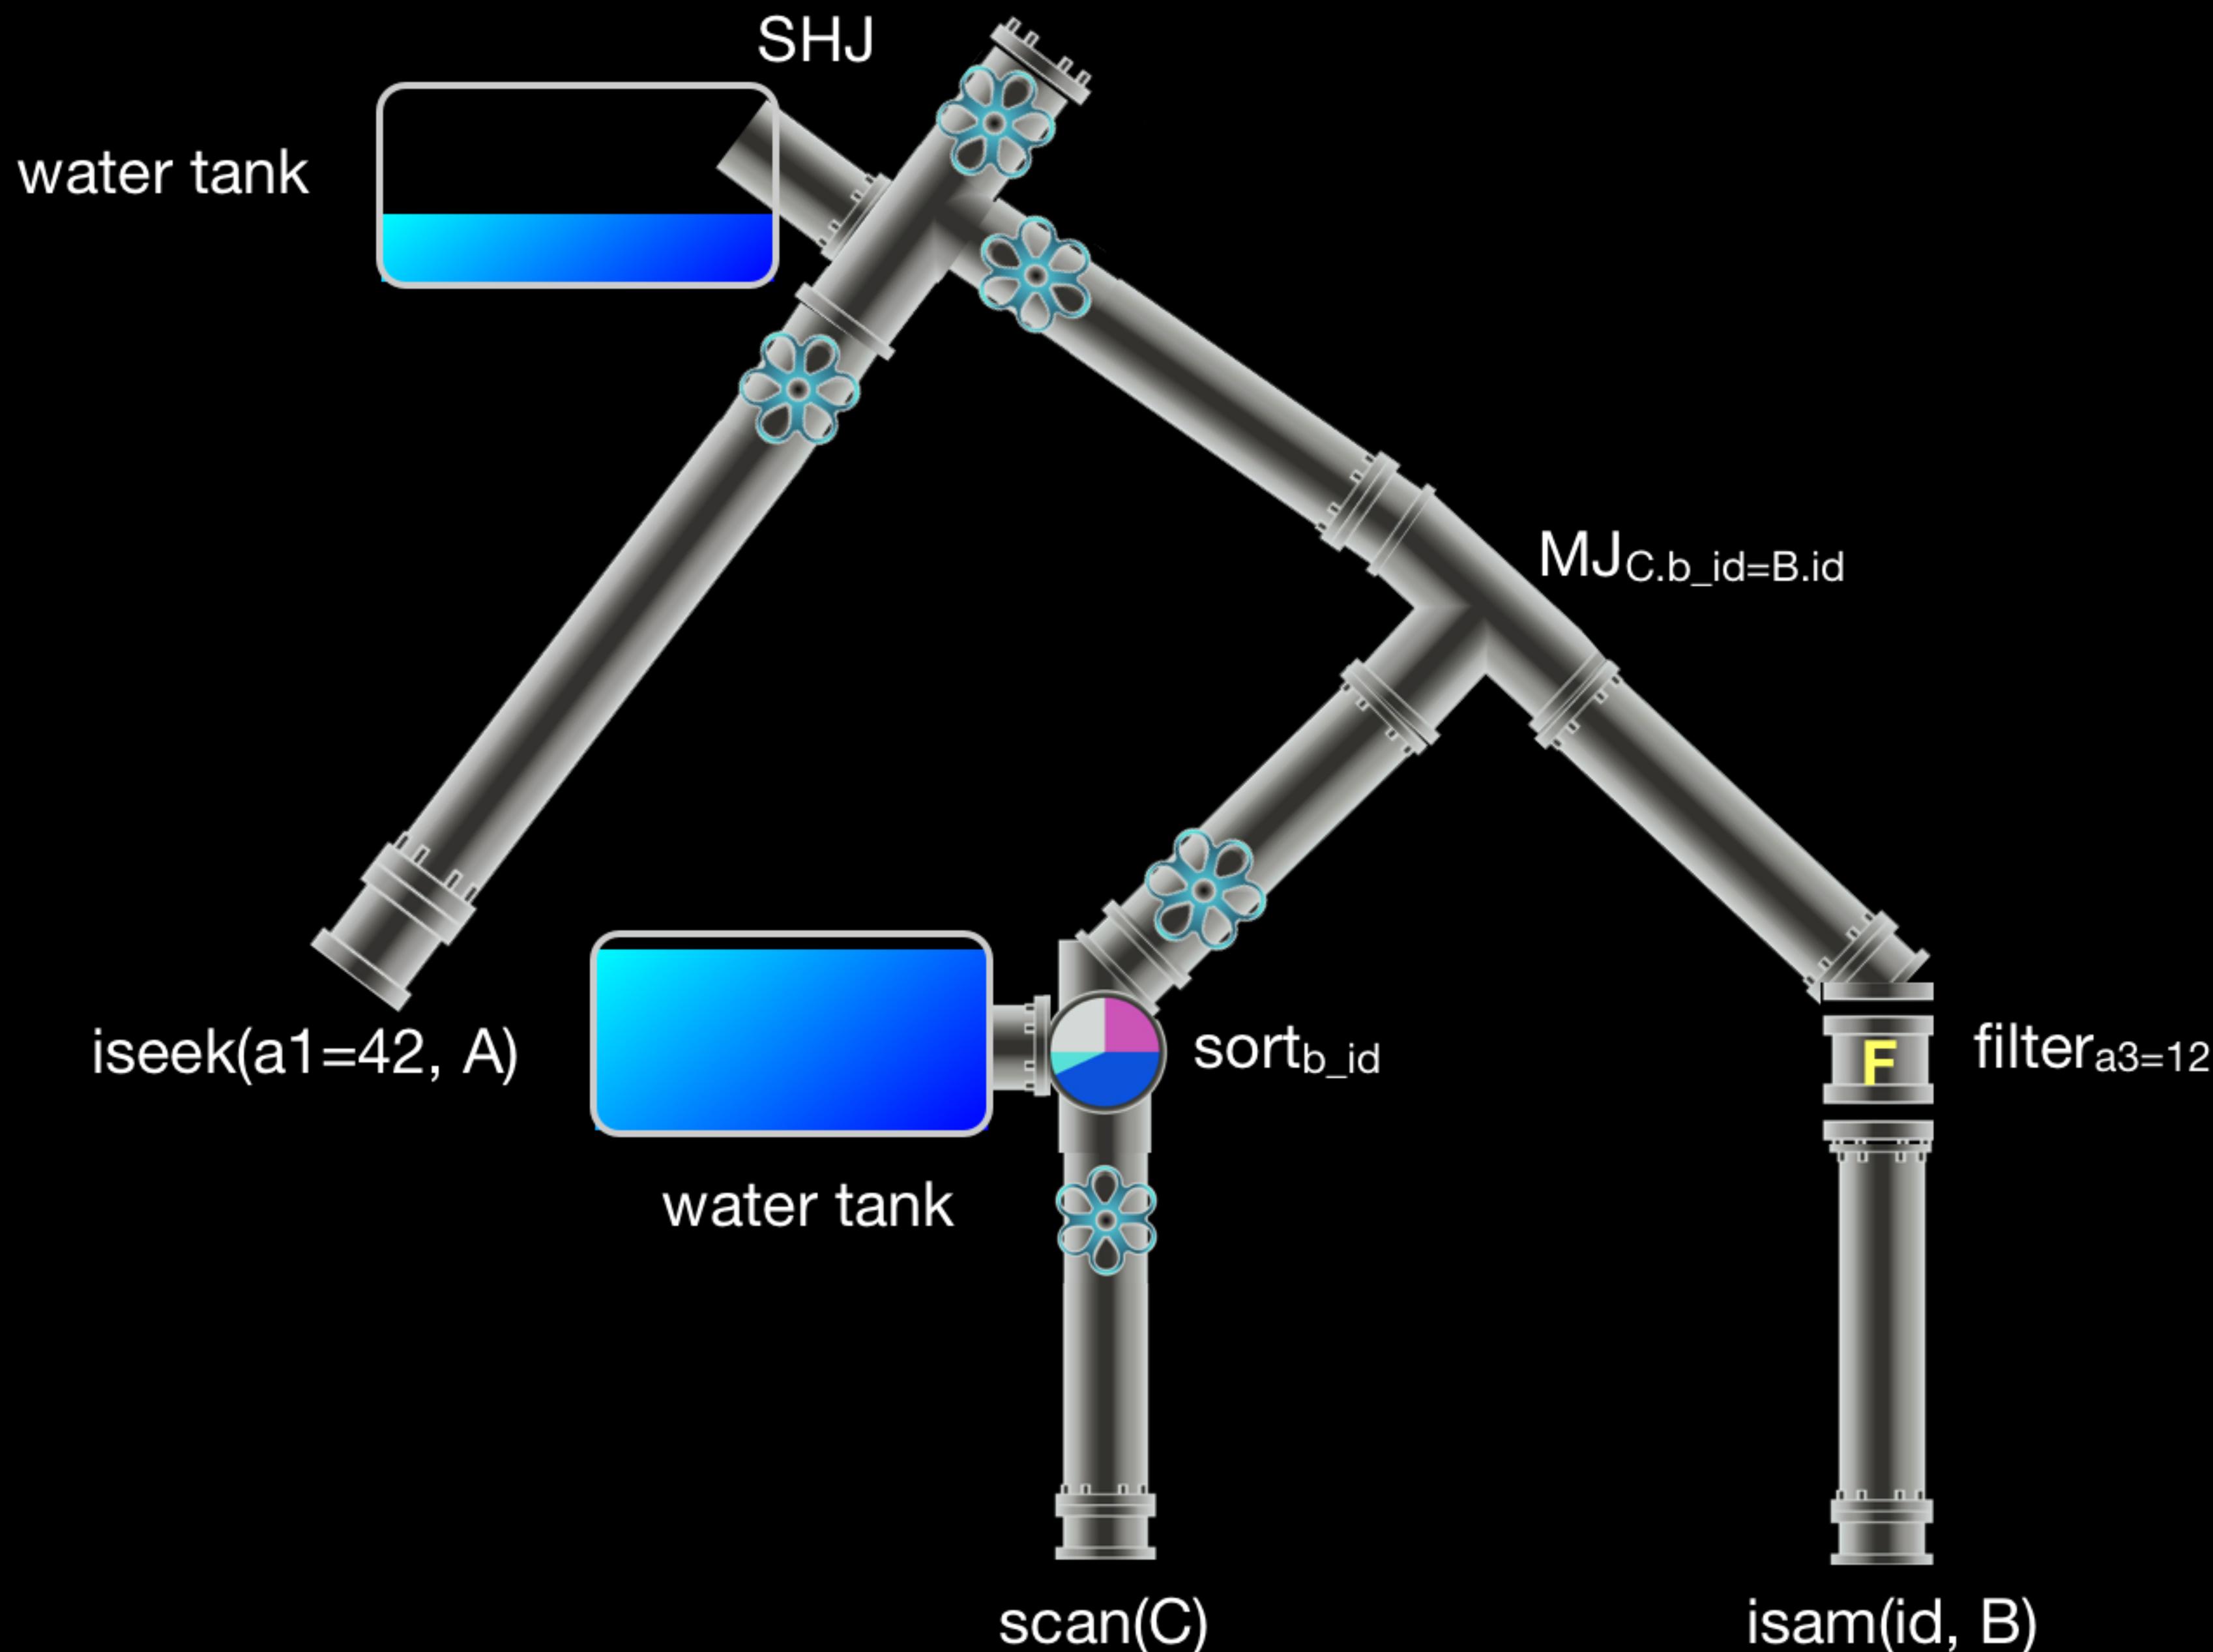
grep <string> | grep ... > out.txt

↑ pipe symbol

Approach 2: Pipelining



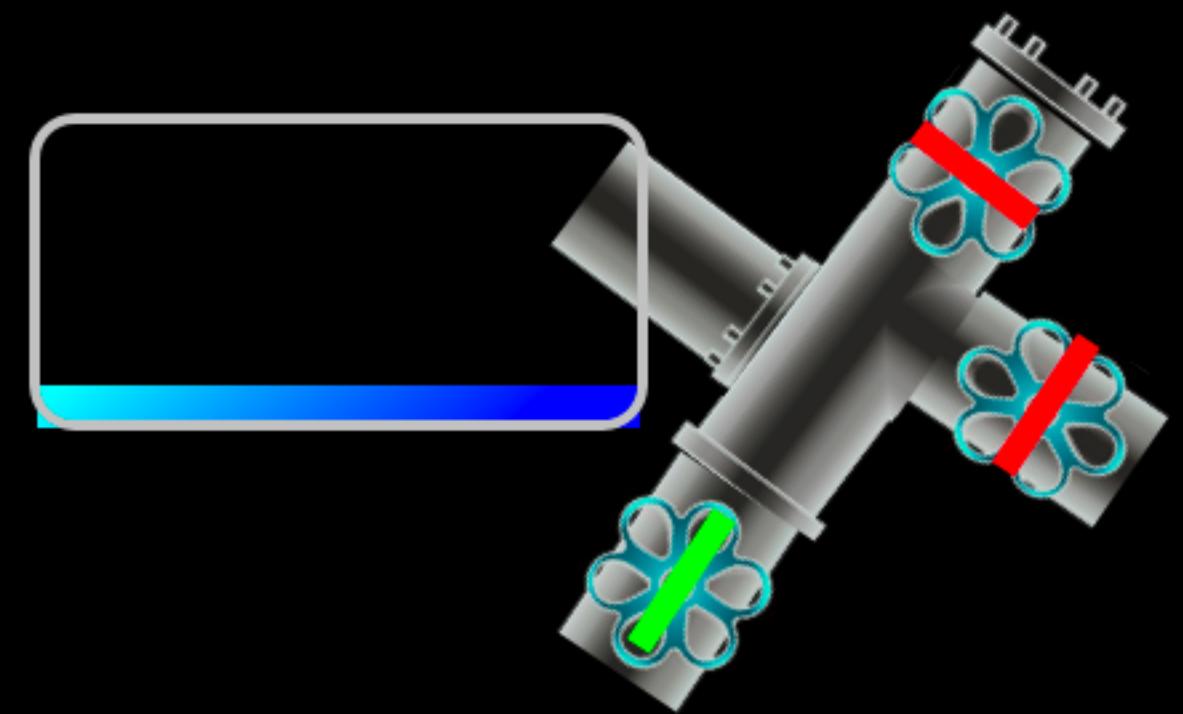
Pipeline Breakers?



Stages of SHJ/INLJ

blocking

(1) reading

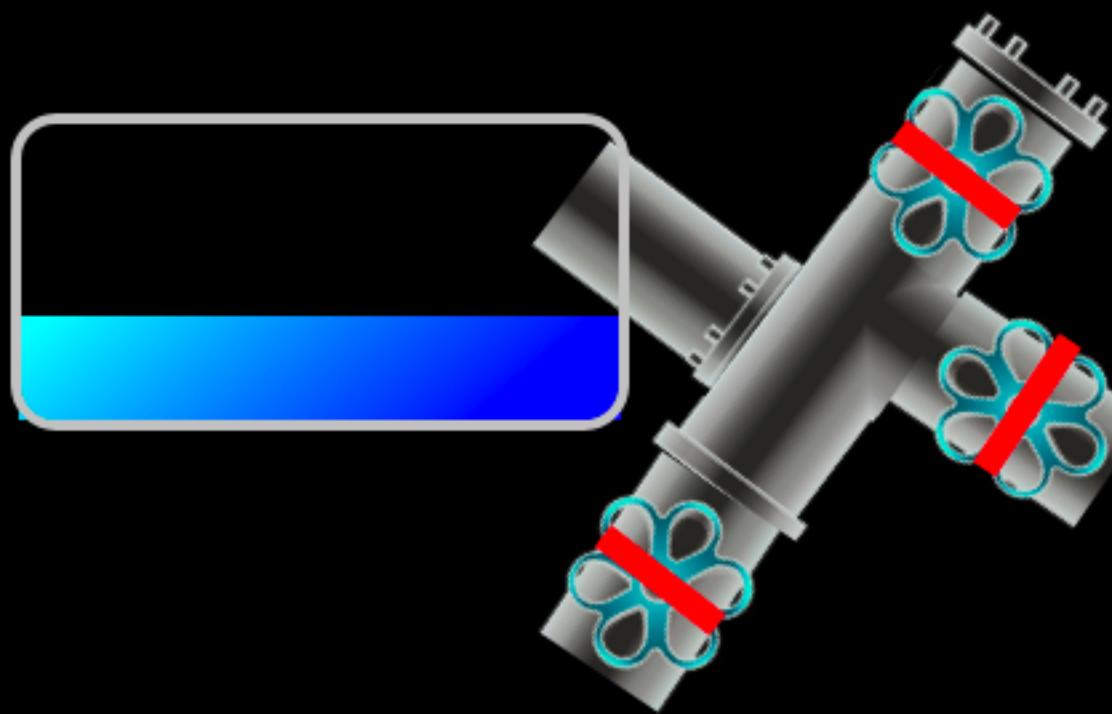


(A) just buffers

(B) tuple-wise insert

blocking

(2) left input exhausted

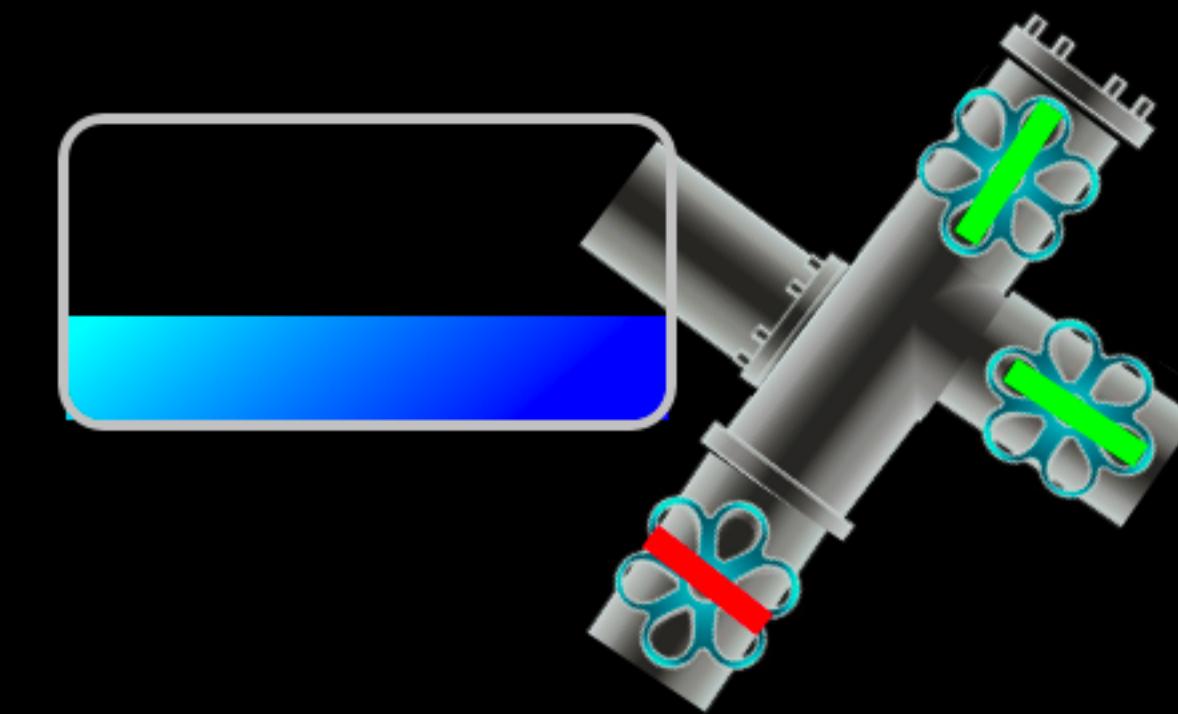


(A) bulkload index

(B) Done.

non-blocking

(3) Looping and Probing

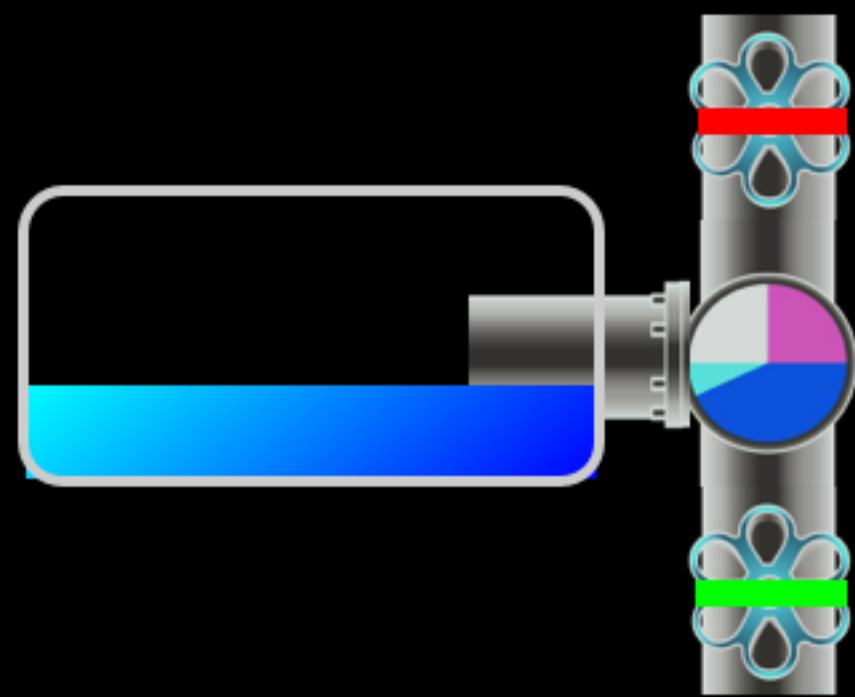


(A+B) index probe

Stages of Quick-Sorting

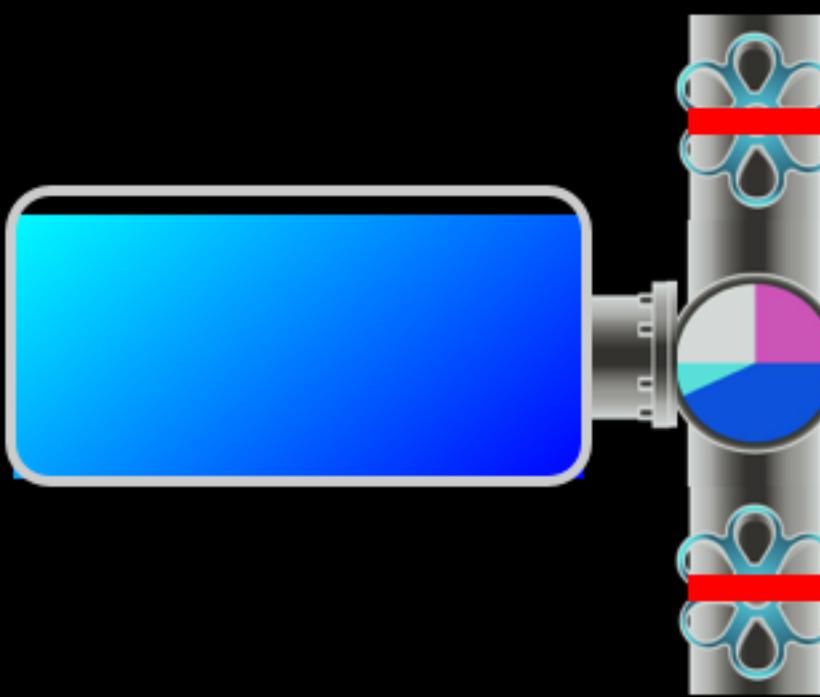
blocking

(1) reading



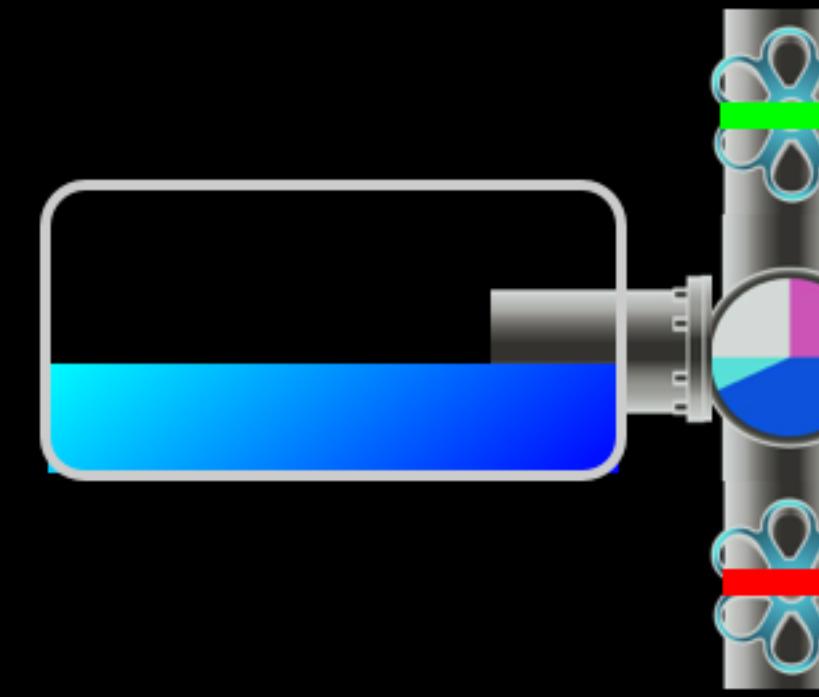
blocking

(2) sorting



non-blocking

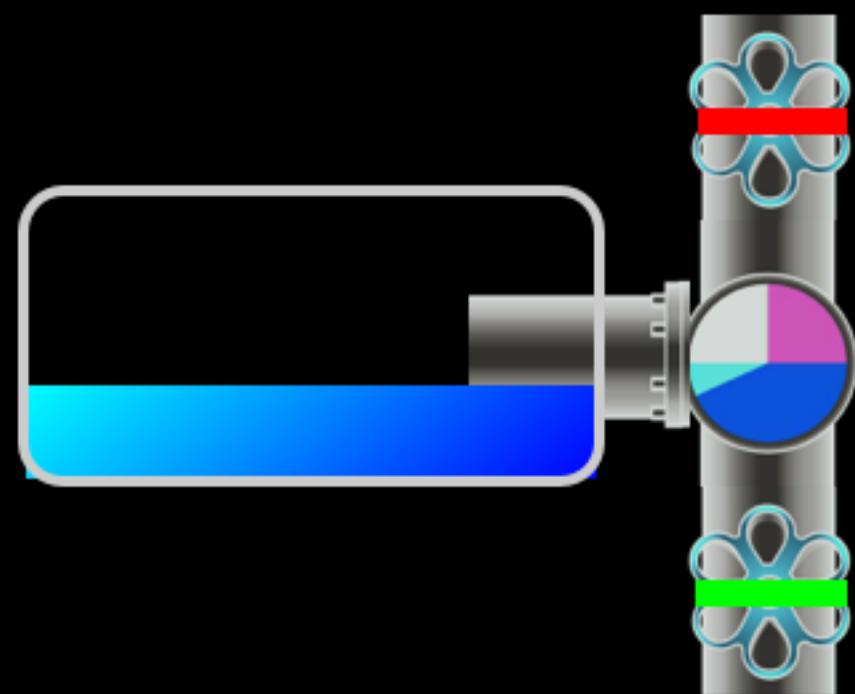
(3) outputting



Stages of External Merge Sort

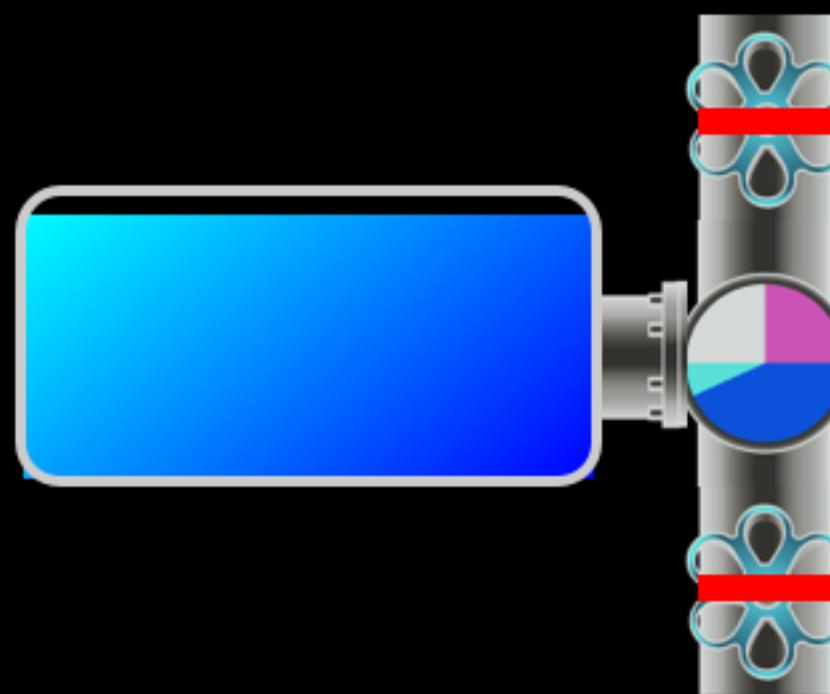
block-
by

- (1) reading
(with run generation)



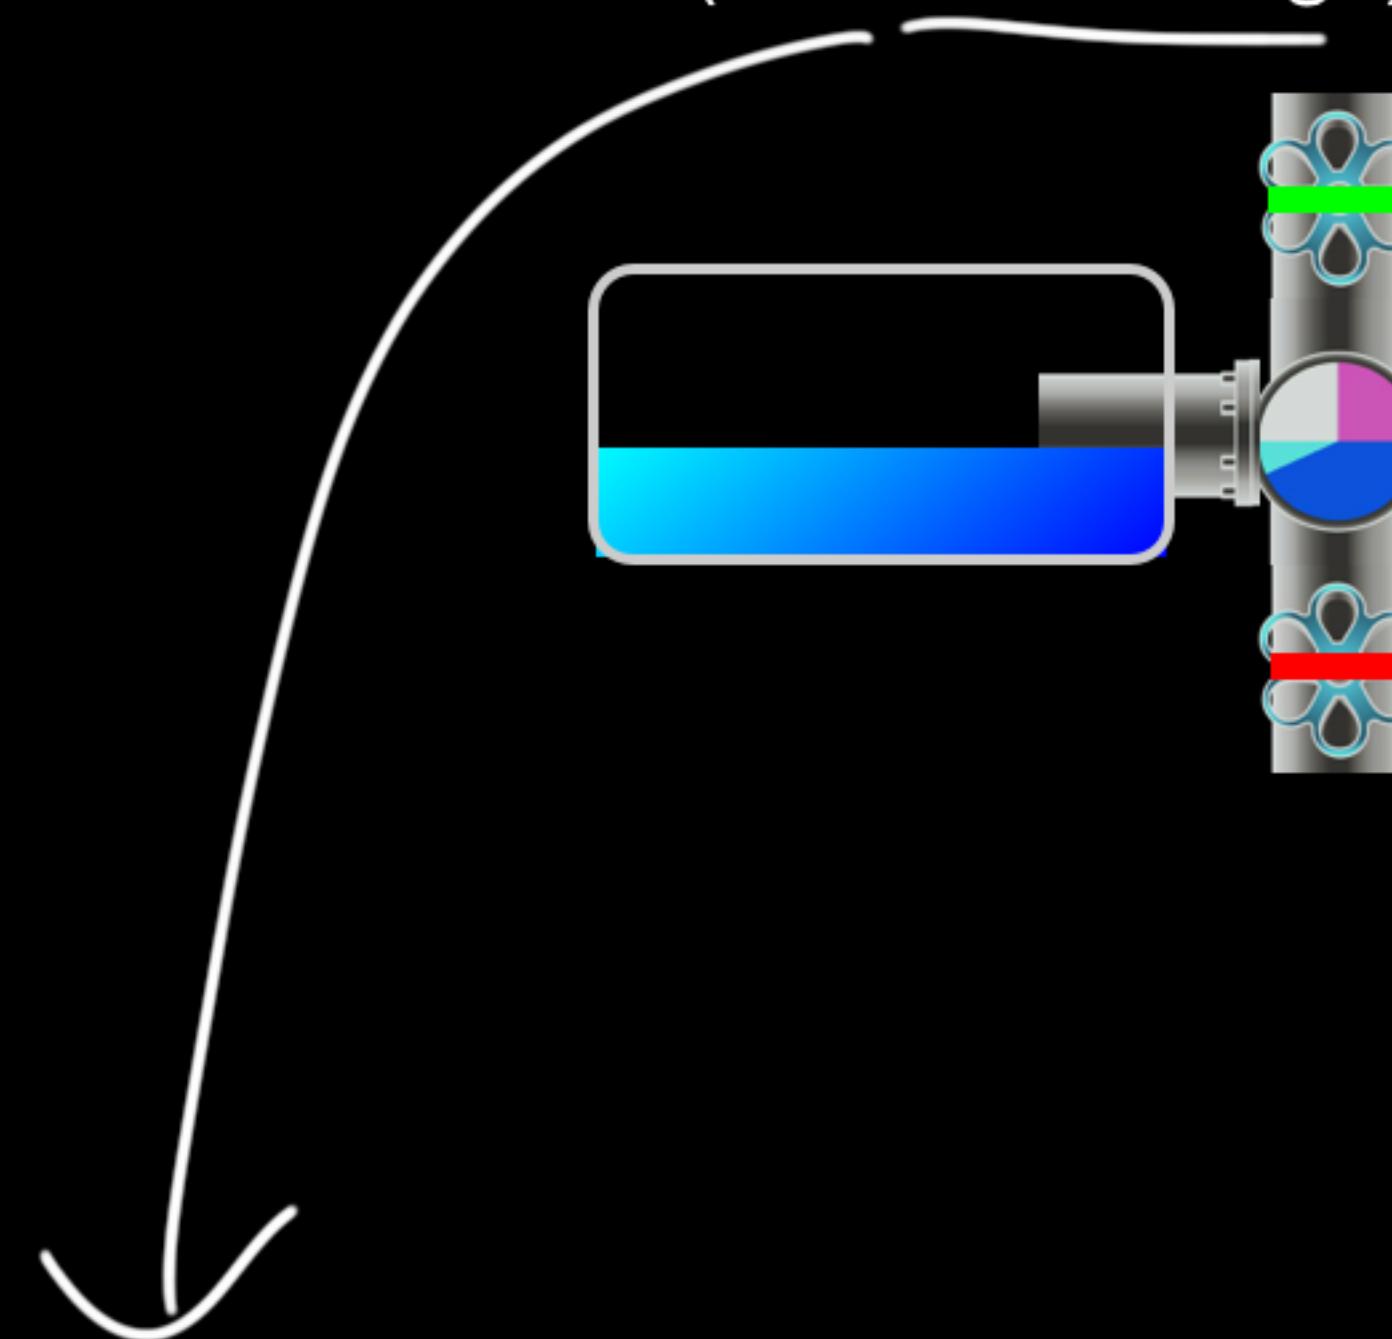
block-
by

- (2) merging
(without final merge)



hol-
block-
by

- (3) outputting
(with final merge)

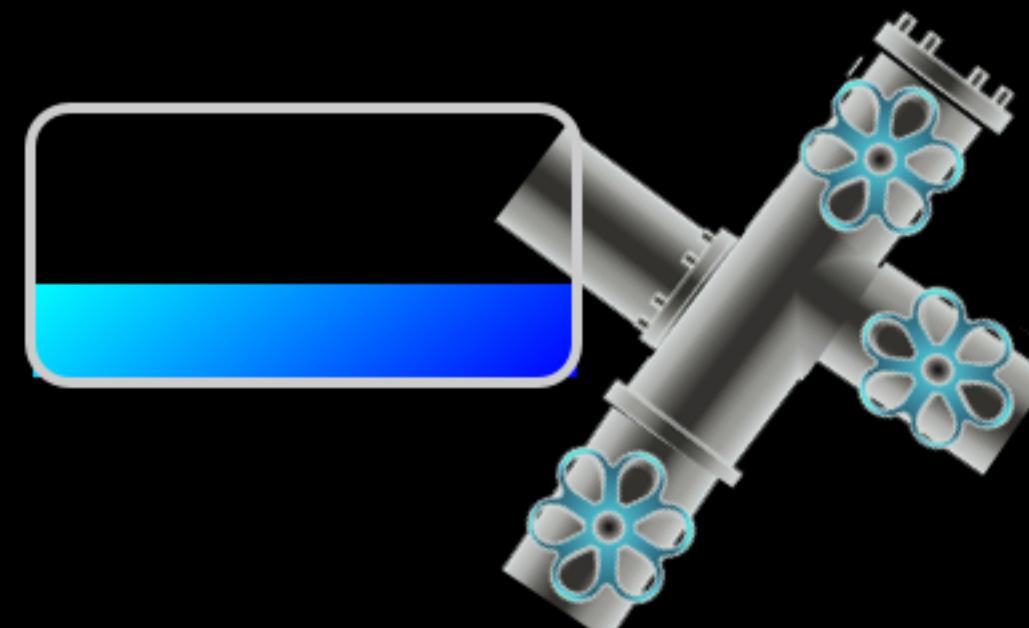


online-
merge

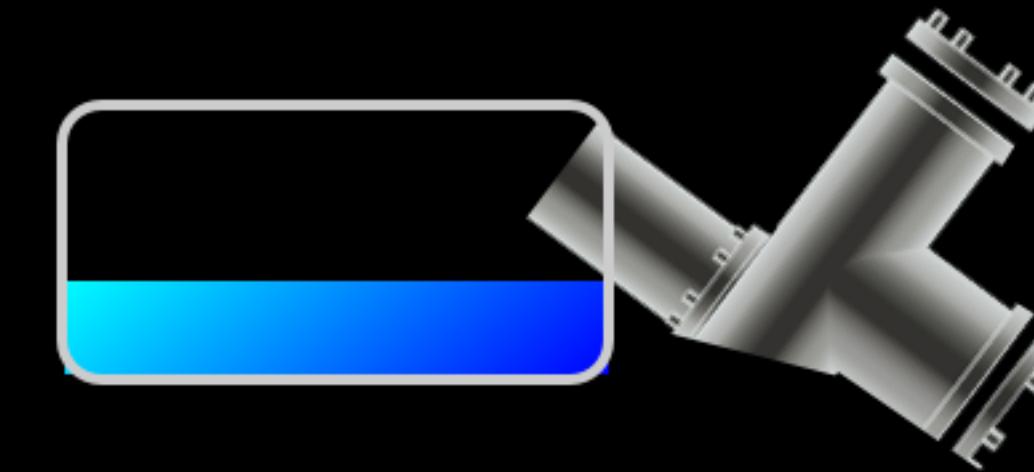
Blocking

vs

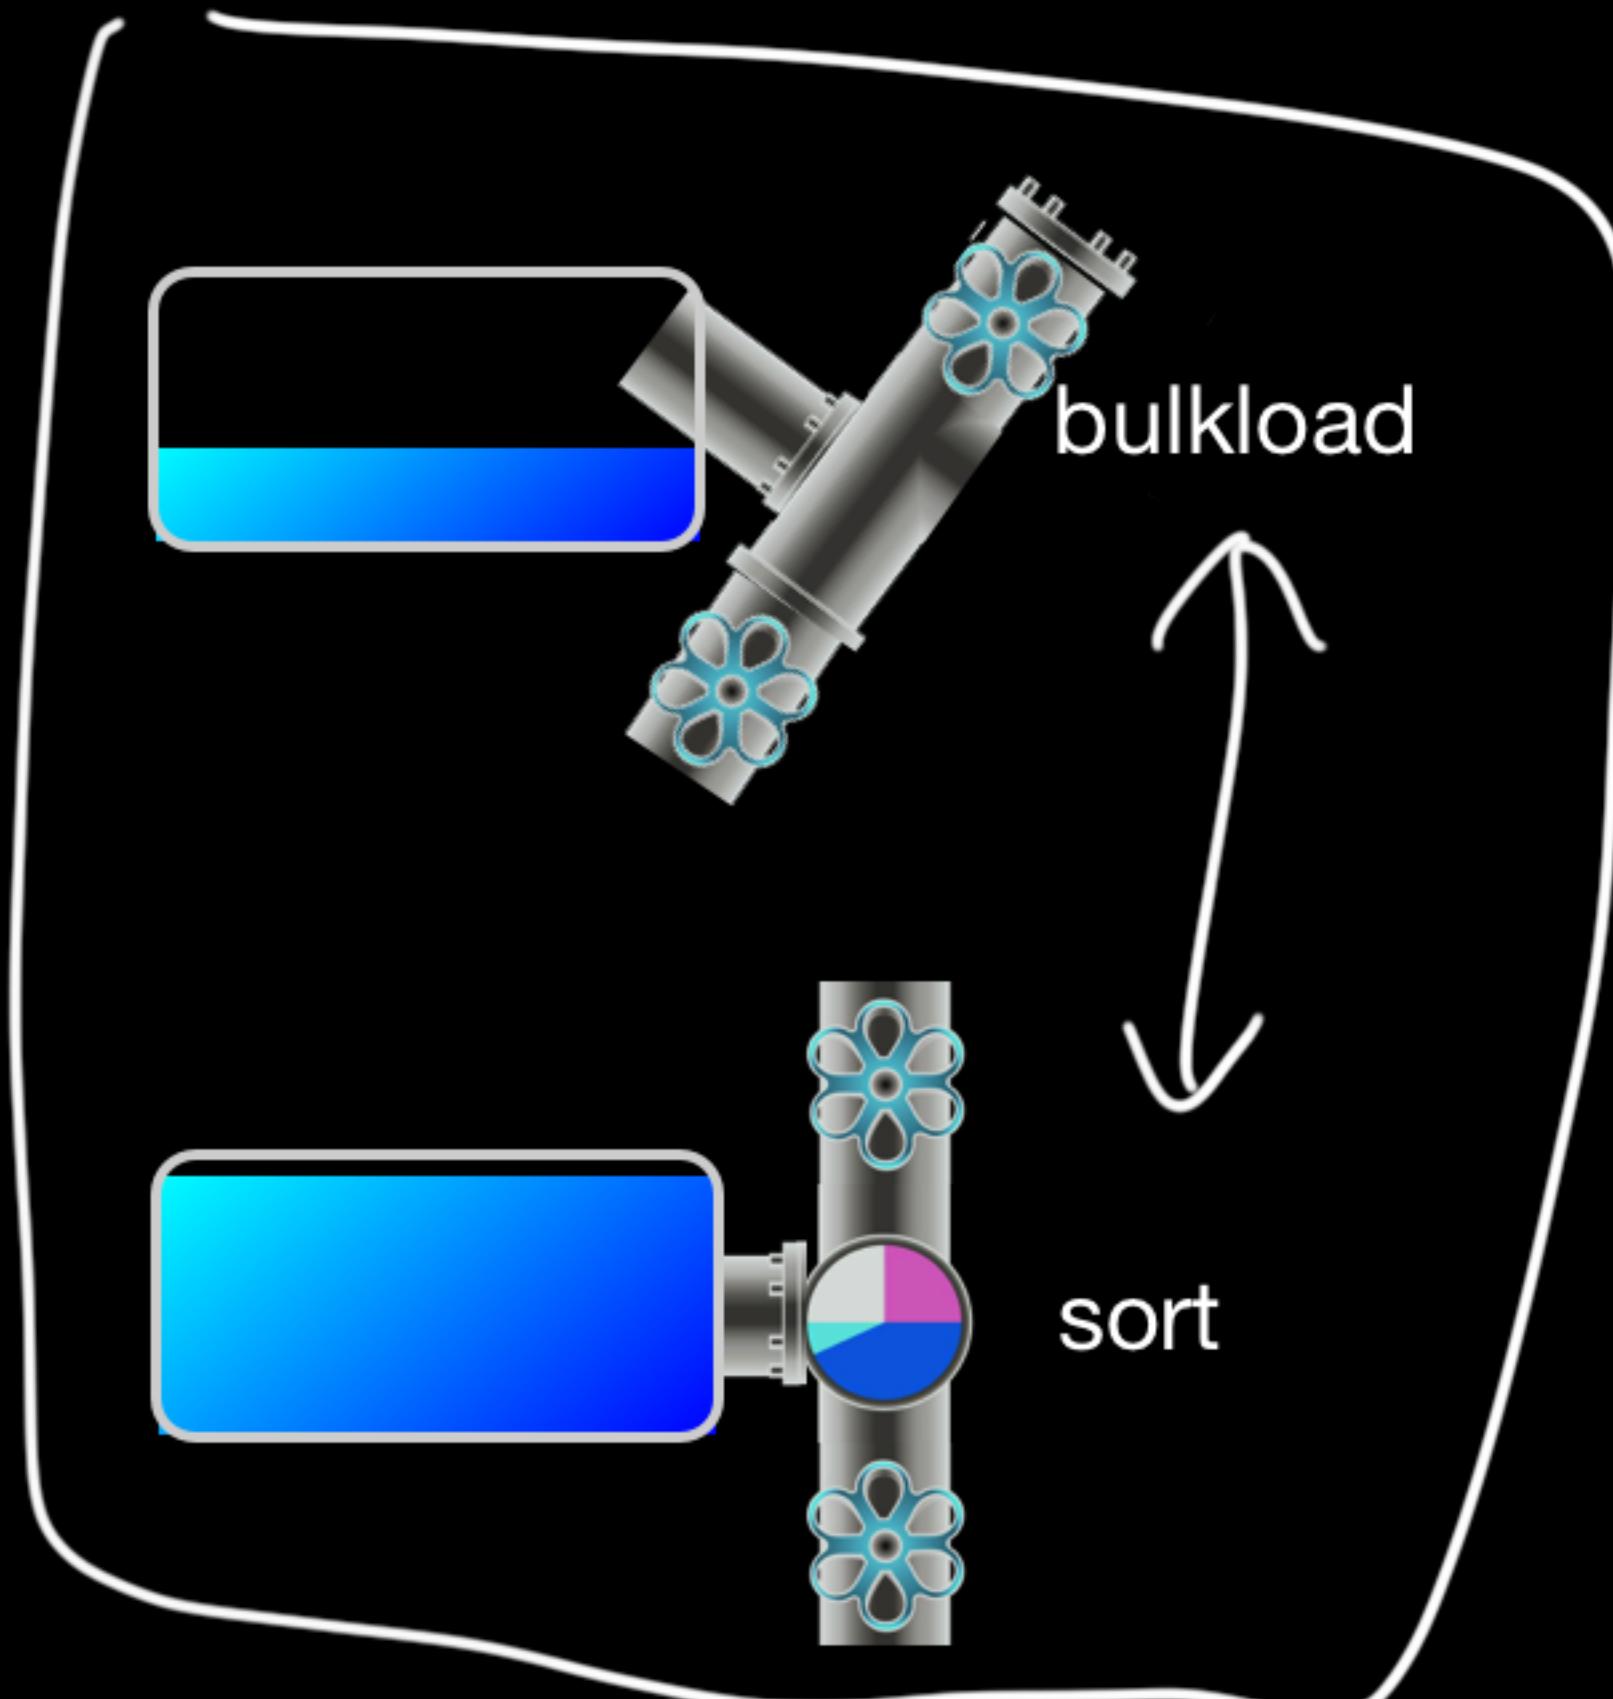
Non-Blocking Algorithms



INLJ/SHJ
(with bulkload)



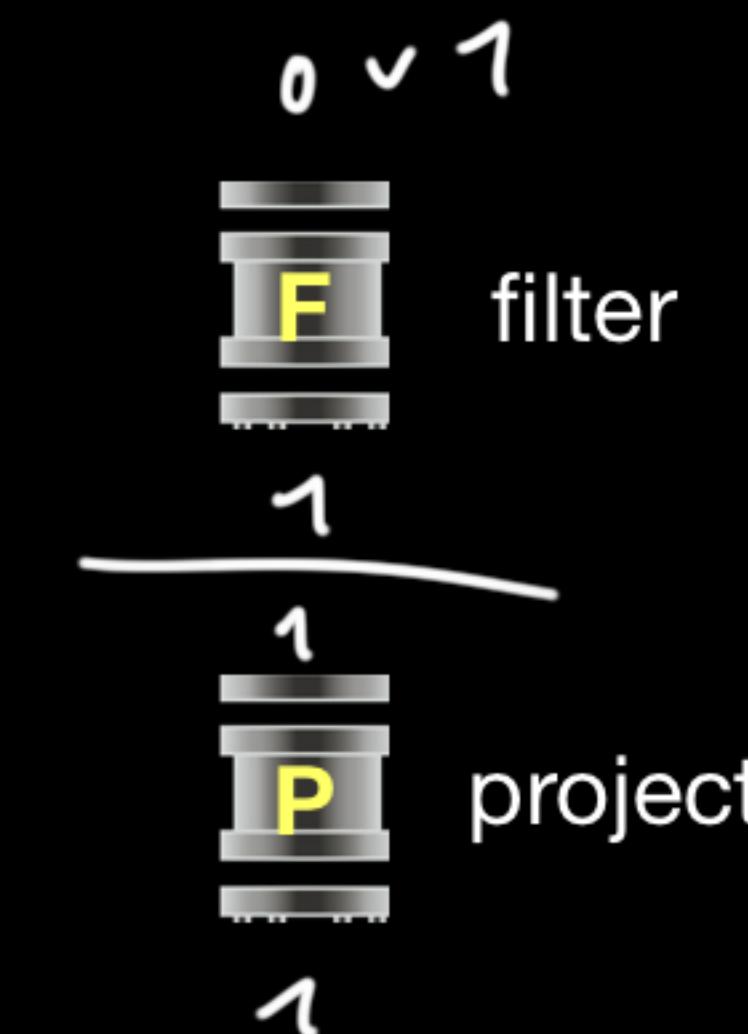
INLJ/SHJ
(without bulkload)
- Index Probe



INLJ / SHJ
(with bulkload)

= bulkload

+ INLJ / SHJ
(without bulkload) ← index probe



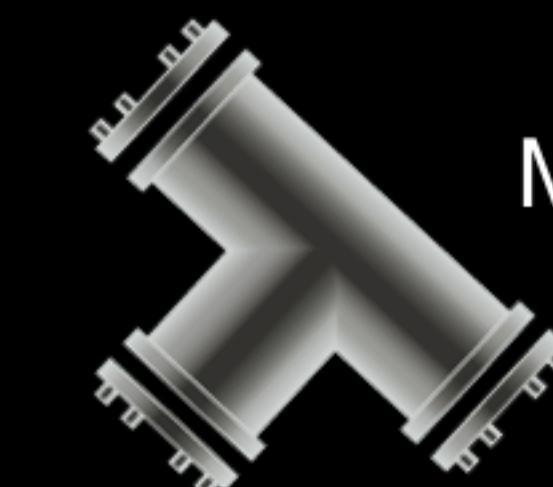
filter

project

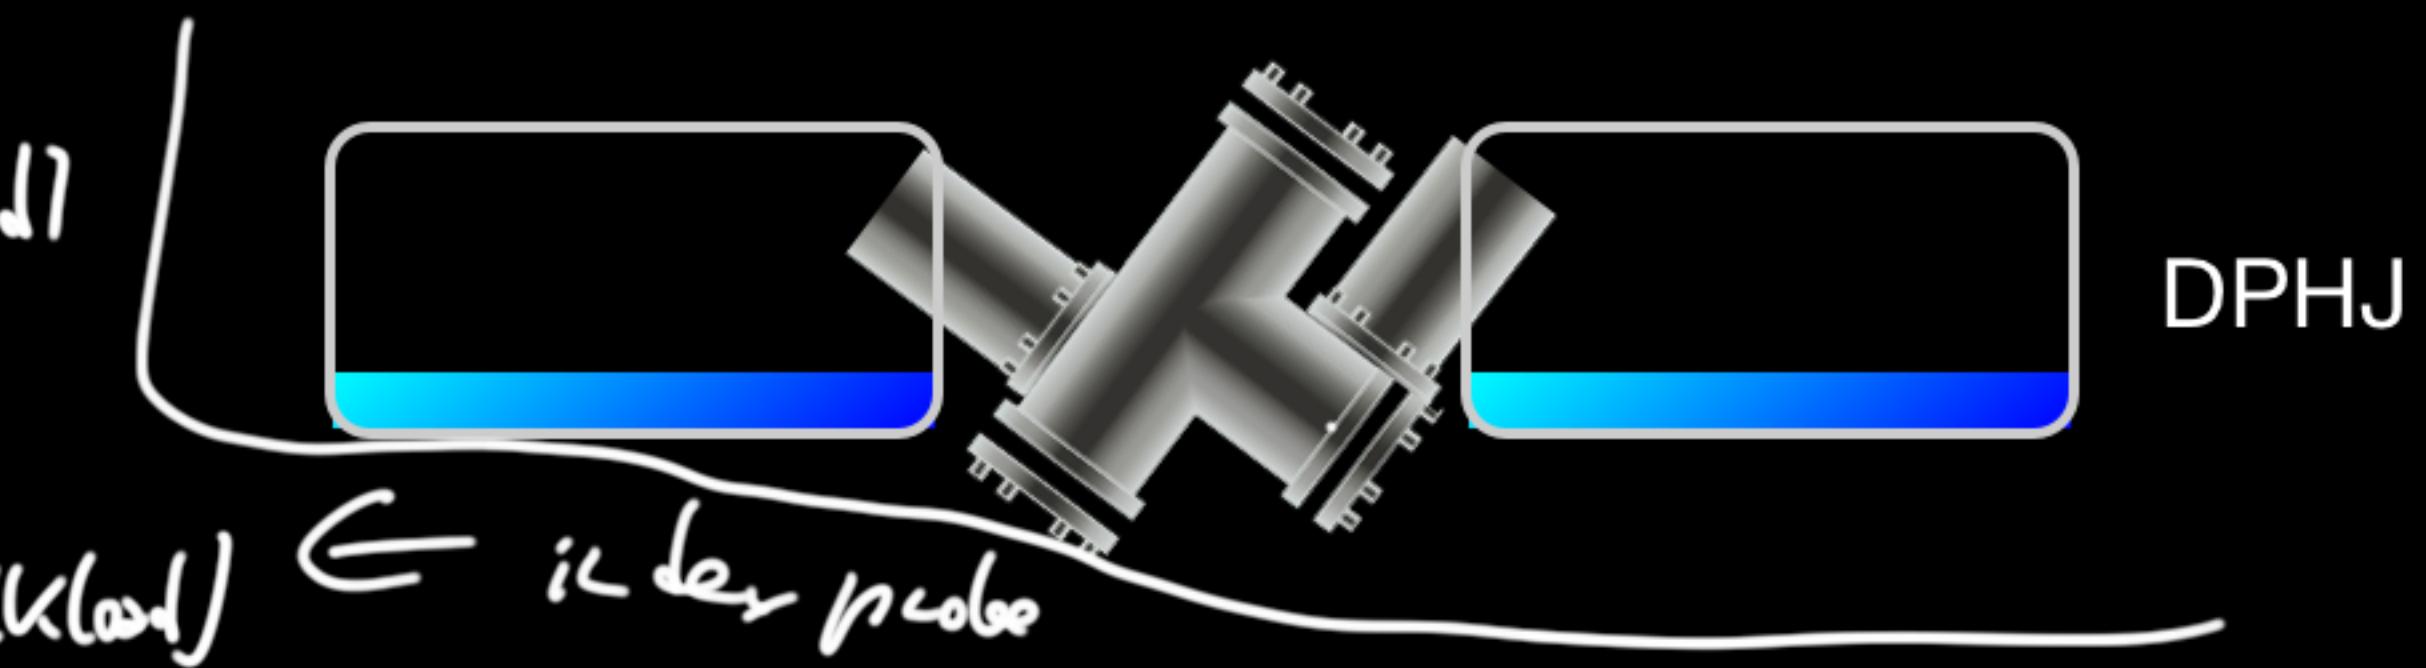


Merge

here:
 $F=2$
 $F \geq 2$



MJ



DPHJ

Copyrights and Credits

© iStock.com:

Horned_Rat