

A graphic of two long rows of server racks, one on the left and one on the right, receding towards a central perspective point. The racks are dark grey with many drive bays visible. They sit on a bright blue rectangular base.

SQL Beyond Structure: Text, Documents, and Key-Value Pairs

Daniel Tahara, Yale University and Hadapt
Daniel Abadi, Yale University and Hadapt

Important Trends

- Dramatic increase in utilization of “NoSQL” databases such as MongoDB, CouchDB, Riak, Cassandra, Hbase, etc.
- Key value and JSON are increasingly popular ways to represent data
 - See MongoDB’s \$150 million “largest database funding ever”
- Performing analysis inside NoSQL databases not a major use-case
 - Poor analytics primitives
 - Poor integration with other data sources (data silos)
 - Poor scan performance
 - Poor integration with analytical tools

A Motivating Example

Relational Datastore

userID	name	email	preferred
1	Walter White	heisenberg@bb.com	false
2	Hank Schrader	mineral@rock.com	true
3	Jesse Pinkman	pink@thenewblack.ca	true
4	Gustavo Fring	gus@lospollos.com	true

Text Datastore

userID	Email
1	What is your return policy?
2	Do you have bulk discounts?
4	I can't figure out your Website. You should fire your designer!

What if you want to send an email to all customers who indicate a potential desire to buy in bulk, with a deal on items they looked at?

JSON Datastore

JSON Documents
{"userID": 1, "action": "view", "productName": "door", "price": 12.50, "color": "green", "tags": ["home", "green"]}
{"userID": 3, "action": "purchase", "price": 57.12}
{"userID": 2, "action": "view", "productName": "window", "color": "orange", "price": 52.50, "tags": ["home", "orange"]}

Option 1: User Code

Relational Datastore

	name	email	preferred
1	Walter White	heisenberg@bb.com	false
2	Hank Schrader	mineral@rock.com	true
3	Jesse Pinkman	pink@thenewblack.ca	true
4	Gustavo Fring	gus@lospollos.com	true

Text Datastore

userID	Email
1	What is your return policy?
2	Do you have bulk discounts?
4	I can't figure out your Website. You should fire your designer!

Extract data from each relevant store, perform join and further query analysis in user code.

```
#!/bin/sh  
  
if ['some  
    condition']  
then  
    'do  
        something'  
fi  
  
'more  
complicated  
code'
```

JSON Datastore

JSON Documents
{"userID": 1, "action": "view", "productName": "door", "price": 12.50, "color": "green", "tags": ["home", "green"]}
{"userID": 3, "action": "purchase", "price": 57.12}
{"userID": 2, "action": "view", "productName": "window", "color": "orange", "price": 52.50, "tags": ["home", "orange"]}

Option 2: ETL

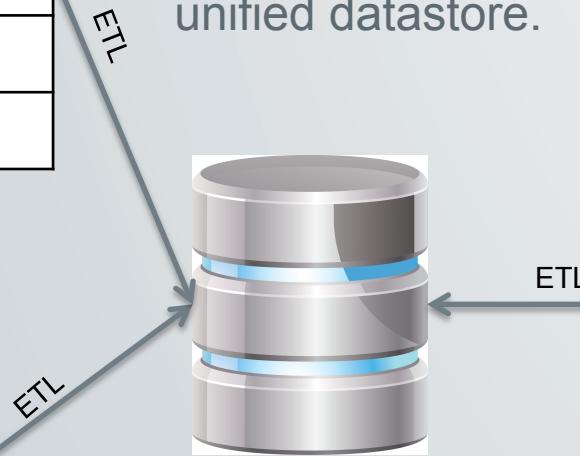
Relational Datastore

userID	name	email	preferred
1	Walter White	heisenberg@bb.com	false
2	Hank Schrader	mineral@rock.com	true
3	Jesse Pinkman	pink@thenewblack.ca	true
4	Gustavo Fring	gus@lospollos.com	true

Text Datastore

userID	Email
1	What is your return policy?
2	Do you have bulk discounts?
4	I can't figure out your Website. You should fire your designer!

Use ETL to combine data from separate backends into one unified datastore.



JSON Datastore

JSON Documents
{“userID”: 1, “action”：“view”, “productName”: “door”, “price”: 12.50, “color”：“green”, “tags”: [“home”, “green”]}
{“userID”: 3, “action”：“purchase”, “price”: 57.12}
{“userID”: 2, “action”：“view”, “productName”: “window”, “color”：“orange”, “price”: 52.50, “tags”: [“home”, “orange”]}

A “Unified” Datastore

userID	name	emailAddr	preferred	json	emails
1	Walter White	heisenberg@bb.com	false	{"eventID": 1, "action":"view", "productName": "door", "price": 12.50,"color": "green", "tags": ["home", "green"]}	What is your return policy?
2	Hank Schrader	mineral@rock.com	true	{"eventID": 2, "action": "view", "productName": "window", "color": "orange", "price": 52.50, "tags": ["home", "orange"]}	Do you have bulk discounts?
3	Jesse Pinkman	pink@thenewblack.ca	true	{"eventID": 3, "action": "purchase", "price": 57.12}	
4	Gustavo Fring	gus@lospollos.com	true		I can't figure out your Website. You should fire your designer!

```
SELECT name,  
       emailAddr,  
       regex_value(json, “productName”: \s*(“[^”]*?”), 1, ‘i’)  
FROM T  
WHERE regex_match(json, “action”: \s*”view”, ‘i’)  
      AND emails like ‘%bulk%’;
```

A Better Approach?

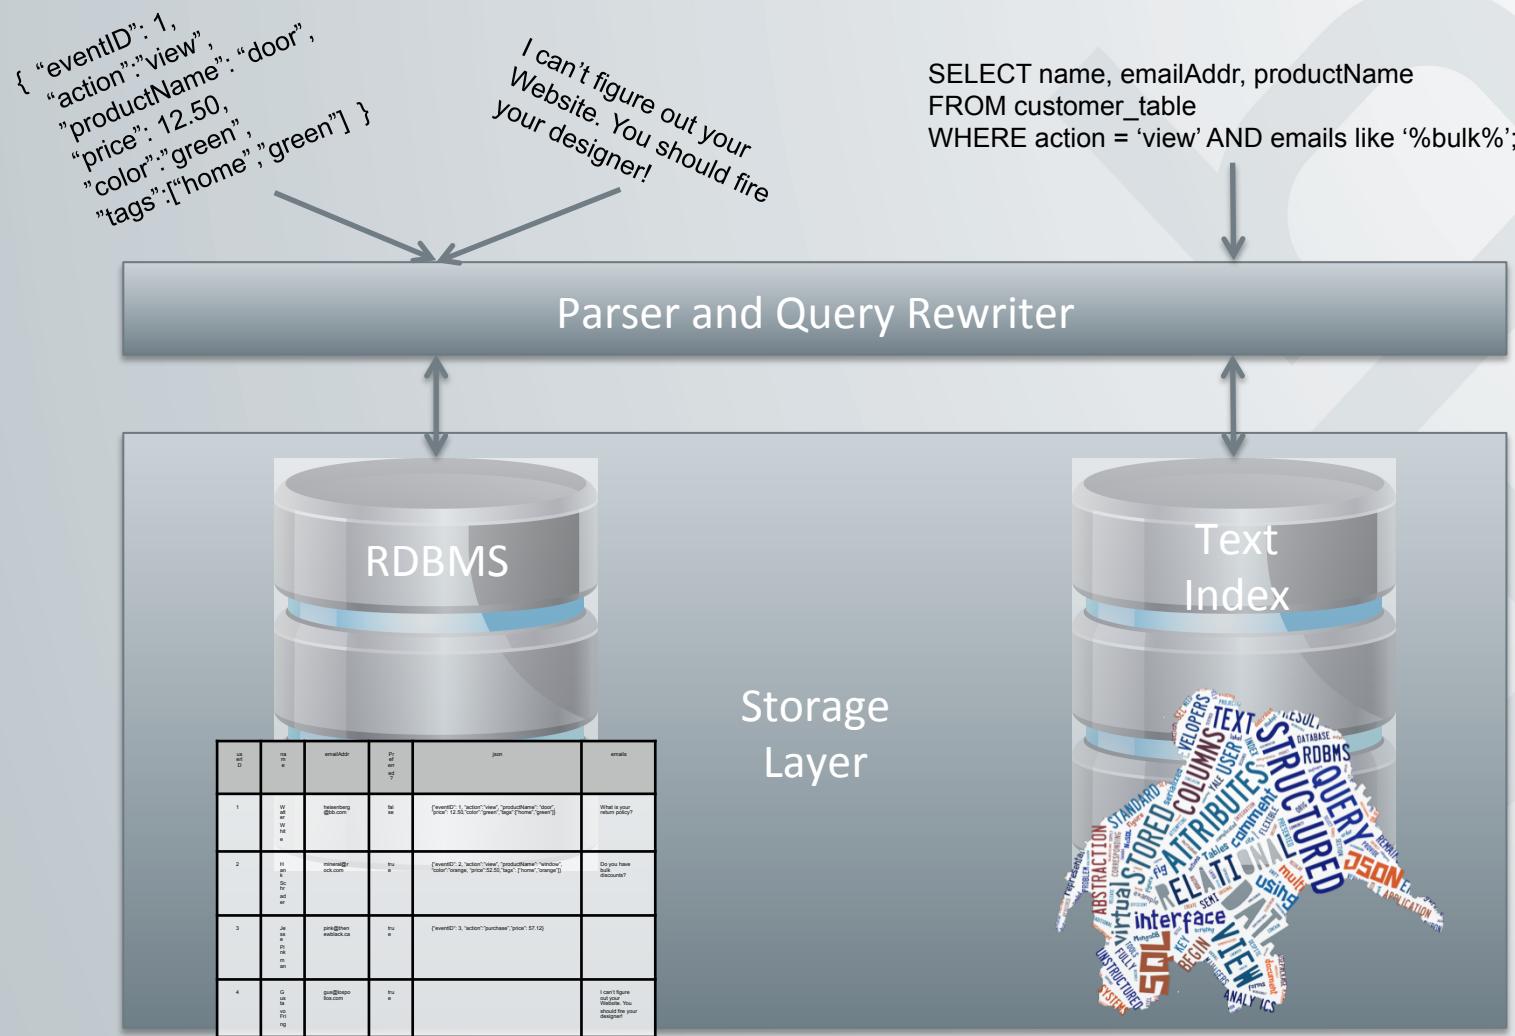
userID	name	emailAddr	preferred	product Name	price	tags	color	action	emails
1	Walter White	heisenberg@bb.com	no	door	12.50	home green	green	view	What is your return policy?
2	Hank Schrader	mineral@rock.com	yes	window	52.50	home orange	orange	view	Do you have bulk discounts?
3	Jesse Pinkman	pink@thenewblack.ca	yes		57.12			purchase	
4	Gustavo Fring	gus@lospollos.com	yes						I can't figure out your Website. You should fire your designer!

```
SELECT name,  
       emailAddr,  
       productName  
  FROM T  
 WHERE action = 'view'  
   AND emails like '%bulk%';
```

Goal: Multistructured Tables

- Store data with various levels of structure side-by-side in an RDBMS
- Remove the ETL step in the processing pipeline
- Maintain a standard SQL interface to the data - allow elements to be manipulated as relational attributes irrespective of original format
- Bonus: Support for non-relational operations such as full text search

Implementation: Overview



Implementation: Details

- System does not store data according to the virtual, “universal” view
- Store key-value data in a serialized, key-value format in order balance competing needs:
 - Limit table width
 - Speed up attribute extraction
 - Identify which attributes exist
- Inverted indexes over all keys and text columns
- Query rewriter transforms:
 - references to serialized, virtual columns into key extractions
 - predicates over virtual columns into queries to the text index, applied as a filter set on the attribute ids

Flexible Schema

json
{"userID": 4, "action":"view", "productName": "door", "price": 12.50,"color":"green","tags": ["home","green"], "size":"standard"}

userID	name	emailAddr	Preferred ?	product Name	price	tags	color	action	size	emails
1	Walter White	heisenberg@bb.com	no	door	12.50	home green	green	view		What is your return policy?
2	Hank Schrader	mineral@rock.com	yes	window	52.50	home orange	orange	view		Do you have bulk discounts?
3	Jesse Pinkman	pink@thenewblack.ca	yes		57.12			purchase		
4	Gustavo Fring	gus@lospollos.com	yes						standard	I can't figure out your Website. You should fire your designer!

Challenges with the Current Design

- Relying on a text index can be slow if the result set is large
- Can't utilize B-tree indexes
- RDBMS optimizer often produces suboptimal query plans

Dynamic Column Creation/Deletion

- Specifications:
 - Create physical columns for most commonly used or frequently appearing keys.
 - Ongoing process. Columns can be created or deleted at any point.
- Benefits:
 - can create indices on physical columns
 - optimizer has better statistics on frequently queried data
 - dynamic physical schema allows the system to adapt to the evolving nature of datasets

“Smart” Query Rewrite and UDFs

- Specification:
 - Define an extraction function and a text index query function for use within the RDBMS
 - *Selectively* execute text predicates in the RDBMS or push them out to the text index
- Benefits:
 - More code running within the RDBMS
 - Take into account system knowledge of the data to produce globally optimal execution

Alternatives

- Entity-Attribute-Value
 - “flatten” key-value data types
 - stored in 3-column tables
- PostgreSQL 9.3
 - use JSON text data type
 - extraction operators
- Column Store
 - store in the form of universal relation

Conclusions

- Bevy of options for formatting data
- As long as SQL remains a standard of data analytics, there will be a need for unified solution
- Tighter RDBMS integration correlates with improved performance

Questions?

- Email: daniel.tahara@yale.edu
- Further Resources:
 - Hadapt Schemaless SQL:
 - <http://hadapt.com/schemaless-sql-overview/>
 - Yale Database Group:
 - <http://db.cs.yale.edu/>
 - My Website:
 - <http://danieltahara.com/>