

OpenMP Overview 1

# Getting Up To Speed On OpenMP 4.0

Ruud van der Pas  
ORACLE®  
Distinguished Engineer  
Architecture and Performance  
SPARC Microelectronics, Santa Clara, CA, USA

HPCS 2015  
Amsterdam, The Netherlands  
July 20, 2015

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

OpenMP Overview 2

## Agenda

- The One Hour Elevator Ride
- Using OpenMP
  - Directives
  - Environment Variables
  - Run time Functions
- Additional Functionality
- Performance Tuning Case Studies

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

OpenMP Overview 3

## The One Hour Elevator Ride

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

OpenMP Overview 4

## What Is OpenMP ?

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

ATLAS 3.10.2 errata

math-atlas.sourceforge.net/errata.html#omp1n

When linking ATLAS's testers, I'm getting a bunch of undefined BLAS symbols (eg. `dgemm_`)

.... Building using OpenMP In general, this is a bad idea, since OpenMP tends to much slower than pthreads for ATLAS use .....

..... In ATLAS, thread affinity is the main reason pthreads wins against OpenMP .....

Building using OpenMP In general, this is a bad idea, since OpenMP tends to much slower than pthreads for ATLAS use. However, if your own application uses OpenMP, sometimes pthread usage can slow down your own threads, making it worthwhile to damage ATLAS performance in order to improve your OpenMP performance. In ATLAS, thread affinity is the main reason pthreads wins against OpenMP, and in OS-X (and probably FreeBSD), which don't support real affinity, are the platforms where it makes sense to use OpenMP regardless. To force ATLAS to use OpenMP rather than pthreads, you must add the following flags to your configure line:

`-fopenmp -fopenmp -fopenmp`

<http://math-atlas.sourceforge.net/errata.html#omp1n>

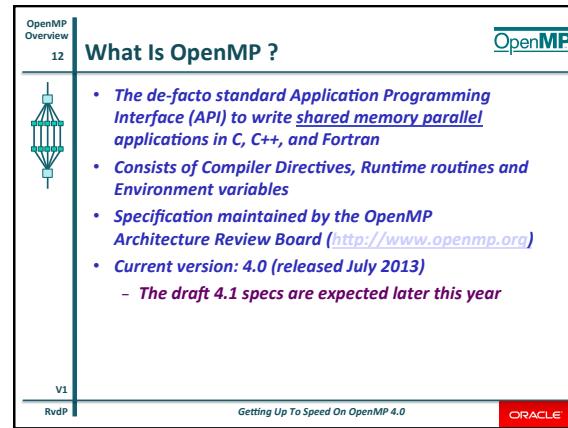
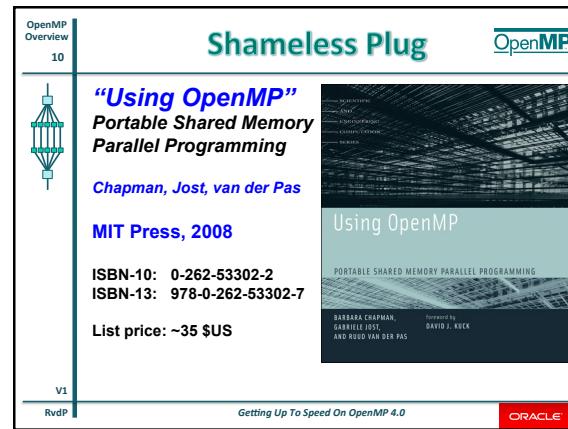
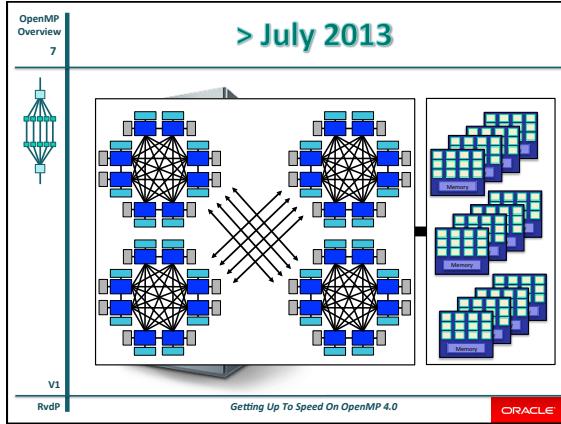
Help and architectural defaults for 64-bit Windows

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

OpenMP Overview 6

## 1997 - 2013

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE



**OpenMP Overview** 13

**Members**

Permanent Members:

- AMD (Greg Stier)
- ARM (Chris Adde)
- Cray (Luis DeFla)
- Fujitsu (Eiji Yamada)
- HP (Sujay Saran)
- IBM (Mike Lai)
- Intel (Xommin Thakur)
- Micron (Kirby Collins)
- NEC (Kazuhiro Kusano)
- NVIDIA (Jeff Laudat)
- Oracle (Steve Barton)
- Red Hat (Matt Newcome)
- Texas Instruments (Andy Fritsch)

Auxiliary Members of the ARB:

- ANL (Kalyan Kumar)
- ASCLL (Michael R. de Supinski)
- BSC (Mariano González)
- cOMPunity (Barbara Chapman/Yonghong Yan)
- EPCC (Mark Bull)
- LANL (David Montoya)
- LBNL (Alice Koniges/Helen He)
- NASA (John Minn)
- ORNL (Oscar Hernandez)
- RWTH Aachen University (Dieter an Mey)
- SNL-Sandia National Lab (Stephen Olivier)
- Texas Advanced Computing Center (Kent Mittlefeld)
- University of Houston (Barbara Chapman/Deepak Easempati)

V1 RvdP

Getting Up To Speed On OpenMP 4.0

ORACLE

**OpenMP Overview** 14

**When To Consider OpenMP ?**

**OpenMP**

- Using an automatically parallelizing compiler
  - The compiler can not find the parallelism
  - The data dependence analysis is not able to determine whether it is safe to parallelize
  - The granularity is not high enough, but you know better
  - Compiler lacks the information to parallelize at a higher level
- Not using an automatically parallelizing compiler
  - No other choice than to parallelize yourself
  - Compilers can still help (e.g. auto-scoping, warnings, etc)

V1 RvdP

Getting Up To Speed On OpenMP 4.0

ORACLE

**OpenMP Overview** 15

**The Advantages Of OpenMP**

**OpenMP**

- De-facto and mature standard
  - Strong focus on backward compatibility
  - Continues to evolve and stays up to date
- Good performance and scalability
  - If you do it right ....
- Portability
  - Supported by a large number of compilers
- Requires little programming effort
  - But, ....
- Allows the program to be parallelized incrementally
  - But, ....

V1 RvdP

Getting Up To Speed On OpenMP 4.0

ORACLE

**OpenMP Overview** 16

**OpenMP And Multicore**

**OpenMP**

**OpenMP is ideally suited for multicore architectures**

**Memory and threading model map naturally**

**Lightweight**

**Mature**

**Widely available and used**

V1 RvdP

Getting Up To Speed On OpenMP 4.0

ORACLE

**OpenMP Overview** 17

**Directive Format**

**OpenMP**

- ❑ C: directives are case sensitive
  - Syntax: #pragma omp directive [clause [clause] ...]
- ❑ Continuation: use \ in pragma
- ❑ Conditional compilation: \_OPENMP macro is set
- ❑ Fortran: directives are case insensitive
  - Syntax: sentinel directive [clause [,] clause]...
  - The sentinel is one of the following:
    - !\$OMP or C\$OMP or \$OMP (fixed format)
    - !\$OMP (free format)
- ❑ Continuation: follows the language syntax
- ❑ Conditional compilation: !\$ or C\$ -> 2 spaces

V1 RvdP

Getting Up To Speed On OpenMP 4.0

ORACLE

**OpenMP Overview** 18

**Defining Parallelism In OpenMP**

**OpenMP**

**A parallel region is a block of code executed by all threads in the team**

```
#pragma omp parallel [clause[, clause] ...]
{
  "this code is executed in parallel"
} // End of parallel section (note: implied barrier)
```

```
!$omp parallel [clause[, clause] ...]
  "this code is executed in parallel"
!$omp end parallel (note: implied barrier)
```

V1 RvdP

Getting Up To Speed On OpenMP 4.0

ORACLE

**OpenMP Overview** 19

### Parallel Region - An Example/1

```
#include <stdlib.h>
#include <stdio.h>

int main(int argc, char *argv[]) {
    printf("Hello World\n");
    return(0);
}
```

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**OpenMP Overview** 20

### Parallel Region - An Example/2

```
#include <stdlib.h>
#include <stdio.h>

int main(int argc, char *argv[]) {
    #pragma omp parallel
    {
        printf("Hello World\n");
    } // End of parallel region
    return(0);
}
```

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**OpenMP Overview** 21

### Parallel Region - An Example/3

```
$ cc -fopenmp -fast hello.c
$ export OMP_NUM_THREADS=2
$ ./a.out
Hello World
Hello World
$ export OMP_NUM_THREADS=4
$ ./a.out
Hello World
Hello World
Hello World
Hello World
Hello World
$
```

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**OpenMP Overview** 22

### The If Clause

**if (scalar expression)**

- ✓ Only execute in parallel if expression evaluates to true
- ✓ Otherwise, execute serially

```
#pragma omp parallel if (n > some_threshold) \
    shared(n,x,y) private(i)
{
    #pragma omp for
    for (i=0; i<n; i++)
        x[i] += y[i];
} /*-- End of parallel region --*/
```

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**OpenMP Overview** 23

### A More Elaborate Example

```
#pragma omp parallel if (n>limit) default(none) \
    shared(n,m,a,b,c,x,y,z) private(f,i,scale)
{
    f = 1.0;
    #pragma omp for nowait
    for (i=0; i<n; i++)
        z[i] = x[i] + y[i];
    #pragma omp for nowait
    for (i=0; i<n; i++)
        a[i] = b[i] + c[i];
    ...
    #pragma omp barrier
    scale = sum(a,0,m) + sum(z,0,n) + f;
}
/*-- End of parallel region --*/
```

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**OpenMP Overview** 24

### Nested Parallelism

Master Thread

3-way parallel

9-way parallel

3-way parallel

Outer parallel region

Nested parallel region

Outer parallel region

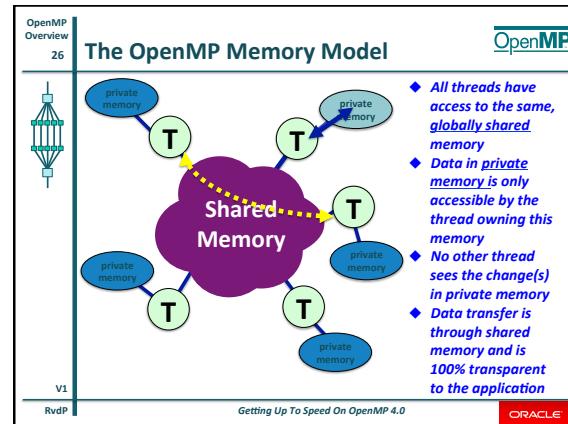
Note: Nesting level can be arbitrarily deep

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**Components Of OpenMP**

Directives	Environment variables	Runtime functions
<ul style="list-style-type: none"> <li>Worksharing</li> <li>Tasking</li> <li>Affinity</li> <li>Accelerators</li> <li>Cancellation</li> <li>Synchronization</li> </ul>	<ul style="list-style-type: none"> <li>Thread Settings</li> <li>Thread Controls</li> <li>Work Scheduling</li> <li>Affinity</li> <li>Accelerators</li> <li>Cancellation</li> <li>Operational</li> </ul>	<ul style="list-style-type: none"> <li>Thread Management</li> <li>Work Scheduling</li> <li>Tasking</li> <li>Affinity</li> <li>Accelerators</li> <li>Cancellation</li> <li>Locking</li> </ul>

RvdP | Getting Up To Speed On OpenMP 4.0 | ORACLE



**Data-Sharing Attributes**

- In an OpenMP program, data needs to be "labeled"
- Essentially there are two basic types:
  - Shared - There is only one instance of the data
    - Threads can read and write the data simultaneously unless protected through a specific construct
    - All changes made are visible to all threads
      - But not necessarily immediately, unless enforced .....
  - Private - Each thread has a copy of the data
    - No other thread can access this data
    - Changes only visible to the thread owning the data

V1 | RvdP | Getting Up To Speed On OpenMP 4.0 | ORACLE

**The Private And Shared Clauses**

**private (list)**

- No storage association with original object
- All references are to the local object
- Values are undefined on entry and exit

**shared (list)**

- Data is accessible by all threads in the team
- All threads access the same address space

V1 | RvdP | Getting Up To Speed On OpenMP 4.0 | ORACLE

**About Storage Association**

- Private variables are undefined on entry and exit of the parallel region
- A private variable within a parallel region has no storage association with the same variable outside of the region
- Use the **firstprivate** and **lastprivate** clauses to override this behavior
- We illustrate these concepts with an example

V1 | RvdP | Getting Up To Speed On OpenMP 4.0 | ORACLE

**The Firstprivate And Lastprivate Clauses**

**firstprivate (list)**

- All variables in the list are initialized with the value the original object had before entering the parallel construct

**lastprivate (list)**

- The thread that executes the sequentially last iteration or section updates the value of the objects in the list

V1 | RvdP | Getting Up To Speed On OpenMP 4.0 | ORACLE

**OpenMP Overview** 31

### Example Firstprivate

```
n = 2; indx = 4;
#pragma omp parallel default(private) private(i,TID)\n    firstprivate(indx) shared(n,a)
{ TID = omp_get_thread_num();
  indx = indx + n*TID;
  for(i=indx; i<indx+n; i++)
    a[i] = TID + 1;
} /*-- End of parallel region --*/
```

	0	1	2	3	4	5	6	7	8	9
index	0	1	2	3	4	5	6	7	8	9
value	1	1	2	2	3	3				

TID = 0      TID = 1      TID = 2

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**OpenMP Overview** 32

### Example Lastprivate

```
#pragma omp parallel for lastprivate(a)
for (int i=0; i<n; i++)
{
  .....
  a = i + 1;
  .....
} // End of parallel region
```

b = 2 \* a; // value of b is 2\*n

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**OpenMP Overview** 33

### The Default Clause

<b>C/C++</b>	<b>Fortran</b>
<b>default ( none   shared )</b>	
<b>default (none   shared   private   threadprivate )</b>	
<b>none</b>	
✓ No implicit defaults; have to scope all variables explicitly	
<b>shared</b>	
✓ All variables are shared	
✓ The default in absence of an explicit "default" clause	
<b>private</b>	
✓ All variables are private to the thread	
✓ Includes common block data, unless THREADPRIVATE	
<b>firstprivate</b>	
✓ All variables are private to the thread; pre-initialized	

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**OpenMP Overview** 34

### Data Scoping – Gotcha's

- Need to get this right
  - Part of the learning curve
- Private data is undefined on entry and exit
  - Can use firstprivate and lastprivate to address this
- Each thread has its own temporary view on the data
  - Applicable to shared data only
  - Means different threads may temporarily not see the same value for the same variable ...
  - Let me explain

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**OpenMP Overview** 35

### The Flush Directive

Thread A

X = 0  
.....  
X = 1

Thread B

while (X == 0)  
{  
 "wait"  
}

If shared variable X is kept within a register, the modification may not be made visible to the other thread(s)

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**OpenMP Overview** 36

### Example Of The Flush Directive

```
void wait_read(int i)
{
  #pragma omp flush
  while ( execution_state[i] != READ_FINISHED )
  {
    system("sleep 1");
    #pragma omp flush
  }
} /*-- End of wait_read --*/
```

This fragment is from the source code for the pipeline example from "Using OpenMP" (overlaps I/O with processing)

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

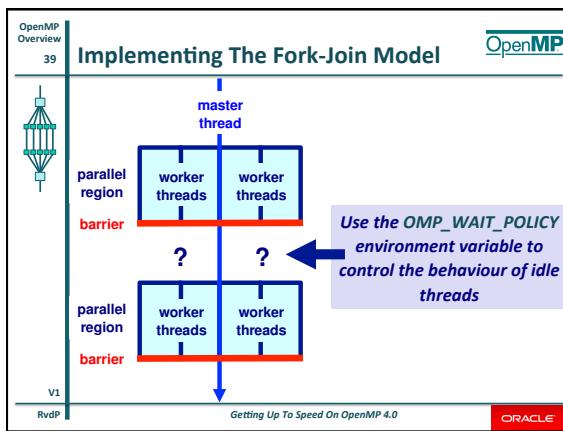
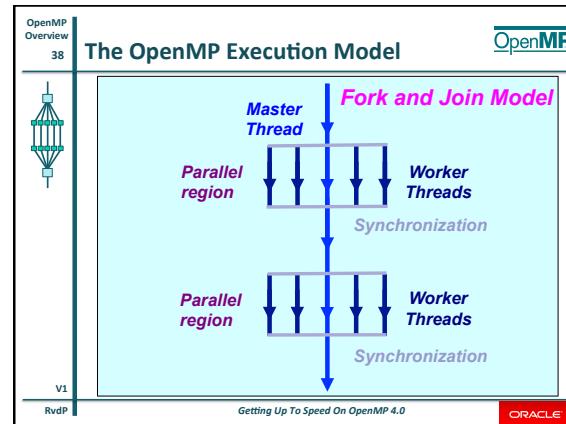
**About The Flush Construct**

- Do not use the flush directive with a list
  - Could give very subtle interactions with compilers
  - If you insist on still doing so, be prepared to face the OpenMP language lawyers in court
- The flush is included in many constructs
  - A Good Thing - This is your safety net
  - Check the specifications for the specific constructs

```
#pragma omp flush [(list)]
!$omp flush [(list)]
```

**Don't use the list**

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE



**The Barrier/1**

Suppose we run each of these two loops in parallel over i:

```
for (i=0; i < n; i++)
  a[i] = b[i] + c[i];
```

```
for (i=0; i < n; i++)
  d[i] = a[i] + b[i];
```

This may give us a wrong answer (one day)

**Why ?**

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**The Barrier/2**

We need to have updated all of a[] first, before using a[\*]

```
for (i=0; i < n; i++)
  a[i] = b[i] + c[i];
```

wait!

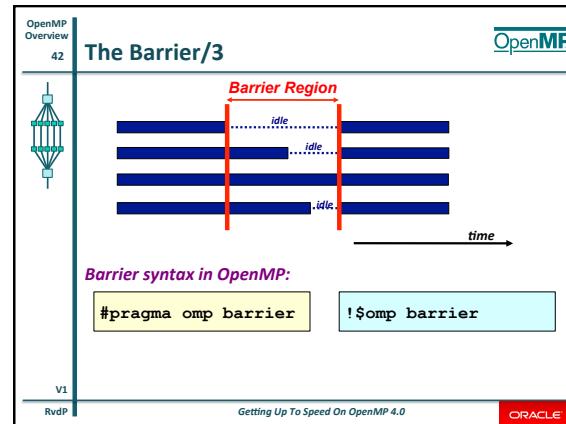
```
for (i=0; i < n; i++)
  d[i] = a[i] + b[i];
```

barrier

All threads wait at the barrier point and only continue when all threads have reached the barrier point

*\*) If there is the guarantee that the mapping of iterations onto threads is identical for both loops, there will not be a data race in this case*

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE



**OpenMP Overview** 43

## When To Use Barriers ?

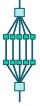


- In case data is updated asynchronously and data integrity is at risk
- Examples:
  - Between parts in the code that read and write the same section of memory
  - After one timestep/iteration in a solver
- Unfortunately, barriers tend to be expensive and also may not scale to a large number of processors
- Therefore, use them with care

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**OpenMP Overview** 44

## The OpenMP Barrier



- Several constructs have an implied barrier
  - This is another safety net (has implied flush by the way)
- In some cases, the implied barrier can be left out through the nowait clause
- This can help fine tuning the application
  - But you'd better know what you're doing
- The explicit barrier comes in quite handy then

#pragma omp barrier

!\$omp barrier

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**OpenMP Overview** 45

## The Nowait Clause



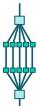
- To minimize synchronization, some directives support the optional nowait clause
  - If present, threads do not synchronize/wait at the end of that particular construct
- In C, it is one of the clauses on the pragma
- In Fortran, it is appended at the closing part of the construct

```
#pragma omp for nowait !$omp do
{
    :
}
    :
    :
    !$omp end do nowait
```

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**OpenMP Overview** 46

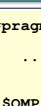
## OpenMP Directives (An Important Subset)




V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**OpenMP Overview** 47

## The Worksharing Constructs



```
#pragma omp for
{
    :
}
    :
    :
    !$OMP DO
    :
    :
    !$OMP END DO
```

```
#pragma omp sections
{
    :
}
    :
    :
    !$OMP SECTIONS
    :
    :
    !$OMP END SECTIONS
```

```
#pragma omp single
{
    :
}
    :
    :
    !$OMP SINGLE
    :
    :
    !$OMP END SINGLE
```

- The work is distributed over the threads
- Must be enclosed in a parallel region
- Must be encountered by all threads in the team, or none at all
- No implied barrier on entry; implied barrier on exit (unless the nowait clause has been specified)
- A work-sharing construct does not launch any new threads

RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**OpenMP Overview** 48

## The Fortran Workshare Construct



Fortran has a fourth worksharing construct:

```
!$OMP WORKSHARE
<array syntax>
!$OMP END WORKSHARE [NOWAIT]
```

Example:

```
!$OMP WORKSHARE
A(1:M) = A(1:M) + B(1:M)
!$OMP END WORKSHARE NOWAIT
```

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**The Omp-For/Do Directive**

```
#pragma omp for [clauses]
for (....)
{
    <code-block>
}

!$omp do [clauses]
do ...
    <code-block>
end do
!$omp end do[nowait]
```

*The iterations of the loop are distributed over the threads*

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**The Omp-For Directive – An Example**

```
#pragma omp parallel default(None) \
    shared(n,a,b,c,d) private(i)
{
    #pragma omp for nowait
    for (i=0; i<n-1; i++)
        b[i] = (a[i] + a[i+1])/2;
    #pragma omp for nowait
    for (i=0; i<n; i++)
        d[i] = 1.0/c[i];
} /*-- End of parallel region --*/ \
(implied barrier)
```

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**The Reduction Clause - Example**

```
sum = 0.0
 !$omp parallel default(None) &
 !$omp shared(n,x) private(i)
 !$omp do reduction (+:sum)
 do i = 1, n
     sum = sum + x(i)
 end do
 !$omp end do
 !$omp end parallel
 print *,sum
```

Variable SUM is a shared variable

- ✓ Care needs to be taken when updating shared variable SUM
- ✓ With the reduction clause, the OpenMP compiler generates code such that a race condition is avoided

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**The Reduction Clause**

C/C++ Fortran

- ✓ Reduction variable(s) must be shared variables
- ✓ A reduction is defined as:

<b>Fortran</b> $x = x \text{ operator expr}$ $x = \text{expr operator } x$ $x = \text{intrinsic } (x, \text{expr\_list})$ $x = \text{intrinsic } (\text{expr\_list}, x)$	<b>C/C++</b> $x = \text{operator expr}$ $x = \text{expr operator } x$ $x++, ++x, x--, --x$ $x <\text{binop}> = \text{expr}$ "min" and "max" intrinsic
--	--

- ✓ Note that the value of a reduction variable is undefined from the moment the first thread reaches the clause till the operation has completed
- ✓ The reduction can be hidden in a function call

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**User Defined Reductions (UDR)**

- In addition to the pre-defined reduction operators supported, one can now write a user specific reduction operator
  - Combines ease of use of the reduction clause with application specific reductions
- Restrictions apply (please read the fine print on this)

```
#pragma omp declare reduction ....
!$omp declare reduction ....
```

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**Loop Collapse**

- Allows parallelization of perfectly nested loops without using nested parallelism
- The collapse clause on for/do loop indicates how many loops should be collapsed
- Compiler forms a single loop and then parallelizes it

```
!$omp parallel do collapse(2) ...
    do i = il, iu, is
        do j = jl, ju, js
            do k = kl, ku, ks
                ....
            end do
        end do
    end do
!$omp end parallel do
```

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**SIMD Support/1**

OpenMP Overview 55

**SIMD = Single Instruction Multiple Data**

- Also called 'micro-vectorization'
- A single instruction performs the same operation on multiple data elements in parallel

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**SIMD Support/2**

OpenMP Overview 56

- The **simd** construct can be applied to:
  - A serial loop (only SIMD instructions generated)
  - A parallel loop (two level parallelism)

```
#pragma omp simd [clause(s)]
#pragma omp for simd [clause(s)]
```

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**Parallel Sections**

OpenMP Overview 57

```
#pragma omp sections [clauses]
{
  #pragma omp section
  {...}
  #pragma omp section
  {...}
  ...
}

 !$omp sections [clauses]
 !$omp section
 {...}
 !$omp section
 {...}
 ...
 !$omp end sections [nowait]
```

Individual section blocks are executed in parallel

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**The Sections Directive – An Example**

OpenMP Overview 58

```
#pragma omp parallel default(None)\n    shared(n,a,b,c,d) private(i)\n{\n    #pragma omp sections nowait\n    {\n        #pragma omp section\n        for (i=0; i<n-1; i++)\n            b[i] = (a[i] + a[i+1])/2;\n\n        #pragma omp section\n        for (i=0; i<n; i++)\n            d[i] = 1.0/c[i];\n    } /* End of sections */\n\n} /* End of parallel region */
```

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**Overlap I/O And Processing/1**

OpenMP Overview 59

Input Thread	Processing Thread(s)	Output Thread
0		
1	0	
2	1	0
3	2	1
4	3	2
5	4	3
	5	4
		5

Time ↓

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**Overlap I/O And Processing/2**

OpenMP Overview 60

```
#pragma omp parallel sections\n{\n    #pragma omp section\n    {\n        for (int i=0; i<N; i++) {\n            (void) read_input(i);\n            (void) signal_read(i);\n        }\n    }\n    #pragma omp section\n    {\n        for (int i=0; i<N; i++) {\n            (void) wait_read(i);\n            (void) process_data(i);\n            (void) signal_processed(i);\n        }\n    }\n    #pragma omp section\n    {\n        for (int i=0; i<N; i++) {\n            (void) wait_processed(i);\n            (void) write_output(i);\n        }\n    }\n} /* End of parallel sections */
```

Input Thread

Processing Thread(s)

Output Thread

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**OpenMP Overview 61 Single Processor Region/1**

**Original Code**

```
"read A[0..N-1];  
....
```

**Only one thread executes the single region**

```
#pragma omp parallel \  
    shared (A)  
{  
    ....  
    #pragma omp single nowait  
    {"read A[0..N-1];"}  
    ....  
    #pragma omp barrier  
    "use A"  
}
```

**Parallel Version**

**ideally suited for I/O or initializations**

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**OpenMP Overview 62 Single Processor Region/2**

**single processor region**

**Other threads wait in the barrier here**

time

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**OpenMP Overview 63 The Single Directive**

**Only one thread in the team executes the code enclosed**

```
#pragma omp single [private][firstprivate] \  
    [copyprivate][nowait]  
{  
    <code-block>  
}  
  
!$omp single [private][firstprivate]  
    <code-block>  
!$omp end single [copyprivate][nowait]
```

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**OpenMP Overview 64 Worksharing Shortcuts**

<b>Single PARALLEL loop</b>	<b>Single WORKSHARE loop</b>
<code>#pragma omp parallel #pragma omp for for (...)</code>	<code>#pragma omp parallel for for (...)</code>
<code>!\$omp parallel !\$omp do ... !\$omp end do !\$omp end parallel</code>	<code>!\$omp parallel do ... !\$omp end parallel do</code>
<b>Single PARALLEL sections</b>	<b>Single WORKSHARE sections</b>
<code>#pragma omp parallel #pragma omp sections { ... }</code>	<code>#pragma omp parallel sections { ... }</code>
<b>Single PARALLEL sections</b>	<b>Single WORKSHARE sections</b>
<code>!\$omp parallel !\$omp sections ... !\$omp end sections !\$omp end parallel</code>	<code>!\$omp parallel sections ... !\$omp end parallel sections</code>

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**OpenMP Overview 65 The Critical Region/1**

**If sum is a shared variable, this loop can not be run in parallel by simply using a "#pragma omp for"**

```
for (i=0; i < n; i++){  
    ....  
    sum += a[i];  
    ....  
}  
  
#pragma omp parallel for  
for (i=0; i < n; i++){  
    ....  
    #pragma omp critical  
    {sum += a[i];}  
    ....  
}
```

All threads execute the update, but now only one at a time will do so

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**OpenMP Overview 66 The Critical Region/2**

- Very useful to avoid a race condition, or to perform I/O
  - But the order of execution is still undefined
- Be aware there is a performance cost associated with a critical region

critical region

time

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**The Critical Construct**

OpenMP Overview 67

**Very useful to avoid data races**

```
#pragma omp critical [(name)]
{<code-block>}
```

```
!$omp critical [(name)]
<code-block>
!$omp end critical [(name)]
```

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

OpenMP Overview 68

**The OpenMP Environment Variables**

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**Environment Variables/1**

OpenMP Overview 69

OpenMP Environment Variable	Category
OMP_DISPLAY_ENV	Diagnostics
OMP_NUM_THREADS	Thread Management
OMP_THREAD_LIMIT	Thread Management
OMP_DYNAMIC {true  false}	Thread Management
OMP_NESTED {true  false}	Parallelism
OMP_MAX_ACTIVE_LEVELS	Parallelism
OMP_STACKSIZE "size [b k m g]"	Operational

✓ The names have to be in uppercase; the values are case insensitive  
✓ Be careful when relying on defaults (they are compiler dependent)

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**About The Stack**

OpenMP Overview 70

Variable Alocal is in private memory, managed by the thread owning it, and stored on the so-called stack

```
#omp parallel shared(Aglobal)
{
    (void) myfunc(&Aglobal);
}
```

```
void myfunc(float *Aglobal)
{
    int Alocal;
    .....
}
```

Aglobal

Thread Alocal

Thread Alocal

Thread Alocal

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**Environment Variables/2**

OpenMP Overview 71

OpenMP Environment Variable	Category
OMP_PLACES	Affinity
OMP_PROC_BIND	Affinity
OMP_DEFAULT_DEVICE	Accelerators
OMP_CANCELLATION	Thread Management
OMP_WAIT_POLICY [active   passive]	Thread Management
OMP_SCHEDULE "schedule, [chunk]"	Work Distribution

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

OpenMP Overview 72

**The OpenMP Runtime Functions**

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

OpenMP Overview 73

## OpenMP Runtime Functions

- OpenMP provides a set of runtime functions
- They all start with "omp\_"
- These functions can be used to:
  - Query for a specific feature or setting
    - For example, the current iteration scheduling policy for loops
    - Change the setting
    - For example, to change the iteration scheduling policy for loops
  - A special category consists of the locking functions

C/C++ : Need to include file <omp.h>  
Fortran : Add "use omp\_lib" or include file "omp.lib.h"

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

OpenMP Overview 74

## Runtime Functions/1

Function Name	Category
<code>omp_get_num_threads/omp_set_num_threads</code>	Thread Management
<code>omp_get_max_threads</code>	Thread Management
<code>omp_get_thread_num</code>	Thread Management
<code>omp_get_dynamic/omp_set_dynamic</code>	Thread Management
<code>omp_in_parallel</code>	Parallelism
<code>omp_get_nested/omp_set_nested</code>	Parallelism

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

OpenMP Overview 75

## Runtime Functions/2

Function Name	Category
<code>omp_get_num_procs</code>	Operational
<code>omp_get_wtime</code>	Operational
<code>omp_get_wtick</code>	Operational
<code>omp_init_lock/omp_init_nest_lock</code>	Locking
<code>omp_set_lock/omp_set_nest_lock</code>	Locking
<code>omp_unset_lock/omp_unset_nest_lock</code>	Locking
<code>omp_test_lock/omp_test_nest_lock</code>	Locking
<code>omp_destroy_lock/omp_destroy_nest_lock</code>	Locking

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

OpenMP Overview 76

## Runtime Functions/3

Function Name	Category
<code>omp_get_thread_limit</code>	Thread Management
<code>omp_get_schedule/omp_set_schedule</code>	Work Distribution
<code>omp_get_max_active_levels/omp_set_max_active_levels</code>	Parallelism
<code>omp_get_level</code>	Parallelism
<code>omp_get_active_level</code>	Parallelism
<code>omp_get_ancestor_thread_num</code>	Parallelism
<code>omp_get_team_size</code>	Parallelism

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

OpenMP Overview 77

## Runtime Functions/4

Function Name	Category
<code>omp_in_final</code>	Tasking
<code>omp_get_proc_bind</code>	Thread Affinity
<code>omp_get_cancellation</code>	Cancellation
<code>omp_get_num_teams</code>	Accelerators
<code>omp_get_team_num</code>	Accelerators
<code>omp_get_default_device</code>	Accelerators
<code>omp_set_default_device</code>	Accelerators
<code>omp_get_num_devices</code>	Accelerators
<code>omp_is_initial_device</code>	Accelerators

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

OpenMP Overview 78

## Using OpenMP



V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**Using OpenMP**

- The overview given so far can already be used to parallelize applications
- It is now time to look at some features in more detail
- And introduce some additional features

**Load Balancing**

**The Schedule Clause**

```
schedule ( static | dynamic | guided | auto [, chunk] )
schedule ( runtime )
```

**static [, chunk]**

- Distribute iterations in blocks of size "chunk" over the threads in a round-robin fashion
- In absence of "chunk", each thread executes approx.  $N/P$  chunks for a loop of length  $N$  and  $P$  threads
  - Details are implementation defined
- Under certain conditions, the assignment of iterations to threads is the same across multiple loops in the same parallel region

**Example Of A Static Schedule**

*A loop of length 16 using 4 threads*

Thread	0	1	2	3
no chunk *	1-4	5-8	9-12	13-16
chunk = 2	1-2	3-4	5-6	7-8
	9-10	11-12	13-14	15-16

*\*) The precise distribution is implementation defined*

**Loop Workload Scheduling Choices**

**dynamic [, chunk]**

- Fixed portions of work; size is controlled by the value of chunk
- When a thread finishes, it starts on the next portion of work

**guided [, chunk]**

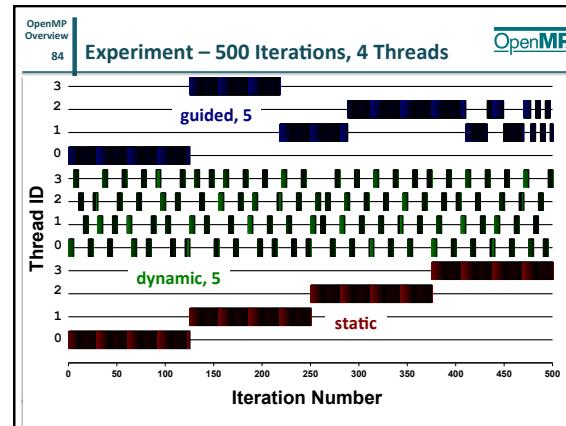
- Same dynamic behavior as "dynamic", but size of the portion of work decreases exponentially

**auto**

- The compiler (or runtime system) decides what is best to use; choice could be implementation dependent

**runtime**

- Iteration scheduling scheme is set at runtime through environment variable `OMP_SCHEDULE`



**The Master Directive**

*Only the master thread executes this code block*

```
#pragma omp master
{<code-block>}
```

```
!$omp master
<code-block>
!$omp end master
```

*There is no implied barrier on entry or exit!*

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**The Ordered Directive**

*The enclosed block of code within a parallel loop is executed in the order in which iterations would be executed sequentially:*

```
#pragma omp ordered
{<code-block>}
```

```
!$omp ordered
<code-block>
!$omp end ordered
```

*May introduce serialization (could be expensive)*

*Note: Need to use the ordered clause on the parallel for/do loop*

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**Atomic Operations**

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**The Atomic Construct/1**

*An atomic operation can be executed without any other thread being able to read or change state that is read or changed during the operation*

```
#pragma omp atomic update
{
    x++;
}
```

*Source: [http://wiki.osdev.org/Atomic\\_operation](http://wiki.osdev.org/Atomic_operation)*

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**The Atomic Construct/2**

*Lightweight*

```
#pragma omp atomic [clause]
```

```
!$omp atomic [clause]
```

*Clauses on the atomic: read, write, update, capture, seq\_cst*

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**Example – The Capture Clause**

*This captures the original (or final) value, stored in a designated variable*

```
#pragma omp atomic capture
{
    capture-statement
    write-statement
}
```

```
#pragma omp atomic capture
{
    old_value = *p;
    (*p)++;
}
// Can use old_value here
```

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**OpenMP Overview** 91

## Sequentially Consistent Atomics

"...the results of any execution is the same as if the operations of all the processors were executed in some sequential order, and the operations of each individual processor appear in this sequence in the order specified by its program." – Leslie Lamport

Specify through the `seq_cst` clause on the atomic construct

This implies a flush without a list

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**OpenMP Overview** 92

## Orphaning

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**OpenMP Overview** 93

## Orphaning

```
#pragma omp parallel
{
    :
    (void) dowork()
    :
} // End of parallel
```

```
void dowork()
{
    :
    #pragma omp for
    for (int i=0;i<n;i++)
    {
        :
    }
}
```

The OpenMP specification does not restrict worksharing and synchronization directives (`omp for`, `omp single`, `critical`, `barrier`, etc.) to be within the lexical extent of a parallel region. These directives can be orphaned.

That is, they can appear outside the lexical extent of a parallel region

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**OpenMP Overview** 94

## More On Orphaning

```
(void) dowork(); !- Sequential FOR
#pragma omp parallel
{
    (void) dowork(); !- Parallel FOR
}
```

```
void dowork()
{
    #pragma omp for
    for (i=0;....)
    {
        :
    }
}
```

When an orphaned worksharing or synchronization directive is encountered in the sequential part of the program (outside the dynamic extent of any parallel region), it is executed by the master thread only. In effect, the directive will be ignored

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**OpenMP Overview** 95

## Example - Parallelizing Bulky Loops

```
for (i=0; i<n; i++) /* Parallel loop */
{
    a = ...
    b = ... a ...
    c[i] = ....
    ....
    for (j=0; j<m; j++)
    {
        <a lot more work in this loop>
    }
    ....
}
```

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**OpenMP Overview** 96

## Step 1: "Outlining"

```
for (int i=0; i<n; i++) /* Parallel loop */
{
    (void) FuncPar(i,m,c,...)
}
```

- ✓ Still a sequential program
- ✓ Should behave identically
- ✓ Easy to test for correctness
- ✓ But, parallel by design

```
void FuncPar(i,m,c,...)
{
    float a, b; /* Private data */
    int j;
    a = ...
    b = ... a ...
    c[i] = ....
    ....
    for (j=0; j<m; j++)
    {
        <a lot more work in this loop>
    }
    ....
}
```

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**Step 2: Parallelize**

```
#pragma omp parallel for private(..) shared(..)
for (int i=0; i<n; i++) /* Parallel loop */
{
    (void) FuncPar(i,m,c,...)
} /*-- End of parallel for --*/

✓ Minimal scoping required
✓ Less error prone
```

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

OpenMP Overview 98

**Global Data**

```
float a, b; /* Private data */
int j;
a = ...
b = ...
c[i] = ...
.....
for (j=0; j<m; j++)
{
    a lot more work in this loop
}
.....
```

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**Global Data – An Example/1**

```
program global_data
use mod_global_data
...
!$omp parallel do private(j)
do j = 1, n
    call suba(j)
end do
!$omp end parallel do
.....
module mod_global_data
implicit none
integer, parameter:: m= ..., n= ...
integer :: a(m,n), b(m)
end module mod_global_data
```

Arrays "a" and "b" are shared

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

OpenMP Overview 100

**Global Data – An Example/2**

```
subroutine suba(j)
...
use mod_global_data
...
do i = 1, m
    b(i) = j
end do
do i = 1, m
    a(i,j) = func_call(b(i))
end do
return
end
```

Data Race!

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**Global Data - A Data Race!**

Thread 1	Thread 2
call suba(1)	call suba(2)
subroutine suba(j=1)	subroutine suba(j=2)
do i = 1, m b(i) = 1 end do	do i = 1, m b(i) = 2 end do
do i = 1, m a(i,1)=func_call(b(i)) end do	do i = 1, m a(i,2)=func_call(b(i)) end do

Shared

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

OpenMP Overview 102

**Global Data – A Solution/1**

```
program global_data
use mod_global_data
...
!$omp parallel do private(j)
do j = 1, n
    call suba(j)
end do
!$omp end parallel do
.....
module mod_global_data
implicit none
integer, parameter:: m= ..., n= ...
integer, parameter:: nthreads = ...
integer :: a(m,n), b(m,nthreads)
end module mod_global_data
```

Make array "b" 2-dimensional

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**OpenMP Overview** 103

## Global Data – A Solution/2

```

subroutine suba(j)
  .....
  use omp_lib
  use mod_global_data
  .....
  TID = omp_get_thread_num() + 1
  do i = 1, m
    b(i,TID) = j
  end do

  do i = 1, m
    a(i,j) = func_call(b(i), TID)
  end do

  return
end

```

A lot of work and not very portable

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**OpenMP Overview** 104

## About Global Data

- Global data is shared and requires special care
- A problem may arise in case multiple threads access the same memory section simultaneously:
  - Read-only data is no problem
  - Updates have to be checked for race conditions
- It is your responsibility to deal with this situation
- In general one can do the following:
  - Split the global data into a part that is accessed in serial parts only and a part that is accessed in parallel
  - Manually create thread private copies of the latter
  - Use the thread ID to access these private copies
- Alternative: Use OpenMP's threadprivate directive !

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**OpenMP Overview** 105

## The Threadprivate Directive

```

#pragma omp threadprivate (list)
!$omp threadprivate (/cb/ [./cb/] ...)

• Thread private copies of the designated global variables and common blocks are created
• Several restrictions and rules apply when doing this:
  - The number of threads has to remain the same for all the parallel regions (i.e. no dynamic threads)
  - Some implementations support changing the number of threads
  - Initial data is undefined, unless copyin is used
  - .....
• Check the specifications when using threadprivate !

```

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**OpenMP Overview** 106

## Global Data – The Preferred Solution

```

program global_data
  ...
  use mod_global_data
  ...
  !$omp parallel do private(j)
  do j = 1, n
    call suba(j)
  end do
  !$omp end parallel do

  module mod_global_data
    implicit none
    integer, parameter:: m=..., n=...
    integer :: a(m,n), b(m)
    !$omp threadprivate(b)
  end module mod_global_data

```

Only add the "threadprivate" directive to the module file; no other changes needed !

This solution also automatically adapts to the number of threads used

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**OpenMP Overview** 107

## The Copyin Clause

copyin (list)

- ✓ Applies to THREADPRIVATE common blocks only
- ✓ At the start of the parallel region, data of the master thread is copied to the thread private copies

Example:

```

common /cblock/velocity
common /fields/xfield, yfield, zfield
! create thread private common blocks
!$omp threadprivate (/cblock/, /fields/)

!$omp parallel
!$omp default (private) &
!$omp copyin (/cblock/, zfield )

```

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**OpenMP Overview** 108

## C++ And Threadprivate

- As of OpenMP 3.0, it has been clarified where/how threadprivate objects are constructed and destructed
- Allow C++ static class members to be threadprivate

```

class T {
public:
  static int i;
  #pragma omp threadprivate(i)
  ...
}

```

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

OpenMP Overview 109

## A Locking Example

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

OpenMP Overview 110

## OpenMP Locking Routines

- Locks provide greater flexibility over critical sections and atomic updates:
  - Possible to implement asynchronous behavior
  - Not block structured
- The so-called lock variable, is a special variable:
  - C/C++: type `omp_lock_t` and `omp_nest_lock_t` for nested locks
  - Fortran: type `INTEGER` and of a `KIND` large enough to hold an address
- Lock variables should be manipulated through the API only
- It is illegal, and behavior is undefined, in case a lock variable is used without the appropriate initialization

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

OpenMP Overview 111

## Nested Locking

- Simple locks: may not be locked if already in a locked state
- Nestable locks: may be locked multiple times by the same thread before being unlocked
- In the remainder, we discuss simple locks only
- The interface for functions dealing with nested locks is

Simple locks	Nestable locks
<code>omp_init_lock</code>	<code>omp_init_nest_lock</code>
<code>omp_destroy_lock</code>	<code>omp_destroy_nest_lock</code>
<code>omp_set_lock</code>	<code>omp_set_nest_lock</code>
<code>omp_unset_lock</code>	<code>omp_unset_nest_lock</code>
<code>omp_test_lock</code>	<code>omp_test_nest_lock</code>

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

OpenMP Overview 112

## OpenMP Locking Example

- The protected region contains the update of a shared variable
- One thread acquires the lock and performs the update
- Meanwhile, the other thread performs some other work
- When the lock is released again, the other thread performs the update

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

OpenMP Overview 113

## Locking Example - The Code

```

Program Locks
...
Call omp_init_lock (LCK)
!$omp parallel shared(LCK)
  Do While ( omp_test_lock (LCK) .EQV. .FALSE. )
    Call Do_Something_Else()
  End Do
  Call Do_Work()
  Call omp_unset_lock (LCK)
!$omp end parallel
  Call omp_destroy_lock (LCK)
Stop
End

```

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

OpenMP Overview 114

## Example Output Using 2 Threads

```

TID: 1 at 09:07:27 => entered parallel region
TID: 1 at 09:07:27 => ready to do the parallel work
TID: 1 at 09:07:27 => this will take about 18 seconds
TID: 0 at 09:07:27 => entered parallel region
TID: 0 at 09:07:27 => WAIT for lock - will do something else for 5 seconds
TID: 0 at 09:07:32 => ready to do the parallel work
TID: 0 at 09:07:37 => WAIT for lock - will do something else for 5 seconds
TID: 0 at 09:07:42 => WAIT for lock - will do something else for 5 seconds
TID: 1 at 09:07:45 => done with my work
TID: 1 at 09:07:45 => ready to leave the parallel region
TID: 1 at 09:07:45 => done with lock - released the lock
TID: 0 at 09:07:45 => ready to leave the parallel region
TID: 0 at 09:07:47 => ready to do the parallel work
TID: 0 at 09:07:47 => this will take about 18 seconds
TID: 0 at 09:08:05 => done with work loop - released the lock
TID: 0 at 09:08:05 => ready to leave the parallel region
Done at 09:08:05 - value of SUM is 1100

```

Used to check the answer

Note: program has been instrumented to get this information

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

OpenMP Overview 115

## Additional Functionality

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE



OpenMP Overview 117

## Additional Functionality

- OpenMP continues to evolve*
  - Increasingly flexible to support more algorithms*
  - Support for modern parallel architectures*
  - cc-NUMA and Accelerators*
- In recent years, several key features have been added to significantly enhance the functionality*
- The following features are discussed next:*
  - Thread Cancellation*
  - Accelerator Support*
  - Tasking*
  - Thread Affinity*

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

OpenMP Overview 118

## Thread Cancellation

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

OpenMP Overview 119

## Thread Cancellation – Basic Idea

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

OpenMP Overview 120

## About Thread Cancellation

- Cancellation – Abort an OpenMP region*
  - Implicit or explicit tasks proceed to the end of the canceled region*
- Cancellation point – Point where tasks check if cancellation has been requested*
  - Take action accordingly (abort or continue)*
- Cancellation points are:*
  - Implicit barriers, barrier regions, cancel regions, cancellation point regions*

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**Cancellation Directives**

*Activate cancellation of the innermost enclosing region  
(if cancellation has been enabled)*

```
#pragma omp cancel <construct> [if-clause]
```

*Define an explicit cancellation point*

```
#pragma omp cancellation point <construct>
```

*Supported constructs:*

```
parallel
sections
for
taskgroup
```

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

OpenMP Overview 122

**Accelerator Support**

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**Accelerator Support/1**

**Host System**

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

OpenMP Overview 124

**Accelerator Support/2**

**“Host Device”**

The “target” construct is used to execute code on a specific device

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**Accelerator Support/3**

```
#pragma omp target device(2)
{
    "execute this code on target device #2"
}
```

- User needs to explicitly manage input and output data using the map clause
- This includes “to”, “from” and “tofrom” syntax

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

OpenMP Overview 126

**Tasking**

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**Tasking In OpenMP**



- Tasking was introduced in OpenMP 3.0
- Until then, it required a heroic effort to efficiently implement certain types of parallelism
  - Divide and Conquer algorithms
  - Recursive algorithms
  - Linked lists
  - ...
- The initial tasking functionality in 3.0 was very simple
  - The idea was to augment tasking as we collectively gain more insight and experience
  - Successive specifications have added functionality

V1 RvdP

Getting Up To Speed On OpenMP 4.0 ORACLE

**What Is An OpenMP Task ?**



- Loosely said, a task is a block of code that can be executed simultaneously with other tasks
- It is upon the programmer to identify tasks in the code
- From there on, the compiler and run time system handle the creation and execution of the tasks

```
#pragma omp task [clauses]
!$omp task [clauses]
```

Defines a task

V1 RvdP

Getting Up To Speed On OpenMP 4.0 ORACLE

**Tasking Example**



```
#pragma omp parallel
{
  #pragma omp task
  {printf("I am a task\n");}
} // End of parallel region
```

All threads execute this task

V1 RvdP

Getting Up To Speed On OpenMP 4.0 ORACLE

**Tasking - Terminology**



A **task** is a specific instance of executable code and its data environment, generated when a thread encounters a task construct or a parallel construct.

A **task region** is a region consisting of all code encountered during the execution of a task.

The **data environment** consists of all the variables associated with the execution of a given task. The data environment for a given task is constructed from the data environment of the generating task at the time the task is generated.

Note: The data environment is not covered here. Come to our tutorial at SC'15 to learn more about this!

V1 RvdP

Getting Up To Speed On OpenMP 4.0 ORACLE

**Tasking - Who Does What ?**



- When a thread encounters a task construct, a new explicit task is generated
  - But not necessarily executed yet
- Execution of tasks is handled by the run time system
  - Tasks are assigned to the threads in the current team
  - This is subject to the thread's availability and thus could be immediate, or deferred until later
- Threads are allowed to suspend the current task region at a task scheduling point in order to execute a different task

Note: Tasks can be nested (not for the faint of heart)

V1 RvdP

Getting Up To Speed On OpenMP 4.0 ORACLE

**Tasking – And When ?**



- Task completion occurs when the end of the structured block associated with the construct that generated the task is reached
- Completion of a subset of all explicit tasks bound to a given parallel region may be specified through the use of task synchronization constructs
  - Completion of all explicit tasks bound to the implicit parallel region is guaranteed by the time the program exits
- A task synchronization construct is a taskwait, taskgroup or a barrier construct

V1 RvdP

Getting Up To Speed On OpenMP 4.0 ORACLE

**OpenMP Overview** 133 **OpenMP**

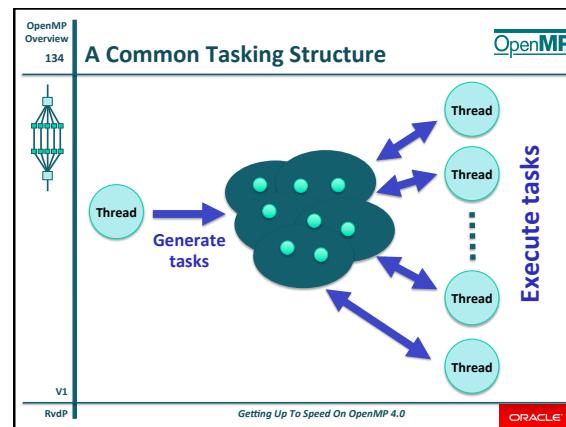
### How To Enforce Task Completion ?

```
#pragma omp taskwait
!$omp flush taskwait
```

*Wait on completion of all child tasks \**

*\* This includes direct children only, not descendant tasks. Use the taskgroup construct for the latter*

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE



**OpenMP Overview** 135 **OpenMP**

### A Simple Plan

**Your Task for Today:**

**Write a program that prints either "A race car" or "A car race" and maximize the parallelism**

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**OpenMP Overview** 136 **OpenMP**

### Tasking Example/1

```
$ cc -fast hello.c
$ ./a.out
A race car
```

```
#include <stdlib.h>
#include <stdio.h>

int main(int argc, char *argv[]) {
    printf("A ");
    printf("race ");
    printf("car ");

    printf("\n");
    return(0);
}
```

**What will this program print ?**

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**OpenMP Overview** 137 **OpenMP**

### Tasking Example/2

```
#include <stdlib.h>
#include <stdio.h>

int main(int argc, char *argv[]) {
    #pragma omp parallel
    {
        printf("A ");
        printf("race ");
        printf("car ");

    } // End of parallel region
    printf("\n");
    return(0);
}
```

**What will this program print using 2 threads ?**

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**OpenMP Overview** 138 **OpenMP**

### Tasking Example/3

```
$ cc -xopenmp -fast hello.c
$ export OMP_NUM_THREADS=2
$ ./a.out
A race car A race car
```

**Note that this program could (for example) also print "A A race race car car" or "A race A car race car", or "A race A race car car", or .....**

**But I have not observed this (yet)**

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

OpenMP Overview 139

### Tasking Example/4

**What will this program print using 2 threads ?**

```
#include <stdlib.h>
#include <stdio.h>

int main(int argc)
{
    #pragma omp parallel
    {
        #pragma omp single
        {
            printf("A ");
            printf("race ");
            printf("car ");
        } // End of parallel region
        printf("\n");
        return(0);
    }
}
```

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

OpenMP Overview 140

### Tasking Example/5

```
$ cc -xopenmp -fast hello.c
$ export OMP_NUM_THREADS=2
$ ./a.out
A race car
```

**But of course now only 1 thread executes .....**

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

OpenMP Overview 141

### Tasking Example/6

**What will this program print using 2 threads ?**

```
int main(int argc, char *argv[])
{
    #pragma omp parallel
    {
        #pragma omp single
        {
            printf("A ");
            #pragma omp task
            {printf("race ");}
            #pragma omp task
            {printf("car ");}
        } // End of parallel region
        printf("\n");
        return(0);
    }
}
```

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

OpenMP Overview 142

### Tasking Example/7

```
$ cc -xopenmp -fast hello.c
$ export OMP_NUM_THREADS=2
$ ./a.out
A race car
$ ./a.out
A race car
$ ./a.out
A car race
$
```

**Tasks can be executed in arbitrary order**

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

OpenMP Overview 143

### Another Simple Plan

**You did well and quickly, so here is a final task to do**

**Have the sentence end with "is fun to watch"**  
(hint: use a print statement)

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

OpenMP Overview 144

### Tasking Example/8

**What will this program print using 2 threads ?**

```
int main(int argc, char *argv[])
{
    #pragma omp parallel
    {
        #pragma omp single
        {
            printf("A ");
            #pragma omp task
            {printf("race ");}
            #pragma omp task
            {printf("car ");}
            printf("is fun to watch ");
        } // End of parallel region
        printf("\n");
        return(0);
    }
}
```

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**Tasking Example/9**

```
$ cc -xopenmp -fast hello.c
$ export OMP_NUM_THREADS=2
$ ./a.out
A is fun to watch race car
$ ./a.out
A is fun to watch race car
$ ./a.out
A is fun to watch car race
$
```

**Tasks are executed at a task execution point**

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**Tasking Example/10**

```
int main(int argc, char **argv)
{
    #pragma omp parallel
    {
        #pragma omp single
        {
            printf("A ");
            #pragma omp task
            {printf("car ");}
            #pragma omp task
            {printf("race ");}

            #pragma omp taskwait
            printf("is fun to watch ");
        }
    } // End of parallel region

    printf("\n");
}
```

**What will this program print using 2 threads ?**

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**Tasking Example/11**

```
$ cc -xopenmp -fast hello.c
$ export OMP_NUM_THREADS=2
$ ./a.out
A car race is fun to watch
$ ./a.out
A car race is fun to watch
$ ./a.out
A race car is fun to watch
$
```

**Tasks are executed first now**

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**More About Tasking**

- Several features have not been covered yet
- These are beyond the scope of this tutorial
  - but it doesn't mean they're not useful
- Please check the specifications for much more information on tasking
  - Or attend our tutorial at SC'15!

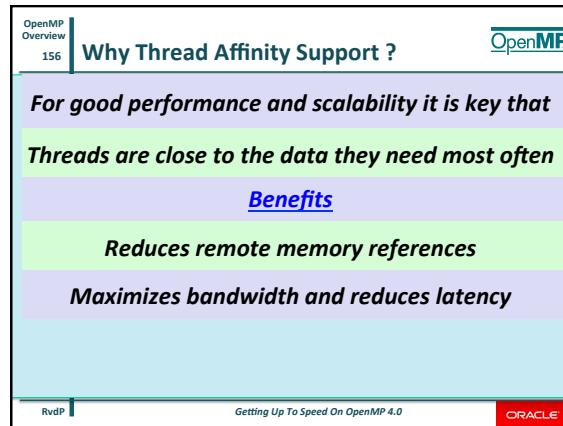
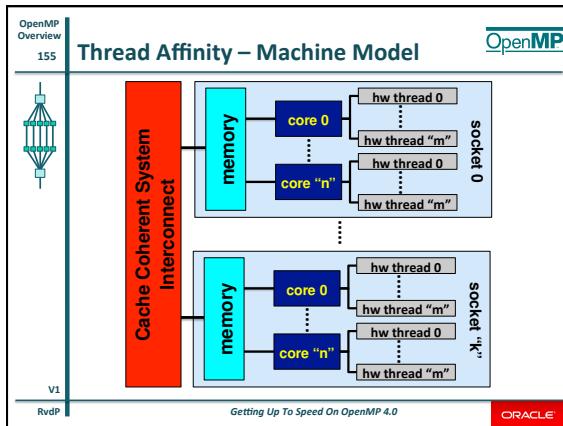
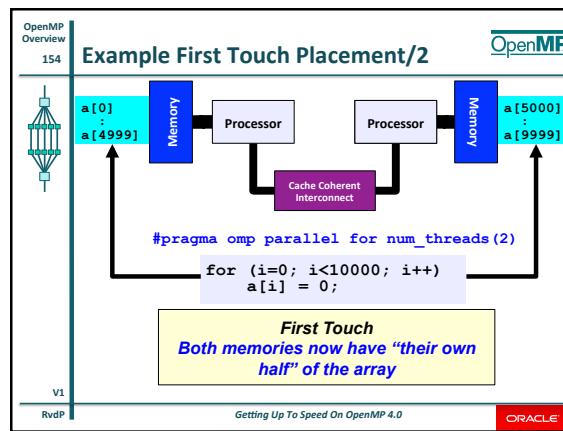
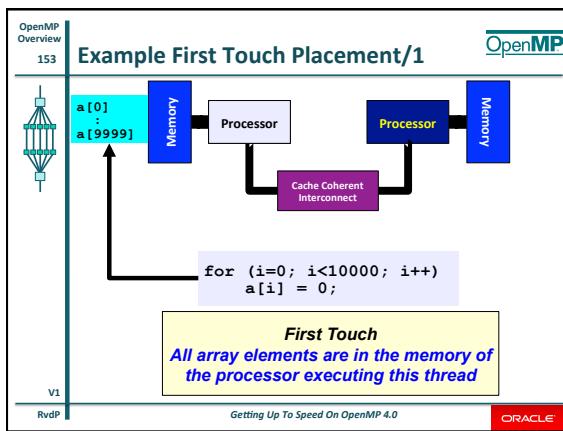
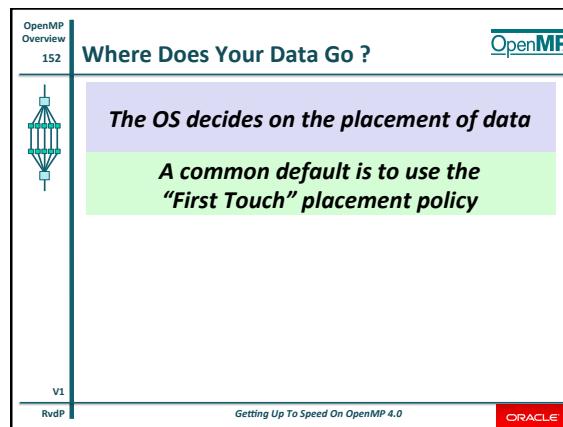
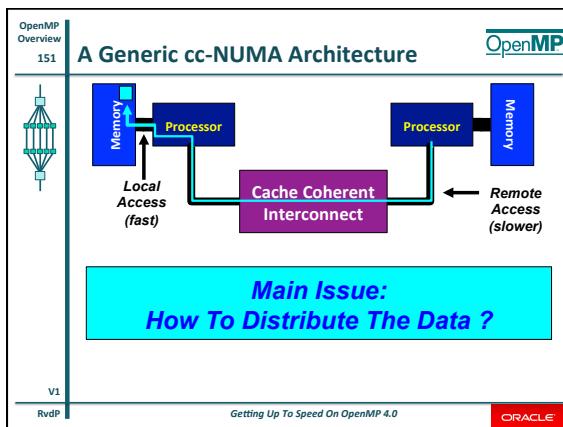
V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**Thread Affinity**

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

**Why Worry ?**

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE



**OpenMP Overview** 157 **OpenMP**

## Basic Affinity Philosophy



*Data is wherever it may be*

*Threads are moved to the data they need most*

V1 RvdP

Getting Up To Speed On OpenMP 4.0 ORACLE

**OpenMP Overview** 158 **OpenMP**

## Two Key Concepts



*The Place List*

*The Thread Affinity Policy*

V1 RvdP

Getting Up To Speed On OpenMP 4.0 ORACLE

**OpenMP Overview** 159 **OpenMP**

## The Place List



V1 RvdP

Getting Up To Speed On OpenMP 4.0 ORACLE

**OpenMP Overview** 160 **OpenMP**

## Places – Definition And Notation



*A place consists of a (set of) numbers*

*Each number represents a scheduling unit*

*That is, something a thread can run on*

*For example, a hardware thread*

V1 RvdP

Getting Up To Speed On OpenMP 4.0 ORACLE

**OpenMP Overview** 161 **OpenMP**

## Example Multicore System



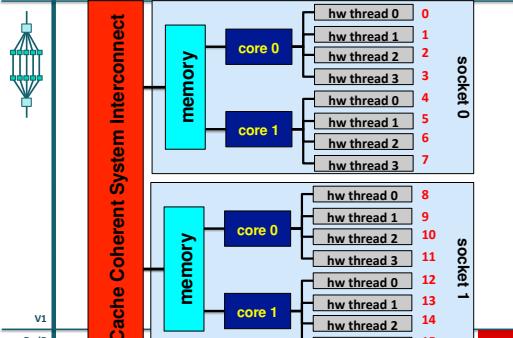
Component	Count	Total
Sockets	2	2
Cores/socket	2	4
Threads/core	4	16

V1 RvdP

Getting Up To Speed On OpenMP 4.0 ORACLE

**OpenMP Overview** 162 **OpenMP**

## Example System Architecture



*Cache Coherent System Interconnect*

*memory*

*core 0*

*core 1*

*hw thread 0 0*  
*hw thread 1 1*  
*hw thread 2 2*  
*hw thread 3 3*  
*hw thread 0 4*  
*hw thread 1 5*  
*hw thread 2 6*  
*hw thread 3 7*

*socket 0*

*memory*

*core 0*

*core 1*

*hw thread 0 8*  
*hw thread 1 9*  
*hw thread 2 10*  
*hw thread 3 11*  
*hw thread 0 12*  
*hw thread 1 13*  
*hw thread 2 14*  
*hw thread 3 15*

*socket 1*

V1 RvdP

Getting Up To Speed On OpenMP 4.0 ORACLE

OpenMP Overview 163

### Places - Example

$\{0,1,2,3\}$  identifies the threads in socket 0, core 0

The diagram illustrates memory access. A red vertical bar labeled "memory" has arrows pointing to two cores: "core 0" and "core 1". "core 0" is connected to four hardware threads (hw thread 0, 1, 2, 3) which are grouped under "socket 0". "core 1" is connected to four hardware threads (hw thread 4, 5, 6, 7) which are grouped under "socket 1".

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

OpenMP Overview 164

### Places - Example

$\{0,1,2,3\}$  identifies threads in socket 0, core 0

Easy interval notation:  $\{0 : 4 : 1\}$

start : count : stride

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

OpenMP Overview 165

### Places – The Order Does Not Matter

The order of the numbers within a place does not matter

$\{0,1,2,3\}$  is the same as  $\{3,2,1,0\}$

*Assumption*  
No preference regarding memory access time

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

OpenMP Overview 166

### The Place List - Definition

The Place List consists of a comma separated list of places

For example:  $\{0,1,2,3\}, \{8,9,10,11\}$

The Order Of The Places In The List DOES Matter

$\{0,1,2,3\}, \{8,9,10,11\} \neq \{8,9,10,11\}, \{0,1,2,3\}$

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

OpenMP Overview 167

### The Place List – How To Set It

Environment variable OMP\_PLACES is used to define the place list

Example: `OMP_PLACES=“{0,1,2,3}, {8,9,10,11}”`

The interval notation is very convenient

Example: `OMP_PLACES=“{0:4:1}, {8:4:1}”`

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

OpenMP Overview 168

### The Place List – Abstract Names

Three abstract names are always available:

**sockets, cores, threads**

Example: `OMP_PLACES=cores`

Example: `OMP_PLACES=“cores(4)”`

Note: Implementation can add names

V1 RvdP Getting Up To Speed On OpenMP 4.0 ORACLE

OpenMP Overview 169

### Example – OMP\_PLACES

**OMP\_PLACES=cores**  
Equivalent To  
**OMP\_PLACES=" {0,1,2,3},{4,5,6,7},{8,9,10,11},{12,13,14,15}"**  
Equivalent To  
**OMP\_PLACES=" {0:4:1},{4:4:1},{8:4:1},{12:4:1}"**  
Equivalent To  
**OMP\_PLACES = " {0:4:1}:4:4"**

RvdP | Getting Up To Speed On OpenMP 4.0 | ORACLE

OpenMP Overview 170



### The Affinity Policy

V1

RvdP | Getting Up To Speed On OpenMP 4.0 | ORACLE

OpenMP Overview 171

### Thread Affinity Policies

**The Affinity Policy defines which places to use**  
**This is defined in a symbolic way:**  
**master, close, or spread**  
**Each parallel region has such an affinity policy**  
**Thread binding to a place is implied**

V1 | Getting Up To Speed On OpenMP 4.0 | ORACLE

OpenMP Overview 172

### Setting The Thread Affinity Policy

**Defined through OMP\_PROC\_BIND**  
**Example: OMP\_PROC\_BIND="spread,close"**  
**Can also use the "proc\_bind" clause**  
**Applies to the current parallel region only**

V1 | Getting Up To Speed On OpenMP 4.0 | ORACLE

OpenMP Overview 173

### Recap Places And Affinity

**The Place List defines what is available  
(fixed for the duration of the program)**  
**The Affinity Policy defines thread placement  
(can be set for each parallel region)**

V1 | Getting Up To Speed On OpenMP 4.0 | ORACLE

OpenMP Overview 174

### Example – The Spread Policy

**OMP\_PLACES=cores  
(this means 4 places in the place list)**  
**OMP\_PROC\_BIND=spread**  
**OMP\_NUM\_THREADS=4**  
**Result: One OpenMP thread per place**

V1 | Getting Up To Speed On OpenMP 4.0 | ORACLE

